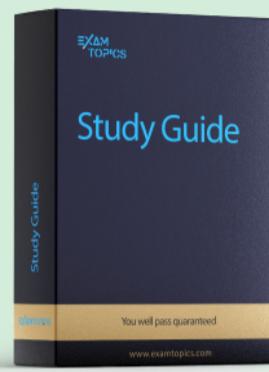




- Expert Verified, Online, Free.

Prepare for your Associate Cloud Engineer exam with additional products

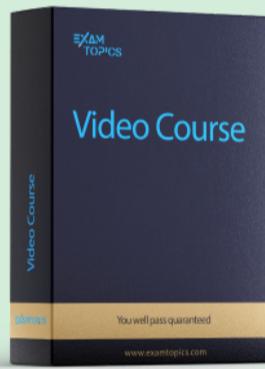


Study Guide

849 PDF Pages

\$19.99

[Buy Now](#)



Video Course

234 Lectures

\$19.99

[Buy Now](#)

[Custom View Settings](#)

Every employee of your company has a Google account. Your operational team needs to manage a large number of instances on Compute Engine. Each member of this team needs only administrative access to the servers. Your security team wants to ensure that the deployment of credentials is operationally efficient and must be able to determine who accessed a given instance. What should you do?

- A. Generate a new SSH key pair. Give the private key to each member of your team. Configure the public key in the metadata of each instance.
- B. Ask each member of the team to generate a new SSH key pair and to send you their public key. Use a configuration management tool to deploy those keys on each instance.
- C. Ask each member of the team to generate a new SSH key pair and to add the public key to their Google account. Grant the `compute.osAdminLogin` role to the Google group corresponding to this team.
- D. Generate a new SSH key pair. Give the private key to each member of your team. Configure the public key as a project-wide public SSH key in your Cloud Platform project and allow project-wide public SSH keys on each instance.

Correct Answer: D

Reference:

<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>

Community vote distribution

C (98%)

✉️  **dan80**  3 years, 4 months ago

C is correct - <https://cloud.google.com/compute/docs/instances/managing-instance-access>
upvoted 67 times

✉️  **adedj99** 2 years, 11 months ago

We recommend collecting users with the same responsibilities into groups and assigning IAM roles to the groups rather than to individual users. For example, you can create a "data scientist" group and assign appropriate roles to enable interaction with BigQuery and Cloud Storage. When a new data scientist joins your team, you can simply add them to the group and they will inherit the defined permissions. You can create and manage groups through the Admin Console.

upvoted 12 times

✉️  **zakhili**  3 years, 4 months ago

Send private key to users is not safe, i think it's C

upvoted 20 times

✉️  **amitsingh4444** 1 month ago

Please share full dump with me at amitsingh4444@gmail.com. Thanks a lot in advance :)
upvoted 2 times

✉️  **anandvpm** 1 year, 1 month ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my anandrajagopal@gmail.com address. Thanks in advance!
upvoted 3 times

✉️  **tulsi122** 3 months, 2 weeks ago

no use of asking for questions because we need to pay huge amount to download pdf and small amount to see full questions but not able to download.any more doubt welcome on tulsigcp@gmail
upvoted 2 times

✉️  **Andy_23** 3 months, 3 weeks ago

If you got the questions can you share the questions with me also.
anujnarayan11@gmail.com
Thanks in advance.
upvoted 1 times

✉️  **karim1321** 3 months, 4 weeks ago

Hi, Could you please share me the question on rifkikarimr@gmail.com ?, Thanks in advance!
upvoted 1 times

✉️  **ROSHANgcp** 3 months, 2 weeks ago

how can anyone send questions, because we can download that in ete format and can only view from login on this portal only
upvoted 1 times

✉️  **Dude1204** 3 months, 3 weeks ago

Hi, same here. Please share the questions to af165e@gmail.com
Thanks in advance

upvoted 1 times

✉ **pat92** Most Recent 1 day ago

Hi Guys.

I request you please send me exam practice test questions and answers(Google Associate Cloud Engineer). I have exam on 15 october.

Email id: ghosh.pratik92@gmail.com

Thank you in advance.

upvoted 1 times

✉ **irfan_4141** 4 days, 18 hours ago

I am preparing for GCP Associate engineer exam. If anyone has the entire questions set, please share them at muhammad.irfan.tahir93@gmail.com

upvoted 1 times

✉ **pavanyennu** 1 week, 2 days ago

I am preparing for GCP Associate engineer exam and it's very near. I was able to access 92 questions only, if anyone has the entire questions please share them to my pavanyennu1@gmail.com address. Thanks in advance!

upvoted 1 times

✉ **Sushilshr** 1 week, 2 days ago

d is correct

upvoted 1 times

✉ **osm0z** 1 week, 3 days ago

Selected Answer: D

d is correct

<https://cloud.google.com/compute/docs/instances/access-overview>

upvoted 1 times

✉ **anmolM** 1 week, 3 days ago

C is correct...

upvoted 2 times

✉ **Sara_jan86** 2 weeks ago

Hi Guys.

I request you please send me exam practice test questions and answers(Google Associate Cloud Engineer). I have exam on 29 September.

Email id: famb86@gmail.com

Thank you in advance.

upvoted 1 times

✉ **Shivangi30** 2 weeks, 2 days ago

I gave exam yesterday ,cleared the exam and almost all ques were from here. Thanks

upvoted 1 times

✉ **alfo23** 2 weeks, 4 days ago

hi, will be appriciate for dpf itisstuffforme@gmai.com .TYA

upvoted 1 times

✉ **karthikesavan** 3 weeks, 4 days ago

I am preparing for GCP Associate engineer exam. I was able to access 92 ques, only. IF anyone have pdf kindly send me to my email srvnkmrj@gmail.com. Thanks in advance!

upvoted 1 times

✉ **pavanyennu** 1 week, 2 days ago

Can you share the dumps if you got any to pavanyennu1@gmail.com

Thanks

upvoted 1 times

✉ **JPBaria** 3 weeks, 5 days ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my contato.jpbaria@gmail.com address. Thanks in advance!

upvoted 1 times

✉ **jovialv** 4 weeks, 1 day ago

Hi everyone - Kindly send me the entire pdf question list on jovial.vinay@gmail.com

upvoted 1 times

✉ **Musfik41** 1 month ago

Hi, same here. Please share the questions to mushfiqueazam72@gmail.com

Thanks in advance

upvoted 1 times

✉ **kkfiner** 1 month ago

I am preparing for the GCP Associate Engineer exam. If anyone has the entire questions set, please send them to kkfiner@gmail.com

upvoted 1 times

✉ **Captain1212** 1 month, 1 week ago

Selected Answer: C

C is the correct answer, it gives you the required access

upvoted 1 times

You need to create a custom VPC with a single subnet. The subnet's range must be as large as possible. Which range should you use?

- A. 0.0.0.0/0
- B. 10.0.0.0/8
- C. 172.16.0.0/12
- D. 192.168.0.0/16

Correct Answer: A

Community vote distribution



Carls 3 years, 6 months ago

I got this question in one of the trainings I did in udemy: <https://www.udemy.com/course/google-cloud-associate-engineer-exam-practice-tests/?referralCode=810D02D4A159FC3E36CC>

B is correct. Pay attention to the question, it is talking about custom VPC subnet and is not mentioning you will use automatic subnet mode creation. If you set subnet to custom, the minimum size is /8.

upvoted 69 times

passnow 3 years, 6 months ago

B is legit as fuck!

upvoted 34 times

hems4all 2 years, 11 months ago

B is correct

Use 10.0.0.0/8 CIDR range. is the right answer.

The private network range is defined by IETF (Ref: <https://tools.ietf.org/html/rfc1918>) and adhered to by all cloud providers. The supported internal IP Address ranges are

1. 24-bit block 10.0.0.0/8 (16777216 IP Addresses)
2. 20-bit block 172.16.0.0/12 (1048576 IP Addresses)
3. 16-bit block 192.168.0.0/16 (65536 IP Addresses)

10.0.0.0/8 gives you the most extensive range - 16777216 IP Addresses.

upvoted 17 times

FunkyB 1 year, 3 months ago

Language. :-(

upvoted 6 times

irfan_4141 4 days, 18 hours ago

I am preparing for GCP Associate engineer exam. If anyone has the entire questions set, please share them at muhammad.irfan.tahir93@gmail.com

upvoted 1 times

tkaur3547 2 weeks, 1 day ago

Hi guys,

I am preparing for ACE. Can someone please send me the pdf file for the question bank?

upvoted 1 times

YourCloudGuru 2 weeks, 3 days ago

B is the correct answer: 10.0.0.0/8.

This is the largest subnet range that you can use in a custom VPC. It has 16,777,216 addresses, which is more than enough for most organizations.

The other options are smaller subnet ranges:

0.0.0.0/0 is the entire IPv4 address space. It is not recommended to use this range for a subnet, because it would give you too many IP addresses to manage.

172.16.0.0/12 has 1,048,576 addresses.

upvoted 1 times

✉ **papel** 1 month, 1 week ago

Hi Guys. I request you please send me exam practice test questions and answers(Google Associate Cloud Engineer) with the discussion pdf file. Email: papellarbie@gmail.com

upvoted 1 times

✉ **Captain1212** 1 month, 1 week ago

b is the correct answer

upvoted 1 times

✉ **goshubh** 1 month, 1 week ago

I am preparing for GCP Associate engineer exam. If anyone has the entire questions set, please share them to go.shubh3@gmail.com
upvoted 1 times

✉ **rrope** 1 month, 3 weeks ago

B is correct ,

upvoted 1 times

✉ **rejar** 2 months ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my rejaryoga099@gmail.com address. Thanks in advance!

upvoted 1 times

✉ **timeto** 1 month, 1 week ago

Hii Rejar or anyone having complete set of questions PDF, can you please send it to thestrongestmetalintheuniverse@gmail.com please
upvoted 1 times

✉ **ETP** 2 months, 1 week ago

Hi Guys.

I request you please send me exam practice test questions and answers(Google Associate Cloud Engineer) with the discussion pdf file.
Email id- eknath.patil2009@gmail.com

Thank you in advance

upvoted 1 times

✉ **Mogamat** 2 months, 2 weeks ago

Answer B

upvoted 1 times

✉ **georgesouzafarias** 3 months, 2 weeks ago

Selected Answer: B

B is correct.

upvoted 1 times

✉ **zwill157** 3 months, 3 weeks ago

Selected Answer: B

B Correta

upvoted 1 times

✉ **cpnath** 4 months ago

B is correct

upvoted 1 times

✉ **JohnPhan** 5 months ago

Selected Answer: B

B is correct

10.0.0.0/8

upvoted 1 times

✉ **Mamoon2023** 5 months ago

I am preparing for GCP Associate engineer exam. If anyone has the entire questions set, please share them to mamoon92@gmail.com
upvoted 1 times

You want to select and configure a cost-effective solution for relational data on Google Cloud Platform. You are working with a small set of operational data in one geographic location. You need to support point-in-time recovery. What should you do?

- A. Select Cloud SQL (MySQL). Verify that the enable binary logging option is selected.
- B. Select Cloud SQL (MySQL). Select the create failover replicas option.
- C. Select Cloud Spanner. Set up your instance with 2 nodes.
- D. Select Cloud Spanner. Set up your instance as multi-regional.

Correct Answer: A

Reference:

<https://cloud.google.com/sql/docs/mysql/backup-recovery/restore>

Community vote distribution

A (100%)

✉ **YashBindlish** Highly Voted 3 years, 5 months ago

A is Correct. You must enable binary logging to use point-in-time recovery. Enabling binary logging causes a slight reduction in write performance. <https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

upvoted 45 times

✉ **ryumada** 1 year, 2 months ago

In this link below, the docs explains clearly that point-in-time recovery requires binary logging.

<https://cloud.google.com/sql/docs/mysql/backup-recovery/pitr#disk-usage>

upvoted 4 times

✉ **Bharathy** Highly Voted 3 years, 6 months ago

A is correct, as Binary Logging enables Point in Recovery in Cloud SQL

upvoted 10 times

✉ **YourCloudGuru** Most Recent 2 weeks, 3 days ago

Correct answer is A. Cloud SQL is a fully-managed relational database service that supports MySQL, PostgreSQL, and SQL Server. It offers high availability, automatic backups, and point-in-time recovery. By enabling binary logging, you can restore your database to a specific point in time. Cloud SQL is a cost-effective solution for small sets of operational data in one geographic location. It is also a good choice for developers who are familiar with MySQL.

upvoted 1 times

✉ **dilchan** 3 weeks, 6 days ago

I am preparing for the GCP Associate Engineer exam. If anyone has the entire questions set, please send them to dilip.cmaurya@gmail.com

upvoted 1 times

✉ **kkfiner** 1 month ago

I am preparing for the GCP Associate Engineer exam. If anyone has the entire questions set, please send them to kkfiner@gmail.com

upvoted 1 times

✉ **Captain1212** 1 month, 1 week ago

for point in recovery - binary logging a is the correct answer

upvoted 1 times

✉ **Mogamat** 2 months, 2 weeks ago

Answe A

upvoted 1 times

✉ **TSEMSOUS** 4 months ago

A is the correct answer Answer A: small set of operational data in one geographic location - CloudSQL (Cloud Spanner generally to hold large data and global)

upvoted 2 times

✉ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: A

Option A:

<https://cloud.google.com/sql/docs/mysql/backup-recovery/pitr>

Perform the point-in-time recovery using binary log positions

While we recommend you perform point-in-time recovery using timestamps as described in the previous procedure, you can also perform point-in-time recovery by providing a specific binary log position in a binary log file.

upvoted 1 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: A

Binary logging and Cloud SQL is correct fit
upvoted 1 times

 **nikila99** 7 months, 4 weeks ago

Selected Answer: A

For a small set of operational data in one geographic location with the need to support point-in-time recovery, the most cost-effective solution on Google Cloud Platform would be option A, "Select Cloud SQL (MySQL). Verify that the enable binary logging option is selected."

Cloud SQL is a fully-managed relational database service that provides high availability and automatic backups, making it an ideal choice for small data sets that require point-in-time recovery. By enabling binary logging in Cloud SQL, you can perform point-in-time recovery and roll back changes to your database to a specific time in the past

upvoted 3 times

 **Buruguduystunstugudunstuy** 9 months, 2 weeks ago

Selected Answer: A

Option A is correct because Cloud SQL (MySQL) is a fully-managed, cloud-native database service that supports point-in-time recovery through the use of binary logs. Enabling binary logging allows you to recover your database at any point in time within the preceding 7 days.

This is a cost-effective solution because Cloud SQL (MySQL) is fully managed, which means you don't have to worry about patching, backups, or scaling.

<https://cloud.google.com/sql/docs/mysql>

upvoted 1 times

 **leogor** 11 months, 3 weeks ago

A, small amount of operation data doesnt need cloud spanner
upvoted 1 times

 **eledu1985** 11 months, 3 weeks ago

the most accurate link is:
<https://cloud.google.com/sql/docs/mysql/backup-recovery/pitr#disk-usage>
upvoted 1 times

 **Sgegmund** 1 year ago

Selected Answer: A

PiT Recovery : A
upvoted 1 times

 **GayuSundar** 1 year, 1 month ago

Answer A: small set of operational data in one geographic location - CloudSQL (Cloud Spanner generally to hold large data and global)
Point-in-time recovery means binary logging should be enabled.
upvoted 1 times

 **HILMY** 1 year, 2 months ago

-cost effective (eliminating point C and D)
-point-in-time recovery (Binary Logging)

So the answer is A

upvoted 2 times

You want to configure autohealing for network load balancing for a group of Compute Engine instances that run in multiple zones, using the fewest possible steps.

You need to configure re-creation of VMs if they are unresponsive after 3 attempts of 10 seconds each. What should you do?

- A. Create an HTTP load balancer with a backend configuration that references an existing instance group. Set the health check to healthy (HTTP)
- B. Create an HTTP load balancer with a backend configuration that references an existing instance group. Define a balancing mode and set the maximum RPS to 10.
- C. Create a managed instance group. Set the Autohealing health check to healthy (HTTP)
- D. Create a managed instance group. Verify that the autoscaling setting is on.

Correct Answer: D

Community vote distribution

C (100%)

 **ReyBan**  3 years, 5 months ago

C, Agreed

reference : <https://cloud.google.com/compute/docs/tutorials/high-availability-autohealing>

Pro Tip: Use separate health checks for load balancing and for autohealing. Health checks for load balancing detect unresponsive instances and direct traffic away from them. Health checks for autohealing detect and recreate failed instances, so they should be less aggressive than load balancing health checks. Using the same health check for these services would remove the distinction between unresponsive instances and failed instances, causing unnecessary latency and unavailability for your users.

upvoted 85 times

 **ashrafh** 2 years, 1 month ago

I also vote for C

go to gcp console create a httpa load balancer and in the health check settings take your mouse to question mark it says

"Ensures that requests are sent only to instances that are up and running"

so its not recreating, if the vm not working it redirect to one which work.

go to gpc console create MIG and check the questions mark of Autohealing health check settings it says

"Autohealing allows recreating VM instances when needed. You can use a health check to recreate a VM instance if the health check finds it unresponsive. If you don't select a health check, Compute Engine will recreate VM instances only when they're not running."

hope this help :)

upvoted 20 times

 **bryanchew**  3 years, 6 months ago

A. Create an HTTP load balancer with a backend configuration that references an existing instance group. Set the health check to healthy(HTTP)

This is a possible answer. This answer assumes that the existing backend is configured correctly.

B. Create an HTTP load balancer with a backend configuration that references an existing instance group. Define a balancing mode and set the maximum RPS to 10.

This is a possible answer. This answer assumes that the existing backend is configured correctly. This answer adds an additional step over answer A.

C. Create a managed instance group. Set the Autohealing health check to healthy(HTTP)

This is only a partial solution. The default configuration is auto scaling enabled. You still need to create the HTTP Load Balancer.

D. Create a managed instance group. Verify that the auto scaling setting is on.

This is only a partial solution. Creating a Managed Instance Group with Auto Scaling is required, but you still need to create the HTTP Load Balancer.

Therefore the best answer is A in my opinion.

upvoted 21 times

 **Ridhanya** 1 year, 10 months ago

it cannot be option A because as you said, load balancer with the health check is already present and now the problem is simply auto healing. so we need to focus only on recreation which can happen only if option C is correct

upvoted 3 times

 **koniec** 2 years, 11 months ago

It's A.

Managed group already exists so create a LB with health checks.

If you go for C, you will have to create a LB anyway so it's more steps to achieve the goal

upvoted 1 times

✉ **tavva_prudhvi** 2 years, 6 months ago

https://www.youtube.com/watch?v=dT7xDtALPQ&list=PLIivdWyY5sqIij_cgINUHZDMnGjVx3rx&index=36

step-1: go to the instance group

step-2: click edit

step-3: scroll down you will see auto-healing off by default change to ON

step-4: create a health check saying 10 seconds as CHECK INTERVAL and UNHEALTHY THRESHOLD=3

upvoted 6 times

✉ **pYWORLD** 2 years, 2 months ago

I hope you saw that in the questions is stated: "network load balancing"

upvoted 1 times

✉ **DickDastardly** 2 years, 7 months ago

It can't be A as a load balancer does not re-create unhealthy instances, as per the requirement.

Has to be C

<https://cloud.google.com/compute/docs/instance-groups>

upvoted 6 times

✉ **YourCloudGuru** **Most Recent** 2 weeks, 3 days ago

The correct answer is C.

Managed instance groups are groups of homogeneous Compute Engine instances that are managed as a single entity. They can be used to distribute traffic across multiple instances and to provide high availability.

Autohealing is a feature of managed instance groups that automatically replaces instances that fail health checks. You can configure autohealing to recreate instances if they are unresponsive after a certain number of attempts.

To configure autohealing for network load balancing, you need to create a managed instance group and set the Autohealing health check to healthy (HTTP). The health check will periodically probe the instances in the group to see if they are responding. If an instance fails the health check, autohealing will recreate it.

upvoted 1 times

✉ **Captain1212** 1 month, 1 week ago

C is the right answer

upvoted 1 times

✉ **alexiscloud** 1 month, 3 weeks ago

I am preparing for GCP Associate engineer exam. If anyone has the entire questions set, please share them to rootkitlucifer@gmail.com

upvoted 1 times

✉ **230198** 1 month, 2 weeks ago

Hi, I am preparing for the same if you got the entire questions please share them at nandininadu23@gmail.com. Thank you so much.

upvoted 1 times

✉ **rejar** 2 months ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my rejaryoga099@gmail.com address. Thanks in advance!

upvoted 1 times

✉ **Mogamat** 2 months, 2 weeks ago

Answe C

upvoted 1 times

✉ **rosh199** 2 months, 3 weeks ago

I also vote for C

upvoted 1 times

✉ **Mamoon2023** 5 months ago

I am preparing for GCP Associate engineer exam. If anyone has the entire questions set, please share them to mamoon92@gmail.com

upvoted 1 times

✉ **Naresg** 6 months, 2 weeks ago

I am preparing for GCP ACE, able to access only 92 questions

Would really appreciate, If anyone can share the full set of questions to naresh.gownolla@gmail.com

upvoted 2 times

✉ **nooneknows** 6 months, 1 week ago

I can sell you by 20 USD!

upvoted 1 times

✉ **ALAMEEN1** 6 months ago

guaranteed to pass?

upvoted 1 times

✉ **Partha117** 6 months, 3 weeks ago

Selected Answer: C

That is the only option with autohealing which will recreate instance, just doing health check using http load balancer will not help to recreate, that can only redirect traffic

upvoted 3 times

 **Jackey0117** 7 months, 3 weeks ago

Selected Answer: C

A.B.D all don't enable Autohealing to recreate unresponsive VMs, so answer should be C.

upvoted 2 times

 **nikila99** 7 months, 4 weeks ago

Option A, "Create an HTTP load balancer with a backend configuration that references an existing instance group. Set the health check to healthy (HTTP)," is a valid solution but requires more steps to be performed to configure autohealing. Option B, "Create an HTTP load balancer with a backend configuration that references an existing instance group. Define a balancing mode and set the maximum RPS to 10," does not configure autohealing for unresponsive instances. Option D, "Create a managed instance group. Verify that the autoscaling setting is on," is not relevant to the autohealing of unresponsive instances.

upvoted 1 times

 **Buruguduystunstugudunstuy** 9 months, 2 weeks ago

Selected Answer: C

Option C is correct because creating a managed instance group allows you to use autohealing to automatically recreate VMs that are unresponsive after 3 attempts of 10 seconds each. You can set the Autohealing health check to healthy (HTTP) to specify the health check that determines whether the instances are considered healthy or not. If an instance becomes unresponsive, Autohealing will recreate the instance and attach it to the managed instance group.

<https://cloud.google.com/compute/docs/instance-groups/autohealing-instance-groups>

upvoted 2 times

 **Shivangi30** 3 months, 3 weeks ago

The link url is invalid

upvoted 1 times

 **sutharson4u** 9 months, 3 weeks ago

Selected Answer: C

C is right answer

upvoted 1 times

 **Suresh290587** 10 months, 1 week ago

A is the right answer

upvoted 2 times

 **saranaws22** 10 months, 1 week ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my sarankrishnan87@gmail.com address. Thanks in advance!

upvoted 1 times

You are using multiple configurations for gcloud. You want to review the configured Kubernetes Engine cluster of an inactive configuration using the fewest possible steps. What should you do?

- A. Use gcloud config configurations describe to review the output.
- B. Use gcloud config configurations activate and gcloud config list to review the output.
- C. Use kubectl config get-contexts to review the output.
- D. Use kubectl config use-context and kubectl config view to review the output.

Correct Answer: D

Reference:

<https://medium.com/google-cloud/kubernetes-engine-kubectl-config-b6270d2b656c>

Community vote distribution

D (90%) 10%

✉ **zukko78** Highly Voted 3 years, 5 months ago

D is correct

upvoted 41 times

✉ **nhusain** 2 years, 5 months ago

<https://medium.com/google-cloud/kubernetes-engine-kubectl-config-b6270d2b656c>

explains it well

upvoted 7 times

✉ **poogcp** Highly Voted 3 years, 4 months ago

C is correct , Use kubectl config get-contexts to review the output : shows the clusters and the configurations and based on the output we can identify the inactive configurations

upvoted 25 times

✉ **fracila** 10 months, 3 weeks ago

kubectl config get-contexts displays a list of contexts as well as the clusters that use them. Here's a sample output.

upvoted 2 times

✉ **Gurnoor** 3 years, 3 months ago

This is wrong get-contexts does not show clusters it only shows contexts.

upvoted 9 times

✉ **jilly** 3 years, 3 months ago

True .

Will give only below results

kubectl config get-contexts

CURRENT NAME CLUSTER AUTHINFO NAMESPACE

* white white dazwilkin

black black dazwilkin

upvoted 1 times

✉ **VanitaGinoya** Most Recent 2 weeks ago

if any one have all questions PDF then can you please send me on vanitasandeepginoya@gmail.com

upvoted 1 times

✉ **YourCloudGuru** 2 weeks, 3 days ago

Selected Answer: C

The correct answer is C. Use kubectl config get-contexts to review the output.

```
(awesome-wares-392903)$ kubectl config get-contexts
CURRENT NAME CLUSTER AUTHINFO NAMESPACE
```

The `kubectl config get-contexts` command lists all of the available contexts, including both active and inactive contexts. To review the configured Kubernetes Engine cluster of an inactive configuration, you can use this command to list all of the contexts, and then identify the inactive context that you want to review.

The other options are not as good. D is the closest one but it requires an additional step.

upvoted 1 times

✉ **CarlosMarin** 1 month ago

Selected Answer: C

I think is C, due to "using the fewest possible steps". In the D option we have 2 steps.

Both C and D are valid, but C is shorter

upvoted 1 times

✉ **Mogamat** 2 months, 2 weeks ago

Answe C

upvoted 1 times

✉ **Raikar123** 6 months ago

if anyone need full qtn pdf ping me 2 yashwanthrai94@gmail.com

upvoted 8 times

✉ **rahul2023** 1 month, 2 weeks ago

Hello Yashwanth, could you please send me the pdf to rahul201220132014@gmail.com.

Thanks in Advance.

upvoted 1 times

✉ **venkaveera95** 1 month, 3 weeks ago

Hello Yashwanth, could you please send me the pdf to venkaveera95@gmail.com. Thank you in advance.

upvoted 1 times

✉ **Shivamwalia** 1 month, 3 weeks ago

Hello Yashwanth, could you please send me the pdf to shivamwalia1698@gmail.com

upvoted 1 times

✉ **Zepopo** 1 month, 4 weeks ago

Hello Yashwanth, could you please send me the pdf to klicker.shou@gmail.com

upvoted 1 times

✉ **kavithaprasanth** 6 months, 3 weeks ago

Completed the exam of 17 Mar 2023. All questions were from this site. Thank you...

upvoted 6 times

✉ **ManvendraYadav** 5 months, 3 weeks ago

Hello, I'm preparing for GCP ACE certification. If anyone have the complete question set, could you please email it to yadav.manav1992@gmail.com

upvoted 2 times

✉ **ALAMEEN1** 6 months ago

u still have the questions, i will pay for it?

upvoted 1 times

✉ **Amrita_2601** 6 months, 1 week ago

Hello, Many congratulations, Have u practiced only questions till 93, or u have practiced a complete set of questions? As not able to access questions after 93.

upvoted 1 times

✉ **Amrita_2601** 6 months, 1 week ago

I have one ques also, in this site "most voted" and "correct answer" is diff for most of the ques, which one we should follow, pls guide me.

upvoted 1 times

✉ **AIMK** 6 months, 1 week ago

I recommend you to take "most voted" as a reference, several of the "correct answers" given by the site are in fact not correct. So after you think about what'd be your chosen answer, I advise you to go to the discussion and see what the community says, there's usually detailed information and sources of why A, B, C or D are the correct answers.

upvoted 1 times

✉ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: D

Answer A: Using `gcloud config configurations described` will only show you the details of the current configuration, not the Kubernetes Engine cluster of an inactive configuration.

Answer B: Using `gcloud config configurations activate` and `gcloud config list` to review the output will only show you the list of configurations and activate one of them, but it won't provide you with the details of the Kubernetes Engine cluster of an inactive configuration.

Answer C: Using `kubectl config get-contexts` will only list the available contexts, including their clusters, but it won't provide you with the details of the Kubernetes Engine cluster of an inactive configuration.

upvoted 2 times

✉ **abmohamed** 8 months, 3 weeks ago

The right answer is both C and D!

1. Show all cluster contexts:
\$ kubectl config get-contexts

2. Use the right context:
\$ kubectl config set-contexts <CONTEXT NAME>

3. Show the context details:
\$ kubectl config view

upvoted 1 times

✉ **abirroy** 1 year, 2 months ago

Selected Answer: D

Right answer is D

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

D is right

upvoted 1 times

✉ **guruanytime** 1 year, 4 months ago

D is right

upvoted 1 times

✉ **kiemkhach** 1 year, 5 months ago

D is correct. A lot details mentioned in this group. Here I only say about eliminating answers. As we go down to between C and D. The question is want to review a inactive configure. So, to me, C is viewing info about all configure while D is apply a specific config and viewing it. So I eliminate C to go with D.

upvoted 6 times

✉ **janukaz** 1 year, 6 months ago

Selected Answer: D

Answer bouncing among the C and D. Here I explain original results from C code,

If we executed "kubectl config get-contexts" code, Output will be,

CURRENT NAME | NAMESPACE | CLUSTER | AUTHINFO

upvoted 2 times

✉ **pranavdhopey** 1 year, 8 months ago

Selected Answer: D

D is correct

upvoted 2 times

✉ **NYF** 1 year, 8 months ago

Selected Answer: D

kubectl config view -o jsonpath='{.users[].name}' # display the first user

kubectl config view -o jsonpath='{.users[*].name}' # get a list of users

kubectl config get-contexts # display list of contexts

kubectl config current-context # display the current-context

kubectl config use-context my-cluster-name # set the default context to my-cluster-name

<https://kubernetes.io/docs/reference/kubectl/cheatsheet/>

upvoted 7 times

Your company uses Cloud Storage to store application backup files for disaster recovery purposes. You want to follow Google's recommended practices. Which storage option should you use?

- A. Multi-Regional Storage
- B. Regional Storage
- C. Nearline Storage
- D. Coldline Storage

Correct Answer: D

Reference:

<https://cloud.google.com/storage/docs/storage-classes#nearline>

Community vote distribution

D (78%) C (17%)

✉️  **poogcp**  3 years, 4 months ago

Best Answer is " Archive Storage "

<https://cloud.google.com/storage/docs/storage-classes>

But as per the given option next best solution is " Coldline Storage"

upvoted 54 times

✉️  **Mutune** 2 years, 7 months ago

Perfectly stated

upvoted 5 times

✉️  **zukko78**  3 years, 5 months ago

D is correct,

Coldline Storage COLDLINE 90 days

99.95% in multi-regions and dual-regions

99.9% in regions

upvoted 9 times

✉️  **YourCloudGuru**  2 weeks, 3 days ago

Selected Answer: D

Archive storage class is actually the recommended choice as mentioned in the Google Cloud Docs. but it is not an option here so the correct answer should be D. Coldline Storage as the cost is cheaper than Nearline.

Source: Google Cloud Doc: <https://cloud.google.com/storage/docs/storage-classes#archive>

upvoted 1 times

✉️  **Captain1212** 1 month, 1 week ago

Cold line storage

upvoted 1 times

✉️  **NoCrapEva** 1 month, 1 week ago

Selected Answer: C

Google document about setting up and configuring Google Backup and DR:

"Step 5 - Choose a default storage class and click Continue. Use nearline when retention is 30 days or less or coldline when retention is 90 days or more. If retention is between 30 and 90 days then consider using coldline."

<https://cloud.google.com/backup-disaster-recovery/docs/deployment/deployment-prep#expandable-2>

We are not told the retention policy (and cannot assume anything) - C gets my vote...

upvoted 1 times

✉️  **CalShark** 1 month, 3 weeks ago

D - Google docs state "Coldline storage is a very-low-cost, highly durable storage service for storing infrequently accessed data" - I think Archive Storage would be a better option but Coldline is the next closest

upvoted 1 times

✉️  **TSEMSOUS** 4 months ago

for disaster recovery i will choose Coldline storage

upvoted 1 times

✉️  **MrZed** 4 months ago

To follow Google's recommended practices for storing application backup files for disaster recovery purposes, the appropriate storage option would be:

B. Regional Storage.

Explanation:

Google recommends using Regional Storage for backup files because it provides both durability and lower latency for regional disaster recovery scenarios.

Regional Storage ensures that your data is stored redundantly across multiple zones within a specific region, providing high availability and resilience against zone-level failures.

It offers a balance between durability, availability, and cost-effectiveness, making it suitable for backup and disaster recovery purposes.

Regional Storage is typically more cost-effective than Multi-Regional Storage while still offering a high level of reliability.

Therefore, option B is the recommended choice when following Google's best practices for storing application backup files for disaster recovery purposes.

upvoted 2 times

 **mjbillings** 5 months ago

Cold line for disaster recovery

upvoted 1 times

 **AIMK** 6 months, 1 week ago

Selected Answer: D

there's some disagreement in the comments about this one. Going through some links about cloud storage I found this:

Cloud Storage Coldline: a low-latency storage class for long-term archiving

Coldline is a new Cloud Storage class designed for long-term archival and disaster recovery. Coldline is perfect for the archival needs of big data or multimedia content, allowing businesses to archive years of data. Coldline provides fast and instant (millisecond) access to data and changes the way that companies think about storing and accessing their cold data. Based on this, D is correct.

Refer to this link for more information: <https://cloud.google.com/blog/products/gcp/introducing-coldline-and-a-unified-platform-for-data-storage>

it explains the functions of each one of the options (Multi-regional, Regional, Nearline and Coldline)

upvoted 3 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: D

Correct choice after Archival

upvoted 1 times

 **smanoj85** 6 months, 3 weeks ago

Coldline Storage is the best answer for disaster recovery. When you create the GCS bucket in the GCP console, you can see the quote under the Coldline storage that, "Best for disaster recovery and data accessed less than once a quarter."

upvoted 1 times

 **Bedmed** 7 months ago

Nearline Storage is recommended for backup

upvoted 2 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: D

Answer D, Coldline Storage, is a good option for storing backup files for disaster recovery purposes for several reasons: Durability, Cost-effective, and Disaster recovery readiness.

upvoted 1 times

 **Vladimir_Sakhonchik** 7 months, 3 weeks ago

Google recommends using Cloud Storage's "Regional Storage" option for storing application backup files for disaster recovery purposes

upvoted 2 times

 **nikila99** 7 months, 4 weeks ago

Answer is A ,

Coldline Storage may be less expensive than other storage options, but it has higher latency and a longer retrieval time, which may not be suitable for disaster recovery purposes where data needs to be quickly retrieved and restored. For disaster recovery purposes, it is recommended to use Multi-Regional Storage which provides high availability and durability of data.

upvoted 1 times

 **JoniMONI** 8 months, 3 weeks ago

Selected Answer: C

The correct answer would depend on the specific requirements of the backup files.

If the backup files need to be accessed infrequently and data retrieval within several hours or days is acceptable, then Coldline Storage would be the most cost-effective option.

If the backup files need to be accessed within seconds and the cost is not a concern, then Nearline Storage would be a suitable option.

If the backup files need to be accessible at all times and the cost is not a concern, then Multi-Regional Storage would be the best option as it provides high durability and availability, as well as low latency access to your data.

It is important to consider the specific requirements of your backup files and choose the storage option that best meets those needs.

upvoted 3 times

Several employees at your company have been creating projects with Cloud Platform and paying for it with their personal credit cards, which the company reimburses. The company wants to centralize all these projects under a single, new billing account. What should you do?

- A. Contact cloud-billing@google.com with your bank account details and request a corporate billing account for your company.
- B. Create a ticket with Google Support and wait for their call to share your credit card details over the phone.
- C. In the Google Platform Console, go to the Resource Manager and move all projects to the root Organization.
- D. In the Google Cloud Platform Console, create a new billing account and set up a payment method.

Correct Answer: D

Reference:

<https://www.whizlabs.com/blog/google-cloud-interview-questions/>

Community vote distribution

D (77%) C (23%)

✉️  **samcat84**  3 years, 4 months ago

C is incomplete. Moving projects under an organization doesn't change their linked billing project.
<https://cloud.google.com/resource-manager/docs/migrating-projects-billing>

Note: The link between projects and billing accounts is preserved, irrespective of the hierarchy. When you move your existing projects into the organization they will continue to work and be billed as they used to before the migration, even if the corresponding billing account has not been migrated yet.

D is incomplete as well, after setting the billing account in the organization you need to link the projects to the new billing account.
upvoted 48 times

✉️  **sarahf** 2 years, 9 months ago

I agree that neither C or D is correct. I did the cert a month ago and this question was not on it. Although a similar question about how to change the payment method from your own card in your project to the company's "card". So they might have removed this one.
upvoted 2 times

✉️  **ehizren** 2 years, 8 months ago

What's was the answer you chose for your particular exam question?
upvoted 7 times

✉️  **GokulVelusaamy** 11 months, 1 week ago

We need to add a new payment method and need to set that as Primary, post that we need to remove the previous one
"If you want to remove a payment method, you should add a new payment method first."
Refer : <https://cloud.google.com/billing/docs/how-to/payment-methods>
upvoted 2 times

✉️  **RazOr** 1 year, 8 months ago

The given answers make D the only possible solution. C can not be right, you all need to look it up here:
https://cloud.google.com/resource-manager/docs/project-migration#change_billing_account
upvoted 3 times

✉️  **ryumada** 1 year, 2 months ago

This link explains clearly that move a project won't affect billing.
<https://cloud.google.com/resource-manager/docs/project-migration#permissions-billing>
upvoted 3 times

✉️  **poogcp**  3 years, 4 months ago

C is correct Answer. there will be 1 billing account for the organization and all projects under that organization are linked to single billing account.
upvoted 16 times

✉️  **Neha_Pallavi** 3 months, 2 weeks ago

The question is under the organization different projects are maintained the different cloud platforms.all the different project should single corporate bill account instead of the employee billing account. So try to update the corporate bill account details and mark it as primary for the all projects, post that employee account details need to removed. So suitable recommended option is D
upvoted 1 times

✉️  **arathefu** 1 year, 6 months ago

https://cloud.google.com/resource-manager/docs/project-migration#change_billing_account
"Moving a project from one organization to another won't impact billing, and charges will continue against the old billing account."
upvoted 6 times

✉️  **Goutom**  6 days, 2 hours ago

Hello, can anyone send me these questions in pdf or word file in my email at goutom825@gmail.com. I will be forever grateful.
upvoted 1 times

 **YourCloudGuru** 2 weeks, 3 days ago

Selected Answer: D

The correct answer is D. You have to follow the complete steps to successfully:

1. Create a new Billing Account
2. Move the existing projects into the new billing account
3. Cancel the earlier billing accounts of individual projects

This would meet all the requirements in the question - to centrally have all the projects under a single billing account.

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: D

D : as per question new account , d is the right answer

upvoted 1 times

 **Praxii** 5 months, 2 weeks ago

Selected Answer: C

I think this is multiple choice. C and D together make sense but otherwise they are both incomplete.

upvoted 1 times

 **aramisrocha** 9 months, 1 week ago

Selected Answer: D

D is correct, C is incorrect, because all accounts already is in root

upvoted 2 times

 **Buruguduystunstugudunstuy** 9 months, 1 week ago

Selected Answer: D

Option A is incorrect because you cannot request a corporate billing account by emailing cloud-billing@google.com. This email address is for general billing inquiries and support.

Option B is incorrect because you cannot create a ticket with Google Support to share your credit card details over the phone. To set up a payment method for a billing account, you must do it through the Google Cloud Platform Console.

Option C is incorrect because moving projects to the root organization will not create a new billing account. You must first create a new billing account and then move the projects to the root organization to ensure that they are all billed to the same billing account.

Therefore, the correct answer is Option D.

https://cloud.google.com/billing/docs/how-to/manage-billing-account#create_a_new_billing_account

upvoted 2 times

 **leogor** 11 months, 3 weeks ago

D, as other options are not correct

upvoted 2 times

 **hiromi** 11 months, 3 weeks ago

Selected Answer: D

D is more complete than C

upvoted 2 times

 **AwesomeGCP** 12 months ago

Selected Answer: D

It's pretty straight-forward: you should establish a new billing account with a company-based payment method. Then set all the projects to use that new billing account.

upvoted 2 times

 **Rajagopal** 1 year ago

I think "D" is the right answer (https://cloud.google.com/resource-manager/docs/project-migration#change_billing_account)

upvoted 3 times

 **Cornholio_LMC** 1 year ago

had this question today

upvoted 2 times

 **ale_brd_** 1 year, 1 month ago

Selected Answer: D

D is the correct one, even tho could be written in a better way.

upvoted 1 times

 **Qureshizaid64** 1 year, 1 month ago

Selected Answer: C

If we move all projects under the root organization hierarchy, they still need to modify to use a billing account within the organization.

Ref: https://cloud.google.com/resource-manager/docs/migrating-projects-billing#top_of_page Note: The link between projects and billing accounts is preserved, irrespective of the hierarchy. When you move your existing projects into the organization, they will continue to work and be billed as they used to before the migration, even if the corresponding billing account has not been migrated yet.

all its projects which is a Google recommended practice. So this is the better of the two options.

Ref: <https://cloud.google.com/billing/docs/concepts>

upvoted 2 times

 **akg001** 1 year, 1 month ago

Selected Answer: D

D. In the Google Cloud Platform Console, create a new billing account and set up a payment method.

upvoted 1 times

 **habros** 1 year, 2 months ago

D

<https://cloud.google.com/billing/docs/concepts>

upvoted 1 times

You have an application that looks for its licensing server on the IP 10.0.3.21. You need to deploy the licensing server on Compute Engine. You do not want to change the configuration of the application and want the application to be able to reach the licensing server. What should you do?

- A. Reserve the IP 10.0.3.21 as a static internal IP address using gcloud and assign it to the licensing server.
- B. Reserve the IP 10.0.3.21 as a static public IP address using gcloud and assign it to the licensing server.
- C. Use the IP 10.0.3.21 as a custom ephemeral IP address and assign it to the licensing server.
- D. Start the licensing server with an automatic ephemeral IP address, and then promote it to a static internal IP address.

Correct Answer: A

Community vote distribution

A (100%)

 **Khaled_Rashwan**  3 years, 4 months ago

A

IP 10.0.3.21 is internal by default, and to ensure that it will be static non-changing it should be selected as static internal ip address.
upvoted 25 times

 **riccamini** 10 months, 3 weeks ago

How do you know it is internal by default?

upvoted 1 times

 **yanlingmedal71** 10 months, 1 week ago

<https://cloud.google.com/vpc/docs/subnets#valid-ranges>

upvoted 4 times

 **zukko78**  3 years, 5 months ago

it's obvious, A

upvoted 23 times

 **YourCloudGuru**  2 weeks, 3 days ago

Selected Answer: A

The correct answer is A. Reserve the IP 10.0.3.21 as a static internal IP address using gcloud and assign it to the licensing server.

To reserve a static internal IP address, you can use the gcloud command-line tool. For example, to reserve the IP address 10.0.3.21, you would run the following command:

gcloud compute addresses reserve 10.0.3.21

Once you have reserved the static internal IP address, you can assign it to the licensing server by running the following command:

gcloud compute instances set-address licensing-server --address 10.0.3.21

Once you have assigned the static internal IP address to the licensing server, the application will be able to reach it using that IP address.
upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Internal IP address

upvoted 1 times

 **SanjeevKumar1983** 1 month, 2 weeks ago

Selected Answer: A

Correct answer is A

upvoted 1 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: A

Considering app is hosted in gcp internal ip address can be set as static

upvoted 1 times

 **Buruguduystunstugudunstuy** 9 months, 2 weeks ago

Selected Answer: A

The correct answer is Option A.

To deploy the licensing server on Compute Engine and ensure that the application can reach it, you should reserve the IP 10.0.3.21 as a static internal IP address and assign it to the licensing server.

By reserving IP 10.0.3.21 as a static internal IP address, you can ensure that the application can reach the licensing server at that IP address without changing the application's configuration.

To reserve the IP 10.0.3.21 as a static internal IP address and assign it to a Compute Engine instance using gcloud, you can use the following command:

```
gcloud compute instances create [INSTANCE_NAME] --address [IP_ADDRESS] --no-address
```

Replace [INSTANCE_NAME] with the name of the Compute Engine instance that you want to create, and [IP_ADDRESS] with the desired static internal IP address (in this case, 10.0.3.21).

The --no-address flag specifies that the instance should not be assigned a public IP address.

upvoted 4 times

 **DLink** 10 months ago

Selected Answer: A

Option A is right among choices

upvoted 1 times

 **riccamini** 10 months, 3 weeks ago

It's not stated whether the app is deployed on GCP. A or B are both correct in my opinion.

upvoted 2 times

 **FeeRoX** 8 months, 1 week ago

ip that starts with 10. is not a public IP. B = wrong.

upvoted 2 times

 **ashit44244** 1 year ago

Its not mentioned anywhere that the application is deployed in GCP or its external

upvoted 4 times

 **mennahibi** 5 months ago

Me too I voted for B and not A, because maybe the license server is open to the internet

upvoted 1 times

 **Chris_91_DE** 11 months, 3 weeks ago

I went into the same trap.

upvoted 2 times

 **AzureDP900** 1 year, 3 months ago

Static IP point to internal .. Option A is right

upvoted 1 times

 **rajsharma2108** 1 year, 4 months ago

A

obvious answer

upvoted 1 times

 **akshaychavan7** 1 year, 4 months ago

If they haven't mentioned that the application has been deployed internally(inside GCP) or externally then how can we determine if the server IP address should be internal or external? -_-

upvoted 5 times

 **knsujith** 1 year, 4 months ago

Selected Answer: A

A, is correct option

upvoted 1 times

 **Alejondri** 1 year, 5 months ago

Selected Answer: A

A, obvio

upvoted 1 times

 **pranavdhopey** 1 year, 8 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **pranavdhopey** 1 year, 8 months ago

Selected Answer: A

A is correct

upvoted 1 times

You are deploying an application to App Engine. You want the number of instances to scale based on request rate. You need at least 3 unoccupied instances at all times. Which scaling type should you use?

- A. Manual Scaling with 3 instances.
- B. Basic Scaling with min_instances set to 3.
- C. Basic Scaling with max_instances set to 3.
- D. Automatic Scaling with min_idle_instances set to 3.

Correct Answer: D

Reference:

<https://cloud.google.com/appengine/docs/standard/python/how-instances-are-managed>

Community vote distribution

D (80%) B (20%)

✉ **zukko78** Highly Voted 3 years, 5 months ago

D is correct.

App Engine supports the following scaling types, which controls how and when instances are created:

Automatic

Basic

Manual

You specify the scaling type in your app's app.yaml.

Automatic scaling

Automatic scaling creates instances based on request rate, response latencies, and other application metrics. You can specify thresholds for each of these metrics, as well as a minimum number instances to keep running at all times.

upvoted 50 times

✉ **vincent2023** 3 weeks, 1 day ago

<https://cloud.google.com/appengine/docs/legacy/standard/python/how-instances-are-managed>

upvoted 1 times

✉ **Finger41** Highly Voted 2 years, 5 months ago

D is correct : <https://cloud.google.com/appengine/docs/standard/go/config/appref>

"App Engine calculates the number of instances necessary to serve your current application traffic based on scaling settings such as target_cpu_utilization and target_throughput_utilization. Setting min_idle_instances specifies the number of instances to run in addition to this calculated number. For example, if App Engine calculates that 5 instances are necessary to serve traffic, and min_idle_instances is set to 2, App Engine will run 7 instances (5, calculated based on traffic, plus 2 additional per min_idle_instances)." upvoted 15 times

✉ **irfan_4141** Most Recent 4 days, 11 hours ago

Can anyone please share the complete list of question & Answers at muhammad.irfan.tahir93@gmail.com

upvoted 1 times

✉ **YourCloudGuru** 2 weeks, 3 days ago

Selected Answer: D

The correct answer is D. Automatic Scaling with min_idle_instances set to 3. By setting min_idle_instances to 3, you can ensure that there are always at least 3 instances available to handle new requests. The other options are not as good:

A. Manual Scaling requires you to manually adjust the number of instances running your application.

B. Basic Scaling is a simpler version of Automatic Scaling. It automatically scales the number of instances based on request rate, but it does not allow you to specify the minimum number of idle instances. This means that there is no guarantee that there will always be at least 3 instances available to handle new requests.

C. The max_instances setting specifies the maximum number of instances to keep running. By setting max_instances to 3, you are limiting the number of instances that your application can scale to. This is not ideal, especially if your application experiences sudden spikes in traffic.

upvoted 1 times

✉ **Captain1212** 1 month, 1 week ago

D, as per request it should be automatic right , so d is the right answer

upvoted 1 times

✉ **Ashi347** 5 months, 1 week ago

Can some one please email the complete set of questions and answers to shatakshi347@gmail.com. My exam is on May 9 would be a huge help! thanks!

upvoted 1 times

✉ **SharmaShubham** 5 months, 1 week ago

I have GCP ACE exam scheduled in one week. Request you please send me complete question and answer of this exam

Email - darkbolt.1991@gmail.com

upvoted 1 times

 **ShubhamSR** 3 months, 2 weeks ago

Pls send question to me have exam tomorrow.

rathorshubham68@gmail.com

upvoted 1 times

 **Rat98** 4 months, 3 weeks ago

Can someone send the full list of questions. I have exam on 23rd May

upvoted 1 times

 **Karthik_s_prof** 5 months, 2 weeks ago

I have GCP ACE exam scheduled on 29 April 2023. Request you please send me complete question and answer of this exam

Email - skarthik2997@gmail.com

upvoted 1 times

 **arnika98** 4 months, 3 weeks ago

did you complete the exam? did you pass?

upvoted 1 times

 **Ashi347** 5 months, 1 week ago

Did u get the complete list of questions?

upvoted 1 times

 **manishsingh213** 5 months, 2 weeks ago

I have GCP ACE exam scheduled on 27 April 2023. Request you please send me complete question and answer of this exam

Email - manishsingh213@outlook.com

upvoted 2 times

 **thiago286** 5 months, 3 weeks ago

Can anyone please share the complete list of question and answers at thiagogcp12@gmail.com? please!!

upvoted 1 times

 **SeeMakbul** 8 months, 1 week ago

Can anyone please share the complete list of question and answers at sweetsimmi1989@gmail.com?

upvoted 2 times

 **JoniMONI** 8 months, 3 weeks ago

Selected Answer: D

D. Automatic Scaling with min_idle_instances set to 3.

Automatic scaling adjusts the number of instances based on the request rate, while maintaining a minimum number of instances available. By setting min_idle_instances to 3, you ensure that at least 3 instances are running and available to handle requests, even when the request rate is low.

Manual scaling allows you to set a fixed number of instances, but does not automatically adjust based on the request rate. Basic scaling adjusts the number of instances based on the request rate, but does not allow you to set a minimum number of idle instances.

In order to keep at least 3 instances running and ready to handle requests, Automatic scaling with min_idle_instances set to 3 is the correct option.

upvoted 4 times

 **Buruguduystunstugudunstuy** 9 months, 2 weeks ago

Selected Answer: D

The correct answer is Option D.

To scale the number of instances based on request rate and ensure that there are always at least 3 unoccupied instances, you should use Automatic Scaling with min_idle_instances set to 3.

Automatic Scaling automatically scales the number of instances based on request rate and other metrics, such as CPU and memory utilization. By setting min_idle_instances to 3, you can ensure that the instance group maintains at least 3 idle instances at all times, ready to handle incoming requests.

upvoted 1 times

 **Saswato** 10 months ago

Hi , if anyone have all set of questions please send in the mail id : saswatoseal8@gmail.com

upvoted 1 times

 **glanshima** 10 months ago

Selected Answer: D

automatic scaling scales based on request rates, and you giveset min instances that you can run at all times

upvoted 1 times

 **cslince** 10 months ago

Selected Answer: B

B is correct

upvoted 2 times

 **jrlsl1991** 8 months, 2 weeks ago

Basic Scaling can't take min_instances as parameter.

upvoted 2 times

 **Sammydp2020** 8 months, 2 weeks ago

True. Reference: https://cloud.google.com/appengine/docs/legacy/standard/java/config/appref#scaling_elements

upvoted 2 times

 **sleoer** 10 months ago

Hi, am preparing for ACE exam which is scheduled on 15th Dec. 2022. If any one have all these set of questions, please send to stanin.leo21@gmail.com

upvoted 1 times

You have a development project with appropriate IAM roles defined. You are creating a production project and want to have the same IAM roles on the new project, using the fewest possible steps. What should you do?

- A. Use gcloud iam roles copy and specify the production project as the destination project.
- B. Use gcloud iam roles copy and specify your organization as the destination organization.
- C. In the Google Cloud Platform Console, use the 'create role from role' functionality.
- D. In the Google Cloud Platform Console, use the 'create role' functionality and select all applicable permissions.

Correct Answer: B

Reference:

<https://cloud.google.com/sdk/gcloud/reference/iam/roles/copy>*Community vote distribution*

A (100%)

coldpar Highly Voted 3 years, 6 months ago

Correct Answer is A not B
upvoted 37 times

Agents89 Highly Voted 3 years, 5 months ago

Correct answer is A
upvoted 23 times

Acielana Most Recent 1 week, 4 days ago

Selected Answer: A
Correct answer is A
upvoted 1 times

YourCloudGuru 2 weeks, 3 days ago

Selected Answer: A
The correct answer is A. Use gcloud iam roles copy and specify the production project as the destination project.

The gcloud iam roles copy command copies a role from one project to another. To use this command, you will need to know the name of the role that you want to copy and the name of the destination project.

For example, to copy the role roles/compute.instanceAdmin from the project my-dev-project to the project my-prod-project, you would run the following command:

```
gcloud iam roles copy roles/compute.instanceAdmin my-dev-project my-prod-project
```

This command will copy the role roles/compute.instanceAdmin to the project my-prod-project. The role will have the same permissions in the production project as it does in the development project.

upvoted 2 times

Captain1212 1 month, 1 week ago

Yes, A is the right answer
upvoted 1 times

rosh199 2 months, 3 weeks ago

Option A is the most appropriate approach to have the same IAM roles on the new production project as the existing development project, using the fewest possible steps.
upvoted 1 times

Praxii 5 months, 2 weeks ago

Correct Answer is A.
Reference:- <https://cloud.google.com/sdk/gcloud/reference/iam/roles/copy>
upvoted 2 times

thiago286 5 months, 3 weeks ago

Can anyone please share the complete list of question and answers at thiagogcp12@gmail.com? please!!
upvoted 1 times

dobberzoon 6 months, 1 week ago

Selected Answer: A
Correct Answer is A
upvoted 1 times

ch2023 7 months ago

Selected Answer: A

Correct answer is A
upvoted 2 times

✉ **ashtonez** 7 months ago

Selected Answer: A

A is correct, the command gcloud iam roles copy is for copying roles
Ref:
<https://cloud.google.com/sdk/gcloud/reference/iam/roles/copy>
upvoted 1 times

✉ **ast3citos** 7 months, 2 weeks ago

Selected Answer: A

To create a copy of an existing role into a project with PROJECT_ID, run:

```
gcloud iam roles copy --source="roles/ROLE" --destination=CustomSpannerDbAdmin --dest-project=PROJECT_ID
```

upvoted 2 times

✉ **Buruguduystunstugudunstuy** 9 months, 2 weeks ago

Selected Answer: A

The correct answer is Option A.

To create the same IAM roles in a production project as in a development project, using the fewest possible steps, you can use the gcloud iam roles copy command and specify the production project as the destination project.

The `gcloud iam roles copy` command allows you to copy IAM roles between projects or organizations. By specifying the production project as the destination project, you can copy the IAM roles from the development project to the production project.

Option B is incorrect because specifying your organization as the destination organization will copy the IAM roles to all projects within the organization, which is not what you want.

upvoted 9 times

✉ **processor** 9 months, 2 weeks ago

Selected Answer: A

Once again the site gave a wrong answer. we are asked to copy roles from a project to another. Not on the whole organisation !
upvoted 1 times

✉ **Onlinenow** 10 months, 1 week ago

Selected Answer: A

A should be correct. It's at project level and Iam role copy
upvoted 2 times

✉ **ChristN** 10 months, 1 week ago

Selected Answer: A

A could be possible if we were talking about organization in the question, But here, it's clearly specified "*project*"
"You have a development project with appropriate IAM roles defined. You are creating a production project and want to have the same IAM roles on the new project, using the fewest possible steps."
from google doc: <https://cloud.google.com/sdk/gcloud/reference/iam/roles/copy>

EXAMPLES

To create a copy of an existing role spanner.databaseAdmin into an organization with 1234567, run:

```
gcloud iam roles copy --source="roles/spanner.databaseAdmin" --destination=CustomViewer --dest-organization=1234567
```

To create a copy of an existing role spanner.databaseAdmin into a project with PROJECT_ID, run:

```
gcloud iam roles copy --source="roles/spanner.databaseAdmin" --destination=CustomSpannerDbAdmin --dest-project=PROJECT_ID
```

upvoted 4 times

✉ **romeaga2** 10 months, 2 weeks ago

Selected Answer: A

Correct the answer please
upvoted 1 times

✉ **Neha_Pallavi** 3 months, 1 week ago

B. roles need to be copied from project not from the organization.

upvoted 1 times

You need a dynamic way of provisioning VMs on Compute Engine. The exact specifications will be in a dedicated configuration file. You want to follow Google's recommended practices. Which method should you use?

- A. Deployment Manager
- B. Cloud Composer
- C. Managed Instance Group
- D. Unmanaged Instance Group

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/instances/>

Community vote distribution

A (94%) 6%

✉ **Buruguduystunstugudunstuy** Highly Voted 9 months, 2 weeks ago

Selected Answer: A

The correct answer is Option A - Deployment Manager. Deployment Manager is a configuration management tool that allows you to define and deploy a set of resources, including Compute Engine VMs, in a declarative manner. You can use it to specify the exact specifications of your VMs in a configuration file, and Deployment Manager will create and manage those VMs for you. Deployment Manager is recommended by Google as a way to automate and manage the deployment of resources on the Google Cloud Platform.

<https://cloud.google.com/deployment-manager/docs/>
upvoted 12 times

✉ **Captain1212** Most Recent 1 month, 1 week ago

Yes, deployment manager is the right answer, it helps you to configure resources as per your file
upvoted 1 times

✉ **sthapit** 2 months ago

Please someone email me all these questions : ssthaphit977@gmail.com

TIA

upvoted 1 times

✉ **Shivamwalia** 1 month, 3 weeks ago

Dis u get the pdf?
upvoted 1 times

✉ **snehachandru** 2 weeks, 5 days ago

Hello, Can you mail me the pdf to this id sneha.chandlal@gmail.com
upvoted 1 times

✉ **SunTzet** 1 month, 1 week ago

Hello, can you also email me the pdf? forsyth.kirell@hotmail.com
upvoted 1 times

✉ **oscardev04** 1 week, 6 days ago

Hello, can you also email me the pdf? oscar.dev04@gmail.com
upvoted 1 times

✉ **mmierke** 2 months, 1 week ago

Selected Answer: C

I would go with C. A has a broader scope, but C is the correct answer for VM's.
upvoted 2 times

✉ **Rkraj** 2 months, 2 weeks ago

Selected Answer: A

correct is A
upvoted 1 times

✉ **ExamsFR** 2 months, 3 weeks ago

Selected Answer: A

The correct answer is : A
upvoted 1 times

✉ **rosh199** 2 months, 3 weeks ago

The correct answer is:

A. Deployment Manager
upvoted 1 times

✉ **himesh046** 5 months, 2 weeks ago

Can anyone please share the complete list of questions at - himeshyadav1994@gmail.com
Thank you!!!
upvoted 1 times

✉ **thiago286** 5 months, 3 weeks ago

Can anyone please share the complete list of question and answers at thiagogcp12@gmail.com? please!!
upvoted 1 times

✉ **AdamCaster** 5 months, 4 weeks ago

Selected Answer: A

A for sure

upvoted 1 times

✉ **esqandares** 6 months, 2 weeks ago

any one can share why is C
upvoted 3 times

✉ **sakdip66** 6 months ago

Managed Instance Group (option C) and Unmanaged Instance Group (option D) are Compute Engine features that allow you to group related VM instances and manage them as a single entity. However, they do not provide a dynamic way of provisioning VMs based on a configuration file like Deployment Manager does.

upvoted 5 times

✉ **Eshan_18** 6 months, 2 weeks ago

Hello everyone, Can anyone please share full questions list to lokeshhurts943@gmail.com, it will very helpful for me, thanks in advance!
upvoted 1 times

✉ **[Removed]** 4 months ago

Hello everyone, Can anyone please share full questions list to samruddhijagtap23@gmail.com, it will very helpful for me, thanks in advance!
upvoted 1 times

✉ **kanavpeer** 6 months, 2 weeks ago

Hello everyone, Can anyone please share full questions list to kanavrz.peer@gmail.com, it will very helpful for me, thanks in advance!
upvoted 1 times

✉ **shubhamu571** 6 months ago

Hello everyone, Can anyone please share full questions list to shubhamu575@gmail.com, it will very helpful for me, thanks in advance!
upvoted 1 times

✉ **23_7k** 9 months ago

Hi which answer we need to select? from discussion or website answer? could you please tell me?
upvoted 2 times

✉ **InigoGutierrez** 8 months ago

Discussion is normally the correct one
upvoted 2 times

✉ **ninjaasmoke** 10 months, 1 week ago

A.
Explained here: <https://cloud.google.com/deployment-manager/docs/configuration/create-basic-configuration>
upvoted 1 times

✉ **ChristN** 10 months, 1 week ago

A is the answer. We are talking about a dedicated config file. <https://cloud.google.com/deployment-manager/docs>
upvoted 2 times

✉ **Omkarshingare** 8 months, 1 week ago

Hi Chris, Can u please share full questions list to warcrycreature@gmail.com, it will very helpful for me, thanks in advance!
upvoted 1 times

✉ **romega2** 10 months, 2 weeks ago

Selected Answer: A

A is okay

upvoted 2 times

✉ **leogor** 11 months, 3 weeks ago

A. deployment manager,
instance template / group is wrong
upvoted 1 times

You have a Dockerfile that you need to deploy on Kubernetes Engine. What should you do?

- A. Use kubectl app deploy <dockerfilename>.
- B. Use gcloud app deploy <dockerfilename>.
- C. Create a docker image from the Dockerfile and upload it to Container Registry. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.
- D. Create a docker image from the Dockerfile and upload it to Cloud Storage. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.

Correct Answer: C

Reference -

<https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app>

Community vote distribution

C (97%)

 **Agents89**  3 years, 5 months ago

C is correct

upvoted 40 times

 **Buruguduystunstugudunstuy**  9 months, 2 weeks ago

Selected Answer: C

The correct answer is Option C. To deploy a Docker container on Kubernetes Engine, you should first create a Docker image from the Dockerfile and push it to Container Registry, which is a fully-managed Docker container registry that makes it easy for you to store, manage, and deploy Docker container images. Then, you can create a Deployment YAML file that specifies the image to use and other desired deployment options, and use the kubectl command-line tool to create the deployment based on the YAML file.

Option A is incorrect because kubectl app deploy is not a valid command.

Option B is incorrect because gcloud app deploy is used to deploy applications to App Engine, not Kubernetes Engine.

Option D is incorrect because it involves storing the image in Cloud Storage rather than Container Registry.

<https://cloud.google.com/kubernetes-engine/docs/how-to/deploying-a-container>

upvoted 16 times

 **ast3citos** 7 months, 2 weeks ago

The link you provided has stopped working :_(

upvoted 3 times

 **dookiecloud**  1 month ago

C is correct.

We need to build docker image then push to Container Registry and setup yaml deployment in GKE to pull our registry image.
upvoted 1 times

 **Captain1212** 1 month, 1 week ago

c is correct , , you have to first upload it in the continer registry
upvoted 1 times

 **thiago286** 5 months, 3 weeks ago

Can anyone please share the complete list of question and answers at thiagogcp12@gmail.com? please!!
upvoted 1 times

 **sajity** 6 months, 4 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

 **ChristN** 10 months, 1 week ago

Selected Answer: C

C is correct.

upvoted 1 times

 **leogor** 11 months, 3 weeks ago

c, u need to upload to Container Registry before deploying
upvoted 1 times

 **hiromi** 11 months, 3 weeks ago

Selected Answer: C

C is correct, cuz you need to upload to Container Registry at first
upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Answer is C
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

C is right
upvoted 1 times

 **Biksthrow** 1 year, 3 months ago

C is correct answer
upvoted 1 times

 **wjtb** 1 year, 4 months ago

Selected Answer: C

C is correct.
A can be eliminated because kubectl app * is not a valid command
B can be eliminated because gcloud app deploy deploys on app engine, not on kubernetes (also it still requires a config file pointing to the image).
D is not correct, since you cannot deploy a container image directly from GCS
upvoted 5 times

 **haroldbenites** 1 year, 4 months ago

Go for C
upvoted 1 times

 **Prassanth** 1 year, 4 months ago

Selected Answer: C

C is correct
upvoted 1 times

 **Chile** 1 year, 5 months ago

C is correct
upvoted 1 times

 **pgb54** 1 year, 7 months ago

Selected Answer: C

C is correct
upvoted 1 times

Your development team needs a new Jenkins server for their project. You need to deploy the server using the fewest steps possible. What should you do?

- A. Download and deploy the Jenkins Java WAR to App Engine Standard.
- B. Create a new Compute Engine instance and install Jenkins through the command line interface.
- C. Create a Kubernetes cluster on Compute Engine and create a deployment with the Jenkins Docker image.
- D. Use GCP Marketplace to launch the Jenkins solution.

Correct Answer: D

Reference:

<https://cloud.google.com/solutions/using-jenkins-for-distributed-builds-on-compute-engine>

Community vote distribution

D (100%)

 **Agents89**  3 years, 5 months ago

D is correct

upvoted 32 times

 **Buruguduystunstugudunstuy**  9 months, 2 weeks ago

Selected Answer: D

The correct answer is Option D. By using GCP Marketplace to launch the Jenkins solution, you can quickly deploy a Jenkins server with minimal steps.

Option A involves deploying the Jenkins Java WAR to App Engine Standard, which requires more steps and may not be suitable for your requirements.

Option B involves creating a new Compute Engine instance and manually installing Jenkins, which also requires more steps.

Option C involves creating a Kubernetes cluster and creating a deployment with the Jenkins Docker image, which again involves more steps and may not be the most efficient solution.

upvoted 11 times

 **YourCloudGuru**  2 weeks ago

The correct answer is D

This is the fastest and easiest way to deploy a Jenkins server on GCP. GCP Marketplace provides pre-configured and pre-packaged applications that you can launch with just a few clicks.

To deploy Jenkins from GCP Marketplace, follow these steps:

- >Go to the GCP Marketplace.
- >Search for "Jenkins".
- >Click the "Jenkins" app listing.
- >Click the "Launch" button.
- >Configure the Jenkins deployment options, such as the machine type and region.
- >Click the "Deploy" button.

GCP Marketplace will deploy a Jenkins server to your GCP project. Once the deployment is complete, you can access the Jenkins web UI at the URL provided in the deployment details.

upvoted 1 times

 **elviskimutai** 3 weeks, 4 days ago

D. Use GCP Marketplace to launch the Jenkins solution: This is the most straightforward and efficient method to deploy Jenkins on Google Cloud Platform

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

D is the correct answer , fewest steps possible

upvoted 1 times

 **vivekvyj** 5 months, 3 weeks ago

Selected Answer: D

D IS CORRECT

upvoted 1 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: D

Marketplace gives jenkins image to use

upvoted 1 times

 **dkorzel** 7 months, 1 week ago

Selected Answer: D

D is correct

upvoted 1 times

 **sk12345** 8 months, 3 weeks ago

D is correct answer

upvoted 1 times

 **anand1210** 9 months, 4 weeks ago

D is correct

upvoted 1 times

 **geedhey** 10 months ago

D is the correct answer

upvoted 1 times

 **cslince** 10 months, 1 week ago

D is correct

upvoted 1 times

 **ChristN** 10 months, 1 week ago

Selected Answer: D

D is correct.

Market place makes things easier.

upvoted 1 times

 **Tx4free** 10 months, 3 weeks ago

D is correct since Jenkins is a produce on GCP

upvoted 1 times

 **leogor** 11 months, 3 weeks ago

D, marketplace has do it for u

upvoted 1 times

 **PKookNN** 12 months ago

Selected Answer: D

D is correct one

upvoted 2 times

 **Cornholio_LMC** 1 year ago

had this question today

upvoted 2 times

You need to update a deployment in Deployment Manager without any resource downtime in the deployment. Which command should you use?

- A. gcloud deployment-manager deployments create --config <deployment-config-path>
- B. gcloud deployment-manager deployments update --config <deployment-config-path>
- C. gcloud deployment-manager resources create --config <deployment-config-path>
- D. gcloud deployment-manager resources update --config <deployment-config-path>

Correct Answer: B

Reference:

<https://cloud.google.com/sdk/gcloud/reference/deployment-manager/deployments/update>

Community vote distribution

B (100%)

 [Removed]  3 years ago

B is correct Additional tip, update and create resource is not even a command under deployment management service.
upvoted 43 times

 Agents89  3 years, 5 months ago

B is correct
upvoted 39 times

 YourCloudGuru  2 weeks ago

Selected Answer: B

The correct answer is B.

This command updates an existing deployment with the configuration specified in the deployment-config-path file. Deployment Manager automatically determines whether resources need to be created, updated, or deleted in order to apply the changes.

Here is an example of how to use the gcloud deployment-manager deployments update command to update a deployment without any resource downtime:

```
gcloud deployment-manager deployments update my-deployment --config deployment.yaml
```

This command will update the my-deployment deployment with the configuration specified in the deployment.yaml file. Deployment Manager will automatically determine whether resources need to be created, updated, or deleted in order to apply the changes.
upvoted 1 times

 Captain1212 1 month, 1 week ago

b is the answer, clearly
upvoted 1 times

 sakdip66 6 months ago

I agree B is the 
upvoted 1 times

 Partha117 6 months, 3 weeks ago

Selected Answer: B
gcloud deployment-manager deployments update --config is correct
upvoted 1 times

 smg1989 7 months, 4 weeks ago

B will be correct answer
upvoted 1 times

 owliguess 8 months, 1 week ago

Selected Answer: B
should not use 'create', so B or D is correct.... keywords 'resource' is invalid, so B is the ans.
upvoted 1 times

 Buruguduystunstugudunstuy 9 months, 2 weeks ago

Selected Answer: B
The correct answer is Option B: `gcloud deployment-manager deployments update --config <deployment-config-path>`. This command updates an existing deployment with the configuration specified in the `deployment-config-path` file. It allows you to make changes to the deployment without any downtime in the resources.

<https://cloud.google.com/sdk/gcloud/reference/deployment-manager/>

upvoted 5 times

 **cslince** 10 months, 1 week ago

B is correct

upvoted 1 times

 **leogor** 11 months, 3 weeks ago

should not use 'create', so B or D is correct
and keyword 'resource' is invalid, so B is the ans

upvoted 4 times

 **PKookNN** 12 months ago

Selected Answer: B

B is correct

upvoted 2 times

 **eloyus** 1 year ago

Selected Answer: B

B is correct

upvoted 1 times

 **krnagg** 1 year, 2 months ago

B is correct. C and D are invalid commands and A is to create a new deployment.

upvoted 2 times

 **RanjithK** 1 year, 3 months ago

Answer is B

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is right to update the deployment.

upvoted 1 times

 **Baan** 1 year, 4 months ago

B is correct

upvoted 1 times

You need to run an important query in BigQuery but expect it to return a lot of records. You want to find out how much it will cost to run the query. You are using on-demand pricing. What should you do?

- A. Arrange to switch to Flat-Rate pricing for this query, then move back to on-demand.
- B. Use the command line to run a dry run query to estimate the number of bytes read. Then convert that bytes estimate to dollars using the Pricing Calculator.
- C. Use the command line to run a dry run query to estimate the number of bytes returned. Then convert that bytes estimate to dollars using the Pricing Calculator.
- D. Run a select count (*) to get an idea of how many records your query will look through. Then convert that number of rows to dollars using the Pricing Calculator.

Correct Answer: B

Reference:

<https://cloud.google.com/bigquery/docs/estimate-costs>

Community vote distribution

B (100%)

 **ESP_SAP** Highly Voted 3 years, 1 month ago

Correct Answers is (B):

On-demand pricing

Under on-demand pricing, BigQuery charges for queries by using one metric: the number of bytes processed (also referred to as bytes read). You are charged for the number of bytes processed whether the data is stored in BigQuery or in an external data source such as Cloud Storage, Drive, or Cloud Bigtable. On-demand pricing is based solely on usage.

https://cloud.google.com/bigquery/pricing#on_demand_pricing

upvoted 35 times

 **Agents89** Highly Voted 3 years, 5 months ago

B is Correct

upvoted 34 times

 **YourCloudGuru** Most Recent 2 weeks ago

Selected Answer: B

The correct answer is B.

This option is the most accurate way to estimate the cost of your query, because it takes into account the actual number of bytes that will be processed by BigQuery.

Here is an example of how to run a dry run query in BigQuery:

```
bq dry run --query "SELECT * FROM <dataset>. <table> WHERE <condition>"
```

This command will print the estimated number of bytes that will be processed by BigQuery. You can then use the Pricing Calculator to convert this bytes estimate to dollars.

Once you have estimated the cost of your query, you can decide whether or not to proceed with running it. If you decide to proceed, you can monitor the cost of your query using the BigQuery Monitoring Console.

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

yes B is the correct answer, big query charges for queries by using one metric

upvoted 1 times

 **certified28** 4 months ago

Selected Answer: B

B is Correct

upvoted 1 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: B

Calculation should be on bytes read

upvoted 3 times

 **BlueJay20** 8 months, 1 week ago

Selected Answer: B

Calculation on bytes read.

upvoted 1 times

✉ **aramisrocha** 9 months, 1 week ago

Selected Answer: B

B, You need check bytes read

upvoted 2 times

✉ **presi** 4 months ago

Obviously ,the best answer here is B and not c.

upvoted 1 times

✉ **Buruguduystunstugudunstuy** 9 months, 2 weeks ago

Selected Answer: B

The correct answer is Option B. Running a dry run query using the command line can estimate the number of bytes read by the query, which can then be used to estimate the cost of running the query using the Pricing Calculator. To estimate the cost of a BigQuery query, you can use the `bq` command-line tool to run a dry-run query.

Option A: Arranging to switch to Flat-Rate pricing will not help you estimate the cost of running the query using on-demand pricing.

Option C: Estimating the number of bytes returned by the query will not give you an accurate estimate of the cost of running the query using on-demand pricing.

Option D: Estimating the number of rows that the query will look through will not give you an accurate estimate of the cost of running the query using on-demand pricing.

<https://cloud.google.com/bigquery/docs/estimate-costs>

upvoted 5 times

✉ **Purshartha** 10 months, 1 week ago

Selected Answer: B

number of bytes read

upvoted 2 times

✉ **leogor** 11 months, 3 weeks ago

B, READ but not RETURNED

upvoted 2 times

✉ **RanjithK** 1 year, 3 months ago

Answer is B

upvoted 2 times

✉ **AzureDP900** 1 year, 3 months ago

B is right . number bytes read instead of number of byte returned.

upvoted 2 times

✉ **haroldbenites** 1 year, 4 months ago

Go for B

upvoted 3 times

✉ **Alejondri** 1 year, 5 months ago

Selected Answer: B

B Is correct

upvoted 2 times

✉ **DanielB96** 1 year, 5 months ago

Selected Answer: B

B is correct :)

upvoted 1 times

✉ **naresh136** 1 year, 7 months ago

B is correct answer.

upvoted 1 times

You have a single binary application that you want to run on Google Cloud Platform. You decided to automatically scale the application based on underlying infrastructure CPU usage. Your organizational policies require you to use virtual machines directly. You need to ensure that the application scaling is operationally efficient and completed as quickly as possible. What should you do?

- A. Create a Google Kubernetes Engine cluster, and use horizontal pod autoscaling to scale the application.
- B. Create an instance template, and use the template in a managed instance group with autoscaling configured.
- C. Create an instance template, and use the template in a managed instance group that scales up and down based on the time of day.
- D. Use a set of third-party tools to build automation around scaling the application up and down, based on Stackdriver CPU usage monitoring.

Correct Answer: B

Community vote distribution

B (100%)

 **mohdaifiuddin**  2 years, 9 months ago

I'll take a simple and logical approach for answering this.
Let's first break down the question into key requirements -

1. automatically scale the application based on underlying infrastructure CPU usage.
2. use virtual machines directly.

A. Not feasible because VMs are not used directly here.
B. This is the correct answer.
C. Time of Day... Easy elimination because this does not scale on CPU usage and time of day is mentioned NOWHERE.
D. Third Party Tools.... Nobody would use GCP if they needed third party tools to do something as simple as scaling based on CPU usage.
all popular cloud providers have native solutions for this including GCP.

upvoted 62 times

 **kopper2019** 2 years, 6 months ago

and also D is out because why would I use a third party tool when is a GCP exam

upvoted 15 times

 **RM0000** 2 years ago

If the resource/solution is not available. It's a possibility.

upvoted 4 times

 **coldpar**  3 years, 6 months ago

correct is B as you have to use VM instances directly.

upvoted 53 times

 **YourCloudGuru**  2 weeks ago

Selected Answer: B

The correct answer is B.

This option is the most efficient way to scale your application based on CPU usage, because it uses Google Cloud's built-in autoscaling capabilities. Autoscaling allows you to specify a minimum and maximum number of instances, and Google Cloud will automatically add or remove instances as needed to maintain your desired CPU utilization.

Options A, C, and D are not as efficient, because they require more manual intervention to scale your application.

Here are the steps to create a managed instance group with autoscaling configured:

1. Create an instance template.
2. Create a managed instance group from the instance template.
3. Configure autoscaling for the managed instance group.

Once you have configured autoscaling, the managed instance group will automatically add or remove instances as needed to maintain your desired CPU utilization.

upvoted 1 times

 **tombatlkap** 3 weeks, 4 days ago

Selected Answer: B

It sounds appropriate.

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

B, is the right answer because, in b you can use the vm's directly and it will be autoconfigured
upvoted 1 times

 **Paras_vohrA** 2 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times

✉ **BlueJay20** 8 months, 1 week ago

Selected Answer: B

B correct answer.
upvoted 2 times

✉ **Buruguduystunstugudunstuy** 9 months, 2 weeks ago

Selected Answer: B

The correct answer is Option B. Creating an instance template and using it in a managed instance group with autoscaling configured will allow you to automatically scale the application based on underlying infrastructure CPU usage and will be operationally efficient and completed quickly.

Option A is incorrect because it involves using Kubernetes, which is not required in this scenario.

Option C is incorrect because it involves scaling based on the time of day, which is not specified as a requirement.

Option D involves using third-party tools and is not necessary for this scenario.
upvoted 5 times

✉ **_adiii** 10 months ago

B, as MIG with autoscaling is best choice
upvoted 1 times

✉ **cslince** 10 months, 1 week ago

correct is B
upvoted 1 times

✉ **ChristN** 10 months, 1 week ago

Selected Answer: B
automatic scale
upvoted 1 times

✉ **leogor** 11 months, 3 weeks ago

B, managed instance group (VM instances) with autoscaling
upvoted 1 times

✉ **gcp_world123** 1 year, 1 month ago

Our requirements are as per the question

1. Use Virtual Machines directly (i.e. not container-based)
2. Scale Automatically
3. Scaling is efficient & is quick

B is correct

Managed instance groups offer autoscaling capabilities that let you automatically add or delete instances from a managed instance group based on increases or decreases in load (CPU Utilization in this case). Autoscaling helps your apps gracefully handle increases in traffic and reduce costs when the need for resources is lower. You define the autoscaling policy and the autoscaler performs automatic scaling based on the measured load (CPU Utilization in this case). Autoscaling works by adding more instances to your instance group when there is more load (upscale), and deleting instances when the need for instances is lowered (downscale).

Ref: <https://cloud.google.com/compute/docs/autoscaler>
upvoted 3 times

✉ **RanjithK** 1 year, 3 months ago

Answer is B
upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

B is perfect for given scenario.
upvoted 1 times

✉ **jagan_cloud** 1 year, 4 months ago

option B is correct
upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago

Go for B
upvoted 1 times

You are analyzing Google Cloud Platform service costs from three separate projects. You want to use this information to create service cost estimates by service type, daily and monthly, for the next six months using standard query syntax. What should you do?

- A. Export your bill to a Cloud Storage bucket, and then import into Cloud Bigtable for analysis.
- B. Export your bill to a Cloud Storage bucket, and then import into Google Sheets for analysis.
- C. Export your transactions to a local file, and perform analysis with a desktop tool.
- D. Export your bill to a BigQuery dataset, and then write time window-based SQL queries for analysis.

Correct Answer: D

Community vote distribution

D (100%)

✉  **mohdafiuddin**  2 years, 9 months ago

Solving this by first eliminating the options that don't suit us. By breaking down the question into the key requirements-

1. Analyzing Google Cloud Platform service costs from three separate projects.
2. Using standard query syntax. -> (Relational data and SQL)

A. 'Cloud Storage bucket'.....'Cloud Bigtable'. Not feasible, mainly because cloud BigTable is not good for Structured Data (or Relational Data on which we can run SQL queries as per the question's requirements). BigTable is better suited for Semi Structured data and NoSQL data.
 B. 'Cloud Storage bucket'....'Google Sheets'. Not Feasible because there is no use of SQL in this option, which is one of the requirements.
 C. Local file, external tools... this is automatically eliminated because the operation we need is simple, and there has to be a GCP native solution for this. We shouldn't need to rely on going out of the cloud for such a simple thing.
 D. 'BigQuery'....'SQL queries' -> This is the right answer.

upvoted 93 times

✉  **ryumada** 1 year, 2 months ago

Cloud billing data can only be exported to a JSON local file and to Bigquery. So, using Cloud Storage to export cloud billing data is not possible to do.

<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>

upvoted 14 times

✉  **cesar7816**  3 years, 7 months ago

Agreed, BigQuery

upvoted 18 times

✉  **yurstev** 2 years, 9 months ago

the key is standard query syntax

upvoted 3 times

✉  **YourCloudGuru**  2 weeks ago

Selected Answer: D

The correct answer is D.

BigQuery is a fully-managed, petabyte-scale analytics data warehouse that enables businesses to analyze all their data very quickly. It is also a good choice for analyzing cost data because it can handle large amounts of data and can perform complex queries quickly.

Time window-based SQL queries allow you to analyze data over a specific time period. For example, you could write a query to calculate the total cost of a particular service for each day of the month.

Here is an example of a BigQuery query that you could use to calculate the total cost of a particular service for each day of the month:

```
sql
SELECT
cost,
DATE(usage_start_time) AS date
FROM
`[PROJECT_ID].billing.dataset`
WHERE
service_id = 'YOUR_SERVICE_ID'
GROUP BY
date
ORDER BY
date
```

This query will return a table with two columns: `cost` and `date`. The `cost` column will contain the total cost of the service for each day of the month. The `date` column will contain the date for each row.

upvoted 1 times

✉  **Captain1212** 1 month, 1 week ago

yes D is the correct answer because we only use bigquery for semi or structured data

upvoted 1 times

✉ **Ashish_Tayal** 6 months, 1 week ago

Selected Answer: D

In GCP, Always use Big Query to export Billing. And Big Query is best for Analyzing

upvoted 3 times

✉ **Partha117** 6 months, 3 weeks ago

Selected Answer: D

Big Query will allow sql analysis

upvoted 2 times

✉ **Buruguduystunstugudunstuy** 9 months, 2 weeks ago

Selected Answer: D

The correct answer is Option D. Exporting the bill to a BigQuery dataset allows you to use SQL queries to analyze the data and create service cost estimates by service type, daily and monthly, for the next six months. This is an efficient and effective way to analyze the data, especially if you are familiar with SQL syntax.

Option A, importing the bill into Cloud Bigtable, may be more complex and may not offer the same level of flexibility as using SQL queries in BigQuery.

Option B, importing the bill into Google Sheets, may be more suitable for simple analysis, but may not be as efficient for more complex analysis.

Option C, exporting the transactions to a local file and using a desktop tool, may not be as efficient or effective as using a cloud-based solution like BigQuery.

<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>

<https://cloud.google.com/bigquery/docs/reference/standard-sql/>

upvoted 5 times

✉ **Kopy** 11 months, 1 week ago

Selected Answer: D

Agreed, BigQuery

upvoted 1 times

✉ **leogor** 11 months, 3 weeks ago

D, using standard query syntax so use BQ

upvoted 1 times

✉ **Haarish** 1 year, 2 months ago

Agreed, BigQuery

Ans: D

upvoted 1 times

✉ **RanjithK** 1 year, 3 months ago

Answer is D

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

D is the right answer.

upvoted 1 times

✉ **denkyira** 1 year, 3 months ago

I agreed, BigQuery Answer D is correct

upvoted 1 times

✉ **jagan_cloud** 1 year, 4 months ago

answer is Option D

upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago

Go for D

upvoted 1 times

✉ **LaxmanTiwari** 1 year, 4 months ago

Agreed , answer is D

upvoted 1 times

✉ **shawnkkk** 1 year, 10 months ago

D. Export your bill to a BigQuery dataset, and then write time window-based SQL queries for analysis.

upvoted 2 times

You need to set up a policy so that videos stored in a specific Cloud Storage Regional bucket are moved to Coldline after 90 days, and then deleted after one year from their creation. How should you set up the policy?

- A. Use Cloud Storage Object Lifecycle Management using Age conditions with SetStorageClass and Delete actions. Set the SetStorageClass action to 90 days and the Delete action to 275 days (365 - 90)
- B. Use Cloud Storage Object Lifecycle Management using Age conditions with SetStorageClass and Delete actions. Set the SetStorageClass action to 90 days and the Delete action to 365 days.
- C. Use gsutil rewrite and set the Delete action to 275 days (365 - 90).
- D. Use gsutil rewrite and set the Delete action to 365 days.

Correct Answer: A

Community vote distribution

B (71%)	A (29%)
---------	---------

 **Sammigbo**  3 years, 4 months ago

Answer is B. There should be no reason to recalculate the time needed to delete after a year.
upvoted 54 times

 **JKRowlings** 2 years, 10 months ago

The correct ans is A.
upvoted 4 times

 **yvinisiupacuando** 2 years, 5 months ago

Right answer is clearly B, "A" does not make any sense.
upvoted 11 times

 **cloudenthu01**  3 years, 3 months ago

Correct is B.
You only re-calculate expiry date when objects are re-written using re-write option to another storage class in which case creation date is rest.
But in this case objects is moveed to Coldline class after 90 days and then we want to delete the object after 365 days.
upvoted 40 times

 **T_T_M** 3 years, 1 month ago

You can change the storage class of an existing object either by rewriting the object or by using Object Lifecycle Management...Since Object Life cycle management was used there was no need to recalculate the expiration date and delete action still remains 365 days.

<https://cloud.google.com/storage/docs/storage-classes>
upvoted 14 times

 **YourCloudGuru**  2 weeks ago

Selected Answer: B

The correct answer is B.

Cloud Storage Object Lifecycle Management is a feature that allows you to automatically transition objects to different storage classes or delete them based on user-defined rules.

To set up a lifecycle management policy to move videos to Coldline after 90 days and then delete them after one year, you would create a rule with the following conditions and actions:

- * Condition: Age is greater than 90 days
- * Action: Set storage class to Coldline
- * Condition: Age is greater than 365 days
- * Action: Delete

This policy will ensure that your videos are automatically moved to Coldline after 90 days, where they will be stored at a lower cost. After one year, the videos will be automatically deleted.

Here is an example of how to create a lifecycle management policy using the gcloud command-line tool:

```
gcloud storage lifecycle management policies set my-bucket my-policy --action-set-storage-class coldline --condition-age-days 90 --action-delete --condition-age-days 365  
upvoted 2 times
```

 **CarlosMarin** 1 month ago

Selected Answer: A

"... and then deleted after one year FROM THEIR CREATION". I vote for A.

upvoted 1 times

 **ExamsFR** 2 months, 3 weeks ago

Selected Answer: B

Answer is B

upvoted 1 times

 **rosh199** 2 months, 3 weeks ago

Answer is B

upvoted 1 times

 **KerolesKhaliil** 4 months, 1 week ago

Answer is B .

Age condition :

The age condition is satisfied when a resource reaches the specified age (in days). Age is measured from the resource's creation time. For objects, the creation time is the time when the object is successfully written to the bucket, such as when an upload completes.

<https://cloud.google.com/storage/docs/lifecycle#age>

upvoted 1 times

 **Ashish_Tayal** 6 months, 1 week ago

Selected Answer: B

In GCP Storage, any rule either life cycle management or retention policy applies based on the creation date and time of the object. Even in the question they mention " from their creation". So ans is B.

upvoted 2 times

 **Jelly_Wang** 6 months, 2 weeks ago

The answer is B. When using Age, it is calculated from when the object is written to the bucket. The age do not get affected when you change the storage class .<https://cloud.google.com/storage/docs/lifecycle#age>

upvoted 3 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: B

NO need to recalculate as it is from object original creation date

upvoted 3 times

 **lights4life** 7 months, 2 weeks ago

Answer is B.

There is term called AGE, always GCP refers AGE by default which is since creation of the object.

upvoted 3 times

 **PB78** 7 months, 2 weeks ago

Answer is A as The age condition is satisfied when a resource reaches the specified age (in days). Age is measured from the resource's creation time.

For objects, the creation time is the time when the object is successfully written to the bucket, such as when an upload completes. The age of an object is unaffected by the object becoming a noncurrent version.

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B

To set up the policy to move videos stored in a specific Cloud Storage Regional bucket to Coldline after 90 days, and then delete them after one year from their creation, you can use Cloud Storage Object Lifecycle Management.

You can configure the lifecycle rule to move objects to Coldline storage class after 90 days by setting an Age condition of 90 days and a SetStorageClass action to Coldline. Then, you can configure the rule to delete objects after one year (365 days) from their creation by setting an Age condition of 365 days and a Delete action. Therefore, Answer B is the correct answer.

<https://cloud.google.com/storage/docs/lifecycle>

<https://cloud.google.com/storage/docs/lifecycle#age>

<https://cloud.google.com/storage/docs/lifecycle#actions>

upvoted 5 times

 **ElonGPT** 7 months ago

Thank you for your detailed explanation with desc and links

upvoted 1 times

 **prasadbishnu09** 7 months, 3 weeks ago

The Correct answer is A, because in question it is mentioned "A YEAR OF THE CREATION"

upvoted 1 times

 **Pksen** 8 months ago

Straightforward the answer is B, as we are using Life Cycle Management.

upvoted 1 times

 **Bobbybash** 8 months, 1 week ago

Selected Answer: A

The correct answer is A. Lets assume the object was created January 1st 2022, after 90 days (March 31st,2022), its moved to coldline storage, and later deleted one year after creation (31st December,2022), which is April 1st to December 31st (275 days). I hope this clears the misunderstanding

upvoted 3 times

 **lifelonglearning** 8 months, 2 weeks ago

Selected Answer: B

Correct answer is B. See the official docs: <https://cloud.google.com/storage/docs/lifecycle-configurations#:~:text=To%20move%20objects,of%20365>.

upvoted 1 times

You have a Linux VM that must connect to Cloud SQL. You created a service account with the appropriate access rights. You want to make sure that the VM uses this service account instead of the default Compute Engine service account. What should you do?

- A. When creating the VM via the web console, specify the service account under the 'Identity and API Access' section.
- B. Download a JSON Private Key for the service account. On the Project Metadata, add that JSON as the value for the key compute-engine-service-account.
- C. Download a JSON Private Key for the service account. On the Custom Metadata of the VM, add that JSON as the value for the key compute-engine-service-account.
- D. Download a JSON Private Key for the service account. After creating the VM, ssh into the VM and save the JSON under `~/.gcloud/compute-engine-service-account.json`.

Correct Answer: A

Reference:

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances>

Community vote distribution

A (56%)	C (39%)	5%
---------	---------	----

✉  **Agents89**  3 years, 5 months ago

A is correct
upvoted 48 times

✉  **ashrafh** 2 years, 1 month ago

I vote A
<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances>

Changing the service account and access scopes for an instance

If you want to run the VM as a different identity, or you determine that the instance needs a different set of scopes to call the required APIs, you can change the service account and the access scopes of an existing instance. For example, you can change access scopes to grant access to a new API, or change an instance so that it runs as a service account that you created, instead of the Compute Engine default service account. However, Google recommends that you use the fine-grained IAM policies instead of relying on access scopes to control resource access for the service account.

To change an instance's service account and access scopes, the instance must be temporarily stopped. To stop your instance, read the documentation for Stopping an instance. After changing the service account or access scopes, remember to restart the instance. Use one of the following methods to change service account or access scopes of the stopped instance.

Hope this helps :)
upvoted 16 times

✉  **boof** 2 years ago

A seems legit, the answer is worded poorly but is the most correct.

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#changeserviceaccountandscopes>

"To change an instance's service account and access scopes, the instance must be temporarily stopped ... After changing the service account or access scopes, remember to restart the instance." So we can stop the instance, change the service account, then start it up again.

upvoted 4 times

✉  **ready2rock** 2 years, 4 months ago

How can this be? It says you HAVE a VM, meaning it's already created. A cannot be the solution.
upvoted 10 times

✉  **jiniguez** 1 year, 10 months ago

As the comment says:

"To change an instance's service account and access scopes, the instance must be temporarily stopped ... After changing the service account or access scopes, remember to restart the instance." So we can stop the instance, change the service account, then start it up again.

upvoted 3 times

✉  **jabrrJ68w02ond1**  1 year, 10 months ago

Either the question or the answers are wrong. The question says that we HAVE a Linux VM, so we should strike all the answers that include "when creating the VM.." - on the other hand, adding JSON Tokens to VM metadata is terrible because it's readable in clear-text for everyone. So, what do we need to do here?

upvoted 10 times

✉  **YourCloudGuru**  2 weeks ago

Selected Answer: D

The correct answer is D.

This is the recommended approach, because it allows you to specify the service account that you want to use without having to modify the VM's metadata.

The other options are not as good:

Option A is not as good, because it requires you to specify the service account when creating the VM. This can be inconvenient if you need to update the service account later.

Option B is not as good, because it requires you to modify the VM's metadata. This can be complex and error-prone.

Option C is not as good, because it requires you to modify the VM's custom metadata. This is not a recommended approach, because custom metadata is intended for use by custom applications.

upvoted 1 times

 **vinodthakur49** 1 month, 2 weeks ago

Selected Answer: C

we have to use the newly created account rather VM default/attached SA.

upvoted 1 times

 **ExamsFR** 2 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 3 times

 **rosh199** 2 months, 3 weeks ago

A is correct

upvoted 2 times

 **geeroylenkins** 2 months, 3 weeks ago

Selected Answer: A

A is correct. No idea why you'd add anything to metadata of an instance <https://cloud.google.com/compute/docs/metadata/overview>

The SA can be specified in the web console during creation of the VM and also if the VM is stopped. This SA will then be used for everything that VM does. Therefore, A is correct.

upvoted 3 times

 **Neha_Pallavi** 2 months, 3 weeks ago

C is the possible correct answer. Already VM instance created.

upvoted 2 times

 **Paras_vohrA** 2 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 3 times

 **alany2000** 4 months, 1 week ago

Selected Answer: A

compute-engine-service-account is not a valid metadata key, therefore its A

upvoted 3 times

 **Kyle1776** 4 months, 2 weeks ago

Selected Answer: D

To ensure that a Linux VM uses a specific service account instead of the default Compute Engine service account when connecting to Cloud SQL, you should follow option D: Download a JSON Private Key for the service account. After creating the VM, ssh into the VM and save the JSON under `~/.gcloud/compute-engine-service-account.json`.

upvoted 2 times

 **Shweta2jun** 4 months, 4 weeks ago

Selected Answer: A

A is correct

upvoted 3 times

 **krishnu123** 5 months ago

Considering that VM is already there so A cant be the answer. And it can be any of B, C or D.

I think D is the right answer here because option D adding the private key to the VM by saving it under `~/.gcloud/compute-engine-service-account.json` by ssh into VM which looks more secure than B or C.

Not B because metadata service is public.

Not C because it can be error prone as metadata key must be specified correctly.

So I vote D.

upvoted 2 times

 **PREETHI26** 5 months, 2 weeks ago

I think c is correct

upvoted 1 times

 **LoveExams** 5 months, 3 weeks ago

Selected Answer: C

A is correct
upvoted 3 times

 **MahAli** 6 months ago

Selected Answer: C

the VM is already created you can use compute-engine-service-account key in the Instance metadata to configure the service account
upvoted 1 times

 **hanweiCN** 6 months, 1 week ago

A is right,
the authentication json web key is configurate in meta server - " service account / identity " key
<https://cloud.google.com/compute/docs/metadata/default-metadata-values>
upvoted 1 times

You created an instance of SQL Server 2017 on Compute Engine to test features in the new version. You want to connect to this instance using the fewest number of steps. What should you do?

- A. Install a RDP client on your desktop. Verify that a firewall rule for port 3389 exists.
- B. Install a RDP client in your desktop. Set a Windows username and password in the GCP Console. Use the credentials to log in to the instance.
- C. Set a Windows password in the GCP Console. Verify that a firewall rule for port 22 exists. Click the RDP button in the GCP Console and supply the credentials to log in.
- D. Set a Windows username and password in the GCP Console. Verify that a firewall rule for port 3389 exists. Click the RDP button in the GCP Console, and supply the credentials to log in.

Correct Answer: B

Reference:

<https://medium.com/falafel-software/sql-server-in-the-google-cloud-a17e8a1f11ce>

Community vote distribution

B (56%) D (44%)

vnxt **Highly Voted** 3 years, 5 months ago

I would say B is correct. RDP is enabled by default when you create a Windows instance (no need to check on it). Just make sure you install an RDP client (chrome ext or RDP) and set windows password.

upvoted 49 times

pas77 2 years, 3 months ago

Obviously, B is not the answer because you have to install an RDP client which is an extra step. D is the answer because you can connect directly using the RDP button in the GCP console.

upvoted 13 times

UtsavDM 2 years, 1 month ago

No, we can't connect using RDP directly in the GCP console. When we click on it, it asks us to install RDP client. So ultimately, B is more accurate.

upvoted 19 times

jabrrJ68w02ond1 1 year, 10 months ago

Tested it myself. At least on my Machines, I was asked to First install a RDP Client.

upvoted 17 times

ryumada 1 year, 2 months ago

The firewall rule for port 3389 is created by default if you create windows server on Compute Engine. So, no need to verify it.

<https://cloud.google.com/compute/docs/instances/connecting-to-windows#before-you-begin>

upvoted 5 times

Gregwaw 4 weeks, 1 day ago

You need firewall ingress rule for port 3389 in your VPC. It is not created by default.

upvoted 1 times

eledu1985 11 months, 3 weeks ago

In the link you provided the first step is to verify firewall rule was created, even if it is the default option!, so D is the most accurate even by the link.

upvoted 6 times

ankit89 **Highly Voted** 3 years, 5 months ago

D seems more correct

upvoted 33 times

obeythefist 1 year, 7 months ago

I tested this on on Compute Engine today by deploying a new instance. D is not correct. When you click the RDP button, you are asked to install a client or use the Windows RDP client if you are running Windows. There is no option to enter credentials or get an RDP session through the web interface.

upvoted 14 times

YourCloudGuru **Most Recent** 2 weeks ago

Selected Answer: D

The correct answer is D.

This option does not require you to install any additional software on your desktop, and it allows you to connect to the instance with just a few clicks.

The other options are not as good:

Option A requires you to install a RDP client on your desktop.

Option B requires you to install a RDP client on your desktop and set a Windows username and password in the GCP Console.

Option C requires you to set a Windows password in the GCP Console, verify that a firewall rule for port 22 exists, and then click the RDP button in the GCP Console and supply the credentials to log in.

upvoted 1 times

□ **Captain1212** 1 month, 1 week ago

Selected Answer: B

B seems more correct because you cannot use rdp by click you have to first install it

upvoted 1 times

□ **Abi143** 1 month, 2 weeks ago

Selected Answer: B

Bi scorrect

I just tested it By default port will be active

Just set password login with RDP as it will also be present in default on every pc

upvoted 1 times

□ **certified28** 3 months, 2 weeks ago

Selected Answer: B

you cannot connect using RDP from GCP console , you need to install a client first so B is the correct answer, verified it in sandbox

upvoted 2 times

□ **Trivris** 3 months ago

Answer D is correct. Even if you are asked to install rdp client, the easiest options is to navigate it from Google console rather finding it somewhere. Most of the cloud providers ask these questions for the sake of asking you to use their environment to do everything rather going through third party to install anything.

upvoted 1 times

□ **krishnu123** 5 months ago

D is the correct answer

upvoted 2 times

□ **Shenannigan** 5 months ago

Selected Answer: B

Hi all, for those that are selecting D please read or re-read this document: <https://cloud.google.com/compute/docs/instances/connecting-to-windows#remote-desktop-connection-app>

It clearly states that you need a RDP Client

You can search the document for "To connect to the remote desktop of a Windows instance, use one of the following procedures." and review the table of options

upvoted 1 times

□ **Praxii** 5 months, 2 weeks ago

Selected Answer: D

The correct answer is D. I tested this today.

I clicked on 'RDP' button and this is the message that came.

"Use a remote desktop protocol (RDP) client to connect to this instance. If you are running Windows on your local machine, use Remote Desktop Connection. Other operating systems might require you to use third-party software. The first time you connect, enter the username and password that you provided when you created the instance.

Note: You must configure the network firewall to open TCP port 3389 to enable RDP access"

So veryfing 3389 is important.

Next, those saying cant connect because third party connection message is coming. Just download the rdp file and it will connect using the app. So, either way you can connect using the button. Don't get confused.

upvoted 1 times

□ **Az900Exam2021** 1 month, 1 week ago

This is a tricky question.

B - "Install RDP on desktop" is redundant, since by default, it's included in windows

D - "Verify 3389 firewall rule" is redundant as well, since 3389/22 port are open by default

There is no perfect answer here.

upvoted 1 times

□ **ManvendraYadav** 5 months, 3 weeks ago

I would say B is correct. RDP is enabled by default when you crate a Windows instance (no need to chek on it). Just make sure you install an RDP client (chrome ext or RDP) and set windows password.

upvoted 2 times

□ **ManvendraYadav** 5 months, 3 weeks ago

I would say B is correct. RDP is enabled by default when you crate a Windows instance (no need to chek on it). Just make sure you install an RDP client (chrome ext or RDP) and set windows password.

upvoted 1 times

 **Ashish_Tayal** 6 months, 1 week ago

Selected Answer: D

As in production we always use custom VPC and in Custom VPC we have to enable RDP manually.

So Port 3389 must be allowed in the firewall rule, so need to check the firewall rule point as well.

upvoted 1 times

 **KC_go_reply** 6 months, 2 weeks ago

Selected Answer: B

I first thought it was D. Then, I tried it out in Compute Engine. After you create a Windows VM instance, port 3389 should be open. You then must set user name and password for RDP.

The actual point here is the RDP button. When you click it, you just get an explanation that you should use a RDP client to connect. So you can't directly connect to the instance with the RDP button. Therefore, D is wrong, and B is the correct answer.

upvoted 4 times

 **chikorita** 5 months, 3 weeks ago

buying this answer

upvoted 1 times

 **red_panda** 7 months, 1 week ago

Selected Answer: D

D is correct.

Why would anyone want to install specific software on their laptop, when we are taking a certification on cloud computing and you can achieve the same thing, in half the time?

upvoted 1 times

 **lights4life** 7 months, 2 weeks ago

Answer: D

upvoted 1 times

 **Buruguduystunstugudunsty** 7 months, 3 weeks ago

Selected Answer: D

Answer D is the correct answer. This method allows you to set a Windows username and password in the GCP Console, verify that the firewall rule for RDP port 3389 exists, and use the RDP button in the GCP Console to connect to the SQL Server instance with the provided credentials. This method is convenient because it does not require the installation of an RDP client on your desktop.

<https://cloud.google.com/compute/docs/instances/connecting-to-windows>

upvoted 4 times

 **Nazz1977** 8 months, 3 weeks ago

Selected Answer: D

I think it is D

upvoted 1 times

You have one GCP account running in your default region and zone and another account running in a non-default region and zone. You want to start a new Compute Engine instance in these two Google Cloud Platform accounts using the command line interface. What should you do?

- A. Create two configurations using gcloud config configurations create [NAME]. Run gcloud config configurations activate [NAME] to switch between accounts when running the commands to start the Compute Engine instances.
- B. Create two configurations using gcloud config configurations create [NAME]. Run gcloud configurations list to start the Compute Engine instances.
- C. Activate two configurations using gcloud configurations activate [NAME]. Run gcloud config list to start the Compute Engine instances.
- D. Activate two configurations using gcloud configurations activate [NAME]. Run gcloud configurations list to start the Compute Engine instances.

Correct Answer: A*Community vote distribution*

A (96%)

  **leba** Highly Voted 3 years, 5 months ago

Correct answer is A as you can create different configurations for each account and create compute instances in each account by activating the respective account. Refer GCP documentation - Configurations Create & Activate Options B, C & D are wrong as gcloud config configurations list does not help create instances. It would only lists existing named configurations.

upvoted 32 times

  **coldpar** Highly Voted 3 years, 6 months ago

A is the correct option

upvoted 28 times

  **YourCloudGuru** Most Recent 1 week, 5 days agoSelected Answer: A

The correct answer is A

This option allows you to create and activate different configurations for your GCP accounts. This way, you can easily switch between accounts and run commands without having to re-enter your credentials.

The other options are not as good:

- * Option B does not specify how to switch between accounts when running the commands to start the Compute Engine instances.
- * Option C does not specify how to create configurations for your GCP accounts.
- * Option D does not specify how to start the Compute Engine instances.

upvoted 1 times

  **Captain1212** 1 month, 1 week ago

Yes A seems to be more correct

upvoted 1 times

  **ExamsFR** 2 months, 3 weeks agoSelected Answer: A

A is correct

upvoted 1 times

  **Neha_Pallavi** 2 months, 3 weeks ago

Yes A is the correct Answer. using the gcloud configurations, we can create the we need to run the another compute engine instance upvoted 2 times

  **[Removed]** 4 months agoSelected Answer: A

A is the correct answer uncontested. The rest state 'list' which is inaccurate.

upvoted 1 times

  **Partha117** 6 months, 3 weeks agoSelected Answer: A

A is correct

upvoted 1 times

  **raselsys** 7 months, 1 week agoSelected Answer: A

Correct answer is A.

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A

Answer A is the correct answer.

To start a new Compute Engine instance in different GCP accounts using the command line interface, you can create two configurations using gcloud config configurations create [NAME], where [NAME] is the name of the configuration for each account. Then, you can switch between accounts using gcloud config configurations and activate [NAME] before running the gcloud compute instances create command to create the instance.

Answer B is incorrect because running gcloud configurations list only lists the available configurations and does not start Compute Engine instances.

Answer C is incorrect because activating configurations using gcloud configurations activate [NAME] only switches the current configuration, but does not start Compute Engine instances.

Answer D is incorrect because running gcloud configurations list only lists the available configurations and does not start Compute Engine instances.

upvoted 7 times

 **ChristN** 10 months, 1 week ago

Selected Answer: A

A is correct.

upvoted 1 times

 **PSS387** 11 months, 3 weeks ago

A is correct answer

upvoted 1 times

 **PKookNN** 12 months ago

Selected Answer: A

a is correct

upvoted 1 times

 **VarunGo** 1 year ago

Selected Answer: A

A is correct bro!

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Answer is A

upvoted 1 times

 **Alok17** 1 year, 3 months ago

A is corect

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right, similar question is part of Tutorials Dojo questions.

upvoted 1 times

You significantly changed a complex Deployment Manager template and want to confirm that the dependencies of all defined resources are properly met before committing it to the project. You want the most rapid feedback on your changes. What should you do?

- A. Use granular logging statements within a Deployment Manager template authored in Python.
- B. Monitor activity of the Deployment Manager execution on the Stackdriver Logging page of the GCP Console.
- C. Execute the Deployment Manager template against a separate project with the same configuration, and monitor for failures.
- D. Execute the Deployment Manager template using the --preview option in the same project, and observe the state of interdependent resources.

Correct Answer: D

Reference:

<https://cloud.google.com/deployment-manager/docs/deployments/updating-deployments>

Community vote distribution

D (100%)

 **YashBindlish**  3 years, 5 months ago

Correct answer is D as Deployment Manager provides the preview feature to check on what resources would be created
upvoted 43 times

 **Buruguduystunstugudunstuy**  7 months, 3 weeks ago

Selected Answer: D

Answer D is the most appropriate choice for getting rapid feedback on changes to a Deployment Manager template.

The preview command in Deployment Manager creates a preview deployment of the resources defined in the configuration, without actually creating or modifying any resources. This allows you to quickly test and validate changes to the template before committing them to the project. During the preview, you can observe the state of interdependent resources and ensure that their dependencies are properly met. This provides rapid feedback on your changes, without actually creating any resources or incurring any costs.

upvoted 11 times

 **YourCloudGuru**  1 week, 5 days ago

Selected Answer: D

The correct option is D

The `--preview` option will preview the changes that will be made to the deployment without actually making them. This allows you to see how the changes will affect the deployment and to identify any potential problems.

The other options are not as good:

* Option A is not as good, because it requires you to add logging statements to your Deployment Manager template. This can be time-consuming and error-prone.

* Option B is not as good, because it requires you to monitor the Stackdriver Logging page of the GCP Console. This can be difficult to do, especially for complex deployments.

* Option C is not as good, because it requires you to create a separate project with the same configuration. This can be time-consuming and expensive.

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

D is the right answer as deployment manager provides the preview option to check the templates

upvoted 1 times

 **ExamsFR** 2 months, 3 weeks ago

Selected Answer: D

D is right

upvoted 1 times

 **Neha_Pallavi** 2 months, 3 weeks ago

D is correct. Execute the Deployment Manager template using the --preview option in the same project, and observe the state of interdependent resources.

upvoted 1 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: D

D as it gives preview option to check

upvoted 1 times

 **ChristN** 10 months, 1 week ago

Selected Answer: D

Preview an updated configuration

You can preview the update you want to make before committing any changes, with the Google Cloud CLI or the API. The Deployment Manager service previews the configuration by expanding the full configuration and creating "shell" resources.

Deployment Manager does not instantiate any actual resources when you preview a configuration, giving you the opportunity to see the deployment before committing to it.

gcloud

API

With the Google Cloud CLI, make an update request with the --preview parameter:

```
gcloud deployment-manager deployments update example-deployment \
--config configuration-file.yaml \
--preview
    upvoted 3 times
```

✉ **Kopy** 11 months, 1 week ago

Selected Answer: D

deployment Manager provides the preview feature

upvoted 1 times

✉ **leogor** 11 months, 3 weeks ago

D Preview mode

upvoted 1 times

✉ **RanjithK** 1 year, 3 months ago

Answer is D

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

D is right.. Preview mode is best to verification

upvoted 2 times

✉ **jagopi9381** 1 year, 3 months ago

Selected Answer: D

D FOR SURE

upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago

Go for D

upvoted 1 times

✉ **tigerbaer** 1 year, 8 months ago

Selected Answer: D

Preview feature is available

upvoted 2 times

✉ **vishnukumartr** 1 year, 10 months ago

D. Execute the Deployment Manager template using the --preview option in the same project, and observe the state of interdependent resources.

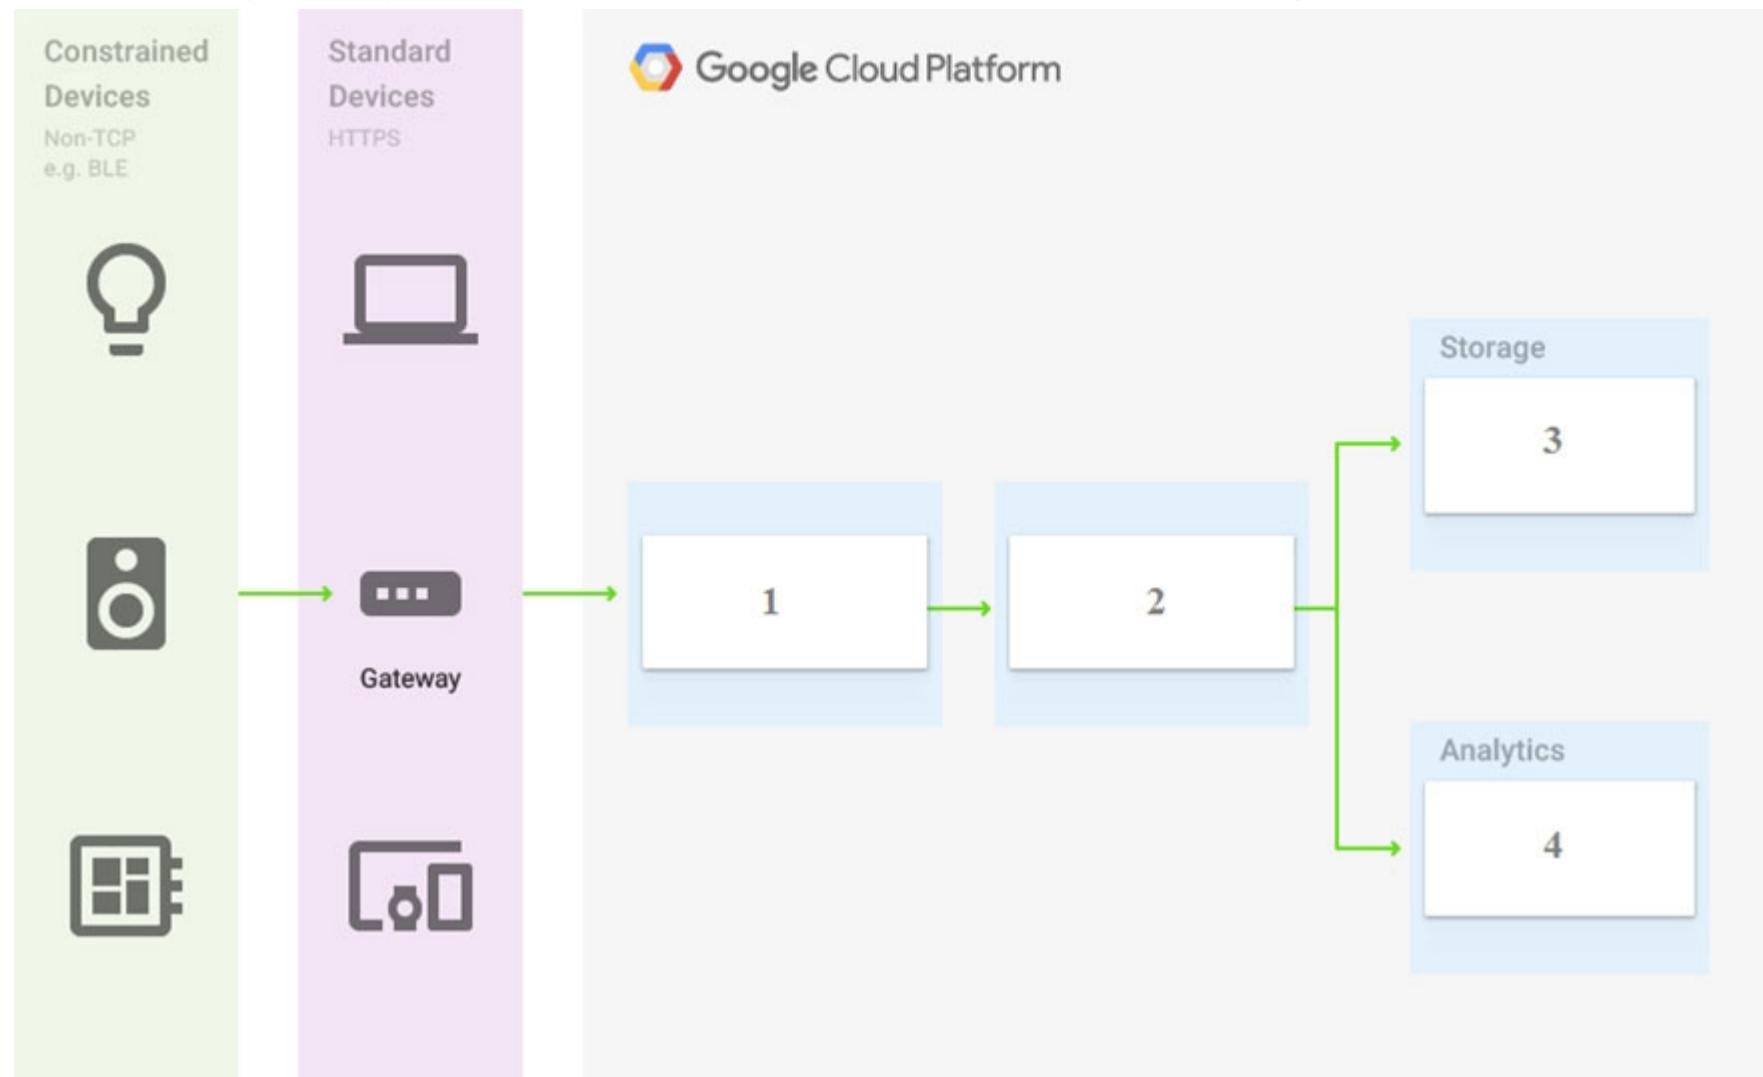
upvoted 1 times

✉ **Jaira1256** 1 year, 10 months ago

D is correct

upvoted 1 times

You are building a pipeline to process time-series data. Which Google Cloud Platform services should you put in boxes 1,2,3, and 4?



- A. Cloud Pub/Sub, Cloud Dataflow, Cloud Datastore, BigQuery
- B. Firebase Messages, Cloud Pub/Sub, Cloud Spanner, BigQuery
- C. Cloud Pub/Sub, Cloud Storage, BigQuery, Cloud Bigtable
- D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery

Correct Answer: D

Reference:

<https://cloud.google.com/solutions/correlating-time-series-dataflow>

Community vote distribution

D (100%)

✉ **cloudenthu01** Highly Voted 3 years, 3 months ago

Without a doubt D.

Whenever we want to process timeseries data look for BigTable.

Also you want to perform analytics in Box 4 ..look for BigQuery

Only D provides this option.

upvoted 59 times

✉ **vlodia** 3 years, 3 months ago

Speaker also looks like an IoT device so D not A

upvoted 2 times

✉ **adedj99** 2 years, 10 months ago

are we considering bigtable as storage in here , since they expecting some storage

upvoted 2 times

✉ **Buruguduystunstugudunstuy** Highly Voted 7 months, 3 weeks ago

Selected Answer: D

The correct process for building a pipeline to process time-series data. Here's how each of the components is used:

1. Cloud Pub/Sub: receives and distributes time-series data from different sources.
2. Cloud Dataflow: processes the data by applying transformations and analytics.
3. Cloud Bigtable: stores and manages the processed data as a NoSQL database.
4. BigQuery: provides a SQL-like interface to analyze the data and extract insights.

By combining these components, you can create a scalable and reliable pipeline to process and analyze time-series data in real time.
upvoted 25 times

✉ **kenrichy** 6 months, 3 weeks ago

Many thanks Buruguduy.. for the extensive explanation you always give to your choice of answers. It's really helpful to understand the concept.

upvoted 4 times

 **YourCloudGuru** Most Recent 1 week, 5 days ago

Selected Answer: D

The correct answer is D

This diagram shows a typical pipeline for processing time-series data:

1. **Cloud Pub/Sub:** A messaging service that allows you to send and receive messages between independent applications.
2. **Cloud Dataflow:** A fully-managed service for transforming and processing data streams.
3. **Cloud Bigtable:** A wide-column, distributed NoSQL database that is optimized for storing and analyzing large amounts of data.
4. **BigQuery:** A fully-managed, petabyte-scale analytics data warehouse that enables businesses to analyze all their data very quickly.

In this pipeline, the following happens:

1. Time-series data is sent to Cloud Pub/Sub.
2. Cloud Dataflow reads the data from Cloud Pub/Sub and performs any necessary transformations or processing.
3. Cloud Dataflow writes the transformed data to Cloud Bigtable.
4. BigQuery queries the data in Cloud Bigtable to generate insights.

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

big table for 3 , so d is the correct answer

upvoted 1 times

 **Neha_Pallavi** 2 months, 3 weeks ago

D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery

upvoted 1 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: D

Cloud Bigtable for time series data storage and BigQuery for analysis

upvoted 1 times

 **cslince** 10 months ago

Selected Answer: D

Option D

upvoted 1 times

 **dennydream** 11 months, 2 weeks ago

I know the answer is D, but it's misleading. "Storage" would suggest cloud storage, not BigTable.

upvoted 2 times

 **akshaydoifode88** 10 months, 4 weeks ago

You can't store timeseries data in cloud storage. Even if you store it how you are gonna access it?

upvoted 1 times

 **ale_brd_** 1 year, 1 month ago

Selected Answer: D

whenever IOT is present go with DataStore

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Answer is D

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

D is right

upvoted 1 times

 **Pleaseletmein** 1 year, 3 months ago

D is the right answer

upvoted 1 times

 **bamosk** 1 year, 3 months ago

Selected Answer: D

Look the picture posted on this page here

<https://cloud.google.com/dataflow#section-7>

upvoted 2 times

 **bamosk** 1 year, 3 months ago

IoT = Unstructured data -> eliminated Datastore + Timeseries requirement = BigTable

Ingestion point = Pub/Sub (Firebase messaging is a push notification service to client)

Process data + realtime or batch = Data flow

Analytics = BigQuery

Hence D.

upvoted 5 times

 **mplibunao** 1 year, 4 months ago

Selected Answer: D

Answer is D

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for D

upvoted 1 times

 **vishnukumartr** 1 year, 10 months ago

D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery

upvoted 1 times

You have a project for your App Engine application that serves a development environment. The required testing has succeeded and you want to create a new project to serve as your production environment. What should you do?

- A. Use gcloud to create the new project, and then deploy your application to the new project.
- B. Use gcloud to create the new project and to copy the deployed application to the new project.
- C. Create a Deployment Manager configuration file that copies the current App Engine deployment into a new project.
- D. Deploy your application again using gcloud and specify the project parameter with the new project name to create the new project.

Correct Answer: A
Community vote distribution

A (85%)	C (15%)
---------	---------

 **coldpar** Highly Voted 3 years, 6 months ago

Correct is A.

Option B is wrong as the option to use gcloud app cp does not exist.

Option C is wrong as Deployment Manager does not copy the application, but allows you to specify all the resources needed for your application in a declarative format using yaml

Option D is wrong as gcloud app deploy would not create a new project. The project should be created before usage

upvoted 68 times

 **manu2020** 3 years, 3 months ago

you're missing one thing. D isn't about using deployment manager to copy the configuration, instead, using the configuration file to copy the configuration from test project.

upvoted 6 times

 **AdelElagawany** 2 weeks, 4 days ago

A is correct since the documentation here explicitly mentioned the roles of external editors https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors

upvoted 1 times

 **leba** Highly Voted 3 years, 5 months ago

Correct answer is A as gcloud can be used to create a new project and the gcloud app deploy can point to the new project. Refer GCP documentation - GCloud App Deploy.

Option B is wrong as the option to use gcloud app cp does not exist

.Option C is wrong as Deployment Manager does not copy the application, but allows you to specify all the resources needed for your application in a declarative format using yaml

Option D is wrong as gcloud app deploy would not create a new project. The project should be created before usage.

upvoted 14 times

 **YourCloudGuru** Most Recent 1 week, 5 days ago

Selected Answer: A

The correct answer is A.

This is the simplest and most straightforward way to create a new production environment. It is also the most efficient way, because it does not require you to copy the deployed application to the new project.

The other options are not as good:

* Option B is not as good, because it requires you to copy the deployed application to the new project. This can be time-consuming and error-prone.

* Option C is not as good, because it requires you to create a Deployment Manager configuration file. This can be complex and time-consuming.

* Option D is not as good, because it does not allow you to create a new project.

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: A

option A seems more correct as in B copy command don't exist

upvoted 1 times

 **Bb_master** 5 months, 3 weeks ago

Selected Answer: C

As this is test on cloud infra then definitely the question is about creating instance with all the configuration which was present in the development environment. I think C is correct because we can copy the existing configuration from deployment manager to the new project.

upvoted 2 times

 **Buruguduystunstugudunsty** 7 months, 3 weeks ago

Selected Answer: A

The correct answer is A. Use gcloud to create the new project, and then deploy your application to the new project.

Answer B, copying the deployed application to the new project, is not a recommended method for creating a new project since it can lead to inconsistencies and compatibility issues between the old and new projects.

Answer C, using a Deployment Manager configuration file to copy the current App Engine deployment to a new project, is also not recommended. Deployment Manager is a service that automates the deployment of infrastructure and applications, but it's not intended to copy a deployed application to a new project.

Answer D, deploying your application again using gcloud and specifying the project parameter with the new project name, is not recommended. This method will create a new deployment of your application in the same project as the development environment, rather than in a separate project for the production environment.

upvoted 7 times

 **leogor** 11 months, 3 weeks ago

A, one project can only have one deployment

upvoted 2 times

 **ale_brd_** 1 year, 1 month ago

Selected Answer: A

cloud can be used to create a new project and the gcloud app deploy can point to the new project.

upvoted 1 times

 **gcp_world123** 1 year, 1 month ago

A is correct

You can deploy to a different project by using -project flag.

By default, the service is deployed the current project configured via:

\$ gcloud config set core/project PROJECT

To override this value for a single deployment, use the -project flag:

\$ gcloud app deploy ~/my_app/app.yaml --project=PROJECT

Ref: <https://cloud.google.com/sdk/gcloud/reference/app/deploy>

upvoted 2 times

 **AzureDP900** 1 year, 3 months ago

A is right

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

go for A

upvoted 1 times

 **[Removed]** 1 year, 9 months ago

Selected Answer: A

Think A also.

upvoted 1 times

 **onyb** 1 year, 10 months ago

Although the answer is not worded correctly, I think C is the right answer . Command is
gcloud app deploy ~/my_app/app.yaml --project=PROJECT (<https://cloud.google.com/sdk/gcloud/reference/app/deploy>)
upvoted 3 times

 **shawnkkk** 1 year, 10 months ago

A. Use gcloud to create the new project, and then deploy your application to the new project.

upvoted 1 times

 **vishnukumartr** 1 year, 10 months ago

A. Use gcloud to create the new project, and then deploy your application to the new project.

upvoted 1 times

 **Jaira1256** 1 year, 10 months ago

Correct is A

upvoted 1 times

 **vmart** 2 years, 4 months ago

A is correct answer

upvoted 1 times

You need to configure IAM access audit logging in BigQuery for external auditors. You want to follow Google-recommended practices. What should you do?

- A. Add the auditors group to the 'logging.viewer' and 'bigQuery.dataViewer' predefined IAM roles.
- B. Add the auditors group to two new custom IAM roles.
- C. Add the auditor user accounts to the 'logging.viewer' and 'bigQuery.dataViewer' predefined IAM roles.
- D. Add the auditor user accounts to two new custom IAM roles.

Correct Answer: C

Reference:

<https://cloud.google.com/iam/docs/roles-audit-logging>

Community vote distribution

A (78%) 13% 9%

✉ **coldpar** Highly Voted 3 years, 6 months ago

Correct is A.

As per google best practices it is recommended to use predefined roles and create groups to control access to multiple users with same responsibility

upvoted 73 times

✉ **droogie** 3 years, 3 months ago

You assume Auditors Group = External Auditors only. Auditors Group may contain both Internal and External Auditors.

upvoted 4 times

✉ **robor97** 2 years, 10 months ago

The question literally says - External Auditors

upvoted 12 times

✉ **adeice** 2 years, 6 months ago

I can create External group and Internal group Auditors

upvoted 2 times

✉ **JavierCorrea** Highly Voted 3 years, 2 months ago

Correct answer is A as per:

https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors

upvoted 43 times

✉ **YourCloudGuru** Most Recent 1 week, 5 days ago

Selected Answer: A

The correct answer is A.

This option follows Google-recommended practices, because it allows you to grant auditors access to view audit logs without granting them access to other resources in your project.

The other options are not as good:

- * Option B is not as good, because it requires you to create two new custom IAM roles. This can be complex and time-consuming.
- * Option C is not as good, because it grants auditors access to all audit logs in your project, including audit logs for resources that they do not need access to.
- * Option D is not as good, because it grants auditors access to all data in your BigQuery datasets, including data that they do not need access to.

upvoted 2 times

✉ **Captain1212** 1 month, 1 week ago

Selected Answer: A

Google Recommended Practice A is the correct Answer add the users in the group then grant them the access

upvoted 2 times

✉ **sthapit** 2 months ago

C

Option A, which suggests adding the auditors group to predefined roles, might not be as appropriate as using individual auditor user accounts. It's generally a best practice to assign permissions to specific users rather than groups, as it provides better granularity and control over access.

upvoted 1 times

✉ **ExamsFR** 2 months, 3 weeks ago

Selected Answer: A

Correct answer is A

upvoted 2 times

✉ **vinodthakur49** 2 months, 3 weeks ago

Selected Answer: C

There is no group created already, so C is the right answer.

upvoted 1 times

✉ **smanoj85** 6 months, 3 weeks ago

Correct Answer is B

By creating a custom IAM role, you can specify the exact permissions that the auditors need, and avoid granting them unnecessary permissions that come with predefined IAM roles. In this case, you can create two custom IAM roles: one for 'logging.viewer' and one for 'bigQuery.dataViewer', and grant the corresponding permissions to each role. Then, you can add the auditors group to these custom roles to give them access to the required logs and data.

upvoted 4 times

✉ **smanoj85** 6 months, 3 weeks ago

Correct Answer is B

Option A is incorrect because the logging.viewer and bigQuery.dataViewer roles only grant read access to logs and data in BigQuery, respectively. These roles do not provide audit logging capabilities.

Option C is incorrect because it suggests adding individual user accounts to the roles, whereas the question specifically asks for adding an auditors group. In addition, adding individual user accounts can be difficult to manage and does not scale well as the number of auditors increases. It is generally recommended to use groups for managing access whenever possible.

Option D suggests adding the auditor user accounts to two new custom IAM roles, which could work. However, the question specifically asks for following Google-recommended practices. The recommended practice is to use predefined roles over custom roles whenever possible. Therefore, option B, which suggests adding the auditors group to two new custom IAM roles, is not recommended.

upvoted 1 times

✉ **red_panda** 6 months, 3 weeks ago

Selected Answer: B

B is correct. Is a best practice to prefer custom role, specially for external users.

upvoted 1 times

✉ **asallo** 7 months ago

A is the most appropriate Answer

upvoted 1 times

✉ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B

I would say that Answer A is not the correct answer. While it is true that adding the auditor's group to the 'logging.viewer' and 'bigQuery.dataViewer' roles would allow them to view the logs and data in BigQuery, it does not enable IAM access audit logging.

The correct answer is Answer B - Add the auditors group to two new custom IAM roles. You should create custom IAM roles with the necessary permissions to view IAM audit logs in BigQuery and assign those roles to the auditor's group. This follows the Google-recommended practice of using custom roles to grant least privilege access to resources.

Answer C is incorrect because you should not add users' accounts to predefined IAM roles like logging.viewer or bigQuery.dataViewer. Predefined roles are meant to provide a general set of permissions for common use cases, and adding users or groups to them may grant them unnecessary access.

Answer D is not the best practice as it is better to create separate custom IAM roles for each type of user rather than combining them.

upvoted 2 times

✉ **Bobbybash** 7 months, 4 weeks ago

Selected Answer: B

B.... The recommended practice for configuring IAM access audit logging in BigQuery is to create two custom IAM roles for auditors: one with the bigquery.datasets.get permission, and the other with the bigquery.tables.getData permission. You should then add the auditors group to these custom IAM roles. This will allow auditors to view metadata about datasets and access data within tables, while preventing them from performing other operations on the BigQuery resources. Therefore, option B is the correct answer.

upvoted 3 times

✉ **processor** 9 months, 2 weeks ago

Selected Answer: A

Once again, the "correct" answer is wrong. (Regarding google best practices). How could you hope someone gives money to get wrong answers. A is the good one.

upvoted 1 times

✉ **jrlsl1991** 8 months, 2 weeks ago

Honestly, I think some answers are "wrong" on purpose because if they all were right, the website could be taken as a "cheat" instead of help to study.

Correct answer is A.

upvoted 1 times

✉ **glanshima** 10 months ago

Correct A

The organization creates a Google group for these external auditors and adds the current auditor to the group. This group is monitored and is typically granted access to the dashboard application.

see: https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors

upvoted 2 times

 **vijay456** 11 months, 2 weeks ago

Selected Answer: A

<https://cloud.google.com/iam/docs/job-functions/auditing>

upvoted 2 times

 **vijay456** 11 months, 2 weeks ago

Selected Answer: C

there is no group created and no option says create group too, so provided option C is suitable best answer

upvoted 3 times

 **leogor** 11 months, 3 weeks ago

A, auditors group

upvoted 1 times

You need to set up permissions for a set of Compute Engine instances to enable them to write data into a particular Cloud Storage bucket. You want to follow Google-recommended practices. What should you do?

- A. Create a service account with an access scope. Use the access scope 'https://www.googleapis.com/auth/devstorage.write_only'.
- B. Create a service account with an access scope. Use the access scope '<https://www.googleapis.com/auth/cloud-platform>'.
- C. Create a service account and add it to the IAM role 'storage.objectCreator' for that bucket.
- D. Create a service account and add it to the IAM role 'storage.objectAdmin' for that bucket.

Correct Answer: D

Community vote distribution

C (97%)

✉  **coldpar**  3 years, 6 months ago

As per as the least privilege recommended by google, C is the correct Option, A is incorrect because the scope doesn't exist. B incorrect because it will give him full of control

upvoted 50 times

✉  **robor97** 2 years, 10 months ago

The scope does exist - <https://download.huihoo.com/google/gdgdevkit/DVD1/developers.google.com/compute/docs/api/how-tos/authorization.html>

upvoted 2 times

✉  **peter77** 2 years ago

No it doesn't. You have read-only, read-write, full-control and others... but "write-only" is not a thing.

<https://cloud.google.com/storage/docs/authentication>

upvoted 4 times

✉  **gielda211** 1 year, 6 months ago

it doesn't exist. show us this on official google website

upvoted 2 times

✉  **johnconnor** 1 year, 3 months ago

Check here, it is A-> <https://cloud.google.com/storage/docs/authentication>
<https://cloud.google.com/storage/docs/authentication>

upvoted 1 times

✉  **karim1321** 3 months, 3 weeks ago

In the Document, 'write -only' does not exist. Just read-only

upvoted 2 times

✉  **CVGCP** 4 months, 2 weeks ago

There is no scope called write-only, as per the reference document.

upvoted 1 times

✉  **Bedmed** 9 months, 2 weeks ago

In the Document, it includes https://www.googleapis.com/auth/devstorage.read_write scope

upvoted 1 times

✉  **XRiddlerX**  3 years, 3 months ago

In reviewing this, it looks to be a multiple answer question. According to Best Practices in this Google Doc (https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#best_practices) you grant the instance the scope and the permissions are determined by the IAM roles of the service account. In this case, you would grant the instance the scope and the role (storage.objectCreator) to the service account.

Ans B and C

Role from GCP Console:

ID = roles/storage.objectCreator

Role launch stage = General Availability

Description = Access to create objects in GCS.

3 assigned permissions

resourcemanager.projects.get

resourcemanager.projects.list

storage.objects.create

upvoted 17 times

✉  **ryumada** 1 year, 2 months ago

Reading the second point of the best practice. You should grant your VM the https://www.googleapis.com/auth/cloud-platform scope to allow access to most of Google Cloud APIs.

So, that the IAM permissions are completely determined by the IAM roles you granted to the service account.

The conclusion is you should not mess up with the VM scopes to grant access to Google Services, instead you should grant the access via IAM roles of the service account you attached to the VM.

https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#best_practices

upvoted 2 times

 **nickyshil** 1 year, 2 months ago

There are many access scopes available to choose from, but a best practice is to set the cloud-platform access scope, which is an OAuth scope for most Google Cloud services, and then control the service account's access by granting it IAM roles..you have an app that reads and writes files on Cloud Storage, it must first authenticate to the Cloud Storage API. You can create an instance with the cloud-platform scope and attach a service account to the instance

<https://cloud.google.com/compute/docs/access/service-accounts>

upvoted 1 times

 **YourCloudGuru** Most Recent 1 week, 4 days ago

Selected Answer: C

The correct answer is C.

The other options are not accurate and go against the principle of giving least required access.

A is incorrect as there is no role as write_only

B is not a good option as it gives full control of google cloud services where as we are looking for write data into a particular cloud storage bucket

D. is not a good option as it gives full control over objects

Sources:

<https://cloud.google.com/storage/docs/authentication>

<https://cloud.google.com/storage/docs/access-control/iam-roles>

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: C

c seems more correct, you need to go iam to provide the permissions , b and d will give it more or full access

upvoted 1 times

 **Neha_Pallavi** 2 months, 3 weeks ago

Associate Cloud Engineer exam booked very soon. kindly share the all the questions and any other support exam to clear this
upvoted 2 times

 **Shubha1** 2 months, 1 week ago

Hi Neha, Please let me know how your exam was? I am taking the exam soon. Thanks

upvoted 1 times

 **ExamsFR** 2 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

 **rosh199** 2 months, 3 weeks ago

C is correct

upvoted 1 times

 **CVGCP** 4 months, 2 weeks ago

Selected Answer: C

C is correct answer

upvoted 3 times

 **trainingexam** 4 months, 2 weeks ago

Selected Answer: C

The ask is how the "Compute Engine instances to enable them to write data into a particular Cloud Storage bucket". A service account is a special kind of account used by an application or compute workload, rather than a person. When you set up an instance to run as a service account, you determine the level of access the service account has by the IAM roles that you grant to the service account. If the service account has no IAM roles, then no resources can be accessed using the service account on that instance.

The best Practice suggested by Google is refer in this link: https://cloud.google.com/compute/docs/access/service-accounts#scopes_best_practice <https://cloud.google.com/storage/docs/access-control/iam-roles> shows that storage.objectCreator is best choice of the role for this problem statement.

upvoted 1 times

 **Praxii** 5 months, 2 weeks ago

Selected Answer: C

The correct answer is C.

There is no role as write only its read only hence A is incorrect.

upvoted 1 times

 **Ashish_Tayal** 6 months, 1 week ago

Selected Answer: C

IAM Work on Principal of least privilege,
upvoted 1 times

 **smanoj85** 6 months, 3 weeks ago

Option C is the correct answer. To grant a set of Compute Engine instances permissions to write data to a particular Cloud Storage bucket, you should create a service account and add it to the IAM role 'storage.objectCreator' for that bucket. This IAM role allows the service account to create new objects in the bucket, but it does not allow it to modify or delete existing objects. Option A is incorrect because the access scope 'https://www.googleapis.com/auth/devstorage.write_only' does not exist. Option B is incorrect because the access scope 'https://www.googleapis.com/auth/cloud-platform' grants permissions for all Google Cloud Platform services, which is overly broad and not recommended. Option D is incorrect because the IAM role 'storage.objectAdmin' provides full control over the bucket, which is more access than necessary to allow the Compute Engine instances to write data to the bucket.

upvoted 2 times

 **red_panda** 6 months, 3 weeks ago

Selected Answer: C

According to least privileges, the correct answer is C
upvoted 1 times

 **ch2023** 6 months, 4 weeks ago

Selected Answer: A

should be A
https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors
upvoted 1 times

 **ch2023** 6 months, 1 week ago

C is the correct answer
upvoted 1 times

 **matiascoca** 7 months ago

Selected Answer: C

C is the correct answer
upvoted 2 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: C

To allow a Compute Engine instance to write to a Cloud Storage bucket, you can create a service account and add it to the IAM role 'storage.objectCreator' for that bucket.

This will grant the service account permission to create objects in the bucket. Therefore, Answer C is correct.
upvoted 4 times

 **nanhukumari** 9 months, 3 weeks ago

- A. There is no scope as devstorage.write_only'
 - B. cloud-platform gives full access
 - C. correct
 - D. gives full access
- upvoted 1 times

You have sensitive data stored in three Cloud Storage buckets and have enabled data access logging. You want to verify activities for a particular user for these buckets, using the fewest possible steps. You need to verify the addition of metadata labels and which files have been viewed from those buckets. What should you do?

- A. Using the GCP Console, filter the Activity log to view the information.
- B. Using the GCP Console, filter the Stackdriver log to view the information.
- C. View the bucket in the Storage section of the GCP Console.
- D. Create a trace in Stackdriver to view the information.

Correct Answer: A

Community vote distribution

B (50%)

A (50%)

 **iamgcp**  3 years, 4 months ago

A is correct. As mentioned in the question, data access logging is enabled. I tried to download a file from a bucket and was able to view this information in Activity tab in console

upvoted 39 times

 **vito9630** 3 years, 4 months ago

data access logging don't provide information about addition of metadata, so B is correct

upvoted 21 times

 **MEHDIGRB** 1 year ago

activity log is deprecated:

<https://cloud.google.com/compute/docs/logging/activity-logs>

upvoted 2 times

 **Rog_444** 7 months, 2 weeks ago

Yes, it is deprecated. However, it became the audit log which is exactly what this question is referring to. Option A is correct in my opinion.

upvoted 2 times

 **barathgdkrish** 10 months ago

You need to see here, <https://cloud.google.com/compute/docs/logging/audit-logging>. Admin activity audit logs.

upvoted 1 times

 **RegisFTM** 1 year, 9 months ago

I did all the configuration enabling data access logging but I still not able to see the logs when uploading or downloading a file. Does someone here has done it with a different result?

upvoted 1 times

 **ryumada** 1 year, 2 months ago

I agree with liyux21 and vito9630. In this reference link below says:

In the Activity page, where the identity performing logged actions is redacted from the audit log entry, User (anonymized) is displayed.

Because of this, I think you can't verify the addition of metadata labels through Activity Logs.

<https://cloud.google.com/logging/docs/audit#view-activity>

upvoted 1 times

 **eliteone11**  2 years, 10 months ago

Answer is A. Activity log does indeed show information about metadata.

I agree with Eshkrkrkr based on <https://cloud.google.com/storage/docs/audit-logs> Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.

upvoted 14 times

 **injarapu** 1 year, 7 months ago

'Admin activity logs' capture metadata modification, but its different from 'Data Access logging', right ?

upvoted 2 times

 **YourCloudGuru**  1 week, 4 days ago

Selected Answer: A

The correct answer is A.

Steps:

- 1 Go to the GCP Console
- 2 Click on the Hamburger menu in the top left corner of the page

- 3 Click on Logging
- 4 Click on the Activity log tab
- 5 Click on the Filter button
- 6 In the Resource field, enter the names of the three Cloud Storage buckets
- 7 In the User field, enter the name of the user whose activities you want to verify
- 8 Click on the Apply button

The Activity log will display all of the activities that have been performed on the three Cloud Storage buckets by the specified user. You can then review the log to identify the addition of metadata labels and which files have been viewed.

The other options are not as efficient

B because it requires you to create a trace in Stackdriver. This can be time-consuming and error-prone
C&D because it does not allow you to filter the results by user or resource

upvoted 2 times

 **HemrajLodhi** 1 week, 4 days ago

A is correct, the activity log has metadata info as well

upvoted 1 times

 **YomanB** 1 month, 1 week ago

The best approach among the given options is to use Option B: Using the GCP Console, filter the Stackdriver log to view the information. This approach would allow you to efficiently filter and analyze the logs related to the Cloud Storage buckets to verify the desired user activities.

upvoted 1 times

 **axantroff** 1 month, 2 weeks ago

I think it's all about Data Access audit logs, you check it - <https://cloud.google.com/storage/docs/audit-logging>

upvoted 1 times

 **Nxt_007** 2 months, 1 week ago

Selected Answer: B

Option B is Correct

By filtering the Stackdriver logs in the GCP Console, you can easily track and verify activities related to the addition of metadata labels and viewed files for the specified user across the three Cloud Storage buckets.

Option A (filtering the Activity log) is not as appropriate because the Stackdriver log provides more detailed and comprehensive logging and monitoring capabilities.

upvoted 1 times

 **Neha_Pallavi** 2 months, 2 weeks ago

Option A suggests filtering the Activity log, but the Activity log only shows user activity within a GCP project, not across projects or resources. In this case, the buckets are likely in different projects or locations, so the Activity log may not show all the relevant activity. So B Is correct

upvoted 1 times

 **Neha_Pallavi** 2 months, 2 weeks ago

Using the GCP Console, Filter the activity log to view the information

upvoted 1 times

 **ExamsFR** 2 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 2 times

 **SHERLOCK2213** 4 months, 3 weeks ago

Selected Answer: A

A is the answer

upvoted 3 times

 **Vamshi_Krishna** 5 months, 1 week ago

B is Correct. Using the GCP Console, filter the Stackdriver log to view the information.

Stackdriver Logging provides an audit trail of the user activity for Cloud Storage. It is a centralized service to search, analyze, and alert on logs generated from GCP services. You can filter the logs to get the details for a particular user and search for specific activity such as metadata updates or file views.

Option A suggests filtering the Activity log, but the Activity log only shows user activity within a GCP project, not across projects or resources. In this case, the buckets are likely in different projects or locations, so the Activity log may not show all the relevant activity.

upvoted 3 times

 **VEEDOL** 5 months, 1 week ago

Both A and B are correct, but using stackdriver it will be faster.

upvoted 1 times

 **tracyonair** 5 months, 4 weeks ago

Selected Answer: B

Filtering Audit Logs

Stackdriver provides both basic and advanced logs filters. Basic log filters allows you to filter the results displayed in the feed by user, resource type and date/time.

entries:

from specific logs or log services
within a given time range
that satisfy conditions on metadata or user-defined fields
that represent a sampling percentage of all log entries
upvoted 1 times

 **hanweiCN** 6 months ago

Selected Answer: B
You can see a high-level overview of all your audit logs on the Cloud Console Activity page. By default, this feed does not display data access logs. To enable them, go to the Filter configuration panel and select the "Data Access" field under Categories.
upvoted 1 times

 **ACEqa** 6 months, 1 week ago

b IS CORRECT
upvoted 1 times

 **ACEqa** 6 months, 1 week ago

Stackdriver Logging provides access and security logs for Cloud Storage in one centralized location, including audit logs for data access and changes. It allows you to filter logs by user identity, activity type, and bucket name to quickly find relevant logs.
upvoted 1 times

You are the project owner of a GCP project and want to delegate control to colleagues to manage buckets and files in Cloud Storage. You want to follow Google-recommended practices. Which IAM roles should you grant your colleagues?

- A. Project Editor
- B. Storage Admin
- C. Storage Object Admin
- D. Storage Object Creator

Correct Answer: B

Community vote distribution

B (100%)

 **ESP_SAP**  3 years, 1 month ago

Correct Answer is (B):

Storage Admin (roles/storage.admin) Grants full control of buckets and objects.

When applied to an individual bucket, control applies only to the specified bucket and objects within the bucket.

```
firebase.projects.get
resourcemanager.projects.get
resourcemanager.projects.list
storage.buckets.*
storage.objects.*
```

upvoted 46 times

 **iambatmanadarkknight** 1 year, 12 months ago

why not storage object admin?

upvoted 3 times

 **TenshiD** 1 year, 11 months ago

Because the object admin don't have control over buckets and you need it

upvoted 17 times

 **Raz0r** 1 year, 9 months ago

Exactly, you want to give someone right to edit storages not just objects. Google does this kind of answers to confuse us.

upvoted 3 times

 **dang1986** 1 year, 7 months ago

Question states "Buckets and Objects"

upvoted 2 times

 **Agents89**  3 years, 5 months ago

B is correct

upvoted 14 times

 **YourCloudGuru**  1 week, 4 days ago

Selected Answer: B

The correct answer is B

This role allows users to create, manage, and delete buckets and files in Cloud Storage. It also allows users to set permissions on buckets and files.

The other options are not as good:

A gives users too much power, as it allows them to manage all resources in a project, including Cloud Storage buckets and files

C gives users too much power, as it allows them to manage all objects in a bucket, including the permissions on those objects

D does not give users enough power, as it does not allow them to manage buckets or set permissions on buckets and objects

Steps to grant Storage Admin IAM role:

1 Go to the Google Cloud Console

2 Click on the IAM & Admin menu

3 Click on the Roles tab

4 Click on the Storage Admin role

5 Click on the Add members button

6 Type the email addresses of your colleagues in the Members field

7 Click on the Add button

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

B is more correct, as it give you the both access

upvoted 1 times

 **chem1** 1 month, 1 week ago

Selected Answer: B

B is correct

upvoted 1 times

 **Neha_Pallavi** 2 months, 2 weeks ago

B.

Storage Admin (roles/storage.admin) - Grants full control of buckets and objects.

<https://cloud.google.com/storage/docs/access-control/iam-roles>

upvoted 1 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: B

Correct option B

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B

Answer B, "Storage Admin," is the correct answer because it grants permissions to manage Cloud Storage resources at the project level, including creating and deleting buckets, changing bucket settings, and assigning permissions to buckets and their contents. This role also includes the permissions of the "Storage Object Admin" and "Storage Object Creator" roles, which allow managing objects and uploading new ones.

Answer A, "Project Editor," is a higher-level role that includes permissions to manage not only Cloud Storage but also other GCP services in the project. Granting this role may not be appropriate if the colleagues only need to manage Cloud Storage resources.

Answers C and D may not be sufficient if the colleagues need to create or delete buckets or change their settings.

upvoted 4 times

 **kkozlow2** 10 months, 1 week ago

Selected Answer: B

Storage Admin (roles/storage.admin)

Grants full control of buckets and objects.

When applied to an individual bucket, control applies only to the specified bucket and objects within the bucket.

While

Storage Object Admin (roles/storage.objectAdmin)

Grants full control over objects, including listing, creating, viewing, and deleting objects.

upvoted 2 times

 **leogor** 11 months, 3 weeks ago

B. Storage Admin

upvoted 1 times

 **Jesse20323645** 11 months, 3 weeks ago

Selected Answer: B

According to the question, your colleagues need to manage "buckets" in Cloud Storage(storage.objects.* permission), so (B) is correct. (C) doesn't have control over the buckets.

Cloud document:

<https://cloud.google.com/storage/docs/access-control/iam-roles#standard-roles>

-->Storage Object Admin (Cannot find "storage.buckets.*" Permission)

---->Storage Admin (Has "storage.buckets.*" Permission)

upvoted 1 times

 **leogor** 1 year ago

Selected Answer: B

the ans is definitively B

upvoted 1 times

 **sandipk91** 1 year, 1 month ago

Selected Answer: B

Storage Admin option A - As we are supposed to create buckets as well

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Answer is B

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

Storage Admin is right.. B is correct

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for B

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

<https://cloud.google.com/storage/docs/access-control/iam-roles>

upvoted 1 times

 **gielda211** 1 year, 6 months ago

Selected Answer: B

<https://cloud.google.com/storage/docs/access-control/iam-roles>

upvoted 3 times

You have an object in a Cloud Storage bucket that you want to share with an external company. The object contains sensitive data. You want access to the content to be removed after four hours. The external company does not have a Google account to which you can grant specific user-based access privileges. You want to use the most secure method that requires the fewest steps. What should you do?

- A. Create a signed URL with a four-hour expiration and share the URL with the company.
- B. Set object access to 'public' and use object lifecycle management to remove the object after four hours.
- C. Configure the storage bucket as a static website and furnish the object's URL to the company. Delete the object from the storage bucket after four hours.
- D. Create a new Cloud Storage bucket specifically for the external company to access. Copy the object to that bucket. Delete the bucket after four hours have passed.

Correct Answer: A*Community vote distribution*

A (100%)

JJ_ME Highly Voted 3 years ago

A.
Signed URLs are used to give time-limited resource access to anyone in possession of the URL, regardless of whether they have a Google account.
<https://cloud.google.com/storage/docs/access-control/signed-urls>
upvoted 37 times

Agents89 Highly Voted 3 years, 5 months ago

A is correct
upvoted 9 times

YourCloudGuru Most Recent 1 week, 4 days ago

Selected Answer: A
The correct answer is A

To do this, follow these steps:

- 1 Go to the Google Cloud Console
- 2 Click on the Storage menu
- 3 Click on the Browser tab
- 4 Click on the name of the bucket containing the object that you want to share
- 5 Click on the name of the object that you want to share
- 6 Click on the Create signed URL button
- 7 In the Expires field, enter 4
- 8 Click on the Create button

The company will be able to access the object using the signed URL for four hours. After four hours, the signed URL will expire and the company will no longer be able to access the object.

The other options are not as secure or efficient:

- B not as secure because it makes the object accessible to anyone who has the URL of the object
- C requires you to configure the storage bucket as a static website and to delete the object after four hours
- D requires you to create a new Cloud Storage bucket and to copy the object to that bucket

<https://cloud.google.com/storage/docs/access-control/signed-urls>

upvoted 1 times

Captain1212 1 month, 1 week ago

A seems more legit
upvoted 1 times

Neha_Pallavi 2 months, 2 weeks ago

- A. This page provides an overview of signed URLs, which give time-limited access to a specific Cloud Storage resource. Anyone in possession of the signed URL can use it while it's active, regardless of whether they have a Google account
- upvoted 1 times

ExamsFR 2 months, 3 weeks ago

Selected Answer: A
A is correct
upvoted 1 times

Buruguduystunstugudunstuy 7 months, 3 weeks ago

Selected Answer: A

Answer A is correct. Creating a signed URL with a short expiration time is a secure way to share objects in a Cloud Storage bucket with an external party, especially when the external company does not have a Google account.

Answer B is incorrect because setting object access to 'public' could potentially expose sensitive data to anyone who knows the object URL.

Answer C is incorrect because configuring the storage bucket as a static website would make the entire bucket contents available publicly, which is not secure.

Answer D is incorrect because creating a new Cloud Storage bucket for the external company to access, and then deleting the bucket after four hours, is a cumbersome and unnecessary approach. It also adds additional complexity to your environment.

upvoted 3 times

 **owliguess** 8 months, 1 week ago

Selected Answer: A

A minimal... B C D too much steps

upvoted 1 times

 **ChristN** 10 months, 1 week ago

Selected Answer: A

A. Signed URL for a certain period of time.

upvoted 1 times

 **leogor** 11 months, 3 weeks ago

A. signed URL

upvoted 1 times

 **Jesse20323645** 11 months, 3 weeks ago

Selected Answer: A

Go for (A)

Anyone in possession of the signed URL can use it while it's active, regardless of whether they have a Google account. It gives time-limited access to a specific Cloud Storage resource.

Reference:

<https://cloud.google.com/storage/docs/access-control/signed-urls>

upvoted 1 times

 **Vidhey** 11 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Answer is A

upvoted 1 times

 **sandeshsegu** 1 year, 3 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **Alok17** 1 year, 3 months ago

A is correct

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right.

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

go for A

upvoted 1 times

You are creating a Google Kubernetes Engine (GKE) cluster with a cluster autoscaler feature enabled. You need to make sure that each node of the cluster will run a monitoring pod that sends container metrics to a third-party monitoring solution. What should you do?

- A. Deploy the monitoring pod in a StatefulSet object.
- B. Deploy the monitoring pod in a DaemonSet object.
- C. Reference the monitoring pod in a Deployment object.
- D. Reference the monitoring pod in a cluster initializer at the GKE cluster creation time.

Correct Answer: B

Community vote distribution

B (100%)

 **JackGlemins**  2 years, 7 months ago

B is right: <https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>

Some typical uses of a DaemonSet are:

running a cluster storage daemon on every node
 running a logs collection daemon on every node
 running a node monitoring daemon on every node
 upvoted 39 times

 **Agents89**  3 years, 5 months ago

B is correct

upvoted 25 times

 **YourCloudGuru**  1 week, 4 days ago

Selected Answer: B

The correct answer is B

DaemonSets ensure that one running instance of the pod is scheduled on every node in the cluster. This means that even if the cluster autoscaler scales the cluster up or down, the monitoring pod will continue to run on each node.

The other options are not as good:

A This is not necessary for the monitoring pod, and it can make it more difficult to scale the cluster
 C This is not necessary for the monitoring pod, and it can make it more difficult to scale the cluster
 D This is not the best way to deploy the monitoring pod, because it makes it difficult to update the monitoring pod and it does not ensure that the monitoring pod is running on every node in the cluster.

<https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

B seems more correct

upvoted 1 times

 **sthapit** 2 months ago

Hi if you have the complete pdf please share it to suspence.suraj@gmail.com

upvoted 1 times

 **Neha_Pallavi** 2 months, 2 weeks ago

Correct Answer is B .Hi if you have the complete pdf please share it to javac.krishnan@gmail.com , thanks

upvoted 1 times

 **beth_pheonix** 3 months, 1 week ago

Hi if you have the complete pdf please share it to bethpheonix8@gmail.com , thanks

upvoted 1 times

 **gksardar** 5 months, 1 week ago

If you have the complete pdf please share it to gayatri.sardar03@gmail.com... Thank you so much in advance

upvoted 2 times

 **Suder** 4 months, 3 weeks ago

Hi if you have got complete pdf , please share it to suderson.ravi89@gmail.com.Thanks in advance

upvoted 1 times

 **senatori** 6 months, 4 weeks ago

Selected Answer: B

daemonset- makes sure that a fixed number of pods is running in every moment in every node of the cluster

upvoted 3 times

 **Buruguduystunstugudunstu** 7 months, 3 weeks ago

Selected Answer: B

Answer B, Deploy the monitoring pod in a DaemonSet object, is the correct answer.

A DaemonSet ensures that all (or some) nodes in a cluster run a copy of a specific Pod. By deploying the monitoring pod in a DaemonSet object, a copy of the pod will run on each node of the cluster. This ensures that the metrics for all containers running on each node will be sent to the third-party monitoring solution.

A. StatefulSet is used for stateful applications that require unique network identifiers, stable storage, and ordered deployment and scaling.

C. A Deployment object is used to manage a set of replica Pods in a declarative way.

D. Cluster initializers are deprecated and no longer recommended for use in Kubernetes.

upvoted 13 times

 **leogor** 11 months, 3 weeks ago

B DaemonSet

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is correct

upvoted 1 times

 **pfabio** 1 year, 4 months ago

Selected Answer: B

DaemonSets attempt to adhere to a one-Pod-per-node model, either across the entire cluster or a subset of nodes. As you add nodes to a node pool, DaemonSets automatically add Pods to the new nodes as needed.

upvoted 7 times

 **haroldbenites** 1 year, 4 months ago

go for B

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for B

upvoted 1 times

 **adityachowta** 1 year, 5 months ago

Selected Answer: B

B is correct

upvoted 1 times

 **obeythefist** 1 year, 7 months ago

"Every Node" is the keyword here, which is what DaemonSet is used for

upvoted 2 times

You want to send and consume Cloud Pub/Sub messages from your App Engine application. The Cloud Pub/Sub API is currently disabled. You will use a service account to authenticate your application to the API. You want to make sure your application can use Cloud Pub/Sub. What should you do?

- A. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.
- B. Rely on the automatic enablement of the Cloud Pub/Sub API when the Service Account accesses it.
- C. Use Deployment Manager to deploy your application. Rely on the automatic enablement of all APIs used by the application being deployed.
- D. Grant the App Engine Default service account the role of Cloud Pub/Sub Admin. Have your application enable the API on the first connection to Cloud Pub/Sub.

Correct Answer: A*Community vote distribution*

A (100%)

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (A)

Quickstart: using the Google Cloud Console

This page shows you how to perform basic tasks in Pub/Sub using the Google Cloud Console.

Note: If you are new to Pub/Sub, we recommend that you start with the interactive tutorial.

Before you begin

Set up a Cloud Console project.

Set up a project

Click to:

Create or select a project.

Enable the Pub/Sub API for that project.

You can view and manage these resources at any time in the Cloud Console.

Install and initialize the Cloud SDK.

Note: You can run the gcloud tool in the Cloud Console without installing the Cloud SDK. To run the gcloud tool in the Cloud Console, use Cloud Shell .

<https://cloud.google.com/pubsub/docs/quickstart-console>

upvoted 31 times

✉  **Bharathy**  3 years, 4 months ago

We need to enable the pub/sub API, if we are going to use it in your project... then APP engine can able to access it with required ServiceAccount

upvoted 23 times

✉  **Captain1212**  1 month, 1 week ago

Yes A, is more correct as first you need to enable the API itself

upvoted 1 times

✉  **bidyut123** 4 months, 1 week ago

Selected Answer: A

ANSWER SHOULD BE A.

upvoted 2 times

✉  **Buruguduystunstugudunsty** 7 months, 3 weeks ago

Selected Answer: A

Answer A is correct. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.

Since the Cloud Pub/Sub API is currently disabled, the first step is to enable it. This can be done through the API Library on the GCP Console. Once the API is enabled, the service account can be used to authenticate the App Engine application to the Cloud Pub/Sub API.

Answer B is incorrect because there is no automatic enablement of APIs when a service account accesses them. The API needs to be enabled manually in the API Library or through the command-line interface.

Answer C is incorrect because enabling APIs through Deployment Manager requires that the APIs be enabled in the project before Deployment Manager can use them.

Answer D is incorrect because granting the App Engine Default service account the Cloud Pub/Sub Admin role could be a security risk, and it is not necessary to enable the API.

upvoted 7 times

✉  **AzureDP900** 1 year, 3 months ago

A is right
upvoted 1 times

 **sedado77** 1 year, 4 months ago

Selected Answer: A
Yup its A
upvoted 3 times

 **haroldbenites** 1 year, 4 months ago

go for A
upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for A
upvoted 1 times

 **Schladde** 1 year, 4 months ago

Selected Answer: A
A is correct, although there are options to activate APIs programmatically (gcloud, curl) <https://cloud.google.com/service-usage/docs/enable-disable>
B and C are incorrect because "rely on ..." is never a good option
D could be possible, but too much permissions are given to the app
upvoted 2 times

 **gielda211** 1 year, 6 months ago

Selected Answer: A
A is correct
upvoted 1 times

 **flashlight01** 1 year, 6 months ago

Selected Answer: A
Enable the Cloud Pub/Sub API in the API Library on the GCP Console.
upvoted 1 times

 **obeythefist** 1 year, 7 months ago

This question seems to be asking "Can you automatically enable the API or do you have to do it manually?" and I think the answer is that there's no automatic enablement of the API.
upvoted 2 times

 **shawnkkk** 1 year, 10 months ago

A. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.
upvoted 2 times

 **vishnukumatr** 1 year, 10 months ago

A. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.
upvoted 1 times

 **Jaira1256** 1 year, 10 months ago

A is correct
upvoted 1 times

 **vikram___** 2 years, 1 month ago

Enabling API is the first step before using it, Ans A
upvoted 1 times

You need to monitor resources that are distributed over different projects in Google Cloud Platform. You want to consolidate reporting under the same Stackdriver

Monitoring dashboard. What should you do?

- A. Use Shared VPC to connect all projects, and link Stackdriver to one of the projects.
- B. For each project, create a Stackdriver account. In each project, create a service account for that project and grant it the role of Stackdriver Account Editor in all other projects.
- C. Configure a single Stackdriver account, and link all projects to the same account.
- D. Configure a single Stackdriver account for one of the projects. In Stackdriver, create a Group and add the other project names as criteria for that Group.

Correct Answer: D

Community vote distribution

C (83%) D (17%)

 **sahedge**  2 years, 9 months ago

First of all D is incorrect, Groups are used to define alerts on set of resources(such as VM instances, databases, and load balancers). FYI tried adding Two projects into a group it did not allowed me as the "AND"/"OR" criteria for the group failed with this combination of resources.

C is correct because,

When you intially click on Monitoring(Stackdriver Monitoring) it creates a workspace(a stackdriver account) linked to the ACTIVE(CURRENT) Project from which it was clicked.

Now if you change the project and again click onto Monitoring it would create an another workspace(a stackdriver account) linked to the changed ACTIVE(CURRENT) Project, we don't want this as this would not consolidate our result into a single dashboard(workspace/stackdriver account).

If you have accidentally created two diff workspaces merge them under Monitoring > Settings > Merge Workspaces > MERGE.

If we have only one workspace and two projects we can simply add other GCP Project under Monitoring > Settings > GCP Projects > Add GCP Projects.

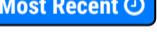
In both of these cases we did not create a GROUP, we just linked GCP Project to the workspace(stackdriver account).

upvoted 85 times

 **coldpar**  3 years, 6 months ago

C is correct not D

upvoted 33 times

 **Yad_datatonic**  1 month ago

Selected Answer: C

Option C is the recommended approach because it allows you to configure a single Stackdriver account and link all your projects to this account. This way, you can centralise monitoring, create custom dashboards, set up alerts, and gain a unified view of your resources distributed across different projects in a more straightforward and consolidated manner. It provides a single point of access and management for monitoring across all projects, which is typically the desired outcome for multi-project environments.

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: C

C is more correct, just link them into one

upvoted 1 times

 **SanjeevKumar1983** 1 month, 2 weeks ago

Selected Answer: D

<https://cloud.google.com/monitoring/settings>

Best practices for scoping projects

We recommend that you use a new Google Cloud project or one without resources as the scoping project when you want to view metrics for multiple Google Cloud projects or AWS accounts.

When a metrics scope contains monitored projects, to chart or monitor only those metrics stored in the scoping project, you must specify filters that exclude metrics from the monitored projects. The requirement to use filters increases the complexity of chart and alerting policy, and it increases the possibility of a configuration error. The recommendation ensures that these scoping projects don't generate metrics, so there are no metrics in the projects to chart or monitor.

upvoted 1 times

 **shreykul** 2 months, 2 weeks ago

Selected Answer: C

upvoted 1 times

✉ **shubham126** 6 months, 3 weeks ago

C is correct because you can monitor multiple GCP projects from a single stackdriver account

upvoted 1 times

✉ **ranit12** 7 months, 1 week ago

c should be correct

upvoted 1 times

✉ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: C

Answer C. Configure a single Stackdriver account and link all projects to the same account.

By linking all projects to the same Stackdriver account, you can consolidate monitoring and reporting for all resources across multiple projects. This can be done by following the instructions provided by Stackdriver on how to add projects to an account.

Answer A is not the best solution since Shared VPC does not consolidate the monitoring and reporting of resources distributed over different projects.

Answer B is not an efficient solution as it requires creating multiple Stackdriver accounts, which can lead to confusion and inefficiency.

Answer D is not the best solution as it only groups the projects under one account, but it does not consolidate monitoring and reporting.

upvoted 5 times

✉ **Kaushal24** 9 months ago

if anyone has contributor access please share the pdf of all questions with answers to - thedeltaproduction.16@gmail.com

upvoted 1 times

✉ **Ramanvita** 9 months, 3 weeks ago

if anyone has contributor access please share the pdf of all questions with answers to - meeeetram@gmail.com

upvoted 2 times

✉ **Nmrmmr** 7 months ago

If you have the complete pdf please share it to nmrmanasa@gmail.com

upvoted 1 times

✉ **[Removed]** 10 months, 1 week ago

Have anybody all questions and want top share? Thanks!

upvoted 1 times

✉ **Sanay** 11 months, 1 week ago

If anyone is the contributor please share the dumps on shah.sana@northeastern.edu

upvoted 1 times

✉ **leogor** 11 months, 3 weeks ago

C. Configure a single Stackdriver account, and link all projects to the same account.

upvoted 1 times

✉ **KravitejaK** 11 months, 3 weeks ago

if any one have contributor access please share all the questions pdf to raviteja922.kotha@outlook.com. I have scheduled the exam tomorrow

upvoted 1 times

✉ **Priam** 1 year ago

if anyone has contributor access please share the pdf of all questions with answers to - devappy07@gmail.com

upvoted 1 times

✉ **sandipk91** 1 year, 1 month ago

Selected Answer: C

Option C

upvoted 1 times

You are deploying an application to a Compute Engine VM in a managed instance group. The application must be running at all times, but only a single instance of the VM should run per GCP project. How should you configure the instance group?

- A. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 1.
- B. Set autoscaling to Off, set the minimum number of instances to 1, and then set the maximum number of instances to 1.
- C. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 2.
- D. Set autoscaling to Off, set the minimum number of instances to 1, and then set the maximum number of instances to 2.

Correct Answer: A

Community vote distribution

A (55%)

B (45%)

✉️  **XRiddlerX**  3 years, 3 months ago

In my GCP console, I created a managed instance group for each answer. For each answer I deleted the instance that was created as a simple test to prove or disprove each answer.

In answer A, another instance was created after I deleted the instance
 In answer B, no other instance was created after I deleted the instance
 In answer C, another instance was created after I deleted the instance
 In answer D, no other instance was created after I deleted the instance

My observation is A is the correct Answer.

A - Correct - It correctly solves the problem with only a single instance at one time
 B - Incorrect - Does not fit the requirement because AFTER the deletion of the instance, no other instance was created
 C - Incorrect - It creates another instance after the delete HOWEVER it 2 VM's could be created even if the target is exceeded
 D - Incorrect - Does not fit the requirement because AFTER the deletion of the instance, no other instance was created
 upvoted 153 times

✉️  **KerolesKhalil** 4 months, 1 week ago

Thank you for taking time to test all options

But B is the correct answer

Min and Max instances arguments are separated from autoscaling .

if VM crashed , autohealing will start a new instance once the running VM become irresponsible .

in your test, you deleted the instance that's why autohealing didn't initiate a creation of another instance

=====

Actions not initiated by the MIG, such as the following:

Preemption of a Spot VM.

Infrastructure maintenance events when the VM instance is not set to live migrate.

Deleting a VM in the group using the instances.delete method or the gcloud compute instances delete command.

<https://cloud.google.com/compute/docs/instance-groups/about-repair>

upvoted 6 times

✉️  **xaqanik** 7 months, 3 weeks ago

why we need recreation VM with autoscale if we have only instance ? in this case autohealing recreates vm if vm crashes.
 if we set max and min 1 instance to run, we dont need autoscale. we only need autohealing in case of crashing.

upvoted 6 times

✉️  **ESP_SAP** 3 years, 1 month ago

thanks so much for your time to test every option.

This kind of collaboration really clarify the doubts!

upvoted 11 times

✉️  **minnnnn** 10 months ago

Thanks so much

upvoted 1 times

✉️  **khaliqn**  3 years, 3 months ago

We want the application running at all times. If the VM crashes due to any underlying hardware failure, we want another instance to be added so we need autoscaling ON

Correct answer is A

upvoted 28 times

✉️  **[Removed]** 2 years, 6 months ago

Read this warning message: "The minimum number of instances is equal to maximum number of instances. This means the autoscaler cannot add or remove instances from the instance group. Make sure this is the correct setting."
 So when minimum is equal to maximum, it does not matter whether autoscaling is on or off. So now the question is who takes care of running the MINIMUM instances: MIG itself.

upvoted 5 times

✉️  **roaming_panda** 9 months, 2 weeks ago

nice explanation
upvoted 1 times

✉  **Gurnoor** 3 years, 3 months ago

A is wrong - What you are talking about is Autohealing. Autoscale will not rebuild the VM on the crash.
upvoted 6 times

✉  **Dips_** 2 years, 9 months ago

won't auto healing take time to recreate that VM? and there should be one VM running all time
upvoted 2 times

✉  **Ozymandiax** 2 years, 9 months ago

As said in other coment, you can fix this with Autohealing, autoscaling means more machines, autohealing means re creating of VM's
upvoted 1 times

✉  **GCP_Student1** 2 years, 8 months ago

Even if I agree with your response, it still does not meet the requirements asked in the question, which is "the application should be running all the time", because in your case when the VM instance for whatever reason stops, that mean the application will experience a downtime.
upvoted 2 times

✉  **Acielana** Most Recent 1 week, 2 days ago

Selected Answer: A

It correctly solves the problem with only a single instance at one time
upvoted 2 times

✉  **HemrajLodhi** 1 week, 4 days ago

A is correct because if you set autoscaling off then the purpose of running the application continuously is defeated ie.e no downtime.
upvoted 1 times

✉  **Jay2215** 1 week, 4 days ago

Selected Answer: A

If autoscaling remains off, it wont create another instance if existing gets deleted
upvoted 1 times

✉  **Jorgeg95** 1 week, 6 days ago

b is correct
upvoted 1 times

✉  **AdelElagawany** 2 weeks, 2 days ago

In my opinion B is the correct answer since no need for scaling if the required number of machines is 1 whatever the load is. In case of failure the autohealing will take care of creating a new instance to keep the max and the minimum number of machines equal to 1. I also tested the question using ChatGPT and give me the correct answer is (B).
upvoted 1 times

✉  **joao_01** 3 weeks ago

Selected Answer: B

I'll go with B. Since we have to scale related to some metric (e.g CPU), and the question only states to have 1 VM all the time, even if we set scaling ON and it crashed, the scaling will not do anything cause there were no CPU increase, for example. So, the scaling option here doesn't matter, what matters is the min and max options that are related to auto healing options. So the option B sets the scaling to OFF and min and max to 1. Seems legit to me.
upvoted 1 times

✉  **joao_01** 3 weeks ago

Selected Answer: B

It must be B. In the option A when we enable autoscaling, we must have different values for max and min, otherwise we will get a error.
upvoted 1 times

✉  **Yad_datatonic** 1 month ago

Selected Answer: B

To ensure that only a single instance of the VM runs per GCP project while keeping the application running at all times, you should configure the instance group as follows:

Option B: Set autoscaling to Off, set the minimum number of instances to 1, and then set the maximum number of instances to 1.

Here's why this option is the correct choice:

Setting autoscaling to Off ensures that the group does not automatically add or remove instances based on load or other conditions. This means the group will always have one instance running unless manually changed.

By setting both the minimum and maximum number of instances to 1, you ensure that there can only be one instance of the VM running at any given time. Even if the instance fails or gets terminated for some reason, the group will automatically restart it to maintain a single instance.

This configuration aligns with your requirement to have a single instance of the VM running continuously in the GCP project.

upvoted 1 times

✉  **respawn** 1 month, 1 week ago

Selected Answer: B

B is the answer, we do not need autoscaling to maintain the instance count in a mig because autohealing will take care of it.
upvoted 1 times

 **Sk2498** 2 months, 2 weeks ago

So the autoscaling need not be enabled if the max and min instances are same ?
upvoted 1 times

 **ExamsFR** 2 months, 3 weeks ago

Selected Answer: A
A is the correct answer
upvoted 1 times

 **shadow1812** 4 months ago

The correct configuration for the instance group in this scenario would be option B: Set autoscaling to Off, set the minimum number of instances to 1, and then set the maximum number of instances to 1.

By setting autoscaling to Off, you ensure that the instance group won't automatically create or remove instances based on scaling policies. This means that the group will always have a fixed number of instances.

Setting both the minimum and maximum number of instances to 1 ensures that only a single instance will be running at all times. The minimum number of instances prevents the group from scaling down to zero instances, while the maximum number of instances prevents the group from scaling up beyond one instance.

This configuration guarantees that the application will always be running, while also ensuring that only a single instance of the VM is active within the Google Cloud Platform (GCP) project.

upvoted 1 times

 **MrZed** 4 months ago

Selected Answer: B
To ensure that only a single instance of the VM runs per GCP project in a managed instance group, you should configure the instance group as follows:

B. Set autoscaling to Off, set the minimum number of instances to 1, and then set the maximum number of instances to 1.

Explanation:

By setting autoscaling to Off, you disable the automatic scaling of the managed instance group.

Setting the minimum number of instances to 1 ensures that at least one instance is always running.

Setting the maximum number of instances to 1 ensures that only one instance can be running at any given time.

Therefore, option B (Set autoscaling to Off, set the minimum number of instances to 1, and then set the maximum number of instances to 1) is the correct configuration to ensure that only a single instance of the VM runs per GCP project in the managed instance group.

upvoted 1 times

 **KerolesKhalil** 4 months, 1 week ago

Selected Answer: B
But B is the correct answer
Min and Max instances arguments are separated from autoscaling .
if VM crashed , autohealing will start a new instance once the running VM become irresponsive .
upvoted 1 times

 **Jelly_Wang** 5 months, 1 week ago

Selected Answer: A
A is correct. Set autoscaling of max and min number of instance of 1 to lock the instance number at 1. And take advantage of autoscaling to create a new instance if the old one crash. For those who vote for B, come on, you can't set max or min number of instance at all. Those are part of autoscaling. <https://cloud.google.com/sdk/gcloud/reference/compute/instance-groups/managed/set-autoscaling>
upvoted 2 times

You want to verify the IAM users and roles assigned within a GCP project named my-project. What should you do?

- A. Run gcloud iam roles list. Review the output section.
- B. Run gcloud iam service-accounts list. Review the output section.
- C. Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.
- D. Navigate to the project and then to the Roles section in the GCP Console. Review the roles and status.

Correct Answer: C

Community vote distribution

C (100%)

 **coldpar**  3 years, 6 months ago

Correct answer is C as IAM section provides the list of both Members and Roles. Option A is wrong as it would provide information about the roles only. Option B is wrong as it would provide only the service accounts. Option D is wrong as it would provide information about the roles only.

upvoted 63 times

 **Agents89**  3 years, 5 months ago

C is the correct answer
upvoted 10 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: C
c seems more legit
upvoted 1 times

 **Shweta2jun** 5 months ago

Selected Answer: C
C is correct answer
upvoted 1 times

 **Buruguduystunstugudunsty** 7 months, 3 weeks ago

Selected Answer: C
Answer C is the correct answer.

To verify the IAM users and roles assigned within a GCP project, you can navigate to the project and then to the IAM section in the GCP Console. In the IAM section, you can review the members and roles assigned to the project. This will allow you to see who has what level of access to the project resources.

Answer A is incorrect because it lists the roles available in the project, but it does not show the IAM users and roles assigned to those roles.

Answer B is incorrect because it lists the service accounts in the project, but it does not show the IAM users and roles assigned to those service accounts.

Answer D is incorrect because it lists the roles available in the project, but it does not show the IAM users and roles assigned to those roles.

upvoted 2 times

 **leogor** 11 months, 3 weeks ago

C. Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.
upvoted 1 times

 **SaiSaiA** 1 year, 2 months ago

Selected Answer: C
C is the only logical answer. If you go IAM & Admin > IAM: You can see Principals and Roles. Users, groups, service accounts are Principals
upvoted 1 times

 **NOOGler** 1 year, 3 months ago

Selected Answer: C
C is correct
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

C is right! A & B doesn't satisfy the requirement.
upvoted 1 times

 **Himadhar1997** 1 year, 4 months ago

Correct answer is C as IAM section provides the list of both Members and Roles. Option A is wrong as it would provide information about the roles only. Option B is wrong as it would provide only the service accounts. Option D is wrong
upvoted 2 times

 **haroldbenites** 1 year, 4 months ago

go for C

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 1 times

 **HarshaKumar** 1 year, 5 months ago

Selected Answer: C

C is correct

upvoted 2 times

 **obeythefist** 1 year, 7 months ago

I chose C after experimenting in the console.

A. Wrong, this just shows you all the roles, not users etc. It's not useful.

B. Wrong, This will just show service accounts, not users and roles like the question asks

C. Correct. I logged onto console and followed the steps and was able to see all the assigned users and roles.

D. Wrong. We need to see how the roles are used. I couldn't even see a "Roles" option directly, you need to browse to IAM then to the Roles subsection, which is not useful anyway.

upvoted 6 times

 **shawnkkk** 1 year, 10 months ago

C. Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.

upvoted 2 times

 **vishnukumartr** 1 year, 10 months ago

C. Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.

upvoted 2 times

 **Jaira1256** 1 year, 10 months ago

C is correct

upvoted 2 times

You need to create a new billing account and then link it with an existing Google Cloud Platform project. What should you do?

- A. Verify that you are Project Billing Manager for the GCP project. Update the existing project to link it to the existing billing account.
- B. Verify that you are Project Billing Manager for the GCP project. Create a new billing account and link the new billing account to the existing project.
- C. Verify that you are Billing Administrator for the billing account. Create a new project and link the new project to the existing billing account.
- D. Verify that you are Billing Administrator for the billing account. Update the existing project to link it to the existing billing account.

Correct Answer: B

Community vote distribution

B (68%)	A (19%)	11%
---------	---------	-----

 **coldpar**  3 years, 6 months ago

D is correct as the project is already created
upvoted 45 times

 **Shubham0072** 1 year, 8 months ago

The Question clearly says "Create a new Billing account" and a Bill Administrator cannot create A billing account so according to the question the clear Answer is B.
upvoted 21 times

 **CVGCP** 4 months, 2 weeks ago

Only options which says Create New billing account is B. This is what asked in the question. However Billing Account Administrator cannot create a new Billing account --> <https://cloud.google.com/billing/docs/how-to/billing-access>
Billing account Administrator --> Manage billing accounts (but not create them).

None of the answers are correct.

upvoted 3 times

 **spudleymcdudley** 3 years, 3 months ago

Listen to this guy. I spent 10 minutes and came to the same. <https://cloud.google.com/billing/docs/how-to/billing-access>
upvoted 8 times

 **Ixgywil** 2 years, 5 months ago

This documentation is literally saying:

"Billing Account Administrator (roles/billing.admin) - Manage billing accounts (but not create them)."

This role has nothing to do with what is required, while Project Billing Manager allows to link/unlink the project to/from a billing account.

upvoted 10 times

 **UtsavDM** 2 years, 1 month ago

It says Billing Account Admin can also link/ unlink projects from the billing account.

Here -

Billing Account Administrator:

This role is an owner role for a billing account. Use it to manage payment instruments, configure billing exports, view cost information, link and unlink projects and manage other user roles on the billing account.

upvoted 3 times

 **bhavik7899** 2 years, 3 months ago

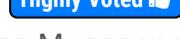
billing administration can not create billing account
so A

upvoted 3 times

 **afooh** 1 year, 8 months ago

B is the correct answer. Existing project with new billing account.

upvoted 6 times

 **Agents89**  3 years, 5 months ago

Project Billing Manager cannot create a billing account, there is nothing like Billing Administrator it is Billing Account Administrator. Both Project Billing Manager and Billing Account Administrator cannot create a billing account. A is the only answer that make sense. We have to assume the the billing account is already created

upvoted 28 times

 **spudleymcdudley** 3 years, 3 months ago

'Verify that you are Billing Administrator for the billing account' - aka Billing Account Administrator. A and B do not have privilege to make a new billing account (<https://cloud.google.com/billing/docs/how-to/billing-access>). C is correct, this has been incorrectly up voted

unvoted 7 times

 **GokulVelusaamy** 8 months, 3 weeks ago

The document says - <https://cloud.google.com/billing/docs/how-to/billing-access>

Billing Account Administrator:

By default, the person who creates the Cloud Billing account is a Billing Account Administrator for the Cloud Billing account.

Can link and unlink projects and manage other user roles on the billing account.

So the Answer is clearly C

upvoted 2 times

 **GokulVelusaamy** 8 months, 3 weeks ago

Also, If you are a billing administrator on only one Cloud Billing account, new projects you create are automatically linked to your existing Cloud Billing account.

Which clearly means, Billing Account Admin has the access to create a project

upvoted 1 times

 **jcloud965** 2 years, 2 months ago

Wrong because "You need to create a new billing account". You can't assume the billing account is already created.

upvoted 7 times

 **nightflyer** 2 years, 9 months ago

The answer is A, yes

upvoted 5 times

 **ankatsu2010** 2 years ago

Project Billing Manager can't create a billing account, still, A is the only feasible answer.

upvoted 1 times

 **AdelElagawany** Most Recent 2 weeks, 2 days ago

I'm quite confused since the question is explicitly mentioning "Creating New Billing Account" however neither "Project billing manager" nor "Billing Administrator" has the permission to create a billing account. Only the billing creator can!!

upvoted 1 times

 **joao_01** 3 weeks ago

Selected Answer: A

I think its A. C and D refers to Billing Administrator and there is only Billing Account Administrator, so C and D are out.

Only left is A and B. Both Project Billing Manager cannot create billing accounts, he can only link and unlink billing accounts to projects, NOT create them.

Thus, the option left is A, and it makes sense, cause that role can link projects to billing accounts and update the existing project with that new billing account.

upvoted 1 times

 **joao_01** 3 weeks ago

***to create billing accounts the role required is Billing Account Creator.

Link: <https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: B

b seems more legit

upvoted 1 times

 **Nxt_007** 2 months, 1 week ago

Selected Answer: B

Option B is the accurate approach. As a Project Billing Manager, you have the authority to create a new billing account and link it to an existing Google Cloud Platform project. This way, you maintain control over both the project and its associated billing.

Option A involves updating the existing project to link it to the existing billing account, which is not the correct sequence for creating a new billing account and linking it to the project.

upvoted 1 times

 **Neha_Pallavi** 2 months, 1 week ago

B is correct answer. Bill Administrator cannot create A billing account

upvoted 1 times

 **findsidd** 2 months, 1 week ago

Project Billing Manager

(roles/billing.projectManager) Link/unlink the project to/from a billing account. Organization, folder, or project. When granted in combination with the Billing Account User role, the Project Billing Manager role allows a user to attach the project to the billing account, but does not grant any rights over resources. Project Owners can use this role to allow someone else to manage the billing for the project without granting them resource access.

upvoted 1 times

 **ExamsFR** 2 months, 3 weeks ago

Selected Answer: B

B is the correct answer

upvoted 1 times

KerolesKhalil 4 months, 1 week ago

Selected Answer: D

D is the correct answer

As per GCP docs neither the Billing Administrator nor Project Billing Manager can create billing account.

so by assuming that account is already created , we only have to link it

Project Billing Manager role has to be combined with the Billing Account User role to be able to link a project to billing account so the correct answer must be D .

<https://cloud.google.com/billing/docs/how-to/billing-access#relationships-between-resources>

upvoted 2 times

trainingexam 4 months, 1 week ago

Selected Answer: D

Billing Administrator and Project Billing Manager both can link and unlink projects to/from a billing account. Billing Administrator does not have permission to create new billing account. hence it can be linked to existing account only. However, Project Billing Manager needs additional role as "Billing Account User" to attach the project to the billing account.

the question says "You need to create a new billing account" but neither Billing Administrator and Project Billing Manager have the permission to create new billing account. Therefore best matching option is that user should have Billing administrator role and use existing account to link to existing projects.

upvoted 1 times

Shweta2jun 5 months ago

The answer is B. If you read all options carefully, the only option talk about new billing account is B.

upvoted 2 times

hanweiCN 6 months ago

Selected Answer: D

A wrong, in order to link/unlink project to billing account, project billing manager must combine with billing account user.

B wrong , only Billing account creator can create a new billing account.

C wrong, Billing Administrator can not create project.

D workable, the billing account could be created by Billing account Creator ahead.

refer to :

<https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 1 times

centralelille 6 months, 1 week ago

I think that there is an issue with this question. Neither Billing Acc Admininistrator nor Billing Acc Manager have the permission to create a new billing account:

<https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 1 times

lights4life 7 months, 2 weeks ago

Answer: B, just common sense is enough

upvoted 2 times

Buruguduystunstugudunstuy 7 months, 3 weeks ago

Selected Answer: B

Answer B is correct. Verify that you are Project Billing Manager for the GCP project. Create a new billing account and link the new billing account to the existing project.

Answer A is incorrect because updating an existing project to link it to a new billing account requires that you already have the new billing account created.

Answer C is incorrect because creating a new project will not link the project to an existing billing account.

Answer D is incorrect because you need to have the Billing Account Administrator role on the billing account to update a project to link it to the billing account.

upvoted 4 times

chikorita 5 months, 3 weeks ago

you're always correct, my lord!

upvoted 1 times

Nazz1977 8 months ago

Selected Answer: D

I think it is D

upvoted 1 times

You have one project called proj-sa where you manage all your service accounts. You want to be able to use a service account from this project to take snapshots of VMs running in another project called proj-vm. What should you do?

- A. Download the private key from the service account, and add it to each VM's custom metadata.
- B. Download the private key from the service account, and add the private key to each VM's SSH keys.
- C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.
- D. When creating the VMs, set the service account's API scope for Compute Engine to read/write.

Correct Answer: C

Community vote distribution

C (100%)

 **jackdbd** Highly Voted 2 years, 2 months ago

C is the correct answer.

It took me a while to figure it out because I didn't understand how service accounts work across project. This article made it clear for me.
<https://gtseres.medium.com/using-service-accounts-across-projects-in-gcp-cf9473fef8f0>

You create the service account in proj-sa and take note of the service account email, then you go to proj-vm in IAM > ADD and add the service account's email as new member and give it the Compute Storage Admin role.

upvoted 33 times

 **JelloMan** 1 year, 6 months ago

As of now, service accounts may be impersonated (new-term). AKA, you can create a service account in one project and then impersonate it in others. Essentially, it involves the same steps as what the medium article suggests (create a service account in the principal (main) project and then add the email of the main project to the project you want to impersonate)
<https://cloud.google.com/iam/docs/impersonating-service-accounts#impersonate-sa-level>

upvoted 1 times

 **SaiSaiA** 1 year, 2 months ago

I have tried C, it doesn't work. Also, this refers to a different Principal (user) impersonating a Service Account which is a different case from what is in the question.

upvoted 1 times

 **kishoredeena** Highly Voted 3 years, 3 months ago

Option C is the right one

upvoted 22 times

 **Captain1212** Most Recent 1 month, 1 week ago

Selected Answer: C

C seems more correct, because you want to use it, you need access for it

upvoted 1 times

 **sthapit** 2 months ago

C is the answer

upvoted 1 times

 **findsidd** 2 months, 1 week ago

C is the correct answer.

Compute Storage Admin (roles/compute.storageAdmin) has permissions to create, modify, and delete disks, images, and snapshots.

For example, if your company has someone who manages project images and you don't want them to have the editor role on the project, then grant this role to their account on the project.

The most common way to let an application authenticate as a service account is to attach a service account to the resource running the application. For example, you can attach a service account to a Compute Engine instance so that applications running on that instance can authenticate as the service account. Then, you can grant the service account IAM roles to let the service account—and, by extension, applications on the instance—access Google Cloud resources.

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: C

Answer C is correct. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.

To take snapshots of VMs running in another project, you need to grant the service account that will take the snapshots the necessary IAM role to perform the action. In this case, granting the service account in the proj-sa project the Compute Storage Admin role in the proj-vm project will allow it to take snapshots of VMs running in that project.

Answers A and B are incorrect because they involve downloading and adding the private key of the service account to each VM, which is not necessary and potentially risky.

Answer D is also incorrect because setting the service account's API scope for Compute Engine to read/write only grants it permission to perform actions on resources within the same project.

<https://cloud.google.com/iam/docs/creating-managing-service-accounts>

<https://cloud.google.com/iam/docs/granting-roles-to-service-accounts>

upvoted 4 times

✉ **leogor** 11 months, 3 weeks ago

C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.

upvoted 1 times

✉ **habros** 1 year, 1 month ago

Safe to eliminate any options that demand transferring of private keys. NOT SAFE

Hence, C.

upvoted 3 times

✉ **theBestStudent** 1 year ago

highly agree with this thoughts! transferring private keys is a big no no here.

upvoted 1 times

✉ **RanjithK** 1 year, 3 months ago

Answer is C

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

C. is the correct answer

Compute Storage Admin

(roles/compute.storageAdmin)

Permissions to create, modify, and delete disks, images, and snapshots.

For example, if your company has someone who manages project images and you don't want them to have the editor role on the project, then grant this role to their account on the project.

Lowest-level resources where you can grant this role:

Disk

Image

Snapshot Beta

upvoted 3 times

✉ **haroldbenites** 1 year, 4 months ago

go for C

upvoted 1 times

✉ **somenick** 1 year, 7 months ago

Selected Answer: C

<https://cloud.google.com/compute/docs/access/iam#compute.storageAdmin>

upvoted 2 times

✉ **RealEL40** 1 year, 10 months ago

When a service account is in one project, and it accesses a resource in another project, you usually must enable the API for that resource in both projects. For example, if you have a service account in the project my-service-accounts and a Cloud SQL instance in the project my-application, you must enable the Cloud SQL API in both my-service-accounts and my-application.

upvoted 2 times

✉ **shawnkkk** 1 year, 10 months ago

C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.

upvoted 1 times

✉ **vishnukumartr** 1 year, 10 months ago

C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.

upvoted 1 times

✉ **Jaira1256** 1 year, 10 months ago

C is correct

upvoted 1 times

✉ **cuongnd** 2 years, 3 months ago

C. Compute Storage Admin role has this: compute.snapshots.*

upvoted 4 times

You created a Google Cloud Platform project with an App Engine application inside the project. You initially configured the application to be served from the us-central region. Now you want the application to be served from the asia-northeast1 region. What should you do?

- A. Change the default region property setting in the existing GCP project to asia-northeast1.
- B. Change the region property setting in the existing App Engine application from us-central to asia-northeast1.
- C. Create a second App Engine application in the existing GCP project and specify asia-northeast1 as the region to serve your application.
- D. Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application.

Correct Answer: C*Community vote distribution*

D (96%)

✉️  **Bharathy**  3 years, 6 months ago

Option D is correct, as there can be only one App Engine application inside a project . C is incorrect, as GCP can't have two app engine applications..

upvoted 83 times

✉️  **jcloud965** 2 years, 2 months ago

Yes, and you can't change an App Engine application region once created

upvoted 19 times

✉️  **saurabh1805**  3 years, 4 months ago

Each Cloud project can contain only a single App Engine application, and once created you cannot change the location of your App Engine application.

<https://cloud.google.com/appengine/docs/flexible/nodejs/managing-projects-apps-billing#create>

upvoted 33 times

✉️  **Captain1212**  1 month, 1 week ago**Selected Answer: D**

D seems more correct , as a project can only have a single app engine application

upvoted 2 times

✉️  **bobthebuilder_karkedikhayenge** 1 month, 1 week ago**Selected Answer: D**

Each Google Cloud project can contain only a single App Engine application, and once created you cannot change the location of your App Engine application

upvoted 2 times

✉️  **sthapit** 2 months ago

D as you cannot have more than one APP engine

upvoted 1 times

✉️  **Neha_Pallavi** 2 months, 1 week ago

Option D is correct.

There can be only one App Engine application inside a project

Besides, you cannot change an app's region after you set it.

<https://cloud.google.com/appengine/docs/standard/locations>

upvoted 1 times

✉️  **findsidd** 2 months, 1 week ago

Option D is correct.

There can be only one App Engine application inside a project

Besides, you cannot change an app's region after you set it.

<https://cloud.google.com/appengine/docs/standard/locations>

upvoted 1 times

✉️  **rosh199** 2 months, 2 weeks ago

Option D is correct

upvoted 1 times

✉️  **Kyle1776** 4 months, 2 weeks ago**Selected Answer: B**

To change the region for serving your existing App Engine application from the us-central region to the asia-northeast1 region, you should:

B. Change the region property setting in the existing App Engine application from us-central to asia-northeast1.

Here's why this option is the correct choice:

Region Property Setting: The region property determines the geographical location where your App Engine application is served from. By changing the region property setting from us-central to asia-northeast1, you instruct App Engine to serve your application from the desired region.

Existing App Engine Application: Since you already have an existing App Engine application, there's no need to create a new project or a new application. You can simply modify the configuration of the existing application to switch the serving region.

By following option B and changing the region property setting in the existing App Engine application from us-central to asia-northeast1, you ensure that your application is served from the desired region.

upvoted 2 times

 **efar_cloud** 4 months, 2 weeks ago

Selected Answer: D

Definitely D is correct

<https://cloud.google.com/appengine/docs/flexible/managing-projects-apps-billing#create>

upvoted 2 times

 **Bb_master** 5 months, 3 weeks ago

Selected Answer: D

D is all over the internet and on the official docs

upvoted 2 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: D

one project can contain only a single app engine

upvoted 2 times

 **chat_gpt** 7 months, 3 weeks ago

Selected Answer: D

Cannot create multiple app engine application inside one project

upvoted 2 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: D

The answer is D: Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application.

Therefore, answer D is the correct answer as it involves creating a new GCP project and App Engine application with the desired region specified from the start.

Answer A is incorrect because changing the default region property setting in the existing GCP project does not automatically change the region in which your App Engine application is served.

Answer B is incorrect because changing the region property setting in the existing App Engine application from us-central to asia-northeast1 only changes the default region for the App Engine app, but it does not actually move the app to the new region. The app would continue to be served from the original region until it is moved to the new region.

Answer C is incorrect because creating a second App Engine application in the existing GCP project does not change the region in which the original App Engine application is served.

upvoted 4 times

 **Mike_SG** 4 months, 3 weeks ago

<https://cloud.google.com/appengine/docs/standard/locations>

You cannot change an app's region after you set it.

upvoted 2 times

 **giuly992** 8 months, 3 weeks ago

Selected Answer: D

It's D because each cloud project can contain only one App Engine application so C is wrong!

And also it's not possible to change the location of an existing app engine.

upvoted 2 times

 **Ratna_15** 11 months ago

I am getting confuse which answer i need to choose correct answer or community vote distribution

upvoted 1 times

 **Ratna_15** 11 months ago

For few questions we have option c or d then which to choose correct answer or community vote distribution

upvoted 1 times

 **leogor** 11 months, 3 weeks ago

D. Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application

upvoted 1 times

You need to grant access for three users so that they can view and edit table data on a Cloud Spanner instance. What should you do?

- A. Run gcloud iam roles describe roles/spanner.databaseUser. Add the users to the role.
- B. Run gcloud iam roles describe roles/spanner.databaseUser. Add the users to a new group. Add the group to the role.
- C. Run gcloud iam roles describe roles/spanner.viewer - --project my-project. Add the users to the role.
- D. Run gcloud iam roles describe roles/spanner.viewer - --project my-project. Add the users to a new group. Add the group to the role.

Correct Answer: A

Community vote distribution

B (92%) 8%

 **yasu**  3 years, 6 months ago

I think it should be B, setup a group first are suggested way from Google.
upvoted 69 times

 **Agents89**  3 years, 5 months ago

B is the correct option
upvoted 25 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: B
b seems more legit as it will add in the group and they need edit access also
upvoted 1 times

 **Fajmayor** 1 month, 2 weeks ago

Selected Answer: B
Setup group and add role to it
upvoted 1 times

 **sthapit** 2 months ago

I go with B but TO have more controlled access, A is correct as well
upvoted 1 times

 **findsidd** 2 months, 1 week ago

Selected Answer: B
Google groups can help you manage users at scale. Each member of a Google group inherits the Identity and Access Management (IAM) roles granted to that group. This inheritance means that you can use a group's membership to manage users' roles instead of granting IAM roles to individual users.

<https://cloud.google.com/iam/docs/groups-in-cloud-console#:~:text=To%20add%20members%3A%20Click%20person,add%20them%20to%20the%20group.>
upvoted 1 times

 **Ash_34** 2 months, 3 weeks ago

Selected Answer: B
B is the correct option as spanner users are grouped into a single group and can be added to the IAM role. Easy for management work.
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B
Answer C is incorrect because the `roles/spanner.viewer` role only allows read-only access to Spanner instances, whereas the question asks for users to be able to view and edit table data.

Answer D is also incorrect for the same reason as answer C. The `roles/spanner.viewer` role does not provide the necessary permissions for editing table data.

Therefore, answers A and B are the only options that provide the `roles/spanner.databaseUser` role, which includes the necessary permissions to view and edit table data on a Cloud Spanner instance.

However, answer B is arguably better since it involves creating a new group and adding the users to that group, which can simplify the management of permissions in the future.

upvoted 7 times

 **librairica** 8 months, 3 weeks ago

Selected Answer: A
There's no mention of the 3 users being related, so why create a group? A seems best.
upvoted 2 times

 **jrlsl1991** 8 months, 2 weeks ago

Because the best practice is to create groups and assign the role to the group, not to the users directly.
upvoted 2 times

 **librairica** 8 months, 3 weeks ago

There's no mention of the 3 users being related, so why create a group? A seems best.
upvoted 1 times

 **giuly992** 8 months, 3 weeks ago

Selected Answer: B
I think it's B, because GC best practices suggest to create groups and associate users to it.
upvoted 1 times

 **jrlsl1991** 8 months, 2 weeks ago

You know what? I also thought it was B but there's no mention about best practices, and Google Groups can only be created if there's a google workspace organization. If you go to Google > Groups, you can't even create groups and the platform asks to choose an organization (from Google workspace). That has got me confused.
upvoted 1 times

 **ChristN** 10 months, 1 week ago

Selected Answer: B
Bis the correct answer
upvoted 1 times

 **PKookNN** 11 months, 3 weeks ago

Selected Answer: B
B is the most correct based on this
<https://cloud.google.com/spanner/docs/iam#spanner.databaseUser>
upvoted 1 times

 **GaneshSurwase** 12 months ago

Selected Answer: B
Ans is B
upvoted 1 times

 **Robertolo** 1 year ago

Selected Answer: A
Yeah, option B if we involve best practices. But the question does not say anything about any group. Thus option A is the correct one
upvoted 1 times

 **danny19g** 1 year, 1 month ago

Option A and B are eliminated because it talks about looking for viewer role and we want view + edit. Option B is correct as we should create a new group and add the group to the role.
upvoted 1 times

 **sandipk91** 1 year, 1 month ago

Selected Answer: B
Correct answer is option B - as per IAM best practices
upvoted 2 times

You create a new Google Kubernetes Engine (GKE) cluster and want to make sure that it always runs a supported and stable version of Kubernetes. What should you do?

- A. Enable the Node Auto-Repair feature for your GKE cluster.
- B. Enable the Node Auto-Upgrades feature for your GKE cluster.
- C. Select the latest available cluster version for your GKE cluster.
- D. Select Container-Optimized OS (cos) as a node image for your GKE cluster.

Correct Answer: B

Community vote distribution

B (100%)

 **Lush**  3 years, 5 months ago

The answer is B
<https://cloud.google.com/kubernetes-engine/versioning-and-upgrades>
 upvoted 39 times

 **4bsolut**  3 years, 3 months ago

"Creating or upgrading a cluster by specifying the version as <latest> does not provide automatic upgrades. Enable automatic node upgrades to ensure that the nodes in your cluster are up-to-date with the latest stable version." --source:
<https://cloud.google.com/kubernetes-engine/versioning-and-upgrades>

-Correct answer: B
 upvoted 29 times

 **Captain1212**  1 month, 1 week ago

b is current, as auto updates provide the more stable version
 upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B
 Answer B is correct. Google Kubernetes Engine (GKE) supports multiple versions of Kubernetes, and new versions are regularly released. To ensure that your GKE cluster runs a supported and stable version of Kubernetes, it is recommended to enable the Node Auto-Upgrades feature. This feature automatically upgrades the Kubernetes version of each node in the cluster to the latest stable version.

Answer A, enabling the Node Auto-Repair feature, is focused on repairing or replacing nodes in case they become unresponsive, but it doesn't address the need for running a supported and stable version of Kubernetes.

Answer C, selecting the latest available cluster version, may not always be the best option as new versions may have bugs or issues that have not yet been identified.

Answer D, selecting Container-Optimized OS (cos) as a node image, is focused on using a lightweight and secure operating system optimized for running containers, but it doesn't address the need for running a supported and stable version of Kubernetes.
 upvoted 4 times

 **sthapit** 2 months ago

True. B is the correct choice
 upvoted 1 times

 **leogor** 11 months, 3 weeks ago

Selected Answer: B
 B. Enable the Node Auto-Upgrades
 upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is correct.
 Creating or upgrading a cluster by specifying the version as latest does not provide automatic upgrades. Enable node auto-upgrades to ensure that the nodes in your cluster are up-to-date with the latest stable version.
 upvoted 2 times

 **pfabio** 1 year, 4 months ago

Selected Answer: B
 Node auto-upgrades help you keep the nodes in your cluster up-to-date with the cluster control plane version when your control plane is updated on your behalf. When you create a new cluster or node pool with Google Cloud console or the gcloud command, node auto-upgrade is enabled by default.
 upvoted 2 times

 **Harbeeb** 1 year, 5 months ago

Selected Answer: B

<https://cloud.google.com/kubernetes-engine/versioning-and-upgrades>

upvoted 1 times

✉ **shawnkkk** 1 year, 10 months ago

B. Enable the Node Auto-Upgrades feature for your GKE cluster.

upvoted 4 times

✉ **vishnukumartr** 1 year, 10 months ago

B. Enable the Node Auto-Upgrades feature for your GKE cluster.

upvoted 1 times

✉ **Jaira1256** 1 year, 10 months ago

Ans is B

upvoted 1 times

✉ **Rameshfm** 1 year, 11 months ago

Ans B- With auto-upgrades, GKE automatically ensures that security updates are applied and kept up to date. Ease of use: Provides a simple way to keep your nodes up to date with the latest Kubernetes features.

upvoted 2 times

✉ **viswanand** 2 years, 5 months ago

I was trying to create GKE but I dont see that option in Console

upvoted 1 times

✉ **mcaromit** 2 years, 5 months ago

B is correct

upvoted 1 times

✉ **[Removed]** 2 years, 6 months ago

B is correct. Enable the Node Auto-Upgrades feature for your GKE cluster.

upvoted 1 times

✉ **yuvি69** 2 years, 6 months ago

answer is B . <https://cloud.google.com/kubernetes-engine/versioning> in 'specifying cluster version' section

upvoted 1 times

✉ **sumanshu** 2 years, 6 months ago

Vote for B

upvoted 1 times

You have an instance group that you want to load balance. You want the load balancer to terminate the client SSL session. The instance group is used to serve a public web application over HTTPS. You want to follow Google-recommended practices. What should you do?

- A. Configure an HTTP(S) load balancer.
- B. Configure an internal TCP load balancer.
- C. Configure an external SSL proxy load balancer.
- D. Configure an external TCP proxy load balancer.

Correct Answer: A

Reference:

<https://cloud.google.com/load-balancing/docs/https/>*Community vote distribution*A (100%) **Gini**  3 years, 6 months ago

According to the documentation of SSL Proxy Load Balancing on Google, "SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing." in my opinion A should be the most suitable choice.

upvoted 59 times

 **yvinisiupacuando** 2 years, 5 months ago

Agree with you but, A is not the most suitable choice, it is the only choice, as the other Load Balancers cannot route HTTP(S) traffic.

upvoted 15 times

 **Agents89**  3 years, 5 months ago

For HTTP(s) Load balancer, the client SSL session terminates at the load balancer. A is the correct option.

upvoted 27 times

 **Captain1212**  1 month, 1 week ago Selected Answer: A

google recommend https for terminate the ssl session so A seems more legit

upvoted 1 times

 **schapit** 2 months ago

Question mentions HTTPS, SO A is the correct answer.

upvoted 2 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago Selected Answer: A

Answer A is correct. Google recommends using an HTTP(S) load balancer for terminating SSL sessions and load-balancing traffic to an instance group serving a public web application over HTTPS.

Answer B is incorrect because it is an internal load balancer, which is not suitable for serving public web applications. Internal load balancers are used for private/internal applications.

Answer C is incorrect because SSL proxy load balancers do not terminate the SSL session, instead they pass the SSL traffic directly to the backends without decrypting it. SSL proxy load balancers are used when you need to ensure that SSL is used end-to-end between the client and the backend, and when you want to offload SSL processing from the backends.

Answer D is incorrect because TCP proxy load balancers do not terminate SSL sessions. TCP proxy load balancers are used for non-HTTP traffic and can balance traffic at the TCP layer, but they do not have the ability to terminate SSL sessions.

upvoted 14 times

 **[Removed]** 6 months, 3 weeks ago

When using External SSL Proxy Load Balancing for your SSL traffic, user SSL (TLS) connections are terminated at the load balancing layer, and then proxied to the closest available backend instances by using either SSL (recommended) or TCP. For the types of backends that are supported, see Backends.

<https://cloud.google.com/load-balancing/docs/ssl/>

upvoted 1 times

 **chikorita** 5 months, 3 weeks ago

yo never fail us, my lord!

upvoted 1 times

 **leogor** 11 months, 3 weeks ago Selected Answer: A

HTTP(S) load balancer.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

I will go with A
upvoted 1 times

 **Aakash7** 1 year, 4 months ago

A is correct.
According to this guide for setting up an HTTP (S) load balancer in GCP: The client SSL session terminates at the load balancer. Sessions between the load balancer and the instance can either be HTTPS (recommended) or HTTP.
upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for C.
It dont say Global balancer.
upvoted 1 times

 **lmttt** 1 year, 5 months ago

Selected Answer: A
Correct A
upvoted 2 times

 **karim_chaara** 1 year, 10 months ago

Correct answer is A - HTTP LB
for SSL proxy, its doest support HTTP traffic
"SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing."
<https://cloud.google.com/load-balancing/docs/ssl>
upvoted 5 times

 **shawnkkk** 1 year, 10 months ago

A. Configure an HTTP(S) load balancer.
upvoted 1 times

 **vishnukumartr** 1 year, 10 months ago

A. Configure an HTTP(S) load balancer.
upvoted 1 times

 **Jaira1256** 1 year, 10 months ago

Ans is A
upvoted 1 times

 **Rameshfm** 1 year, 11 months ago

A- please refer the this link - <https://cloud.google.com/load-balancing/docs/choosing-load-balancer>
upvoted 2 times

 **sidharthwader** 2 years, 3 months ago

No doubt its answer is (A) in Azure its application gateway. These both are layer 7 traffic load balancer
upvoted 2 times

 **contaexamtopics** 2 years, 4 months ago

A or c?
upvoted 1 times

You have 32 GB of data in a single file that you need to upload to a Nearline Storage bucket. The WAN connection you are using is rated at 1 Gbps, and you are the only one on the connection. You want to use as much of the rated 1 Gbps as possible to transfer the file rapidly. How should you upload the file?

- A. Use the GCP Console to transfer the file instead of gsutil.
- B. Enable parallel composite uploads using gsutil on the file transfer.
- C. Decrease the TCP window size on the machine initiating the transfer.
- D. Change the storage class of the bucket from Nearline to Multi-Regional.

Correct Answer: B*Community vote distribution*

B (100%)

leba Highly Voted 3 years, 5 months ago

Correct answer is B as the bandwidth is good and its a single file, gsutil parallel composite uploads can be used to split the large file and upload in parallel. Refer GCP documentation - Transferring Data to GCP & upvoted 41 times

berezinsn Highly Voted 3 years, 4 months ago

Truly B is absolutely correct
upvoted 16 times

kenjaixv Most Recent 4 weeks ago

The best option to upload the file is B. Enable parallel composite uploads using gsutil on the file transfer. This is because parallel composite uploads can speed up the upload of large files by dividing them into chunks upload time.

The other options are not as effective or feasible as option B:

Option A. Use the GCP Console to transfer the file instead of gsutil is not a good choice because the GCP Console has a limit of 5 GB per file upload.

Option C. Decrease the TCP window size on the machine initiating the transfer is not advisable because it would reduce the amount of data that can be sent before receiving an acknowledgment, which could lead to lower throughput and higher latency.

Option D. Change the storage class of the bucket from Nearline to Multi-Regional is not relevant to the upload speed, as it only affects the availability and cost of storing and accessing the data.

upvoted 2 times

Captain1212 1 month, 1 week agoSelected Answer: B

b is legit correct as it helps you more to increase the speed.
upvoted 1 times

sthapit 2 months ago

Parallel composite is the right ans
upvoted 1 times

Partha117 6 months, 3 weeks agoSelected Answer: B

B is correct
upvoted 1 times

Buruguduystunstugudunstuy 7 months, 3 weeks agoSelected Answer: B

Answer B is correct. Enable parallel composite uploads using gsutil on the file transfer.

The most efficient way to upload the large file to Nearline Storage bucket using a single WAN connection rated at 1 Gbps is to enable parallel composite uploads using gsutil. By default, gsutil uses a single thread to upload a single object. But with parallel composite uploads, gsutil will split the file into smaller chunks and upload these chunks in parallel using multiple threads. This will allow the file to be uploaded faster and more efficiently.

<https://cloud.google.com/storage/docs/parallel-composite-uploads>
upvoted 8 times

Neo29 8 months, 1 week agoSelected Answer: B

B is correct Answer
upvoted 1 times

 **HiddenClouds** 9 months ago

Selected Answer: B

Just passed my exam and this question is on the exam, the correct answer is B
upvoted 3 times

 **leogor** 11 months, 3 weeks ago

Selected Answer: B

parallel composite uploads
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is right
<https://cloud.google.com/storage/docs/parallel-composite-uploads>
upvoted 3 times

 **lmttt** 1 year, 5 months ago

Selected Answer: B

b is right
upvoted 2 times

 **Ery** 1 year, 6 months ago

Selected Answer: B

<youtube.com/watch?v=NlevtGlo-E0> slice upload elephant file
upvoted 3 times

 **alaahakim** 1 year, 10 months ago

Ans : B
upvoted 2 times

 **shawnkkk** 1 year, 10 months ago

. Enable parallel composite uploads using gsutil on the file transfer.
upvoted 2 times

 **vishnukumartr** 1 year, 10 months ago

B. Enable parallel composite uploads using gsutil on the file transfer.
upvoted 1 times

 **Jaira1256** 1 year, 10 months ago

B is correct
upvoted 1 times

You've deployed a microservice called myapp1 to a Google Kubernetes Engine cluster using the YAML file specified below:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp1-deployment
spec:
  selector:
    matchLabels:
      app: myapp1
  replicas: 2
  template:
    metadata:
      labels:
        app: myapp1
    spec:
      containers:
        - name: main-container
          image: gcr.io/my-company-repo/myapp1:1.4
          env:
            - name: DB_PASSWORD
              value: "t0ugh2guess!"
          ports:
            - containerPort: 8080
```

You need to refactor this configuration so that the database password is not stored in plain text. You want to follow Google-recommended practices. What should you do?

- A. Store the database password inside the Docker image of the container, not in the YAML file.
- B. Store the database password inside a Secret object. Modify the YAML file to populate the DB_PASSWORD environment variable from the Secret.
- C. Store the database password inside a ConfigMap object. Modify the YAML file to populate the DB_PASSWORD environment variable from the ConfigMap.
- D. Store the database password in a file inside a Kubernetes persistent volume, and use a persistent volume claim to mount the volume to the container.

Correct Answer: C

Community vote distribution

B (100%)

✉  **rramani7**  3 years, 4 months ago

it is good practice to use Secrets for confidential data (like API keys) and ConfigMaps for non-confidential data (like port numbers). B is correct.

upvoted 64 times

✉  **saurabh1805**  3 years, 4 months ago

B is correct answer

<https://cloud.google.com/kubernetes-engine/docs/concepts/secret>

upvoted 35 times

✉  **hjyhf** 2 years, 2 months ago

"Storing sensitive data in Secrets is more secure than in plaintext ConfigMaps or in Pod specifications"

upvoted 8 times

✉  **bearfromoso**  2 weeks, 1 day ago

Storing database passwords, or any sensitive credentials, inside a ConfigMap is not recommended from a security standpoint. "B" it is!

upvoted 1 times

✉  **Captain1212** 1 month, 1 week ago

Selected Answer: B

b is correct as it good pracits to use secrrets for the passwords

upvoted 1 times

✉  **bobthebuilder_karkedikhayenge** 1 month, 1 week ago

Selected Answer: B

correct answer is B, as secrets are used to store credentials and configmap is used to store the configuration

upvoted 1 times

✉  **sthapit** 2 months ago

B is the right approach

upvoted 1 times

✉️ **ExamsFR** 2 months, 3 weeks ago

Selected Answer: B

B is correct.

upvoted 1 times

✉️ **raselsys** 7 months ago

B is the correct answer as configmap is configurations non confidential.

upvoted 1 times

✉️ **PB78** 7 months, 2 weeks ago

A common use case for a service is to use ConfigMaps to separate application code from configuration. ConfigMap is similar to Secret except that you use a Secret for sensitive information and you use a ConfigMap to store non-sensitive data such as connection strings, public credentials, hostnames, and URLs.

upvoted 1 times

✉️ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B

Answer B is the correct choice as it recommends storing the database password inside a Secret object, which is designed to securely store sensitive data like passwords, and then modifying the YAML file to populate the DB_PASSWORD environment variable from the Secret.

Storing sensitive information such as passwords in plain text inside configuration files is not secure and violates Google's security best practices. Instead, secrets should be stored separately and securely. In Kubernetes, secrets are designed to store sensitive information such as passwords, API keys, and tokens. Secrets are encrypted and can be used to pass sensitive data to containers in a safe manner.

To implement this in the given configuration, you can create a secret object and store the database password as a key-value pair. Then, modify the YAML file to populate the DB_PASSWORD environment variable from the secret.

<https://kubernetes.io/docs/concepts/configuration/secret/>

upvoted 1 times

✉️ **Di4sa** 8 months, 1 week ago

B is the correct answer

A Secret is an object that contains a small amount of sensitive data such as a password, a token, or a key. Such information might otherwise be put in a Pod specification or in a container image. Using a Secret means that you don't need to include confidential data in your application code.

upvoted 1 times

✉️ **glanshima** 8 months, 2 weeks ago

Selected Answer: B

B not C read here that configmap don't encrypt

<https://kubernetes.io/docs/concepts/configuration/configmap/>

upvoted 1 times

✉️ **giuly992** 8 months, 3 weeks ago

Selected Answer: B

It's B, for sensitive data it's better to use secrets

upvoted 1 times

✉️ **abmohamed** 8 months, 3 weeks ago

Selected Answer: B

Secrets used for confidential data.

upvoted 1 times

✉️ **redwing** 10 months ago

Selected Answer: B

B is correct

password using secret

upvoted 1 times

✉️ **leogor** 11 months, 3 weeks ago

Selected Answer: B

Secret object

upvoted 1 times

✉️ **Robertolo** 1 year ago

Selected Answer: B

Passwords should ALWAYS be saved as k8s secrets. There's no doubt

upvoted 1 times

You are running an application on multiple virtual machines within a managed instance group and have autoscaling enabled. The autoscaling policy is configured so that additional instances are added to the group if the CPU utilization of instances goes above 80%. VMs are added until the instance group reaches its maximum limit of five VMs or until CPU utilization of instances lowers to 80%. The initial delay for HTTP health checks against the instances is set to 30 seconds.

The virtual machine instances take around three minutes to become available for users. You observe that when the instance group autoscales, it adds more instances than necessary to support the levels of end-user traffic. You want to properly maintain instance group sizes when autoscaling. What should you do?

- A. Set the maximum number of instances to 1.
- B. Decrease the maximum number of instances to 3.
- C. Use a TCP health check instead of an HTTP health check.
- D. Increase the initial delay of the HTTP health check to 200 seconds.

Correct Answer: D*Community vote distribution*

D (100%)

berezinsn Highly Voted 3 years, 4 months ago

D is correct answer.
upvoted 21 times

professor Highly Voted 3 years, 3 months ago

Ans is D

The virtual machine instances take around three minutes to become available for users.

upvoted 18 times

Abalagu 2 years, 4 months ago

D is correct. The reason is that when you do health check, you want the VM to be working. Do the first check after initial setup time of 3 mins = 180 s < 200 s is reasonable.
upvoted 20 times

Captain1212 Most Recent 1 month, 1 week ago

Selected Answer: D
D is more correct because it gives you time to check the available instance
upvoted 1 times

gcpacepass 1 month, 4 weeks ago

If anyone has complete pdf of these questions, please send me at: hina.rana93@outlook.com
upvoted 1 times

Ashish_Tayal 6 months, 1 week ago

Selected Answer: D
First Health check must be done after proper boot of VM.
upvoted 2 times

Partha117 6 months, 3 weeks ago

Selected Answer: D
Increase delay to check all instances are available
upvoted 1 times

Buruguduystunstugudunstuy 7 months, 3 weeks ago

Selected Answer: D
Answer D is the correct solution to maintain instance group sizes when autoscaling.

When autoscaling is enabled, new instances are added based on a metric or metrics (such as CPU utilization) when certain thresholds are met. When adding new instances, it is important to ensure that only the necessary number of instances are added to the instance group and that the group size is properly maintained to prevent overprovisioning and unnecessary costs.

In this scenario, the instance group is adding more instances than necessary when autoscaling due to the initial delay of HTTP health checks. Increasing the initial delay to 200 seconds will ensure that the health check properly reflects the actual availability of the instances and prevent overprovisioning.

Answers A and B limit the maximum number of instances, which could cause issues when scaling to support higher levels of end-user traffic.

Answer C suggests using a TCP health check instead of an HTTP health check, but it does not address the issue of overprovisioning when autoscaling.

upvoted 5 times

 **leogor** 11 months, 3 weeks ago

Selected Answer: D

Increase the initial delay

upvoted 1 times

 **ale_brd_** 1 year ago

Selected Answer: D

D is correct answer.

upvoted 1 times

 **abirroy** 1 year, 2 months ago

Selected Answer: D

Increase initial delay

upvoted 1 times

 **GouseBasha** 1 year, 3 months ago

D is correct one

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

D is more appropriate for this question

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for D

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for D

upvoted 1 times

 **Ram_krishna** 1 year, 6 months ago

Selected Answer: D

logically first 3 options are out of context, so d is right

upvoted 2 times

 **vishnukumart** 1 year, 10 months ago

D. Increase the initial delay of the HTTP health check to 200 seconds.

upvoted 2 times

 **Jaira1256** 1 year, 10 months ago

Ans is D

upvoted 1 times

You need to select and configure compute resources for a set of batch processing jobs. These jobs take around 2 hours to complete and are run nightly. You want to minimize service costs. What should you do?

- A. Select Google Kubernetes Engine. Use a single-node cluster with a small instance type.
- B. Select Google Kubernetes Engine. Use a three-node cluster with micro instance types.
- C. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.
- D. Select Compute Engine. Use VM instance types that support micro bursting.

Correct Answer: C*Community vote distribution*

C (100%)

✉  **gcper**  2 years, 12 months ago

As everyone has said the answer is C but here is the source for the information. "For example, batch processing jobs can run on preemptible instances. If some of those instances stop during processing, the job slows but does not completely stop. Preemptible instances complete your batch processing tasks without placing additional workload on your existing instances and without requiring you to pay full price for additional normal instances."

srouce: <https://cloud.google.com/compute/docs/instances/preemptible>

upvoted 54 times

✉  **vnxt**  3 years, 5 months ago

I woud say C is the correct answer

upvoted 29 times

✉  **Captain1212**  1 month, 1 week ago

Selected Answer: C

c is the correct answer, use preemptible for the compute engine

upvoted 1 times

✉  **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: C

ANSWER C. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.

Preemptible VM instances offer the lowest cost for batch processing jobs in the Google Cloud Platform. Preemptible VM instances are computed instances that can run for a maximum of 24 hours and provide no availability guarantees. Preemptible VM instances are up to 80% cheaper than standard compute instances, making them an excellent choice for batch-processing workloads that can be interrupted.

The small instance type in a single-node cluster (ANSWER A) would not provide enough resources for batch processing jobs, and the micro instance types in a three-node cluster (ANSWER B) may not provide enough resources for the batch processing jobs to complete within the allotted time. VM instance types that support micro-bursting (ANSWER D) may not provide enough sustained CPU performance to complete batch processing jobs within the desired time frame.

upvoted 6 times

✉  **RAVI321** 1 year, 1 month ago

batch processing jobs can run on preemptible instances. if some of those instances stop during processing, the job slows but does not completely stop. preemptible instances complete your batch processing tasks without placing additional worklods on your existing instances and without reuring you to pay full price for additional normal instances"

upvoted 2 times

✉  **RAVI321** 1 year, 1 month ago

hey guys tell me one important thing i am learning GCP but did not get anything i mean whatever you guys are discussing in this forum
upvoted 3 times

✉  **shykot** 10 months ago

as you said you are learning, it takes time to master

upvoted 2 times

✉  **AzureDP900** 1 year, 3 months ago

C is right .

If your apps are fault-tolerant and can withstand possible instance preemptions, then preemptible instances can reduce your Compute Engine costs significantly. For example, batch processing jobs can run on preemptible instances. If some of those instances stop during processing, the job slows but does not completely stop. Preemptible instances complete your batch processing tasks without placing additional workload on your existing instances and without requiring you to pay full price for additional normal instances.

upvoted 3 times

✉  **sedado77** 1 year, 4 months ago

Selected Answer: C

Yup, C for batch and cost

upvoted 3 times

✉ **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 1 times

✉ **ruben82** 1 year, 9 months ago

I read that C is the right answer, but the question doesn't say that batch can be stopped and restarted.

upvoted 3 times

✉ **VijaySrinivasan** 1 year, 8 months ago

Batch will not be stopped and load will be shifted to another instances.

upvoted 1 times

✉ **vishnukumartr** 1 year, 10 months ago

C. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.

upvoted 3 times

✉ **Jaira1256** 1 year, 10 months ago

Ans is C

upvoted 3 times

✉ **beginnercloud** 2 years, 1 month ago

C - correctly

upvoted 2 times

✉ **dc2016bte0006** 2 years, 3 months ago

Option C is the correct answer

upvoted 3 times

✉ **mcaromit** 2 years, 5 months ago

C is correct...For cost-saving & not immediate fault-tolerant workloads like batch jobs use Preemptible VM instances

upvoted 1 times

✉ **Faizan2991** 2 years, 5 months ago

Might be D , because preemptible VM can be terminate at any time with short notice

upvoted 1 times

✉ **VijaySrinivasan** 1 year, 8 months ago

Even if the VM terminates, load will be shifted to another VM since the batch required resources

upvoted 1 times

You recently deployed a new version of an application to App Engine and then discovered a bug in the release. You need to immediately revert to the prior version of the application. What should you do?

- A. Run gcloud app restore.
- B. On the App Engine page of the GCP Console, select the application that needs to be reverted and click Revert.
- C. On the App Engine Versions page of the GCP Console, route 100% of the traffic to the previous version.
- D. Deploy the original version as a separate application. Then go to App Engine settings and split traffic between applications so that the original version serves 100% of the requests.

Correct Answer: D

Reference:

<https://medium.com/google-cloud/app-engine-project-cleanup-9647296e796a>

Community vote distribution

C (97%)

 **coldpar**  3 years, 6 months ago

correct is C NOT D.

Option A is wrong as gcloud app restore was used for backup and restore and has been deprecated. Option B is wrong as there is no application revert functionality available. Option D is wrong as App Engine maintains version and need not be redeployed.

upvoted 69 times

 **Bharathy**  3 years, 6 months ago

App engine maintains versions and to revert back to previous version, traffic can be set to 100% for the prior version.. hence correct answer is C

upvoted 26 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: C

c is the correct answer

upvoted 1 times

 **YomanB** 1 month, 1 week ago

Correct option is C.

upvoted 1 times

 **RobAlt** 1 month, 1 week ago

Selected Answer: C

App Engine Version page and route 100% to the previous version

upvoted 1 times

 **sthapit** 2 months ago

C is faster. Stick with C

upvoted 1 times

 **Partha117** 6 months, 3 weeks ago

Selected Answer: C

App engine allows versioning

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: C

ANSWER C. On the App Engine Versions page of the GCP Console, route 100% of the traffic to the previous version.

To immediately revert to the prior version of an application in App Engine, you can route 100% of the traffic to the previous version. This can be done through the App Engine Versions page of the GCP Console by selecting the desired version and selecting "Migrate traffic" and moving the slider to 100%. This will ensure that all traffic is directed to the prior version until the bug is fixed and the new version can be safely redeployed.

<https://cloud.google.com/appengine/docs/flexible/migrating-traffic>

ANSWER A (Run gcloud app restore) and ANSWER B (Click Revert on GCP Console) are not valid actions to revert to the prior version of the application. ANSWER D (Deploy the original version as a separate application) is not necessary and would complicate the environment by requiring a split traffic configuration.

upvoted 6 times

 **public_figure** 9 months, 2 weeks ago

Selected Answer: C

NO NEED OF REDEPLOY

upvoted 1 times

 **bhavyach** 9 months, 3 weeks ago

which ans need to take most voted or correct option pls suggest

upvoted 2 times

 **cslince** 10 months, 1 week ago

Selected Answer: D

correct is C

upvoted 1 times

 **leogor** 11 months, 2 weeks ago

Selected Answer: C

route 100% of the traffic to the previous version

upvoted 2 times

 **danny19g** 1 year, 1 month ago

Correct option is C.

AppEngine already creates a version for you. Also you do not create a application as one project is associated with one AppEngine application.

upvoted 2 times

 **joeMP** 1 year, 1 month ago

Selected Answer: C

App engine is a version control tool for your running app

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Answer is C

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

C is right

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 2 times

You deployed an App Engine application using gcloud app deploy, but it did not deploy to the intended project. You want to find out why this happened and where the application deployed. What should you do?

- A. Check the app.yaml file for your application and check project settings.
- B. Check the web-application.xml file for your application and check project settings.
- C. Go to Deployment Manager and review settings for deployment of applications.
- D. Go to Cloud Shell and run gcloud config list to review the Google Cloud configuration used for deployment.

Correct Answer: A

Reference:

<https://cloud.google.com/endpoints/docs/openapi/troubleshoot-aeflex-deployment>

Community vote distribution

D (93%) 7%

✉  **Bharathy**  3 years, 6 months ago

I would opt option D : as it would help to check the config details and Option A is not correct, as app.yaml would have only the runtime and script to run parameters and not the Project details

upvoted 51 times

✉  **alejandrombc** 2 years ago

Why would you choose Cloud Shell if its not even mention on the question? (what if the person did the command on its own computer?, this would not work)

upvoted 2 times

✉  **zaxxon** 1 year, 11 months ago

gcloud app deploy means sdk

upvoted 13 times

✉  **csrazdan** 1 year, 3 months ago

Regardless if you use your computer or cloud shell, you have to use SDK for gcloud command-line interface. gcloud uses a configuration file which contains default project, region and zone details so that command line can omit these parameters and use default.

upvoted 1 times

✉  **ahmed812**  3 years, 6 months ago

Option D - The config list will give the name of the project

C:\GCP\appeng>gcloud config list

[core]

account = xxx@gmail.com

disable_usage_reporting = False

project = my-first-demo-xxxx

upvoted 37 times

✉  **Captain1212**  1 month, 1 week ago

Selected Answer: D

option d as it will give you full information why it dont get deployed to the intended project

upvoted 1 times

✉  **Nxt_007** 2 months, 1 week ago

Selected Answer: D

Option D is the appropriate choice for diagnosing why the App Engine application did not deploy to the intended project. By running gcloud config list in Cloud Shell, you can view the current configuration settings, including the project ID, region, and other relevant settings used for deployment.

Options A and B involve checking the configuration files for the application (app.yaml and web-application.xml), but they may not directly provide information about where the application deployed or why it didn't deploy to the intended project.

Option C involves Deployment Manager, which is a tool for creating, deploying, and managing resources in Google Cloud Platform, but it's not specifically related to App Engine deployments and may not provide the necessary insights in this context.

upvoted 3 times

✉  **Vamshi_Krishna** 5 months, 1 week ago

Selected Answer: D

D is CORRECT

upvoted 1 times

✉  **Zahir1004** 7 months ago

Selected Answer: D

I VOTE FOR D

upvoted 1 times

Buruguduystunstugudunstuy 7 months, 3 weeks ago

Selected Answer: D

ANSWER D. CORRECT. Go to Cloud Shell and run gcloud config list to review the Google Cloud configuration used for deployment.

Running gcloud config list in the Cloud Shell will show the currently active configuration that was used for the deployment. This can help identify if the wrong project was selected or if the configuration was set up incorrectly.

<https://cloud.google.com/sdk/gcloud/reference/config/list>

ANSWER A may be helpful to ensure that the project and deployment settings are correctly specified, but it does not provide information on where the application was actually deployed.

ANSWER B is not relevant for App Engine deployments as this is an XML configuration file typically used in Java web applications deployed to servlet containers.

ANSWER C is also not relevant for App Engine deployments, as Deployment Manager is typically used to create and manage deployments of cloud infrastructure resources such as virtual machines, load balancers, and databases.

upvoted 9 times

bhavyach 9 months, 3 weeks ago

which and i need to prefer correct or voted

upvoted 1 times

cslince 10 months, 1 week ago

Selected Answer: D

option D

upvoted 1 times

leogor 11 months, 2 weeks ago

Selected Answer: D

D. check project setting by gcloud config list

upvoted 1 times

danny19g 1 year, 1 month ago

Both A and D seem correct.

Lots of people mentioned that app.yaml does not contain project id. That is an incorrect statement. Project is contained in app.YAML for "Standard" app engine application but gcloud config list has to be used for "FLEXIBLE" application. Since the questions does not inform us whether it is a standard or flexible app, Option D is correct. Look at the link here.

<https://cloud.google.com/appengine/docs/flexible/python/reference/app-yaml#app-id>.

The doc states:

"In some App Engine standard environment runtimes, you might have specified the Cloud Platform project ID (sometimes called "app ID") in the project's app.yaml file.

However, in the flexible environment, the project ID (app ID) is specified either:

By using gcloud init when you install the Google Cloud CLI. To view the default project ID of the gcloud CLI, run gcloud config list.

By using the gcloud config set project [YOUR_PROJECT_ID] command to set the default project ID of the gcloud CLI.

By using the --project flag when you deploy your app, for example: gcloud app deploy --project [YOUR_PROJECT_ID]

upvoted 8 times

RAVI321 1 year, 1 month ago

Why C is incorrect

upvoted 2 times

Haz993 1 year ago

App Engine Does not use deployment manager

upvoted 1 times

anallava 1 year, 2 months ago

D

<https://cloud.google.com/endpoints/docs/openapi/troubleshoot-gce-deployment>

upvoted 1 times

RanjithK 1 year, 3 months ago

Answer is D

upvoted 1 times

AzureDP900 1 year, 3 months ago

Dis correct

To list the project property in the core section, run:

gcloud config list project

upvoted 1 times

haroldbenites 1 year, 4 months ago

Go for D

upvoted 1 times

haroldbenites 1 year, 4 months ago

Go for D

upvoted 1 times

You want to configure 10 Compute Engine instances for availability when maintenance occurs. Your requirements state that these instances should attempt to automatically restart if they crash. Also, the instances should be highly available including during system maintenance. What should you do?

- A. Create an instance template for the instances. Set the 'Automatic Restart' to on. Set the 'On-host maintenance' to Migrate VM instance. Add the instance template to an instance group.
- B. Create an instance template for the instances. Set 'Automatic Restart' to off. Set 'On-host maintenance' to Terminate VM instances. Add the instance template to an instance group.
- C. Create an instance group for the instances. Set the 'Autohealing' health check to healthy (HTTP).
- D. Create an instance group for the instance. Verify that the 'Advanced creation options' setting for 'do not retry machine creation' is set to off.

Correct Answer: B

Community vote distribution

A (100%)

✉ **lio123** Highly Voted 2 years, 7 months ago

A

<https://cloud.google.com/compute/docs/instances/setting-instance-scheduling-options>

onHostMaintenance: Determines the behavior when a maintenance event occurs that might cause your instance to reboot.

[Default] MIGRATE, which causes Compute Engine to live migrate an instance when there is a maintenance event.
TERMINATE, which stops an instance instead of migrating it.

automaticRestart: Determines the behavior when an instance crashes or is stopped by the system.

[Default] true, so Compute Engine restarts an instance if the instance crashes or is stopped.
false, so Compute Engine does not restart an instance if the instance crashes or is stopped.

upvoted 37 times

✉ **Imdeepak12** Highly Voted 1 year, 12 months ago

Seems like it was a very obvious option i.e. A...Who selected B, I want to know his/her location?

upvoted 23 times

✉ **aldrinzee** 10 months ago

lol, yeah i want to examine their brain as well

upvoted 2 times

✉ **Kickbutowski_** 1 year, 10 months ago

Nikki singh.

upvoted 4 times

✉ **nnecode** Most Recent 3 weeks, 5 days ago

Selected Answer: A

I choose A

upvoted 1 times

✉ **Captain1212** 1 month, 1 week ago

Selected Answer: A

A seems more legit rightt ?

upvoted 1 times

✉ **sthapit** 2 months ago

The correct option is:

A. Create an instance template for the instances. Set the 'Automatic Restart' to on. Set the 'On-host maintenance' to Migrate VM instance. Add the instance template to an instance group.

Explanation:

The scenario requires configuring Compute Engine instances for availability during maintenance while also enabling automatic restart if they crash. Google Cloud Platform provides features and settings to achieve high availability and automatic restart for instances.

Option A outlines the appropriate steps

upvoted 1 times

✉ **Vamshi_Krishna** 5 months, 1 week ago

Selected Answer: A

Create an instance template for the instances. Set the 'Automatic Restart' to on. Set the 'On-host maintenance' to Migrate VM instance. Add the instance template to an instance group.

upvoted 1 times

□ **Technobie** 6 months, 1 week ago

A is correct for sure - yt- the Technobie

upvoted 1 times

□ **MahAli** 6 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

□ **Zahir1004** 7 months ago

Selected Answer: A

I vote for D

upvoted 1 times

□ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A

ANSWER A. CORRECT. Create an instance template for the instances. Set the 'Automatic Restart' to on. Set the 'On-host maintenance' to Migrate VM instance. Add the instance template to an instance group.

To ensure the availability of instances during maintenance and restart instances in the event of a crash, you should create an instance group with an instance template that specifies 'Automatic Restart' to on. This will allow your instance to restart in the event of a crash.

Additionally, you should set the 'On-host maintenance' to Migrate VM instance so that VM instances are live migrated to another host in the event of an infrastructure maintenance event.

The instance group should have a health check configured to verify that the instances are healthy. By using an instance group, you can also take advantage of the autoscaling and load-balancing capabilities that come with instance groups.

upvoted 7 times

□ **kenrichy** 6 months, 3 weeks ago

Hi Buru, as always, amazing explanation. many thanks!!

upvoted 3 times

□ **vonquestor** 9 months, 1 week ago

Selected Answer: A

Should be A

upvoted 1 times

□ **PavanPalla** 9 months, 1 week ago

yes, correct answer is 'A' only...can see options during the Instance Template creation as well.

upvoted 1 times

□ **PavanPalla** 9 months, 1 week ago

I read it as Answer is 'C', during the instance template creation it won't give any option of Auto-restart or Auto healing or health check etc., all options will get during the creation of Instance Group.

So I feel it as 'C'

upvoted 1 times

□ **public_figure** 9 months, 2 weeks ago

Selected Answer: A

ANS: A

upvoted 1 times

□ **bhavyach** 9 months, 3 weeks ago

which one need to prefer most voted or the correct option pls suggest

upvoted 1 times

□ **shykot** 10 months ago

can someone correct this answer or clarify why B correctly

upvoted 1 times

□ **Paulo_Jorge** 10 months ago

Option A:

To configure Compute Engine instances for availability when maintenance occurs and to enable automatic restart, you can follow these steps:

In the Google Cloud Console, navigate to the Compute Engine section and select the "Instances" page.

Select the 10 instances that you want to configure for availability.

From the "Actions" menu, select "Edit instance settings".

In the "Maintenance policy" section, select the "Automatic restart" option. This will cause the instances to automatically restart if they crash.

In the "Availability policy" section, select the "Migrate" option. This will cause the instances to be migrated to another host during system maintenance, ensuring that they remain available.

Save the changes to the instance settings.

After completing these steps, the selected instances will be configured to automatically restart if they crash and will be migrated to another host during system maintenance, ensuring that they remain highly available.

upvoted 5 times

You host a static website on Cloud Storage. Recently, you began to include links to PDF files on this site. Currently, when users click on the links to these PDF files, their browsers prompt them to save the file onto their local system. Instead, you want the clicked PDF files to be displayed within the browser window directly, without prompting the user to save the file locally. What should you do?

- A. Enable Cloud CDN on the website frontend.
- B. Enable 'Share publicly' on the PDF file objects.
- C. Set Content-Type metadata to application/pdf on the PDF file objects.
- D. Add a label to the storage bucket with a key of Content-Type and value of application/pdf.

Correct Answer: C

Community vote distribution

C (100%)

 **berezinsn** Highly Voted 3 years, 4 months ago
C is correct
upvoted 24 times

 **Buruguduystunstugudunstuy** Highly Voted 7 months, 3 weeks ago
Selected Answer: C

ANSWER A, enabling Cloud CDN on the website frontend, is not relevant to displaying PDF files in the browser. Cloud CDN is a content delivery network that caches content at edge locations around the world to reduce latency and improve website performance.

ANSWER B, enabling "Share publicly" on the PDF file objects, only controls whether or not the files are accessible to users without authentication. It does not affect how the files are displayed in the browser.

ANSWER D, adding a label to the storage bucket with a key of Content-Type and value of application/pdf, is not the correct way to set the Content-Type metadata for individual objects. Labels are used for organizing resources, while metadata is used to provide information about the data itself.

Therefore, ANSWER C, setting Content-Type metadata to application/pdf on the PDF file objects, is the correct answer.
upvoted 18 times

 **dnur** 6 months, 3 weeks ago
Many thanks for clear explanations! :)
upvoted 2 times

 **Captain1212** Most Recent 1 month, 1 week ago
Selected Answer: C

answer is c , as other are not relevant
upvoted 1 times

 **YomanB** 1 month, 1 week ago
C. Set Content-Type metadata to application/pdf on the PDF file objects.

Explanation: The Content-Type metadata indicates the media type of the content and helps the browser understand how to handle the file. In this case, by setting the Content-Type metadata of the PDF files to "application/pdf," you're informing the browser that the files are in PDF format, and the browser will attempt to display them directly within the browser window, rather than prompting the user to download them.

upvoted 1 times

 **Nxt_007** 2 months, 1 week ago

Selected Answer: C
C is correct
Setting the Content-Type metadata to application/pdf on the PDF file objects instructs the web browser to treat these files as PDF documents and display them inline, rather than prompting the user to download them.
upvoted 1 times

 **Paulo_Jorge** 10 months ago

Option C:
To display PDF files directly within the browser window on a website hosted on Cloud Storage, you can follow these steps:

In the Google Cloud Console, navigate to the Cloud Storage section and select the "Buckets" page.

Select the bucket that contains the static website and the PDF files.

From the "Actions" menu, select "Edit bucket" and then go to the "Website" tab.

In the "Website Configuration" section, select the "Serve objects with this content type" option and enter "application/pdf" in the text field. This will cause PDF files to be served with the correct content type.

Save the changes to the bucket configuration.

After completing these steps, the PDF files on your website will be served with the correct content type and will be displayed directly within the browser window when clicked, without prompting the user to save the file locally.

upvoted 12 times

 **leogor** 11 months, 2 weeks ago

Selected Answer: C

C. Set Content-Type metadata to application/pdf

upvoted 1 times

 **Untamables** 11 months, 3 weeks ago

Selected Answer: C

FYI

Importance of setting the correct MIME type

https://developer.mozilla.org/en-US/docs/Web/HTTP/Basics_of_HTTP/MIME_Types#importance_of_setting_the_correct_mime_type

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

C is correct. Edit the PDF objects in Cloud Storage and reconfigure their Content-Type metadata into application/pdf.

upvoted 3 times

 **BBB_ACE** 1 year, 4 months ago

Selected Answer: C

Ans: C

upvoted 2 times

 **PLUTONIUM54** 1 year, 8 months ago

C IS CORRECT

upvoted 1 times

 **vishnukumartr** 1 year, 10 months ago

C. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times

 **Jaira1256** 1 year, 10 months ago

Ans - C

upvoted 1 times

 **sunilw** 2 years, 3 months ago

C is correct.

<https://cloud.google.com/storage/docs/metadata#content-type>

upvoted 11 times

 **vmart** 2 years, 4 months ago

c is correct

upvoted 1 times

 **[Removed]** 2 years, 6 months ago

C is correct. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times

 **yuvি69** 2 years, 6 months ago

correct option is C

upvoted 1 times

You have a virtual machine that is currently configured with 2 vCPUs and 4 GB of memory. It is running out of memory. You want to upgrade the virtual machine to have 8 GB of memory. What should you do?

- A. Rely on live migration to move the workload to a machine with more memory.
- B. Use gcloud to add metadata to the VM. Set the key to required-memory-size and the value to 8 GB.
- C. Stop the VM, change the machine type to n1-standard-8, and start the VM.
- D. Stop the VM, increase the memory to 8 GB, and start the VM.

Correct Answer: D

Community vote distribution

D (100%)

✉  **cesar7816** Highly Voted  3 years, 6 months ago

coldpar, why are you getting the people confused? you need to stop teh VM and modify the RAM, that's all
upvoted 60 times

✉  **iambatmanadarkknight** 1 year, 12 months ago

who is coldpar
upvoted 7 times

✉  **spatidar2711** 1 year, 2 months ago

He deleted his comment
upvoted 1 times

✉  **chikorita** 5 months, 3 weeks ago

i dont think we can delete our comments
upvoted 1 times

✉  **CarlS** Highly Voted  3 years, 6 months ago

D is correct. If you pay attention to the question, option C mentions n1-standard-8. That instance type has 8vCPUs and 30 GB RAM, and we only need 8GB. On top of that, it is possible to use custom machine type to adjust current VM RAM to the value we need. Got the answer from this course I did to prepare the exam: <https://www.udemy.com/course/google-cloud-associate-engineer-exam-practice-tests/?couponCode=21CDE6A4C2B95F79BD97>
good luck!

upvoted 42 times

✉  **Veera_Venkata_Satyanarayana** 1 year, 3 months ago

How to use coupon code carls
upvoted 1 times

✉  **Captain1212** Most Recent  1 month, 1 week ago

Selected Answer: D

D is the correct answer
upvoted 1 times

✉  **gpais** 1 month, 3 weeks ago

You can add extended memory only to custom machine types. Predefined machine types are not supported.
upvoted 1 times

✉  **samrat46** 6 months, 2 weeks ago

D is correct.
C.n1 standard8 has 30GB RAM.
A&B- No vm instance stop, Hence can't be updated.
upvoted 2 times

✉  **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: D

ANSWER D is correct because it is the correct process to follow to increase the memory of a virtual machine in the Google Cloud Platform.

To increase the memory of a virtual machine, you need to first stop the VM, since it is not possible to modify the memory of a running VM. Then, you can increase the memory of the VM by editing the machine type and selecting a machine type with more memory. Once you have made the change, you can start the VM again.

ANSWER A is not the best approach as it relies on live migration which can be a risky operation.

ANSWER B is incorrect because adding metadata to the VM will not change the amount of memory allocated to the VM.

ANSWER C is incorrect because changing the machine type to n1-standard-8 would also increase the number of vCPUs to 8, which may not

be necessary and could result in overprovisioning of resources. In addition, changing the machine type would also affect the cost of the VM instance, which may not be desired. Since the primary concern in this scenario is to increase memory.

upvoted 12 times

 **cslince** 10 months, 1 week ago

Selected Answer: D

Option D

upvoted 1 times

 **leogor** 11 months, 2 weeks ago

Selected Answer: D

C is incorrect because the 8 in "n1-standard-8" means 8 cpus instead of 8 gb of ram

upvoted 2 times

 **santhu1039** 1 year, 2 months ago

Selected Answer: D

D is right one

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for D

upvoted 1 times

 **ChiefArch** 1 year, 6 months ago

The reason A is not correct is because live migration simply moves an existing VM between hosts, no attributes or properties are changed otherwise. Hence, you cannot live migrate from 1 VM type to another.

<https://cloud.google.com/compute/docs/instances/live-migration>

upvoted 5 times

 **mchaconr** 1 year, 9 months ago

To be exact... There is no "n1" option that has 2vCPU and 4GB. The closest would be n1-standard-2 with 2vCPU and 7.50GB RAM. This machine is already custom. So, option D.

upvoted 2 times

 **vishnukumartr** 1 year, 10 months ago

D. Stop the VM, increase the memory to 8 GB, and start the VM.

upvoted 5 times

 **Jaira1256** 1 year, 10 months ago

Ans - D

upvoted 1 times

 **Gpjhā** 2 years, 2 months ago

We do not have the option to increase memory in GCP VM directly. Instead, machine type needs to be changed. However, given the choices, D is the best answer here.

upvoted 2 times

 **YAS007** 2 years, 2 months ago

D is correct

https://cloud.google.com/compute/docs/instances/creating-instance-with-custom-machine-type#add_extended_memory_to_an_existing_vm_instance

upvoted 4 times

 **Umesh09** 2 years, 4 months ago

why do we want to stop the VM when live migration is an option I would go for Option A

upvoted 3 times

 **YAS007** 2 years, 2 months ago

I don't find how to manually initiate a live migration, its apparently a google managed process

upvoted 1 times

You have production and test workloads that you want to deploy on Compute Engine. Production VMs need to be in a different subnet than the test VMs. All the VMs must be able to reach each other over Internal IP without creating additional routes. You need to set up VPC and the 2 subnets. Which configuration meets these requirements?

- A. Create a single custom VPC with 2 subnets. Create each subnet in a different region and with a different CIDR range.
- B. Create a single custom VPC with 2 subnets. Create each subnet in the same region and with the same CIDR range.
- C. Create 2 custom VPCs, each with a single subnet. Create each subnet in a different region and with a different CIDR range.
- D. Create 2 custom VPCs, each with a single subnet. Create each subnet in the same region and with the same CIDR range.

Correct Answer: A*Community vote distribution*

JamesBond Highly Voted 3 years, 6 months ago

A is correct
upvoted 33 times

nwk Highly Voted 2 years, 11 months ago

Vote A
<https://cloud.google.com/vpc/docs/using-vpc#subnet-rules>
Primary and secondary ranges for subnets cannot overlap with any allocated range, any primary or secondary range of another subnet in the same network, or any IP ranges of subnets in peered networks.
upvoted 21 times

ranjitsinhgutte Most Recent 2 weeks ago

A is correct
If you create more than one subnet in a VPC, the CIDR blocks of the subnets cannot overlap. For example, if you create a VPC with CIDR block 10.0.0.0/24 , it supports 256 IP addresses. You can break this CIDR block into two subnets, each supporting 128 IP addresses.
upvoted 1 times

certboss 3 weeks, 6 days ago

For anyone new to the business, prod and test networks should never talk to each other.... The requirement in this question (that both envs can reach each other) is completely against best practice and common sense... There should always be complete network isolation between prod and non-prod environments.

upvoted 1 times

fraiacca 1 month ago

Selected Answer: A
I tried to create a VPC with 2 subnets in same region and same CIDR
I got the following error
Operation type [insert] failed with message "Invalid IPCidrRange: 10.0.0.0/28 conflicts with existing subnetwork 'subnet-1' in region 'asia-east1'."
upvoted 1 times

Captain1212 1 month, 1 week ago

A is correct as it help to make sure they have a diffenret subnets
upvoted 1 times

raselsys 7 months ago

Selected Answer: A
A is the correct Answer. People voting for B need to improve their networking knowledge.
upvoted 5 times

Buruguduystunstugudunstuy 7 months, 3 weeks ago

Selected Answer: A
ANSWER A meets the requirement because it creates a single custom VPC with 2 subnets, with each subnet in a different region and with a different CIDR range. This ensures that the production and test VMs are in separate subnets and that they can communicate with each other over Internal IP without creating additional routes. Since the subnets are in different regions, they will also have different internal routing tables, which can help isolate the traffic between the two subnets. This configuration provides the necessary network isolation and connectivity required by the production and test workloads.

ANSWER B suggests creating a single custom VPC with two subnets in the same region and with the same CIDR range. However, the requirement is that production VMs need to be in a different subnet than the test VMs. With the subnets in the same region and with the same CIDR range, it would not be possible to separate the production and test VMs into different subnets. Therefore, ANSWER B does not meet the requirement.

upvoted 11 times

 **NosFerazi** 8 months ago

Selected Answer: A

<https://www.youtube.com/watch?v=XLaFU1t9pM8>

8:15

upvoted 2 times

 **sidharthwader** 8 months, 3 weeks ago

I feel the Answer is A it has to be in the same VPC to talk to each other but on 2 different subnet but is there any sense to have it different region? It should be fine with different CIDR for same region i feel

upvoted 2 times

 **leogor** 11 months, 2 weeks ago

Selected Answer: A

A, same VPC network with different CIDR range

upvoted 1 times

 **PKookNN** 11 months, 3 weeks ago

A is correct

upvoted 1 times

 **Robertolo** 12 months ago

Selected Answer: A

It has to be A: one VPC, two subnets in different regions and different CIDR range. (It would also be valid to have both subnets in the same region.)

What about option B? We know that "each primary or secondary IPv4 range for all subnets in a VPC network must be a unique valid CIDR block" (read here <https://cloud.google.com/vpc/docs/subnets#ipv4-ranges>). Thus, prod and test subnets cannot overlap -> option B is not valid.

Options C and D are not valid neither, because "all the VMs must be able to reach each other" - this will not happen if we distribute the VMs across two VPC.

upvoted 4 times

 **manukoli1986** 1 year ago

Leutenant_Ololo, I test and checked. It is B answer

upvoted 1 times

 **FeaRoX** 8 months, 1 week ago

how do you want to have 2 subnets with same CIDR? Not only in GCP but anywhere...

upvoted 1 times

 **Akhi2022** 1 year, 2 months ago

Option A is incorrect, see what us says "create each subnet in a different region and with a different CIDR range" in this case routing is required due different CIDR range.

Option B is correct, If you create 2 subnets in same CIDR you can communicate over internal IP without additional routing.

upvoted 4 times

 **Leutenant_Ololo** 1 year ago

routers will be created automatically. Just go create a new VPC with 2 subnets and then check routes.

<https://cloud.google.com/vpc/docs/routes#subnet-routes>

upvoted 2 times

 **73173v2** 1 year, 2 months ago

Selected Answer: A

Different regions is something odd, but the main reason why its A is cause the CIDR range.

CIDR is the short for Classless Inter-Domain Routing. So, if we have 2 subnets, they CAN NOT BE the use the same CIDR.

IPv4 subnet ranges

"Each primary or secondary IPv4 range for all subnets in a VPC network must be a unique valid CIDR block. Refer to the per network limits for the number of secondary IP ranges you can define."

<https://cloud.google.com/vpc/docs/vpc>

upvoted 4 times

 **uk_dataguy** 2 weeks, 1 day ago

It seems that it doesn't have to be different regions, but the CIDR range should be different.

That's the key

upvoted 1 times

 **ryumada** 1 year, 2 months ago

You should understand that the VPC is a Global Resource. You can create a VPC and having subnets accross regions. This subnets could communicate to each other through their Private IP.

<https://cloud.google.com/sql/docs/mysql/private-ip#overview>

upvoted 3 times

 **csrazdan** 1 year, 2 months ago

A is correct.

B - is incorrect because 2 subnets cannot have the same CIDR Range.

C and D - are incorrect because creating 2 VPC's will require paring and exchange routes for communication.

upvoted 3 times

You need to create an autoscaling managed instance group for an HTTPS web application. You want to make sure that unhealthy VMs are recreated. What should you do?

- A. Create a health check on port 443 and use that when creating the Managed Instance Group.
- B. Select Multi-Zone instead of Single-Zone when creating the Managed Instance Group.
- C. In the Instance Template, add the label 'health-check'.
- D. In the Instance Template, add a startup script that sends a heartbeat to the metadata server.

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances>

Community vote distribution

A (86%) 14%

 **cesar7816**  3 years, 6 months ago

I'll go with A, MIGs support autohealing, load balancing, autoscaling, and auto-updating. no the Images templates, this is set up in the MIG

upvoted 62 times

 **tanito83**  2 years, 4 months ago

The correct answer is A. Please, modify it.

upvoted 12 times

 **friacca**  1 month ago

Selected Answer: A

Only A answer has some sense

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: A

443 means http A seems more correct

upvoted 1 times

 **shtapit** 2 months ago

C is incomplete. A all the way

upvoted 1 times

 **Nxt_007** 2 months, 1 week ago

Selected Answer: A

Option A is the correct choice because it involves creating a health check specifically on port 443, which is the standard port for HTTPS traffic. This health check will monitor the health of the instances based on their ability to respond to HTTPS requests. When creating the Managed Instance Group, you would configure it to use this health check to determine the health of the instances. Options B, C, and D are not directly related to setting up proper health checks for autoscaling and ensuring unhealthy instances are recreated

upvoted 1 times

 **Backlander** 4 months ago

A for A-Game let's goooo!

upvoted 1 times

 **Vamshi_Krishna** 5 months, 1 week ago

Selected Answer: A

C is definitely incorrect. Adding a label does not recreate unhealthy VMs.

A is CORRECT.

upvoted 1 times

 **Ashish_Tayal** 6 months, 1 week ago

Selected Answer: A

As web application is on Https, so will set port 443 to check health of instance.

upvoted 1 times

 **raselsys** 7 months ago

Selected Answer: A

A is the right answer.

ANSWER C is incorrect, adding a label to the instance template, has no direct impact on the ability to recreate unhealthy VMs.

upvoted 2 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A

ANSWER A. Create a health check on port 443 and use that when creating the Managed Instance Group.

To ensure that unhealthy VMs are recreated, a health check should be created to monitor the instances in the managed instance group. This health check should be configured to check the appropriate endpoint for the web application, which in this case would be port 443 for HTTPS. If an instance is determined to be unhealthy, the instance group will automatically recreate it.

INCORRECT:

ANSWER B is not directly related to recreating unhealthy VMs, but instead ensures that the instance group spans multiple zones for increased availability.

ANSWER C, adding a label to the instance template, has no direct impact on the ability to recreate unhealthy VMs.

ANSWER D, adding a startup script to send a heartbeat to the metadata server, can help detect and recover from application-level failures, but it does not directly ensure that unhealthy VMs are recreated.

upvoted 9 times

 **Anonimos** 10 months ago

I go with A, default port is 443 for https

upvoted 1 times

 **dennydream** 11 months, 2 weeks ago

This shows the importance of reading exactly what it says. The label can throw you off. A is the answer.

upvoted 1 times

 **leogor** 11 months, 2 weeks ago

Selected Answer: A

A obviously, port 443 for HTTPS application

upvoted 1 times

 **Untamables** 11 months, 3 weeks ago

A is correct.

https://cloud.google.com/compute/docs/instance-groups/autohealing-instances-in-migs#setting_up_an_autohealing_policy

upvoted 2 times

 **ravip12345** 1 year ago

Answer is A. Just adding label doesn't do any thing

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Answer is A

upvoted 2 times

Your company has a Google Cloud Platform project that uses BigQuery for data warehousing. Your data science team changes frequently and has few members.

You need to allow members of this team to perform queries. You want to follow Google-recommended practices. What should you do?

- A. 1. Create an IAM entry for each data scientist's user account. 2. Assign the BigQuery jobUser role to the group.
- B. 1. Create an IAM entry for each data scientist's user account. 2. Assign the BigQuery dataViewer user role to the group.
- C. 1. Create a dedicated Google group in Cloud Identity. 2. Add each data scientist's user account to the group. 3. Assign the BigQuery jobUser role to the group.
- D. 1. Create a dedicated Google group in Cloud Identity. 2. Add each data scientist's user account to the group. 3. Assign the BigQuery dataViewer user role to the group.

Correct Answer: D

Reference:

<https://cloud.google.com/bigquery/docs/cloud-sql-federated-queries>

Community vote distribution

C (77%) D (23%)

✉️  **Gini**  3 years, 5 months ago

C is correct because dataViewer does not allow user to perform queries. jobUser can.
upvoted 92 times

✉️  **rbrto** 3 years, 5 months ago

dataviewer can perform queries as well. D is correct
upvoted 6 times

✉️  **borinot** 6 months ago

UPDATES and INSERTS are queries and can not be performed with dataviewer
upvoted 1 times

✉️  **kabu_shawarib** 3 years, 4 months ago

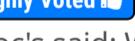
I just checked IAM roles, it doesn't have job.create like jobUser.
upvoted 6 times

✉️  **Raz0r** 1 year, 8 months ago

How about this?
<https://cloud.google.com/bigquery/docs/access-control#:~:text=Provides%20permissions%20to%20run%20jobs%2C%20including%20queries%2C%20within%20the%20project>.
upvoted 4 times

✉️  **nightflyer** 2 years, 9 months ago

incorrect
upvoted 2 times

✉️  **zukko78**  3 years, 5 months ago

C is correct, doc's said: When applied to a dataset, dataViewer provides permissions to:

Read the dataset's metadata and to list tables in the dataset.

Read data and metadata from the dataset's tables.

When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.

upvoted 45 times

✉️  **spudleymcdudley** 3 years, 3 months ago

Ref 'D' data viewer = When applied at the project or organization level, this role can also enumerate all datasets in the project.
Additional roles, however, are necessary to allow the running of jobs.
upvoted 7 times

✉️  **spudleymcdudley** 3 years, 3 months ago

listen to this guy, he's right
upvoted 2 times

✉️  **dartharlinton** 1 year, 3 months ago

according to the principle of least privilege that Google instills on its IAM offerings, answer c would be correct!
upvoted 1 times

✉️  **Captain1212**  1 month, 1 week ago

Selected Answer: C

C is the correct answer as per the , google recommended practise add them into the group then assign the role upvoted 1 times

□ **Ram65453** 1 month, 1 week ago

C is correct.

upvoted 1 times

□ **Neha_Pallavi** 1 month, 2 weeks ago

UPDATES and INSERTS are queries and can not be performed with dataviewer. So C is correct Answer upvoted 1 times

□ **Nxt_007** 2 months, 1 week ago

Selected Answer: C

C is correct

D is incorrect as the dataViewer role provides more access than needed for just performing queries. This role allows users to view data in datasets and tables, which might not be necessary or appropriate for your data science team.

upvoted 1 times

□ **broly1884** 3 months, 2 weeks ago

Selected Answer: C

dataViewer only allows viewing. Pretty obvious!

upvoted 1 times

□ **carlalap** 3 months, 2 weeks ago

The both BigQuery jobUser and dataViewer role can execute queries. But the key here, is that "data science team changes frequently", so the dataViewer role has less permissions. So, D is correct.

upvoted 1 times

□ **kumarts** 4 months ago

Selected Answer: C

C is correct, refer <https://cloud.google.com/bigquery/docs/access-control>

upvoted 1 times

□ **madboss9** 4 months, 2 weeks ago

C is the correct answer

upvoted 1 times

□ **Prashant89** 4 months, 3 weeks ago

Selected Answer: D

I am confused between C or D. I checked in gcp console,permissions for jobUser role are
bigquery.config.get

bigquery.jobs.create

resourcemanager.projects.get

resourcemanager.projects.list

and DataViewer role has more permissions. So I think answer should be D instead of C. Please correct me if I am wrong

upvoted 1 times

□ **JC0926** 7 months, 2 weeks ago

Selected Answer: C

perform queries.

upvoted 1 times

□ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: C

The correct answer is ANSWER C.

Creating a dedicated Google group in Cloud Identity is a good practice because it simplifies user management. Rather than adding individual users to each resource's IAM policy, you can add the group to the resource's IAM policy. This way, you only need to manage the group membership rather than each user's permissions. Also, the BigQuery jobUser role provides the necessary permission to run queries and jobs, which is appropriate for data scientists who need to perform queries.

upvoted 6 times

□ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

ANSWER A is incorrect because assigning the BigQuery jobUser role to individual user accounts would require granting the role to each user separately, which would be time-consuming and difficult to manage, especially for a large data science team. Additionally, the BigQuery jobUser role is primarily intended for running queries and monitoring job status, not for accessing data.

ANSWER B is incorrect because the BigQuery dataViewer role is intended for read-only access to datasets and tables. While data scientists may need to view data, they also need to write queries, which require additional permissions.

ANSWER D is incorrect because while creating a group in Cloud Identity and adding users to the group is a good approach for managing access for a large data science team, assigning the BigQuery dataViewer role to the group would not provide sufficient permissions for writing queries and creating views, which are likely required for data science work.

upvoted 4 times

□ **Yashwant_Aditya** 8 months, 1 week ago

C is the correct answer

BigQuery Job User

(roles/bigquery.jobUser)

Provides permissions to run jobs, including queries, within the project.

<https://cloud.google.com/bigquery/docs/access-control#bigquery.jobUser>

upvoted 1 times

 **nanhukumari** 9 months, 2 weeks ago

C. As for datga viewer: When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs. <https://cloud.google.com/bigquery/docs/access-control#bigquery.dataViewer>

upvoted 1 times

 **ankyt9** 10 months, 1 week ago

c is correct , Data viewer role doesn't allow to query tables/views

upvoted 1 times

 **romea2** 10 months, 2 weeks ago

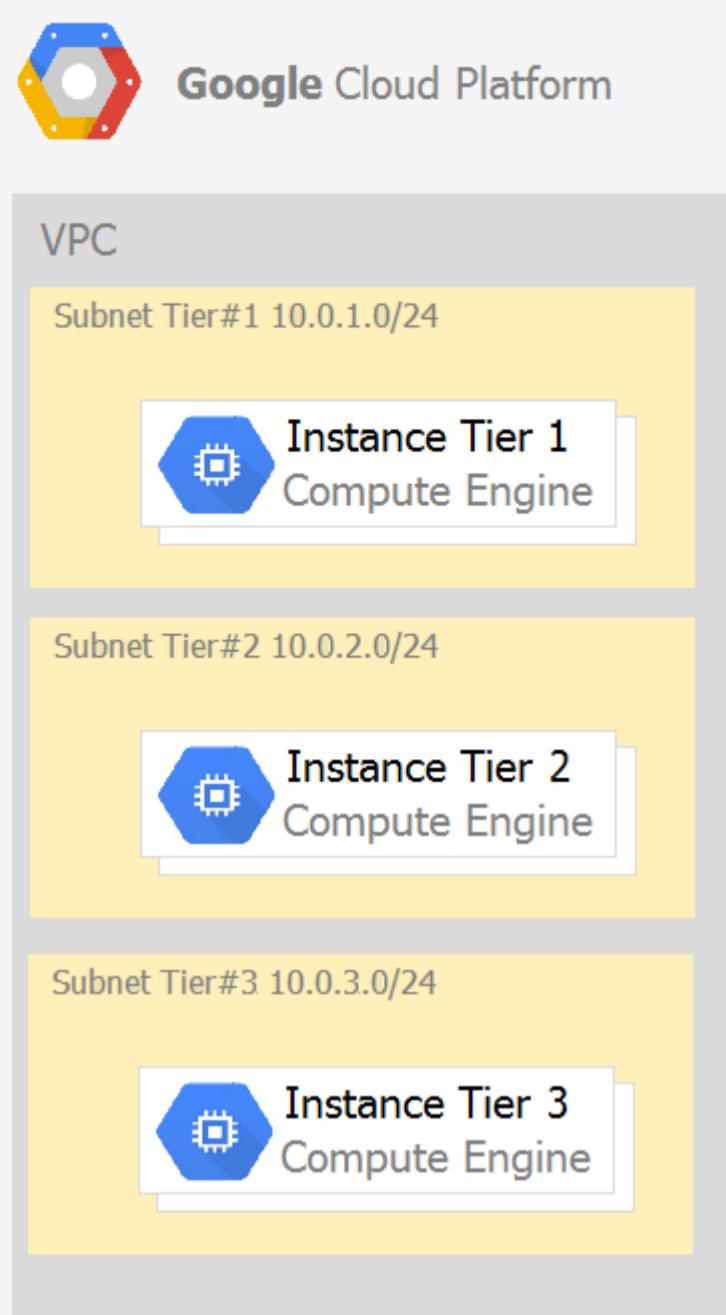
Selected Answer: C

jobUser has bigquery.jobs.create which is querying actually

answer C

upvoted 2 times

Your company has a 3-tier solution running on Compute Engine. The configuration of the current infrastructure is shown below.



Each tier has a service account that is associated with all instances within it. You need to enable communication on TCP port 8080 between tiers as follows:

- * Instances in tier #1 must communicate with tier #2.
- * Instances in tier #2 must communicate with tier #3.

What should you do?

- A. 1. Create an ingress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.2.0/24) Protocols: allow all 2. Create an ingress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.1.0/24) Protocols: allow all
- B. 1. Create an ingress firewall rule with the following settings: Targets: all instances with tier #2 service account Source filter: all instances with tier #1 service account Protocols: allow TCP:8080 2. Create an ingress firewall rule with the following settings: Targets: all instances with tier #3 service account Source filter: all instances with tier #2 service account Protocols: allow TCP: 8080
- C. 1. Create an ingress firewall rule with the following settings: Targets: all instances with tier #2 service account Source filter: all instances with tier #1 service account Protocols: allow all 2. Create an ingress firewall rule with the following settings: Targets: all instances with tier #3 service account Source filter: all instances with tier #2 service account Protocols: allow all
- D. 1. Create an egress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.2.0/24) Protocols: allow TCP: 8080 2. Create an egress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.1.0/24) Protocols: allow TCP: 8080

Correct Answer: B

Community vote distribution

B (100%)

obeythefist 1 year, 6 months ago

This question is designed to waste your time during the exam by making you read all those long answers. Remember that part of exam technique is not about knowing the product at all, but understanding multiple choice questions.

For example when two answers are very similar to each other, this can increase the likelihood that the correct answer is one of those two.

In this case it's an easy process of elimination as all answers are similar, we just need to filter out the wrong ones (and whacking the

wrong answer in an exam is sometimes the best way to find the right one).

Two answers mention port 8080, and two mention all ports. Obviously we just need port 8080, so we can immediately eliminate those two questions that want all ports open. That gives us a 50/50 chance of getting this question right.

Of the remaining answers, one says "ingress" and the other "egress". We know that by default egress is permitted and ingress is not, so that makes "b" the only surviving choice.

upvoted 59 times

 **kenrichy** 6 months, 3 weeks ago

Hi Obey, many thanks for the exam tricks

upvoted 2 times

 **kopper2019**  2 years, 6 months ago

if you see closely, port 8080 and service account is required so B is the answer without reading all answers

upvoted 25 times

 **GCPjunkie** 1 year, 8 months ago

Love they way you think, drill down to the important details.

upvoted 4 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: B

b is the correct answer

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B

ANSWER B is the correct answer because it creates ingress firewall rules that allow communication between the instances in the different tiers on TCP port 8080, based on their associated service accounts. The first rule allows traffic from instances in Tier#1 with the Tier#1 service account to instances in Tier#2 with the Tier#2 service account. The second rule allows traffic from instances in Tier#2 with the Tier#2 service account to instances in Tier#3 with the Tier#3 service account. This ensures that only the appropriate instances can communicate with each other.

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

ANSWER A is incorrect because it creates ingress firewall rules to allow communication between instances based on the IP ranges of their respective subnets. However, this doesn't guarantee that only instances in the desired tiers will be able to communicate with each other. Other instances outside the desired tiers that happen to be in the same subnet ranges will also be able to communicate.

ANSWER C is incorrect because it also allows all protocols for communication between instances in the desired tiers. This may not be desirable from a security standpoint, as it can potentially allow for unauthorized access or communication between instances.

ANSWER D is incorrect because it creates egress firewall rules instead of ingress rules. Egress rules control outbound traffic from instances, whereas ingress rules control inbound traffic. In this case, we need to control inbound traffic to allow communication between tiers on TCP port 8080.

upvoted 4 times

 **cslince** 10 months, 1 week ago

Selected Answer: B

B is the correct answer

upvoted 1 times

 **leogor** 11 months, 2 weeks ago

Selected Answer: B

B is correct obviously

upvoted 1 times

 **abirroy** 1 year, 2 months ago

Selected Answer: B

B is the correct answer

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is right, We need to open firewall rules to allow port 8080 and It shouldn't be wide open... like /24 network.

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

go for B

upvoted 1 times

 **Jerickson** 1 year, 8 months ago

Selected Answer: B

B is correct

upvoted 2 times

 **[Removed]** 1 year, 9 months ago

Selected Answer: B

B is correct
upvoted 1 times

vishnukumartr 1 year, 10 months ago

B. 1. Create an ingress firewall rule with the following settings:
Targets: all instances with tier #2 service account
Source filter: all instances with tier #1 service account
Protocols: allow TCP:8080
2. Create an ingress firewall rule with the following settings:
Targets: all instances with tier #3 service account
Source filter: all instances with tier #2 service account
Protocols: allow TCP: 8080

upvoted 1 times

Jaira1256 1 year, 10 months ago

Ans - B
upvoted 1 times

sunilw 2 years, 3 months ago

B is correct
upvoted 10 times

vmart 2 years, 4 months ago

B is correct
upvoted 2 times

shankymre01 2 years, 4 months ago

B is correct
upvoted 4 times

mcaromit 2 years, 5 months ago

B is correct
upvoted 3 times

You are given a project with a single Virtual Private Cloud (VPC) and a single subnetwork in the us-central1 region. There is a Compute Engine instance hosting an application in this subnetwork. You need to deploy a new instance in the same project in the europe-west1 region. This new instance needs access to the application. You want to follow Google-recommended practices. What should you do?

- A. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.
- B. 1. Create a VPC and a subnetwork in europe-west1. 2. Expose the application with an internal load balancer. 3. Create the new instance in the new subnetwork and use the load balancer's address as the endpoint.
- C. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Use Cloud VPN to connect the two subnetworks. 3. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.
- D. 1. Create a VPC and a subnetwork in europe-west1. 2. Peer the 2 VPCs. 3. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

Correct Answer: A

Community vote distribution

A (100%)

 **Agents89**  3 years, 5 months ago

A is correct

upvoted 37 times

 **LaxmanTiwari**  1 year, 4 months ago

Selected Answer: A

JUST PASS THE EXAM THIS MORING , THIS ONE IS THERE AND I CHOOSE A

upvoted 33 times

 **Backlander**  4 months ago

A-team FTW!

upvoted 1 times

 **miroslav_radulovic** 6 months, 1 week ago

B. 1. Create a VPC and a subnetwork in europe-west1. 2. Expose the application with an internal load balancer. 3. Create the new instance in the new subnetwork and use the load balancer's address as the endpoint.

This option follows Google-recommended practices by creating a new VPC and subnetwork in the region where the new instance will be deployed. The application is exposed using an internal load balancer, which allows the new instance to access the application using the load balancer's private IP address as the endpoint. This approach provides a secure and scalable way to connect instances across regions. Option A is incorrect because it creates a new subnetwork in the same VPC, which may cause issues with network latency and scalability. Option C is incorrect because it uses Cloud VPN, which is typically used for connecting on-premises networks to GCP, and may not be the most efficient option for connecting instances within GCP. Option D is incorrect because it peers two VPCs, which may not be the most efficient option for connecting instances within the same project.

this is what GPT said, does it make sense?

upvoted 2 times

 **romulo_rosa** 5 months, 3 weeks ago

GPT also told me B is correct. If asked "in gcp, can instances in the same vpc but in different subnets communicate using internal ip in different regions?" chatGPT answers "Yes, instances in the same VPC but in different subnets can communicate using internal IP even if they are in different regions. As long as the VPC network is set up properly, the instances can communicate with each other using their internal IP addresses, regardless of the region. However, it's important to note that traffic between regions will incur additional network egress charges, so it's important to consider the cost implications when designing your network architecture."

So I think the correct answer would still be A.

upvoted 1 times

 **arnika98** 4 months, 3 weeks ago

If you trust ChatGPT damn sure you are gonna fail the exam.

upvoted 3 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A

ANSWER A is the correct answer because it follows Google's recommended practices of using a single VPC per project and creating a new subnetwork in the same VPC in the europe-west1 region. This allows the new instance to communicate with the existing instance using its private IP address as the endpoint.

upvoted 7 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

ANSWER B is incorrect because creating a new VPC and subnetwork in the europe-west1 region is not necessary and goes against Google's recommended practices of using a single VPC per project. Additionally, using an internal load balancer to expose the

application is not necessary since the new instance will be in the same project and can communicate directly with the existing instance.

ANSWER C is also incorrect because Cloud VPN is used to establish a secure connection between a VPC and an external network, such as an on-premises data center or another cloud provider. It is not designed to enable communication between subnetworks in the same VPC, especially not across different regions.

ANSWER D is incorrect because VPC peering only works between VPCs in the same region, so it would not be possible to peer the existing VPC in us-central1 with a new VPC in europe-west1.

upvoted 4 times

 **Ary_Almeida_Junior** 9 months, 1 week ago

Answer A is correct.

upvoted 1 times

 **cslince** 10 months, 1 week ago

Selected Answer: A

A is correct

upvoted 1 times

 **leogor** 11 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

 **gcpBeginner** 1 year ago

if A is correct can someone explain what "use the first instance's private address as the endpoint" means? Does it mean to use IP from previous subnet? or does it mean use first IP from new subnet?

upvoted 2 times

 **Leutenant_Ololo** 1 year ago

A is correct. VPC allows you to spawn multiple subnets in different zones. Routing is handled automatically (because Routers are created automatically).

"use the first instance's private address as the endpoint" means that this new instance will be accessing the app via first instance's private IP (so there should be some routing rules created). Question says: "This new instance needs access to the application." ..

upvoted 2 times

 **Cornholio_LMC** 1 year ago

had this question today

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for A

upvoted 2 times

 **Bableves** 1 year, 6 months ago

A cannot be good, I mean guys you are not good at networking, if you have two different subnets, you cannot use an IP from the other subnet, just randomly, you have to "give access" which means you have to connect to the two subnets somehow, it would be better with routing, but VPN does the job...

<https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview>

Cloud VPN securely connects your peer network to your Virtual Private Cloud (VPC) network through an IPsec VPN connection. Traffic traveling between the two networks is encrypted by one VPN gateway and then decrypted by the other VPN gateway. This action protects your data as it travels over the internet. You can also connect two instances of Cloud VPN to each other.

upvoted 1 times

 **JelloMan** 1 year, 5 months ago

In GCP, VPC's are global - and subnets across different regions can be accessed using private IP's (no VPN setup required).

upvoted 19 times

 **AzureDP900** 1 year, 3 months ago

There is no need of setting VPN as you mentioned. AWS is different ..

upvoted 4 times

 **Leutenant_Ololo** 1 year ago

I mean guys you are not good at networking <-- but we are glad you are proficient with it :-D

Routing between subnets in GCP is not the same as in "regular" networking.

<https://cloud.google.com/vpc/docs/routes#subnet-routes> : "When you add a subnet, Google Cloud creates a corresponding subnet route for the subnet's primary IP address range." RTFM..

upvoted 2 times

 **Vidyaji** 1 year, 10 months ago

Selected Answer: A

A is perfect

upvoted 2 times

 **vishnukumartr** 1 year, 10 months ago

A. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

upvoted 1 times

 **Jaira1256** 1 year, 10 months ago

Ans - A

upvoted 2 times

 **Deeska** 1 year, 11 months ago

Subnets are global. A is correct

upvoted 3 times

 **akshaychavan7** 1 year, 4 months ago

subnets are regional!

upvoted 1 times

 **piyu1515** 1 year, 4 months ago

VPC ARE GLOBA,SUBNETS ARE REGIONAL

upvoted 4 times

 **vmart** 2 years, 4 months ago

A is the correct option

upvoted 3 times

Your projects incurred more costs than you expected last month. Your research reveals that a development GKE container emitted a huge number of logs, which resulted in higher costs. You want to disable the logs quickly using the minimum number of steps. What should you do?

- A. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.
- B. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE Cluster Operations resource.
- C. 1. Go to the GKE console, and delete existing clusters. 2. Recreate a new cluster. 3. Clear the option to enable legacy Stackdriver Logging.
- D. 1. Go to the GKE console, and delete existing clusters. 2. Recreate a new cluster. 3. Clear the option to enable legacy Stackdriver Monitoring.

Correct Answer: A

Community vote distribution



Gini Highly Voted 3 years, 5 months ago

The question mentioned that "GKE container emitted a huge number of logs", in my opinion A is correct.
upvoted 49 times

JackGlemins Highly Voted 2 years, 8 months ago

I think A is right.
<https://cloud.google.com/logging/docs/api/v2/resource-list>

GKE Containers have more log than GKE Cluster Operations:

.-GKE Container:

cluster_name: An immutable name for the cluster the container is running in.
namespace_id: Immutable ID of the cluster namespace the container is running in.
instance_id: Immutable ID of the GCE instance the container is running in.
pod_id: Immutable ID of the pod the container is running in.
container_name: Immutable name of the container.
zone: The GCE zone in which the instance is running.

VS

.-GKE Cluster Operations

project_id: The identifier of the GCP project associated with this resource, such as "my-project".
cluster_name: The name of the GKE Cluster.
location: The location in which the GKE Cluster is running.
upvoted 15 times

Captain1212 Most Recent 1 month, 1 week ago

Selected Answer: A
A seems more correct as it also uses the fewes steps also
upvoted 1 times

_Sande 6 months, 1 week ago

Just as a side note. Stackdriver is now Google Cloud Operations Suite.
upvoted 4 times

Buruguduystunstugudunstuy 7 months, 3 weeks ago

Selected Answer: A
The correct answer is A.

To disable the logs quickly using the minimum number of steps, you should go to the Logs ingestion window in Stackdriver Logging and disable the log source for the GKE container resource. This will prevent the GKE container from emitting logs, which will in turn reduce the amount of log data generated and lower the costs incurred.

Answer B is incorrect because disabling the log source for the GKE Cluster Operations resource would not stop the logs emitted by the GKE container, and would not reduce the amount of log data generated.

Answers C and D are incorrect because deleting and recreating the GKE cluster is not a recommended approach for disabling logs. Additionally, clearing the option to enable legacy Stackdriver Logging or Monitoring would not disable the logs emitted by the GKE container.

upvoted 5 times

cslince 10 months, 1 week ago

Selected Answer: A
A is correct
upvoted 1 times

Jagger 11 months, 2 weeks ago

Selected Answer: A

A can do it with fewest steps
upvoted 1 times

 **Untamables** 11 months, 3 weeks ago

Selected Answer: B

Currently B is correct. (Stackdriver Logging has been named Cloud Logging.)
<https://cloud.google.com/stackdriver/docs/solutions/gke/installing#migrating>
upvoted 1 times

 **abirroy** 1 year, 2 months ago

Selected Answer: A

Go for A
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

C,D are straight forward you can eliminate and decide between A & B .. A makes more sense .. logical thinking required to solve it quickly.
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A right
upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for A
upvoted 2 times

 **haroldbenites** 1 year, 4 months ago

Beginning with GKE version 1.15.7, you can configure Cloud Operations for GKE to only capture system logs and not collect application logs.
<https://cloud.google.com/stackdriver/docs/solutions/gke/managing-logs>
upvoted 2 times

 **JelloMan** 1 year, 5 months ago

Selected Answer: A

A is correct. Target GKE containers since that is what is emitting the most logs
upvoted 1 times

 **POOJA3808** 1 year, 6 months ago

Selected Answer: A

Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.
upvoted 2 times

 **Durgesh1997** 1 year, 8 months ago

Selected Answer: A

A is right answer
upvoted 1 times

 **ankit2021** 1 year, 8 months ago

Selected Answer: A

A is ans
upvoted 1 times

 **vishnukumartr** 1 year, 10 months ago

A. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.
upvoted 2 times

 **Jaira1256** 1 year, 10 months ago

A is correct
upvoted 2 times

You have a website hosted on App Engine standard environment. You want 1% of your users to see a new test version of the website. You want to minimize complexity. What should you do?

- A. Deploy the new version in the same application and use the --migrate option.
- B. Deploy the new version in the same application and use the --splits option to give a weight of 99 to the current version and a weight of 1 to the new version.
- C. Create a new App Engine application in the same project. Deploy the new version in that application. Use the App Engine library to proxy 1% of the requests to the new version.
- D. Create a new App Engine application in the same project. Deploy the new version in that application. Configure your network load balancer to send 1% of the traffic to that new application.

Correct Answer: B

Community vote distribution

B (100%)

 **yasu**  3 years, 6 months ago

I will prefer B as the answer.. why we need create new application?

upvoted 53 times

 **YAS007** 2 years, 2 months ago

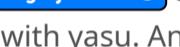
more over, in app engine we cannot create "new application", we have to create a new Project to do that, an app engine projet has 1 application (which can have multiple versions and services)

upvoted 17 times

 **sanhoo** 2 years, 4 months ago

Agree B is correct. creating a new application in the same project for app engine is anyways not possible.

upvoted 8 times

 **Gini**  3 years, 6 months ago

I agree with yasu. And only one app engine can exist in one project. B is the best choice, simple and easy.

upvoted 17 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: B

B is the correct answer

upvoted 2 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B

The correct answer is B.

By using the App Engine's traffic splitting feature, we can easily direct a certain percentage of traffic to a specific version of our application. In this case, we want to send 1% of traffic to the new test version and keep the remaining 99% on the current version. This can be achieved by deploying the new version in the same application and using the `--splits` option to give a weight of 99 to the current version and a weight of 1 to the new version.

Answer A is incorrect because the `--migrate` option is used for migrating traffic to a new version after it has been fully tested and is ready for full deployment.

Answer C is incorrect because it requires additional configuration to proxy requests to the new version, increasing complexity unnecessarily.

Answer D is incorrect because it involves configuring a network load balancer, which is not necessary for this use case and adds unnecessary complexity.

upvoted 8 times

 **Jelly_Wang** 5 months, 2 weeks ago

While I agree with your choice and your explanation of B. I also believe C and D are wrong simply because you can only have one App Engine within a project <https://cloud.google.com/appengine/docs/flexible/managing-projects-apps-billing#:~:text=Important%3A%20Each%20Cloud%20project%20can,of%20your%20App%20Engine%20application.>

upvoted 4 times

 **cslince** 10 months, 1 week ago

Selected Answer: B

B is correct

upvoted 1 times

 **leogor** 11 months, 2 weeks ago

Selected Answer: B

B. deploy new version with --splits option
upvoted 1 times

 **Cornholio_LMC** 1 year ago

had this question today
upvoted 2 times

 **habros** 1 year, 2 months ago

B! A very natural answer... Perfect for switching users over to new version. Imagine creating multiple projects to update App Engine deployments, isn't that logically unnecessary?
upvoted 1 times

 **Madj** 1 year, 3 months ago

Hint:
One app engine per project. So Option C,D eliminated. this hint will help in many similar questions.
Splitting traffic hint will help as well
upvoted 3 times

 **AzureDP900** 1 year, 3 months ago

B is right.
upvoted 1 times

 **mplibunao** 1 year, 4 months ago

Selected Answer: B
b is my answer.
a: --migrate is for enabling gradual traffic migration as opposed to migrating traffic immediately
c & d: no need to create a project. You can split the traffic any time
upvoted 3 times

 **haroldbenites** 1 year, 4 months ago

Go for B
upvoted 1 times

 **Durgesh1997** 1 year, 8 months ago

B is the correct answer as there is only one app engine can be created per project
upvoted 2 times

 **Priyankahere** 1 year, 8 months ago

This question was there , go with community answers.
upvoted 2 times

 **[Removed]** 1 year, 9 months ago

B,
--splits exists for such requirements
upvoted 2 times

 **vishnukumartr** 1 year, 10 months ago

B. Deploy the new version in the same application and use the --splits option to give a weight of 99 to the current version and a weight of 1 to the new version.
upvoted 1 times

 **Jaira1256** 1 year, 10 months ago

Ans - B
upvoted 2 times

You have a web application deployed as a managed instance group. You have a new version of the application to gradually deploy. Your web application is currently receiving live web traffic. You want to ensure that the available capacity does not decrease during the deployment. What should you do?

- A. Perform a rolling-action start-update with maxSurge set to 0 and maxUnavailable set to 1.
- B. Perform a rolling-action start-update with maxSurge set to 1 and maxUnavailable set to 0.
- C. Create a new managed instance group with an updated instance template. Add the group to the backend service for the load balancer. When all instances in the new managed instance group are healthy, delete the old managed instance group.
- D. Create a new instance template with the new application version. Update the existing managed instance group with the new instance template. Delete the instances in the managed instance group to allow the managed instance group to recreate the instance using the new instance template.

Correct Answer: C

Community vote distribution

B (100%)

 **Carls**  3 years, 6 months ago

Correct option is B. We need to ensure the global capacity remains intact, for that reason we need to establish maxUnavailable to 0. On the other hand, we need to ensure new instances can be created. We do that by establishing the maxSurge to 1. Option C is more expensive and more difficult to set up and option D won't meet requirements since it won't keep global capacity intact.

upvoted 89 times

 **yanlingmedal71** 10 months ago

maxSurge- configure how many new instances the MIG can create above its targetSize during an automated update. For example, if you set maxSurge to 5, the MIG uses the new instance template to create up to 5 new instances above your target size. Setting a higher maxSurge value speeds up your update, at the cost of additional instances

upvoted 8 times

 **space_cadet** 7 months, 2 weeks ago

Thanks for this.

And setting it to one makes sense, seeing that we want a gradual update

upvoted 1 times

 **JavierCorrea**  3 years, 1 month ago

I take my own previous comment back. It's definitely B.

upvoted 19 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: B

B is the correct answer

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B

Answer B is the correct answer because it allows for a safe and controlled rolling deployment with zero downtime and without reducing the available capacity during the deployment.

The `maxSurge` parameter controls the maximum number of new instances that can be created above the desired number of instances during the update process. By setting `maxSurge` to 1, the new version of the application can be gradually rolled out while maintaining the same number of available instances.

The `maxUnavailable` parameter controls the maximum number of instances that can be unavailable during the update process. By setting `maxUnavailable` to 0, at least one instance of the previous version will be available at all times, ensuring that there is no decrease in available capacity during the deployment.

By performing a rolling update with `maxSurge` set to 1 and `maxUnavailable` set to 0, the new version of the application can be gradually deployed with zero downtime and no decrease in available capacity.

upvoted 10 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Answer A is incorrect because setting maxSurge to 0 means that no additional instances are created beyond the existing number of instances in the group, which can potentially lead to a decrease in capacity. Also, setting maxUnavailable to 1 means that one instance can be unavailable at any given time, which can potentially lead to some users experiencing downtime.

Answer C is incorrect because creating a new managed instance group would require adding the new group to the backend service, which can take time and potentially cause downtime. Also, deleting the old managed instance group before ensuring that the new group is healthy can cause a decrease in capacity.

may take some time for new instances to be created with the new instance template. Also, the new instances may take time to warm up, which can cause a delay in serving traffic.

upvoted 3 times

✉ **vlodia** 9 months ago

If you do not want any unavailable machines during an update, set the maxUnavailable value to 0 and the maxSurge value to greater than 0. With these settings, Compute Engine removes each old machine only after its replacement new machine is created and running.

https://cloud.google.com/compute/docs/instance-groups/rolling-out-updates-to-managed-instance-groups#max_unavailable

upvoted 3 times

✉ **rajivdutt** 9 months, 1 week ago

If you do not want any unavailable machines during an update, set the maxUnavailable value to 0 and the maxSurge value to greater than 0. With these settings, Compute Engine removes each old machine only after its replacement new machine is created and running.

upvoted 1 times

✉ **Mission94** 9 months, 4 weeks ago

HI all,

if you guys have all the questions and answers please mail it to
untranslatable[dot]character@gmail[dot]com

Thanks in advance.

upvoted 2 times

✉ **Rubankumar** 10 months ago

Selected Answer: B

B is Correct

upvoted 1 times

✉ **Shubamakabra** 10 months, 1 week ago

Selected Answer: B

this this this this this

upvoted 1 times

✉ **Zoze** 10 months, 3 weeks ago

Selected Answer: B

I have seen this question in others websites, and in all of them, the answer is B B B!!.

upvoted 2 times

✉ **leogor** 11 months, 2 weeks ago

Selected Answer: B

B. maxUnavailable set to 0 is the key

upvoted 1 times

✉ **gcpBeginner** 1 year ago

it should be B. if we change template it cause outage and question said no outage allowed.

upvoted 1 times

✉ **ashit4424** 1 year ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!

upvoted 1 times

✉ **king2658** 1 year ago

try searching on google, may be you can find it.

upvoted 2 times

✉ **RanjithK** 1 year, 3 months ago

Answer is B

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

Correct option is B. Perform a rolling-action start-update with maxSurge set to 1 and maxUnavailable set to 0. This is also present on Tutorial Dojo practice questions.

upvoted 1 times

✉ **Backlander** 4 months ago

Sir, this is a GCP certification platform..

upvoted 1 times

✉ **LaxmanTiwari** 1 year, 4 months ago

WAS IN MY EXAM GO WITY THE MAJORITY

upvoted 3 times

✉ **haroldbenites** 1 year, 4 months ago

Go for B

upvoted 1 times

You are building an application that stores relational data from users. Users across the globe will use this application. Your CTO is concerned about the scaling requirements because the size of the user base is unknown. You need to implement a database solution that can scale with your user growth with minimum configuration changes. Which storage solution should you use?

- A. Cloud SQL
- B. Cloud Spanner
- C. Cloud Firestore
- D. Cloud Datastore

Correct Answer: B

Community vote distribution

B (100%)

✉ **Fidget_** Highly Voted 3 years, 2 months ago

B

Cloud SQL for small relational data, scaled manually
Cloud Spanner for relational data, scaled automatically
Cloud Firestore for app-based data(?)
Cloud Datastore for non-relational data

Correct me if i'm wrong

upvoted 94 times

✉ **theBestStudent** 1 year, 4 months ago

Just one detail: Cloud Firestore for non relational data (noSql)

upvoted 10 times

✉ **Backlander** 4 months ago

N E R D

upvoted 4 times

✉ **KC_go_reply** 6 months, 2 weeks ago

'small relational data' as in 3 TB for Shared core or 64 TB for Dedicated core in Cloud SQL

upvoted 1 times

✉ **karol_wu** Highly Voted 3 years, 6 months ago

in my opinion correct is B

upvoted 29 times

✉ **Captain1212** Most Recent 1 month, 1 week ago

Selected Answer: B

B, for large and automatically scaled

upvoted 1 times

✉ **keton** 4 months, 4 weeks ago

Correct ans is B... Focus on two words "Relational" which means option C & D has been eliminated bcz these are non-relational DB.And another word 'Globally' which means Option A also eliminated bcz Cloud Sql does not support global deployments.

upvoted 3 times

✉ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B

The best storage solution for this scenario would be Cloud Spanner. Cloud Spanner is a fully managed, scalable, relational database that is designed to handle global deployments with ease. It can handle large amounts of data and high transactional volumes. It also provides automatic sharding and synchronous replication, ensuring high availability and durability of data. Cloud Spanner supports SQL semantics and provides a familiar relational database experience to developers, which would make it easy to adopt in existing workflows.

Cloud SQL, on the other hand, has limits on scalability and does not support global deployments as well as Cloud Spanner.

Cloud Firestore and Cloud Datastore are NoSQL databases that are better suited for document-based data storage and not optimized for relational data storage.

upvoted 8 times

✉ **cslince** 10 months, 1 week ago

Selected Answer: B

correct is B

upvoted 1 times

✉ **Tmitchelltec919** 10 months, 4 weeks ago

could someone please explain why the answer is not A

upvoted 1 times

✉ **leogor** 11 months, 2 weeks ago

Selected Answer: B

B. Spanner for autoscale

upvoted 1 times

✉ **learn_GCP** 1 year ago

Selected Answer: B

B.

keywords: "global" and "relational"

upvoted 2 times

✉ **gcp_world123** 1 year, 1 month ago

Correct answer is B

Cloud Firestore and Cloud Datastore - can easily be eliminated doesn't fall under Relational DB

upvoted 1 times

✉ **abirroy** 1 year, 2 months ago

Selected Answer: B

B is correct

upvoted 1 times

✉ **theBestStudent** 1 year, 3 months ago

B as many have mentioned already

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

Global is keyword in this question, Cloud Spanner is the right option.. B

upvoted 5 times

✉ **haroldbenites** 1 year, 4 months ago

Go for B

upvoted 1 times

✉ **TheScalator** 1 year, 8 months ago

Selected Answer: B

B is correct

upvoted 2 times

✉ **k8sdr** 1 year, 9 months ago

Selected Answer: B

B is correct Cloud spanner can scale across regions

upvoted 2 times

✉ **look1** 1 year, 10 months ago

Selected Answer: B

B only

upvoted 1 times

You are the organization and billing administrator for your company. The engineering team has the Project Creator role on the organization. You do not want the engineering team to be able to link projects to the billing account. Only the finance team should be able to link a project to a billing account, but they should not be able to make any other changes to projects. What should you do?

- A. Assign the finance team only the Billing Account User role on the billing account.
- B. Assign the engineering team only the Billing Account User role on the billing account.
- C. Assign the finance team the Billing Account User role on the billing account and the Project Billing Manager role on the organization.
- D. Assign the engineering team the Billing Account User role on the billing account and the Project Billing Manager role on the organization.

Correct Answer: D

Community vote distribution

C (55%)

A (44%)

 **Bharathy**  3 years, 6 months ago

Option A is correct, as we don't want the engineering team to link projects to billing account and want only the Finance team. Billing Account User role will help to link projects to the billing account...

upvoted 80 times

 **naveedpk00** 3 years, 1 month ago

Option A is incorrect: Reason-

This role has very restricted permissions, so you can grant it broadly, typically in combination with Project Creator. These two roles allow a user to create new projects linked to the billing account on which the role is granted.

Reference: <https://cloud.google.com/billing/docs/how-to/billing-access>

I will go with option C.

upvoted 37 times

 **Josephsundarraj** 2 months, 1 week ago

Option C gives permission on org level where fin team can modify other projects billing. Question clearly says they should not be able to do that. So I think option A is good here in my opinion.

upvoted 1 times

 **willy_p** 1 year, 10 months ago

The question states that the user should ONLY link projects to billing accounts and nothing more. This is why I think A would be the best answer for this scenario.

upvoted 8 times

 **fishnoodlesoup** 1 year, 9 months ago

The question states that Finance department should ONLY be able to link projects to billing accounts.

If you look at the definition of Project Billing Creator:

Project Billing Manager

(roles/billing.projectManager) Link/unlink the project to/from a billing account.

It also gives permissions to unlink. Hence, A is correct.

upvoted 11 times

 **BobbyFlash** 1 year, 11 months ago

I would also go with A. I would think they are trying to get a quick answer from you as "Billing Administrator": engineering team already has the project creator role; you just would want finance team to link (and only) link projects to billing accounts, nothing else. Maybe the key phrase here is "but they should not be able to make any other changes to projects" and that would include the action of unlinking projects.

upvoted 7 times

 **mwoodm** 3 years, 1 month ago

Option A makes the most sense since Billing Account User can link projects to the billing account and the question reinforces principle of least privilege. Source: <https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 11 times

 **djgodzilla** 2 years, 4 months ago

wrong : you can't link project to billing accounts without Project billing manager.

C is Correct

upvoted 13 times

 **djgodzilla** 2 years, 4 months ago

my bad , it seems A,C are both technically correct

upvoted 1 times

 **measmo**  3 years, 4 months ago

for me is C:

https://cloud.google.com/billing/docs/how-to/modify-project#permissions_required_for_this_task_2

"Roles with adequate permissions to perform this task:

* Project Owner or Project Billing Manager on the project, AND Billing Account Administrator or Billing Account User for the target Cloud Billing account."

upvoted 55 times

 **fracila** 10 months, 3 weeks ago

We are assigning the finance team the Billing Account User role on the billing account, which allows them to create new projects linked to the billing account on which the role is granted. We are also assigning them the Project Billing Manager role on the organization (trickles down to the project as well) which lets them attach the project to the billing account, but does not grant any rights over resources.

upvoted 4 times

 **obeythefist** 1 year, 7 months ago

The question states that the finance group should not be able to make changes to existing projects. Granting the finance team organizational level Billing Account Administrator will allow them to make changes to other projects. C cannot be correct.

upvoted 3 times

 **Robertolo** 12 months ago

Project Billing Manager does not allow to make any changes to projects. It's just about linking+unlinking projects to billing accounts

On the other hand, the single role "billing account user" does not grant any right to view projects. Even less likely to link them to any billing account. (see <https://cloud.google.com/iam/docs/job-functions/billing> "The Billing Account User role gives the service account the permissions to enable billing (associate projects with the organization's billing account for all projects in the organization) and thereby permit the service account to enable APIs that require billing to be enabled."). Thus A is not the correct answer.

The right answer is C, without any kind of doubt

upvoted 3 times

 **Jake500** 5 months, 3 weeks ago

"Project Billing Manager does not allow to make any changes to projects. It's just about linking+unlinking projects to billing accounts"

Correct, but the problem states "... You do not want the engineering team to be able to link projects to the billing account." So in that case, wouldn't it be option A?

upvoted 1 times

 **[Removed]** 11 months, 1 week ago

Are you blind ? you posted link where its clearly stated in billing account user description: (associate projects with the organization's billing account for all projects in the organization)

So you literally posted link with clarification that answer A is correct.

answer C will give finance team additional permission to unlink billing account from projects and question clearly states that finance team should not be able to make any other changes to projects so C without any kind of doubt is wrong.

upvoted 4 times

 **izekc** 9 months, 1 week ago

Billing Account User

Principal: Service account that is used for automating project creation.

It is for service account, so C is correct

upvoted 1 times

 **ekta25** Most Recent 4 days, 21 hours ago

C. Assign the finance team the Billing Account User role on the billing account and the Project Billing Manager role on the organization.

upvoted 1 times

 **raxt** 4 days, 22 hours ago

Selected Answer: A

<https://cloud.google.com/billing/docs/how-to/billing-access#overview-of-cloud-billing-roles-in-cloud-iam>

upvoted 1 times

 **AdelElagawany** 1 week ago

This question is very tricky since there is a confusion in understanding the IAM roles for billing.

As a conclusion: Linking/Unlinking a resource (Project, GKE cluster, Compute instance,) to a "billing account" requires permissions on both the "Billing Account" and the "resource".

The "Billing Account User" has link/unlink permissions on the Billing account [1] but doesn't have permissions on the resources.

The "Project Billing Manager" has permissions on the project which is a resource [2] but doesn't have a permission on the billing account so both roles are required to link/unlink a project to a billing account.

Note: The only role that has permissions on "resources" and "billing accounts" is the Billing Account Administrator [3]

upvoted 1 times

 **AdelElagawany** 1 week ago

So The correct Answer is (C)

[1] Billing Account user role has "billing.resourceAssociations.create" but doesn't have "resourcemanager.projects.createBillingAssignment"

[2] Project billing manager has "resourcemanager.projects.createBillingAssignment" and

"resourcemanager.projects.deleteBillingAssignment" but doesn't have "billing.resourceAssociations.create"

[3] <https://cloud.google.com/billing/docs/how-to/billing-access#billing.admin>

upvoted 1 times

 **raxt** 4 days, 23 hours ago

The finance team should not be able to "unlink" a project. And this permission is granted when using Project Billing Manager. A looks like a better option.

upvoted 1 times

 **raxt** 4 days, 22 hours ago

The only role that can only link a project to a billing account is the Billing Account User.

upvoted 1 times

 **joao_01** 2 weeks, 6 days ago

Selected Answer: C

Its C guys. Look at the documentation:

"When granted in combination with the Billing Account User role, the Project Billing Manager role allows a user to attach the project to the billing account, but does not grant any rights over resources. Project Owners can use this role to allow someone else to manage the billing for the project without granting them resource access."

Link: <https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 1 times

 **Rishuwa** 3 weeks, 3 days ago

Selected Answer: C

<https://cloud.google.com/iam/docs/understanding-roles#billing.projectManager>

<https://cloud.google.com/iam/docs/understanding-roles#billing.user>

both are required

upvoted 1 times

 **RichBui** 3 weeks, 4 days ago

A. Assign the finance team only the Billing Account User role on the billing account.

The Billing Account User role allows users to link projects to billing accounts, but does not grant them permissions to manage projects in any other way. This ensures that the finance team can link a project to a billing account, but cannot make any other changes to projects. The engineering team, with the Project Creator role, can create projects but cannot link them to a billing account.

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: C

option c makes more logic

upvoted 1 times

 **YomanB** 1 month, 1 week ago

Option C grants unnecessary permissions:

Assigning the Project Billing Manager role on the organization gives them more control than needed, as they should only be managing billing for projects.

Correct Answer is : A

upvoted 1 times

 **Ahmed_Y** 1 month, 2 weeks ago

Selected Answer: C

it is C as you need the project billing manager to change the billing of a project.

upvoted 1 times

 **sthapit** 2 months ago

Option C provides the appropriate setup for achieving this

upvoted 2 times

 **Nxt_007** 2 months, 1 week ago

Selected Answer: C

The Billing Account User role allows users to link and unlink projects to a billing account, while the Project Billing Manager role allows users to view and manage the billing settings of projects within an organization.

-> Option A is not appropriate because the Billing Account User role does not grant permissions to manage billing settings for projects within the organization.

upvoted 1 times

 **rosh199** 2 months, 2 weeks ago

for me is C

upvoted 1 times

 **anujithn** 2 months, 2 weeks ago

Selected Answer: C

Obviously, its C

upvoted 1 times

 **ExamsFR** 2 months, 3 weeks ago

Selected Answer: C

Option C.

upvoted 1 times

 **jayjani66** 2 months, 3 weeks ago

Answer C

upvoted 1 times

You have an application running in Google Kubernetes Engine (GKE) with cluster autoscaling enabled. The application exposes a TCP endpoint. There are several replicas of this application. You have a Compute Engine instance in the same region, but in another Virtual Private Cloud (VPC), called gce-network, that has no overlapping IP ranges with the first VPC. This instance needs to connect to the application on GKE. You want to minimize effort. What should you do?

- A. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Set the service's externalTrafficPolicy to Cluster. 3. Configure the Compute Engine instance to use the address of the load balancer that has been created.
- B. 1. In GKE, create a Service of type NodePort that uses the application's Pods as backend. 2. Create a Compute Engine instance called proxy with 2 network interfaces, one in each VPC. 3. Use iptables on this instance to forward traffic from gce-network to the GKE nodes. 4. Configure the Compute Engine instance to use the address of proxy in gce-network as endpoint.
- C. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add an annotation to this service: cloud.google.com/load-balancer-type: Internal 3. Peer the two VPCs together. 4. Configure the Compute Engine instance to use the address of the load balancer that has been created.
- D. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add a Cloud Armor Security Policy to the load balancer that whitelists the internal IPs of the MIG's instances. 3. Configure the Compute Engine instance to use the address of the load balancer that has been created.

Correct Answer: A

Community vote distribution

C (54%)

A (46%)

✉  **someoneinthecloud**  3 years, 2 months ago

I believe it's A. It's never mentioned in the question that traffic cannot go through the Internet but it's mentioned that effort should be minimized. A requires a lot less effort than C to accomplish the same (no VPC peering, per example).

upvoted 54 times

✉  **ShakthiGCP** 2 years, 7 months ago

Ans: A . This sounds correct and avoids unnecessary steps in C. C is also correct but compared to it, A is much easier to achieve. Go over Kubernetes Loadbalancer concepts to get more details. Initially i was thinking C is the Answer. but after putting some time on K8's Network - changed my mind to A.

upvoted 13 times

✉  **AmitKM** 3 years, 1 month ago

Yeah, I feel the same. Nowhere does it say that the traffic has to be internal. But it does say "minimal effort" which I feel is option A.

upvoted 10 times

✉  **pgb54** 1 year, 7 months ago

Totally agree. I had the same thought and looked through the question for any indication that the traffic must be private.

upvoted 2 times

✉  **juancambb**  3 years, 4 months ago

i think C is better solution, the solution A pass trafic trought public internet, also C by internal network and the "no overlap ips" in the statament suggest that.

upvoted 44 times

✉  **ekta25**  4 days, 20 hours ago

C. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add an annotation to this service: cloud.google.com/load-balancer-type: Internal 3. Peer the two VPCs together. 4. Configure the Compute Engine instance to use the address of the load balancer that has been created.

upvoted 1 times

✉  **SinghAnc** 6 days, 19 hours ago

Selected Answer: C

Correct Answer is C

Option A suggests setting the service's externalTrafficPolicy to Cluster. While this is a valid configuration, it's not directly related to the scenario described.

In the given scenario, the goal is to connect a Compute Engine instance from a different VPC to the application running in GKE. This involves networking configurations, peering the VPCs, and potentially setting up a LoadBalancer.

Setting the externalTrafficPolicy to Cluster primarily affects how traffic is balanced across Pods within the cluster, but it doesn't directly address the requirement of connecting an external instance from a different VPC.

upvoted 1 times

✉  **RobAlt** 1 month, 1 week ago

Selected Answer: A

Minimal effort is the point.

upvoted 1 times

ExamsFR 2 months, 3 weeks ago

Selected Answer: C

Option C.

upvoted 1 times

KerolesKhil 4 months, 1 week ago

Selected Answer: C

A is not correct

Because the GKE cluster and the instance are not in the same vpc , so without vpc peering traffic can't be established .

C is the correct answer.

Traffic still internal not exposed to internet , as they mentioned creating internal tcp loadbalancer not public one and created vpc peering . so no additional steps are needed.

upvoted 2 times

tempdir 3 months, 4 weeks ago

C is also correct but as the question states, minima effort. In A you dont need vpc peering since you will be using loadbalancer to expose the application, therefore, traffic can still be established b/n the two.

upvoted 1 times

sana_sree 4 months, 1 week ago

Selected Answer: C

Correct is C

please refer

<https://www.youtube.com/watch?v=qx8PEmxKYzg>

upvoted 3 times

DrLegendgun 6 months, 1 week ago

Selected Answer: C

The answer is C as two VPC needed to Peer first

upvoted 1 times

esqandares 6 months, 1 week ago

Selected Answer: A

same region, but in another Virtual Private Cloud (VPC), called gce-network, that has no overlapping IP ranges with the first VPC.... need to read question again and again

upvoted 1 times

raselsys 7 months ago

peering is lot easier effort but with dependent on not having overlapping IPs and that was clearly stated on the question. So C without any doubt is the correct answer here IMO.

upvoted 1 times

Buruguduystunstugudunstuy 7 months, 3 weeks ago

Selected Answer: A

Answer A is the correct solution.

In Answer A, we can create a Service of the type LoadBalancer in GKE that uses the application's Pods as a backend. This will create a Google Cloud load balancer with an external IP address that can be used to connect to the application. We can set the service's externalTrafficPolicy to Cluster to ensure that traffic is routed only to the nodes running the application. Then we can configure the Compute Engine instance to use the address of the load balancer that has been created.

upvoted 4 times

Shivangi30 3 months, 1 week ago

externalTrafficPolicy is supported for internal LoadBalancer Services (via the TCP/UDP load balancer), but load balancing behavior depends on where traffic originates from and the configured traffic policy. Hence Answer is C as per link:

<https://cloud.google.com/kubernetes-engine/docs/how-to/service-parameters#externalTrafficPolicy>

upvoted 2 times

antivrillee 6 months, 3 weeks ago

Option C requires less effort compared to option A.

In option A, you need to set the service's externalTrafficPolicy to Cluster, which means that the traffic will be load balanced across all nodes in the cluster, including those outside of the VPC network. You will also need to configure the Compute Engine instance to use the address of the load balancer that has been created.

In option C, you only need to add an annotation to the service with the value of "Internal", which will create an internal load balancer that is only accessible from within the VPC network. You will also need to peer the two VPCs together and configure the Compute Engine instance to use the address of the load balancer that has been created.

Therefore, option C requires less effort as it involves fewer steps and less configuration.

upvoted 2 times

Buruguduystunstugudunstuy 7 months, 3 weeks ago

Answer B is not recommended because it requires the creation of an additional instance called a proxy, and the use of iptables to forward traffic from gce-network to the GKE nodes. This solution introduces additional complexity and potential points of failure.

configuration.

Answer D is not recommended because it involves using Cloud Armor to whitelist the internal IPs of the MIG's instances. This solution introduces additional complexity and potential security risks.

Therefore, Answer A is the most straightforward and least complex solution to connect the Compute Engine instance to the application running on GKE.

upvoted 5 times

 **Bobbybash** 7 months, 4 weeks ago

Selected Answer: A

A is correct

Option A is the best solution to minimize effort. In GKE, creating a Service of type LoadBalancer that uses the application's Pods as backend and setting the service's externalTrafficPolicy to Cluster will expose the TCP endpoint of the application with a public IP address. Then, configuring the Compute Engine instance to use the address of the load balancer that has been created will allow it to connect to the application on GKE. Option B requires creating a separate instance as a proxy and using iptables to forward traffic, which adds unnecessary complexity. Option C involves peering the two VPCs together, which may not be desirable or feasible in all cases. Option D adds additional complexity by adding a Cloud Armor Security Policy to the load balancer.

upvoted 3 times

 **David_C_90** 8 months, 1 week ago

Selected Answer: C

Option A exposes an unsecured TCP endpoint on the internet, and there is no mention of the VM IP ranges with the first VPC.

upvoted 2 times

 **kajitsu** 8 months, 2 weeks ago

Selected Answer: A

A answer

upvoted 1 times

 **GS300** 8 months, 3 weeks ago

Selected Answer: A

A has minimal effort

upvoted 1 times

 **abmohamed** 8 months, 3 weeks ago

Selected Answer: C

<https://cloud.google.com/kubernetes-engine/docs/how-to/service-parameters#externalTrafficPolicy>

The externalTrafficPolicy is a standard Service option that defines how and whether traffic incoming to a GKE node is load balanced. Cluster is the default policy, but Local is often used to preserve the source IP of traffic coming into a cluster node. Local effectively disables load balancing on the cluster node so that traffic that is received by a local Pod sees the original source IP address.

There is no need to set an already default value.

upvoted 2 times

Your organization is a financial company that needs to store audit log files for 3 years. Your organization has hundreds of Google Cloud projects. You need to implement a cost-effective approach for log file retention. What should you do?

- A. Create an export to the sink that saves logs from Cloud Audit to BigQuery.
- B. Create an export to the sink that saves logs from Cloud Audit to a Coldline Storage bucket.
- C. Write a custom script that uses logging API to copy the logs from Stackdriver logs to BigQuery.
- D. Export these logs to Cloud Pub/Sub and write a Cloud Dataflow pipeline to store logs to Cloud SQL.

Correct Answer: A

Reference:

<https://cloud.google.com/logging/docs/audit/>

Community vote distribution

 B (95%) 5%

 **yasu**  3 years, 6 months ago

Why not B? cost effective
upvoted 54 times

 **uganeshku** 1 year, 9 months ago

B is correct because Coldline Storage is the perfect service to store audit logs from all the projects and is very cost-efficient as well. Coldline Storage is a very low-cost, highly durable storage service for storing infrequently accessed data.
upvoted 6 times

 **_Sande** 6 months, 1 week ago

That seems to be the one...
upvoted 1 times

 **Ixs** 1 year, 11 months ago

BigQuery data after 90 days has the same cost for storage as Cloud Storage Nearline. Storing it in Cloud Storage adds more costs for data retrieval if the class is i.e archival
upvoted 1 times

 **KerolesKhalil** 4 months, 1 week ago

the options have cold-line storage not nearline.
so B is the cheapest option.
upvoted 1 times

 **Gini**  3 years, 6 months ago

if it is all about cost, B is the best. However, speaking of "audit" you probably need to access the data once in a while, which Coldline storage might not be ideal for this case I guess? I would go for A in the exam though.
upvoted 20 times

 **pas77** 2 years, 2 months ago

The question is clearly saying cost effect. BQ is one of the most expensive services in GCP.
upvoted 6 times

 **Ale1973** 3 years, 1 month ago

Be strong!!! If B is the best, go for B!!!
upvoted 12 times

 **boof** 2 years ago

I would play it safe and interpret the question literally, implying that they will only store the audit logs and not be accessing them a lot.
upvoted 2 times

 **elviskimutai**  2 weeks, 4 days ago

B is correct. coldline storage it is cost-effective and for long-term storage
upvoted 1 times

 **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: B
answer B
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B
Answer B. Create an export to the sink that saves logs from Cloud Audit to a Coldline Storage bucket.

Storage Coldline storage class. The coldline storage class is designed for cold data storage and offers lower storage costs and higher retrieval costs when compared to other storage classes.

Exporting logs from Cloud Audit to a Coldline Storage bucket can be done by creating an export sink. This is a straightforward process that can be done via the Cloud Console, Cloud SDK, or REST API.

upvoted 3 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

INCORRECT:

Answer A is incorrect because while it does store logs in BigQuery, it is a more expensive option than storing logs in Coldline storage.

Answer C is incorrect because writing a custom script to copy logs to BigQuery would be complex and more difficult to maintain compared to using an export sink.

Answer D is incorrect because while logs can be exported to Pub/Sub, writing a Cloud Dataflow pipeline to store logs in Cloud SQL would require additional configuration and might not be as cost-effective as exporting logs to Coldline storage.

upvoted 6 times

 **Mac_1612** 1 month ago

Buru, Thank you so much from whole GCP community for your efforts, you are incredible man !!!

upvoted 3 times

 **cucinareblog** 3 weeks, 5 days ago

I agree thanks a lot!!!!

upvoted 1 times

 **kaisehop** 8 months, 1 week ago

Selected Answer: B

B. Create an export to the sink that saves logs from Cloud Audit to a Coldline Storage bucket. This is the most cost-effective approach for log file retention as Coldline Storage has lower storage costs compared to other storage classes, making it suitable for infrequently accessed data that still needs to be retained for a long time.

upvoted 1 times

 **alex000** 9 months, 1 week ago

B is the most cost effective

upvoted 1 times

 **Neeleshinuk** 9 months, 2 weeks ago

https://cloud.google.com/logging/docs/export/configure_export_v2

I go with B

upvoted 1 times

 **cslince** 10 months, 1 week ago

Selected Answer: B

B is the best

upvoted 1 times

 **Zoze** 10 months, 3 weeks ago

Selected Answer: B

B is the answer, because he wants a cost-effective solution so B is the cheapest option.

upvoted 1 times

 **leogor** 11 months, 2 weeks ago

Selected Answer: B

Coldline Storage bucket

upvoted 1 times

 **dennydream** 11 months, 2 weeks ago

Selected Answer: B

Coldline would be the least expensive.

upvoted 1 times

 **sandipk91** 1 year, 1 month ago

Selected Answer: B

Option B because it talks about cost effective solution, I know BQ has the same cost as Coldline in GCS if data is kept for 90 days but in Cloud Storage we can save more by further moving the class to Archival which is cheaper than Coldline. SO DEFINATELY IT'S OPTION B

upvoted 2 times

 **joeMP** 1 year, 1 month ago

Selected Answer: A

My bad. A is good for bigquery can supports store many data

upvoted 1 times

 **joeMP** 1 year, 1 month ago

Selected Answer: B

Hundreds of projects means many logs. Bigquery is the good Storage.

upvoted 1 times

 **abirroy** 1 year, 2 months ago

Selected Answer: B

Answer is B

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Answer is B

upvoted 1 times

You want to run a single caching HTTP reverse proxy on GCP for a latency-sensitive website. This specific reverse proxy consumes almost no CPU. You want to have a 30-GB in-memory cache, and need an additional 2 GB of memory for the rest of the processes. You want to minimize cost. How should you run this reverse proxy?

- A. Create a Cloud Memorystore for Redis instance with 32-GB capacity.
- B. Run it on Compute Engine, and choose a custom instance type with 6 vCPUs and 32 GB of memory.
- C. Package it in a container image, and run it on Kubernetes Engine, using n1-standard-32 instances as nodes.
- D. Run it on Compute Engine, choose the instance type n1-standard-1, and add an SSD persistent disk of 32 GB.

Correct Answer: B

Community vote distribution

A (76%) B (24%)

✉  **jzh**  3 years, 1 month ago

Go to cloud console and create instance
select Memorystore with Basic tier, select us-central1 and us-central1-a, and capacity 32GB, the cost estimate is \$0.023/GB/hr

select VM instance with custom machine type with 6 vCPUs and 32 GB memory, the same region and zone as Memorystore setting, the cost estimate is \$0.239/hr

Option B will definitely cost more as it adds on CPU usage cost even it uses little in this scenario, but still charge you. So answer is A from real practice example.

upvoted 50 times

✉  **mexblood1** 3 years ago

Using pricing calculator matching 730 hrs per month for both.. Memorystore is 537.28 per month and vm (6 cpus 32 gb memory) is 174.41 per month. So vm is still cheaper even with 6 cpus.

upvoted 7 times

✉  **[Removed]** 3 years ago

I agree its cheaper, but 2 drawbacks, 1 hit latency, 2 you need to install cache plain won't help even if check. So still think A.
upvoted 3 times

✉  **[Removed]** 3 years ago

Typo correct my statements. There are 2 issues If you create a VM how every cheap there will be a hit with latency when communicating with it. Second, you will need to install/implement a caching system on that.

Whereas cloud memorystore for redis is designed for the sole purpose of HTTP caching which has very low latency compared to any other solution we are thinking of doing.

upvoted 3 times

✉  **FenixRa73** 2 years, 9 months ago

\$0.023 * 32 = \$0.736

is it cheaper?

upvoted 4 times

✉  **SSPC** 3 years, 1 month ago

I agree with you

upvoted 1 times

✉  **Rothmansua** 1 year, 11 months ago

and what about HTTP, how are you supporting that with Redis?

upvoted 2 times

✉  **obeythefist** 1 year, 7 months ago

A quick Bing search shows a number of solutions for caching HTTP services with Redis.

upvoted 2 times

✉  **smarty_arce** 1 year, 3 months ago

Who uses Bing at this present day and age?

upvoted 11 times

✉  **RNSS** 10 months, 3 weeks ago

believe me it is very good and clean. When I was doing my research I have used both google and bing. and find bing as more trusted and complete answer.

upvoted 1 times

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer should be A:

The question mention "You want to have a 30-GB in-memory cache, and need an additional 2 GB of memory for the rest of the processes"

What is Google Cloud Memorystore?

Overview. Cloud Memorystore for Redis is a fully managed Redis service for Google Cloud Platform. Applications running on Google Cloud Platform can achieve extreme performance by leveraging the highly scalable, highly available, and secure Redis service without the burden of managing complex Redis deployments.

upvoted 35 times

✉️  **ESP_SAP** 3 years, 1 month ago

Just to complement the answer:
We are looking for "latency-sensitive website"

What it's good for

Memorystore for Redis provides a fast, in-memory store for use cases that require fast, real-time processing of data. From simple caching use cases to real time analytics, Memorystore for Redis provides the performance you need.

Caching: Cache is an integral part of modern application architectures. Memorystore for Redis provides low latency access and high throughput for heavily accessed data, compared to accessing the data from a disk based backend store. Session management, frequently accessed queries, scripts, and pages are common examples of caching.

https://cloud.google.com/memorystore/docs/redis/redis-overview#what_its_good_for

upvoted 18 times

✉️  **[Removed]** 3 years, 1 month ago

I agree with your reasoning. Given that the question stresses that this is for a 'latency sensitive website', that's a clue that Redis is part of the answer. Even if spinning up a similarly sized VM were more cost effective, I can't find any documentation that this would provide sufficiently low latency as a memory cache. Yes, you want to keep costs low, but not if it causes your latency-sensitive website problems. Thus I agree that option A is the answer.

upvoted 2 times

✉️  **magistrum** 2 years, 9 months ago

Agree, don't think anything you create with the Compute Engine will meet the "You want to have a 30-GB in-memory cache" requirement...that's a very different technology

upvoted 2 times

✉️  **ekta25** [Most Recent] 4 days, 20 hours ago

B. Run it on Compute Engine, and choose a custom instance type with 6 vCPUs and 32 GB of memory.

upvoted 1 times

✉️  **Captain1212** 1 month, 1 week ago

Selected Answer: A

A is the correct as redis for low latency

upvoted 1 times

✉️  **marcus021** 4 months ago

Selected Answer: A

Low latency should be A.

upvoted 2 times

✉️  **sana_sree** 4 months, 1 week ago

correct answer is A

<https://www.youtube.com/watch?v=a1p1pB375Ik>

upvoted 1 times

✉️  **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A

ANSWER A is the most cost-effective solution for running a caching HTTP reverse proxy on GCP. Cloud Memorystore for Redis is a managed service that provides an in-memory cache for your applications. It offers a high throughput and low latency access to the Redis protocol. Cloud Memorystore offers an SLA of 99.9% availability and automatic failover for Redis instances. In this case, a 32-GB Redis instance is sufficient to accommodate the 30-GB cache and the additional 2 GB of memory required for the rest of the processes. This solution is highly scalable and allows you to increase the size of the Redis instance as your needs grow.

upvoted 4 times

✉️  **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

INCORRECT:

ANSWER B is not a cost-effective solution since it requires a custom instance type with 6 vCPUs and 32 GB of memory, which is over-provisioned for a caching HTTP reverse proxy.

ANSWER C is also not a cost-effective solution since it uses Kubernetes Engine, which has a higher management overhead and may not be necessary for a single caching HTTP reverse proxy. Additionally, using n1-standard-32 instances as nodes is over-provisioned for the requirements of the caching HTTP reverse proxy.

ANSWER D is not a viable solution since the instance type n1-standard-1 only provides 3.75 GB of memory, which is insufficient for the 30-GB cache and the additional 2 GB of memory required for the rest of the processes. Adding an SSD persistent disk of 32 GB will not provide enough memory for the reverse proxy.

upvoted 4 times

✉️  **kaisehhon** 8 months, 1 week ago

Selected Answer: A

A. Create a Cloud Memorystore for Redis instance with 32-GB capacity is the recommended option. This option provides the required memory and is cost-effective since the proxy requires almost no CPU. Cloud Memorystore for Redis is designed specifically for in-memory caching, making it the best choice for your use case.

upvoted 1 times

 **cslince** 10 months, 1 week ago

Selected Answer: A

answer is A

upvoted 1 times

 **fragment137** 10 months, 2 weeks ago

Selected Answer: B

While Redis is definitely the easiest and best solution for a latency sensitive workload, the question is worded in such a way to emphasize the requirement of cost. "You want to have a 30-GB in-memory cache, and need an additional 2 GB of memory for the rest of the processes. You want to minimize cost". Given this, the answer has to be B, even if that's not the best technical solution for the problem.

upvoted 1 times

 **Zoze** 10 months, 3 weeks ago

Selected Answer: A

A is correct, he only wants to have memory capacity, and doesn't care about CPU at all. In addition Memory-store is already configured to use is a cache memory.

upvoted 1 times

 **dennydream** 11 months, 2 weeks ago

How do you figure the correct answer here? The votes overwhelmingly say one thing, but the correct answer is another.

upvoted 1 times

 **PKookNN** 11 months, 3 weeks ago

Selected Answer: B

I changed my mind - cost effective is B (while A is easiest)

upvoted 1 times

 **ashit44244** 1 year ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Selected Answer: A

Answer is A

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right ..Memorystore for Redis provides a fast.

upvoted 1 times

 **kohsiangyu** 1 year, 3 months ago

Selected Answer: B

MemoryStore Pricing is \$0.023/GB/hr, for 32GB means \$0.736/hr compares to \$0.239/hr. The question states that it needs additional 2GB for its process, which mean if you choose A, you will need another vm with 2GB ram either.

upvoted 3 times

You are hosting an application on bare-metal servers in your own data center. The application needs access to Cloud Storage. However, security policies prevent the servers hosting the application from having public IP addresses or access to the internet. You want to follow Google-recommended practices to provide the application with access to Cloud Storage. What should you do?

- A. 1. Use nslookup to get the IP address for storage.googleapis.com. 2. Negotiate with the security team to be able to give a public IP address to the servers. 3. Only allow egress traffic from those servers to the IP addresses for storage.googleapis.com.
- B. 1. Using Cloud VPN, create a VPN tunnel to a Virtual Private Cloud (VPC) in Google Cloud. 2. In this VPC, create a Compute Engine instance and install the Squid proxy server on this instance. 3. Configure your servers to use that instance as a proxy to access Cloud Storage.
- C. 1. Use Migrate for Compute Engine (formerly known as Velostrata) to migrate those servers to Compute Engine. 2. Create an internal load balancer (ILB) that uses storage.googleapis.com as backend. 3. Configure your new instances to use this ILB as proxy.
- D. 1. Using Cloud VPN or Interconnect, create a tunnel to a VPC in Google Cloud. 2. Use Cloud Router to create a custom route advertisement for 199.36.153.4/30. Announce that network to your on-premises network through the VPN tunnel. 3. In your on-premises network, configure your DNS server to resolve *.googleapis.com as a CNAME to restricted.googleapis.com.

Correct Answer: C

Community vote distribution

D (100%)

 **poogcp**  3 years, 4 months ago

D is the correct one as per Ref: <https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid>
upvoted 53 times

 **obeythefist**  1 year, 7 months ago

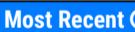
What messy answers! I chose D and here is my reasoning per answer.

- A. It's bad practice to use nslookup to try find a permanent IP address because IPs can change. That's what DNS is for! Also, the security team aren't going to budge... this is just a silly answer.
 - B. We're getting warmer. Any time a question mentions on-prem and cloud, Google wants you to think about Cloud VPN. This solution might even work, but installing Squid? This is a messy solution to a more simple problem.
 - C. Talk about using a sledge hammer to swat a mosquito. I think this could work, but migrating servers to cloud to solve a simple networking problem?
 - D. Once more Google's favorite Cloud VPN is in the answer. I'm not sure about the networking component of this question.
- upvoted 16 times

 **obeythefist** 1 year, 7 months ago

Edit: Of course the reason D: is correct is because 199.36.153.4/30 is the network segment that you can direct traffic to if you want to use Google services "internally". So your on prem servers will resolve storage.googleapis.com to something in this 199.36.153.4/30 range. Then they will route using Cloud Router and your VPN tunnel into Google Cloud privately.

upvoted 9 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: D

D as per google recommended practices

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: D

ANSWER D is the recommended solution because it provides a secure and direct connection to Cloud Storage without requiring internet access or exposing the servers to public IP addresses.

* By setting up a VPN or Interconnect tunnel, the on-premises servers can access Google Cloud resources over a private and encrypted connection.

* The custom route advertisement for 199.36.153.4/30 ensures that traffic is routed correctly between the on-premises network and Google Cloud.

* Configuring the DNS server to resolve *.googleapis.com as a CNAME to restricted.googleapis.com ensures that requests are directed to Google Cloud over the VPN or Interconnect tunnel.

upvoted 6 times

 **warrior9000** 9 months, 1 week ago

D but anyone wanna try to explain how the hell you can have a VPN connection without accessing the public internet? The only option for D should be using Interconnect for a direct private wire from your data center to GCP. VPN doesn't make any sense.

upvoted 5 times

 **ast3citos** 7 months ago

The machine hosting the application cannot access directly the public internet. So to go to Google Cloud it must go through a VPN.

upvoted 1 times

✉ **xaqanik** 9 months ago

<https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid>

upvoted 2 times

✉ **cslince** 10 months, 1 week ago

Selected Answer: D

D is the correct

upvoted 1 times

✉ **gcpBeginner** 1 year ago

correct answer is D. why will cx migrate it env. to GCP. easiest and faster approach is to have Cloud VPN setup and advertise route o cloud router

upvoted 1 times

✉ **gcpBeginner** 1 year ago

in my above comment I meant advertise routes to local router used on-prem

upvoted 1 times

✉ **Cornholio_LMC** 1 year ago

had this question today

upvoted 3 times

✉ **RAVI321** 1 year, 1 month ago

how you all know that everybody

mere to upar se nikal raha hai sab kuch pls help guys

upvoted 5 times

✉ **Souvik_79** 11 months, 3 weeks ago

Me too bro

upvoted 1 times

✉ **aakash4Chaudhary** 1 year ago

tu akela nhi hai bro

upvoted 2 times

✉ **abirroy** 1 year, 2 months ago

Selected Answer: D

Answer is D

upvoted 1 times

✉ **RanjithK** 1 year, 3 months ago

Selected Answer: D

Answer is D

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

D is most appropriate, This is part of Tutorial Dojo practice questions.

upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago

Go for D

upvoted 1 times

✉ **luciorifa** 1 year, 7 months ago

Selected Answer: D

D is the correct answer

upvoted 2 times

✉ **Durgesh1997** 1 year, 8 months ago

Selected Answer: D

D is practically good to implement so D is the answer

upvoted 2 times

✉ **[Removed]** 1 year, 9 months ago

D makes most sense

upvoted 3 times

✉ **hellorf20** 1 year, 9 months ago

Selected Answer: D

D is Correct

upvoted 3 times

You want to deploy an application on Cloud Run that processes messages from a Cloud Pub/Sub topic. You want to follow Google-recommended practices. What should you do?

- A. 1. Create a Cloud Function that uses a Cloud Pub/Sub trigger on that topic. 2. Call your application on Cloud Run from the Cloud Function for every message.
- B. 1. Grant the Pub/Sub Subscriber role to the service account used by Cloud Run. 2. Create a Cloud Pub/Sub subscription for that topic. 3. Make your application pull messages from that subscription.
- C. 1. Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint.
- D. 1. Deploy your application on Cloud Run on GKE with the connectivity set to Internal. 2. Create a Cloud Pub/Sub subscription for that topic. 3. In the same Google Kubernetes Engine cluster as your application, deploy a container that takes the messages and sends them to your application.

Correct Answer: D

Community vote distribution

C (72%)

B (28%)

✉  **Meix**  3 years, 4 months ago

C looks right for me as per <https://cloud.google.com/run/docs/tutorials/pubsub#integrating-pubsub>
upvoted 51 times

✉  **ChrisBelt5** 2 years, 2 months ago

great doc, its' C
upvoted 6 times

✉  **Bhagirathi**  2 years, 9 months ago

why c ?
explained>>
You can use Pub/Sub to push messages to the endpoint of your Cloud Run service, where the messages are subsequently delivered to containers as HTTP requests. You cannot use Pub/Sub pull subscriptions because Cloud Run only allocates CPU during the processing of a request.
upvoted 36 times

✉  **hanweiCN**  5 months, 1 week ago

Selected Answer: C

it is explicated recommended use " push":
Note: Google recommends using push subscriptions to consume messages from a Pub/Sub topic on Cloud Run. Although it is possible to use Pub/Sub pull subscriptions, pull subscriptions require you to monitor message delivery latency and manually scale the number of instances to maintain a healthy delivery latency. If you want to use pull subscriptions, use the CPU always allocated setting along with a number of minimum instances.
<https://cloud.google.com/run/docs/triggering/pubsub-push>
upvoted 1 times

✉  **SMR123** 5 months, 3 weeks ago

what is the actual answer people who are voting or actual answer?
upvoted 1 times

✉  **dobberzoon** 6 months, 1 week ago

Selected Answer: C

ExamPro also explains it's C.
upvoted 2 times

✉  **due** 7 months ago

Selected Answer: B

"processes messages from a Cloud Pub/Sub topic"
should be pull subscription
i will go with B
upvoted 2 times

✉  **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B

Answer B, I would say is the correct answer for me.

When you want to deploy an application on Cloud Run that processes messages from a Cloud Pub/Sub topic, you should follow Google-recommended practices to ensure that your application can securely and reliably process the messages.

Answer A is not the correct answer because using a Cloud Function to call your application on Cloud Run for every message adds additional complexity and potential points of failure.

upvoted 3 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Answer C is not the correct answer ALTHOUGH it is possible to create a subscription that uses Cloud Run as the push endpoint, it requires an additional setup that is not necessary for this use case. Additionally, this approach requires your Cloud Run service to be publicly accessible.

Answer D is not the correct answer because it requires deploying a container in a Kubernetes cluster to handle the Pub/Sub messages and forward them to Cloud Run. This adds additional complexity to the architecture, and Cloud Run can directly subscribe to a Pub/Sub topic without needing a proxy service in between.

Therefore, the recommended approach is to grant the Pub/Sub Subscriber role to the service account used by Cloud Run and create a Cloud Pub/Sub subscription for that topic. This approach is straightforward and aligns with Google's best practices for integrating Pub/Sub with Cloud Run.

upvoted 1 times

 **space_cadet** 7 months, 2 weeks ago

The advantage of Cloud Run scaling to zero might be a better practice like Shwom suggested.
I also thought B cus of it's straightforward approach, but C seems better now.

upvoted 1 times

 **Shwom** 8 months, 1 week ago

Selected Answer: C

C is correct

upvoted 1 times

 **Anirudha077** 10 months ago

C is the right answer

upvoted 1 times

 **Pho_King** 10 months, 1 week ago

Ans is C (push subscription).

I think B (pull subscription) can work, but it does not follow best practices. With pull subscription, you cannot take advantage of scale-to-zero for Cloud Run. Why is this the case? Because you need to keep your subscriber client (Cloud Run in option B) running, otherwise you don't know if there are new messages. (Whereas if you use push subscription instead, you can scale to zero and rest assured that Cloud Run will be "notified" of new messages).

See 3:15 and 4:10 of this video: <https://www.youtube.com/watch?v=KObjkda4ZfY>

Remember this:

Push subscription -> Pub/Sub server initiates a request to your subscriber client to deliver messages

Pull subscription -> Your subscriber client initiates requests to a Pub/Sub server to retrieve messages

Documentation: <https://cloud.google.com/pubsub/docs/subscriber>

upvoted 5 times

 **axantroff** 1 week, 2 days ago

Explained well enough. Recommended

upvoted 1 times

 **Pho_King** 10 months, 1 week ago

Ans is C (push subscription).

I think B (pull subscription) can work, but it does not follow best practices. You cannot take advantage of scale-to-zero for Cloud Run. See 3:15 and 4:10 of this video: <https://www.youtube.com/watch?v=KObjkda4ZfY>

upvoted 1 times

 **leogor** 11 months, 2 weeks ago

C obviously

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Selected Answer: C

Answer is C

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

C is right .. Look for Meix explanation...

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 1 times

 **gielda211** 1 year, 6 months ago

Selected Answer: C

C. is the correct answer:

Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint

upvoted 3 times

 **gielda211** 1 year, 6 months ago

C. is the correct answer:

Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint

upvoted 1 times

You need to deploy an application, which is packaged in a container image, in a new project. The application exposes an HTTP endpoint and receives very few requests per day. You want to minimize costs. What should you do?

- A. Deploy the container on Cloud Run.
- B. Deploy the container on Cloud Run on GKE.
- C. Deploy the container on App Engine Flexible.
- D. Deploy the container on GKE with cluster autoscaling and horizontal pod autoscaling enabled.

Correct Answer: B*Community vote distribution*A (100%)

 **Gurnoor**  3 years, 3 months ago

A should be cheapest as no infra needed.
upvoted 65 times

 **spudleymcdudley** 3 years, 3 months ago

Listen to this guy. Google says "Cloud Run abstracts away all infrastructure management by automatically scaling up and down from zero almost instantaneously—depending on traffic. Cloud Run only charges you for the exact resources you use."
upvoted 28 times

 **ESP_SAP**  3 years, 1 month ago

Correct Answer should be A:

Cloud Run takes any container images and pairs great with the container ecosystem: Cloud Build, Artifact Registry, Docker. ... No infrastructure to manage: once deployed, Cloud Run manages your services so you can sleep well. Fast autoscaling. Cloud Run automatically scales up or down from zero to N depending on traffic.

<https://cloud.google.com/run>
upvoted 26 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: A
A, as it does not include the infra services and its cheaper
upvoted 1 times

 **sthapit** 2 months ago

Should be A
upvoted 1 times

 **yash_901** 4 months, 2 weeks ago

Can anyone confirm these questions are the exact questions that would come in the exam? I can see there are 201 questions mentioned but I got only 197. Also, with each question, its mentioned topic 1. Does that mean there are more topics for ACE exam?
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A
Definitely the correct answer is A, which is deploying the container on Cloud Run, is the most cost-effective option to deploy a container image that exposes an HTTP endpoint and receives very few requests per day.

Cloud Run is a fully managed serverless platform that allows you to run stateless containers and automatically scales up or down based on incoming requests.

With Cloud Run, you only pay for the actual usage, so if your application receives very few requests per day, your costs will be minimal.
upvoted 5 times

 **cslince** 10 months, 1 week ago

Selected Answer: A
Correct Answer should be A
upvoted 1 times

 **Zoze** 10 months, 4 weeks ago

Selected Answer: A
A is correct, it's scaled automatically and service-less and more cheaper than GKE.
upvoted 2 times

 **leogor** 11 months ago

Selected Answer: A

Serverless option: Cloud Run is the cheapest
upvoted 1 times

✉ **Cornholio_LMC** 1 year ago
had this question today
upvoted 7 times

✉ **xharf** 1 year, 2 months ago
Don't forget that cloud run invoking rules
"Cloud Run redirects all HTTP requests to HTTPS but terminates TLS before they reach your web service."
in last case when i use cloud run with POST request and HTTP method, it will generate bug which redirect the request to GET https header :)
<https://cloud.google.com/run/docs/triggering/https-request#>
upvoted 1 times

✉ **orous** 1 year, 2 months ago

Selected Answer: A
What does It mean, "Cloud Run on GKE"?
B is not correct at all, there is no such thing – "Cloud Run on GKE". Or you deploy on "Cloud Run" Or on "GKE".
upvoted 4 times

✉ **glanshima** 8 months, 2 weeks ago
Possible to deploy cloud run on GCP. Thee docs says so.
upvoted 1 times

✉ **David_C_90** 9 months ago
It is possible to use Cloud Run on Anthos (kubernetes): <https://cloud.google.com/anthos/run>
upvoted 1 times

✉ **RanjithK** 1 year, 3 months ago

Selected Answer: A
Answer is A. CloudRun is the cheapest option of all
upvoted 3 times

✉ **AzureDP900** 1 year, 3 months ago
This is no brainier question.. If you know the GCP services well straight forward answer is Cloud Run to reduce the cost.. Option A is right!
upvoted 1 times

✉ **Tirthankar17** 1 year, 4 months ago

Selected Answer: A
A is correct as In B you will be paying for underlying infra
upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago
Go for A
upvoted 1 times

✉ **pfabio** 1 year, 4 months ago

Selected Answer: A
Cloud Run takes any container images and pairs great with the container ecosystem: Cloud Build, Artifact Registry, Docker.
upvoted 2 times

Your company has an existing GCP organization with hundreds of projects and a billing account. Your company recently acquired another company that also has hundreds of projects and its own billing account. You would like to consolidate all GCP costs of both GCP organizations onto a single invoice. You would like to consolidate all costs as of tomorrow. What should you do?

- A. Link the acquired company's projects to your company's billing account.
- B. Configure the acquired company's billing account and your company's billing account to export the billing data into the same BigQuery dataset.
- C. Migrate the acquired company's projects into your company's GCP organization. Link the migrated projects to your company's billing account.
- D. Create a new GCP organization and a new billing account. Migrate the acquired company's projects and your company's projects into the new GCP organization and link the projects to the new billing account.

Correct Answer: D

Reference:

<https://cloud.google.com/resource-manager/docs/migrating-projects-billing>

Community vote distribution

A (70%) B (23%) 7%

✉️  **GunjGupta**  3 years, 3 months ago

To me, A looks correct. projects are linked to another organization as well in the acquired company so migrating would need google cloud support. we can not do ourselves. however, we can link other company projects to an existing billing account to generate total cost.
<https://medium.com/google-cloud/google-cloud-platform-cross-org-billing-41c5db8fefa6>

upvoted 51 times

✉️  **uganeshku** 1 year, 9 months ago

A is correct because linking all projects of the acquired organization to the main organization's billing account will generate a single bill for all projects.
D is incorrect because there is no need to create a new organization for this.

upvoted 5 times

✉️  **spudleymcdudley** 3 years, 3 months ago

Listen to this guy. It's 'A' as moving projects can take some time from Google. There's no need to create a new organisation and other options don't make any sense

upvoted 9 times

✉️  **lxgywil** 2 years, 5 months ago

You're saying it as if "moving projects" was a viable option. What about B?

upvoted 2 times

✉️  **ryumada** 1 year, 2 months ago

I think B is not make sense. You don't want to do statistical analytic to the billing data. You want to consolidate all the costs as of tomorrow. So, the costs as of tomorrow should be billed in one billing account. That's what I've understand from the question.

upvoted 3 times

✉️  **XRiddlerX**  3 years, 2 months ago

I could be missing something but where does it say in the question that the two orgs want to migrate projects? I believe the question and key points are "consolidate all GCP costs" and "consolidate all costs as of tomorrow". With that said, C and D would not be a 24 hour task and seems a bit cumbersome to perform for something simple as "creating a single invoice" AND that's a migration and not a consolidation of cost. With A, I can't find anywhere in GCP docs that this is a best practice, only a medium.com blog. IMHO, I won't go down this route because "Just because you can do something, doesn't mean you should." and I would consult GCP support for best practices on A before I do something like that.

That leaves B which is to export both detailed billing to BigQuery and create a invoice/report. This would be a temporary solution until you migrate Organizations. IMHO

I go with B.

upvoted 46 times

✉️  **myuniqueusername** 2 years ago

How can Bigquery consolidate into 1 invoice? It's for analytics, not invoicing. B is therefore incorrect, and A must be right.

upvoted 14 times

✉️  **RazOr** 1 year, 8 months ago

You can by querying the data and create output that groups costs. BigQuery is not analytics, it can create the datasets you need from given data.

upvoted 1 times

✉️  **obeythefist** 1 year, 6 months ago

That will show you the numbers, but that does not create an invoice.

upvoted 8 times

✉ **ninjaasmoke** 10 months, 1 week ago

What does exporting data to BigQuery have to do with creating an Invoice?

upvoted 4 times

✉ **TAvenger** 2 years, 7 months ago

I am not sure that exporting some statistical data to BigQuery means anything for Google who creates the invoice. With "A" you are right, that is not the best practice, but the key word "for tomorrow" allows this custom approach. So the answer is "A"

upvoted 17 times

✉ **zaxma** 1 year, 6 months ago

I will go with A in the exam as well, but just wondering, they are two different organisations, how can you link all projects from org2 to org1's billing account without the help of GCP support??

upvoted 1 times

✉ **eBooKz** 8 months, 2 weeks ago

Cloud Billing accounts can be used across organization resources. However, organization resource moves often also include a requirement to move to a new billing account. To get the permissions that you need to change the project's billing account, ask your administrator to grant you the following IAM roles:

Billing Account User (roles/billing.user) on the destination billing account
Project billing manager (roles/billing.projectManager) on the project

<https://cloud.google.com/resource-manager/docs/project-migration#permissions-billing>

upvoted 1 times

✉ **ashrafh** 2 years, 1 month ago

I also vote B,

why?

agree with this technical explanation and my finance team not gonna pay some newly acquired company bill by tomorrow :)

upvoted 5 times

✉ **Captain1212** [Most Recent] 1 month, 1 week ago

Selected Answer: A

A is the correct answer

upvoted 1 times

✉ **sthapit** 2 months ago

To consolidate all GCP costs of both organizations onto a single invoice, you would need to set up a new GCP organization and a new billing account. This new billing account would be used to aggregate the costs of both companies' projects.

OPTION D

upvoted 1 times

✉ **jayjani66** 2 months, 3 weeks ago

Correct answer is B. By configuring both billing accounts to export their billing data into the same BigQuery dataset, you can perform cost analysis and generate consolidated reports. This will enable you to have a comprehensive view of the costs incurred by both organizations and facilitate the process of managing expenses and making financial decisions.

upvoted 1 times

✉ **Vamshi_Krishna** 5 months, 1 week ago

Selected Answer: A

A would be the quickest

upvoted 1 times

✉ **yash_901** 4 months, 2 weeks ago

Hi Krishna: I have a query:

Could you confirm if these questions are the exact questions that would come in the exam? I can see there are 201 questions mentioned but I got only 197. Also, with each question, its mentioned topic 1. Does that mean there are more topics for ACE exam?

upvoted 1 times

✉ **sakdip66** 6 months ago

option D is not the best option we can have since it involves additional effort and resources, .i.e migrating all projects from both companies which is a hassle and time consuming.

our best option is A since we acquire new company we just add it to our "company" for the purpose of putting it onto single invoice

upvoted 2 times

✉ **lazyboy** 6 months, 1 week ago

Hello everyone

upvoted 1 times

✉ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A

ANSWER "A" would be the quickest way to consolidate all costs onto a single invoice as of tomorrow. As in "tomorrow" or else your BOSS will fire you for not providing within 24 hours. LOL! Just kidding. By linking the acquired company's projects to your company's billing account, you can ensure that all costs for both GCP organizations are billed to a single billing account.

ANSWER "C" would be the proper way to consolidate the GCP organizations in the long term, but it would take more time to migrate the acquired company's projects into your company's GCP organization and link the projects to your company's billing account.

ANSWER "B" suggests configuring both billing accounts to export billing data into the same BigQuery dataset. While this approach can help consolidate billing data and provide a more comprehensive view of costs, it does not provide a single invoice for all costs as the question requires.

upvoted 3 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

ANSWER "D" involves creating a new GCP organization and a new billing account and migrating both companies' projects into the new organization and billing account. While this approach can consolidate costs onto a single invoice, it is a more complex and time-consuming process that may not be feasible to complete within 24 hours as the question requires. It is also not ideal if the two companies' projects have different access control requirements or billing policies that need to be maintained separately.

Therefore, in the context of the given question and the requirement to consolidate costs as of tomorrow onto a single invoice, ANSWER "A", linking the acquired company's projects to the existing billing account, is the most appropriate and practical solution.

upvoted 2 times

 **kaisehhop** 8 months, 1 week ago

Selected Answer: C

C. Migrate the acquired company's projects into your company's GCP organization. Link the migrated projects to your company's billing account.

To consolidate costs onto a single invoice, you need to bring all projects under the same GCP organization and link them to the same billing account. The best way to achieve this is to migrate the acquired company's projects into your company's GCP organization and link them to your company's billing account. This will ensure that all costs are consolidated onto a single invoice and can be easily managed and monitored from a single location.

upvoted 2 times

 **cslince** 10 months, 1 week ago

Selected Answer: A

Correct Answer should be A

upvoted 1 times

 **fragment137** 10 months, 2 weeks ago

Selected Answer: A

I'll say A. I've dealt with a few companies that have purchased other organizations and have done this for simplified billing.

upvoted 3 times

 **Robertolo** 12 months ago

Selected Answer: B

A. Can hundreds of projects be linked for tomorrow to the company's billing account? I highly doubt it.

B. Exporting to BigQuery seems a reasonable urgent solution to consolidate costs ("as of tomorrow").

C. Can hundreds of projects be migrated + linked for tomorrow? Same as with A, I doubt it.

D. Same as A and C: I doubt that we can create a new GCP organization, migrate all new and old projects and link them to the new billing account in just a single day. This is madness :-D

upvoted 2 times

 **sjagkoo7** 12 months ago

i will go with option D . As question states single invoice for all projects means all should be linked to one billing account. Re-organising the only option this case..

upvoted 1 times

 **ashit44244** 1 year ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!

upvoted 1 times

 **orious** 1 year, 2 months ago

Selected Answer: B

Cost consolidation is an Analytical Task (Account task), you rely on analytical tools, therefore you use BigQuery. So, there is no need for complexity.

-- "Consolidation Expenses means, for any period and with respect to any Person, the aggregate of all expenditures by such Person and its Subsidiaries on a consolidated basis during such period related to the consolidation of Station". (

Consolidation Expenses Definition | Law Insiderhttps://www.lawinsider.com › dictionary › consolidation-e...)

upvoted 2 times

 **AzureDP900** 1 year, 3 months ago

A is correct

upvoted 2 times

You built an application on Google Cloud that uses Cloud Spanner. Your support team needs to monitor the environment but should not have access to table data.

You need a streamlined solution to grant the correct permissions to your support team, and you want to follow Google-recommended practices.

What should you do?

- A. Add the support team group to the roles/monitoring.viewer role
- B. Add the support team group to the roles/spanner.databaseUser role.
- C. Add the support team group to the roles/spanner.databaseReader role.
- D. Add the support team group to the roles/stackdriver.accounts.viewer role.

Correct Answer: B

Community vote distribution

A (100%)

 **poogcp**  3 years, 4 months ago

its A, As you need to monitor only
upvoted 44 times

 **WindDriver** 2 years, 2 months ago

A, right, correct answer.
B and C are incorrect because allow to read data.
D also incorrect: Not for monitoring. roles/stackdriver.accounts.viewer Stackdriver Accounts Viewer:
Read-only access to get and list information about Stackdriver account structure (resourcemanager.projects.get, resourcemanager.projects.list and stackdriver.projects.get)
upvoted 14 times

 **WindDriver** 2 years, 2 months ago

<https://cloud.google.com/iam/docs/understanding-roles>
upvoted 3 times

 **Gurnoor**  3 years, 3 months ago

A is correct as user should not have any access to data, so B and C cant be used in this scenario.
upvoted 18 times

 **ekta25**  4 days ago

A. Add the support team group to the roles/monitoring.viewer role
upvoted 1 times

 **axantroff** 1 week, 2 days ago

Selected Answer: A
Makes sense for me
upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: A
A as you only need the monitor access
upvoted 1 times

 **sakdip66** 6 months ago

the goal of support team is to MONITOR the environment only. therefore roles/monitoring.viewer role is the best option we have
<https://cloud.google.com/spanner/docs/iam#roles>
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A
Answer A, adding the support team group to the roles/monitoring.viewer role, is the CORRECT answer. This role grants read-only access to monitoring data for all resources in a project, which allows the support team to monitor the environment but not access the table data.
Answer B, adding the support team group to the roles/spanner.databaseUser role, grants read and write access to all tables in the specified database, which is NOT required for the support team to monitor the environment.

Answer C, adding the support team group to the roles/spanner.databaseReader role, grants read-only access to all tables in the specified database, which would give the support team access to the table data.

Answer D, adding the support team group to the roles/stackdriver.accounts.viewer role, grants permissions to view Stackdriver data for all resources in a project, which is NOT directly related to monitoring the Cloud Spanner environment.

upvoted 6 times

 **cslince** 10 months, 1 week ago

Selected Answer: A

A is correct
upvoted 1 times

 **Zoze** 10 months, 3 weeks ago

Selected Answer: A

A is correct, the team need to monitor the environment not read the data.
upvoted 1 times

 **Cornholio_LMC** 1 year ago

had this question today
upvoted 3 times

 **raghu09** 1 year ago

Selected Answer: A

B is wrong because it grants write access also we only need monitoring access.
upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Selected Answer: A

A. This is the only role that provides read-only access to get and list information about all monitoring data and configurations.
upvoted 2 times

 **AzureDP900** 1 year, 3 months ago

You only need to monitor so A is correct!
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

roles/monitoring.viewer

Monitoring Viewer Grants read-only access to Monitoring in the Google Cloud console and API.

upvoted 1 times

 **wolfie09** 1 year, 4 months ago

Selected Answer: A

A is correct
upvoted 2 times

 **haroldbenites** 1 year, 4 months ago

Go for A
upvoted 1 times

 **Priyankahere** 1 year, 8 months ago

This was there in exam, go with community answers.
upvoted 8 times

 **NoniGeorge** 1 year, 9 months ago

Selected Answer: A

You only need to monitor so A is correct!
upvoted 1 times

For analysis purposes, you need to send all the logs from all of your Compute Engine instances to a BigQuery dataset called platform-logs. You have already installed the Cloud Logging agent on all the instances. You want to minimize cost. What should you do?

- A. 1. Give the BigQuery Data Editor role on the platform-logs dataset to the service accounts used by your instances. 2. Update your instances' metadata to add the following value: logs-destination: bq://platform-logs.
- B. 1. In Cloud Logging, create a logs export with a Cloud Pub/Sub topic called logs as a sink. 2. Create a Cloud Function that is triggered by messages in the logs topic. 3. Configure that Cloud Function to drop logs that are not from Compute Engine and to insert Compute Engine logs in the platform-logs dataset.
- C. 1. In Cloud Logging, create a filter to view only Compute Engine logs. 2. Click Create Export. 3. Choose BigQuery as Sink Service, and the platform-logs dataset as Sink Destination.
- D. 1. Create a Cloud Function that has the BigQuery User role on the platform-logs dataset. 2. Configure this Cloud Function to create a BigQuery Job that executes this query: `INSERT INTO dataset.platform-logs (timestamp, log) SELECT timestamp, log FROM compute.logs WHERE timestamp > DATE_SUB(CURRENT_DATE(), INTERVAL 1 DAY)` 3. Use Cloud Scheduler to trigger this Cloud Function once a day.

Correct Answer: C

Community vote distribution



✉ **sumanshu** 2 years, 6 months ago
vote for "C"

https://cloud.google.com/logging/docs/export/configure_export_v2
upvoted 20 times

✉ **vmart** 2 years, 4 months ago
I vote for C
upvoted 7 times

✉ **axantroff** 1 week, 2 days ago
C, it's simple enough
upvoted 1 times

✉ **Captain1212** 1 month, 1 week ago

Selected Answer: C
C , it minimizes the cost
upvoted 1 times

✉ **jayjani66** 2 months, 3 weeks ago

Answer: C. Option C allows you to create a log export from Cloud Logging to BigQuery with minimal setup and cost. By creating a filter to view only Compute Engine logs, you ensure that only the relevant logs are exported to BigQuery, reducing unnecessary data transfer and storage costs.
upvoted 1 times

✉ **Kyle1776** 4 months, 1 week ago

Selected Answer: B
B looks like the most cost effective option since filtering out only the logs you need will reduce storage and data transfer costs.
upvoted 1 times

Correction - C

upvoted 1 times

✉ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: C
The most cost-effective and recommended solution to send logs from Compute Engine instances to BigQuery is to use the Cloud Logging agent with a sink that streams the logs to BigQuery.

Answer C is the most appropriate solution. In Cloud Logging, create a filter to view only Compute Engine logs. Click Create Export. Choose BigQuery as Sink Service, and the platform-logs dataset as Sink Destination. This will allow all Compute Engine instance logs to be exported to BigQuery with minimal complexity and cost.

upvoted 1 times

✉ **David_Esteban** 10 months ago

Selected Answer: C
Mi vote is for "c"
upvoted 1 times

 **cslince** 10 months, 1 week ago

Selected Answer: C

vote for "C"

upvoted 1 times

 **Charumathi** 1 year ago

C. is correct, Sinks control how Cloud Logging routes logs. Using sinks, you can route some or all of your logs to supported destinations.

upvoted 3 times

 **AzureDP900** 1 year, 3 months ago

I will go with C..

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 1 times

 **somenick** 1 year, 6 months ago

Selected Answer: C

Outdated question. It's now about Cloud Sink, but C is the closest option

upvoted 5 times

 **NoniGeorge** 1 year, 9 months ago

I think it's C as all the other ones seem to get logs from everywhere not just Compute Engine!

upvoted 4 times

 **jaykumarjk99** 1 year, 10 months ago

Selected Answer: C

I vote for C

upvoted 4 times

 **vishnukumartr** 1 year, 10 months ago

C. 1. In Cloud Logging, create a filter to view only Compute Engine logs. 2. Click Create Export. 3. Choose BigQuery as Sink Service, and the platform-logs dataset as Sink Destination.

upvoted 2 times

 **Nivesh93** 2 years, 1 month ago

c is correct as it uses filter so it basically reduces the cost of operation .

upvoted 2 times

You are using Deployment Manager to create a Google Kubernetes Engine cluster. Using the same Deployment Manager deployment, you also want to create a DaemonSet in the kube-system namespace of the cluster. You want a solution that uses the fewest possible services. What should you do?

- A. Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.
- B. Use the Deployment Manager Runtime Configurator to create a new Config resource that contains the DaemonSet definition.
- C. With Deployment Manager, create a Compute Engine instance with a startup script that uses kubectl to create the DaemonSet.
- D. In the cluster's definition in Deployment Manager, add a metadata that has kube-system as key and the DaemonSet manifest as value.

Correct Answer: C

Reference:

<https://cloud.google.com/kubernetes-engine/docs/how-to/cluster-access-for-kubectl>

Community vote distribution

A (82%) D (18%)

 **ESP_SAP**  3 years, 1 month ago

Correct Answer is (A):

Adding an API as a type provider

This page describes how to add an API to Google Cloud Deployment Manager as a type provider. To learn more about types and type providers, read the Types overview documentation.

A type provider exposes all of the resources of a third-party API to Deployment Manager as base types that you can use in your configurations. These types must be directly served by a RESTful API that supports Create, Read, Update, and Delete (CRUD).

If you want to use an API that is not automatically provided by Google with Deployment Manager, you must add the API as a type provider.

<https://cloud.google.com/deployment-manager/docs/configuration/type-providers/creating-type-provider>

upvoted 69 times

 **magistrum** 2 years, 9 months ago

very good find, sounds like you hit the nail in the head
upvoted 7 times

 **kishoredeena**  3 years, 3 months ago

Option A is the right answer
upvoted 10 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: A
A is the correct , bcoz it help you contact directly to the gke cluster to create daemon
upvoted 1 times

 **sthapit** 2 months ago

Should have been D
upvoted 1 times

 **sakdip66** 6 months ago

Selected Answer: A
option A is the right answer because it lets you directly interact with the Kubernetes API to create the Daemonset using the same deployment Manager Deployment
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A
I would say both Answer A and Answer D are valid solutions, and it depends on your preference and requirements.

Answer A involves adding the cluster's API as a new Type Provider in Deployment Manager and using the new type to create the DaemonSet. This solution would allow you to create and manage the DaemonSet and the cluster in the same Deployment Manager deployment.

Answer D involves adding a metadata block to the Deployment Manager deployment of the cluster, which will create the DaemonSet in the kube-system namespace of the cluster. This solution would allow you to create the DaemonSet in a simple way and avoid the need to create a new Type of Provider.

In conclusion, I would choose Answer A to be considered the answer that uses the fewest possible services, as it only involves adding the cluster's API as a new Type Provider in Deployment Manager, which is a lightweight solution.

upvoted 4 times

 **Bobbybash** 7 months, 4 weeks ago

Selected Answer: D

D. In the cluster's definition in Deployment Manager, add a metadata that has kube-system as key and the DaemonSet manifest as value.

This approach involves adding the DaemonSet manifest directly as a metadata entry in the cluster's definition in Deployment Manager. When the cluster is created, the DaemonSet is automatically created in the kube-system namespace. This approach is the simplest and requires the fewest number of services. Option A is also a viable solution but requires more work to set up a Type Provider. Option B is not suitable because it involves a separate service (Runtime Configurator). Option C is also not recommended because it involves creating a Compute Engine instance and using kubectl to create the DaemonSet, which is more complicated and less efficient than the other options.

upvoted 3 times

 **vkamlesh0205** 9 months, 2 weeks ago

Selected Answer: A

Option A is the right answer

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Selected Answer: A

Answer is A.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

go with A as per ESP_SAP explanations..

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for A

upvoted 1 times

 **luciorifa** 1 year, 7 months ago

Selected Answer: A

A is the correct answe, the API need to be added as a type provider

upvoted 1 times

 **ArunTaneja** 1 year, 8 months ago

Selected Answer: A

A should be correct one

upvoted 1 times

 **thuvh** 1 year, 9 months ago

Selected Answer: A

<https://medium.com/google-cloud/cloud-deployment-manager-kubernetes-2dd9b8124223>

upvoted 2 times

 **axantroff** 1 week, 2 days ago

Good reference. Thanks for it. Recomended

upvoted 1 times

 **MCMS** 1 year, 10 months ago

Selected Answer: A

Correct Answer is (A)

upvoted 2 times

 **PR0704** 1 year, 10 months ago

couldn't be more confusing

upvoted 5 times

 **vishnukumartr** 1 year, 10 months ago

A. Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.

upvoted 1 times

You are building an application that will run in your data center. The application will use Google Cloud Platform (GCP) services like AutoML. You created a service account that has appropriate access to AutoML. You need to enable authentication to the APIs from your on-premises environment. What should you do?

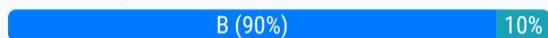
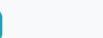
- A. Use service account credentials in your on-premises application.
- B. Use gcloud to create a key file for the service account that has appropriate permissions.
- C. Set up direct interconnect between your data center and Google Cloud Platform to enable authentication for your on-premises applications.
- D. Go to the IAM & admin console, grant a user account permissions similar to the service account permissions, and use this user account for authentication from your data center.

Correct Answer: B

Reference:

<https://cloud.google.com/vision/automl/docs/before-you-begin>

Community vote distribution

 B (90%)  10%

 **ESP_SAP**  3 years, 1 month ago

Correct answer should be (B):

To use a service account outside of Google Cloud, such as on other platforms or on-premises, you must first establish the identity of the service account. Public/private key pairs provide a secure way of accomplishing this goal.

<https://cloud.google.com/iam/docs/creating-managing-service-account-keys>
upvoted 47 times

 **Mcsniper** 1 month ago

Can you please share all the GCP questions with me at mcsniper976@gmail.com
upvoted 1 times

 **SIGMAAYAN** 2 months, 3 weeks ago

Can you share the all the GCP questions to me at samalayan@gmail.com
upvoted 1 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: B

B is the correct answer, as to access the out side the google cloud , you need the key
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: B

The recommended approach for enabling authentication from an on-premises environment to Google Cloud Platform (GCP) services like AutoML is to use a service account and generate a JSON key file for the service account. This key file can then be used to authenticate and authorize API calls from your on-premises environment to GCP.

Therefore, the correct answer is B. Use gcloud to create a key file for the service account that has appropriate permissions.
upvoted 4 times

 **Bobbybash** 7 months, 4 weeks ago

Selected Answer: A

A. Use service account credentials in your on-premises application.

To enable authentication to GCP services from your on-premises environment, you can use service account credentials in your on-premises application. This involves creating a service account that has appropriate access to the required GCP services, downloading the service account key file, and using the key file to authenticate the API requests in your on-premises application. This is a secure way to authenticate to GCP services as it does not require direct access to your GCP project or credentials from your on-premises environment.
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Cloud Security/Auditor doesn't like Answer "A". Using service account credentials in your on-premises application could be a security risk if the credentials are compromised. If the key file is stolen or leaked, an attacker could use it to access your GCP resources, potentially causing data breaches, service disruptions, or financial losses.

I would select Answer "B". Use gcloud to create a key file for the service account that has appropriate permissions and let Security Auditor stay away from my back. Never-ending "You cannot do this, you cannot do that" on Answer A.
upvoted 3 times

 **srasam** 9 months, 2 weeks ago

I am able to access 96 questions only, if anyone has the entire questions please share them on gcp1on1@gmail.com

upvoted 1 times

✉ **u422628** 9 months ago

u must pay

upvoted 5 times

✉ **cslince** 10 months, 1 week ago

Selected Answer: B

B it is.

upvoted 1 times

✉ **mvk2022** 10 months, 2 weeks ago

Selected Answer: B

B it is.

upvoted 1 times

✉ **Kopy** 11 months, 1 week ago

Selected Answer: B

Correct answer should be (B):

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

B is right

To use a service account from outside of Google Cloud, such as on other platforms or on-premises, you must first establish the identity of the service account. Public/private key pairs provide a secure way of accomplishing this goal. When you create a service account key, the public portion is stored on Google Cloud, while the private portion is available only to you. For more information about public/private key pairs, see Service account keys.

upvoted 4 times

✉ **haroldbenites** 1 year, 4 months ago

Go for B

upvoted 1 times

✉ **NoniGeorge** 1 year, 9 months ago

Selected Answer: B

Even thought A and B seem to be doing the same thing the best practice is to create a key so B is the right answer !

upvoted 1 times

✉ **vishnukumartr** 1 year, 10 months ago

B. Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 3 times

✉ **shawnkkk** 1 year, 10 months ago

B. Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 3 times

✉ **Vivekvkt123** 2 years ago

Why not A? Aren't A and B getting the same key file?

upvoted 2 times

✉ **jabrrJ68w02ond1** 1 year, 10 months ago

A is not really telling you the steps to accomplish the task, it's only telling you the result of it (creating a SA with sufficient permissions and then use Console / gcloud to create a JSON token for it)

upvoted 1 times

✉ **sunilw** 2 years, 3 months ago

B is correct.

Creating service account keys

To use a service account from outside of Google Cloud, such as on other platforms or on-premises, you must first establish the identity of the service account. Public/private key pairs provide a secure way of accomplishing this goal. When you create a service account key, the public portion is stored on Google Cloud, while the private portion is available only to you. For more information about public/private key pairs, see Service account keys.

upvoted 3 times

✉ **arsh1916** 2 years, 4 months ago

B is correct

upvoted 1 times

✉ **mcaromit** 2 years, 5 months ago

B is correct as a 1st step...the key file is to be referenced in the env variable GOOGLE_APPLICATION_CREDENTIALS which would then provide access to on-prem application using ADC library

upvoted 1 times

You are using Container Registry to centrally store your company's container images in a separate project. In another project, you want to create a Google Kubernetes Engine (GKE) cluster. You want to ensure that Kubernetes can download images from Container Registry. What should you do?

- A. In the project where the images are stored, grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes.
- B. When you create the GKE cluster, choose the Allow full access to all Cloud APIs option under 'Access scopes'.
- C. Create a service account, and give it access to Cloud Storage. Create a P12 key for this service account and use it as an imagePullSecrets in Kubernetes.
- D. Configure the ACLs on each image in Cloud Storage to give read-only access to the default Compute Engine service account.

Correct Answer: C

Community vote distribution

A (92%) 8%

✉️ **ESP_SAP** Highly Voted 3 years, 1 month ago

Correct Answer (A):

IAM permissions

IAM permissions determine who can access resources. All users, service accounts, and other identities that interact with Container Registry must have the appropriate Cloud Storage permissions.

By default, Google Cloud uses default service accounts to interact with resources within the same project. For example, the Cloud Build service account can both push and pull images when Container Registry is in the same project.

You must configure or modify permissions yourself if:

You are using a service account in one project to access Container Registry in a different project

You are using a default service account with read-only access to storage, but you want to both pull and push images

You are using a custom service account to interact with Container Registry

<https://cloud.google.com/container-registry/docs/access-control>

upvoted 58 times

✉️ **[Removed]** 3 years ago

A is correct, practical implementation in video <https://www.youtube.com/watch?v=R16z7Sjrkxs>

upvoted 11 times

✉️ **XRiddlerX** Highly Voted 3 years, 1 month ago

A is correct...

Container Registry uses Cloud Storage buckets as the underlying storage for container images. You control access to your images by granting appropriate Cloud Storage permissions to a user, group, service account, or other identity.

If the service account needs to access Container Registry in another project, you must grant the required permissions in the project with Container Registry.

Reference:

<https://cloud.google.com/container-registry/docs/access-control#permissions>

upvoted 18 times

✉️ **Captain1212** Most Recent 1 month, 1 week ago

Selected Answer: A

A is the correct answer , as granting this role allow to download the image

upvoted 1 times

✉️ **sakdip66** 6 months ago

Selected Answer: A

Grating Storage Object Viewer IAM Role to the service account used by Kubernetes nodes allow the nodes to download the images from Container registry.

upvoted 1 times

✉️ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A

Answer A. In the project where the images are stored, grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes.

To ensure that Kubernetes can download container images from Container Registry, you need to grant the necessary permissions to the service account used by the Kubernetes nodes. In this case, you would need to grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes in the project where the images are stored. This role allows the service account to read objects from Cloud Storage buckets, including the container images in Container Registry.

upvoted 3 times

 **jrlsl1991** 8 months, 1 week ago

Selected Answer: A

Definitely A seems more practical and accurate.

upvoted 1 times

 **GaneshSurwase** 11 months, 4 weeks ago

CORRET ANS is A

upvoted 1 times

 **ravip12345** 1 year ago

Answer A is correct.

Storage Object Viewer (roles/storage.objectViewer) -- Grant the role on the registry storage bucket.

<https://cloud.google.com/container-registry/docs/access-control>

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

Storage Object viewer is enough, A is right.

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for A.

the option A is specific with the role that will be used. In GCP , the recommendations is using the specific permission. The others options not are specific and are not correct .

upvoted 1 times

 **somenick** 1 year, 6 months ago

Selected Answer: A

<https://cloud.google.com/container-registry/docs/access-control>

upvoted 2 times

 **luciorifa** 1 year, 7 months ago

Selected Answer: A

A is the correct answe because it follows google recommended practices to grant permissions to service accounts, also the object viewer is the appropiate role for the GKE to pull the image

upvoted 2 times

 **jdx000** 1 year, 8 months ago

Selected Answer: A

A is the most obvious

upvoted 1 times

 **ArunTaneja** 1 year, 8 months ago

A should be correct

https://cloud.google.com/container-registry/docs/pushing-and-pulling#pulling_images_from_a_registry

<https://cloud.google.com/storage/docs/access-control/iam-roles#standard-roles>

upvoted 2 times

 **NoniGeorge** 1 year, 9 months ago

It's A as you need Storage Object Viewer IAM role in order to have access to the images!

upvoted 1 times

 **DiegoCG** 1 year, 9 months ago

Selected Answer: R

ES A Y PÚNTO

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: R "as in" RONG. LOL!

upvoted 2 times

 **Anirudha077** 10 months ago

What??

upvoted 1 times

 **Pret** 1 year, 10 months ago

Option A:

Pulling images from a registry

Pulling an image requires the Storage Object Viewer for the registry storage bucket, or a role with the same permissions.

To pull from Container Registry, use the command:

`docker pull HOSTNAME/PROJECT-ID/IMAGE:TAG`

<https://cloud.google.com/container-registry/docs/pushing-and-pulling>

upvoted 1 times

You deployed a new application inside your Google Kubernetes Engine cluster using the YAML file specified below.

```
apiVersion: apps/v1           apiVersion: v1
kind: Deployment               kind: Service
metadata:
  name: myapp-deployment     metadata:
                                name: myapp-service
spec:
  selector:
    matchLabels:
      app: myapp
  replicas: 2
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
        - name: myapp
          image: myapp:1.1
        ports:
          - containerPort: 80
ports:
  - port: 8000
    targetPort: 80
    protocol: TCP
selector:
  app: myapp
```

You check the status of the deployed pods and notice that one of them is still in PENDING status:

```
kubectl get pods -l app=myapp
NAME                           READY   STATUS    RESTART   AGE
myapp-deployment-58ddbbb995-lp86m   0/1     Pending   0          9m
myapp-deployment-58ddbbb995-qjpkd   1/1     Running   0          9m
```

You want to find out why the pod is stuck in pending status. What should you do?

- A. Review details of the myapp-service Service object and check for error messages.
- B. Review details of the myapp-deployment Deployment object and check for error messages.
- C. Review details of myapp-deployment-58ddbbb995-lp86m Pod and check for warning messages.
- D. View logs of the container in myapp-deployment-58ddbbb995-lp86m pod and check for warning messages.

Correct Answer: C

Reference:

<https://cloud.google.com/run/docs/gke/troubleshooting>

Community vote distribution

C (88%)

13%

 **spudleymcdudley**  3 years, 3 months ago

It's C - <https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/#debugging-pods>
upvoted 30 times

 **someoneinthecloud**  3 years, 2 months ago

Answer is C - You can't view logs of a pod that isn't deployed, so D is incorrect.
C allows you to check the pod deployment messages and look for errors
upvoted 20 times

 **sidharthwader** 2 years, 3 months ago

What u said is incorrect you can view pod's log even in pending state.
kubectl logs <pod-name> -n <namespace>
upvoted 3 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: C

C is correct, as its help you to check the error
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: C

Answer C. Review details of myapp-deployment-58ddbbb995-lp86m Pod and check for warning messages.

To find out why a pod is stuck in pending status, you can review the details of the pod and check for any warning messages. Answer C is the correct answer because it suggests reviewing the details of the specific pod that is stuck in pending status. You can use the kubectl

describe pod <pod-name> command to view detailed information about the pod, including any warning messages that might indicate why the pod is not scheduled.

upvoted 2 times

✉ **AKSHAT09jain** 9 months ago

D : we first check logs

upvoted 1 times

✉ **srasam** 9 months, 2 weeks ago

I am able to access 96 questions only, if anyone has the entire questions please share them on gcp1on1@gmail.com

upvoted 1 times

✉ **Zoze** 10 months, 3 weeks ago

Selected Answer: C

I vote C; because if we imagine that we will go to a main menu that display the errors of the all deployment object as hole, we will surely navigate thin to the pod menu ! so the C option will direct take us to the second menu.

upvoted 1 times

✉ **Charumathi** 1 year ago

C is correct,

Debugging Pods

The first step in debugging a Pod is taking a look at it. Check the current state of the Pod and recent events with the following command:

```
kubectl describe pods ${POD_NAME}
```

upvoted 1 times

✉ **Letahrgicbeagle** 1 year ago

Selected Answer: C

Definitely

upvoted 1 times

✉ **Dheeraj1986** 1 year, 1 month ago

Selected Answer: B

I guess it's B. its deployment that creates the pod and it has the information why it is not able to create. it shows the information if you describe the deployment (kubectl describe deployment)

upvoted 1 times

✉ **abirroy** 1 year, 2 months ago

Selected Answer: C

Answer is C

upvoted 1 times

✉ **RanjithK** 1 year, 3 months ago

Selected Answer: C

Answer is C

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

C is right

upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 1 times

✉ **sharmaishu** 1 year, 7 months ago

C is the correct answer: The first step in debugging a Pod is taking a look at it. Check the current state of the Pod and recent events with the following command:

```
kubectl describe pods ${POD_NAME}
```

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/>

upvoted 1 times

✉ **sharmaishu** 1 year, 7 months ago

C is the correct answer.

If a Pod is stuck in Pending it means that it can not be scheduled onto a node. Generally this is because there are insufficient resources of one type or another that prevent scheduling. Look at the output of the kubectl describe ... command above. There should be messages from the scheduler about why it can not schedule your Pod.

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/>

upvoted 2 times

✉ **vishnukumartr** 1 year, 10 months ago

C. Review details of myapp-deployment-58ddbbb995-lp86m Pod and check for warning messages.

upvoted 1 times

You are setting up a Windows VM on Compute Engine and want to make sure you can log in to the VM via RDP. What should you do?

- A. After the VM has been created, use your Google Account credentials to log in into the VM.
- B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.
- C. When creating the VM, add metadata to the instance using 'windows-password' as the key and a password as the value.
- D. After the VM has been created, download the JSON private key for the default Compute Engine service account. Use the credentials in the JSON file to log in to the VM.

Correct Answer: D

Community vote distribution

B (100%)

 **John_Iam**  3 years, 4 months ago

Correct Answer is B.

B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.

<https://cloud.google.com/sdk/gcloud/reference/beta/compute/reset-windows-password>
upvoted 58 times

 **voler** 3 years, 3 months ago

Yes! "If the Windows account does not exist, this command will cause the account to be created and the password for that new account will be returned."

upvoted 3 times

 **ragu123** 3 years, 1 month ago

Correct answer is B.

gcloud beta compute reset-windows-password allows a user to reset and retrieve a password for a Windows virtual machine instance. If the Windows account does not exist, this command will cause the account to be created and the password for that new account will be returned.

upvoted 3 times

 **dan80** 3 years, 4 months ago

did you even look at the link you provide ? it clearly say gcloud beta compute reset-windows-password my-instance and not gcloud compute reset-windows-password. D is correct - <https://cloud.google.com/iam/docs/creating-managing-service-account-keys>

upvoted 4 times

 **Ixgywil** 2 years, 5 months ago

Oh yes? Then what about this link (for non-beta command)?

<https://cloud.google.com/sdk/gcloud/reference/compute/reset-windows-password>

"If the Windows account does not exist, this command will cause the account to be created and the password for that new account will be returned."

The answer is obviously B. Just test it and it'll become very clear

upvoted 7 times

 **dan80** 3 years, 4 months ago

nobody talk on reset the password but how to access the Windows - best way - Service Account

upvoted 1 times

 **KerolesKhalil** 4 months ago

Service accounts shouldn't be used for RDP , they are used to machine authentication with services.

upvoted 1 times

 **KerolesKhalil** 4 months ago

Also how you will RDP with Service account and private key

You need username and password , it is not ssh

upvoted 1 times

 **Buruguduystunstugudunstuy**  7 months, 3 weeks ago

Selected Answer: B

CORRECT:

Answer B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.

INCORRECT:

Answer A is not correct because Google Account credentials cannot be used to log in to Windows VMs.

Answer C is not correct because metadata can be used to specify some settings for a VM, but the 'windows-password' metadata key is not used for specifying the login password for a Windows VM.

Answer D is not correct because the JSON private key for the default Compute Engine service account is not used for logging in to a Windows VM.

upvoted 5 times

 **Captain1212** Most Recent 1 month, 1 week ago

Selected Answer: B

Yes , B is the correct answer

upvoted 1 times

 **sakdip66** 6 months ago

Selected Answer: B

after creating Windows VM on COmpute Engine it has a local user account as well. This acct is used to login to the VM via RDP. If you forget the password you can use gcloud compute reset-windows-password to reset it. This command generates a new password and sets it for the user account on the VM.

upvoted 1 times

 **Partha117** 6 months, 2 weeks ago

Selected Answer: B

Correct option B

upvoted 1 times

 **smanoj85** 6 months, 3 weeks ago

Option B is correct. After creating a Windows VM instance on Compute Engine, you need to use the gcloud compute reset-windows-password command to retrieve the login credentials for the VM. This command generates a new Windows password and displays it in the output of the command. You can then use this password to log in to the VM via RDP.

Option A is incorrect because logging in to the VM using your Google Account credentials is not a supported method for Windows VM instances.

Option C is also incorrect because 'windows-password' is not a recognized metadata key for Windows VM instances.

Option D is incorrect because you cannot use the JSON private key for the default Compute Engine service account to log in to a Windows VM instance via RDP.

upvoted 2 times

 **cslince** 10 months ago

Selected Answer: B

Correct Answer is B.

B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM

upvoted 1 times

 **ale_brd_** 1 year ago

Selected Answer: B

After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is perfect for this question..

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for B

<https://cloud.google.com/sdk/gcloud/reference/compute/reset-windows-password>

upvoted 2 times

 **Himadhar1997** 1 year, 4 months ago

B is correct

after VM creating , you can reset the password

upvoted 1 times

 **Harbeeb** 1 year, 5 months ago

Selected Answer: B

After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM

upvoted 2 times

 **dinesh198728** 1 year, 6 months ago

Selected Answer: B

b is correct

upvoted 1 times

 **luciorifa** 1 year, 7 months ago

Selected Answer: B

B is the correct answer, the command is the correct one

upvoted 1 times

 **ttttttttt** 1 year, 8 months ago

Test. Yyyyy

upvoted 1 times

 **xianpxian** 1 year, 9 months ago

Yes, It's B

upvoted 1 times

 **Zufair** 1 year, 9 months ago

Selected Answer: B

Correct Answer is B. Tested it.

upvoted 1 times

You want to configure an SSH connection to a single Compute Engine instance for users in the dev1 group. This instance is the only resource in this particular Google Cloud Platform project that the dev1 users should be able to connect to. What should you do?

- A. Set metadata to enable-oslogin=true for the instance. Grant the dev1 group the compute.osLogin role. Direct them to use the Cloud Shell to ssh to that instance.
- B. Set metadata to enable-oslogin=true for the instance. Set the service account to no service account for that instance. Direct them to use the Cloud Shell to ssh to that instance.
- C. Enable block project wide keys for the instance. Generate an SSH key for each user in the dev1 group. Distribute the keys to dev1 users and direct them to use their third-party tools to connect.
- D. Enable block project wide keys for the instance. Generate an SSH key and associate the key with that instance. Distribute the key to dev1 users and direct them to use their third-party tools to connect.

Correct Answer: D

Reference:

<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>

Community vote distribution

A (94%) 6%

✉️  **poogcp**  3 years, 4 months ago

A correct one
upvoted 44 times

✉️  **spudleymcdudley** 3 years, 3 months ago

For further evidence... <https://cloud.google.com/compute/docs/instances/managing-instance-access>
upvoted 8 times

✉️  **nithinpb180** 3 years, 3 months ago

Agree with that
upvoted 3 times

✉️  **student002**  2 years, 11 months ago

Pure from logic thinking: A can't be right. If the group get access to that instance with enable-oslogin=true, then they could have access to every instance that has enable-oslogin=true. Or do I miss something?
upvoted 11 times

✉️  **akshaychavan7** 1 year, 4 months ago

Note the sentence "Set metadata to enable-oslogin=true for the instance." This means the metadata for oslogin has been set to that particular instance only, and not for all.
upvoted 4 times

✉️  **magistrum** 2 years, 9 months ago

I'm convinced with this logic
upvoted 1 times

✉️  **bgallet** 1 year, 10 months ago

clearly, question say "the only ressource they need to access in this project"
as you said, all ressources will be available if we set the role
upvoted 2 times

✉️  **jrlsl1991** 8 months, 1 week ago

That's not necessarily true - <https://cloud.google.com/compute/docs/oslogin/set-up-oslogin>. The doc says "If you want enable OS Login for all VMs in a project, set the metadata at the project-level. If you want to enable OS Login for a single VM, set the metadata at the instance-level."

That means you can do it at the instance level, so there shouldn't be a problem with following A.

upvoted 1 times

✉️  **Captain1212**  1 month, 1 week ago

Selected Answer: A

A is correct as , it gives the only specific access
upvoted 1 times

✉️  **KerolesKhalil** 4 months ago

Selected Answer: A

<https://cloud.google.com/compute/docs/oslogin/set-up-oslogin>

upvoted 1 times

✉ **sakdip66** 6 months ago

Selected Answer: A

Enabling OSLogin allow user to login to Google Cloud credentials to authenticate to instance, instead of SSH key. Granting 'compute.osLogin' to the dev1 lets them login using OSLogin to the resourcee

BD are incorrect because "block project-wide SSH keys" is an advance security features taht is used for high secured environment where more granular control over SSH us required

C is hassle because it manually ditribute keys to each user in Dev1 which is time consuming

upvoted 1 times

✉ **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A

Answer B is incorrect because setting the service account to no service account has no impact on SSH access to the VM instance.

Answer C is incorrect because generating an SSH key for each user in the dev1 group and distributing them is cumbersome and not scalable, especially if you have many users.

Answer D is incorrect because generating a single SSH key and distributing it to multiple users undermines security, as it means any of the users in possession of the key can access the VM instance.

upvoted 1 times

✉ **FeaRoX** 8 months, 1 week ago

In my opinion A would be best, but they have to use this and only this 1 instance. You don't know if any other instances has this metadata set up - if they do, dev1 team has also access to this instances, what invalidates the answer. To make sure they are using only this 1 instance, I'd say D.

upvoted 1 times

✉ **jrisl1991** 8 months, 1 week ago

Selected Answer: A

Based on <https://cloud.google.com/compute/docs/oslogin/set-up-oslogin> I'd go for A.

upvoted 1 times

✉ **alex000** 9 months ago

Selected Answer: C

The dev1 users should be able to connect only to this VM instance

upvoted 1 times

✉ **cslince** 10 months ago

Selected Answer: A

A correct one

upvoted 1 times

✉ **Charumathi** 1 year ago

A is the correct answer

Granting OS Login IAM roles

After you enable OS Login on one or more instances in your project, those VMs accept connections only from user accounts that have the necessary IAM roles in your project or organization.

roles/compute.osLogin, which doesn't grant administrator permissions

upvoted 2 times

✉ **Cornholio_LMC** 1 year ago

had this question today

upvoted 3 times

✉ **RanjithK** 1 year, 3 months ago

Selected Answer: A

Answer is A

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

A is correct..

upvoted 1 times

✉ **LaxmanTiwari** 1 year, 4 months ago

Selected Answer: A

For further evidence... <https://cloud.google.com/compute/docs/instances/managing-instance-access>

upvoted 2 times

✉ **pfabio** 1 year, 5 months ago

Selected Answer: A

A is correct and recommended option.

D is incorrect because block project-wide restrict access to this instance, evidence:

<https://cloud.google.com/compute/docs/connect/restrict-ssh-keys>
<https://cloud.google.com/compute/docs/connect/restrict-ssh-keys>

upvoted 1 times

 **juliusali** 1 year, 5 months ago

D. Enable block project wide keys for the instance. Generate an SSH key and associate the key with that instance. Distribute the key to dev1 users and direct them to use their third-party tools to connect.

upvoted 1 times

You need to produce a list of the enabled Google Cloud Platform APIs for a GCP project using the gcloud command line in the Cloud Shell. The project name is my-project. What should you do?

- A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.
- B. Run gcloud init to set the current project to my-project, and then run gcloud services list --available.
- C. Run gcloud info to view the account value, and then run gcloud services list --account <Account>.
- D. Run gcloud projects describe <project ID> to verify the project value, and then run gcloud services list --available.

Correct Answer: A

Community vote distribution

A (88%)

13%

✉️  **dan80**  3 years, 4 months ago

A is the correct answer, log to gcloud and run the commands, doesn't make sense to run cloud init and gcloud services list --available gives you the full services that are available.

upvoted 48 times

✉️  **raffiq** 3 years, 3 months ago

Yes, Answer A correct. it shows only enabled services of API

upvoted 6 times

✉️  **lxgywil**  2 years, 5 months ago

"A" is correct.

For those, who have doubts:

`gcloud services list --available` returns not only the enabled services in the project but also services that CAN be enabled. Therefore, option B is incorrect.

<https://cloud.google.com/sdk/gcloud/reference/services/list#--available>

upvoted 12 times

✉️  **squishy_fishy** 2 years, 4 months ago

Best answer!

upvoted 1 times

✉️  **sabrinakloud**  5 months, 3 weeks ago

Selected Answer: A

Answer: A

upvoted 1 times

✉️  **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: A

Let's do a side-by-side analysis of Answer A and Answer D to clear our doubts:

Answer A: Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.

This option first retrieves the project ID by running the gcloud projects list command.

Then, it uses the gcloud services list command with the --project flag to list the enabled APIs for the specified project.

upvoted 3 times

✉️  **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Answer D: Run gcloud projects describe <project ID> to verify the project value, and then run gcloud services list --available.

This option uses the gcloud projects describe command with the project ID to retrieve information about the specified project, including the project ID.

Then, it uses the gcloud services list command with the --available flag to list all available APIs, not just the ones that are enabled for the specified project.

Based on the scenario presented in the question, we want to produce a list of the enabled APIs for a GCP project, NOT a list of all available APIs. Therefore, Answer A is more appropriate because it specifically lists the enabled APIs for the specified project.

Answer D lists all available APIs which may include APIs that are not enabled in the project, which could cause confusion or unnecessary information.

Therefore, Answer A is the correct option in this case.

upvoted 3 times

✉️  **Bobbybash** 7 months, 4 weeks ago

Selected Answer: A

D. Run gcloud projects describe <project ID> to verify the project value, and then run gcloud services list --available.

To produce a list of enabled Google Cloud Platform APIs for a GCP project using the gcloud command line, you can first run gcloud projects describe <project ID> to verify the project ID for the project in question. Then, you can run gcloud services list --available to list all the available APIs and see which ones are enabled for the project. This command shows the full list of services and their status, including whether they are enabled, disabled, or ready for use. Option A is incorrect because it lists all the available services, regardless of whether they are enabled or not. Option B is incorrect because it lists only the available services, which might not be enabled in the project. Option C is incorrect because it shows account information and not service information.

upvoted 1 times

 **Charumathi** 1 year ago

A is correct answer,

Run the following command to list the enabled APIs and services in your current project:

gcloud services list

whereas, Run the following command to list the APIs and services available to you in your current project:

gcloud services list --available

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is the correct answer

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for A

upvoted 1 times

 **luciorifa** 1 year, 7 months ago

Selected Answer: A

A is the correct answer

upvoted 2 times

 **sharmaishu** 1 year, 7 months ago

A is the correct answer.

<https://cloud.google.com/sdk/gcloud/reference/services/list#--available>

--available

Return the services available to the project to enable. This list will include any services that the project has already enabled.

To list the services the current project has enabled for consumption, run:

gcloud services list --enabled

To list the services the current project can enable for consumption, run:

gcloud services list --available

upvoted 2 times

 **[Removed]** 1 year, 9 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **vishnukumartr** 1 year, 10 months ago

A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.

upvoted 1 times

 **shawnkkk** 1 year, 10 months ago

A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.

upvoted 1 times

 **Litan** 2 years ago

A Correct

upvoted 2 times

 **mcaromit** 2 years, 5 months ago

A is correct

upvoted 2 times

 **Rabbitfur** 2 years, 5 months ago

A is correct answer --project option available in gcloud command

upvoted 1 times

 **[Removed]** 2 years, 6 months ago

A is correct. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.

upvoted 2 times

You are building a new version of an application hosted in an App Engine environment. You want to test the new version with 1% of users before you completely switch your application over to the new version. What should you do?

- A. Deploy a new version of your application in Google Kubernetes Engine instead of App Engine and then use GCP Console to split traffic.
- B. Deploy a new version of your application in a Compute Engine instance instead of App Engine and then use GCP Console to split traffic.
- C. Deploy a new version as a separate app in App Engine. Then configure App Engine using GCP Console to split traffic between the two apps.
- D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

Correct Answer: D

Community vote distribution

D (100%)

 **SIX**  3 years, 4 months ago

Correct answer is D
upvoted 53 times

 **mohdafiuddin**  2 years, 9 months ago

Splitting the question to the key requirements

1. new version of an application hosted in an App Engine environment.
2. test the new version with 1% of users

App engine supports versioning and traffic splitting so no need to involve anything else
(source - <https://cloud.google.com/appengine#all-features>)

- A.'Google Kubernetes Engine'.... - No need to involve GKE. Not the right option
- B.'Compute Engine instance'.... - No need to involve Compute Engine.
- C.'Separate app in App Engine'....- No need to deploy as a separate app. versioning is supported already. Not the right option.
- D. This is the right answer.

upvoted 29 times

 **akshaychavan7** 1 year, 4 months ago

Just to add, for option C you cannot have two applications deployed inside an app engine project. In order to do so, you need to create the application inside a new project.
So, we just eliminate option C.

upvoted 5 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: D

D is correct , because you cannot create a sepearte app in the same app engine
upvoted 1 times

 **Neha_Pallavi** 1 month, 2 weeks ago

D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly
upvoted 1 times

 **trainingexam** 3 months, 1 week ago

Selected Answer: D

App engine provides out of box functionality to split the traffic between multiple versions
upvoted 1 times

 **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: D

Answer D is the correct answer
upvoted 1 times

 **ashtonez** 7 months ago

Selected Answer: D

D correct, no more than 1 app engine per project GKE or CE doesnt make sense, the right approach is deploying new version and splitting traffic
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 3 weeks ago

Selected Answer: D

Answer D is the correct answer as App Engine is designed to allow you to deploy multiple versions of the same application and route traffic between them. To test the new version with 1% of users, you can deploy the new version alongside the current version and then use

the App Engine Traffic Splitting feature to gradually increase the percentage of users who are routed to the new version. This can be done easily using GCP Console.

Answers A and B are not the optimal solutions as Kubernetes Engine and Compute Engine do not offer the same level of built-in traffic splitting and routing features as App Engine.

Answer C is also POSSIBLE but may not be the best approach since deploying a separate app requires additional configuration and maintenance. It is simpler to deploy multiple versions of the same application and use App Engine Traffic Splitting to route traffic between them.

upvoted 1 times

 **cslince** 10 months ago

Selected Answer: D

Correct answer is D

upvoted 1 times

 **cslince** 10 months ago

Selected Answer: D

Correct answer is D

upvoted 1 times

 **leogor** 11 months ago

Selected Answer: D

D. Deploy a new version of your application in App Engine

upvoted 1 times

 **gcpBeginner** 1 year ago

without any doubt answer is D

upvoted 1 times

 **learn_GCP** 1 year ago

Selected Answer: D

D is correct

A. no need to use GKE

B. Compute engine not-relevant.

C. cant deploy 2 applications in one project.

D. --splits & --split-by flags are available for gcloud app deploy.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

D is correct

upvoted 1 times

 **mplibunao** 1 year, 4 months ago

Selected Answer: D

D is the correct answer

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for D

upvoted 1 times

 **vishnukumartr** 1 year, 10 months ago

D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 3 times

You need to provide a cost estimate for a Kubernetes cluster using the GCP pricing calculator for Kubernetes. Your workload requires high IOPs, and you will also be using disk snapshots. You start by entering the number of nodes, average hours, and average days. What should you do next?

- A. Fill in local SSD. Fill in persistent disk storage and snapshot storage.
- B. Fill in local SSD. Add estimated cost for cluster management.
- C. Select Add GPUs. Fill in persistent disk storage and snapshot storage.
- D. Select Add GPUs. Add estimated cost for cluster management.

Correct Answer: C

Reference:

<https://cloud.google.com/products/calculator#tab=container>

Community vote distribution

A (95%) 5%

✉ **dan80** Highly Voted 3 years, 4 months ago

This one is Tricky, local SSD is require for High IOPS - <https://cloud.google.com/compute/docs/disks/local-ssd> , but it say using disk snapshots. A is correct.

upvoted 58 times

✉ **poogcp** Highly Voted 3 years, 4 months ago

A is correct .

upvoted 16 times

✉ **Captain1212** Most Recent 1 month, 1 week ago

Selected Answer: A

A is correct , as ssd requires the High IOPS

upvoted 1 times

✉ **trainingexam** 3 months, 1 week ago

Selected Answer: A

A is correct

upvoted 1 times

✉ **sakdip66** 6 months ago

Selected Answer: A

since the requirement is high IOPs Local SSD is our best option. that makes:

C and D as irrelevant.

B Add estimated cost for cluster management is not related to storage requirement mentioned in the scenario

upvoted 1 times

✉ **Zahir1004** 7 months ago

Dear friends,

Great responses in question.

Can someone with contributor access, send me the remaining questions to this email rihazs@gmail.com

upvoted 2 times

✉ **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: A

To provide a cost estimate for a Kubernetes cluster using the GCP pricing calculator for Kubernetes, after entering the number of nodes, average hours, and average days, you should fill in the required storage and snapshot details.

Given that your workload requires high IOPs and will also be using disk snapshots, the appropriate option would be;

- A. Fill in local SSD. Fill in persistent disk storage and snapshot storage.

upvoted 2 times

✉ **chikorita** 5 months, 3 weeks ago

my lord! you're always right

upvoted 1 times

✉ **NosFerazi** 8 months ago

Selected Answer: B

Why do I need to fill out persistent disk storage and snapshot storage, it is already populated.
filling out local ssd should suffice. going with B

upvoted 1 times

✉ **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

If persistent disk storage and snapshot storage are already populated in the GCP pricing calculator, you do not need to fill them out again. In that case, selecting local SSD and adding an estimated cost for cluster management, as suggested in Answer B, would be sufficient.

However, it is important to note that the cost estimate may not be accurate if any of the details in the GCP pricing calculator are incorrect or do not match your requirements. Therefore, it is always a good practice to review all the details and ensure that they are accurate and up to date before finalizing the cost estimate.

In summary, if persistent disk storage and snapshot storage are already populated, and you only require local SSD and an estimated cost for cluster management, then Answer B is a valid choice. Still, I go with Answer A my friend. :)

upvoted 3 times

 **cslince** 10 months ago

Selected Answer: A

A is correct .

upvoted 1 times

 **leogor** 11 months ago

A. local SSD for high iops

upvoted 1 times

 **Racinely** 11 months, 4 weeks ago

The question is about storage then A is the correct answer.

upvoted 1 times

 **PKookNN** 12 months ago

Selected Answer: A

for high IOPS needs local ssd and snapshot wording is clear enough so A is correct

upvoted 1 times

 **gcpBeginner** 1 year ago

we should not use GPU in this case. SSD is correct solution for storage . correct answer is A

upvoted 2 times

 **sandipk91** 1 year, 1 month ago

Selected Answer: A

Option A is correct

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is correct, Local SSD provides high performance.

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for A

upvoted 1 times

 **LaxmanTiwari** 1 year, 4 months ago

Selected Answer: A

This one is Tricky, local SSD is require for High IOPS - <https://cloud.google.com/compute/docs/disks/local-ssd> , but it say using disk snapshots. A is correct.

upvoted 3 times

You are using Google Kubernetes Engine with autoscaling enabled to host a new application. You want to expose this new application to the public, using HTTPS on a public IP address. What should you do?

- A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.
- B. Create a Kubernetes Service of type ClusterIP for your application. Configure the public DNS name of your application using the IP of this Service.
- C. Create a Kubernetes Service of type NodePort to expose the application on port 443 of each node of the Kubernetes cluster. Configure the public DNS name of your application with the IP of every node of the cluster to achieve load-balancing.
- D. Create a HAProxy pod in the cluster to load-balance the traffic to all the pods of the application. Forward the public traffic to HAProxy with an iptable rule. Configure the DNS name of your application using the public IP of the node HAProxy is running on.

Correct Answer: A

Reference:

<https://cloud.google.com/kubernetes-engine/docs/tutorials/http-balancer>

Community vote distribution

A (100%)

 **arsav**  2 years, 7 months ago

HAProxy is HTTP only, doesn't support HTTPS, so you can reject option D
<https://www.haproxy.org/#desc>

Cluster IP - is an internal IP, you cannot expose public externally. reject option B

out of option A and C

C, port 443 is https but public DNS is not going to give you a load balancing
A is the right choice,
kubernets ingress exposes HTTPS
<https://kubernetes.io/docs/concepts/services-networking/ingress/>

and cloud load balancer is the right choice which will help to expose the app to public
upvoted 44 times

 **NoniGeorge** 1 year, 9 months ago

Pretty sure that option D works more from on premise then cloud because with cloud you pretty much don't have to configure your ip tables !
upvoted 1 times

 **dan80**  3 years, 4 months ago

A is correct.

upvoted 33 times

 **magistrum** 2 years, 9 months ago

Saw this which provides good context <https://medium.com/google-cloud/kubernetes-nodeport-vs-loadbalancer-vs-ingress-when-should-i-use-what-922f010849e0>
upvoted 18 times

 **nitinz** 2 years, 7 months ago

you nailed it.
upvoted 2 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: A

option A is correct as you need load balancing and in option c dns will not give you load balancing
upvoted 1 times

 **frantishk** 1 month, 3 weeks ago

I didnt know, that ClusterIP is an internal IP and you cannot expose public externally..

Thanks !

upvoted 1 times

 **trainingexam** 3 months, 1 week ago

Selected Answer: A

upvoted 1 times

✉ **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: A

To expose a new application hosted on Google Kubernetes Engine with autoscaling enabled to the public using HTTPS on a public IP address, the most appropriate option would be;

A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 1 times

✉ **GS300** 8 months, 3 weeks ago

Selected Answer: A

A works and is correct, but service type should be ClusterIP

upvoted 1 times

✉ **cslince** 10 months ago

Selected Answer: A

A is correct .

upvoted 2 times

✉ **gcpBeginner** 1 year ago

Correct answer is A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 1 times

✉ **Cornholio_LMC** 1 year ago

had this question today

upvoted 1 times

✉ **NOOGLeR** 1 year, 3 months ago

A is the correct

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

This is A without any second thought.

upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago

Go for A

upvoted 1 times

✉ **BalajiMBala** 1 year, 5 months ago

option A is an correct answer

upvoted 1 times

✉ **cysteine** 1 year, 6 months ago

Selected Answer: A

A is correct. Because, you want to access Global HTTP(s) Load Balancing with GKE services, you must use Ingress Object.

<https://cloud.google.com/kubernetes-engine/docs/concepts/service>

upvoted 3 times

✉ **dev_sami** 1 year, 6 months ago

Selected Answer: A

option B - WRONG --> cluster IP is an internal IP, so we cannt expose publically.

option C - WRONG-->port 443 is HTTPS but public DNS is not provide load balancing.

option D - WRONG -->HAProxy is HTTP only not HTTPS

so

A is right option.

upvoted 9 times

✉ **vishnukumartr** 1 year, 10 months ago

A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

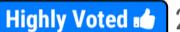
upvoted 2 times

You need to enable traffic between multiple groups of Compute Engine instances that are currently running two different GCP projects. Each group of Compute Engine instances is running in its own VPC. What should you do?

- A. Verify that both projects are in a GCP Organization. Create a new VPC and add all instances.
- B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.
- C. Verify that you are the Project Administrator of both projects. Create two new VPCs and add all instances.
- D. Verify that you are the Project Administrator of both projects. Create a new VPC and add all instances.

Correct Answer: B*Community vote distribution*

B (100%)

-  **dan80**  3 years, 4 months ago
B - <https://cloud.google.com/vpc/docs/shared-vpc>
upvoted 33 times
-  **glam**  2 years, 12 months ago
B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.
upvoted 9 times
-  **trainingexam**  3 months, 1 week ago
Selected Answer: B
shared-vpc is right option
upvoted 1 times
-  **sabrinakloud** 5 months, 3 weeks ago
Selected Answer: B
Use shared VPC
upvoted 1 times
-  **Partha117** 6 months, 2 weeks ago
Selected Answer: B
Shared VPC is the correct choice here
upvoted 1 times
-  **Buruguduystunstugudunstuy** 7 months, 2 weeks ago
Selected Answer: B
To enable traffic between multiple groups of Compute Engine instances running in different VPCs of different GCP projects, the best option would be;
B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.
upvoted 3 times
-  **hems4all** 9 months, 3 weeks ago
Selected Answer: B
B correct
upvoted 1 times
-  **cslince** 10 months ago
Selected Answer: B
B is correc
upvoted 1 times
-  **Zoze** 10 months, 3 weeks ago
Selected Answer: B
B is correct because is the concept of the shared VPC.
upvoted 1 times
-  **darcal95** 1 year ago
ok, is B, but that means that the VMs in the "other project" have to change their ip?
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

It is shared VPC concept. Go with B
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

Shared VPC allows an organization to connect resources from multiple projects to a common Virtual Private Cloud (VPC) network, so that they can communicate with each other securely and efficiently using internal IPs from that network. When you use Shared VPC, you designate a project as a host project and attach one or more other service projects to it. The VPC networks in the host project are called Shared VPC networks. Eligible resources from service projects can use subnets in the Shared VPC network

upvoted 6 times

 **mplibunao** 1 year, 4 months ago

Selected Answer: B

shared vpc is the answer
upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for B
upvoted 1 times

 **JohnnyStudy** 1 year, 6 months ago

B. You use "Shared VPC Network" to share a network across several projects in your GCP organization. You designate a project as a host project, and attach one or more other service projects to it. Only works within the same organization. Only works across projects.
upvoted 2 times

 **vishnukumartr** 1 year, 10 months ago

B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.
upvoted 1 times

 **shawnkkk** 1 year, 10 months ago

B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.
upvoted 1 times

 **Litan** 2 years ago

B is correct
upvoted 1 times

You want to add a new auditor to a Google Cloud Platform project. The auditor should be allowed to read, but not modify, all project items. How should you configure the auditor's permissions?

- A. Create a custom role with view-only project permissions. Add the user's account to the custom role.
- B. Create a custom role with view-only service permissions. Add the user's account to the custom role.
- C. Select the built-in IAM project Viewer role. Add the user's account to this role.
- D. Select the built-in IAM service Viewer role. Add the user's account to this role.

Correct Answer: C

Reference:

<https://cloud.google.com/resource-manager/docs/access-control-project>

Community vote distribution



✉️ **cloudenthu01** Highly Voted 3 years, 3 months ago

C is correct
roles/Viewer role provides access to all resources under the projects but do not alter the state of these resources
upvoted 41 times

✉️ **mav3r1ck** 1 year, 2 months ago

It should be A.

https://cloud.google.com/iam/docs/faq#when_would_i_use_basic_roles
When would I use basic roles?

You can use basic roles in development and test environments, where it might be appropriate for some principals to have wide-ranging permissions. Avoid basic roles in production environments.

upvoted 3 times

✉️ **jrlsl1991** 8 months, 1 week ago

But in this case we're not asked to follow any best practices. Besides, the help article says "In production environments, do not grant basic roles unless there is no alternative.", and in this case there's no alternative since we need to grant access to all resources.
upvoted 1 times

✉️ **mav3r1ck** 1 year, 2 months ago

Principle of least privilege
upvoted 1 times

✉️ **creativenets** 3 months, 3 weeks ago

i disagree.
upvoted 1 times

✉️ **glam** Highly Voted 2 years, 12 months ago

C. Select the built-in IAM project Viewer role. Add the user's account to this role.
upvoted 13 times

✉️ **jayjani66** Most Recent 2 months, 3 weeks ago

Answer is C. Select the built-in IAM project Viewer role. Add the user's account to this role.

The IAM project Viewer role is a built-in role in Google Cloud that provides read-only access to all resources within a project. This role allows users to view project items, configurations, and metadata but does not grant any permission to modify or make changes to the resources.

upvoted 1 times

✉️ **trainingexam** 3 months, 1 week ago

Selected Answer: A

with principle of leastprivilege should be A
Also, question is asking to set permission on single project. Basic principles grants permissions on all project.
upvoted 1 times

✉️ **Jeevan4433** 5 months, 3 weeks ago

It is option A. I just referred here!
<https://cloud.google.com/iam/docs/roles-overview>
Caution: Basic roles include thousands of permissions across all Google Cloud services. In production environments, do not grant basic roles unless there is no alternative. Instead, grant the most limited predefined roles or custom roles that meet your needs
upvoted 2 times

✉️ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: A

upvoted 1 times

✉ **Jelly_Wang** 5 months, 4 weeks ago

Selected Answer: C

Go for C.

The debate is between A and C. From auditor accessibility perspective they are the same, but from practical perspective C is the only option. For people who vote for A, you must never work with auditors in an enterprise level project. There are hundred if not thousands of permission you need to set one by one if you create custom role by yourself. And they will come to you and ask for permission every single day. And this is an "there's no alternative" situation where using Basic role is practical.

upvoted 6 times

✉ **ashtonez** 7 months ago

Selected Answer: C

I would go with C, A and C are equally correct, with principle of leastprivilege should be A, with recomendation of not using custom roles becasue they are not maintained by gcp it should be C, since its not stating its a production env its a little bit ambiguous

upvoted 1 times

✉ **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: C

To allow the new auditor to read, but not modify, all project items in a Google Cloud Platform project, the best option would be;

C. Select the built-in IAM project Viewer role. Add the user's account to this role.

upvoted 3 times

✉ **xaqanik** 7 months, 2 weeks ago

Selected Answer: A

Google recommends:

Basic roles include thousands of permissions across all Google Cloud services. In production environments, do not grant basic roles unless there is no alternative. Instead, grant the most limited predefined roles or custom roles that meet your needs.

upvoted 1 times

✉ **Vladimir_Sakhonchik** 7 months, 2 weeks ago

I didn't find in question about production

upvoted 1 times

✉ **xaqanik** 7 months ago

auditor works on a production project. nothing to do them in test and development projects

upvoted 4 times

✉ **Nazz1977** 8 months, 1 week ago

Selected Answer: A

Avoid basic roles in production environments.

upvoted 2 times

✉ **jrisl1991** 8 months, 1 week ago

Selected Answer: C

It must be C. Primitive/Basic roles are meant to be used as well when universal access is required. Creating a custom role would imply an excessive amount of work to make sure the auditor can view all the resources.

upvoted 2 times

✉ **Nazz1977** 8 months, 2 weeks ago

Selected Answer: A

Avoid using basic roles except when absolutely necessary. These roles are very powerful, and include a large number of permissions across all Google Cloud services.

<https://cloud.google.com/resource-manager/docs/access-control-project>

upvoted 1 times

✉ **riccamini** 10 months, 1 week ago

C is the one the makes the most sense. But I did not find a pre-defined "projectViewer" role. Probably is referring to the basic role "Viewer" applied at the project level (which is still not encouraged, but better than creating a custom role)

upvoted 2 times

✉ **skyblue07** 10 months, 1 week ago

For those that say A is correct? Does the "project view-only" (or similar) exists?

upvoted 3 times

✉ **leogor** 11 months ago

Selected Answer: C

built-in IAM project Viewer role

upvoted 1 times

✉ **PKookNN** 11 months ago

A is more correct (best practice is to avoid basic roles)

upvoted 2 times

You are operating a Google Kubernetes Engine (GKE) cluster for your company where different teams can run non-production workloads. Your Machine Learning (ML) team needs access to Nvidia Tesla P100 GPUs to train their models. You want to minimize effort and cost. What should you do?

- A. Ask your ML team to add the `accelerator: gpu` annotation to their pod specification.
- B. Recreate all the nodes of the GKE cluster to enable GPUs on all of them.
- C. Create your own Kubernetes cluster on top of Compute Engine with nodes that have GPUs. Dedicate this cluster to your ML team.
- D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the `cloud.google.com/gke-accelerator: nvidia-tesla-p100` nodeSelector to their pod specification.**

Correct Answer: B

Community vote distribution

D (100%)

✉ **John_lam** Highly Voted 3 years, 4 months ago

D is the correct answer.

<https://cloud.google.com/kubernetes-engine/docs/how-to/gpus>
upvoted 46 times

✉ **tablet444** 3 years, 2 months ago

the documentation states "Limitations
Before using GPUs on GKE, keep in mind the following limitations:

You cannot add GPUs to existing node pools.
GPU nodes cannot be live migrated during maintenance events."
upvoted 10 times

✉ **nightflyer** 2 years, 10 months ago

In this case it is about adding a GPU enabled node pool not a GPU to an existing node-pool
upvoted 15 times

✉ **fragment137** 10 months, 2 weeks ago

You're correct that D says that, except that the question also says to use the most cost-effective method. Two node-pools would be more expensive than rebuilding the current one with GPU enabled.
upvoted 1 times

✉ **Gulithor** 9 months ago

It also says to minimize effort, wouldn't recreating all the pools take way longer than just adding 1?
upvoted 1 times

✉ **glam** Highly Voted 2 years, 12 months ago

D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the `cloud.google.com/gke-accelerator: nvidia-tesla-p100` nodeSelector to their pod specification.
upvoted 14 times

✉ **trainingexam** Most Recent 3 months, 1 week ago

Selected Answer: D

D is the correct answer.
upvoted 1 times

✉ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: D

Option D is good
upvoted 1 times

✉ **sakdip66** 6 months ago

Selected Answer: D

Creating new node pool w/ GPU-enabled instances is cost - saving solution. This way ML team workload will GPU instance for their ML and other team workload will run smoothly
upvoted 1 times

✉ **Prat25200607** 6 months, 1 week ago

Selected Answer: D

D makes more cost effective
upvoted 1 times

Buruguduystunstugudunstuy 7 months, 2 weeks ago

Selected Answer: D

Answer D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the cloud.google.com/gke-accelerator: nvidia-tesla-p100 nodeSelector to their pod specification.

Adding a new node pool with GPUs is the best option because it allows for a separate set of nodes that can be specifically allocated to workloads that require GPU acceleration, such as the Machine Learning (ML) team's workloads. This approach will not affect other workloads running on the original nodes, keeping the costs low and the overall cluster performance stable.

upvoted 4 times

cslince 10 months ago

Selected Answer: D

D is the correct answer.

upvoted 1 times

leogor 11 months ago

D is correct

upvoted 1 times

raghu09 1 year ago

Selected Answer: D

D is correct

upvoted 1 times

iadarsh 1 year ago

Selected Answer: D

D is correct

Because if you create entirely new node pool then its not cost efficient and also the pods which not require that much high GPU is get scheduled into it. So instead of that add a new node pool with GPU and in the pod YAML file mention the node affinity to get scheduled into the GPU enabled node pool.

upvoted 2 times

AzureDP900 1 year, 3 months ago

By looking at all answers first 3 can be eliminated without any second thought. D is correct.

upvoted 2 times

haroldbenites 1 year, 4 months ago

Go for D

upvoted 1 times

LaxmanTiwari 1 year, 4 months ago

Selected Answer: D

In this case it is about adding a GPU enabled node pool not a GPU to an existing node-pool

upvoted 2 times

luciorifa 1 year, 7 months ago

Selected Answer: D

D is the correct answer

upvoted 1 times

jaffarali 1 year, 10 months ago

Selected Answer: D

D would be the right option when there is possibility to add GPUs without recreating the nodes.

upvoted 3 times

ME_MYSELF 1 year, 10 months ago

I have a doubt. Will these questions repeat in ace exam or not??Please answer if u have written the exam...it will be very helpful

upvoted 5 times

ahsangh 1 year, 9 months ago

from past experience yes, haven't written this one yet, will revert after writing in a few days.

upvoted 3 times

Your VMs are running in a subnet that has a subnet mask of 255.255.255.240. The current subnet has no more free IP addresses and you require an additional 10 IP addresses for new VMs. The existing and new VMs should all be able to reach each other without additional routes. What should you do?

- A. Use gcloud to expand the IP range of the current subnet.
- B. Delete the subnet, and recreate it using a wider range of IP addresses.
- C. Create a new project. Use Shared VPC to share the current network with the new project.
- D. Create a new subnet with the same starting IP but a wider range to overwrite the current subnet.

Correct Answer: C

Community vote distribution



✉ **JustLearning** Highly Voted 3 years, 4 months ago

A: Expand the existing subnet.
<https://cloud.google.com/sdk/gcloud/reference/compute/networks/subnets/expand-ip-range>
 upvoted 52 times

✉ **glam** Highly Voted 2 years, 12 months ago

A. Use gcloud to expand the IP range of the current subnet.
 upvoted 11 times

✉ **Captain1212** Most Recent 1 month, 1 week ago

Selected Answer: A
 A is correct , as you need to expand it
 upvoted 1 times

✉ **trainingexam** 3 months, 1 week ago

Selected Answer: A
`gcloud compute networks subnets expand-ip-range NAME --prefix-length=PREFIX_LENGTH [--region=REGION] [GCLOUD_WIDE_FLAG ...]`
 upvoted 1 times

✉ **dasgcp** 6 months, 2 weeks ago

Selected Answer: C
 You can't expand the subnet because the question states "The current subnet has no more free IP addresses", correct answer is C.
 upvoted 3 times

KerolesKhaili 4 months ago

that's exactly why you expand the subnet =D
 upvoted 3 times

✉ **Jimut** 6 months, 3 weeks ago

Selected Answer: A
 Use gcloud to expand the IP range of the current subnet.
 upvoted 1 times

✉ **Mobin92** 7 months, 1 week ago

Selected Answer: A
 A. Use gcloud to expand the IP range of the current subnet.
 upvoted 1 times

✉ **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: A
 Answer A (Use gcloud to expand the IP range of the current subnet): This option is correct because it allows you to expand the primary IP range of the existing subnet to accommodate the additional 10 IP addresses required for the new VMs. This can be done without deleting or recreating the subnet, which saves time and avoids disrupting any existing resources that are using the subnet.
 upvoted 3 times

✉ **Bobbybash** 7 months, 4 weeks ago

Selected Answer: A
 A
 To get an additional 10 IP addresses for new VMs running in a subnet that has no more free IP addresses and where existing and new VMs should be able to reach each other without additional routes, you should use gcloud to expand the IP range of the current subnet. This can be achieved by modifying the IP range of the subnet with the gcloud command-line tool. You can use the following command to increase the range of the subnet:

```
gcloud compute networks subnets expand [SUBNET_NAME] --range=[NEW_IP_RANGE]
```

Make sure to specify the name of the subnet you want to modify, and choose a new IP range that includes the current range and 10 additional IP addresses.

upvoted 2 times

✉ **ChillinBoy** 10 months ago

Selected Answer: A

A 100%

upvoted 1 times

✉ **cslince** 10 months ago

Selected Answer: A

A: Expand the existing subnet.

upvoted 1 times

✉ **leogor** 11 months ago

Selected Answer: A

expand the IP range of the current subnet

upvoted 1 times

✉ **dennydream** 11 months, 2 weeks ago

Why in the world would C be the answer? A is def the answer.

upvoted 1 times

✉ **Charumathi** 1 year ago

A. is correct,

gcloud compute networks subnets expand-ip-range - expand the IP range of a Compute Engine subnetwork

gcloud compute networks subnets expand-ip-range NAME --prefix-length=PREFIX_LENGTH [--region=REGION] [GCLOUD_WIDE_FLAG ...]

upvoted 1 times

✉ **gcpBeginner** 1 year ago

correct answer is A. Use gcloud to expand the IP range of the current subnet.

upvoted 1 times

✉ **sandipk91** 1 year, 1 month ago

Selected Answer: A

Option A is the answer

upvoted 1 times

✉ **Shlok27_gcloud** 1 year ago

Hello can anyone send me the questions to my mail : avinash.shlok0@gmail.com I have exam in next week.

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

A is right answer

upvoted 1 times

Your organization uses G Suite for communication and collaboration. All users in your organization have a G Suite account. You want to grant some G Suite users access to your Cloud Platform project. What should you do?

- A. Enable Cloud Identity in the GCP Console for your domain.
- B. Grant them the required IAM roles using their G Suite email address.
- C. Create a CSV sheet with all users' email addresses. Use the gcloud command line tool to convert them into Google Cloud Platform accounts.
- D. In the G Suite console, add the users to a special group called cloud-console-users@yourdomain.com. Rely on the default behavior of the Cloud Platform to grant users access if they are members of this group.

Correct Answer: B

Reference:

<https://cloud.google.com/resource-manager/docs/creating-managing-organization>

Community vote distribution

B (100%)

 **austinl**  3 years, 4 months ago

B is correct

upvoted 30 times

 **Ciumela**  3 years, 3 months ago

B is correct: To actively adopt the Organization resource, the G Suite or Cloud Identity super admins need to assign the Organization Administrator Cloud IAM role to a user or group

upvoted 20 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: B

B seems more correct as per thhe google pratcices

upvoted 2 times

 **trainingexam** 3 months, 1 week ago

Selected Answer: B

B is correct

upvoted 1 times

 **zwwdplay** 6 months, 1 week ago

Dear friends,

Great responses in question.

Can someone with contributor access, send me the remaining questions to this email: zwwdplay@hotmail.com

upvoted 1 times

 **Zahir1004** 7 months ago

Dear friends,

Great responses in question.

Can someone with contributor access, send me the remaining questions to this email rihazs@gmail.com

upvoted 2 times

 **Buruguduystunstugudunsty** 7 months, 2 weeks ago

Selected Answer: B

Answer B. Grant them the required IAM roles using their G Suite email address.

To grant G Suite users access to a Cloud Platform project, you should use their G Suite email addresses to grant them the required IAM roles.

Answer A is incorrect because enabling Cloud Identity is not necessary for granting G Suite users access to a Cloud Platform project. Cloud Identity provides a centralized identity management system for G Suite and Cloud Platform, but it is not required for this use case.

Answer C is incorrect because there is no need to convert G Suite email addresses into Google Cloud Platform accounts. G Suite users already have Google accounts and can be granted access to Cloud Platform using their G Suite email addresses.

Answer D is incorrect because there is no default behavior in the Cloud Platform to grant access to users who are members of a particular group. Access to Cloud Platform resources is granted based on IAM roles and policies, not group membership.

upvoted 6 times

 **leogor** 11 months ago

Selected Answer: B

Grant them the required IAM roles using their G Suite email address

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is right

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for B

upvoted 1 times

 **Majkl93** 1 year, 7 months ago

Selected Answer: B

B as per the comments

upvoted 2 times

 **vishnukumartr** 1 year, 10 months ago

B. Grant them the required IAM roles using their G Suite email address.

upvoted 2 times

 **alaahakim** 1 year, 10 months ago

B is Correct

upvoted 1 times

 **bubblegumbeach** 1 year, 12 months ago

B is correct

upvoted 2 times

 **mcaromit** 2 years, 4 months ago

B is correct

upvoted 2 times

 **Christiank** 2 years, 6 months ago

B. Grant them the required IAM roles using their G Suite email address.

upvoted 2 times

 **[Removed]** 2 years, 6 months ago

B is correct. Grant them the required IAM roles using their G Suite email address.

upvoted 3 times

You have a Google Cloud Platform account with access to both production and development projects. You need to create an automated process to list all compute instances in development and production projects on a daily basis. What should you do?

- A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.
- B. Create two configurations using gsutil config. Write a script that sets configurations as active, individually. For each configuration, use gsutil compute instances list to get a list of compute resources.
- C. Go to Cloud Shell and export this information to Cloud Storage on a daily basis.
- D. Go to GCP Console and export this information to Cloud SQL on a daily basis.

Correct Answer: A*Community vote distribution*

A (100%)

-  **cludenthu01** Highly Voted 3 years, 3 months ago
A is correct
upvoted 38 times
 -  **glam** Highly Voted 2 years, 12 months ago
A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.
upvoted 14 times
 -  **Captain1212** Most Recent 1 month, 1 week ago
Selected Answer: A
A is correct , first list , then activate it
upvoted 2 times
 -  **Neha_Pallavi** 1 month, 2 weeks ago
A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.
upvoted 1 times
 -  **trainingexam** 3 months, 1 week ago
Selected Answer: A
Activate each config and list the instances
upvoted 1 times
 -  **Partha117** 6 months, 2 weeks ago
Selected Answer: A
A is correct
upvoted 1 times
 -  **Buruguduystunstugudunstuy** 7 months, 2 weeks ago
Selected Answer: A
Answer A is the correct answer.
- The most straightforward way to list all compute instances in development and production projects is to use gcloud compute instances list command. However, since the account has access to both production and development projects, it's necessary to create two configurations with different project IDs.
- Answer B is incorrect because gsutil is used for object storage operations and not compute instances. (DISTRACTOR)
- Answers C and D are incorrect because they do not provide a straightforward solution for listing compute instances in multiple projects.
upvoted 4 times
-  **Shwom** 8 months, 1 week ago
Selected Answer: A
A is correct
upvoted 1 times
 -  **cslince** 10 months ago
Selected Answer: A
A is correct
upvoted 1 times
 -  **leogor** 11 months ago

Selected Answer: A

gcloud instead of gsutil

upvoted 1 times

 **[Removed]** 1 year ago

Selected Answer: A

"gcloud" can create and manage Google Cloud resources while "gsutil" cannot do so. "gsutil" can manipulate buckets, bucket's objects and bucket ACLs on GCS(Google Cloud Storage) while "gcloud" cannot do so

upvoted 3 times

 **Shlok27_gcloud** 1 year ago

Which one is correct in all cases , suggested one or community one . Im confused totally for all questions

upvoted 2 times

 **habros** 1 year, 2 months ago

A. Gsutil is used for cloud storage bucket

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right, This is part of Tutorial Dojo practice questions

upvoted 4 times

 **haroldbenites** 1 year, 4 months ago

Go for A

upvoted 1 times

 **cysteine** 1 year, 6 months ago

Selected Answer: A

Obviously A is correct

upvoted 1 times

 **RegisFTM** 1 year, 9 months ago

This looks to be a multiple-choice question. The answer A is correct, and the C completes the task... Does it make sense?

upvoted 1 times

 **Zufair** 1 year, 9 months ago

C is not automated task

A is correct

upvoted 3 times

 **vishnukumartr** 1 year, 10 months ago

A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.

upvoted 2 times

You have a large 5-TB AVRO file stored in a Cloud Storage bucket. Your analysts are proficient only in SQL and need access to the data stored in this file. You want to find a cost-effective way to complete their request as soon as possible. What should you do?

- A. Load data in Cloud Datastore and run a SQL query against it.
- B. Create a BigQuery table and load data in BigQuery. Run a SQL query on this table and drop this table after you complete your request.
- C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.
- D. Create a Hadoop cluster and copy the AVRO file to NDFS by compressing it. Load the file in a hive table and provide access to your analysts so that they can run SQL queries.

Correct Answer: C*Community vote distribution*

C (100%)

✉  **mohdafiuddin**  2 years, 9 months ago

Breaking down the question into key points -

1. 5-TB AVRO file stored in a Cloud Storage bucket.
2. Analysts are proficient only in SQL
3. cost-effective way to complete their request as soon as possible

A.Load data in Cloud Datastore... (Not Correct because Cloud Datastore is not a good option to run SQL Queries)

B. ...Load data in BigQuery.... (Not Cost Effective because loading the data which is already present in the bucket into BigQuery again is expensive)

C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.

(This is the right answer as it meets all the requirements from the question)

D. Create a Hadoop cluster and copy the AVRO file to NDFS by compressing it. Load the file in a hive table and provide access to your analysts so that they can run SQL queries.

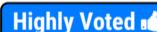
(Too roundabout and indirect. Not the right option)

upvoted 139 times

✉  **pondai** 2 years, 6 months ago

listem this guy

upvoted 14 times

✉  **Ciumela**  3 years, 3 months ago

C is correct: <https://cloud.google.com/bigquery/external-data-sources>

upvoted 23 times

✉  **Captain1212**  1 month, 1 week ago

Selected Answer: C

C is the right answer , as

upvoted 1 times

✉  **Neha_Pallavi** 1 month, 2 weeks ago

C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.

upvoted 1 times

✉  **Neha_Pallavi** 1 month, 2 weeks ago

C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.

upvoted 1 times

✉  **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: C

Answer C is the most cost-effective and efficient way to provide analysts access to the data stored in the 5-TB AVRO file in Cloud Storage.

Here's why:

You can create external tables in BigQuery that point to the 5-TB AVRO file stored in Cloud Storage. External tables allow you to query data stored in Cloud Storage without the need to load the data into BigQuery. This is a cost-effective way to provide your analysts' access to the data they need, and it is also an efficient solution since you can run SQL queries on the data directly in BigQuery.

upvoted 4 times

✉ **D_ee28** 1 month, 4 weeks ago

Hi, if someone has PDF for all the Q/A, please send it to deek2805@gmail.com.
Thank you in advance.

upvoted 1 times

✉ **Emmanski08** 9 months, 1 week ago

External tables in BigQuery

upvoted 1 times

✉ **cslince** 10 months ago

Selected Answer: C

C is correct: <https://cloud.google.com/bigquery/external-data-sources>

upvoted 1 times

✉ **leogor** 11 months ago

Selected Answer: C

external tables in BigQuery

upvoted 1 times

✉ **Untamables** 11 months, 3 weeks ago

Selected Answer: C

Similar to Athena

upvoted 1 times

✉ **Charumathi** 1 year ago

C. is correct,
An external data source is a data source that you can query directly from BigQuery, even though the data is not stored in BigQuery storage.

BigQuery supports the following external data sources:

Amazon S3
Azure Storage
Cloud Bigtable
Cloud Spanner
Cloud SQL
Cloud Storage
Drive

upvoted 2 times

✉ **fifi1907** 1 year ago

answer is c

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

mohdafiuddin explanation is very detailed .. C is right

upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 2 times

✉ **JelloMan** 1 year, 5 months ago

Selected Answer: C

@mohdafiuddin covered it exactly but heres more info:
<https://cloud.google.com/bigquery/docs/loading-data-cloud-storage-avro>

upvoted 2 times

✉ **luciorifa** 1 year, 7 months ago

Selected Answer: C

C is the correct answer

upvoted 2 times

✉ **ericyev** 1 year, 10 months ago

C is good. <https://cloud.google.com/bigquery/external-data-cloud-storage>

upvoted 1 times

You need to verify that a Google Cloud Platform service account was created at a particular time. What should you do?

- A. Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.
- B. Filter the Activity log to view the Configuration category. Filter the Resource type to Google Project.
- C. Filter the Activity log to view the Data Access category. Filter the Resource type to Service Account.
- D. Filter the Activity log to view the Data Access category. Filter the Resource type to Google Project.

Correct Answer: D

Community vote distribution

A (100%)

 **John_lam**  3 years, 4 months ago

Correct Answer is A.

Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.

upvoted 76 times

 **mlantonis** 3 years, 4 months ago

I agree A

upvoted 9 times

 **shafiqeee1**  3 years, 2 months ago

A - I reproduced in my project.

upvoted 26 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: A

A is the correct answer, you can see it in configuration category

upvoted 2 times

 **Neha_Pallavi** 1 month, 2 weeks ago

A. Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: A

Answer A. Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.

The Activity log is the primary tool for viewing and analyzing activity within a Google Cloud project, including all Service Account-related activity. By filtering the Activity log to view the Configuration category and filtering the Resource type to Service Account, you can see when a Service Account was created, updated, or deleted, along with other related metadata such as the user who performed the action and the IP address from which the action was performed.

upvoted 3 times

 **nishant7290** 9 months, 2 weeks ago

Selected Answer: A

Correct Answer is A.

upvoted 1 times

 **ChillinBoy** 10 months ago

Selected Answer: A

A 100%

upvoted 1 times

 **cslince** 10 months ago

Selected Answer: A

Correct Answer is A.

upvoted 1 times

 **ArjunKennedy** 10 months, 1 week ago

Thank God, we have a discussion community and not just the solutions from this site lol.

upvoted 9 times

 **Zoze** 10 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

 **leogor** 11 months ago

A. Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.
upvoted 1 times

 **sausalito** 11 months ago

Selected Answer: A

Answer is A

upvoted 1 times

 **anolive** 11 months, 1 week ago

Selected Answer: A

for sure, i simulated here.

upvoted 2 times

 **alleinallein** 1 year ago

Selected Answer: A

A key word is configuration

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Selected Answer: A

Answer is A. Tried and tested

upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Answer is A

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right ..

upvoted 1 times

You deployed an LDAP server on Compute Engine that is reachable via TLS through port 636 using UDP. You want to make sure it is reachable by clients over that port. What should you do?

- A. Add the network tag allow-udp-636 to the VM instance running the LDAP server.
- B. Create a route called allow-udp-636 and set the next hop to be the VM instance running the LDAP server.
- C. Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.
- D. Add a network tag of your choice to the instance running the LDAP server. Create a firewall rule to allow egress on UDP port 636 for that network tag.

Correct Answer: C*Community vote distribution*

C (100%)

✉  **kishoredeena**  3 years, 3 months ago

Option C is the right one
upvoted 32 times

✉  **cloudenthalu01**  3 years, 3 months ago

C is correct
You tag the instances ,then create ingress firewall rules to allow udp on desired port for target-tags name applied to instances
upvoted 23 times

✉  **Captain1212**  1 month, 1 week ago

Selected Answer: C
C is correct bcoz of ingress
upvoted 1 times

✉  **Partha117** 6 months, 2 weeks ago

Selected Answer: C
Firewall rule for ingress is correct
upvoted 1 times

✉  **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: C
Answer C is correct: Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.

To make sure that the LDAP server is reachable by clients over port 636 using UDP, you need to allow ingress traffic on that port. You can achieve this by adding a network tag to the instance running the LDAP server and then creating a firewall rule that allows ingress traffic on UDP port 636 for that network tag.

upvoted 2 times

✉  **ratnesh_uk01** 9 months, 2 weeks ago

can anyone please suggest why D is not correct? thanks
upvoted 1 times

✉  **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Answer D is incorrect because adding a network tag of your choice to the instance running the LDAP server and creating a firewall rule to allow egress traffic on UDP port 636 for that network tag would not allow incoming traffic on that port. You need to create a firewall rule that allows ingress traffic on that port.

upvoted 1 times

✉  **leogor** 11 months ago

Selected Answer: C
allow ingress
upvoted 1 times

✉  **GA2022** 1 year, 1 month ago

In this site I can see only 92 questions , can someone have access to 93-195 questions ?
Please mail me gajare.amar2022@gmail.com
upvoted 2 times

✉  **AzureDP900** 1 year, 3 months ago

C is right.
upvoted 1 times

✉  **haroldbenites** 1 year, 4 months ago

upvoted 1 times

upvoted 1 times

✉ **Fayi** 1 year, 8 months ago

You are developing a new web application that will be deployed on Google Cloud Platform. As part of your release cycle, you want to test updates to your application on a small portion of real user traffic. The majority of the users should still be directed towards a stable version of your application. What should you do?

- A. Deploy me application on App Engine For each update, create a new version of the same service Configure traffic splitting to send a small percentage of traffic to the new version
- B. Deploy the application on App Engine For each update, create a new service Configure traffic splitting to send a small percentage of traffic to the new service.
- C. Deploy the application on Kubernetes Engine For a new release, update the deployment to use the new version
- D. Deploy the application on Kubernetes Engine For a now release, create a new deployment for the new version Update the service e to use the now deployment.

upvoted 4 times

✉ **akhun** 8 months ago

A is correct

upvoted 1 times

✉ **shank910** 1 year, 8 months ago

A create new version

upvoted 1 times

✉ **akshaychavan7** 1 year, 4 months ago

A without any doubt.

upvoted 1 times

✉ **Teyo** 1 year, 8 months ago

Selected Answer: C

C is correct as you can use tags and then set firewall rules for instances with such tag.

upvoted 2 times

✉ **Vidyaji** 1 year, 10 months ago

C IS PERFECT

upvoted 2 times

✉ **vishnukumartr** 1 year, 10 months ago

C. Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.

upvoted 2 times

✉ **alaahakim** 1 year, 10 months ago

C is Correct

upvoted 2 times

✉ **devopsbatch** 2 years, 4 months ago

C A tag is simply a character string added to a tags field in a resource, such as Compute Engine virtual machine (VM) instances or instance templates. A tag is not a separate resource, so you cannot create it separately. All resources with that string are considered to have that tag. Tags enable you to make firewall rules and routes applicable to specific VM instances.

upvoted 2 times

✉ **mcaromit** 2 years, 4 months ago

C is correct

upvoted 1 times

You need to set a budget alert for use of Compute Engine services on one of the three Google Cloud Platform projects that you manage. All three projects are linked to a single billing account. What should you do?

- A. Verify that you are the project billing administrator. Select the associated billing account and create a budget and alert for the appropriate project.
- B. Verify that you are the project billing administrator. Select the associated billing account and create a budget and a custom alert.
- C. Verify that you are the project administrator. Select the associated billing account and create a budget for the appropriate project.
- D. Verify that you are project administrator. Select the associated billing account and create a budget and a custom alert.

Correct Answer: B*Community vote distribution*

A (100%)

✉  **kishoredeena**  3 years, 4 months ago

I think the answer is A, You can rely on default alert. No need for custom alert
upvoted 45 times

✉  **BobbyFlash** 1 year, 11 months ago

I agree. If I'm not wrong, project admin doesn't have billing permissions so C and D discarded. Between A and B, option B looks like it works but we would be creating a budget and alert receiving info about billing as a whole; so A delimits billing for the project you want to get info from.

upvoted 7 times

✉  **alan9999** 3 years ago

Right its not asking to set custom alert
upvoted 1 times

✉  **Eshkrkrkr** 2 years, 11 months ago

One point - there is no such role as Project Billing Administrator - it should be Project Billing Manager but he can't create budgets, the only one who can - Billing Account Administrator. Nor Project Administrator exists. Very tricky question, maybe the option a wrong, hope smb will catch it on exam and pass some light on real variants.
<https://cloud.google.com/iam/docs/understanding-roles#billing-roles>

upvoted 6 times

✉  **Wachy** 2 years, 9 months ago

Eshkrkrkr read the question calmly. The role there is Billing Administrator. Not Project Billing Administrator.

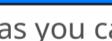
It's more like: "Verify you are the project; billing administrator"

upvoted 11 times

✉  **ryumada** 1 year, 2 months ago

more like: ""Verify you are the project's billing administrator""

upvoted 6 times

✉  **Ciumela**  3 years, 3 months ago

A is correct, as you can set a default alert also on a single project: <https://cloud.google.com/billing/docs/how-to/budgets>
upvoted 17 times

✉  **mwoodm** 3 years ago

Agreed. Per the link included: "To create a budget for your Cloud Billing account, you must be a Billing Account Administrator on the Cloud Billing account." So that eliminates C & D. Then no need for custom alert, eliminating B. The answer is A.
upvoted 3 times

✉  **nickyshil** 1 year, 2 months ago

why nobody is talking about "set a budget alert for use of Compute Engine services" only.. why not custom alert ?how default alert ?
upvoted 2 times

✉  **Captain1212**  1 month, 1 week ago

Selected Answer: A

A is correct because, there is default alert, no need of custom alert

upvoted 2 times

✉  **Neha_Pallavi** 1 month, 2 weeks ago

A. Verify that you are the project billing administrator. Select the associated billing account and create a budget and alert for the appropriate project.

upvoted 1 times

✉  **jayjani66** 2 months, 3 weeks ago

Answer is B.

Explanation: In Google Cloud, budget alerts are associated with billing accounts, not individual projects. Since all three projects are linked to a single billing account, you need to be the billing administrator to set up a budget and alert for that billing account.

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: A

Answer A. Verify that you are the project billing administrator. Select the associated billing account and create a budget and alert for the appropriate project.

In this scenario, you need to create a budget alert for the use of Compute Engine services on a specific project. Since all three projects are linked to a single billing account, you need to make sure that you are the billing administrator for that account. Once verified, you can create a budget and alert for the specific project by selecting the associated billing account and setting the budget and alert for the appropriate project.

upvoted 4 times

 **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

INCORRECT

Answer B is incorrect because a custom alert is not necessary for this scenario. A budget alert alone is sufficient to notify you when your spending reaches a certain threshold.

Answer C is incorrect because, while a project administrator can create a budget for the project, they cannot set a budget alert. Only a billing administrator has the necessary permissions to create a budget alert.

Answer D is incorrect because a project administrator cannot create a custom alert on the associated billing account. Custom alerts can only be created by billing administrators.

upvoted 1 times

 **jrls1991** 8 months, 1 week ago

Selected Answer: A

I don't even think there's an option for custom budget alert since all budget alerts are kind of the same and we can only customize (with the actual word "customize") the recipients. A should be correct.

upvoted 1 times

 **[Removed]** 10 months, 1 week ago

I think it's B because you can choose how to get alerts, by mail, slack, phone, etc.

upvoted 1 times

 **SK0710** 10 months, 2 weeks ago

Compute Engineer services, please read as Compute Engine services. Ans is A

upvoted 1 times

 **Zoze** 10 months, 3 weeks ago

I will choice B, It's clear in the question that he only wants to have a specific budget for the VMs instance and have alert of the cost of these instances go out the budget.

upvoted 1 times

 **Charumathi** 1 year ago

A. is correct,

You can define the scope of the budget. For example, you can scope the budget to apply to the spend in an entire Cloud Billing account, or narrow the scope to one or more projects, and/or one or more services, and/or other budget filters applicable to your Cloud Billing account.

upvoted 1 times

 **Adeem1981** 1 year, 2 months ago

Selected Answer: A

Billing Administrator is in between organization and projects, So he can chose and select appropriate project.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

You need billing admin not a project admin .. A is right

upvoted 1 times

 **Aksher** 1 year, 4 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for A

upvoted 1 times

 **oaP12** 1 year, 5 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **browneyes1985** 1 year, 6 months ago

Selected Answer: A

A is correct

upvoted 1 times

You are migrating a production-critical on-premises application that requires 96 vCPUs to perform its task. You want to make sure the application runs in a similar environment on GCP. What should you do?

- A. When creating the VM, use machine type n1-standard-96.
- B. When creating the VM, use Intel Skylake as the CPU platform.
- C. Create the VM using Compute Engine default settings. Use gcloud to modify the running instance to have 96 vCPUs.
- D. Start the VM using Compute Engine default settings, and adjust as you go based on Rightsizing Recommendations.

Correct Answer: B*Community vote distribution*

A (95%) 5%

 **dan80** Highly Voted 3 years, 4 months ago

A is correct - <https://cloud.google.com/compute/docs/machine-types>
upvoted 50 times

 **Ahmed_Y** 1 month, 3 weeks ago

Indeed, there is a n1-standard-96 machine type in the machine types list here <https://cloud.google.com/compute/docs/general-purpose-machines>
upvoted 1 times

 **glam** Highly Voted 2 years, 12 months ago

A. When creating the VM, use machine type n1-standard-96.
upvoted 12 times

 **Rajeshpaspi** Most Recent 1 week, 3 days ago

A is the correct answer
upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: A
A is correct, use machine type n1-standard-96 while creating the VM
upvoted 2 times

 **sakdip66** 6 months ago

Selected Answer: A
the goal is to have an equivalent of this app in GCP. therefore A is the best shot we have
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: A
A. When creating the VM, use machine type n1-standard-96.

Answer A is the correct answer as it directly addresses the requirement to have 96 vCPUs by selecting the n1-standard-96 machine type. This machine type offers 96 vCPUs, 360 GB of memory, and up to 2,400 GB of local SSD storage.

<https://cloud.google.com/compute/docs/machine-resource>

Answer B is incorrect because selecting a CPU platform alone will not guarantee the availability of the required number of vCPUs.

Answer C is incorrect because it is not possible to modify a running Compute Engine instance to add vCPUs. vCPUs can only be added or removed during instance creation or by stopping the instance first.

Answer D is incorrect because while Rightsizing Recommendations can help optimize compute resources, they will not guarantee that the application has the required 96 vCPUs to function properly.

upvoted 5 times

 **cslince** 10 months ago

Selected Answer: A
A is correct - <https://cloud.google.com/compute/docs/machine-types>
upvoted 2 times

 **fragment137** 10 months, 2 weeks ago

the instance name for 96 vcpu N1 is "n1-highcpu-96", not n1-standard-96.

Possible that has been updated since this question came out?

upvoted 3 times

Selected Answer: A

keyword: n1-standard-96

upvoted 1 times

 **Charumathi** 1 year ago

A. is correct,
N1
* CPU types --> Skylake, Broadwell, Haswell, Sandy Bridge, and Ivy Bridge
* Architecture --> x86
* vCPUs 1 to 96
upvoted 1 times

 **Sivaprakash674** 1 year, 2 months ago

Go with A, It is the Correct one .

upvoted 1 times

 **Adeem1981** 1 year, 2 months ago

Selected Answer: B

SKYLAKE IS 96VCPU. B is correct.

upvoted 1 times

 **vividso** 1 year, 1 month ago

<https://cloud.google.com/compute/docs/machine-types>

N1 machine series have up to 96 vCPUs, 6.5 GB of memory per vCPU, and are available on Intel Sandy Bridge, Ivy Bridge, Haswell, Broadwell, and Skylake CPU platforms.

upvoted 3 times

 **RanjithK** 1 year, 3 months ago

Selected Answer: A

Go with A

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

<https://cloud.google.com/compute/docs/general-purpose-machines> - A is correct

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for A

<https://cloud.google.com/compute/all-pricing?hl=es>

upvoted 2 times

 **LaxmanTiwari** 1 year, 4 months ago

Selected Answer: A

Why not update to A .. doesn't look like a rocket science to me.

upvoted 1 times

 **Lowballed** 1 year, 8 months ago

A is correct.

upvoted 2 times

You want to configure a solution for archiving data in a Cloud Storage bucket. The solution must be cost-effective. Data with multiple versions should be archived after 30 days. Previous versions are accessed once a month for reporting. This archive data is also occasionally updated at month-end. What should you do?

- A. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Coldline Storage.
- B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.
- C. Add a bucket lifecycle rule that archives data from regional storage after 30 days to Coldline Storage.
- D. Add a bucket lifecycle rule that archives data from regional storage after 30 days to Nearline Storage.

Correct Answer: B

Reference:

<https://cloud.google.com/storage/docs/managing-lifecycles>

Community vote distribution

B (100%)

 **neelesh88**  3 years, 3 months ago

B is correct

upvoted 32 times

 **ESP_SAP**  3 years, 1 month ago

Correct Answer (B):

NumberOfNewerVersions

The NumberOfNewerVersions condition is typically only used in conjunction with Object Versioning. If the value of this condition is set to N, an object version satisfies the condition when there are at least N versions (including the live version) newer than it. For a live object version, the number of newer versions is considered to be 0. For the most recent noncurrent version, the number of newer versions is 1 (or 0 if there is no live object version), and so on.

Important: When specifying this condition in a .json configuration file, you must use numNewerVersions instead of NumberOfNewerVersions.

<https://cloud.google.com/storage/docs/lifecycle#numberofnewerversions>

upvoted 24 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: B

B is the right answer, because of data is accessing infrequently and nearline storage is good for it

upvoted 2 times

 **SanjeevKumar1983** 1 month, 1 week ago

Selected Answer: B

B is correct

upvoted 1 times

 **jayjani66** 2 months, 3 weeks ago

Correct ans is A.

Explanation: In this scenario, you need to archive data after 30 days, which implies that the data with multiple versions is considered for archiving. Since you need to access previous versions once a month for reporting, using Coldline Storage is the most cost-effective option.
upvoted 1 times

 **Partha117** 6 months, 2 weeks ago

Selected Answer: B

since accessed frequently it will be nearline

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: B

Answer B, adding a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage, is the correct answer for this scenario.

Nearline Storage is designed for data that is accessed less frequently, such as for backup and archival purposes. It has a minimum storage duration of 30 days, which makes it suitable for archiving data that needs to be kept for a long time but is accessed infrequently. Additionally, Nearline Storage has lower storage costs than Coldline Storage, making it more cost-effective for this use case.

By adding a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage, you can ensure that the data is automatically moved to a more cost-effective storage class while still being easily accessible for reporting purposes.

upvoted 6 times

 **dnur** 6 months, 2 weeks ago

You're incorrect. Coldline storage has a lower costs than Nearline Storage. <https://cloud.google.com/storage/docs/storage-classes>.
upvoted 1 times

 **chikorita** 5 months, 3 weeks ago

just FYI that my lord, @Buruguduystunstugudunstuy, is always right!!!!1
upvoted 3 times

 **eaakgul** 6 months, 2 weeks ago

The question tells us that the previous versions are accessed once a month for reporting. So, nearline makes more sense in this case. 'Buruguduystunstugudunstuy' has mentioned that nearline has lower storage costs for only this 'use case'
upvoted 1 times

 **leogor** 11 months ago

Selected Answer: B

archives data with newer versions after 30 days to Nearline Storage.
upvoted 1 times

 **kadc** 1 year ago

Selected Answer: B

B should be correct:

Nearline has min storage of 30 days, while Coldline has 90 days.

Since "archive data is also occasionally updated at month-end", updating object before min storage period is allowed but causes early deletion fees as if the object was stored for the min duration, so using Coldline will always charge for 90 days and not likely to save cost.

<https://cloud.google.com/storage/pricing#early-delete>

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is right and straight forward.
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

regional storage after 30 days to Nearline Storage option is trick you :)
upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for B
upvoted 1 times

 **Rukman** 1 year, 6 months ago

Selected Answer: B

B is correct
upvoted 1 times

 **Vidyaji** 1 year, 10 months ago

B is perfect
upvoted 2 times

 **vishnukumartr** 1 year, 10 months ago

B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.
upvoted 1 times

 **shawnkkk** 1 year, 10 months ago

B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.
upvoted 1 times

 **Chotebhaisahab** 2 years, 1 month ago

Agree B is the correct
upvoted 2 times

 **Rahul183** 2 years, 1 month ago

B- <https://cloud.google.com/storage/docs/storage-classes>
upvoted 1 times

Your company's infrastructure is on-premises, but all machines are running at maximum capacity. You want to burst to Google Cloud. The workloads on Google Cloud must be able to directly communicate to the workloads on-premises using a private IP range. What should you do?

- A. In Google Cloud, configure the VPC as a host for Shared VPC.
- B. In Google Cloud, configure the VPC for VPC Network Peering.
- C. Create bastion hosts both in your on-premises environment and on Google Cloud. Configure both as proxy servers using their public IP addresses.
- D. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.

Correct Answer: D

Community vote distribution

D (100%)

✉  **SIX**  3 years, 4 months ago

I believe D is the right answer
upvoted 55 times

✉  **dan80** 3 years, 4 months ago

B is correct - <https://cloud.google.com/solutions/best-practices-vpc-design>. this answer also on all machines are running at maximum capacity.
upvoted 2 times

✉  **xharf** 1 year, 2 months ago

"Google Cloud VPC Network Peering allows internal IP address connectivity across two Virtual Private Cloud (VPC) networks regardless of whether they belong to the same project or the same organization."
<https://cloud.google.com/vpc/docs/vpc-peering>

while

"Cloud Interconnect provides low latency, high availability connections that enable you to reliably transfer data between your on-premises and Google Cloud Virtual Private Cloud (VPC) networks."
<https://cloud.google.com/network-connectivity/docs/interconnect/concepts/overview>

and

"HA VPN is a high-availability (HA) Cloud VPN solution that lets you securely connect your on-premises network to your VPC network through an IPsec VPN connection in a single region."
<https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview>

so, cloud vpn is the best answer for the question requirement

upvoted 5 times

✉  **JustLearning** 3 years, 4 months ago

vpc network peering does not connect to on-prem. Cloud VPN is the correct solution.
<https://cloud.google.com/vpn/docs/concepts/overview>
upvoted 25 times

✉  **mlantonis** 3 years, 4 months ago

You need VPN, so D is the correct. VPC network peering is between VPCs.
upvoted 14 times

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (D):

Access internal IPs directly

Your VPC network's internal (RFC 1918) IP addresses are directly accessible from your on-premises network with peering, no NAT device or VPN tunnel required.

Hybrid made easy

Today's business climate demands flexibility. Connecting your on-premises resources to your cloud resources seamlessly, with minimum latency or interruption, is a business-critical requirement. The speed and reliability of Cloud Interconnect lets you extend your organization's data center network into Google Cloud, simply and easily, while options such as Cloud VPN provide flexibility for all your workloads. This unlocks the potential of hybrid app development and all the benefits the cloud has to offer.

In the graphic below: What GCP Connection is right for you? shows clearly what is the method for extend your on premise network (IP Private communication).

What GCP Connection is right for you?

<https://cloud.google.com/hybrid-connectivity>

upvoted 34 times

✉  **Captain1212**  1 month, 1 week ago

Selected Answer: D

D is the right answer as they need the private range and the machine are also on high working load
upvoted 2 times

 **rosapersiani** 4 months, 2 weeks ago

Selected Answer: D

d is right
upvoted 1 times

 **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: D

VPN to connect your on-premise network to the cloud
upvoted 1 times

 **Partha117** 6 months, 2 weeks ago

Selected Answer: D

VPN for on premise connection to GCP
upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: D

Answer D. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.

To burst into Google Cloud from the on-premises infrastructure, a VPN connection can be established between the on-premises network and Google Cloud. VPN provides a secure, private tunnel to transfer data between on-premises infrastructure and Google Cloud. Cloud VPN would allow workloads on Google Cloud to communicate with workloads on-premises over private IP addresses, making it a suitable option for this scenario.

Answer A (Shared VPC) and Answer B (VPC Network Peering) do not address the requirement of communicating over a private IP range between on-premises and Google Cloud.

Answer C (bastion hosts) involves the use of public IP addresses, which may not be suitable for a private, secure connection.
upvoted 4 times

 **cslince** 10 months ago

Selected Answer: D

D is the right answer
upvoted 1 times

 **leogor** 11 months ago

Selected Answer: D

Cloud VPN
upvoted 1 times

 **VaneA** 1 year ago

Selected Answer: D

It is the answer
upvoted 1 times

 **RanjithK** 1 year, 3 months ago

Selected Answer: D

Go with D
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

D is right answer
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

Cloud VPN is way to establish connection between on prem to cloud . D is correct.
upvoted 2 times

 **Uqqasha** 1 year, 8 months ago

Selected Answer: D

D is the right answer.
upvoted 2 times

 **ARVII** 1 year, 9 months ago

On-premise -> GCP
There are 2 ways
1. Cloud VPN
2. Interconnect
Since we have VPN as an option, others is not recommended
upvoted 3 times

 **Vidyaji** 1 year, 10 months ago

D is perfect

upvoted 1 times

 **vishnukumartr** 1 year, 10 months ago

D. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.

upvoted 1 times

You want to select and configure a solution for storing and archiving data on Google Cloud Platform. You need to support compliance objectives for data from one geographic location. This data is archived after 30 days and needs to be accessed annually. What should you do?

- A. Select Multi-Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Coldline Storage.
- B. Select Multi-Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Nearline Storage.
- C. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Nearline Storage.
- D. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Coldline Storage.

Correct Answer: D*Community vote distribution*

D (100%)

✉  **SIX**  3 years, 4 months ago

D

Google Cloud Coldline is a new cold-tier storage for archival data with access frequency of less than once per year. Unlike other cold storage options, Nearline has no delays prior to data access, so now it is the leading solution among competitors.

upvoted 45 times

✉  **dan80** 3 years, 4 months ago

C is correct - This data is archived after 30 days - Nearline Storage 30 days , Coldline Storage 90 days

<https://cloud.google.com/storage/docs/storage-classes>

upvoted 12 times

✉  **lxgywil** 2 years, 5 months ago

The answer is D.

The main thing here is how often the data is retrieved. The question is saying that data needs to be accessed annually - i.e. once a year. Therefore, you should choose Coldline Storage, as it implies less frequent access than Nearline. (Archival Storage would fit even better but there's no such option)

From the link you provided:

"Nearline Storage is ideal for data you plan to read or modify on average once per month or less."

and

"Coldline Storage is ideal for data you plan to read or modify at most once a quarter. "

upvoted 18 times

✉  **nithinpb180** 3 years, 3 months ago

That is the minimum storage duration. I would go with D. Coldline storage is more suitable for infrequent data access.

upvoted 8 times

✉  **mlantonis** 3 years, 4 months ago

dan80 is right

upvoted 1 times

✉  **JustLearning** 3 years, 4 months ago

D is correct. Coldline is a better choice.

upvoted 9 times

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (D):

<https://cloud.google.com/storage/docs/storage-classes>

Nearline Storage

Nearline Storage is a low-cost, highly durable storage service for storing infrequently accessed data. Nearline Storage is a better choice than Standard Storage in scenarios where slightly lower availability, a 30-day minimum storage duration, and costs for data access are acceptable trade-offs for lowered at-rest storage costs.

Nearline Storage is ideal for data you plan to read or modify on average once per month or less. For example, if you want to continuously add files to Cloud Storage and plan to access those files once a month for analysis, Nearline Storage is a great choice.

Nearline Storage is also appropriate for data backup, long-tail multimedia content, and data archiving. Note, however, that for data accessed less frequently than once a quarter, Coldline Storage or Archive Storage are more cost-effective, as they offer lower storage costs.

<https://cloud.google.com/storage/docs/storage-classes#nearline>

upvoted 16 times

✉  **ESP_SAP** 3 years, 1 month ago

CORRECTION.
Correct Answer is (D):

The Real description is about Coldline storage Class:

Coldline Storage

Coldline Storage is a very-low-cost, highly durable storage service for storing infrequently accessed data. Coldline Storage is a better choice than Standard Storage or Nearline Storage in scenarios where slightly lower availability, a 90-day minimum storage duration, and higher costs for data access are acceptable trade-offs for lowered at-rest storage costs.

Coldline Storage is ideal for data you plan to read or modify at most once a quarter. Note, however, that for data being kept entirely for backup or archiving purposes, Archive Storage is more cost-effective, as it offers the lowest storage costs.

<https://cloud.google.com/storage/docs/storage-classes#coldline>

upvoted 13 times

 Captain1212 Most Recent 1 month, 1 week ago

Selected Answer: D

D is the correct answer, as the data is accessed only once a year
upvoted 2 times

 sabrinakloud 5 months, 3 weeks ago

Selected Answer: D

"This data is archived after 30 days and needs to be accessed annually"
ideally archive; coldline is the closest.
upvoted 1 times

 Elya 6 months, 1 week ago

The best option would be to select Regional Storage and add a bucket lifecycle rule that archives data after 30 days to Nearline Storage. Nearline Storage is designed for data that is accessed less frequently, but still needs to be readily available when accessed. It has a lower storage cost than Regional Storage, and retrieval costs are lower than those of Coldline Storage.
upvoted 1 times

 inbalinbal 6 months, 1 week ago

Selected Answer: D
D is correct
upvoted 2 times

 Buruguduystunstugudunstuy 7 months, 2 weeks ago

Selected Answer: D

Answer D is the CORRECT answer. The scenario mentioned in the question requires archiving data after 30 days and accessing it annually. As per the Cloud Storage documentation, Coldline storage is ideal for data that is accessed at most once a quarter. Hence, selecting regional storage and adding a bucket lifecycle rule that archives data after 30 days to Coldline Storage is the best solution to meet the compliance objectives and cost-effectiveness requirements.
upvoted 3 times

 Buruguduystunstugudunstuy 7 months, 2 weeks ago

INCORRECT:

Answer A, selecting Multi-Regional Storage and adding a bucket lifecycle rule that archives data after 30 days to Coldline Storage, is not a good fit for this scenario because Multi-Regional Storage is more expensive than Regional Storage and it does not provide a clear advantage for this use case.

Answer B, selecting Multi-Regional Storage and adding a bucket lifecycle rule that archives data after 30 days to Nearline Storage, is also not the best option because Nearline Storage is more appropriate for data that is accessed less than once a month, while in this scenario, the data needs to be accessed at least once a year.

Answer C, selecting Regional Storage and adding a bucket lifecycle rule that archives data after 30 days to Nearline Storage, is not ideal because Nearline Storage is more suitable for data that is accessed less than once a month. If the data is accessed only once a year, it might be more cost-effective to choose Coldline Storage instead.
upvoted 1 times

 cslince 10 months ago

Selected Answer: D

The answer is D.
upvoted 1 times

 leogor 11 months ago

Selected Answer: D

Regional Storage, Coldline Storage.
upvoted 1 times

 biren111 11 months, 3 weeks ago

As in question it is asking for "one geographic location". So multi region options A & B is eliminated. And Between C & D "D is correct" as data will be accessed once a year.
upvoted 1 times

 gcpBeginner 1 year ago

It's C because archive data is 30 days and Nearline storage supports that. <https://cloud.google.com/storage/docs/storage-classes>

upvoted 1 times

 **Cornholio_LMC** 1 year ago

had this question today

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

D is for sure.

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for D

"... and needs to be accessed annually"

Cold line is the better choice.

upvoted 1 times

 **Rukman** 1 year, 6 months ago

Selected Answer: D

D is the right answer.

upvoted 1 times

 **Uqqasha** 1 year, 8 months ago

Selected Answer: D

D is the right answer.

upvoted 1 times

 **Vidyaji** 1 year, 10 months ago

D is perfect

upvoted 2 times

Your company uses BigQuery for data warehousing. Over time, many different business units in your company have created 1000+ datasets across hundreds of projects. Your CIO wants you to examine all datasets to find tables that contain an employee_ssn column. You want to minimize effort in performing this task.

What should you do?

- A. Go to Data Catalog and search for employee_ssn in the search box.
- B. Write a shell script that uses the bq command line tool to loop through all the projects in your organization.
- C. Write a script that loops through all the projects in your organization and runs a query on INFORMATION_SCHEMA.COLUMNS view to find the employee_ssn column.
- D. Write a Cloud Dataflow job that loops through all the projects in your organization and runs a query on INFORMATION_SCHEMA.COLUMNS view to find employee_ssn column.

Correct Answer: D

Community vote distribution

A (100%)

 **poogcp**  3 years, 4 months ago

Its A.

upvoted 38 times

 **filco72**  3 years, 2 months ago

Correct is A.

I tested on my account following this procedure: <https://cloud.google.com/bigquery/docs/quickstarts/quickstart-web-ui?authuser=4>
I created a data set and through Data Catalog I easily and effortlessly searched for the column name "gender"

upvoted 30 times

 **Captain1212**  1 month, 1 week ago

Selected Answer: A

Answer is A, as it requires the less effort and other options are more time consuming and error prone

upvoted 1 times

 **deadsong** 5 months ago

The most efficient approach to identify tables that contain an employee_ssn column in BigQuery would be to query the INFORMATION_SCHEMA.COLUMNS view, which provides metadata about all columns in all tables in a given dataset. Therefore, options C and D are both possible solutions.

Option A, searching for the column name in Data Catalog, may not be efficient if there are too many datasets to search through manually.

Option B, writing a shell script to loop through all the projects in your organization, may work, but it would require more effort and time than options C and D. Also, it would be more error-prone since the script would need to handle authentication and authorization, handle exceptions and errors, and collect the results.

Therefore, options C and D are better choices, but option D, using Cloud Dataflow, might be overkill for this specific task. Option C, looping through all projects and querying INFORMATION_SCHEMA.COLUMNS view, is the simplest and most effective solution to minimize effort.

upvoted 1 times

 **AARUSHBATHINA** 5 months, 1 week ago

Can anyone send the entire question I am able to see till question 96 send at aarush.bathina@gmail.com

upvoted 1 times

 **Astro_123** 5 months, 2 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 96 questions only, if anyone has the entire questions please share them with my bindashish315@gmail.com address. I have exam on next week, pls share Thanks in advance. I would be forever grateful.

upvoted 1 times

 **Neeyo** 5 months, 2 weeks ago

Hi All, I have my GCP ACE exam scheduled for tomorrow. However, I am only being able to access 96 questions. Can anyone kindly share the entire list of questions as I have hardly anytime left before my exam. oniyi6@yahoo.com. Thank you all so much

upvoted 1 times

 **arnika98** 5 months, 1 week ago

Did you pass the exam? If so any questions from here came? Please let us know so that it will be helpful

upvoted 1 times

 **ahtasham007** 6 months ago

Can anyone send the entire question I am able to see till question 96 send at ahtashamfirdausi07@gmail.com

upvoted 1 times

 **jinks006** 6 months, 3 weeks ago

I have my exam tomorrow and I could only access 96 questions. If someone has the whole set of questions, can you please send it to ajmajinkya06@gmail.com?

upvoted 1 times

 **R4F** 7 months, 1 week ago

C is IMO the correct answer (<https://stackoverflow.com/questions/68746567/big-query-find-all-column-name-containing-surname-across-all-tables>)

upvoted 1 times

 **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: A

Answer A is the correct answer. Go to Data Catalog and search for employee_ssn in the search box.

Data Catalog is a fully managed and scalable metadata management service that allows you to discover, understand, and manage your data. It provides search functionality that allows you to search for datasets, tables, columns, and other metadata across your organization. Therefore, you can simply go to Data Catalog and search for "employee_ssn" in the search box to find all datasets that contain this column. This is the most efficient and straightforward solution to the problem.

Answers B, C, and D are not ideal solutions.

Answer B requires writing a shell script and using the bq command line tool to loop through all the projects, which is time-consuming and error-prone.

Answer C requires writing a script that loops through all the projects and runs a query on INFORMATION_SCHEMA.COLUMNS view, which is also time-consuming and error-prone.

Answer D involves writing a Cloud Dataflow job, which is unnecessary and OVERKILL for this simple task.

upvoted 5 times

 **Kirangm** 8 months, 1 week ago

The answer is A

upvoted 1 times

 **Ushnishm** 8 months, 3 weeks ago

Hello, I have my GCP ACE exam scheduled early next week. However, I am only being able to access 96 questions. Can anyone kindly share the entire list of questions as I have hardly anytime left before my exam.

upvoted 2 times

 **Rajiv467** 9 months ago

I am able to access 96 questions only, if anyone has the entire questions please share them on rajiv09467@gmail.com

upvoted 2 times

 **srasam** 9 months, 2 weeks ago

I am able to access 96 questions only, if anyone has the entire questions please share them on gcp1on1@gmail.com

upvoted 2 times

 **Alison** 8 months, 1 week ago

you should pay buddy.

upvoted 2 times

 **Vinaykas** 11 months ago

I am taking the exam soon, could someone please send me entire questions to my email id - vinaykastwar@gmail.com ?

upvoted 3 times

 **leogor** 11 months ago

Selected Answer: A

Data Catalog

upvoted 1 times

You create a Deployment with 2 replicas in a Google Kubernetes Engine cluster that has a single preemptible node pool. After a few minutes, you use kubectl to examine the status of your Pod and observe that one of them is still in Pending status:

```
$ kubectl get pods -l app=myapp
NAME                 READY   STATUS    RESTART   AGE
myapp-deployment-58ddbbb995-1p86m   0/1     Pending   0          9m
myapp-deployment-58ddbbb995-qjpkq   1/1     Running   0          9m
```

What is the most likely cause?

- A. The pending Pod's resource requests are too large to fit on a single node of the cluster.
- B. Too many Pods are already running in the cluster, and there are not enough resources left to schedule the pending Pod.
- C. The node pool is configured with a service account that does not have permission to pull the container image used by the pending Pod.
- D. The pending Pod was originally scheduled on a node that has been preempted between the creation of the Deployment and your verification of the Pods' status. It is currently being rescheduled on a new node.

Correct Answer: B

Community vote distribution

B (51%)

D (49%)

✉ **ESP_SAP** Highly Voted 3 years, 1 month ago

Correct Answer is (B):

Reasons for a Pod Status Pending:

Troubleshooting Reason #1: Not enough CPU

Troubleshooting Reason #2: Not enough memory

Troubleshooting Reason #3: Not enough CPU and memory

<https://managedkube.com/kubernetes/k8sbot/troubleshooting/pending/pod/2019/02/22/pending-pod.html>

upvoted 47 times

✉ **Linus11** 2 years, 5 months ago

The real crux of this question is the mention about "Pre-emptible Node pool". That need to take into consider while determining the answer. If we choose B, then the importance of "Pre-emptible node pool" is not there. Whether the node pool is pre-emptible or not, resource scarcity can lead to pending pods.

When we consider the mention of "Pre-emptible Node Poll" , then the answer is obviously D. if a pre-emptible Node get pre-empted there will be a delay in cluster to sync it.

Answer is D.

upvoted 26 times

✉ **Finger41** 2 years, 4 months ago

This is to throw you off, when there is insufficient resources for a Pod to stand up, then the status will equal pending :

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/#:~:text=If%20a%20Pod%20is%20stuck,be%20scheduled%20onto%20a%20node.&text=You%20don't%20have%20enough,new%20nodes%20to%20your%20cluster.>

upvoted 3 times

✉ **nssrr** 7 months, 2 weeks ago

if you look at snapshot the pod age is 9m and it is still pending. Hence, the correct answer is B.

upvoted 3 times

✉ **alexgrig** 2 years ago

Questions says "Single Node" at that case the second pod can't be in running state.

upvoted 4 times

✉ **MidhunJose** 1 year, 8 months ago

It says a single node pool, not a single node. Meaning there can be multiple nodes, right?

upvoted 9 times

✉ **brvinod** 1 year, 8 months ago

A node can have multiple pods. So that is not a problem.

upvoted 2 times

✉ **SSPC** 3 years, 1 month ago

I agree with you. The correct answer is B

upvoted 5 times

✉ **[Removed]** 3 years ago

D gives you the reason why the resource could not be available as it was preempted

upvoted 3 times

✉ **Finger41** 2 years, 4 months ago

Its in a deployment, the pod will be recreated. There is insufficient resources in the node, not because its preemptible but because there is no memory/cpu.....

upvoted 2 times

✉ **tavva_prudhvi** 2 years, 6 months ago

If it was preempted, then it has to be restarted right? then it will show its failing not pending, check the articles mentioned by ESP_SAP

upvoted 3 times

✉ **ashtonez** 7 months ago

No, it will show as pending initially while reallocating

upvoted 1 times

✉ **cloudenthalu01** Highly Voted 3 years, 3 months ago

D is correct as the node on which pod was scheduled to run was preempted & now this pod is scheduled to run on different preemptible node from the node-pool

upvoted 21 times

✉ **myuniqueusername** 2 years ago

Incorrect. There is a single preemptible instance, if it was preempted then both pods would show as 'Pending'. B is correct.

upvoted 6 times

✉ **ashtonez** 7 months ago

No, because one of the pods may run on another node that its still up

upvoted 1 times

✉ **obeythefist** 1 year, 6 months ago

> There is a single preemptible instance

Where does it say that? It doesn't. Don't make things up. There's a single pre-emptible node pool. A single pool is not the same as a single node.

upvoted 8 times

✉ **Captain1212** Most Recent 1 month, 1 week ago

Selected Answer: B

B , seems more correct as it dont have enough resources

upvoted 2 times

✉ **respawn** 1 month, 1 week ago

Selected Answer: B

D is a decoy meant to confuse you, answer is B

upvoted 2 times

✉ **jayjani66** 2 months, 3 weeks ago

option D: The pending Pod was originally scheduled on a node that has been preempted between the creation of the Deployment and your verification of the Pods' status. It is currently being rescheduled on a new node.

upvoted 1 times

✉ **geeroylenkins** 2 months, 3 weeks ago

Selected Answer: D

I'm going D

upvoted 1 times

✉ **trainingexam** 4 months, 1 week ago

Selected Answer: D

I bet my answer on "preemptible node pool" keyword on the problem statement.

upvoted 1 times

✉ **vivekvj** 5 months, 2 weeks ago

Selected Answer: B

Most likely because of unavailability of resources. Look at the age 9m. If the pod was deleted and being rescheduled, it will not take 9m.

upvoted 1 times

✉ **hanweiCN** 5 months, 3 weeks ago

Selected Answer: B

the "age" the same with the running pod and "restart" both are 0 , means, the containers in both pod never been restarted, the "pending" status pod is the created at the same time with the "running" status pod. they r both NO preempted or they r both preempted, in second scenario, the "running" status pod should have 1 in "restart" . so i will go with B

upvoted 1 times

✉ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: B

<https://kubernetes.io/docs/tasks/debug/debug-application/debug-pods/#debugging-pods>

answer B

upvoted 1 times

✉ **inbalinbal** 6 months, 1 week ago

Selected Answer: D

they have mention preemptible node pool

upvoted 1 times

✉ **ashtonez** 7 months ago

Selected Answer: D

D is Right , I think a lot of people discussing here has very low k8s experience .

Node pool is a term for referring one or more (usually 4 or more nodes) nodes, you can view nodes as servers that run containers (remember pods are one or more containers) .

When one of the nodes from the node pool is down because it's a preemptible, all of their pods go down and the kube controller (which is an element running on control plane) reallocates the pods to other nodes from their node pools and then those pods will show as running (if the nodes accomplish the resource quotas) .

So basically what happened here is that we just did a kubectl get pods command in the mid time where one of the nodes from the node pool went down and the cluster tried to reposition the pods into other nodes

upvoted 4 times

✉ **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Both Answer B and Answer D could be valid answers depending on the specific circumstances of the cluster and deployment.

If the node pool is not experiencing any preemptions, then Answer B is likely the correct answer, indicating that there are not enough resources left in the cluster to schedule the pending Pod. In this case, you could try increasing the resources available in the cluster or scaling down other workloads to free up resources for the pending Pod.

If preemptions are occurring in the node pool, then Answer D may be the correct answer, indicating that the pending Pod was originally scheduled on a preempted node and is being rescheduled on a new node. In this case, you may want to consider using a different type of node pool or adjusting your workload to better handle node preemptions.

upvoted 3 times

✉ **chikorita** 5 months, 3 weeks ago

agreed, my lord!

but i believe they particularly want to test us on the role of "preemptible nodes" in a GKE cluster from that pov; D weighs more

upvoted 2 times

✉ **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

So, which option is more likely to be the cause? It's hard to say without more information because GOOGLE loves to confuse us. If the cluster is running many Pods and the nodes are close to maximum capacity, Answer B is more likely. If the cluster is not heavily utilized and preemptible nodes are being used, Answer D is more likely.

upvoted 1 times

✉ **akhun** 8 months ago

Selected Answer: B

If a Pod is stuck in Pending it means that it can not be scheduled onto a node. Generally this is because there are insufficient resources of one type or another that prevent scheduling.

You don't have enough resources: You may have exhausted the supply of CPU or Memory in your cluster, in this case you need to delete Pods, adjust resource requests, or add new nodes to your cluster.

upvoted 1 times

✉ **Nazz1977** 8 months, 1 week ago

Selected Answer: B

There is a single preemptible instance, if it was preempted then both pods would show as 'Pending'. B is correct.

upvoted 1 times

✉ **Di4sa** 8 months, 1 week ago

Selected Answer: D

D is the correct answer.

A single preemptible node pool can have multiple nodes

<https://cloud.google.com/blog/products/containers-kubernetes/cutting-costs-with-google-kubernetes-engine-using-the-cluster-autoscaler-and-preemptible-vms>

upvoted 2 times

✉ **Di4sa** 8 months, 1 week ago

D is the correct answer.

A single preemptible node pool can have multiple nodes

<https://cloud.google.com/blog/products/containers-kubernetes/cutting-costs-with-google-kubernetes-engine-using-the-cluster-autoscaler-and-preemptible-vms>

upvoted 1 times

You want to find out when users were added to Cloud Spanner Identity Access Management (IAM) roles on your Google Cloud Platform (GCP) project. What should you do in the GCP Console?

- A. Open the Cloud Spanner console to review configurations.
- B. Open the IAM & admin console to review IAM policies for Cloud Spanner roles.
- C. Go to the Stackdriver Monitoring console and review information for Cloud Spanner.
- D. Go to the Stackdriver Logging console, review admin activity logs, and filter them for Cloud Spanner IAM roles.

Correct Answer: B*Community vote distribution*

D (100%)

✉️  **samvegas**  3 years, 2 months ago

Answer = D, I have simple rule; if metrics then Monitoring, if Auditing then Logging.

upvoted 74 times

✉️  **Meix**  3 years, 4 months ago

I think the answer is D

upvoted 44 times

✉️  **Anand2608** 2 years, 11 months ago

As per the Cloud Audit logs documentation.

upvoted 3 times

✉️  **Captain1212**  1 month, 1 week ago

Selected Answer: D

D , seems more correct as it shows you the history also

upvoted 1 times

✉️  **rosh199** 2 months, 2 weeks ago

hi guys! to those who have complete set of questions, maybe you can find time to share it with me please...ginare.roshni@gmail.com
upvoted 1 times

✉️  **sana_sree** 4 months, 1 week ago

option D is correct

refer:

<https://www.exam-answer.com/google/ace/question95>

upvoted 1 times

✉️  **IraLetiy** 4 months, 2 weeks ago

I am taking the exam soon, could someone please send me entire questions to my email id - ira2901@gmail.com ?

upvoted 1 times

✉️  **3ana** 5 months, 3 weeks ago

hi guys! to those who have complete set of questions, maybe you can find time to share it with me please...anashela.03@gmail.com
upvoted 1 times

✉️  **sonia_mola** 5 months, 3 weeks ago

Selected Answer: D

Answer is D

upvoted 1 times

✉️  **Buruguduystunstugudunstuy** 7 months, 2 weeks ago

Selected Answer: D

Answer A is incorrect because the Cloud Spanner console only shows configurations related to Cloud Spanner instances and databases, but not IAM roles.

Answer B is partially correct in that the IAM & admin console is where IAM policies can be viewed and edited. However, it does not show a history of when users were added to Cloud Spanner IAM roles.

Answer C is incorrect because Stackdriver Monitoring is used to monitor the performance of Google Cloud resources and applications, and does not provide information about IAM role changes.

Overall, the best answer is D, as Stackdriver Logging provides a comprehensive history of all administrative activity logs, including when users were added to Cloud Spanner IAM roles.

upvoted 4 times

✉️  **narenderttn** 9 months, 2 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 96 questions only, if anyone has the entire questions please share them with my narenderpnwr123@gmail.com address. Thanks in advance. I would be forever grateful.

upvoted 1 times

 **u422628** 9 months ago

u must pay hahahah

upvoted 5 times

 **HiddenClouds** 9 months, 3 weeks ago

Selected Answer: D

This was on the exam 12/16/2022

upvoted 5 times

 **leogor** 11 months ago

Selected Answer: D

Stackdriver Logging console > admin activity logs > Cloud Spanner IAM role

upvoted 1 times

 **klyaal** 11 months, 2 weeks ago

I'm taking the exam next week. I can't access all questions. Can someone please send me the full version to klacewicz@gmail.com
Thanks a lot!

upvoted 1 times

 **Rajehs** 11 months, 2 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 96 questions only, if anyone has the entire questions please share them with my rajeshkannafeb6@gmail.com address. Thanks in advance. I would be forever grateful.

upvoted 1 times

 **Charumathi** 1 year ago

D is correct, Activity logs captures the time when the users were given the IAM roles for Cloud Spanner

upvoted 1 times

 **Priyanka109** 1 year ago

You need to see when not what so D.

upvoted 1 times

 **TinkalS** 1 year ago

ANS- D

upvoted 1 times

Your company implemented BigQuery as an enterprise data warehouse. Users from multiple business units run queries on this data warehouse. However, you notice that query costs for BigQuery are very high, and you need to control costs. Which two methods should you use? (Choose two.)

- A. Split the users from business units to multiple projects.
- B. Apply a user- or project-level custom query quota for BigQuery data warehouse.
- C. Create separate copies of your BigQuery data warehouse for each business unit.
- D. Split your BigQuery data warehouse into multiple data warehouses for each business unit.
- E. Change your BigQuery query model from on-demand to flat rate. Apply the appropriate number of slots to each Project.

Correct Answer: BE

Community vote distribution

BE (73%)	D (23%)	5%
----------	---------	----

✉️  **CPBach**  3 years, 3 months ago

I'd say B and E. So either you do B or E to reduce costs.
upvoted 43 times

✉️  **KC_go_reply** 6 months, 2 weeks ago

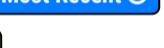
The question says 'which two methods should you use', implying that you want to use *both*. Using quotas together with the flat-rate pricing doesn't make any sense.
Besides that, E is wrong imo. Why? Because flat-rate pricing is very expensive, you pay a fixed high price for something that likely won't be used enough by the average business unit. You need to allocate different 'slots' which is inflexible and complex. It would make more sense to rely on quotas, which means you have an upper limit for costs, but don't necessarily pay the maximum.

upvoted 1 times

✉️  **Jignesh_Gamdha**  2 years, 11 months ago

B & E
Refer below link - first of all you can define quotas on project or user level and 2nd one is you can change from on demand to flat rate model
and define the parameters based on your requirement ---

<https://cloud.google.com/bigquery/docs/custom-quotas>
https://cloud.google.com/bigquery/pricing#flat_rate_pricing
upvoted 23 times

✉️  **Captain1212**  1 month, 1 week ago

Selected Answer: BE
BE seems more legit
upvoted 1 times

✉️  **Neha_Pallavi** 1 month, 2 weeks ago

BE is correct Answer
upvoted 1 times

✉️  **FCADAM** 1 month, 2 weeks ago

B, after reading the specs of BigQuery pricing. C and E doesn't meet the cost requirements. E, Google is no longer offering the flat-rate anymore. E, doesn't meet the criteria and B offers the option to set User-level and project-level custom cost controls.
upvoted 2 times

✉️  **danieerll** 2 months ago

Flat rate is no longer available for Big Query.
upvoted 2 times

✉️  **kasperm** 2 months, 2 weeks ago

BigQuery Flat-Rate Model is no longer sold as of July 5, 2023
"Only available to customers with fixed priced flat-rate prior to end of sale"
<https://cloud.google.com/bigquery/pricing>
upvoted 3 times

✉️  **jayjani66** 2 months, 2 weeks ago

need the pdf as well, jayjani66@gmail.com. Please let me know if you already got it. Thanks.
upvoted 1 times

✉️  **tanha_26** 3 months, 1 week ago

Could someone please send all the questions with answers to tanhapatel2508@gmail.com
upvoted 1 times

✉️  **_F4LLEN_** 4 months, 3 weeks ago

upvoted 2 times

✉ **_F4LLEN_** 5 months, 3 weeks ago

Inorder to get the rest of questions you can type " Google cloud Ace exam question 97" in Google and get it.till 197 is there I guess
upvoted 2 times

✉ **pritampanda1988** 5 months, 2 weeks ago

That's not right. I have got contributor access and I can assure that this is not the way you can get all the questions and answers.
upvoted 1 times

✉ **varun2asses** 5 months, 3 weeks ago

if anyone have got the full access to the PDF please do send it to askvarunn@gmail.com thank you in advance
upvoted 2 times

✉ **_F4LLEN_** 5 months, 3 weeks ago

Message me @dive_into_dev Instagram)
upvoted 1 times

✉ **AkilaSaminathan** 6 months, 1 week ago

Hi, If anyone has complete set of questions, could you please email to akila.saminathan49@gmail.com
upvoted 1 times

✉ **Clam09** 6 months, 2 weeks ago

Hello Everyone, I'm preparing for GCP ACE certification. If anyone have the complete question set, could you please email it to clambor09@gmail.com
Thanks in advance.
upvoted 2 times

✉ **SoumyaRM** 6 months, 2 weeks ago

BE BOTH ARE CORRECT
upvoted 1 times

✉ **RIZI23** 6 months, 2 weeks ago

Hi guys, I am sitting this exam next Monday, Is it possible that anyone can email me the rest of the questions - rizwanfarooq14@hotmail.com I would really appreciate it!!!
upvoted 1 times

✉ **Cecilia_33** 7 months, 2 weeks ago

Dear friends, I am wondering if anyone could kindly send the rest of the questions as pdf to my email:
jing.cecilia.liao@outlook.com
Thank you!
upvoted 1 times

✉ **sk_0007** 7 months, 1 week ago

need the pdf as well, zhaoriches@gmail.com. Please let me know if you already got it. Thanks.
upvoted 1 times

✉ **rizzzzzzz** 6 months, 3 weeks ago

If you guys received it, Please send it to ritwiksingh1605@gmail.com
upvoted 1 times

✉ **artemB** 6 months, 2 weeks ago

Same question here. If anyone can send the rest of questions to the following email I'd be really appreciate it:
chilly.artem@yahoo.com
upvoted 1 times

✉ **iyah** 4 months, 2 weeks ago

May I also have it? ghoziyah2@gmail.com
Thankss!
upvoted 1 times

You are building a product on top of Google Kubernetes Engine (GKE). You have a single GKE cluster. For each of your customers, a Pod is running in that cluster, and your customers can run arbitrary code inside their Pod. You want to maximize the isolation between your customers' Pods. What should you do?

- A. Use Binary Authorization and whitelist only the container images used by your customers' Pods.
- B. Use the Container Analysis API to detect vulnerabilities in the containers used by your customers' Pods.
- C. Create a GKE node pool with a sandbox type configured to gvisor. Add the parameter runtimeClassName: gvisor to the specification of your customers' Pods.
- D. Use the cos_containerd image for your GKE nodes. Add a nodeSelector with the value cloud.google.com/gke-os-distribution: cos_containerd to the specification of your customers' Pods.

Correct Answer: C

Reference:

<https://cloud.google.com/kubernetes-engine/sandbox/>

Community vote distribution

C (100%)

✉  **akshaychavan7**  1 year, 4 months ago

Let me be honest, I did not have any clue to answer this question. However, I spotted the keyword, 'isolation', from the question and a keyword, 'sandbox' from the answers and guessed the answer which turned out to be correct.

So, yes it is C!

upvoted 18 times

✉  **Sac3433**  1 year, 5 months ago

Correct answer is C: You can enable GKE Sandbox on your cluster to isolate untrusted workloads in sandboxes on the node. GKE Sandbox is built using gVisor, an open source project: https://cloud.google.com/kubernetes-engine/docs/concepts/security-overview?hl=en#protecting_nodes_from_untrusted_workloads

upvoted 7 times

✉  **lov75**  9 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

✉  **mattcl** 11 months, 1 week ago

GKE Sandbox <https://cloud.google.com/kubernetes-engine/docs/concepts/sandbox-pods>

upvoted 3 times

✉  **theBestStudent** 1 year, 2 months ago

Selected Answer: C

As it has been mentioned already: <https://cloud.google.com/kubernetes-engine/docs/how-to/sandbox-pods?hl=en>

https://cloud.google.com/kubernetes-engine/docs/how-to/sandbox-pods?hl=en#working_with

upvoted 2 times

✉  **AzureDP900** 1 year, 3 months ago

gVisor is the way to isolate. Those who already preparing for CKS can answer this question without even thinking further. C is right

upvoted 3 times

✉  **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 1 times

✉  **PAUGURU** 1 year, 5 months ago

Selected Answer: C

https://cloud.google.com/kubernetes-engine/docs/concepts/security-overview?hl=en#protecting_nodes_from_untrusted_workloads

upvoted 2 times

Your customer has implemented a solution that uses Cloud Spanner and notices some read latency-related performance issues on one table. This table is accessed only by their users using a primary key. The table schema is shown below.

```
CREATE TABLE Persons (
    person_id INT64 NOT NULL,      // sequential number based on number of registration
    account_creation_date DATE,   // system date
    birthdate DATE,               // customer birthdate
    firstname STRING (255),       // first name
    lastname STRING (255),        // last name
    profile_picture BYTES (255)  // profile picture
) PRIMARY KEY (person_id)
```

You want to resolve the issue. What should you do?

- A. Remove the profile_picture field from the table.
- B. Add a secondary index on the person_id column.
- C. Change the primary key to not have monotonically increasing values.

- D. Create a secondary index using the following Data Definition Language (DDL):

```
CREATE INDEX person_id_ix
ON Persons (
    person_id,
    firstname,
    lastname
) STORING (
    profile_picture
)
```

Correct Answer: D

Community vote distribution

C (63%)	D (28%)	9%
---------	---------	----

✉  **BenKenGo**  1 year, 1 month ago

Selected Answer: D

Create a secondary index using the following Data Definition.
If we watch the next video, he talks about a change to monotonically when we insert rows.
Finally when we talk about read and we have a performance issues, we must create a index.
<https://www.youtube.com/watch?v=r6uj0HMNQNQ>
upvoted 8 times

✉  **SilNilanjan** 8 months, 1 week ago

Adding index for faster retrieval is a basic DBMS concept but why do we need the index on firstname and lastname as per D?
upvoted 2 times

✉  **Captain1212**  1 month, 1 week ago

Selected Answer: C

C seems more correct
upvoted 1 times

✉  **dasgcp** 6 months, 2 weeks ago

How is this a GCP question?
upvoted 3 times

✉  **Mariuselul** 6 months, 1 week ago

Spanner and distribution of primary key
upvoted 1 times

✉  **temple1305** 6 months, 3 weeks ago

Selected Answer: C

PK already has index by default, so not B. D - index by 3 fields. but users use person_id for acces, so D is wrong.
So C - because monotonically increasing fields is not good candidate for PK(because index degradation)
upvoted 2 times

✉  **Spiff** 7 months, 1 week ago

Selected Answer: C

Based on the supplied video by others; <https://www.youtube.com/watch?v=r6uj0HMNQNQ>, we can see at time 1:47 that due to the slitting of the rows, a sequential primary key will create hotspots. Therefor we need a non-sequential key; e.g. hash-based key
upvoted 1 times

✉  **Dohohoh** 7 months, 4 weeks ago

Selected Answer: B

B. is correct...Add a secondary index on the person_id column.

Adding a secondary index on the person_id column would help resolve the read latency-related performance issues on this table. Since the table is accessed using only the primary key, creating a secondary index on the person_id column would allow Cloud Spanner to retrieve the data using the index, rather than scanning the entire table. This can significantly reduce the read latency for queries that access this table.

Removing the profile_picture field or changing the primary key to not have monotonically increasing values may not necessarily resolve the performance issues related to read latency. Creating a secondary index is a more targeted solution to address the specific issue at hand.

Option D is incomplete and does not provide enough information to assess its correctness.

upvoted 3 times

 **Kopy** 10 months, 1 week ago

Selected Answer: C

Create a secondary index using the following Data Definition.

If we watch the next video, he talks about a change to monotonically when we insert rows.

Finally when we talk about read and we have a performance issues, we must create a index.

<https://www.youtube.com/watch?v=r6uj0HMNQNQ>

upvoted 2 times

 **Nazz1977** 10 months, 1 week ago

Selected Answer: C

I think it is C

upvoted 1 times

 **romega2** 10 months, 1 week ago

Selected Answer: C

It's definitely C, D doesn't make sense here

upvoted 1 times

 **babu85** 11 months ago

Choose a primary key to prevent hotspots

As mentioned in Schema and data model, you should be careful when choosing a primary key to not accidentally create hotspots in your database. One cause of hotspots is having a column whose value monotonically increases as the first key part, because this results in all inserts occurring at the end of your key space. This pattern is undesirable because Spanner divides data among servers by key ranges, which means all your inserts will be directed at a single server that will end up doing all the work.

upvoted 2 times

 **mattcl** 11 months, 1 week ago

C <https://cloud.google.com/spanner/docs/schema-design#primary-key-prevent-hotspots>

upvoted 4 times

 **Sozan** 11 months, 1 week ago

Selected Answer: C

C is the right answer. Why? "This table is accessed only by their users using a primary key." So adding additional indexes on firstname and lastname won't help.

upvoted 3 times

 **Aninina** 11 months, 1 week ago

Selected Answer: C

you should be careful when choosing a primary key to not accidentally create hotspots in your database. One cause of hotspots is having a column whose value monotonically increases as the first key part, because this results in all inserts occurring at the end of your key space. This pattern is undesirable because Spanner divides data among servers by key ranges, which means all your inserts will be directed at a single server that will end up doing all the work.

upvoted 3 times

 **Erie** 11 months, 3 weeks ago

Selected Answer: C

C: this is to avoid having hotspots. If the PK is monotonic, then there is a higher chance of requests being routed to the same spanner server and thus overloading it.

D: Primary key is an index already. The question states that the users only accesses this table by PK.

upvoted 2 times

 **leliodesouza** 12 months ago

Selected Answer: D

The answer is D.

upvoted 1 times

 **theBestStudent** 12 months ago

Selected Answer: C

Answer C: Schema design best practice #1: Do not choose a column whose value monotonically increases or decreases as the first key part for a high write rate table.

<https://cloud.google.com/spanner/docs/schema-design#postgresql>

upvoted 2 times

 **alleinallein** 1 year ago

Selected Answer: C

IMO C

<https://cloud.google.com/spanner/docs/schema-design>

upvoted 2 times

Your finance team wants to view the billing report for your projects. You want to make sure that the finance team does not get additional permissions to the project. What should you do?

- A. Add the group for the finance team to roles/billing user role.
- B. Add the group for the finance team to roles/billing admin role.
- C. Add the group for the finance team to roles/billing viewer role.
- D. Add the group for the finance team to roles/billing project/Manager role.

Correct Answer: C*Community vote distribution*

C (100%)

 **miniso8153** Highly Voted 3 years, 4 months ago

C

"Billing Account Viewer access would usually be granted to finance teams, it provides access to spend information, but does not confer the right to link or unlink projects or otherwise manage the properties of the billing account."

<https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 51 times

 **dan80** Highly Voted 3 years, 4 months ago

Answer is C - Billing Account Viewer access would usually be granted to finance teams, it provides access to spend information, but does not confer the right to link or unlink projects or otherwise manage the properties of the billing account.

upvoted 21 times

 **krsourav** 2 years, 8 months ago

Hey, look at this:)

upvoted 4 times

 **Captain1212** Most Recent 1 month, 1 week ago

Selected Answer: C

C is the correct answer

upvoted 1 times

 **anujithn** 2 months, 2 weeks ago

Selected Answer: C

all other will give additional permission

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

It is no brainer question, C is right

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 1 times

 **ElenaL** 1 year, 9 months ago

Selected Answer: C

C - the only role appropriate answer to view and not change anything in the project is the billing viewer role.

upvoted 2 times

 **Vidyaji** 1 year, 10 months ago

C is perfect

upvoted 1 times

 **vishnukumartr** 1 year, 10 months ago

C. Add the group for the finance team to roles/billing viewer role.

upvoted 1 times

 **Chotebhaisahab** 2 years, 1 month ago

agree c is the correct option. needs only billing viewer role.

upvoted 1 times

 **Katheri7689** 2 years, 4 months ago

C is the correct answer

upvoted 2 times

 **mcaromit** 2 years, 4 months ago

C is correct

upvoted 1 times

 **[Removed]** 2 years, 6 months ago

C is correct. Add the group for the finance team to roles/billing viewer role.

upvoted 1 times

 **cloud_guru** 2 years, 7 months ago

Obvious choice here is C-- viewer to make sure they dont have additional access.

Source: this is my job

upvoted 1 times

 **Vic1043** 2 years, 7 months ago

ANS - C

upvoted 1 times

 **GCP_Student1** 2 years, 7 months ago

C. Add the group for the finance team to roles/billing viewer role.

upvoted 2 times

 **swatititame** 2 years, 10 months ago

• C. Add the group for the finance team to roles/billing viewer role.

upvoted 2 times

Your organization has strict requirements to control access to Google Cloud projects. You need to enable your Site Reliability Engineers (SREs) to approve requests from the Google Cloud support team when an SRE opens a support case. You want to follow Google-recommended practices. What should you do?

- A. Add your SREs to roles/iam.roleAdmin role.
- B. Add your SREs to roles/accessapproval.approver role.
- C. Add your SREs to a group and then add this group to roles/iam.roleAdmin.role.
- D. Add your SREs to a group and then add this group to roles/accessapproval.approver role.

Correct Answer: B

Community vote distribution

D (100%)

 **reinocd21** Highly Voted 3 years, 4 months ago

D. Add your SREs to a group and then add this group to roles/accessapproval approver role.
-Google recommendation.
upvoted 55 times

 **Priyankahere** Highly Voted 1 year, 8 months ago

This was there in exam, go with community answers.
upvoted 16 times

 **Stella_1999** 1 year, 8 months ago

how was your exam? is this website qts useful?
upvoted 2 times

 **Captain1212** Most Recent 1 month, 1 week ago

Selected Answer: D
D seems more correct
upvoted 1 times

 **SanjeevKumar1983** 1 month, 1 week ago

D is correct
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

D is right ..
upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for D
upvoted 1 times

 **hiranfilho** 1 year, 5 months ago

Selected Answer: D
Answers C and D are correct, but it doesn't say if the SRE already has a group and as it is Google's recommendation to make a group to add users and privileges to the group, the right one is D
upvoted 3 times

 **WTY** 1 year, 5 months ago

It mentioned more than one SRE, so adding the user to group is most suitable approach, Answer is D.
upvoted 2 times

 **ErenYeager** 1 year, 6 months ago

Passed my exams today. Not because just because of the questions I practiced here, but because of you guys, your knowledge and experience and breakdown of questions. Too bad this site can't go legit. It such an wholesome resource.

Some final words... KEEP MOVING FORWARD UNTIL ALL THE QUESTIONS ARE DESTROYED TATAKAE!!!!!!
upvoted 14 times

 **obeythefist** 1 year, 7 months ago

I've seen about 5 questions which are like this, always asking how to grant access and "follow Google best practice", and every time it's just making sure you know to use a group to control access to resources for users, and not adding users directly to objects.

Remember that keyword, "Google best practice" means "make sure you use a group"
upvoted 10 times

 **Mhkgupta** 1 year, 7 months ago

D. Add your SREs to a group and then add this group to roles/accessapproval.approver role.
upvoted 3 times

 **pnVino27** 1 year, 10 months ago

Selected Answer: D

D is Correct
upvoted 3 times

 **Vidyaji** 1 year, 10 months ago

D is perfect
upvoted 3 times

 **alaahakim** 1 year, 10 months ago

Selected Answer: D

D Correct
upvoted 4 times

 **shawnkkk** 1 year, 10 months ago

D. Add your SREs to a group and then add this group to roles/accessapproval.approver role.
upvoted 2 times

 **TenshiD** 1 year, 10 months ago

Selected Answer: D

D is correct
upvoted 3 times

 **Litan** 2 years ago

D correct
upvoted 3 times

You need to host an application on a Compute Engine instance in a project shared with other teams. You want to prevent the other teams from accidentally causing downtime on that application. Which feature should you use?

- A. Use a Shielded VM.
- B. Use a Preemptible VM.
- C. Use a sole-tenant node.
- D. Enable deletion protection on the instance.

Correct Answer: D

Community vote distribution

C (54%) D (46%)

 **ESP_SAP**  3 years, 1 month ago

Correct Answer is (D):

Preventing Accidental VM Deletion

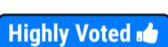
This document describes how to protect specific VM instances from deletion by setting the `deletionProtection` property on an `Instance` resource. To learn more about VM instances, read the Instances documentation.

As part of your workload, there might be certain VM instances that are critical to running your application or services, such as an instance running a SQL server, a server used as a license manager, and so on. These VM instances might need to stay running indefinitely so you need a way to protect these VMs from being deleted.

By setting the `deletionProtection` flag, a VM instance can be protected from accidental deletion. If a user attempts to delete a VM instance for which you have set the `deletionProtection` flag, the request fails. Only a user that has been granted a role with `compute.instances.create` permission can reset the flag to allow the resource to be deleted.

<https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>

upvoted 56 times

 **professor**  3 years, 3 months ago

Agree with D

You can enable Termination protection

upvoted 14 times

 **ekta25**  5 days, 1 hour ago

D. Enable deletion protection on the instance.

upvoted 1 times

 **Blessiq** 2 weeks, 6 days ago

Selected Answer: C

C. "Use a sole-tenant node" allows you to have dedicated hardware for your VM instances, providing isolation from other workloads. This isolation can help prevent other teams' actions from impacting your application's availability.

upvoted 1 times

 **scanner2** 1 month ago

Selected Answer: D

Answer is D

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: C

option c is correct because, it gives you isolation from other projects, where option d is only prevent if from delete only

upvoted 1 times

 **Neha_Pallavi** 1 month, 2 weeks ago

C. Use a sole-tenant node.

Explanation:

By using a sole-tenant node in Compute Engine, you can isolate your application from other teams sharing the same project. A sole-tenant node is a physical Compute Engine server that is dedicated exclusively to your project. This ensures that no other project can run their instances on the same physical server.

upvoted 2 times

 **shtapit** 2 months ago

Option D (Enable deletion protection): Enabling deletion protection helps prevent accidental deletion of the instance, but it does not provide isolation from other teams' activities or ensure application availability.

Option C

upvoted 1 times

 **Nxt_007** 2 months ago

Selected Answer: C

Correct Answer should be (C):

A sole-tenant node is a physical Compute Engine server that is dedicated exclusively to your project. By using a sole-tenant node, you ensure that your application runs in isolation from other workloads belonging to other teams. This isolation provides enhanced security and performance guarantees, as well as protection against resource contention.

-> Option (D) is not suitable because

Enabling deletion protection prevents the accidental deletion of the instance, but it does not prevent other teams from affecting the application's performance or availability.

upvoted 1 times

 **rosh199** 2 months, 2 weeks ago

C. Use a sole-tenant node.

Explanation:

By using a sole-tenant node in Compute Engine, you can isolate your application from other teams sharing the same project. A sole-tenant node is a physical Compute Engine server that is dedicated exclusively to your project. This ensures that no other project can run their instances on the same physical server.

Using a sole-tenant node helps to prevent other teams from accidentally causing downtime on your application. It ensures that your application runs in isolation, and any potential performance issues or disruptions caused by other projects won't affect your application hosted on the dedicated node.

upvoted 2 times

 **pet1er9** 3 months, 3 weeks ago

Selected Answer: D

I chose D

upvoted 1 times

 **Morenothing** 5 months, 2 weeks ago

Selected Answer: D

D is the answer

upvoted 1 times

 **Shweta2jun** 5 months, 2 weeks ago

The answer is clearly C.

Explanation -Through a configurable host maintenance policy, you can control the behavior of sole-tenant VMs while their host is undergoing maintenance. You can specify when maintenance occurs, and whether the VMs maintain affinity with a specific physical server or are moved to other sole-tenant nodes within a node group.

Reference - <https://cloud.google.com/compute/docs/nodes/sole-tenant-nodes>

upvoted 1 times

 **Hdjduidiidkdn** 6 months ago

C - A sole-tenant node is a physical Compute Engine server that is dedicated exclusively to your project. When you use a sole-tenant node, you have complete control over the underlying hardware and can ensure that no other project or user can affect your application's performance.

upvoted 1 times

 **adityanarayan** 6 months, 1 week ago

Sole-tenant does not prevent you from stopping a VM in GCP. Sole-tenant nodes let you have exclusive access to a physical Compute Engine server that is dedicated to hosting only your project's VMs¹. You can use sole-tenant nodes to keep your VMs physically separated from VMs in other projects, or to group your VMs together on the same host hardware¹.

upvoted 1 times

 **Vladimir_Sakhonchik** 7 months, 2 weeks ago

Selected Answer: D

Using a sole-tenant node is not the best solution in this scenario, as sole-tenant nodes are designed to provide dedicated hardware resources to a single tenant or application. While this can help with security and isolation, it does not specifically address the issue of accidental deletion by other teams in a shared project.

upvoted 2 times

 **xaqanik** 7 months, 2 weeks ago

Selected Answer: C

what about if someone stops VM instead of deleting? both can cause downtime.

so that is why deletion protection does not make a sense.

i vote for C

upvoted 5 times

 **[Removed]** 6 months ago

Sole-tenant node are even less relevant to this question. Because the project is shared with other teams. D makes the most sense here

upvoted 2 times

Your organization needs to grant users access to query datasets in BigQuery but prevent them from accidentally deleting the datasets. You want a solution that follows Google-recommended practices. What should you do?

- A. Add users to roles/bigquery user role only, instead of roles/bigquery dataOwner.
- B. Add users to roles/bigquery dataEditor role only, instead of roles/bigquery dataOwner.
- C. Create a custom role by removing delete permissions, and add users to that role only.
- D. Create a custom role by removing delete permissions. Add users to the group, and then add the group to the custom role.

Correct Answer: B

Community vote distribution

D (60%)

A (37%)

✉  **someoneinthecloud**  3 years, 2 months ago

I believe the key part is the "following Google Best Practices" phrase.

- A - Works, but doesn't follow GCP best practices
- B - Doesn't work as the role grants permission to delete datasets
- C - Works, but is more complicated than A and doesn't follow Google best practices
- D - Correct, more complicated than A, but it follows Google Best Practices.

upvoted 66 times

✉  **YuvarajK** 2 years, 3 months ago

I think A is the Answer and it follows GCP best practices.

<https://cloud.google.com/iam/docs/understanding-roles#bigquery-roles>

We do have the role - BigQuery User which does the below permissions

When applied to a project, this role also provides the ability to run jobs, including queries, within the project.

bigquery.datasets.create

bigquery.datasets.get

bigquery.datasets.getIamPolicy

upvoted 9 times

✉  **Abhi00754** 6 months, 1 week ago

bigquery.datasets.create allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets so he can delete these created datasets

upvoted 2 times

✉  **Abhi00754** 6 months, 1 week ago

D seems correct

upvoted 1 times

✉  **kyo** 2 years, 2 months ago

I don't think A works properly.

roles/bigquery.user has bigquery.datasets.create. And the documentation states:

> Additional, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets.

If bigquery.user creates a new dataset, it's likely that bigquery.user will get permission to delete that dataset. This means that bigquery.user may have permission to delete data.

<https://cloud.google.com/bigquery/docs/access-control>

upvoted 12 times

✉  **brvinod** 1 year, 8 months ago

A bigquery.user will get a "data owner" role on the datasets he creates. That means he can delete those data sets he created. In that sense A fails to that extent.

upvoted 5 times

✉  **Bossam** 1 year, 2 months ago

See the question carefully "accidentally deleting the datasets" it is saying not to delete "the" datasets which means original dataset which existed before his creation .So answer is A.

upvoted 1 times

✉  **Zina12** 10 months ago

The only way a user can accidentally delete a dataset is if they have the delete permission anyway. So brvinod and kyo's points still stand

upvoted 2 times

✉  **TAvenger** 2 years, 7 months ago

Read description carefully "prevent from accidentally deleting the datasets". Not tables, datasets! option B does not allow to delete datasets either.

Check dataset permissions in the roles/bigquery.dataEditor:

bigrquery.datasets.get
bigrquery.datasets.getIamPolicy
bigrquery.datasets.updateTag
You CANNOT delete dataset with option "B"
upvoted 11 times

✉ **afooh** 1 year, 2 months ago
But it means you will have to add the users one by one which doesn't follow Google best practices...
upvoted 3 times

✉ **Bableves** 1 year, 6 months ago
Neither with A.
upvoted 1 times

✉ **ESP_SAP** Highly Voted 3 years, 1 month ago
Correct Answer is (D):

The proper answer regarding to bigrquery roles is the listed in the options, the proper rol that resolve this requirement is:
roles/bigrquery.dataViewer
https://cloud.google.com/bigquery/docs/access-control#custom_roles

on the other hand, the question explicitly is asking to use the GCP best practices on IAM :
GCP Best Practices explain clearly these rules:

Policy management

- Set organization-level IAM policies to grant access to all projects in your organization.
- Grant roles to a Google group instead of individual users when possible. It is easier to add members to and remove members from a Google group instead of updating an IAM policy to add or remove users.
- If you need to grant multiple roles to allow a particular task, create a Google group, grant the roles to that group, and then add users to that group.

https://cloud.google.com/iam/docs/using-iam-securely#policy_management

upvoted 53 times

✉ **JackGlemins** 2 years, 7 months ago
Other best practice is use predefine roles over custom roles. Maybe A is correct
upvoted 8 times

✉ **JackGlemins** 2 years, 7 months ago
I correct myself: <https://cloud.google.com/iam/docs/understanding-custom-roles>
Key Point: Custom roles enable you to enforce the principle of least privilege, ensuring that the user and service accounts in your organization have only the permissions essential to performing their intended functions.
upvoted 6 times

✉ **prashuG** 3 years, 1 month ago
Answer is A: roles/bigrquery.user is a BigQuery User role which when applied to a project provides the ability to run jobs, including queries, within the project. A member with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project.

Ref: <https://cloud.google.com/iam/docs/understanding-roles#bigquery-roles>
upvoted 5 times

✉ **j1c4b** 2 years, 6 months ago
you can create data set with bigrquery.user role because it has bigrquery.datasets.create permissions. And if a user has bigrquery.datasets.create permissions, when that user creates a dataset, they are granted bigrquery.dataOwner access to it. So A is NOT a choice
upvoted 5 times

✉ **ekta25** Most Recent 5 days, 1 hour ago
A. Add users to roles/bigrquery user role only, instead of roles/bigrquery dataOwner.
upvoted 1 times

✉ **scanner2** 1 month ago
Selected Answer: A
Answer is A
upvoted 1 times

✉ **NoCrapEva** 1 month ago
Selected Answer: A
GCP Best Practices are to use pre-defined roles and the default permissions for the BigQuery.User role are sufficient...
<https://cloud.google.com/iam/docs/understanding-roles#bigquery.user>
upvoted 1 times

✉ **Captain1212** 1 month, 1 week ago
Selected Answer: D
D is more correct as per the google recommended practices
upvoted 1 times

✉ **SanjeevKumar1983** 1 month, 1 week ago
D is correct
upvoted 1 times

 Neha_Pallavi 1 month, 2 weeks ago

D. Create a custom role by removing delete permissions. Add users to the group, and then add the group to the custom role.
upvoted 1 times

 sthapit 2 months ago

C is the answer
upvoted 1 times

 Praxii 5 months, 1 week ago

Selected Answer: D

Correct answer is D. User Role will grant them permission to CREATE new datasets and the datasets created by the role will make the user OWNER of those datasets thereby giving them permission to delete.

If we want to follow google's best practices, correct answer is D.

upvoted 1 times

 sabrinakloud 5 months, 2 weeks ago

Selected Answer: A

I believe it's A:

"When applied to a project, this role also provides the ability to run jobs, including queries, within the project. A principal with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project."

upvoted 1 times

 sabrinakloud 5 months, 2 weeks ago

bigquery.bireservations.get

bigquery.capacityCommitments.get

bigquery.capacityCommitments.list

bigquery.config.get

bigquery.datasets.create

bigquery.datasets.get

bigquery.datasets.getIamPolicy

bigquery.jobs.create

bigquery.jobs.list

bigquery.models.list

bigquery.readsessions.*

bigquery.reservationAssignments.list

bigquery.reservationAssignments.search

bigquery.reservations.get

bigquery.reservations.list

bigquery.routines.list

bigquery.savedqueries.get

bigquery.savedqueries.list

bigquery.tables.list

bigquery.transfers.get

bigquerymigration.translation.translate

resourcemanager.projects.get

resourcemanager.projects.list

upvoted 1 times

 sabrinakloud 5 months, 2 weeks ago

no delete permissions with this role

upvoted 1 times

 sabrinakloud 5 months, 2 weeks ago

<https://cloud.google.com/bigquery/docs/access-control>

upvoted 1 times

 sabrinakloud 5 months, 2 weeks ago

however "the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets." so they have delete permissions > i think D is correct

upvoted 1 times

- ✉ **Abhi00754** 6 months, 1 week ago
<https://cloud.google.com/bigquery/docs/access-control#bigquery.user>
A is correct , user has bigquery.jobs.create,bigquery.jobs.list. bigquery.jobs.create - Run jobs (including queries) within the project.
upvoted 1 times
- ✉ **Abhi00754** 6 months, 1 week ago
but bigquery.datasets.create allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets so he can delete these created datasets.
D seems correct
upvoted 1 times
- ✉ **dasgcp** 6 months, 2 weeks ago
Selected Answer: A
A. Look up the IAM definition for the role.
bigquery.datasets.create

bigquery.datasets.get

bigquery.datasets.getIamPolicy
upvoted 1 times
- ✉ **xaqanik** 8 months ago
i added this question to CHATGPT,, answer is B.
The recommended practice to prevent accidental deletion of datasets in BigQuery is to add users to the roles/bigquery.dataEditor role only, instead of the roles/bigquery.dataOwner role. The roles/bigquery.dataEditor role provides users with the ability to query datasets, but does not allow them to modify or delete datasets. This helps to ensure that your data remains protected, even if users make accidental changes or deletions.

Alternatively, you could also create a custom role by removing delete permissions and add users to that role, but it is easier and recommended to simply use the roles/bigquery.dataEditor role for this purpose.
upvoted 1 times
- ✉ **NosFerazi** 8 months ago
chatgpt is biased towards using predefined role
upvoted 2 times
- ✉ **Nazz1977** 8 months, 1 week ago
Selected Answer: D
Your organization needs to grant users access to query datasets in BigQuery but prevent them from accidentally deleting the datasets.
How about edit options?
So I think D is the best answer, since the users need read and edit .
upvoted 1 times
- ✉ **jrlsl1991** 8 months, 1 week ago
Selected Answer: D
D seems to me like the best option. Following the least privilege best practice, we can remove anything that's not required to query datasets. They don't need to create new datasets, so option A would give them more access than necessary perform the task.
upvoted 1 times
- ✉ **gishonia** 8 months, 3 weeks ago
Selected Answer: D
Sole-tenancy means that VM just runs on isolated hardware, nothing to do with permissions and access
upvoted 1 times

You have a developer laptop with the Cloud SDK installed on Ubuntu. The Cloud SDK was installed from the Google Cloud Ubuntu package repository. You want to test your application locally on your laptop with Cloud Datastore. What should you do?

- A. Export Cloud Datastore data using gcloud datastore export.
- B. Create a Cloud Datastore index using gcloud datastore indexes create.
- C. Install the google-cloud-sdk-datastore-emulator component using the apt get install command.
- D. Install the cloud-datastore-emulator component using the gcloud components install command.

Correct Answer: D

Community vote distribution

C (51%)

D (49%)

✉  **someoneinthecloud**  3 years, 2 months ago

I believe answer is C

<https://cloud.google.com/sdk/docs/downloads-apt-get>

The question is not about the datastore command itself but from where we should run the update command on the Ubuntu to install the component.

upvoted 42 times

✉  **XRiddlerX** 3 years, 2 months ago

I agree with this comment. The answer is C.

If you installed the SDK from the Ubuntu repo and try to do the following:

\$ gcloud components install cloud-datastore-emulator

You will receive this message:

ERROR: (gcloud.components.install)

You cannot perform this action because the Cloud SDK component manager is disabled for this installation. You can run the following command to achieve the same result for this installation:

sudo apt-get install google-cloud-sdk-datastore-emulator

upvoted 45 times

✉  **myuniqueusername** 2 years ago

absolutely insane if that question comes up during the associate exam, who on earth would know that off the top of their heads?

upvoted 31 times

✉  **Ale1973** 3 years, 1 month ago

WOW!!! Today I have learned a new and interesting thing thanks to you...

upvoted 7 times

✉  **stepkurniawan** 3 years, 1 month ago

it says that in your Ubuntu, you have Cloud SDK installed already. So it should be able to run the command in D

upvoted 4 times

✉  **Ale1973** 3 years, 1 month ago

Yes, but it says that "The Cloud SDK was installed from the Google Cloud Ubuntu package repository", then to install datastore emulator you should use the command in Option C.

upvoted 14 times

✉  **professor**  3 years, 3 months ago

Ans is D

<https://cloud.google.com/datastore/docs/tools/datastore-emulator>

upvoted 23 times

✉  **Eshkrkrkr** 2 years, 11 months ago

Wrong! The answer is C! When you install SDK using apt Cloud SDK Component Manager is disabled and you need to install extra packages again using apt.

https://cloud.google.com/sdk/docs/components#managing_cloud_sdk_components

Note: These instructions will not work if you have installed Cloud SDK using a package manager such as APT or yum because Cloud SDK Component Manager is disabled when using that method of installation.

upvoted 26 times

✉  **SWObaby** 2 years, 10 months ago

I believe the answer is C...

It is a tricky question!! The question states, "The Cloud SDK was installed from the Google Cloud Ubuntu package repository." For those, who aren't that familiar with Debian/Ubuntu, D seems like an attractive answer. It works as a way to install Datastore...but it

I recommend looking back to G Cloud SDK installation (Debian/Ubuntu): <https://cloud.google.com/sdk/docs/install#deb>

Read the "Installation Steps" in the documentation. In Step 3, "sudo apt-get update && sudo apt-get install google-cloud-sdk". Then, Step 4 is additionally adding other components, such as "sudo apt-get install google-cloud-sdk-datastore-emulator".

Proving C the correct answer.

upvoted 23 times

✉ **ShakthiGCP** 2 years, 7 months ago

Go With 'C' ... just tried creating a ubuntu server and verified these.. Dont worry about any other options.
<https://cloud.google.com/sdk/docs/quickstart#deb> check this link .

upvoted 5 times

✉ **ekta25** **Most Recent** 5 days, 1 hour ago

D. Install the cloud-datastore-emulator component using the gcloud components install command.

upvoted 1 times

✉ **the_rizzo** 3 weeks ago

Selected Answer: D

Answer is D https://cloud.google.com/datastore/docs/tools/datastore-emulator#installing_the_emulator

upvoted 1 times

✉ **scanner2** 1 month ago

Selected Answer: D

Answer is D

upvoted 1 times

✉ **scanner2** 1 month ago

Selected Answer: D

Correct answer is D.

Command = gcloud components install cloud-datastore-emulator

https://cloud.google.com/datastore/docs/tools/datastore-emulator#installing_the_emulator

upvoted 1 times

✉ **Khalid209** 1 month ago

"The Cloud SDK was installed from the Google Cloud Ubuntu package repository." For those, who aren't that familiar with Debian/Ubuntu, D seems like an attractive answer. It works as a way to install Datastore...but it does NOT fit the context of the question. I recommend looking back to G Cloud SDK installation (Debian/Ubuntu): <https://cloud.google.com/sdk/docs/install#deb> Read the "Installation Steps" in the documentation. In Step 3, "sudo apt-get update && sudo apt-get install google-cloud-sdk". Then, Step 4 is additionally adding other components, such as "sudo apt-get install google-cloud-sdk-datastore-emulator".

Proving C the correct answer

upvoted 1 times

✉ **Captain1212** 1 month, 1 week ago

Selected Answer: C

the answer is C

upvoted 1 times

✉ **Neha_Pallavi** 1 month, 2 weeks ago

C.Install the google-cloud-sdk-datastore-emulator component using the apt get install command

upvoted 1 times

✉ **someone1224** 2 months, 1 week ago

Selected Answer: D

The Datastore emulator is a component of the Google Cloud CLI's gcloud CLI. Use the gcloud components install command to install the Datastore emulator:

gcloud components install cloud-datastore-emulator

<https://cloud.google.com/datastore/docs/tools/datastore-emulator>

upvoted 1 times

✉ **rosh199** 2 months, 2 weeks ago

The answer is C.

upvoted 1 times

✉ **Shweta2jun** 5 months, 2 weeks ago

Correct answer is C.

Look at the command entered to perform the same.

<https://github.com/googleapis/google-cloud-datastore/issues/223>

upvoted 1 times

✉ **dobberzoon** 6 months, 1 week ago

Selected Answer: C

The question states, "The Cloud SDK was installed from the Google Cloud Ubuntu package repository."

The google-cloud-sok-app-engine-java component can be installed as follows:

`sudo apt-get install google-cloud-sdk app-engine-java`

Installing Cloud SDK

Read the "Installation Steps" in the documentation. In Step 3, "sudo apt-get update && sudo apt-get install google-cloud-sak". Then, Step 4 is additionally adding other components, such as "sudo apt-get install google-cloud-sdk-datastore-emulator"

upvoted 1 times

✉ **Ark88** 6 months, 4 weeks ago

Selected Answer: D

<https://cloud.google.com/datastore/docs/tools/datastore-emulator>

upvoted 1 times

✉ **bbgcp** 7 months ago

Selected Answer: C

Verified on Ubuntu. 100% C

upvoted 2 times

✉ **Bobbybash** 7 months, 4 weeks ago

Selected Answer: C

C. Install the google-cloud-sdk-datastore-emulator component using the apt get install command.

To test your application locally on your laptop with Cloud Datastore, you should install the google-cloud-sdk-datastore-emulator component using the apt get install command. This component provides a local development server that can emulate the Cloud Datastore service. You can then use the Cloud Datastore client libraries to interact with the emulator as if it were the actual Cloud Datastore service.

Exporting Cloud Datastore data using gcloud datastore export or creating a Cloud Datastore index using gcloud datastore indexes create is not relevant for testing your application locally with Cloud Datastore.

Installing the cloud-datastore-emulator component using the gcloud components install command is also not relevant for testing your application locally with Cloud Datastore on Ubuntu. The google-cloud-sdk-datastore-emulator component is the one you need to install for this purpose.

upvoted 1 times

✉ **Nazz1977** 8 months, 1 week ago

Selected Answer: D

It is D

Installing the emulator:

The Datastore emulator is a component of the Google Cloud CLI's gcloud CLI. Use the gcloud components install command to install the Datastore emulator:

gcloud components install cloud-datastore-emulator

https://cloud.google.com/datastore/docs/tools/datastore-emulator#installing_the_emulator

upvoted 1 times

Your company set up a complex organizational structure on Google Cloud. The structure includes hundreds of folders and projects. Only a few team members should be able to view the hierarchical structure. You need to assign minimum permissions to these team members, and you want to follow Google-recommended practices. What should you do?

- A. Add the users to roles/browser role.
- B. Add the users to roles/iam.roleViewer role.
- C. Add the users to a group, and add this group to roles/browser.
- D. Add the users to a group, and add this group to roles/iam.roleViewer role.

Correct Answer: C*Community vote distribution*

C (100%)

  **SIX**  3 years, 4 months ago

C is the better answer.
upvoted 30 times

  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (C):

We need to apply the GCP Best practices.

roles/browser Browser Read access to browse the hierarchy for a project, including the folder, organization, and IAM policy. This role doesn't include permission to view resources in the project.

<https://cloud.google.com/iam/docs/understanding-roles>

upvoted 28 times

  **scanner2**  1 month ago

Selected Answer: C

C is the correct answer.
<https://cloud.google.com/iam/docs/understanding-roles#browser>
upvoted 1 times

  **Captain1212** 1 month, 1 week ago

Selected Answer: C
C is the correct answer
upvoted 1 times

  **Neha_Pallavi** 1 month, 2 weeks ago

C. Add the users to a group, and add this group to roles/browser.
upvoted 1 times

  **Ash_34** 2 months, 3 weeks ago

Selected Answer: C
C is the answer,
Browser - Read access to browse the hierarchy for a project, including the folder, organization, and allow policy
upvoted 1 times

  **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: C
C is correct
upvoted 1 times

  **Charumathi** 1 year ago

C is the correct answer,
ID
roles/browser
Role launch stage
General Availability
Description
Access to browse GCP resources.

6 assigned permissions
resourcemanager.folders.get
resourcemanager.folders.list
resourcemanager.organizations.get
resourcemanager.projects.get
resourcemanager.projects.getIamPolicy
resourcemanager.projects.list

upvoted 2 times

✉ **theBestStudent** 1 year, 2 months ago

Selected Answer: C

<https://cloud.google.com/iam/docs/understanding-roles>

upvoted 1 times

✉ **ikychugtai** 1 year, 2 months ago

C is the better answer

upvoted 1 times

✉ **bobthebuilder55110** 1 year, 2 months ago

Selected Answer: C

C is correct

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

Correct Answer is (C):

upvoted 1 times

✉ **Tirthankar17** 1 year, 4 months ago

Selected Answer: C

C all the way

upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 1 times

✉ **avik001** 1 year, 6 months ago

C is the better ans

upvoted 1 times

✉ **casokan** 1 year, 8 months ago

Correct Answer: C

upvoted 1 times

✉ **pililili_bed** 1 year, 8 months ago

C is correct

upvoted 2 times

Your company has a single sign-on (SSO) identity provider that supports Security Assertion Markup Language (SAML) integration with service providers. Your company has users in Cloud Identity. You would like users to authenticate using your company's SSO provider. What should you do?

- A. In Cloud Identity, set up SSO with Google as an identity provider to access custom SAML apps.
- B. In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.
- C. Obtain OAuth 2.0 credentials, configure the user consent screen, and set up OAuth 2.0 for Mobile & Desktop Apps.
- D. Obtain OAuth 2.0 credentials, configure the user consent screen, and set up OAuth 2.0 for Web Server Applications.

Correct Answer: A*Community vote distribution*

B (100%)

✉  **poogcp**  3 years, 4 months ago

For me its B option
upvoted 29 times

✉  **PhilipAWS**  2 years, 9 months ago

Only option B make sense to me as per - https://support.google.com/cloudidentity/answer/6262987?hl=en&ref_topic=7558767
upvoted 27 times

✉  **nitinz** 2 years, 7 months ago

you nailed it. B is correct.
upvoted 3 times

✉  **ekta25**  5 days ago

B. In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.
upvoted 1 times

✉  **scanner2** 1 month ago

Selected Answer: B

When you use SSO for Cloud Identity or Google Workspace, your external IdP is the SAML IdP and Google is the SAML service provider.
https://cloud.google.com/architecture/identity/single-sign-on#single_sign-on_process
upvoted 1 times

✉  **Captain1212** 1 month, 1 week ago

Selected Answer: B
B is the correct answer as c and d are not SAML
upvoted 1 times

✉  **goshubh** 1 month, 1 week ago

Chat GPT has suggested option A, here's the reason it gave for why B is not the right option-

setting up SSO with a third-party identity provider with Google as a service provider, is typically used when your organization wants to use an external SSO IdP, not Google Cloud Identity, to authenticate users.
upvoted 1 times

✉  **Neha_Pallavi** 1 month, 2 weeks ago

B. In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.
upvoted 1 times

✉  **Haigk** 3 months, 3 weeks ago

Selected Answer: B
B is correct, only one that makes sense.
C & D are OAuth not SAML, and A says Google as IDP...
upvoted 1 times

✉  **hiromi** 11 months, 1 week ago

Selected Answer: B
For me b is the ans
upvoted 1 times

✉  **Omkarshingare** 7 months, 2 weeks ago

Hi my GCP ACE exam is next week, if anyone has contributor access, please send questions pdf to this mail ID-
warcrycreature@gmail.com, Thanks in advance!
upvoted 1 times

 **leliodesouza** 12 months ago

Selected Answer: B

B is the correct answer

upvoted 2 times

 **AwesomeGCP** 1 year ago

Selected Answer: B

B. In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.

upvoted 1 times

 **Angel_99** 1 year, 1 month ago

Selected Answer: B

Only option B make sense to me

upvoted 1 times

 **habros** 1 year, 2 months ago

Same and OAuth were two different standards! Google Account is an IdP on its own, whereas question calls for external IdP. Definitely B.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

Option B

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for B

upvoted 1 times

 **Terzlightyear** 1 year, 5 months ago

Selected Answer: B

B is the right one

upvoted 2 times

 **luciorifa** 1 year, 7 months ago

Selected Answer: B

B is the correct answer

upvoted 2 times

Your organization has a dedicated person who creates and manages all service accounts for Google Cloud projects. You need to assign this person the minimum role for projects. What should you do?

- A. Add the user to roles/iam.roleAdmin role.
- B. Add the user to roles/iam.securityAdmin role.
- C. Add the user to roles/iam.serviceAccountUser role.
- D. Add the user to roles/iam.serviceAccountAdmin role.

Correct Answer: D

Community vote distribution

D (100%)

 **PhilipAWS** Highly Voted  2 years, 9 months ago

Whoever say C is right answer, please read the question 1000000000 times if not understand - "Your organization has a dedicated person who creates and manages all service accounts for Google Cloud projects. " Dedicated person who creates and manages all service... Now read below;

To allow a user to manage service accounts, grant one of the following roles:

Service Account User (roles/iam.serviceAccountUser): Includes permissions to list service accounts, get details about a service account, and impersonate a service account.

Service Account Admin (roles/iam.serviceAccountAdmin): Includes permissions to list service accounts and get details about a service account. Also includes permissions to create, update, and delete service accounts, and to view or change the IAM policy on a service account.

Now look in which role mentioned "CREATE"?

Obviously - roles/iam.serviceAccountAdmin..... So Answer is????

1M% - D only

upvoted 59 times

 **Jhelum** 1 year, 8 months ago

Calm down Jamal, don't pull out the knife...

upvoted 23 times

 **creativenets** 3 months, 3 weeks ago

He's right. Calm down. Save that bomb for later.

upvoted 2 times

 **SIX** Highly Voted  3 years, 4 months ago

The right answer is D.

upvoted 40 times

 **scanner2** Most Recent  1 month ago

Selected Answer: D

roles/iam.serviceAccountAdmin --> Create and manage service accounts.

<https://cloud.google.com/iam/docs/understanding-roles#iam.serviceAccountAdmin>

roles/iam.serviceAccountUser --> Run operations as the service account.

<https://cloud.google.com/iam/docs/understanding-roles#iam.serviceAccountUser>

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: D

D is the right answer, as the person needs to create also

upvoted 1 times

 **Neha_Pallavi** 1 month, 2 weeks ago

D. Service Account Admin (roles/iam.serviceAccountAdmin): Includes permissions to list service accounts and get details about a service account. Also includes permissions to create, update, and delete service accounts, and to view or change the IAM policy on a service account.

upvoted 2 times

 **AzureDP900** 1 year, 3 months ago

dedicated person who creates and manages all service accounts is key word makes me select D as right answer.

upvoted 1 times

 **haroldbenites** 1 year, 4 months ago

Go for D

upvoted 1 times

 **luciorifa** 1 year, 7 months ago

D is the correct answer
upvoted 1 times

 **Priyankahere** 1 year, 8 months ago

This was there in exam, go with community answers.
upvoted 2 times

 **kped21** 1 year, 9 months ago

D - ServiceAccountUser has no privs on ServiceAccounts so C is wrong.
upvoted 1 times

 **jaffarali** 1 year, 10 months ago

D is the right option
upvoted 1 times

 **sharmamohitkr700** 1 year, 10 months ago

Selected Answer: D

To allow a user to manage service accounts, grant one of the following roles:

Service Account User (roles/iam.serviceAccountUser): Includes permissions to list service accounts, get details about a service account, and impersonate a service account.

Service Account Admin (roles/iam.serviceAccountAdmin): Includes permissions to list service accounts and get details about a service account. Also includes permissions to create, update, and delete service accounts, and to view or change the IAM policy on a service account.

upvoted 3 times

 **alaahakim** 1 year, 10 months ago

Selected Answer: D

D is Correct

upvoted 2 times

 **TenshiD** 1 year, 10 months ago

Selected Answer: D

D is correct.

upvoted 1 times

 **maggieli** 1 year, 11 months ago

D is correct.

Service Account Admin (roles/iam.serviceAccountAdmin): Includes permissions to list service accounts and get details about a service account. Also includes permissions to create, update, and delete service accounts, and to view or change the IAM policy on a service account.

upvoted 1 times

 **mcaromit** 2 years, 4 months ago

D is correct

upvoted 2 times

 **[Removed]** 2 years, 6 months ago

D is correct. Add the user to roles/iam.serviceAccountAdmin role.

upvoted 3 times

You are building an archival solution for your data warehouse and have selected Cloud Storage to archive your data. Your users need to be able to access this archived data once a quarter for some regulatory requirements. You want to select a cost-efficient option. Which storage option should you use?

- A. Cold Storage
- B. Nearline Storage
- C. Regional Storage
- D. Multi-Regional Storage

Correct Answer: A

Nearline, Coldline, and Archive offer ultra low-cost, highly-durable, highly available archival storage. For data accessed less than once a year, Archive is a cost-effective storage option for long-term preservation of data. Coldline is also ideal for cold storage—data your business expects to touch less than once a quarter.

For warmer storage, choose Nearline: data you expect to access less than once a month, but possibly multiple times throughout the year. All storage classes are available across all GCP regions and provide unparalleled sub-second access speeds with a consistent API.

Reference:

<https://cloud.google.com/storage/archival>

Community vote distribution

B (56%)

A (44%)

✉ **Teegongkia** Highly Voted 3 years, 2 months ago

Took ACE last week and the exact question came out. I go with B as i felt A is a trick answer. There is no Cold Storage in GCP.
upvoted 32 times

✉ **ssankar** 3 years, 1 month ago

Hello Teegongkia ,
is the questions are still valid ??
Thanks
upvoted 1 times

✉ **BenAji** 2 years, 5 months ago

Cold data tiering refers to the storage of less frequently, or sporadically accessed data in low cost media such as HDFS (Hadoop Distributed File System) and cloud storage options including Amazon Web Services (AWS), Google Cloud Platform (GCP), and Azure Data Lake Storage (ADLS) that are managed separately from the SAP HANA database, but still accessible at any time.
blogs.sap.com/2018/12/03/what-is-sap-hana-cold-data-tiering/

upvoted 1 times

✉ **droogie** Highly Voted 3 years, 3 months ago

This one is confusing. First, there's no 'Cold' storage. It's Coldline.
Nearline Storage is ideal for data you plan to read or modify on average once per month or less. Coldline Storage is ideal for data you plan to read or modify at most once a quarter.
<https://cloud.google.com/storage/docs/storage-classes>

So with the misspelling of 'Cold' and these guys accessing it every 90 days, I'm leaning towards Nearline

upvoted 16 times

✉ **sarahf** 2 years, 9 months ago

At the page for data archiving (<https://cloud.google.com/storage/archival>) the first paragraph says: "Coldline is also ideal for cold storage—data your business expects to touch less than once a quarter."
So there is such thing as Cold storage according to Google.

Also at (<https://cloud.google.com/storage/docs/storage-classes#archive>) they talk about Cold storage: "Cold data storage - Archived data, such as data stored for legal or regulatory reasons, can be stored at low cost as Archive Storage, yet still be available if you need it."

upvoted 16 times

✉ **sanhoo** 2 years, 4 months ago

Thanks for bringing this up. Really helpfull.

upvoted 1 times

✉ **sanhoo** 2 years, 4 months ago

This line indicates that cold storage term is used for - archival / coldline
"With low latency and a consistent API across Cloud Storage, Archive and Coldline introduce cold storage you can actually use"
<https://cloud.google.com/storage/archival>

upvoted 2 times

✉ **DickDastardly** 2 years, 7 months ago

upvoted 12 times

✉  **Eshkrkrkr** 2 years, 11 months ago

It's a typo. Google wouldn't force to consume knowledge that is a non-best practice from Google. Answer is A.

upvoted 4 times

✉  **Ixgywil** 2 years, 5 months ago

For Google, these exams are just another business.

upvoted 8 times

✉  **TAvenger** 2 years, 7 months ago

I believe the question is old, when Regional and Multi-Regional were also storage classes of the GCS.

Before changes: (Multi-Region, Regional, Nearline, Coldline)

After recent changes we have

- Storage Classes (Standard, Nearline, Coldline, Archive)

- Storage Locations (Regional, Dual-region, Multi-Region)

It's tricky for exam because we don't know to answer according to old version or new version.

For the latest version, costs for 1Gb for storing (3 month) + retrieval

Nearline: $0.01 * 3 + 0.01 = 0.04$

Coldline: $0.004 * 3 + 0.02 = 0.032$

Coldline is more cost effective.

If "Cold" means Coldline (not Archive) the answer is A

If "Cold" means Archive the answer is B

I hope that "Cold" means Coldline. I would try with A

upvoted 2 times

✉  **ri_unhou119** 2 years, 4 months ago

A:

Google Cloud doc:

<https://cloud.google.com/storage/docs/storage-classes#coldline>

upvoted 1 times

✉  **obeythefist** 1 year, 7 months ago

Yes, but the question says the data will be accessed once per quarter, Google's documentation tells us that Coldline is most suitable for data accessed less than once per quarter. This direct part of the question tells us how we must answer.

upvoted 4 times

✉  **ekta25** Most Recent 5 days ago

A. Cold Storage

Coldline storage is designed for data that is accessed roughly less than once a quarter. It offers 99% availability and is cost-efficient compared to other storage classes. This makes it a suitable choice for archival solutions where data needs to be accessed infrequently, such as once a quarter for regulatory requirements

upvoted 1 times

✉  **scanner2** 1 month ago

Selected Answer: B

Most appropriate answer here would be Coldline storage, but since this is not present in the given options, so next best answer is Nearline storage.

Cold storage is not a valid cloud storage class.

upvoted 1 times

✉  **cheny** 1 month, 1 week ago

Standard storage -- Hot storage

Nearline storage -- Warm storage

Coldline & Archive storage --- Cold storage

upvoted 1 times

✉  **Neha_Pallavi** 1 month, 2 weeks ago

Nearline, Coldline, and Archive offer ultra low-cost, highly-durable, highly available archival storage. For data accessed less than once a year, Archive is a cost-effective storage option for long-term preservation of data. Coldline is also ideal for cold storage data your business expects to touch less than once a quarter.

For warmer storage, choose Nearline: data you expect to access less than once a month, but possibly multiple times throughout the year. All storage classes are available across all GCP regions and provide unparalleled sub-second access speeds with a consistent API. A is correct answer

upvoted 1 times

✉  **MilanRajGupta** 2 months, 1 week ago

If "Cold" means Coldline (not Archive) the answer is "A"

If "Cold" means Archive the answer is "B"

upvoted 1 times

✉  **kumarts** 4 months, 1 week ago

Google Cloud Coldline Storage is a public cloud cold storage service for infrequently accessed data. It is specifically designed to store data that organizations access once a quarter.

upvoted 2 times

✉  **Astro_123** 5 months, 3 weeks ago

A is correct answer

upvoted 1 times

 **temple1305** 6 months, 1 week ago

Selected Answer: B

There is not COLD storage inside GCP. Coldline only....

upvoted 1 times

 **dnur** 6 months, 2 weeks ago

The answer should be B.

There is no typo in Cold Storage. Check this link: <https://cloud.google.com/storage/docs/storage-classes#coldline>

Cold data storage - Archived data, such as data stored for legal or regulatory reasons, can be stored at low cost as Archive storage, yet still be available if you need it. Archive storage also has higher costs for data access and operations, as well as a 365-day minimum storage duration. Archive storage is the best choice for data that you plan to access less than once a year.

upvoted 1 times

 **Rog_4444** 7 months, 1 week ago

It has to be option B

Option A is not coldline, it says cold storage which is not a real option

upvoted 1 times

 **JC0926** 7 months, 3 weeks ago

Selected Answer: A

once a quarter for some regulatory requirements

upvoted 1 times

 **Vincenzo_Cassano** 8 months, 1 week ago

Selected Answer: A

Nearline Storage - best for storing infrequently accessed data, like reading or modifying data on average once a month or less

Coldline Storage - meant for reading or modifying data, at most, once every 90 days

upvoted 1 times

 **kumarDeep** 9 months ago

B is the correct answer

upvoted 1 times

 **u422628** 9 months ago

Selected Answer: B

correct answer is B, cold storage doesn't exist, and the nearline is used to access file once per month, instead of COLDLINE storage that is once per year

upvoted 1 times

 **Blitzer** 8 months, 3 weeks ago

minimum storage duration for coldline is 90 days but not year

upvoted 1 times

 **siddy_888** 9 months ago

The correct ans is B -> Nearline storage because as per google.

<https://cloud.google.com/storage/docs/storage-classes#coldline>

"Cold data storage - Archived data, such as data stored for legal or regulatory reasons, can be stored at low cost as Archive storage, yet still be available if you need it."

upvoted 1 times

A team of data scientists infrequently needs to use a Google Kubernetes Engine (GKE) cluster that you manage. They require GPUs for some long-running, non-restartable jobs. You want to minimize cost. What should you do?

- A. Enable node auto-provisioning on the GKE cluster.
- B. Create a VerticalPodAutcaler for those workloads.
- C. Create a node pool with preemptible VMs and GPUs attached to those VMs.
- D. Create a node pool of instances with GPUs, and enable autoscaling on this node pool with a minimum size of 1.

Correct Answer: C

Reference:

<https://cloud.google.com/kubernetes-engine/docs/how-to/gpus>

Community vote distribution

A (61%) D (34%) 5%

✉  **Polok**  3 years, 4 months ago

If you need something for long-running, non-restartable jobs you don't use preemptible VMs

Think answer is D.

upvoted 64 times

✉  **[Removed]**  2 years, 6 months ago

Incorrect options are

B. VerticalPodAutcaler scales PODS based on the app you deploy.

For handle infrequently GPU access, you need infrequently GPU nodes

VerticalAutcaler Pod deployed on a non GPU node it useless,

[We can't have the node always have GPU for infrequent requests]

C. Preemptible VMs can't last long

D. For infrequent access, you don't want to have a permanent homogenous cluster.

The correct option is "A"

auto-provisioning = Attaches and deletes node pools to cluster based on the requirements.

Hence creating a GPU node pool, and auto-scaling would be better

<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>

upvoted 22 times

✉  **JCH760310** 1 year, 9 months ago

they "require GPUs" - so after checking in Udemy practice tests there is similar question there. And the D answer seems to be the best fit for our scenario here.

"This option is the most optimal solution for the requirement. Rather than recreating all nodes, you create a new node pool with GPU enabled. You then modify the pod specification to target particular GPU types by adding node selector to your workload's Pod specification. You still have a single cluster, so you pay Kubernetes cluster management fee for just one cluster, thus minimizing the cost." Still better option than creating new GKE cluster with GPUs.

Ref: <https://cloud.google.com/kubernetes-engine/docs/how-to/gpus>

Ref: <https://cloud.google.com/kubernetes-engine/pricing>

upvoted 7 times

✉  **kimharsh** 1 year, 10 months ago

A is not correct because you can't add a GPU node to an existing GKE cluster

Limitations

Before using GPUs on GKE, keep in mind the following limitations:

You cannot add GPUs to existing node pools.

GPU nodes cannot be live migrated during maintenance events.

GPUs are only supported with general-purpose N1 machine types.

GPUs are not supported in Windows Server node pools

REF: <https://cloud.google.com/kubernetes-engine/docs/how-to/gpus#limitations>

So the answer should be D

upvoted 10 times

✉  **rachee** 1 year, 10 months ago

Your reference says existing "node pools" not GKE cluster. Auto-provisioning creates new "node pools":

<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>

upvoted 5 times

✉  **Ridhanya** 1 year, 10 months ago

but node pools are homogenous, so how can we be sure that option A will create a GPU node pool
upvoted 2 times

✉ **wjtib** 1 year, 4 months ago

<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>
Node auto-provisioning creates node pools based on the following information:

CPU, memory and ephemeral storage resource requests.
GPU requests
Pending Pods' node affinities and label selectors.
Pending Pods' node taints and tolerations.
upvoted 9 times

✉ **kyo** 2 years, 2 months ago

I think using NAP is the correct answer.
→Node Auto Provisioning (NAP a.k.a., Nodepool Auto Provisioning)
There is an introduction of NAP described below on the blog.

>The above recommendations optimize for cost. NAP, for instance, reduces costs by taking down nodes during underutilized periods.

<https://cloud.google.com/blog/products/containers-kubernetes/best-practices-for-creating-a-highly-available-gke-cluster>
upvoted 3 times

✉ **dttnc1** 1 year, 12 months ago

I do agree A is the answer. Since this is for infrequent needs, autoscaling in letter D is not cost effective as it will always run min. of 1 instance. If we need to infrequently use a cluster, the nodes should be able to adjust based on the current need.

"With node auto-provisioning, new node pools are created and deleted automatically." <https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>

upvoted 7 times

✉ **ekta25** Most Recent 5 days ago

D. Create a node pool of instances with GPUs, and enable autoscaling on this node pool with a minimum size of 1.
upvoted 1 times

✉ **raxt** 1 week ago

Selected Answer: D
Option A, which seems to be the most voted so far, is incorrect in my opinion.
From Google's docs: "You cannot add GPUs to existing node pools."
<https://cloud.google.com/kubernetes-engine/docs/how-to/gpus#limitations>

So, answer A implies that GPUs are already used in the current cluster that needs to be autoscaled.
upvoted 1 times

✉ **gloria216** 1 week, 2 days ago

The correct answer should be A
upvoted 1 times

✉ **smilyluv** 4 weeks, 1 day ago

Selected Answer: C
why cant be C?
upvoted 2 times

✉ **scanner2** 1 month ago

Selected Answer: A
node auto-provisioning works in Standard Google Kubernetes Engine (GKE) clusters. Node auto-provisioning automatically manages and scales a set of node pools on the user's behalf. With node auto-provisioning, GKE automatically creates and deletes node pools.
Node auto-provisioning can be configured by using a YAML configuration file. Single or Multiple settings can be specified in a single config file. Some advanced configurations can only be specified by using a configuration file.
Sets resource limits for CPU, memory and GPU. Node auto-provisioning will not create a node if the total size of the cluster exceeds the specified resource limits.
<https://cloud.google.com/kubernetes-engine/docs/concepts/node-auto-provisioning>
<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>
upvoted 1 times

✉ **ExamsFR** 1 month ago

Selected Answer: D
The correct option is "D"
upvoted 1 times

✉ **Captain1212** 1 month, 1 week ago

Selected Answer: A
A seems more correct
upvoted 1 times

✉ **Ahmed_Y** 1 month, 3 weeks ago

It is A
"How node auto-provisioning works
Node auto-provisioning is a mechanism of the cluster autoscaler, which only scales existing node pools. With node auto-provisioning

enabled, the cluster autoscaler can create node pools automatically based on the specifications of unschedulable Pods." reference: <https://cloud.google.com/kubernetes-engine/docs/concepts/node-auto-provisioning>
upvoted 1 times

✉ **sthapit** 2 months ago

C is the answer.
upvoted 1 times

✉ **chiNelo** 2 months, 3 weeks ago

I believe the correct option is A. Note auto-provisioning creates node pools based on the following; CPU, memory and ephemeral storage, GPU requests, pending pods (node affinities and label selectors, node taints and toleration). - <https://cloud.google.com/kubernetes-engine/docs/concepts/node-auto-provisioning>
upvoted 1 times

✉ **Jelly_Wang** 5 months ago

Selected Answer: A
D works for most part except for "infrequently". With option D you keep node running all the time which is no necessary. with A, node pool is created and deleted automatically based on workload. Auto-provisioning enable auto scaling as well, which you can specify min number of instance. so A include D, but more than D.

<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>
upvoted 2 times

✉ **vivekvi** 5 months, 2 weeks ago

Selected Answer: A
I thought D but read about Node auto provisioning feature of GKE and got convinced A is correct.
upvoted 1 times

✉ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: A
to minimize the cost
upvoted 1 times

✉ **Nazz1977** 8 months, 1 week ago

Selected Answer: D
I think it is D

Creating an autoscaling GPU node pool:
To take the best, most cost-effective advantage of GPUs on GKE, and to take advantage of cluster autoscaling, we recommend creating separate GPU node pools in your clusters.

https://cloud.google.com/kubernetes-engine/docs/how-to/gpus#gpu_pool
upvoted 1 times

✉ **jrisl1991** 8 months, 1 week ago

Selected Answer: A
I'll go with A - https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning?hl=es_419#operation.

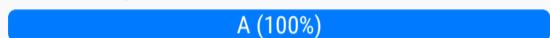
You can request GPUs in autoprovioning. Besides, creating a node pool of instances with GPUs would mean that the GPUs will be there after the data scientists jobs are done, not very costly efficient to me.
upvoted 4 times

Your organization has user identities in Active Directory. Your organization wants to use Active Directory as their source of truth for identities. Your organization wants to have full control over the Google accounts used by employees for all Google services, including your Google Cloud Platform (GCP) organization. What should you do?

- A. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.
- B. Use the cloud Identity APIs and write a script to synchronize users to Cloud Identity.
- C. Export users from Active Directory as a CSV and import them to Cloud Identity via the Admin Console.
- D. Ask each employee to create a Google account using self signup. Require that each employee use their company email address and password.

Correct Answer: A

Reference:

<https://cloud.google.com/solutions/federating-gcp-with-active-directory-introduction>*Community vote distribution* A (100%) **professor**  3 years, 3 months ago

Ans is A

<https://tools.google.com/dlpage/dirsync/>

upvoted 27 times

 **ESP_SAP**  3 years, 1 month ago

Correct Answer (A):

Directory Sync

Google Cloud Directory Sync enables administrators to synchronize users, groups and other data from an Active Directory/LDAP service to their Google Cloud domain directory

<https://tools.google.com/dlpage/dirsync/>

upvoted 19 times

 **scanner2**  1 month ago**Selected Answer: A**<https://support.google.com/a/answer/106368?hl=en><https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-synchronizing-user-accounts><https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-introduction>

upvoted 1 times

 **Captain1212** 1 month, 1 week ago**Selected Answer: A**

A is the correct answer as it help you to synchronize users

upvoted 1 times

 **Neha_Pallavi** 1 month, 2 weeks ago

A. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.

upvoted 1 times

 **Nazz1977** 10 months, 1 week ago**Selected Answer: A**

A....is right

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right, this is part of Tutorials Dojo practice test

upvoted 2 times

 **haroldbenites** 1 year, 4 months ago

Go for A

upvoted 1 times

 **crisyeb** 1 year, 6 months ago**Selected Answer: A**

A is correct

upvoted 1 times

 **LeonNin** 1 year, 9 months ago

Selected Answer: A

A is right

upvoted 3 times

 **alaahakim** 1 year, 10 months ago

Selected Answer: A

A is Correct

upvoted 4 times

 **mcaromit** 2 years, 4 months ago

A is correct

upvoted 3 times

 **[Removed]** 2 years, 6 months ago

A is correct. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.

upvoted 1 times

 **GCP_Student1** 2 years, 7 months ago

- A. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.

upvoted 3 times

 **devscorpio2001** 2 years, 10 months ago

This is A , you can use Google Cloud Sync

upvoted 1 times

 **swatititame** 2 years, 10 months ago

- A. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.

upvoted 1 times

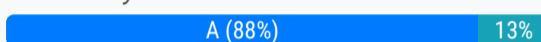
 **RockAJ** 3 years ago

A also for me!

upvoted 1 times

You have successfully created a development environment in a project for an application. This application uses Compute Engine and Cloud SQL. Now you need to create a production environment for this application. The security team has forbidden the existence of network routes between these 2 environments and has asked you to follow Google-recommended practices. What should you do?

- A. Create a new project, enable the Compute Engine and Cloud SQL APIs in that project, and replicate the setup you have created in the development environment.
- B. Create a new production subnet in the existing VPC and a new production Cloud SQL instance in your existing project, and deploy your application using those resources.
- C. Create a new project, modify your existing VPC to be a Shared VPC, share that VPC with your new project, and replicate the setup you have in the development environment in that new project in the Shared VPC.
- D. Ask the security team to grant you the Project Editor role in an existing production project used by another division of your company. Once they grant you that role, replicate the setup you have in the development environment in that project.

Correct Answer: A*Community vote distribution* A (88%)

13%

 **[Removed]**  3 years, 1 month ago

A - correct. Best practice is to create a new project for each environment, such as production and testing. There are no routes between VPCs in these projects by default, so that satisfies the requirement by the security team.
B. Nope. not best practice and allows communication.
C. While this is best practice to create a new project for a different environment, it explicitly breaks the security team's rule of having no path between environments by nature of the shared VPC. The shared VPC allows entities in both VPCs to communicate as if they were in the same VPC. That's definitely wrong.
D. One - not best practice to replicate in the setup in that project. Two - why do they suddenly need the project editor rule? Just a bad answer. Wrong.

upvoted 35 times

 **poogcp**  3 years, 4 months ago

Correct answer is A.

upvoted 35 times

 **pYWORLD** 2 years, 2 months ago

Correct answer!

upvoted 4 times

 **scanner2**  1 month ago**Selected Answer: A**

According to Google recommended practices, you should create a separate project for different environments (dev, test, and prod). Also, the question has forbidden the existence of these environments so shared VPC cannot be used.

upvoted 1 times

 **Captain1212** 1 month, 1 week ago**Selected Answer: A**

A is the correct answer, as it satisfy tyhe requirement of security team , no commiunication , as option c allows coummnication

upvoted 1 times

 **Neha_Pallavi** 1 month, 2 weeks ago

A. Create a new project, enable the Compute Engine and Cloud SQL APIs in that project, and replicate the setup you have created in the development environment.

upvoted 1 times

 **diasporabro** 11 months, 3 weeks ago**Selected Answer: A**

Satisfies requirements by the security team

upvoted 1 times

 **anolive** 12 months ago**Selected Answer: A**

make sense

upvoted 1 times

 **alexandercamachop** 1 year, 2 months ago**Selected Answer: A**

A is definitely the answer.

upvoted 1 times

 **theBestStudent** 1 year, 2 months ago

<https://cloud.google.com/architecture/framework/system-design/resource-management#decouple>
upvoted 1 times

 **andreherwanto** 1 year, 2 months ago

Selected Answer: A

Correct answer is A.
upvoted 1 times

 **csrazdan** 1 year, 3 months ago

Selected Answer: C

Technically we should create a new VPC if the network is not shared. Creating resources in a new project even within a new subnet will not separate unless firewall rules are not explicitly denying the traffic. The best answer is to create a shared VPC where DEV and PROD are service projects.

My Answer is: C
upvoted 1 times

 **theBestStudent** 1 year, 2 months ago

Why you want to share environments? they should be isolated. Therefore Answer should be A.
upvoted 3 times

 **haroldbenites** 1 year, 4 months ago

Go for A
I thought that the correct answer was the C , but the question did not say to communicate both environments.
upvoted 1 times

 **Rukman** 1 year, 6 months ago

Selected Answer: A
Ans: A
Agreed!
upvoted 1 times

 **JieHeng** 2 years, 3 months ago

Should be A
it's a best practice "to have one project per application per environment." - <https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations#project-structure>
upvoted 6 times

 **arsh1916** 2 years, 4 months ago

A answer
upvoted 1 times

 **mcaromit** 2 years, 4 months ago

A is correct
upvoted 1 times

 **EJJ** 2 years, 6 months ago

i will go with A
upvoted 1 times

Your management has asked an external auditor to review all the resources in a specific project. The security team has enabled the Organization Policy called

Domain Restricted Sharing on the organization node by specifying only your Cloud Identity domain. You want the auditor to only be able to view, but not modify, the resources in that project. What should you do?

- A. Ask the auditor for their Google account, and give them the Viewer role on the project.
- B. Ask the auditor for their Google account, and give them the Security Reviewer role on the project.
- C. Create a temporary account for the auditor in Cloud Identity, and give that account the Viewer role on the project.
- D. Create a temporary account for the auditor in Cloud Identity, and give that account the Security Reviewer role on the project.

Correct Answer: C

Community vote distribution

C (80%) A (20%)

✉  **dan80**  3 years, 4 months ago

C - https://cloud.google.com/iam/docs/roles-audit-logging#scenario_external_auditors
upvoted 48 times

✉  **spudleymcdudley** 3 years, 3 months ago

This guy is right!
upvoted 7 times

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (C):

roles/viewer Read access to all resources. Get and list access for all resources.

Using primitive roles

The following table lists the primitive roles that you can grant to access a project, the description of what the role does, and the permissions bundled within that role. Avoid using primitive roles except when absolutely necessary. These roles are very powerful, and include a large number of permissions across all Google Cloud services. For more details on when you should use primitive roles, see the Identity and Access Management FAQ.

IAM predefined roles are much more granular, and allow you to carefully manage the set of permissions that your users have access to. See Understanding Roles for a list of roles that can be granted at the project level. Creating custom roles can further increase the control you have over user permissions.

https://cloud.google.com/resource-manager/docs/access-control-proj#using_primitive_roles

upvoted 20 times

✉  **scanner2**  1 month ago

Selected Answer: C

The Resource Manager provides a domain restriction constraint that can be used in organization policies to limit resource sharing based on domain or organization resource. This constraint allows you to restrict the set of identities that are allowed to be used in Identity and Access Management policies.

Organization policies can use this constraint to limit resource sharing to identities that belong to a particular organization resource.
<https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains>

upvoted 1 times

✉  **Captain1212** 1 month, 1 week ago

Selected Answer: C

C is more correct
upvoted 1 times

✉  **Neha_Pallavi** 1 month, 2 weeks ago

Correct Answer is (C):
upvoted 1 times

✉  **WendyLC** 4 months ago

Selected Answer: C

Correct Answer is (C):

Answer A is wrong because we can't use the the auditor Google account, security team has enabled the Organization Policy specifying only one Cloud Identity domain.

upvoted 1 times

✉  **Shenannigan** 5 months ago

Selected Answer: C

Answer is definitely C

Please review this as it seems to be looked over in the other comments

<https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains>

(a google account that isn't part of the domain will not work unless you specifically allow exceptions at the project level and that was not defined in the answers)

upvoted 1 times

✉ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: C

i believe it is C

upvoted 1 times

✉ **thaliath** 9 months ago

Correct answer is C. A is not correct. You can not ask someone to create a personal google account. He/she has no obligation to do so
upvoted 1 times

✉ **alex000** 9 months ago

Selected Answer: A

From: https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors

"The organization creates a Google group for these external auditors and adds the current auditor to the group. This group is monitored and is typically granted access to the dashboard application.

During normal access, the auditors' Google group is only granted access to view the historic logs stored in BigQuery. If any anomalies are discovered, the group is granted permission to view the actual Cloud Logging Admin Activity logs via the dashboard's elevated access mode. At the end of each audit period, the group's access is then revoked."

upvoted 1 times

✉ **alex000** 9 months, 1 week ago

Selected Answer: A

Correct Answer is (A)

https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors

nowhere I see "temporary account"

upvoted 1 times

✉ **jrlsl1991** 8 months, 1 week ago

Honestly that seems to me a terrible reasoning. Nowhere it says "ask the auditor to give you their Google account" ¯_(ツ)_/¯.

I think it's C, because that same scenario says "The organization creates a Google group for these external auditors and adds the current auditor to the group.". Plus, that specific scenario talks about logs in BigQuery, and in the question they are asking to review all the resources.

upvoted 2 times

✉ **alexandercamachop** 1 year, 2 months ago

Selected Answer: C

C - https://cloud.google.com/iam/docs/roles-audit-logging#scenario_external_auditors

upvoted 1 times

✉ **abirroy** 1 year, 2 months ago

Selected Answer: C

Create a temporary account for the auditor in Cloud Identity, and give that account the Viewer role on the project.

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

C is right

upvoted 1 times

✉ **haroldbenites** 1 year, 4 months ago

Go for C

upvoted 1 times

✉ **Rukman** 1 year, 6 months ago

Selected Answer: C

Ans: C

upvoted 1 times

✉ **Cthakker** 2 years, 2 months ago

C is right

upvoted 2 times

You have a workload running on Compute Engine that is critical to your business. You want to ensure that the data on the boot disk of this workload is backed up regularly. You need to be able to restore a backup as quickly as possible in case of disaster. You also want older backups to be cleaned automatically to save on cost. You want to follow Google-recommended practices. What should you do?

- A. Create a Cloud Function to create an instance template.
- B. Create a snapshot schedule for the disk using the desired interval.
- C. Create a cron job to create a new disk from the disk using gcloud.
- D. Create a Cloud Task to create an image and export it to Cloud Storage.

Correct Answer: B

Community vote distribution

B (100%)

 **ESP_SAP**  3 years, 1 month ago

Correct Answer (B):

Best practices for persistent disk snapshots

You can create persistent disk snapshots at any time, but you can create snapshots more quickly and with greater reliability if you use the following best practices.

Creating frequent snapshots efficiently

Use snapshots to manage your data efficiently.

Create a snapshot of your data on a regular schedule to minimize data loss due to unexpected failure.

Improve performance by eliminating excessive snapshot downloads and by creating an image and reusing it.

Set your snapshot schedule to off-peak hours to reduce snapshot time.

Snapshot frequency limits

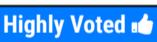
Creating snapshots from persistent disks

You can snapshot your disks at most once every 10 minutes. If you want to issue a burst of requests to snapshot your disks, you can issue at most 6 requests in 60 minutes.

If the limit is exceeded, the operation fails and returns the following error:

<https://cloud.google.com/compute/docs/disks/snapshot-best-practices>

upvoted 24 times

 **DarioFama23**  3 years, 3 months ago

B is correct for this question

upvoted 21 times

 **stepkurniawan** 3 years, 1 month ago

Question: One cannot delete the old disk when using snapshot, right?

upvoted 3 times

 **Ale1973** 3 years, 1 month ago

Snapshots and disks are independent objects on GCP, you could create a snapshot from disk and then delete the disk, the snapshot will stay in place. Actually, you could use this snapshot to create a new disk, assign to another VM, mount it, and use it (all the information that the original disk had at the time of the snapshot will still be there).

upvoted 6 times

 **Ridhanya** 1 year, 10 months ago

In snapshot schedule, there is autodelete and you can specify the days after which auto delete can happen

upvoted 5 times

 **scanner2**  1 month ago

Selected Answer: B

Create a snapshot schedule to regularly and automatically back up your zonal and regional persistent disks. Use snapshot schedules as a best practice to back up your Compute Engine workloads.

A snapshot retention policy defines how long you want to keep your snapshots.

<https://cloud.google.com/compute/docs/disks/scheduled-snapshots>

https://cloud.google.com/compute/docs/disks/scheduled-snapshots#retention_policy

upvoted 1 times

 **Captain1212** 1 month, 1 week ago

Selected Answer: B

B is the correct answer, as you can create the snapshot as per your requirement

upvoted 1 times

 **abirroy** 1 year, 2 months ago

Selected Answer: B

Create a snapshot schedule for the disk using the desired interval.

upvoted 2 times

 **csrazdan** 1 year, 3 months ago

Selected Answer: B

Snapshot is a better option because they are incremental and you can configure them to consolidate and delete snapshots that are not required for recovery. Image can also provide this functionality but the image is full backup which is inefficient in cases where the content of the file system is changing frequently.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is right

upvoted 2 times

 **haroldbenites** 1 year, 4 months ago

Go for B

upvoted 1 times

 **Rukman** 1 year, 6 months ago

Selected Answer: B

Ans: B

upvoted 2 times

 **ryzior** 1 year, 7 months ago

say no more:

<https://cloud.google.com/compute/docs/disks/scheduled-snapshots>

"Use snapshot schedules as a best practice to back up your Compute Engine workloads."

upvoted 4 times

 **alaahakim** 1 year, 10 months ago

The right Ans is : B

upvoted 1 times

 **mcaromit** 2 years, 4 months ago

B is correct

upvoted 1 times

 **[Removed]** 2 years, 6 months ago

B is correct. Create a snapshot schedule for the disk using the desired interval.

upvoted 1 times

 **cloud_guru** 2 years, 6 months ago

B is correct due to snapshots being the perfect solution to the problem question

upvoted 1 times

 **GCP_Student1** 2 years, 7 months ago

B. Create a snapshot schedule for the disk using the desired interval.

upvoted 2 times

 **EABDAJA** 2 years, 7 months ago

B is Correct

upvoted 2 times

 **Vikash211982** 2 years, 8 months ago

B is correct

upvoted 1 times

You need to assign a Cloud Identity and Access Management (Cloud IAM) role to an external auditor. The auditor needs to have permissions to review your Google Cloud Platform (GCP) Audit Logs and also to review your Data Access logs. What should you do?

- A. Assign the auditor the IAM role roles/logging.privateLogViewer. Perform the export of logs to Cloud Storage.
- B. Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.
- C. Assign the auditor's IAM user to a custom role that has logging.privateLogEntries.list permission. Perform the export of logs to Cloud Storage.
- D. Assign the auditor's IAM user to a custom role that has logging.privateLogEntries.list permission. Direct the auditor to also review the logs for changes to Cloud IAM policy.

Correct Answer: C*Community vote distribution*

B (100%)

  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (B):

Background

Google Cloud provides Cloud Audit Logs, which is an integral part of Cloud Logging. It consists of two log streams for each project: Admin Activity and Data Access.

Admin Activity logs contain log entries for API calls or other administrative actions that modify the configuration or metadata of resources. Admin Activity logs are always enabled. There is no charge for your Admin Activity audit logs.

Data Access logs record API calls that create, modify, or read user-provided data. Data Access audit logs are disabled by default because they can be large.

logging.viewer: The logging.viewer role gives the security admin team the ability to view the Admin Activity logs.

logging.privateLogViewer : The logging.privateLogViewer role gives the ability to view the Data Access logs.

upvoted 60 times

  **ESP_SAP** 3 years, 1 month ago

Correct Answer is (B): (Continuation).

Scenario: External auditors

In this scenario, audit logs for an organization are aggregated and exported to a central sink location. A third-party auditor is granted access several

times a year to review the organization's audit logs. The auditor is not authorized to view PII data in the Admin Activity logs.

During normal access, the auditors' Google group is only granted access to view the historic logs stored in BigQuery. If any anomalies are discovered,

the group is granted permission to view the actual Cloud Logging Admin Activity logs via the dashboard's elevated access mode. At the end of each audit period,

the group's access is then revoked.

Data is redacted using Cloud DLP before being made accessible for viewing via the dashboard application.

upvoted 21 times

  **ESP_SAP** 3 years, 1 month ago

Correct Answer is (B): (Continuation).

The table below explains IAM logging roles that an Organization Administrator can grant to the service account used by the dashboard,

as well as the resource level at which the role is granted:

logging.viewer Organization Dashboard service account The logging.viewer role permits the service account to read the Admin Activity logs in Cloud Logging.

bigquery.dataViewer BigQuery dataset Dashboard service account The bigquery.dataViewer role permits the service account used by the dashboard application

to read the exported Admin Activity logs.

upvoted 19 times

  **DarioFama23**  3 years, 3 months ago

for me B is the correct answer..

upvoted 17 times

  **Eshkrkrkr** 2 years, 11 months ago

Yes, B is correct because:

1) Question doesn't ask us to export and store logs for any long period of time.

2) Custom role with only logging.privateLogEntries.list permission won't let the auditor to access Log Explorer at all (https://cloud.google.com/logging/docs/access-control#console_permissions - Minimal read-only access: logging.logEntries.list)

upvoted 6 times

 **scanner2** Most Recent 1 month ago

Selected Answer: B

- The Logs Viewer role (roles/logging.viewer) gives you read-only access to Admin Activity, Policy Denied, and System Event audit logs. If you have just this role, you cannot view Data Access audit logs that are in the _Default bucket.
 - The Private Logs Viewer role(roles/logging.privateLogViewer) includes the permissions contained in roles/logging.viewer, plus the ability to read Data Access audit logs in the _Default bucket.
- Therefore, no need to export logs to Cloud storage explicitly, the _Default bucket sink access is already provided from the above role.
https://cloud.google.com/iam/docs/audit-logging#audit_log_permissions
- upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: B

b is the correct answer
upvoted 1 times

 **Neha_Pallavi** 1 month, 2 weeks ago

B. Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.
upvoted 1 times

 **MilanRajGupta** 2 months, 1 week ago

This answer is similar to answer choice B, but it suggests creating a custom role for the auditor that includes the "logging.privateLogEntries.list" permission. While this would provide the auditor with access to the necessary logs, directing them to also review Cloud IAM policy logs is not relevant to their request. Therefore, this answer is also not correct.
upvoted 1 times

 **MilanRajGupta** 2 months, 1 week ago

Correct Ans: B
upvoted 1 times

 **anjanc** 9 months, 2 weeks ago

I also think B
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is right. Similar practice question in tutorials dojo
upvoted 1 times

 **Rutu_98** 1 year, 4 months ago

Selected Answer: B
B is correct ans
upvoted 1 times

 **luciorifa** 1 year, 7 months ago

Selected Answer: B
B is the correct answer
upvoted 3 times

 **lazyabhi606** 1 year, 9 months ago

Selected Answer: B
Correct Answer is (B)
upvoted 1 times

 **maggieli** 1 year, 11 months ago

Correct Answer is B.
upvoted 1 times

 **ankatsu2010** 2 years ago

A is the correct answer. Exporting logging data to Cloud Storage is ideal, and 'Cloud IAM Policy' is not mentioned in this question.
upvoted 3 times

 **mcaromit** 2 years, 4 months ago

B is correct
upvoted 1 times

 **[Removed]** 2 years, 6 months ago

B is correct. Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.
upvoted 1 times

 **yubi69** 2 years, 6 months ago

answer is B
upvoted 2 times

You are managing several Google Cloud Platform (GCP) projects and need access to all logs for the past 60 days. You want to be able to explore and quickly analyze the log contents. You want to follow Google-recommended practices to obtain the combined logs for all projects. What should you do?

- A. Navigate to Stackdriver Logging and select resource.labels.project_id="*"
- B. Create a Stackdriver Logging Export with a Sink destination to a BigQuery dataset. Configure the table expiration to 60 days.
- C. Create a Stackdriver Logging Export with a Sink destination to Cloud Storage. Create a lifecycle rule to delete objects after 60 days.
- D. Configure a Cloud Scheduler job to read from Stackdriver and store the logs in BigQuery. Configure the table expiration to 60 days.

Correct Answer: B

Reference:

<https://cloud.google.com/blog/products/gcp/best-practices-for-working-with-google-cloud-audit-logging>

Community vote distribution

B (100%)

✉️  **Verve**  3 years, 2 months ago

Its B.

upvoted 25 times

✉️  **[Removed]**  3 years ago

The question is to view log past 60 days. B, c, D talks about deleting an object or truncation of table data

upvoted 11 times

✉️  **[Removed]** 3 years ago

Answer should be A

upvoted 3 times

✉️  **[Removed]** 3 years ago

Also A specifically talks about aggregation

upvoted 4 times

✉️  **[Removed]** 3 years ago

Also by default, you have a lot of flexibility when viewing logging in stack driver , to filter and query.

upvoted 2 times

✉️  **xtian2900** 3 years ago

what about minimum retention is 30 days ? is it true ?

upvoted 3 times

✉️  **[Removed]** 3 years ago

Ur correct so minimally is 30 for data access logs <https://cloud.google.com/logging/quotas>
then B is the way to go.

upvoted 3 times

✉️  **scanner2**  1 month ago

Selected Answer: B

Provides storage of log entries in BigQuery datasets. You can use big data analysis capabilities on the stored logs. Logging sinks stream logging data into BigQuery in small batches, which lets you query data without running a load job.

You can set a default table expiration time at the dataset level, or you can set a table's expiration time when the table is created. A table's expiration time is often referred to as "time to live" or TTL. When a table expires, it is deleted along with all of the data it contains.

https://cloud.google.com/logging/docs/export/configure_export_v2#overview

https://cloud.google.com/bigquery/docs/managing-tables#updating_a_tables_expiration_time

upvoted 1 times

✉️  **Captain1212** 1 month ago

Selected Answer: B

B is the correct answer, we can use bq to get 60 days logs and analyse

upvoted 1 times

✉️  **Neha_Pallavi** 1 month, 2 weeks ago

B. Create a Stackdriver Logging Export with a Sink destination to a BigQuery dataset. Configure the table expiration to 60 days.
upvoted 1 times

✉️  **Prat25200607** 6 months, 1 week ago

Selected Answer: B

<https://cloud.google.com/architecture/security-log-analytics>

upvoted 1 times

 **sai_learner** 1 year, 2 months ago

All options are wrong , they are talking about deletion after 60 days, but questions asks us to analyse logs of past 60 days
upvoted 4 times

 **FeaRoX** 8 months, 1 week ago

You are absolutely wrong - meaning of "past 60 days" is same as "last 60 days" in that sentence.
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is right for sure
upvoted 1 times

 **Tirthankar17** 1 year, 4 months ago

Selected Answer: B

B is the correct answer.
upvoted 2 times

 **dttncl** 1 year, 12 months ago

I believe B is the answer.

All that matters in this scenario is the logs for the past 60 days.

We can use BigQuery to analyze contents so C is incorrect. We need to configure a BQ as the sink for the logs export so we can query and analyze log data in the future. Therefore D is incorrect.
<https://cloud.google.com/logging/docs/audit/best-practices#export-best-practices>

Since we only care about the logs within 60 days, we can set the expiration time to 60 to retain only the logs within that time frame. Once data is beyond 60 days old, it wouldn't be included in future analyzations.

https://cloud.google.com/bigquery/docs/managing-tables#updating_a_tables_expiration_time

upvoted 6 times

 **ryzior** 1 year, 7 months ago

I think here we have the case described in details:

<https://cloud.google.com/architecture/exporting-stackdriver-logging-for-security-and-access-analytics>

upvoted 1 times

 **ankatsu2010** 2 years ago

D should be the correct answer. To 'quickly analyze', you need to use BQ, next, you always need access to the logs 'for past 60days'. This means you have to export logs on a daily basis. You don't want to do this job manually right?

upvoted 1 times

 **ankatsu2010** 2 years ago

My apologies, B is correct... 'Sink' can route logging data to BQ automatically.

upvoted 3 times

 **AD_0525** 2 years, 3 months ago

B is the correct one, option A does not give you the flexibility to analyze.

upvoted 3 times

 **mcaromit** 2 years, 4 months ago

B is correct as analysis of the log contents is a key requirement

upvoted 2 times

 **tifo16** 2 years, 5 months ago

Firstly i thought that A was correct but when considering "default" retention period which is 30 days, i go for B.

<https://cloud.google.com/blog/products/it-ops/best-practices-for-working-with-google-cloud-audit-logging>

Log entries are held in Stackdriver Logging for a limited time known as the retention period. After that, the entries are deleted. To keep log entries longer, you need to export them outside of Stackdriver Logging by configuring log sinks.

upvoted 5 times

 **Crad** 2 years, 5 months ago

I think it's C.

If you set expiration date to 60days on a table then that table will be deleted after 60 days of creation.

Let's say you created some logs on 31st day - they'll be gone within 29 days and not 60.

upvoted 1 times

 **Crad** 2 years, 5 months ago

The output is in JSON so it's still possible to analyze the logs using BigQuery easily as an external source of data.

upvoted 1 times

 **[Removed]** 2 years, 6 months ago

B is correct. Create a Stackdriver Logging Export with a Sink destination to a BigQuery dataset. Configure the table expiration to 60 days.

upvoted 1 times

 **GCP_Student1** 2 years, 6 months ago

B. Create a Stackdriver Logging Export with a Sink destination to a BigQuery dataset. Configure the table expiration to 60 days.

upvoted 4 times

You need to reduce GCP service costs for a division of your company using the fewest possible steps. You need to turn off all configured services in an existing GCP project. What should you do?

- A. 1. Verify that you are assigned the Project Owners IAM role for this project. 2. Locate the project in the GCP console, click Shut down and then enter the project ID.
- B. 1. Verify that you are assigned the Project Owners IAM role for this project. 2. Switch to the project in the GCP console, locate the resources and delete them.
- C. 1. Verify that you are assigned the Organizational Administrator IAM role for this project. 2. Locate the project in the GCP console, enter the project ID and then click Shut down.
- D. 1. Verify that you are assigned the Organizational Administrators IAM role for this project. 2. Switch to the project in the GCP console, locate the resources and delete them.

Correct Answer: C

Community vote distribution

A (100%)

-  **DarioFama23** Highly Voted 3 years, 3 months ago
for me is A the correct answer
upvoted 42 times
 -  **shafiqeee1** Highly Voted 3 years, 2 months ago
A - I reproduced in my project
upvoted 17 times
 -  **scanner2** Most Recent 1 month ago
Selected Answer: A
<https://cloud.google.com/resource-manager/docs/access-control-proj#permissions>
https://cloud.google.com/resource-manager/docs/creating-managing-projects#shutting_down_projects
upvoted 2 times
 -  **Captain1212** 1 month ago
Selected Answer: A
a is the correct answer
upvoted 1 times
 -  **Neha_Pallavi** 1 month, 2 weeks ago
Verify that you are assigned the Project Owners IAM role for this project. 2. Locate the project in the GCP console, click Shut down and then enter the project ID.
upvoted 1 times
 -  **vinodthakur49** 3 months ago
Selected Answer: A
A is the correct answer
upvoted 1 times
 -  **Shenannigan** 5 months ago
Selected Answer: A
Answer is A
<https://support.google.com/googleapi/answer/6251787?hl=en#zippy=%2Cshut-down-a-project>
upvoted 2 times
 -  **sabrinakloud** 5 months, 3 weeks ago
Selected Answer: A
option A
upvoted 1 times
 -  **sabrinakloud** 5 months, 3 weeks ago
i believe it's option A.
- roles/owner Owner All Editor permissions and permissions for the following actions:
Manage roles and permissions for a project and all resources within the project.
Set up billing for a project.
upvoted 1 times
-  **researched answer boi** 8 months, 3 weeks ago

According to "https://cloud.google.com/resource-manager/docs/access-control-org#resourcemanager.organizationAdmin" and "https://cloud.google.com/resource-manager/docs/access-control-proj#basic_roles", only the project owner (and project deleter (roles/resourcemanager.projectDeleter)) can delete a project.
So, answer A is the technically correct one.

upvoted 1 times

✉ **roaming_panda** 9 months, 3 weeks ago

Selected Answer: A

project not org , A all d way

upvoted 2 times

✉ **patashish** 1 year, 3 months ago

A is right

Hint : You need to turn off all configured services in an ***existing GCP project***.

So C and D out from selection

upvoted 1 times

✉ **RanjithK** 1 year, 3 months ago

Selected Answer: A

Tried and tested

upvoted 1 times

✉ **keep_it_on** 1 year, 3 months ago

correct ans is A

upvoted 2 times

✉ **taiyi078** 1 year, 3 months ago

<https://cloud.google.com/run/docs/tutorials/gcloud>

Clean up

In the dialog, type the project ID, and then click Shut down to delete the project.

upvoted 2 times

✉ **AzureDP900** 1 year, 3 months ago

A is right

upvoted 2 times

✉ **Rutu_98** 1 year, 4 months ago

Selected Answer: A

It's A

Because we have to follow the least required permission.

Also here given is that they want to reduce the services for a GCP project so the Project Owner role would be sufficient.

upvoted 3 times

You are configuring service accounts for an application that spans multiple projects. Virtual machines (VMs) running in the web-applications project need access to BigQuery datasets in `crm-databases-proj`. You want to follow Google-recommended practices to give access to the service account in the web-applications project. What should you do?

- A. Give `project owner` for web-applications appropriate roles to `crm-databases-proj`.
- B. Give `project owner` role to `crm-databases-proj` and the web-applications project.
- C. Give `project owner` role to `crm-databases-proj` and `bigrquery.dataViewer` role to web-applications.
- D. Give `bigrquery.dataViewer` role to `crm-databases-proj` and appropriate roles to web-applications.

Correct Answer: C

Reference:

<https://cloud.google.com/blog/products/gcp/best-practices-for-working-with-google-cloud-audit-logging>

Community vote distribution

D (61%)

C (39%)

✉  **ezat**  3 years, 3 months ago

D cuz u just need read for DB at the other project
upvoted 30 times

✉  **[Removed]** 2 years, 9 months ago

Question didn't specify if the required access is Read only or more, its saying "access" which could be write permissions as well. I will go with C
upvoted 2 times

✉  **[Removed]** 2 years, 9 months ago

U r right, it D. why to give "project owner" as stated on C. correct answer is D
upvoted 5 times

✉  **GCPACE2020** 2 years, 3 months ago

but why giving `bigrquery.dataViewer` to `crm-databases-proj`. we should give for web-application.
upvoted 8 times

✉  **S_marquez** 1 year, 6 months ago

You can technically give `bigrquery.dataViewer` to `crm-databases-proj` service account then create a Key and use that key on the VMs, there for making it correct to use D as answer but is way to dumb I would prefer C BUUUUUUT WHY would I give Project Owner to `crm-databases-proj`? they really do not evaluate your knowledge
upvoted 3 times

✉  **dttnc1** 1 year, 12 months ago

It is D because you're right, the question doesn't specify any specific kind of access, however, we need to follow the principle of least-privilege. Hence, we can only assume that read-only access is needed.

`bigrquery.dataViewer` should be assigned to the group of analysts in the `crm-databases-proj` project.

https://cloud.google.com/bigquery/docs/access-control-examples#read_access_to_data_in_a_different_project
upvoted 9 times

✉  **DarioFama23** 3 years, 3 months ago

U re right, D is the correct answee
upvoted 4 times

✉  **tavva_prudhvi** 2 years, 6 months ago

See the option correctly, as the web app needs access to the big query datasets we have to give access to the web app the data viewer role to only read the datasets! Hence, C
upvoted 6 times

✉  **DarioFama23**  3 years, 3 months ago

C is correct..
upvoted 11 times

✉  **BigQuery** 1 year, 6 months ago

THAT SO DUM
upvoted 4 times

✉  **GCPACE2020** 2 years, 3 months ago

But why giving project owner role to `crm-databases-proj`?
upvoted 2 times

✉  **Captain1212**  1 month ago

Selected Answer: D

D is the correct answer, because all other option giveing access to project owner
upvoted 1 times

 **SanjeevKumar1983** 1 month, 1 week ago

Corrcet Answer is D.

Lets just read the options D this way, then it makes sense

Give service account the bigquery.dataViewer role to crm-databases-proj and service account the appropriate roles to web-applications.
upvoted 1 times

 **Ahmed_Y** 1 month, 3 weeks ago

Selected Answer: D

Thanks guys for making that clear for me.

Now simply guys, among all the answers, D is giving to the web-application proj the appropriate role, while giving the crm-databases-proj the least privilege role.

upvoted 1 times

 **Nxt_007** 2 months ago

Selected Answer: D

Correct answer is D

As basic roles (including Owner) should not be used in production environment:

upvoted 1 times

 **yichuan** 2 months, 2 weeks ago

Selected Answer: C

C is correct..

upvoted 1 times

 **creativenets** 3 months, 3 weeks ago

I dont get the question. It says "web-applications project need access to BigQuery datasets in crm-databases-proj"

And all you folks stating C or D is the correct one. Why would we want to give those permissions to the DB? When the question clearly states that the web-app is the one that needs access to the DB?

upvoted 1 times

 **KC_go_reply** 4 months ago

Selected Answer: C

It says 'web-applications project need access to BigQuery datasets in crm-databases-proj'. Therefore, give web-applications the BigQuery Data Viewer role - not the other way around. Why would crm-databases-proj need this role in this situation?

upvoted 1 times

 **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: D

I believe that the correct answer is option D. Although the web application requires the bigquery.dataViewer role, the option D mentions "appropriate roles to web applications" and the appropriate role in this case is indeed bigquery.dataViewer. It is not recommended to give the project owner role to crm-databases-proj as it grants too many permissions. Google's best practice is to minimize the number of permissions granted, so option D aligns with this principle.

upvoted 3 times

 **xaqanik** 7 months, 2 weeks ago

Selected Answer: D

why we should grant 'project owner' role? we only need to give access for BigQuery
i choose Answer D

upvoted 1 times

 **jrlsl1991** 8 months, 1 week ago

Selected Answer: D

Honestly none of the answers make full sense to me but Best Practices for Service Accounts basically ask to keep the roles/permissions as low as possible since there are potential bad scenarios where a service account can be used to jeopardize a project, sometimes without leaving traces. Based on this, none of the Project-owner answers would be correct, so I'm going for D.

upvoted 2 times

 **vivekvj** 5 months, 2 weeks ago

true. none of the options make sense but project owner role for no reason is an absolute NO. So D.

upvoted 1 times

 **SK1990** 9 months, 2 weeks ago

Selected Answer: D

D is the best answer.

upvoted 1 times

 **ChristN** 10 months, 1 week ago

Selected Answer: C

C is the correct answer.

bigquery.dataViewer role to crm-databases-proj doesn't full fill the requirement. only the web application needs this access.

upvoted 2 times

 **fragment137** 10 months, 1 week ago

This question is misleading. the requirements are for the web -application service account to have "access" to the BQ Datasets, but it doesn't specify what they need to do. Principal of Least Privilege would lead you to think they just need Viewer. Problem is the answers are also misleading. C gives Project Owner to the dataset project, and -appropriate- permissions to the web-application service account. It seems to indicate that it's giving project owner to the project itself, which makes no sense. D gives dataViewer to the BQ project, which doesn't make sense either because it's the web-application service account that needs access to BQ.

I think this is a poorly worded question and poorly worded answers personally.

upvoted 6 times

 **Untamables** 11 months, 3 weeks ago

Selected Answer: C

All options are hard to understand.

Only C mentions bigquery.dataViewer for web-applications

upvoted 3 times

 **AwesomeGCP** 12 months ago

Selected Answer: D

No doubt about this Question!!!

Answer is simple and its D. Give bigquery.dataViewer role to crm-databases-proj and appropriate roles to web-applications.

END OF THE DISCUSSION NOW!!!

upvoted 2 times

An employee was terminated, but their access to Google Cloud was not removed until 2 weeks later. You need to find out if this employee accessed any sensitive customer information after their termination. What should you do?

- A. View System Event Logs in Cloud Logging. Search for the user's email as the principal.
- B. View System Event Logs in Cloud Logging. Search for the service account associated with the user.
- C. View Data Access audit logs in Cloud Logging. Search for the user's email as the principal.
- D. View the Admin Activity log in Cloud Logging. Search for the service account associated with the user.

Correct Answer: B

Community vote distribution



✉️ **sanvit** 4 weeks ago

Kindly share the all question (kurhekavita@gmail.com)able to see only 120 questions.

I have planned exam in the next week

upvoted 1 times

✉️ **Captain1212** 1 month ago

Selected Answer: C

Option C is more correct , as data access logs contain API , from this you can check for it

upvoted 1 times

✉️ **Neha_Pallavi** 1 month, 2 weeks ago

option C is correct :

Data Access audit logs contain API calls that read the configuration or metadata of resources, as well as user-driven API calls that create, modify, or read user-provided resource data. Kindly share the all question (javac.krishnan@gmail.com)able to see only 120 questions

upvoted 2 times

✉️ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: C

I think option C is correct :

Data Access audit logs contain API calls that read the configuration or metadata of resources, as well as user-driven API calls that create, modify, or read user-provided resource data.

upvoted 3 times

✉️ **Bobbybash** 7 months, 4 weeks ago

Selected Answer: C

C. View Data Access audit logs in Cloud Logging. Search for the user's email as the principal.

Data Access audit logs record all activity related to accessing or modifying data, including reading, writing, and deleting operations. By searching for the terminated employee's email as the principal, you can see if they accessed any sensitive customer information after their termination. System Event Logs and Admin Activity logs may not have the details of the data accessed, so Data Access audit logs are the most appropriate option in this scenario.

upvoted 1 times

✉️ **mrvergara** 9 months ago

Selected Answer: A

<https://cloud.google.com/logging/docs/audit>.

Data Access audit logs are disabled by default

upvoted 1 times

✉️ **Cornholio_LMC** 1 year ago

had this question today

upvoted 2 times

✉️ **abirroy** 1 year, 2 months ago

Selected Answer: C

View Data Access audit logs in Cloud Logging. Search for the user's email as the principal

upvoted 2 times

✉️ **Rutu_98** 1 year, 4 months ago

Selected Answer: C

ANSWER IS C

As we want to find out whether the user has accessed the data or not , so Data Acess Logs would be correct option to view that

upvoted 2 times

✉️ **akshaychavan7** 1 year, 4 months ago

upvoted 1 times

✉ **bigbenben** 1 year, 5 months ago

Selected Answer: C

C. <https://cloud.google.com/logging/docs/audit#data-access>

upvoted 3 times

✉ **rsuresh27** 1 year, 5 months ago

Selected Answer: C

C is the correct answer. We are trying to find out if any sensitive data was accessed. Data access logs are the only logs that show this. C is the only option that mentions data access logs.

upvoted 4 times

✉ **rtnk22** 1 year, 5 months ago

Shouldn't the correct option be A here?

What does service account have to do here.

upvoted 2 times

✉ **amrith501** 1 year, 5 months ago

It should be C User activity generally come under audit logs

upvoted 3 times

You need to create a custom IAM role for use with a GCP service. All permissions in the role must be suitable for production use. You also want to clearly share with your organization the status of the custom role. This will be the first version of the custom role. What should you do?

- A. Use permissions in your role that use the 'supported' support level for role permissions. Set the role stage to ALPHA while testing the role permissions.
- B. Use permissions in your role that use the 'supported' support level for role permissions. Set the role stage to BETA while testing the role permissions.
- C. Use permissions in your role that use the 'testing' support level for role permissions. Set the role stage to ALPHA while testing the role permissions.
- D. Use permissions in your role that use the 'testing' support level for role permissions. Set the role stage to BETA while testing the role permissions.

Correct Answer: C

Community vote distribution

A (100%)

✉️  **raksteer**  3 years, 2 months ago

You need a custom role with permissions supported in prod and you want to publish the status of the role.

<https://cloud.google.com/iam/docs/custom-roles-permissions-support>

SUPPORTED The permission is fully supported in custom roles.

TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.

NOT_SUPPORTED The permission is not supported in custom roles.

You can't use TESTING as it is not good for prod. And you need first version which should be ALPHA. Answer should be A.

upvoted 88 times

✉️  **passmep1s** 3 years, 1 month ago

good job

upvoted 2 times

✉️  **BigQuery** 1 year, 6 months ago

WAY TO GO. VERY CLEAR EXP INDEED

upvoted 1 times

✉️  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (A):

Testing and deploying

Custom roles include a launch stage, which is stored in the stage property for the role. The launch stage is informational; it helps you keep track of whether each role is ready for widespread use.

Each custom role can have one of the following launch stages:

Launch stages

ALPHA The role is still being developed or tested, or it includes permissions for Google Cloud services or features that are not yet public. It is not ready for widespread use.

BETA The role has been tested on a limited basis, or it includes permissions for Google Cloud services or features that are not generally available.

GA The role has been widely tested, and all of its permissions are for Google Cloud services or features that are generally available.

upvoted 30 times

✉️  **ESP_SAP** 3 years, 1 month ago

Correct Answer is (A): Continuation

Support levels for permissions in custom roles

You can include many, but not all, Identity and Access Management (IAM) permissions in custom roles. Each permission has one of the following support levels:

Support level Description

SUPPORTED The permission is fully supported in custom roles.

TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.

NOT_SUPPORTED The permission is not supported in custom roles.

The first version of the Custom Role is ALPHA then suitable to productions all permissions in "Supported"...

upvoted 13 times

□  **GCP_Student1** 2 years, 7 months ago

ESP_SAP

There is a discrepancy between your first post and the second post. Compare these two sentences;

1st POST - ALPHA The role is still being developed or tested, or it includes permissions for Google Cloud services or features that are not yet public. It is not ready for widespread use.

2nd POST - SUPPORTED The permission is fully supported in custom roles.

Are you still going to go with A ?

upvoted 2 times

□  **learn_GCP** 1 year ago

Here ALPHA is for Google cloud feature, only informational. given to identify whether the feature is fully available as a service. and SUPPORTED -- is for a custom role which is supported by Google cloud, meaning any support is provided by Google cloud

upvoted 2 times

□  **scanner2** Most Recent 1 month ago

Selected Answer: A

A is correct answer.

<https://cloud.google.com/iam/docs/roles-overview#custom-role-supported-permissions>

<https://cloud.google.com/iam/docs/roles-overview#custom-role-testing-deploying>

upvoted 1 times

□  **Captain1212** 1 month ago

Selected Answer: A

Answer is A, as you need for production and you dont neeed testing for it and you need first version , so it will be ALPHA , not beta

upvoted 1 times

□  **Neha_Pallavi** 1 month, 2 weeks ago

You need a custom role with permissions supported in prod and you want to publish the status of the role.

<https://cloud.google.com/iam/docs/custom-roles-permissions-support>

SUPPORTED The permission is fully supported in custom roles.

upvoted 1 times

□  **Tofer2022** 10 months, 3 weeks ago

why not B?

upvoted 1 times

□  **temple1305** 6 months, 3 weeks ago

Because ...FIRST VERSION... is ALPHA.

upvoted 1 times

□  **theBestStudent** 1 year, 2 months ago

Selected Answer: A

It must be suitable for production so Supported permissions only. Plus, it is your first version of the custom role, so you need to check if all is good, then ALPHA.

upvoted 1 times

□  **abirroy** 1 year, 2 months ago

Selected Answer: A

Use permissions in your role that use the 'supported' support level for role permissions. Set the role stage to ALPHA while testing the role permissions.

upvoted 1 times

□  **AzureDP900** 1 year, 3 months ago

A is right

upvoted 1 times

□  **Raz0r** 1 year, 8 months ago

Selected Answer: A

A is the only right solution.

upvoted 2 times

□  **sid0127** 1 year, 10 months ago

Selected Answer: A

A is the correct answer

upvoted 2 times

□  **sudav** 2 years, 1 month ago

A -

SUPPORTED -The permission is fully supported in custom roles.

role stage, the stage transitions from ALPHA -> BETA -> GA

The only option that satisfies "ALPHA" stage with "SUPPORTED" support level is

Use permissions in your role that use the SUPPORTED support level for role permissions. Set the role stage to ALPHA while testing the role permissions

upvoted 3 times

□  **mcaromit** 2 years, 4 months ago

A is correct

upvoted 1 times

 **[Removed]** 2 years, 6 months ago

A is correct. Use permissions in your role that use the "supported" support level for role permissions. Set the role stage to ALPHA while testing the role permissions.

upvoted 1 times

 **yuvি69** 2 years, 6 months ago

answer is A. because it contains SUPPORTED which we must see when creating custom roles and as it first version we must set it to ALPHA

upvoted 1 times

 **cloud_guru** 2 years, 6 months ago

Looks like A to be because the role is the first iteration and still under testing, we need to select 'Alpha' for the custom role. Additionally, we shouldn't use 'testing' permissions in the role after they said not to

upvoted 1 times

 **SSunny** 2 years, 7 months ago

A:

Supporting Points:

a. The status of the custom rule is to be presented to the organization. And since this is yet in testing state, the custom rule state will be Alpha.

b. The Support level, must be SUPPORTED so as to be used in production.

If the support level is at TESTING, then in that case - the custom role must not be used as its results may change and unexpected behaviour may be expected.

SUPPORTED The permission is fully supported in custom roles.

TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.

NOT_SUPPORTED The permission is not supported in custom roles.

<https://cloud.google.com/iam/docs/custom-roles-permissions-support>

upvoted 3 times

Your company has a large quantity of unstructured data in different file formats. You want to perform ETL transformations on the data. You need to make the data accessible on Google Cloud so it can be processed by a Dataflow job. What should you do?

- A. Upload the data to BigQuery using the bq command line tool.
- B. Upload the data to Cloud Storage using the gsutil command line tool.
- C. Upload the data into Cloud SQL using the import function in the console.
- D. Upload the data into Cloud Spanner using the import function in the console.

Correct Answer: B

Reference:

<https://cloud.google.com/solutions/performing-etl-from-relational-database-into-bigquery>

Community vote distribution

B (100%)

✉ **DarioFama23** Highly Voted 3 years, 3 months ago

B looks correct. Key work unstructured data
upvoted 24 times

✉ **obeythefist** 1 year, 7 months ago

Also "different" file formats, this further supports B as the correct choice.
upvoted 1 times

✉ **scanner2** Most Recent 1 month ago

Selected Answer: B

Key term is "unstructured data in different file formats". Except B, remaining options are suitable for structured data. So, correct answer is B.
upvoted 1 times

✉ **Captain1212** 1 month ago

B is correct for Unstructurd DAta its Cloud storage
upvoted 1 times

✉ **shivampriya11** 1 month, 3 weeks ago

why this page ask for contributer access.. i can not access whole questions
upvoted 1 times

✉ **Nxt_007** 2 months ago

Selected Answer: B

Option B is correct
Cloud Storage is a scalable and cost-effective object storage service that can hold unstructured data of various file formats. Before performing ETL (Extract, Transform, Load) transformations, it's often beneficial to store the raw data in a centralized location, like Cloud Storage.
upvoted 2 times

✉ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: B

Answer B
upvoted 1 times

✉ **Untamables** 11 months, 3 weeks ago

Selected Answer: B

Cloud Storage as a datalake
upvoted 1 times

✉ **VietmanOfficiel** 1 year, 1 month ago

Answer B :
"large quantity" : Cloud Storage or BigQuery
"files" a file is nothing but an Object

So Cloud Storage is the better option.

upvoted 3 times

✉ **abirroy** 1 year, 2 months ago

Selected Answer: B

Upload the data to Cloud Storage using the gsutil command line tool.
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B looks correct
upvoted 1 times

 **pfabio** 1 year, 4 months ago

Selected Answer: B

For unstructured data use cloud storage. Use Big Query for analytics, data warehouse with structured data
upvoted 4 times

 **PAUGURU** 1 year, 5 months ago

Selected Answer: B
B looks correct
upvoted 2 times

 **Akash7** 1 year, 6 months ago

Answer is B, Cloud Storage for unstructured data
upvoted 1 times

 **Rukman** 1 year, 6 months ago

Selected Answer: B
Ans: B

We can upload unstructured data to Cloud storage not to bigquery.
upvoted 3 times

 **associatecloudexamuser** 2 years, 2 months ago

B is correct. Since GCS is a datalake
upvoted 2 times

 **mcaromit** 2 years, 4 months ago

B is correct
upvoted 1 times

 **[Removed]** 2 years, 6 months ago

B is correct. Upload the data to Cloud Storage using the gsutil command line tool.
upvoted 2 times

You need to manage multiple Google Cloud projects in the fewest steps possible. You want to configure the Google Cloud SDK command line interface (CLI) so that you can easily manage multiple projects. What should you do?

- A. 1. Create a configuration for each project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.
- B. 1. Create a configuration for each project you need to manage. 2. Use gcloud init to update the configuration values when you need to work with a non-default project
- C. 1. Use the default configuration for one project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.
- D. 1. Use the default configuration for one project you need to manage. 2. Use gcloud init to update the configuration values when you need to work with a non-default project.

Correct Answer: D

Community vote distribution

A (100%)

 **SSunny**  2 years, 7 months ago

A

Cloud SDK comes with a default configuration. To create multiple configurations, use gcloud config configurations create, and gcloud config configurations activate to switch between them.

<https://cloud.google.com/sdk/gcloud/reference/config/set>
upvoted 28 times

 **GCP_Student1**  2 years, 7 months ago

A. 1. Create a configuration for each project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.

upvoted 8 times

 **udehoriaku**  3 weeks, 5 days ago

please send me the dump question pdf here udehoriaku@gmail.com
upvoted 1 times

 **scanner2** 1 month ago

Selected Answer: A

The gcloud config command group lets you set, view and unset properties used by Google Cloud CLI. A configuration is a set of properties that govern the behavior of gcloud and other Google Cloud CLI tools. The initial default configuration is set when gcloud init is run. You can create additional named configurations using gcloud init or gcloud config configurations create. To switch between configurations, use gcloud config configurations activate.

<https://cloud.google.com/sdk/gcloud/reference/config>
<https://cloud.google.com/sdk/gcloud/reference/config/configurations/activate>

upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: A

A is the correct answer, as it comes with the default configuration and you don't need to update it
upvoted 1 times

 **Neha_Pallavi** 1 month, 2 weeks ago

A. 1. Create a configuration for each project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects. KINDLY SHARE COMPLETE QUESTION
upvoted 1 times

 **xaqanik** 8 months ago

Selected Answer: A

Google Cloud SDK allows you to create multiple configurations for different projects, and you can easily switch between these configurations as needed. To manage multiple projects efficiently, you can create a separate configuration for each project and activate the appropriate configuration when you work with each assigned project. The gcloud config configurations create and gcloud config configurations activate commands allow you to create and activate different configurations. By using different configurations, you can ensure that your CLI commands are always executed in the correct context and against the correct project, without the need to manually change the configuration each time you switch projects.

upvoted 5 times

 **VietmanOfficiel** 1 year, 1 month ago

Selected Answer: A

1. Generate your configurations with "gcloud config configurations create <config_id> ..." then activate the one you need according to the project you are working on with "gcloud config activate <config_id>"

upvoted 1 times

 **abirroy** 1 year, 2 months ago

Selected Answer: A

Create a configuration for each project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right.

gcloud config set project PROJECT_ID

<https://cloud.google.com/sdk/gcloud/reference/config/set>

upvoted 1 times

 **Rutu_98** 1 year, 4 months ago

Selected Answer: A

A is correct Answer

upvoted 1 times

 **pfabio** 1 year, 4 months ago

Selected Answer: A

You have to create a config for each project and activate to use it

upvoted 2 times

 **luciorifa** 1 year, 7 months ago

Selected Answer: A

A is the correct answer

upvoted 1 times

 **Raju44** 1 year, 9 months ago

why we need to create config, when project is created so its config .

upvoted 1 times

 **jaffarali** 1 year, 10 months ago

Selected Answer: A

A is the right option. Multiple configurations. Activate to switch between configurations.

upvoted 1 times

 **Ridhanya** 1 year, 10 months ago

A, C seems hectic, B and D surely are eliminated

upvoted 1 times

 **rachee** 1 year, 10 months ago

A - <https://cloud.google.com/sdk/docs/configurations>

upvoted 1 times

Your managed instance group raised an alert stating that new instance creation has failed to create new instances. You need to maintain the number of running instances specified by the template to be able to process expected application traffic. What should you do?

- A. Create an instance template that contains valid syntax which will be used by the instance group. Delete any persistent disks with the same name as instance names.
- B. Create an instance template that contains valid syntax that will be used by the instance group. Verify that the instance name and persistent disk name values are not the same in the template.
- C. Verify that the instance template being used by the instance group contains valid syntax. Delete any persistent disks with the same name as instance names. Set the disks.autoDelete property to true in the instance template.
- D. Delete the current instance template and replace it with a new instance template. Verify that the instance name and persistent disk name values are not the same in the template. Set the disks.autoDelete property to true in the instance template.

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances>

Community vote distribution

A (67%) C (33%)

 **TAvenger**  2 years, 7 months ago

Really tricky question.

Ideal scenario would be

1. create new template, while creating ensure that in the new template disks.autoDelete=true, 3. delete existing persistent disks, 4. make rolling update ...

In order to switch to new template we need "Rolling update". Unfortunately, it is not mentioned.

With current options

C - not correct, we cannot update existing template

D - not correct, we cannot delete existing template when it is in use (just checked in GCP) (We need rolling update)

B - will not solve our problem without Rolling update

A - This is the only option (I know that it can be temporary) that will work without Rolling update according to <https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs>

upvoted 50 times

 **FunkyB** 8 months, 2 weeks ago

Thank you for providing the link.

upvoted 2 times

 **ShakthiGCP** 2 years, 7 months ago

Well reasoned. I'm also going with A.

upvoted 2 times

 **ESP_SAP**  3 years, 1 month ago

Correct Answer is (C):

Your instance template has set the disks.autoDelete option to false for boot persistent disks so that when a VM has been deleted (for example, because of autohealing), the persistent disk was not deleted. When the managed instance group attempted to recreate the VM with the same name, it ran into the same issue where a persistent disk already exists with the same name. Delete the existing persistent disk to resolve the immediate problem and update the instance template to set the disks.autoDelete to true if you would like boot persistent disks to be deleted alongside the instance

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances#troubleshooting>

upvoted 42 times

 **magistrum** 2 years, 9 months ago

Can't update instance templates, see below

upvoted 19 times

 **JackGlemins** 2 years, 7 months ago

https://cloud.google.com/compute/docs/instance-templates#how_to_update_instance_templates

upvoted 12 times

 **DannSecurity**  2 weeks, 6 days ago

correct answer is C. It is not updating an instance template, just validate the syntax. The same instance template was working before, so why would you need to edit it? the issue should be in the disk name.

upvoted 1 times

 **scanner2** 1 month ago

Selected Answer: A

Correct answer is A : Create an instance template that contains valid syntax which will be used by the instance group. Delete any persistent disks with the same name as instance names.

upvoted 1 times

✉ **xaqanik** 8 months ago

why we dont prefer option B? it says be sure names are different. but option A deletes discs. if we can change names why to delete them?

upvoted 1 times

✉ **Kopy** 11 months ago

Selected Answer: A

Answer is A

As many have mentioned here C and D is straight NO, C talks about updating the instance template and D talks about deleting an instance template BUT as per google documentation you cannot delete an instance template nor update it if it is in use. So A is the only one the makes sense here.

https://cloud.google.com/compute/docs/instance-templates#how_to_update_instance_templates

upvoted 3 times

✉ **Untamables** 11 months, 3 weeks ago

Selected Answer: C

Vote C

<https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs>

<https://cloud.google.com/sdk/gcloud/reference/compute/instances/set-disk-auto-delete>

upvoted 2 times

✉ **sneha1607** 1 year, 1 month ago

C is the correct answer.

1. Ensure you don't have any persistent disks with the same name as the VM instance.
2. Ensure the disk autodelete property is turned on (disks.autoDelete set to true).
3. Ensure instance template syntax is valid

upvoted 1 times

✉ **sneha1607** 1 year, 1 month ago

As described in this article, "My managed instance group keeps failing to create a VM. What's going on?"

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances#troubleshooting>

upvoted 1 times

✉ **bobthebuilder55110** 1 year, 2 months ago

Selected Answer: A

Answer : A

As many have mentioned here C and D is straight NO, C talks about updating the instance template and D talks about deleting an instance template BUT as per google documentation you cannot delete an instance template nor update it if it is in use. So A is the only one the makes sense here.

https://cloud.google.com/compute/docs/instance-templates#how_to_update_instance_templates

upvoted 5 times

✉ **gscharly** 1 year, 2 months ago

Selected Answer: A

Correct answer is A

You don't have the option to edit an existing instance.

upvoted 1 times

✉ **rdumitru** 1 year, 2 months ago

Selected Answer: A

Correct answer is A

You don't have the option to edit an existing instance.

upvoted 1 times

✉ **Bumbah** 1 year, 3 months ago

Selected Answer: A

You cannot update a template, need to create a new one.

upvoted 1 times

✉ **patashish** 1 year, 3 months ago

A is correct answer. [Instance template resources are immutable in GCP.]

How to update instance templates

Instance templates are designed to create instances with identical configurations. So you cannot update an existing instance template or change an instance template after you create it.

If you need to make changes to the configuration, create a new instance template. You can create a template based on an existing instance template, or based on an existing instance. You can also override instance template fields when creating a VM instance from an instance template.

Ref : <https://cloud.google.com/compute/docs/instance-templates#:~:text=So%20you%20cannot%20update%20an,base%20on%20an%20existing%20instance.>

upvoted 2 times

✉ **Roro_Brother** 1 year, 3 months ago

Selected Answer: C

C is the correct answer because your instance exists you don't have to recreate it
upvoted 3 times

 **davidsalomon** 10 months ago
an instance template cant be modified.
upvoted 1 times

 **deadlydeb** 1 year, 3 months ago
d IS RIGHT
upvoted 1 times

 **pfabio** 1 year, 4 months ago

Selected Answer: A
Incorrect:
B - incompleted
C - Instance template cannot be modified
D - you cannot delete a instance template while its in use
upvoted 2 times

 **somenick** 1 year, 6 months ago

Selected Answer: C
<https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs>
upvoted 2 times

Your company is moving from an on-premises environment to Google Cloud. You have multiple development teams that use Cassandra environments as backend databases. They all need a development environment that is isolated from other Cassandra instances. You want to move to Google Cloud quickly and with minimal support effort. What should you do?

- A. 1. Build an instruction guide to install Cassandra on Google Cloud. 2. Make the instruction guide accessible to your developers.
- B. 1. Advise your developers to go to Cloud Marketplace. 2. Ask the developers to launch a Cassandra image for their development work.
- C. 1. Build a Cassandra Compute Engine instance and take a snapshot of it. 2. Use the snapshot to create instances for your developers.
- D. 1. Build a Cassandra Compute Engine instance and take a snapshot of it. 2. Upload the snapshot to Cloud Storage and make it accessible to your developers. 3. Build instructions to create a Compute Engine instance from the snapshot so that developers can do it themselves.

Correct Answer: D

Community vote distribution

B (93%) 7%

-  **DarioFama23** Highly Voted 3 years, 3 months ago
B is correct for me.. launch a solution from marketplace
upvoted 27 times
-  **ESP_SAP** Highly Voted 3 years, 1 month ago
Correct Answer is (B):
<https://medium.com/google-cloud/how-to-deploy-cassandra-and-connect-on-google-cloud-platform-with-a-few-clicks-11ee3d7001d1>
upvoted 23 times
-  **xtian2900** 3 years ago
thanks, i always look for your insight
upvoted 5 times
-  **santabanta0112** 10 months, 3 weeks ago
Hi ESP_SAP I have GCP exam schedule this week if you have remaining question that is from 101 to 201 or complete list of question than can you please email me on santabantaahen@gmail.com
upvoted 1 times
-  **nightflyer** 2 years, 9 months ago
But we are moving from on premises to gcp
upvoted 1 times
-  **scanner2** Most Recent 1 month ago
Selected Answer: B
Key term is "move to Google Cloud quickly and with minimal support effort". Right away you can think of Google Cloud Marketplace in such situations.
upvoted 1 times
-  **Captain1212** 1 month ago
Selected Answer: B
B is the correct answer as , it requires the minimal effoort
upvoted 1 times
-  **creativenets** 3 months, 3 weeks ago
Selected Answer: D
D is the right answer not B. Remember the question states moving from on premise to gcloud. Meaning we already have servers build. We would need those image or snapshot so devs can create their own insance.
A and B - start from scratch
C - Doesn't give permissions to devs so you'll have to create the instance (more support effort)
upvoted 1 times
-  **geeroylenkins** 2 months, 4 weeks ago
nah - D is not the right answer. D says build a Cassandra CE instance. The answer does not refer to taking a snapshot of the on-prem instances, it refers to a snapshot of a Compute Engine VM which you'd have to manually build.
B is correct
upvoted 1 times
-  **Untamables** 11 months, 3 weeks ago
Selected Answer: B
B

upvoted 3 times

 **AzureDP900** 1 year, 3 months ago

B absolutely correct, there is no need of manual installs.

upvoted 1 times

 **pfabio** 1 year, 4 months ago

Selected Answer: B

B is correct: You want to move to Google Cloud quickly and with minimal support effort.

<https://cloud.google.com/marketplace>

upvoted 1 times

 **obeythefist** 1 year, 6 months ago

I suspect every time an answer features Google Marketplace, that is the correct answer. In these kinds of exams the purpose is often just to check general product knowledge.

upvoted 4 times

 **luciorifa** 1 year, 7 months ago

Selected Answer: B

B is the correct answer

upvoted 2 times

 **Raju44** 1 year, 9 months ago

Minimum support when you have steps laydown to your developer to do task on their own. So its D.

upvoted 1 times

 **jaffarali** 1 year, 10 months ago

Selected Answer: B

B is the right option.

upvoted 2 times

 **sid0127** 1 year, 10 months ago

Selected Answer: B

B. its easier to launch from the marketplace

upvoted 3 times

 **jabrrJ68w02ond1** 1 year, 10 months ago

For me it's B, on upfront it offers the lowest effort to launch Marketplace solutions.

upvoted 1 times

 **alaahakim** 1 year, 10 months ago

Selected Answer: B

B is Correct

upvoted 1 times

 **sudav** 2 years, 1 month ago

B - You can deploy Cassandra as a Service, called Astra, on the Google Cloud Marketplace. Not only do you get a unified bill for all GCP services, your Developers can now create Cassandra clusters on Google Cloud in minutes and build applications with Cassandra as a database as a service without the operational overhead of managing Cassandra

upvoted 6 times

 **arsh1916** 2 years, 4 months ago

B is correct

upvoted 1 times

You have a Compute Engine instance hosting a production application. You want to receive an email if the instance consumes more than 90% of its CPU resources for more than 15 minutes. You want to use Google services. What should you do?

- A. 1. Create a consumer Gmail account. 2. Write a script that monitors the CPU usage. 3. When the CPU usage exceeds the threshold, have that script send an email using the Gmail account and smtp.gmail.com on port 25 as SMTP server.
- B. 1. Create a Cloud Monitoring Workspace and associate your Google Cloud Platform (GCP) project with it. 2. Create a Cloud Monitoring Alerting Policy that uses the threshold as a trigger condition. 3. Configure your email address in the notification channel.
- C. 1. Create a Cloud Monitoring Workspace and associate your GCP project with it. 2. Write a script that monitors the CPU usage and sends it as a custom metric to Cloud Monitoring. 3. Create an uptime check for the instance in Cloud Monitoring.
- D. 1. In Cloud Logging, create a logs-based metric to extract the CPU usage by using this regular expression: CPU Usage: ([0-9]{1,3})%. 2. In Cloud Monitoring, create an Alerting Policy based on this metric. 3. Configure your email address in the notification channel.

Correct Answer: D

Community vote distribution

B (100%)

 **Bobbybash** 7 months, 4 weeks ago

Selected Answer: B

B

By setting up a Cloud Monitoring Alerting Policy and configuring the notification channel to an email address, you can receive email alerts when the CPU usage exceeds the threshold that you set. This can be done by using Google Cloud's built-in monitoring and alerting features, which can be more reliable and easier to manage than setting up a custom script or using external email services.

upvoted 2 times

 **LDAP_Anand** 8 months, 3 weeks ago

D. 1. In Cloud Logging, create a logs-based metric to extract the CPU usage by using this regular expression: CPU Usage: ([0-9]{1,3})%. 2. In Cloud Monitoring, create an Alerting Policy based on this metric. 3. Configure your email address in the notification channel.

upvoted 1 times

 **u422628** 9 months ago

Selected Answer: B

Agree, B works

upvoted 1 times

 **fragment137** 10 months, 1 week ago

Selected Answer: B

Agree it's B if we're talking about StackDriver.

upvoted 1 times

 **anolive** 12 months ago

Selected Answer: B

is correct

upvoted 1 times

 **ale_brd_** 1 year ago

Selected Answer: B

answer is B, but I would write it this way as stackdriver is deprecated and Operation Suite uses scopes now.

1. Create a Cloud Monitoring metric scope and associate your Google Cloud Platform (GCP) project with it.

2. Create a Cloud Monitoring Alerting Policy that uses the threshold as a trigger condition.

3. Configure your email address in the notification channel.

upvoted 3 times

 **Akso** 1 year, 1 month ago

Selected Answer: B

The dump has been changed some days ago. This answer was the old and best version:

1. Create a Stackdriver Workspace, and associate your Google Cloud Platform (GCP) project with it.
2. Create an Alerting Policy in Stackdriver that uses the threshold as a trigger condition.
3. Configure your email address in the notification channel.

upvoted 4 times

 **thimai** 1 year ago

And the answer is still the same?

upvoted 1 times

You have an application that uses Cloud Spanner as a backend database. The application has a very predictable traffic pattern. You want to automatically scale up or down the number of Spanner nodes depending on traffic. What should you do?

- A. Create a cron job that runs on a scheduled basis to review Cloud Monitoring metrics, and then resize the Spanner instance accordingly.
- B. Create a Cloud Monitoring alerting policy to send an alert to oncall SRE emails when Cloud Spanner CPU exceeds the threshold. SREs would scale resources up or down accordingly.
- C. Create a Cloud Monitoring alerting policy to send an alert to Google Cloud Support email when Cloud Spanner CPU exceeds your threshold. Google support would scale resources up or down accordingly.
- D. Create a Cloud Monitoring alerting policy to send an alert to webhook when Cloud Spanner CPU is over or under your threshold. Create a Cloud Function that listens to HTTP and resizes Spanner resources accordingly.

Correct Answer: D

Community vote distribution

D (100%)

-  **GCP_Student1** Highly Voted 2 years, 7 months ago

D. Create a Cloud Monitoring alerting policy to send an alert to webhook when Cloud Spanner CPU is over or under your threshold. Create a Cloud Function that listens to HTTP and resizes Spanner resources accordingly.

upvoted 24 times
-  **theBestStudent** Highly Voted 1 year, 1 month ago

Selected Answer: D

Without knowing that much, you can discard easily B,C as they don't make any sense. Automation should be a key in this answer. Also you should discard "A" as with a CronJob you won't spann on time as it will be a fixed time checking. So the only one that is left is D, as just creating an alert and sending it to "something else" (in this case a webhook) in an automated way, should be the common sense way of handling this.

upvoted 6 times
-  **FeaRoX** 8 months, 1 week ago

Isn't "fixed time checking" appropriate for quote : very predictable traffic pattern?

upvoted 1 times
-  **space_cadet** 7 months, 2 weeks ago

Crossed my mind too, but why check every time when you can trigger a response when it happens.
Predictability can also be used to determine the threshold.

upvoted 1 times
-  **scanner2** Most Recent 1 month ago

Selected Answer: D

Correct answer is D.

upvoted 1 times
-  **Captain1212** 1 month ago

Selected Answer: D

D is the correct Answer as B or C does not do it automatically, and a doesnot use for long spanning

upvoted 1 times
-  **temppp** 9 months, 2 weeks ago

Why not A is correct as question suggested specific time where as D is like an unpredictable time?

upvoted 3 times
-  **jrlsl1991** 8 months, 1 week ago

Because even though the traffic has a clear pattern, if the traffic changes one day (like a special holiday for ecommerce websites), you wouldn't be able to serve accordingly. It's never a good practice use fixed jobs for time-based traffic issues.

upvoted 3 times
-  **roaming_panda** 9 months, 3 weeks ago

D is definitely correct .. people

upvoted 1 times
-  **abirroy** 1 year, 2 months ago

Selected Answer: D

Correct answer is D

upvoted 1 times
-  **AzureDP900** 1 year, 3 months ago

D is correct. It is part of Tutorialspoint practice test

upvoted 1 times

✉ **rsuresh27** 1 year, 5 months ago

Answer is D. The keyword to look for is "automatically". A, B, C all have steps that are not automatic. Thus, only D is left.

upvoted 3 times

✉ **alaahakim** 1 year, 10 months ago

D is the answer

upvoted 3 times

✉ **Ricky_92** 1 year, 11 months ago

<https://cloud.google.com/architecture/autoscaling-cloud-spanner>

upvoted 5 times

✉ **kukabura** 1 year, 12 months ago

D is the answer

upvoted 3 times

✉ **ankatsu2010** 1 year, 12 months ago

It's a tricky question. The answer is 'A'.

Question says traffic pattern is predictable. This means you can schedule the scalability.

You can achieve this easily by using Cloud Spanner's API.

<https://cloud.google.com/spanner/docs/reference/rest/v1/projects.instances/patch>

upvoted 3 times

✉ **BobbyFlash** 1 year, 11 months ago

Really tricky indeed...But if you think about it and you don't stick religiously to "traffic pattern is predictable" sentence (that sentence is to trick you), you would think that despite the pattern is very predictable, it could vary at some point in time. With a cron job you can schedule a job using a fixed starting point and a a fixed ending, can't you? Well, what would happen if the app traffic suddenly spikes before your cron job starts running? Or after?. At this point, I go for D, making scalability automatic thanks to Cloud Functions.

upvoted 6 times

✉ **erikamrqz** 2 years ago

D 205%

upvoted 5 times

✉ **bubblegumbeach** 1 year, 12 months ago

seriously? 205%?

upvoted 4 times

✉ **jcols** 2 years, 4 months ago

Option D. There's an official repository that does something similar to provide autoscaling to Cloud Spanner.

<https://github.com/cloudspannerecosystem/autoscaler>

upvoted 2 times

✉ **zaxxon** 1 year, 11 months ago

Which uses a Cloud Scheduler so answer A

upvoted 1 times

✉ **ryumada** 1 year, 2 months ago

"...and then resize the Spanner instance accordingly." I think the last sentence says that the resize action is done manually. If so, then D is still the right answer.

upvoted 1 times

✉ **ryumada** 1 year, 2 months ago

For the D option, "Create a Cloud Function that listens to HTTP and resizes Spanner resources accordingly." The listen and resize actions is done by Cloud Function.

upvoted 1 times

✉ **arsh1916** 2 years, 4 months ago

D no doubt

upvoted 2 times

✉ **mcaromit** 2 years, 4 months ago

D is perfect

upvoted 2 times

Your company publishes large files on an Apache web server that runs on a Compute Engine instance. The Apache web server is not the only application running in the project. You want to receive an email when the egress network costs for the server exceed 100 dollars for the current month as measured by Google Cloud.

What should you do?

- A. Set up a budget alert on the project with an amount of 100 dollars, a threshold of 100%, and notification type of email.
- B. Set up a budget alert on the billing account with an amount of 100 dollars, a threshold of 100%, and notification type of email.
- C. Export the billing data to BigQuery. Create a Cloud Function that uses BigQuery to sum the egress network costs of the exported billing data for the Apache web server for the current month and sends an email if it is over 100 dollars. Schedule the Cloud Function using Cloud Scheduler to run hourly.
- D. Use the Cloud Logging Agent to export the Apache web server logs to Cloud Logging. Create a Cloud Function that uses BigQuery to parse the HTTP response log data in Cloud Logging for the current month and sends an email if the size of all HTTP responses, multiplied by current Google Cloud egress prices, totals over 100 dollars. Schedule the Cloud Function using Cloud Scheduler to run hourly.

Correct Answer: A

Community vote distribution



- ✉️ **GCP_Student1** Highly Voted 2 years, 7 months ago
- C. Export the billing data to BigQuery. Create a Cloud Function that uses BigQuery to sum the egress network costs of the exported billing data for the Apache web server for the current month and sends an email if it is over 100 dollars. Schedule the Cloud Function using Cloud Scheduler to run hourly.
upvoted 23 times
- ✉️ **MyName_** Highly Voted 2 years, 6 months ago
- [C]
I think the keyword here is "as measured by Google Cloud". In Answer D you calculate the price yourself, in C you use the billing provided by GCP. Thus I think the Answer is C.
upvoted 12 times
- ✉️ **scanner2** Most Recent 1 month ago
- Selected Answer: C
Correct answer is C.
upvoted 1 times
- ✉️ **Captain1212** 1 month ago
- Selected Answer: C
C is the correct answer , as it gives you everything the question wants
upvoted 1 times
- ✉️ **Shenannigan** 5 months ago
- Selected Answer: A
Answer is A

From this Link:
<https://cloud.google.com/load-balancing/docs/ssl>

it states this:
External SSL proxy load balancers are intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use an external HTTP(S) load balancer.

443 is HTTPS traffic

for those saying 443 isn't https
<https://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml?search=443>
upvoted 1 times
- ✉️ **Shenannigan** 5 months ago
- Disregard this was supposed to be for question 127
upvoted 1 times
- ✉️ **HiddenClouds** 9 months, 3 weeks ago
- Selected Answer: C
Took the exam this week this question is on there, C is the correct answer you need to remember to filter.
upvoted 7 times
- ✉️ **Cherumathi** 1 year ago

Selected Answer: C

C is the correct answer,
You export the bill to BigQuery and filter for the Egress cost for the particular application, and send an email if the cost is over 100 dollars, to send an email you need to use cloud function to monitor and trigger based on the conditions.

upvoted 1 times

✉ **sylva91** 1 year ago

Selected Answer: C

it can only be the C because "it's not the only app running"
upvoted 1 times

✉ **VietmanOfficiel** 1 year, 1 month ago

Selected Answer: D

[D]
Complex but works for resume network egrs cost of ONLY the Apache instance
upvoted 1 times

✉ **theBestStudent** 1 year, 1 month ago

Selected Answer: C

I would say this is the approach without knowing that much:
A & B discarded as they are not resourced oriented. We need to charge for the apache server (so focusing on the VM where it is hosted) in order to charge this server.

D can't be as you are not charged in this case for that Response payload received.
upvoted 1 times

✉ **ryumada** 1 year, 2 months ago

Selected Answer: C

Answer C and D are correct to me. You can calculate the current Google Cloud egress prices using Cloud billing catalog API which can be used in D option.

ref: <https://cloud.google.com/blog/topics/cost-management/introducing-cloud-billing-catalog-api-gcp-pricing-in-real-time>

But, if you choose option C, then you will get the usage cost directly from cloud billing data. You should add a label to the Apache web server in order to select its cost.

D has the more complex step, using more services too which is Cloud Logging to store the logging data of the VM and Cloud Billing Catalog API. The using of more services makes the D option to have more expensive costs.
upvoted 3 times

✉ **andreherwanto** 1 year, 2 months ago

Selected Answer: C

C is best options
upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

C is best options
upvoted 1 times

✉ **PAUGURU** 1 year, 5 months ago

Selected Answer: B

C cannot be the right answer, you export the data to big Query only once. What is the use of a cloud function running every hour on the same data? It doesn't say "you export the data to BigQuery hourly".
upvoted 2 times

✉ **PAUGURU** 1 year, 5 months ago

I was wrong, when you configure an export to Bigquery billing data are updated periodically, so answer is actually C.
<https://cloud.google.com/billing/docs/how-to/export-data-bigquery-tables?hl=en>
upvoted 6 times

✉ **Leo_li_leo** 1 year, 6 months ago

I feel the answer should be B as projects can be specified and also services.
upvoted 2 times

✉ **Pitchino80** 1 year, 8 months ago

Selected Answer: C

It's C
upvoted 2 times

✉ **Ridhanya** 1 year, 10 months ago

C. Even I thought its D initially, but D is a lot of steps while C has minimum steps
upvoted 2 times

✉ **obeythefist** 1 year, 7 months ago

C and D are not correct, because running a project for a week when it has no load is not an accurate way to estimate costs.
upvoted 1 times

You have designed a solution on Google Cloud that uses multiple Google Cloud products. Your company has asked you to estimate the costs of the solution. You need to provide estimates for the monthly total cost. What should you do?

- A. For each Google Cloud product in the solution, review the pricing details on the products pricing page. Use the pricing calculator to total the monthly costs for each Google Cloud product.
- B. For each Google Cloud product in the solution, review the pricing details on the products pricing page. Create a Google Sheet that summarizes the expected monthly costs for each product.
- C. Provision the solution on Google Cloud. Leave the solution provisioned for 1 week. Navigate to the Billing Report page in the Cloud Console. Multiply the 1 week cost to determine the monthly costs.
- D. Provision the solution on Google Cloud. Leave the solution provisioned for 1 week. Use Cloud Monitoring to determine the provisioned and used resource amounts. Multiply the 1 week cost to determine the monthly costs.

Correct Answer: C*Community vote distribution*

A (100%)

✉️  **GCP_user**  2 years, 6 months ago

Yes I agree with A. Makes more sense.

upvoted 28 times

✉️  **BobbyFlash** 1 year, 11 months ago

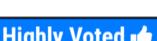
Why not B?? Even though answer A makes sense, they are also stating to provide estimates for "monthly total costs". One would think that it is not only necessary to get estimates from every resource, but also consolidate them to inform the monthly total cost required.

upvoted 2 times

✉️  **BobbyFlash** 1 year, 11 months ago

Ignore the comment folks. Sentence B is missing the little thing about the pricing calculator. I go with A.

upvoted 6 times

✉️  **GCP_Student1**  2 years, 7 months ago

A. For each Google Cloud product in the solution, review the pricing details on the products pricing page. Use the pricing calculator to total the monthly costs for each Google Cloud product.

upvoted 6 times

✉️  **scanner2**  1 month ago

Selected Answer: A

Correct answer is A. Use GCP pricing calculator.

upvoted 1 times

✉️  **Captain1212** 1 month ago

Selected Answer: A

A seems more correct

upvoted 1 times

✉️  **Untamables** 11 months, 3 weeks ago

Selected Answer: A

Vote A

<https://cloud.google.com/free/docs/estimate-costs-google-cloud-platform>

upvoted 2 times

✉️  **Charumathi** 1 year ago

Selected Answer: A

A is the correct answer, use the pricing calculator to estimate the pricing for a month and download the estimate to csv file, or you can share the URL of the pricing calculator or email the estimate to the respective people in the company.

upvoted 1 times

✉️  **VietmanOfficiel** 1 year, 1 month ago

Selected Answer: A

[A]

When "estimate" you need to read "price calculator"

upvoted 1 times

✉️  **AzureDP900** 1 year, 3 months ago

I agree with A

upvoted 1 times

 **dishum** 1 year, 8 months ago

Answer is C

The question says 'cost of solution of the design' means, how it is going to work in practical, means the traffic load, no of users, uploads, downloads, transactions etc etc.

In this case, the more nearer option is to run for a week, then calculate.

Answer A is not correct, becoz A is applicable when the cost of the product is to be determined.

upvoted 4 times

 **SunnyDey** 11 months ago

Question has asked to estimate the costs of the solution i.e. the initial setup cost not the running cost. Therefore, C and D discarded and B is of no use without pricing calculator. Question is testing whether you are aware of thes calculator service from Google.

upvoted 2 times

 **sid0127** 1 year, 10 months ago

Selected Answer: A

answer is a

upvoted 2 times

 **alaahakim** 1 year, 10 months ago

Selected Answer: A

The Right Ans: A

upvoted 3 times

 **TenshiD** 1 year, 10 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **[Removed]** 2 years ago

As per GCP best practices, A makes more sense

upvoted 2 times

 **jcloud965** 2 years, 2 months ago

Correct is A. It is the main purpose of Pricing calculator. You have to check pricing details on products pages to understand what charges apply and how to size it.

B. It works bur it is more complicated than A.

C. It works but you need to wait one week and you pay for this. Moreoever, you won't be in production so somes charges may lack like net egress charges

D. No, because Cloud Monitoring can't be used for billing purpose

upvoted 5 times

 **tanito83** 2 years, 4 months ago

The correct answer is A. Please, modify it.

upvoted 2 times

 **arsh1916** 2 years, 4 months ago

A, is best option

upvoted 1 times

 **mcaromit** 2 years, 4 months ago

A is correct

upvoted 1 times

You have an application that receives SSL-encrypted TCP traffic on port 443. Clients for this application are located all over the world. You want to minimize latency for the clients. Which load balancing option should you use?

- A. HTTPS Load Balancer
- B. Network Load Balancer
- C. SSL Proxy Load Balancer
- D. Internal TCP/UDP Load Balancer. Add a firewall rule allowing ingress traffic from 0.0.0.0/0 on the target instances.

Correct Answer: C

Reference:

<https://cloud.google.com/load-balancing/docs/ssl>

Community vote distribution

C (75%)

A (25%)

✉  Bituz  3 years, 2 months ago

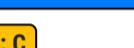
SSL Proxy Load Balancing support for the following ports: 25, 43, 110, 143, 195, 443, 465, 587, 700, 993, 995, 1883, 3389, 5222, 5432, 5671, 5672, 5900, 5901, 6379, 8085, 8099, 9092, 9200, and 9300. When you use Google-managed SSL certificates with SSL Proxy Load Balancing, the frontend port for traffic must be 443 to enable the Google-managed SSL certificates to be provisioned and renewed.

upvoted 27 times

✉  DarioFama23  3 years, 3 months ago

C is correct

upvoted 19 times

✉  scanner2  1 month ago

Selected Answer: C

Correct answer is C.

External proxy load balancer supports global and regional scope. While external passthrough network load balancer supports regional scope.

upvoted 1 times

✉  Captain1212 1 month ago

Selected Answer: C

C is the right answer, read it carefully TCP traffic

upvoted 1 times

✉  juliorevk 2 months, 1 week ago

Selected Answer: A

443 is HTTPS

upvoted 1 times

✉  Shenannigan 5 months ago

Selected Answer: A

Answer is A

From this Link:

<https://cloud.google.com/load-balancing/docs/ssl>

it states this:

External SSL proxy load balancers are intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use an external HTTP(S) load balancer.

443 is HTTPS traffic

for those saying 443 isn't https

<https://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml?search=443>

upvoted 1 times

✉  Jelly_Wang 5 months, 1 week ago

Selected Answer: C

Go for C. We are taking exam so read the question smart. For HTTPS traffic use HTTPS load balancer. For non-HTTP traffic, use SSL Proxy Load Balancer. (<https://cloud.google.com/load-balancing/docs/ssl>). Network Load Balancer (External TCP/UDP network load balancer) for regional pass through traffic (<https://cloud.google.com/load-balancing/docs/network>). Here it mentioned SSL-encrypted TCP traffic all over the world, go for SSL proxy load balancer. Read the keyword "SSL-encrypted TCP" and "all over the world". Keep it simple, don't over complicate yourself

upvoted 6 times

✉  Praxii 5 months, 1 week ago

Selected Answer: C

The answer is C.

"SSL encrypted TCP" traffic is not exactly SSL traffic. SSL encrypted TCP traffic is usually used for non HTTP application. Port 443 is supported by SSL proxy load balancer so this isn't what will differentiate the two for us. The main part here is "SSL encrypted TCP" traffic.

upvoted 1 times

✉ **Bobbybash** 7 months, 4 weeks ago

Selected Answer: C

C. SSL Proxy Load Balancer would be the best option for minimizing latency for the clients, as it terminates SSL traffic and forwards unencrypted traffic directly to the backend instances. This reduces the amount of processing and latency associated with SSL encryption/decryption. Additionally, because the clients are located all over the world, using a global SSL Proxy Load Balancer can distribute traffic to the closest backend instances for the clients, further reducing latency.

upvoted 3 times

✉ **xaqanik** 8 months ago

Selected Answer: A

A. HTTPS Load Balancer

The HTTPS Load Balancer is the best option for minimizing latency for clients located all over the world. HTTPS Load Balancer provides a global solution for load balancing secure (SSL/TLS) traffic, including the ability to distribute traffic to backend instances based on IP address, based on request content, or both. It is designed to handle encrypted traffic and can terminate SSL/TLS connections, making it the optimal choice for an application that receives SSL-encrypted TCP traffic on port 443. Additionally, HTTPS Load Balancer has built-in features to minimize latency, such as support for HTTP/2 and connection multiplexing, which can reduce the number of connections and round trips required to complete a request.

upvoted 3 times

✉ **eBookz** 8 months ago

Selected Answer: A

A:

Although both HTTPS and SSL Proxy load balancers can both be used for global external load balancing, for HTTP(S) traffic, Google recommends that you use HTTP(S) Load Balancing.

<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 2 times

✉ **eBookz** 8 months ago

To add further clarity, Google documentation says:

"External SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing."

Port 443 (mentioned in the question) is for HTTPS

<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 2 times

✉ **Kamngur** 7 months ago

You are assuming that this is HTTPS traffic. But from question we know that it is SSL-encrypted TCP , so we can't use HTTPS load balancer. Moreover we will be missing HTTP data for URL map

upvoted 1 times

✉ **alex000** 9 months, 1 week ago

Selected Answer: A

Correct answer: (A)

Global https load balancer

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer>

upvoted 2 times

✉ **mrvergara** 9 months ago

The question does not mention the layer 7 protocol (HTTP/S), only that is TCP (layer 4). This is way it should be C

upvoted 2 times

✉ **FeaRoX** 8 months, 1 week ago

HTTPs uses both SSL encryption and TCP layer as well. 443 port is HTTPs default port, which suggests A.

upvoted 1 times

✉ **Kamngur** 7 months ago

443 is default port for SSL/TLS communication. I can be HTTPs or it can be somethink else.

upvoted 1 times

✉ **Viggy1212** 10 months, 3 weeks ago

This is a tricky question.

First point to consider is the port. TCP 443 port is used for HTTPS traffic.

Second : SSL Proxy LB is intended for non-HTTPs traffic and for HTTPs traffic, it should be global HTTPs LB.

Answer is A : HTTPS Load Balancer

upvoted 3 times

✉ **Charumathi** 1 year ago

Selected Answer: C

C is the correct Answer,
SSL proxy load balancer with Traffic Type - TCP with SSL offload, for global IPv4, IPv6, external ports for load balancing -
25,43,110,143,195,443,465,587,700,993,995,1883,5222

upvoted 2 times

✉ **abirroy** 1 year, 2 months ago

Selected Answer: C

SSL Proxy Load Balancer
upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

SSL proxy load balancer - C is right
upvoted 1 times

✉ **JelloMan** 1 year, 4 months ago

Selected Answer: C

Network LB = Regional
SSL Proxy = Global <---That alone tells you the answer is C
upvoted 5 times

✉ **fragment137** 10 months, 1 week ago

HTTPS Load Balancer can be global too: <https://cloud.google.com/static/load-balancing/images/choose-lb.svg>

That said, the question calls for TCP traffic with SSL, so based on the above image, it's still C: SSL Proxy

upvoted 1 times

You have an application on a general-purpose Compute Engine instance that is experiencing excessive disk read throttling on its Zonal SSD Persistent Disk. The application primarily reads large files from disk. The disk size is currently 350 GB. You want to provide the maximum amount of throughput while minimizing costs.

What should you do?

- A. Increase the size of the disk to 1 TB.
- B. Increase the allocated CPU to the instance.
- C. Migrate to use a Local SSD on the instance.
- D. Migrate to use a Regional SSD on the instance.

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/disks/performance>

Community vote distribution

C (83%)

Other

✉️  **DarioFama23**  3 years, 3 months ago

C is correct, local SSD has more IOPS

upvoted 34 times

✉️  **RegisFTM** 1 year, 9 months ago

trick question... local-ssd is not persistent. increasing the size of the disk will also increase the iops. A is correct imho.

upvoted 5 times

✉️  **shax910** 1 year, 8 months ago

Local SSDs have higher throughput and lower latency than standard persistent disks or SSD persistent disks. The data that you store on a local SSD persists only until the instance is stopped or deleted.

upvoted 1 times

✉️  **obeythefist** 1 year, 7 months ago

I once thought that A was the correct response because of the persistence problem, but reading the question carefully, we must choose C over A.

The question does not stipulate that the local files must be persistent, and this is the only reason why you would choose "A" over "C".

Also, the question has an important key word: Minimising costs.

1TB of zonal persistent disk costs a huge amount more than 350GB of local disk.

We should choose C.

upvoted 6 times

✉️  **ace_2021_1** 1 year ago

they have been using persistent disk....there was a reason during the initial design.....persistent disk was a req....we should find a solution without affecting that...increasing the size would be the best solution for this scenario.

upvoted 3 times

✉️  **pfabio** 1 year, 4 months ago

Very Nice, agree C is correct

upvoted 1 times

✉️  **TAvenger** 2 years, 7 months ago

Agree. This is also cheaper than having 350 Gb persistent SSD:

Here are calculations (taken from GCP when creating instance)

350 Gb SSD Persistent disk: 59.50\$/month, read IOPS: 10 500 with n1-standard-1

1000 Gb SSD Persistent disk: 170.00\$/month, read IOPS: 15 000 with n1-standard-1

375 Gb Local SSD (NVMe): 30.00\$/month, read IOPS: 170 000 with n1-standard-1

upvoted 12 times

✉️  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (C):

Performance

Standard persistent disks are efficient and economical for handling sequential read/write operations, but they aren't optimized to handle high rates of random input/output operations per second (IOPS). If your apps require high rates of random IOPS, use SSD persistent disks. SSD persistent disks are designed for single-digit millisecond latencies. Observed latency is application specific.

upvoted 16 times

✉ **hogtrough** 1 year, 8 months ago

A local SSD is not the same as an SSD persistent disk.

"Local SSDs are physically attached to the server that hosts your VM instance. Local SSDs have higher throughput and lower latency than standard persistent disks or SSD persistent disks. The data that you store on a local SSD persists only until the instance is stopped or deleted."

<https://cloud.google.com/compute/docs/disks#localssds>

The answer is C.

upvoted 3 times

✉ **hogtrough** 1 year, 8 months ago

Sorry, A not C.

upvoted 1 times

✉ **scanner2** Most Recent 1 month ago

Selected Answer: C

Local SSD provides more throughput than Persistent disks and is cost effective solution. So, correct answer is C.

upvoted 1 times

✉ **Bobbybash** 7 months, 4 weeks ago

Selected Answer: C

C. Migrate to use a Local SSD on the instance. Local SSDs provide higher throughput and lower latency compared to Zonal SSD Persistent Disks, and are optimized for use cases that require high-speed, temporary storage. They are physically attached to the instance, so network latencies are minimized. However, they are not intended for long-term storage and may not provide the same level of durability as persistent disks. Since the application is primarily reading large files from disk and experiencing disk read throttling, using a Local SSD should provide a significant improvement in performance while minimizing costs. Increasing the size of the Zonal SSD Persistent Disk or increasing the allocated CPU to the instance may provide some improvement, but are unlikely to fully address the disk read throttling issue. Migrating to a Regional SSD is also not likely to improve performance significantly, as the disk is still separate from the instance and network latencies can impact performance.

upvoted 3 times

✉ **researched_answer_boi** 8 months, 3 weeks ago

According to the page containing the tables and to the tables "https://cloud.google.com/compute/docs/disks/performance#n1_vms" and "https://cloud.google.com/compute/docs/disks/performance#n2_vms", the number of CPU cores greatly influence the available max. read THROUGHTPUT ("excessive disk read throttling on its Zonal SSD Persistent Disk", "The application primarily reads large files from disk.") on general purpose VMs.

The question also requires minimizing the costs, however, as Local SSDs are EPHEMERAL, they are out of question for the scenario at hand.

So, answer "B" seems to be the correct one.

upvoted 1 times

✉ **AwesomeGCP** 12 months ago

Selected Answer: C

C. Migrate to use a Local SSD on the instance

upvoted 1 times

✉ **Charumathi** 1 year ago

Selected Answer: C

C is the correct Answer,

Local SSDs

Local SSDs are physically attached to the server that hosts your VM instance. Local SSDs have higher throughput and lower latency than standard persistent disks or SSD persistent disks. The data that you store on a local SSD persists only until the instance is stopped or deleted. Each local SSD is 375 GB in size, but you can attach a maximum of 24 local SSD partitions for a total of 9 TB per instance.

Performance

Local SSDs are designed to offer very high IOPS and low latency. Unlike persistent disks, you must manage the striping on local SSDs yourself. Combine multiple local SSD partitions into a single logical volume to achieve the best local SSD performance per instance, or format local SSD partitions individually.

Local SSD performance depends on which interface you select. Local SSDs are available through both SCSI and NVMe interfaces.

upvoted 1 times

✉ **Cornholio_LMC** 1 year ago

had this question today

upvoted 1 times

✉ **gcp_world123** 1 year, 1 month ago

C is correct

Local SSDs are physically attached to the server that hosts your VM instance. Local SSDs have higher throughput and lower latency than standard persistent disks or SSD persistent disks. The performance gains from local SSDs require certain trade-offs in availability, durability, and flexibility. Because of these trade-offs, Local SSD storage isn't automatically replicated and all data on the local SSD might be lost if the instance terminates for any reason.

Ref: <https://cloud.google.com/compute/docs/disks/localssds>

Ref: https://cloud.google.com/compute/docs/disks/performance#type_comparison

upvoted 1 times

✉ **theBestStudent** 1 year, 1 month ago

Selected Answer: B

C can't B for the reasons explained here in the answers (local disk info will be totally deleted if you restart or delete your vm). Then A or B. Google recommends to increase the size of the disk as performance of the disk is linear to the size of the same. But Im wondering if 1TB is too much (it is almost 2 times more the original size of the disk, and that sounds like too much, plus if with 350GB the files already fit there, then increasing the size of the disk "a little bit more" should have been a better approach to test how the performance increases. Then Google Also recommends to add more CPU to get a better IOPS, it might be cheaper than option A if the right machine with more CPU is chosen.

Answer: B

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

Local SSD .. C is right

upvoted 1 times

✉ **Tirthankar17** 1 year, 4 months ago

Selected Answer: C

C is correct.

upvoted 1 times

✉ **Rutu_98** 1 year, 4 months ago

Selected Answer: C

Go with C

As Local SSDs have high IOPS

upvoted 1 times

✉ **Akash7** 1 year, 6 months ago

C for me.

<https://cloud.google.com/compute/docs/disks#performance>

upvoted 1 times

✉ **Soso_zozo** 1 year, 6 months ago

Correct Answer is B

From both tables:

https://cloud.google.com/compute/docs/disks/performance#performance_by_disk_size

<https://cloud.google.com/compute/docs/disks/performance#machine-type-disk-limits>

The general purpose vm limit is less than ssd persistent disk size limit of 250-500

upvoted 3 times

✉ **s_a_t_y_a_m** 1 year, 7 months ago

Selected Answer: C

Increasing Disk size would increase cost while question is asking for minimising the cost

upvoted 2 times

✉ **Mixxer5** 1 year, 8 months ago

Selected Answer: A

We have to assume that persistent type of disk was used on purpose so we can't switch to non-persistent solution. Larger disk will- in itself- increase IOPS as well so A is an answer.

upvoted 1 times

✉ **theBestStudent** 1 year, 1 month ago

But "A" is expensive.

upvoted 1 times

Your Dataproc cluster runs in a single Virtual Private Cloud (VPC) network in a single subnet with range 172.16.20.128/25. There are no private IP addresses available in the VPC network. You want to add new VMs to communicate with your cluster using the minimum number of steps. What should you do?

- A. Modify the existing subnet range to 172.16.20.0/24.
- B. Create a new Secondary IP Range in the VPC and configure the VMs to use that range.
- C. Create a new VPC network for the VMs. Enable VPC Peering between the VMs' VPC network and the Dataproc cluster VPC network.
- D. Create a new VPC network for the VMs with a subnet of 172.32.0.0/16. Enable VPC network Peering between the Dataproc VPC network and the VMs VPC network. Configure a custom Route exchange.

Correct Answer: A*Community vote distribution*

✉ **ESP_SAP** Highly Voted 3 years, 1 month ago

Correction.

Correct Answers is (A):

gcloud compute networks subnets expand-ip-range

NAME

gcloud compute networks subnets expand-ip-range - expand the IP range of a Compute Engine subnetwork
upvoted 46 times

✉ **ccieman2016** 1 year, 1 month ago

I think, you can't expand ip range subnet, if there isn't space in VPC. I read this question a lot, VPC CIDR like with 172.16.20.128/25 and there's only one subnet 172.16.20.128/25 inside this VPC, so you can't expand nothing. for me, there's Letter C and D works, but letter D is necessary extra work. LETTER C is right.

upvoted 3 times

✉ **BigMac666** 3 months, 2 weeks ago

VPC's DO NOT have IP range limitations.

You can only object if 172.16.20.0/25 is in use in the same VPC or in a VPC that this VPC is already peered with.

.128/25 expands to .0/24 (i.e. "backwards") So as long as it's free, you're good.

In a question like this, it's obvious that the simple answer is the right one, i.e. A - Expansion.

upvoted 1 times

✉ **gastonrepetto77** 2 weeks, 1 day ago

- The statement is clear with point 1A and 2A:

1A.- "single Virtual Private Cloud (VPC) network in a single subnet"

2A.- "There are no private IP addresses available in the VPC network."

- Question: How can you expand if there is a single VPC with a single subnet and there are no private IP addresses available in the only VPC network ???

- Result: Yes it has limitation, this question is clear, this is an exam, not the real life, we cannot verify anything else and we have the limitation of the statement.

upvoted 1 times

✉ **FeaRoX** 8 months, 1 week ago

There's no information about VPC CIDR, only subnet. You can't tell that there's no space

upvoted 2 times

✉ **francisco_guerra** Highly Voted 3 years, 2 months ago

I think is A,

upvoted 35 times

✉ **Priyanka109** 1 year ago

No it can't be as you can't modify ip address but can expand. There is no ip in the existing vpc so you have to create a new vpc and connect it using peering.

upvoted 3 times

✉ **passmepls** 3 years, 1 month ago

thank you

upvoted 4 times

✉ **SinghAnc** Most Recent 6 days, 18 hours ago

Selected Answer: C

C is the correct answer

Because By creating a new VPC network for the VMs and enabling VPC Peering, you can ensure communication between the VMs and the Dataproc cluster while avoiding the need to modify existing subnets, which could potentially disrupt existing resources. This approach also helps keep the environments separate, which can be beneficial for security and manageability.

upvoted 1 times

✉ **pessoalsp** 1 week, 1 day ago

A is incorrect because modifying the existing subnet range would require moving all of the existing VMs in the subnet to a new subnet.

upvoted 1 times

✉ **gastonrepetto77** 2 weeks, 1 day ago

- The statement is clear with point 1A and 2A:

1A.- "single Virtual Private Cloud (VPC) network in a single subnet"

2A.- "There are no private IP addresses available in the VPC network."

- Question: How can you expand if there is a single VPC with a single subnet and there are no private IP addresses available in the only VPC network ???

upvoted 1 times

✉ **scanner2** 1 month ago

Selected Answer: A

<https://cloud.google.com/vpc/docs/create-modify-vpc-networks#expand-subnet>

upvoted 1 times

✉ **Captain1212** 1 month ago

Selected Answer: C

C is the correct answer, because in the question its mention that there is no available ip address in the VPC network

upvoted 1 times

✉ **ROSHANgcp** 3 months, 2 weeks ago

Selected Answer: A

This involves modifying the subnet range of the existing VPC network to increase the number of available IP addresses. By changing the subnet range to 172.16.20.0/24, you will have a larger IP address range to allocate to new VMs, allowing them to communicate with the Dataproc cluster.

upvoted 2 times

✉ **kumarts** 3 months, 3 weeks ago

Dataproc Metastore uses private IP only, so no public IP is exposed. This means that only VMs on the provided Virtual Private Cloud (VPC) network or on-premises (connected through Cloud VPN or Cloud Interconnect) can access the Dataproc Metastore service.

Dataproc Metastore leverages VPC network peering to provide IP address connectivity to the Dataproc Metastore service's endpointUri. (Refer <https://cloud.google.com/dataproc-metastore/docs/access-service>)

upvoted 1 times

✉ **kumarts** 3 months, 3 weeks ago

A virtual machine can communicate with a dataproc through SSH only right, please clarify.

upvoted 1 times

✉ **sabrinakloud** 5 months, 2 weeks ago

Selected Answer: C

"no private IP addresses available in the VPC network"

upvoted 1 times

✉ **Y__ash** 5 months, 3 weeks ago

Selected Answer: C

You couldn't extend the existing subnet range because there are no private IP addresses available in the VPC network. You could extend subnet range only using IP addresses from the VPC

upvoted 1 times

✉ **chikorita** 5 months, 3 weeks ago

A cuz subnet expansion is possible and minimum number of steps is required

upvoted 1 times

✉ **JC0926** 7 months, 2 weeks ago

Selected Answer: C

There are no private IP addresses available in the VPC network.

upvoted 2 times

✉ **JC0926** 7 months, 3 weeks ago

Selected Answer: C

There are no private IP addresses available in the VPC network.

upvoted 1 times

✉ **Andoameda9** 8 months ago

Network peering is done with different orgs. it would have been tricky if it was shared vpc. I will go for 'a' for 'minimum steps'.

upvoted 3 times

 **jrls1991** 8 months, 1 week ago

Selected Answer: A

I'm going for A. We can expand the CIDR range (there are more comments doing that). There's only one subnet, so there are no possibilities of having other subnets with overlapping ranges. This would also meet the "minimum number of steps", since option C (the other potential solution) would require many steps more.

upvoted 1 times

You manage an App Engine Service that aggregates and visualizes data from BigQuery. The application is deployed with the default App Engine Service account.

The data that needs to be visualized resides in a different project managed by another team. You do not have access to this project, but you want your application to be able to read data from the BigQuery dataset. What should you do?

- A. Ask the other team to grant your default App Engine Service account the role of BigQuery Job User.
- B. Ask the other team to grant your default App Engine Service account the role of BigQuery Data Viewer.
- C. In Cloud IAM of your project, ensure that the default App Engine service account has the role of BigQuery Data Viewer.
- D. In Cloud IAM of your project, grant a newly created service account from the other team the role of BigQuery Job User in your project.

Correct Answer: B

The Owner, Editor, and Viewer primitive roles include the BigQuery Admin (roles/bigquery.dataOwner), BigQuery Data Editor (roles/bigquery.dataEditor), and

BigQuery Data Viewer (roles/bigquery.dataViewer) roles, respectively. This means the Owner, Editor, and Viewer primitive roles have BigQuery access as defined for the respective BigQuery roles.

Reference:

<https://cloud.google.com/bigquery/docs/access-control>

Community vote distribution

A (57%) B (43%)

✉  **Hjameel**  3 years, 1 month ago

I think B is the answer
upvoted 30 times

✉  **KC_go_reply** 4 months ago

'I think' could you elaborate further please?
upvoted 1 times

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (B):
Sorry, I copied/pasted the wrong statement.
This is the proper explanation regarding to Big Query Data Viewer Role.
The resource that you need to get access is in the other project.

roles/bigquery.dataViewer BigQuery Data Viewer
When applied to a table or view, this role provides permissions to:

Read data and metadata from the table or view.
This role cannot be applied to individual models or routines.

When applied to a dataset, this role provides permissions to:

Read the dataset's metadata and list tables in the dataset.
Read data and metadata from the dataset's tables.
When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.
upvoted 24 times

✉  **[Removed]** 3 years ago

A is correct, data viewer role does not allow you to execute query, that can be done if you are the user role.
upvoted 9 times

✉  **[Removed]** 3 years ago

The question states you want to aggregate and visualize data, that is run aggregate SQL on data before visualizing.
upvoted 2 times

✉  **Captain1212**  1 month ago

Selected Answer: A

A seems more correct , as A give you both required access i.e aggragets and visualizes data from BQ
upvoted 1 times

✉  **Ahmed_Y** 1 month, 2 weeks ago

Selected Answer: B

"BigQuery Data Viewer: When applied to a dataset, this role provides permissions to:
Read the dataset's metadata and list tables in the dataset.
Read data and metadata from the dataset's tables."

"BigQuery Job User: Provides permissions to run jobs, including queries, within the project."

So it is not giving access to read BigQuery Dataset!

<https://cloud.google.com/bigquery/docs/access-control>

upvoted 1 times

 **tczorro** 2 months ago

Selected Answer: B

The answer is B

BigQuery Job User only grant you query execution permission, without Data Viewer permission, you can't really query a dataset.

BigQuery Data Viewer permission allows one to access data inside a table or dataset. You can't query those table, but you can access data from REST api.

upvoted 1 times

 **bond007_userguy** 2 months, 1 week ago

checked by assigning the permission it is JobUser - A

upvoted 1 times

 **Naree** 3 months ago

Selected Answer: A

Refer the following links for the difference:

<https://cloud.google.com/bigquery/docs/access-control#bigquery.jobUser>

<https://cloud.google.com/bigquery/docs/access-control#bigquery.dataViewer>

upvoted 1 times

 **Dmosh** 3 months, 2 weeks ago

Selected Answer: A

I was considering B, but the app has to run queries in order to extract the information, therefore, A.

upvoted 1 times

 **krop** 3 months, 3 weeks ago

Selected Answer: A

Correct answer is A. IN order to get the data, you must be able to execute a query. So "BigQuery Job User" role is required. That role allows you to submit a job, which means execute a query. I have tested this against both roles in question. Only "BigQuery Job User" allows you to get the data as in question.

upvoted 1 times

 **KC_go_reply** 4 months ago

Selected Answer: A

The people voting for B probably haven't really used BigQuery so far.

To use BigQuery 'normally' as a user, you need both the Data Viewer AND the Job User role. In practice, with Data Viewer, you can only list the datasets, tables, views etc. in BigQuery, and view their schemas and metadata.

Now, here in this scenario, it states 'aggregates and visualizes data from BigQuery'. How do you think it aggregates the data? It needs to run some queries against it.

In this case, the Data Viewer would not be required, as the App Engine service specifically targets the data to create visualizations from. In the other hand, Looker Studio etc. also require to be Data Viewer.

So in this case, we do not necessarily need Data Viewer, but definitely need Job User to do anything with the data for visualization purposes.

upvoted 5 times

 **JC0926** 7 months, 3 weeks ago

Selected Answer: A

aggregates and visualizes data from BigQuery.

upvoted 3 times

 **Andoameda9** 8 months ago

Why do we need an application to view the plain datasets, that can be done via bigquery console browser. I have to query some meaningful data from bigquery to aggregate and showcase it through an application. I will go with 'A'.

upvoted 3 times

 **kar_techie** 10 months ago

B is the answer

according to Google's least privilege model and question states that service account needs to have permission to access another project's BigQuery dataset.

(roles/bigquery.dataViewer) - Lowest-level resources where you can grant this role to Table to view data only for specific dataset

(roles/bigquery.jobUser) - Provides permissions to run jobs, including queries, within the project.

if so it has been assigned with jobUser it has privilege to access all datasets in that project and it is access violation

upvoted 1 times

 **romega2** 10 months, 1 week ago

Selected Answer: B

B seems better

upvoted 1 times

✉ **anolive** 10 months, 4 weeks ago

Selected Answer: B

the question said, the data needs to be visualized, it is not asking for query data.

upvoted 1 times

✉ **FeaRoX** 8 months, 1 week ago

How do you think app is getting the data for visualisation?

upvoted 2 times

✉ **Untamables** 11 months, 3 weeks ago

Selected Answer: A

Applications generally get data from BigQuery with querying. It needs BigQuery Job User role to query.

<https://cloud.google.com/bigquery/docs/access-control#bigquery>

I think using the default App Engine service account is recommended.

<https://cloud.google.com/appengine/docs/legacy/standard/python/service-account>

upvoted 2 times

✉ **AwesomeGCP** 12 months ago

Selected Answer: B

you want your application to be able to read data from the BigQuery dataset so it's simple B. Ask the other team to grant your default App Engine Service account the role of BigQuery Data Viewer

upvoted 2 times

You need to create a copy of a custom Compute Engine virtual machine (VM) to facilitate an expected increase in application traffic due to a business acquisition.

What should you do?

- A. Create a Compute Engine snapshot of your base VM. Create your images from that snapshot.
- B. Create a Compute Engine snapshot of your base VM. Create your instances from that snapshot.
- C. Create a custom Compute Engine image from a snapshot. Create your images from that image.
- D. Create a custom Compute Engine image from a snapshot. Create your instances from that image.

Correct Answer: D

A custom image belongs only to your project. To create an instance with a custom image, you must first have a custom image.

Reference:

<https://cloud.google.com/compute/docs/instances/create-start-instance>

Community vote distribution

B (52%)

D (48%)

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (D):

Preparing your instance for an image

You can create an image from a disk even while it is attached to a running VM instance. However, your image will be more reliable if you put the instance in a state that is easier for the image to capture. Use one of the following processes to prepare your boot disk for the image:

Stop the instance so that it can shut down and stop writing any data to the persistent disk.

If you can't stop your instance before you create the image, minimize the amount of writes to the disk and sync your file system.

Pause apps or operating system processes that write data to that persistent disk.

Run an app flush to disk if necessary. For example, MySQL has a FLUSH statement. Other apps might have similar processes.

Stop your apps from writing to your persistent disk.

Run sudo sync.

After you prepare the instance, create the image.

https://cloud.google.com/compute/docs/images/create-delete-deprecate-private-images#create_instance

upvoted 33 times

✉  **pca2b**  2 years, 6 months ago

B:

we just need to make 'a copy' of the VM, B works well for that.

not D: Had the question mentioned more copies, we would need to go the way of images...templates etc. D will work but not needed here.
upvoted 12 times

✉  **djgodzilla** 2 years, 3 months ago

custom images are better a fit if its for a new business workload you just acquired

upvoted 2 times

✉  **wolfie09** 1 year, 4 months ago

What about the answer that says create your instanceS ??

upvoted 1 times

✉  **gpais**  4 days, 13 hours ago

Selected Answer: D

<https://cloud.google.com/compute/docs/images/create-custom#image:create>
~:text=You%20can%20create%20disk%20images%20from%20the%20following%20sources%3A
upvoted 1 times

✉  **NoCrapEva** 1 month ago

Correct answer must be B:

While Option D: will work, there is NO mention of a pre-created snapshot, you cannot assume you already have a snapshot. For Option D: to work, the answer should read: Create a snapshot, then create a custom image from the snapshot, then create instances from that image....

upvoted 1 times

✉  **Ahmed_Y** 1 month, 1 week ago

Selected Answer: B

The correct answer is B.

Print all in one "The correct answer is B" and add a link to the question in the first post in the thread. Then add full

upvoted 2 times

CVGCP 4 months ago

Option D says, Create a custom Compute Engine image from a snapshot. For this we have to create a Snapshot first, we would not like to use an old snapshot. If we have to create a Snapshot, then we can directly use that Snapshot to create a VM, which is option B.

upvoted 2 times

Tarkik 5 months, 1 week ago

Selected Answer: D

use case of snapshot is for DR and backup. Images are more for creating identical VMs. So I would opt for Option D.

upvoted 1 times

rcheekati 5 months, 3 weeks ago

<https://cloud.google.com/compute/docs/instances/create-start-instance>

To quickly create more than one VM with the same boot disk, create a custom image, then create VMs from that image instead of using a snapshot.

ANS:D

upvoted 1 times

nooneknows 5 months, 4 weeks ago

Selected Answer: D

Option A is not recommended because snapshots are not directly usable as images for creating new instances. While snapshots are useful for backup and disaster recovery, they are not optimized for creating new VMs.

Option B is also not recommended because snapshots are not directly usable as instances. While it is possible to create new disks from snapshots, additional configuration steps would be required to create a new instance from a disk.

Option C is a valid option, but it requires additional steps to create an image from a snapshot before creating new images from that image.

Option D is the best option because it involves creating a custom Compute Engine image directly from the snapshot, which can be used to create new instances with minimal additional configuration. This approach is faster and simpler than creating images from images.

Therefore, the correct answer is option D: Create a custom Compute Engine image from a snapshot. Create your instances from that image.

upvoted 1 times

Bobbybash 7 months, 4 weeks ago

Selected Answer: D

D. Create a custom Compute Engine image from a snapshot. Create your instances from that image. This is the best approach to creating a copy of a custom Compute Engine virtual machine (VM). By creating a custom image, you can capture the exact configuration and software stack of the original VM. This image can then be used to create new instances with the same configuration and software stack, thus allowing you to create an identical copy of the original VM. This approach is much more efficient than copying and configuring each component of the original VM individually.

upvoted 1 times

Nazz1977 8 months, 1 week ago

Selected Answer: D

It is D.

If you deploy an VM from an snapshot, you gonna clone the VM, so you network will have many clones machines (they have the same SSID) , If you are in a Microsoft domain, the DC will see only one machine, even you have done the snapshot more than one.

upvoted 1 times

Nazz1977 8 months, 2 weeks ago

Selected Answer: D

D...Absolutely

upvoted 2 times

sidhu524 8 months, 3 weeks ago

Selected Answer: B

can some one explain it is B or D

upvoted 2 times

Nazz1977 8 months, 1 week ago

It is D.

First you have to create an image (from a snapshot of one VM). After you have an image you can deploy new VMs.

upvoted 1 times

Blitzer 8 months, 3 weeks ago

Selected Answer: B

Answer B: Create a VM from a snapshot

To quickly create more than one VM with the same boot disk, create a custom image, then create VMs from that image instead of using a snapshot.

<https://cloud.google.com/compute/docs/instances/create-start-instance#createSnapshot>

upvoted 3 times

ourst1951 9 months, 4 weeks ago

Selected Answer: B

The question says to create A COPY, not multiple copies. If we were doing more than one VM then we would select D for custom image.

<https://cloud.google.com/compute/docs/machine-images>

upvoted 3 times

 **kar_techie** 10 months ago

Correct answer is B

we can create a VM from Snapshots too.. Q states that copy of compute engine which is replica similar to existing running VM to meet the traffic. Why do we need to create a custom image then create a VM ? if it can be created via snapshot with few steps

upvoted 3 times

 **Nazz1977** 10 months, 1 week ago

Selected Answer: D

I have tested it. It is D

It is not B because if you make a snapshot and replicate it ,the new machines will have the same SSID, imagine if machine are in a domain, all the clones machines are viewing from domain controllers are the same. It doesn't work.

upvoted 7 times

You have deployed an application on a single Compute Engine instance. The application writes logs to disk. Users start reporting errors with the application. You want to diagnose the problem. What should you do?

- A. Navigate to Cloud Logging and view the application logs.
- B. Connect to the instance's serial console and read the application logs.
- C. Configure a Health Check on the instance and set a Low Healthy Threshold value.
- D. Install and configure the Cloud Logging Agent and view the logs from Cloud Logging.

Correct Answer: D

Community vote distribution

D (50%)	B (45%)	5%
---------	---------	----

✉  **hiteshrup**  3 years ago

Answer: D

App logs can't be visible to Cloud Logging until we install Cloud Logging Agent on GCE
upvoted 25 times

✉  **hiteshrup** 3 years ago

Continuation of reasoning.

If Problem statement is not having this statement "The application writes logs to disk", then we might assume that application is writing logs on Cloud Logging with google-fluentd agent API library. However, problem statement is clearly mentioned that logs are writing down on disk, we need agent installed on GCE to fetch those logs from disk to Cloud Logging. If that is not desirable, then option B is left

upvoted 5 times

✉  **hiteshrup** 3 years ago

(Correction) Answer is A after rethinking and doing some research by focusing words "App Engine", which has by default enabled Request Logs which has App logs on each request and those logs are enabled for Cloud Logging ..

https://cloud.google.com/appengine/docs/standard/python/logs#request_logs_vs_application_logs

upvoted 3 times

✉  **temple1305** 6 months, 1 week ago

Sorry, but where do you see App Engine?

upvoted 2 times

✉  **abirroy** 1 year, 2 months ago

For compute engine you need to install logging agent. Answer is D

upvoted 1 times

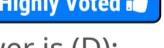
✉  **ashrafh** 2 years, 1 month ago

Hi all

check this document and decide :)

<https://cloud.google.com/logging/docs/agent/logging/installation>

upvoted 3 times

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (D):

In its default configuration, the Logging agent streams logs from common third-party applications and system software to Logging; review the list of default logs. You can configure the agent to stream additional logs; go to Configuring the Logging agent for details on agent configuration and operation.

It is a best practice to run the Logging agent on all your VM instances. The agent runs under both Linux and Windows. To install the Logging agent, go to Installing the agent.

<https://cloud.google.com/logging/docs/agent>

upvoted 18 times

✉  **ESP_SAP** 3 years, 1 month ago

CORRECTION.

Correct Answer is (A):

Activity logging is enabled by default for all Compute Engine projects.

You can see your project's activity logs through the Logs Viewer in the Google Cloud Console:

In the Cloud Console, go to the Logging page.

Go to the Logging page

When in the Logs Viewer, select and filter your resource type from the first drop-down list.

From the All logs drop-down list, select compute.googleapis.com/activity_log to see Compute Engine activity logs.

https://cloud.google.com/compute/docs/logging/activity-logs#viewing_logs

Besides:

Activity logs are provided as part of the Cloud Logging service. For more information about Logging in general, read the Cloud Logging documentation.

<https://cloud.google.com/compute/docs/logging/activity-logs>

upvoted 11 times

✉  **ESP_SAP** 3 years, 1 month ago

Additional information about VM Image for AWS EC2:

The Logging agent streams logs from your VM instances and from selected third-party software packages to Cloud Logging. It is a best practice to run the Logging agent on all your VM instances.

The VM images for Compute Engine and Amazon Elastic Compute Cloud (EC2) don't include the Logging agent, so you must complete these steps to install it on those instances. The agent runs under both Linux and Windows.

If your VMs are running in Google Kubernetes Engine or App Engine, the agent is already included in the VM image, so you can skip this page.

upvoted 3 times

✉  **magistrum** 2 years, 9 months ago

This points to D then

upvoted 3 times

✉  **babusartop17** 2 years, 1 month ago

I feel sorry for the woman in your life.

upvoted 17 times

✉  **DamonSalvatore** 2 years ago

Haha! That was funny

upvoted 2 times

✉  **mexblood1** 3 years, 1 month ago

Activity Logs do not include 3rd party application logs. Activity logs are more related to operations and changes in the infrastructure. This question is tricky, I think it's either D or B, because if it's only an application on a single instance, you can connect to the instance and read the application logs directly and you save the cost of logging agent.

upvoted 9 times

✉  **mexblood1** 3 years, 1 month ago

Maybe I was assuming serial console is the same than system console, technically I guess they're not the same, hence I guess D will be my chosen answer.

upvoted 1 times

✉  **sapguru** 3 years, 1 month ago

Cloud logging enabled by default for compute engine

upvoted 1 times

✉  **csrazdan** 1 year, 3 months ago

Do you mean the logging agent is installed by default? It depends on the OS you decide. For example, it is installed in Ubuntu but not on RedHat or Windows. Besides installing of the agent is not enough. You have to configure and let the agent know where your application is writing the logs on the disk so that it can monitor and stream the log to cloud monitoring. D is the correct answer

upvoted 1 times

✉  **Captain1212** Most Recent 1 month ago

Selected Answer: D

D makes more sense , as application writes logs to disk and to diagnose it we need the cloud logging agent

upvoted 1 times

✉  **Praxii** 5 months, 1 week ago

The line "application writes logs to disk" is crucial. It means logs are not available in cloud logging to yet. Hence we need to install the logging agent to send the logs to Cloud Logging.

Answer is D

upvoted 1 times

✉  **MahAli** 6 months ago

Selected Answer: D

Rethinking it again:

The serial console provides a text-based interface that allows you to view the boot-up process and access the login prompt of the instance. From there, you can troubleshoot issues, change configuration settings or perform other administrative tasks.

The serial console is a useful tool for troubleshooting issues with Compute Engine instances, especially in situations where network connectivity is not available or where the operating system is not functioning properly.

upvoted 1 times

✉  **MahAli** 6 months ago

Selected Answer: B

Errors are already on the disk, just ssh to the instance and read them
upvoted 3 times

✉ **xaqanik** 8 months ago

Selected Answer: D

To configure logging for your GCP VMs, you need to install the Cloud Logging agent on each of your instances. The Cloud Logging agent collects log data from the instance and sends it to the Cloud Logging API, where it is stored and can be viewed in the Cloud Logging console.

upvoted 2 times

✉ **FeaRoX** 8 months, 1 week ago

Selected Answer: B

We need to diagnose it quickly - simply login to instance and check logs. When there's application outage, it's not good idea to install cloud agent and wait until logs appear in the Cloud Logging

upvoted 2 times

✉ **eeeeee123** 8 months, 1 week ago

Serial port is not used for application debugging generally. It is mainly used for boot issues, network connectivity issues, accessing VM not possible by other methods

upvoted 1 times

✉ **alain_maza** 10 months ago

logging agent is deprecated,i find this in documentation

<https://cloud.google.com/logging/docs/agent/logging/installation>

"While this agent is still supported, we recommend that you use the Ops Agent for new Google Cloud workloads and eventually transition your existing Compute Engine VMs to use the new agent. The Ops Agent, which combines the collection of metrics and logging into a single agent, is the eventual replacement for the existing agents."

so the question points to the logging agent not to the ops agent so i think correct answer is A

upvoted 1 times

✉ **romea2** 10 months, 1 week ago

Selected Answer: D

D

Installing the Cloud Logging agent on individual VMs

If your VMs are running in Google Kubernetes Engine or App Engine, the agent is already included in the VM image, so you can skip this page.

upvoted 1 times

✉ **AwesomeGCP** 12 months ago

Selected Answer: D

D. Install and configure the Cloud Logging Agent and view the logs from Cloud Logging

upvoted 1 times

✉ **Charumathi** 1 year ago

Selected Answer: D

D is the correct answer,

Install fluentD logging agent to compute engine and view the logs from cloud logging.

upvoted 1 times

✉ **theBestStudent** 1 year, 1 month ago

Selected Answer: D

The assumption is errors are related to THE APPLICATION, and NOT to the VM instance itself (like booting, starting up, crashing, etc).

That being said, Cloud Logging is not by default installed in the VMS, and we need to install the agent first in order to see the application logs in Cloud Logging.

Of course another option would have been to just go to the file where the logs are located in the VM, but that is not an option in this question.

CORRECT ANSWER: D.

upvoted 2 times

✉ **abirroy** 1 year, 2 months ago

Selected Answer: D

For compute engine you need to install logging agent. Answer is D

upvoted 1 times

✉ **bobthebuilder55110** 1 year, 2 months ago

Selected Answer: B

Answer should be B, Read the question it clearly mentions that "The application writes logs to disk." What is the point of writing the log to disk if you don't want to use them. D does not make sense since why install an agent if you are already writing logs + the errors have already happened, How will the agent go back in time and give an error messages back?

upvoted 5 times

✉ **S00999** 1 year, 3 months ago

Selected Answer: D

Answer: D

Cloud Logging knows nothing about applications installed on the system without an agent collecting logs.

Using the serial console is not a best-practice and is impractical on a large scale.

Exam questions always invite responses in favor of the tools offered by the publisher.

upvoted 4 times

 **SergijF** 1 year, 4 months ago

Selected Answer: D

D is correct. Not B. You can't read anything from the disc with the serial console. You can only see the system events, not the logs of a running application that are already written to the disc. With the agent you are able to access your logs

upvoted 1 times

An application generates daily reports in a Compute Engine virtual machine (VM). The VM is in the project corp-iot-insights. Your team operates only in the project corp-aggregate-reports and needs a copy of the daily exports in the bucket corp-aggregate-reports-storage. You want to configure access so that the daily reports from the VM are available in the bucket corp-aggregate-reports-storage and use as few steps as possible while following Google-recommended practices. What should you do?

- A. Move both projects under the same folder.
- B. Grant the VM Service Account the role Storage Object Creator on corp-aggregate-reports-storage.
- C. Create a Shared VPC network between both projects. Grant the VM Service Account the role Storage Object Creator on corp-iot-insights.
- D. Make corp-aggregate-reports-storage public and create a folder with a pseudo-randomized suffix name. Share the folder with the IoT team.

Correct Answer: A

Reference:

<https://cloud.google.com/billing/docs/onboarding-checklist>*Community vote distribution* B (100%) **ESP_SAP**  3 years, 1 month ago

Correct Answer is (B):

Predefined roles

The following table describes Identity and Access Management (IAM) roles that are associated with Cloud Storage and lists the permissions that are contained in each role. Unless otherwise noted, these roles can be applied either to entire projects or specific buckets.

Storage Object Creator (roles/storage.objectCreator) Allows users to create objects. Does not give permission to view, delete, or overwrite objects.

<https://cloud.google.com/storage/docs/access-control/iam-roles#standard-roles>

upvoted 33 times

 **ESP_SAP** 3 years, 1 month ago

Basically, you are giving the permissions to the VM Service Account to create a copy of the daily report on the bucket that the other team has access.

upvoted 5 times

 **francisco_guerra**  3 years, 2 months ago

i think is B

upvoted 13 times

 **francisco_guerra** 3 years, 2 months ago

Object creator cant see object so i think is D

upvoted 1 times

 **Ixgywil** 2 years, 5 months ago

VM doesn't need to see the objects - just to create them. It's B:

The VM is located in project "corp-iot-insights" - give its SA the Storage Object Creator role for bucket "corp-aggregate-reports-storage" that is located in project "corp-aggregate-reports", where your team operates.

upvoted 4 times

 **Captain1212**  1 month ago**Selected Answer: B**

B is the correct as it gives the service account required access

upvoted 1 times

 **tatyavinchu** 1 month, 3 weeks ago**Selected Answer: B**

Correct Answer is B

upvoted 1 times

 **Naree** 3 months ago**Selected Answer: B**

Just take below sentence from the question which is added just for confusion :)

"Your team operates only in the project corp-aggregate-reports"

upvoted 1 times

 **Stef Johnson** 1 year ago

Correct Answer is (B)

upvoted 1 times

✉ **theBestStudent** 1 year, 1 month ago

Selected Answer: B

If that is the default service Account of the Compute Instance, then we should do nothing. As the role is already included. Either way, we should do nothing as the role is already covered. Also we shouldn't modify Compute instance Service account. But again, I will assume it is not the default.

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

B is right

upvoted 1 times

✉ **deadlydeb** 1 year, 3 months ago

Selected Answer: B

b it is

upvoted 1 times

✉ **[Removed]** 1 year, 8 months ago

Selected Answer: B

You should be able to add a service account to another project:

Create the first service account in project A in the Cloud Console. Activate it using gcloud auth activate-service-account.

In the Cloud Console, navigate to project B. Find the "IAM & admin" > "IAM" page. Click the "Add" button. In the "New members" field paste the name of the service account (it should look like a strange email address) and give it the appropriate role.

Run gcloud commands with --project set to project B. They should succeed (I just manually verified that this will work).

Automatic creation of service accounts is something that we're hesitant to do until we can work through all of the security ramifications.

<https://stackoverflow.com/a/35558464>

upvoted 4 times

✉ **Mixxer5** 1 year, 8 months ago

Selected Answer: B

It's B since bucket names are globally unique so it's enough to refer to them when you've proper role assigned

upvoted 3 times

✉ **arsh1916** 2 years, 4 months ago

B, assign access is less step

upvoted 1 times

✉ **pradr** 2 years, 6 months ago

From stackoverflow: Bucket names are globally unique, so your app will refer to an existing bucket in another project in the same way that it refers to buckets in its own project. Hence the shared VPC is not required to access the bucket. Just the IAM role.

upvoted 3 times

✉ **EABDAJA** 2 years, 7 months ago

B is correct

upvoted 1 times

✉ **GCP_Student1** 2 years, 7 months ago

B - Grant the VM Service Account the role Storage Object Creator on corp-aggregate-reports-storage.

upvoted 2 times

✉ **victory108** 2 years, 8 months ago

B - Grant the VM Service Account the role Storage Object Creator on corp-aggregate-reports-storage.

upvoted 1 times

✉ **Bhagirathi** 2 years, 10 months ago

why not A ?

upvoted 1 times

[lots of steps vs B](https://stackoverflow.com/a/35558464)

upvoted 1 times

You built an application on your development laptop that uses Google Cloud services. Your application uses Application Default Credentials for authentication and works fine on your development laptop. You want to migrate this application to a Compute Engine virtual machine (VM) and set up authentication using Google-recommended practices and minimal changes. What should you do?

- A. Assign appropriate access for Google services to the service account used by the Compute Engine VM.
- B. Create a service account with appropriate access for Google services, and configure the application to use this account.
- C. Store credentials for service accounts with appropriate access for Google services in a config file, and deploy this config file with your application.
- D. Store credentials for your user account with appropriate access for Google services in a config file, and deploy this config file with your application.

Correct Answer: B

Reference:

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances>

Community vote distribution

A (62%)

B (38%)

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (B):

Best practices

In general, Google recommends that each instance that needs to call a Google API should run as a service account with the minimum permissions necessary for that instance to do its job. In practice, this means you should configure service accounts for your instances with the following process:

Create a new service account rather than using the Compute Engine default service account.

Grant IAM roles to that service account for only the resources that it needs.

Configure the instance to run as that service account.

Grant the instance the <https://www.googleapis.com/auth/cloud-platform> scope to allow full access to all Google Cloud APIs, so that the IAM permissions of the instance are completely determined by the IAM roles of the service account.

Avoid granting more access than necessary and regularly check your service account permissions to make sure they are up-to-date.

https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#best_practices

upvoted 52 times

✉  **Ridhanya** 1 year, 10 months ago

you just gave justification for option A which is right

upvoted 1 times

✉  **ryumada** 1 year, 2 months ago

Maybe for the option A you are modifying the default service account because it's not explain which service account used by the VM, is it the default one or the new one?

The best practice is to Create a new service account rather than using the Compute Engine default service account.

B still has the bigger prove here as the answer.

upvoted 4 times

✉  **ryumada** 1 year, 2 months ago

You should read lxgywil comment. His comment explains how authentication works to access Google Services in your application.

a relevant link also:

https://cloud.google.com/storage/docs/reference/libraries#setting_up_authentication

upvoted 1 times

✉  **cRobert** 2 years, 9 months ago

From your quote:

Configure the "instance" to run as that service account.

From answer B:

and configure the "application" to use this account.

You don't add service accounts to applications, ans A

upvoted 19 times

✉  **lxgywil** 2 years, 5 months ago

When you use a GCP service within your app (code), you have to use its client libraries. When you instantiate a client with client libraries you can pass it a Service Account key, which will define on behalf of which SA the client will be acting. That's how you can configure your app to use a particular service account.

upvoted 3 times

✉  **akshaydoifode88** 10 months, 3 weeks ago

In question it's written application uses application default credentials. So taking that as a hint. B is the answer because here we are configuring service account key into the application. Similar approach.

upvoted 1 times

✉  **magistrum** 2 years, 9 months ago

wording is the clue :)

upvoted 1 times

✉  **TAvenger** 2 years, 7 months ago

It's dirty play with words... All understand that we need custom SA, grant required permissions and attach this SA to the VM... Why Google does this?

upvoted 7 times

✉  **filco72**  3 years, 1 month ago

I would choose: A. Assign appropriate access for Google services to the service account used by the Compute Engine VM. as there is no need to create a new service account.

upvoted 20 times

✉  **xaqanik** 7 months, 2 weeks ago

by default a vm uses a default service account. if you grant permission to this service account it will apply to all VMs default service accounts in the project . in this case you need create a new service account and give it appropriate permission

upvoted 1 times

✉  **Hjameel** 3 years, 1 month ago

I agree, there is no need to create a new service account

upvoted 9 times

✉  **Captain1212**  1 month ago

Selected Answer: A

A is the right answer , as after providing the appropriate access to the service account compute Engine

upvoted 1 times

✉  **MilanRajGupta** 2 months ago

The correct answer is "B". In general, Google recommends that each instance that needs to call a Google API should run as a service account with the minimum permissions necessary for that instance to do its job. In practice, this means you should configure service accounts for your instances with the following process: Create a new service account rather than using the Compute Engine default service account. Grant IAM roles to that service account for only the resources that it needs. Configure the instance to run as that service account. Grant the instance the <https://www.googleapis.com/auth/cloud-platform> scope to allow full access to all Google Cloud APIs, so that the IAM permissions of the instance are completely determined by the IAM roles of the service account. Avoid granting more access than necessary and regularly check your service account permissions to make sure they are up-to-date.

upvoted 1 times

✉  **3arle** 2 months, 1 week ago

B

"Many of these Google Cloud services also provide a default service account. Using the default service account is not recommended, because by default the default service account is highly privileged, which violates the principle of least privilege."

<https://cloud.google.com/docs/authentication/provide-credentials-adc#attached-sa>

upvoted 1 times

✉  **krop** 3 months, 3 weeks ago

Selected Answer: A

Based on the documentation here : <https://cloud.google.com/docs/authentication/application-default-credentials> looks like the correct answer is "A".

This is exactly why ADC has been created for. You develop your code on your laptop and you use it in your APP code ADC as a way to authorize to GCP - then depends on where you would like to test your code on - you simply execute `gcloud auth application-default login` in your system to store right credentials in your ADC on your laptop.

When you copy your code into PROD VM, your app without any changes will scan below locations in order to find credentials :

1. GOOGLE_APPLICATION_CREDENTIALS environment variable
2. User credentials set up by using the Google Cloud CLI
3. The attached service account, returned by the metadata server

As you can see above, your app will not find any credentials in 1. and 2. location but it will go to location 3, which is the credential for Service Account assigned to your VM.

Minimal effort here means, you don't need to change your APP to get right credentials to GCP services.

upvoted 2 times

✉  **geeroylenkins** 2 months, 3 weeks ago

This seems the best answer. Using ADC enables the App itself to use the VM's SA after migration. No change to the app is needed.

upvoted 1 times

✉  **eez64480** 4 months, 4 weeks ago

Correct answer is A

Application Default Credentials (ADC) is not a declaration of permissions, but more of a directive to find permissions already granted to the application's environment.

Referencing the link below, you could certainly perform answer B, but the question specifies "minimal changes". We assign appropriate permissions to the instance that will be running the application and the ADC will find them and use them to authenticate. Kind of dumps the permissions responsibility on the cloud engineer to make sure the instance has the proper permissions.

https://google-auth.readthedocs.io/en/master/user-guide.html#application-default-~:text=Applications%20running%20on,compute_engine.Credentials%3A
upvoted 1 times

✉ hanweiCN 5 months, 3 weeks ago

Selected Answer: A

no need to configure app to use service account
upvoted 3 times

✉ xaqanik 8 months ago

Selected Answer: B

go for B. Google recommends use a new service account instead of default VM service account for authendication.
upvoted 2 times

✉ Nazz1977 8 months, 2 weeks ago

Selected Answer: A

I think it is A "You want to migrate this application to a Compute Engine virtual machine (VM) ".....Only A has VM option
upvoted 1 times

✉ researched_answer_boi 8 months, 3 weeks ago

Key phrases: "an application that uses GC services", "migrate ... to Compute Engine VM", "Google-recommended best practices and minimal changes".
According to "<https://download.huihoo.com/google/gdgdevkit/DVD1/developers.google.com/compute/docs/authentication.html>", answer A is the "correctest".
upvoted 1 times

✉ Nazz1977 9 months ago

Selected Answer: A

I think it is A
upvoted 1 times

✉ Untamables 11 months, 3 weeks ago

Selected Answer: A

A service account is able to be attached to VMs, not applications in VM.
upvoted 3 times

✉ PKookNN 11 months, 3 weeks ago

Selected Answer: B

B since I am focusing on recommended practices which is to create a new service account and not to use the default one.
upvoted 1 times

✉ theBestStudent 11 months, 4 weeks ago

Selected Answer: A

- Assuming its not default Service answer is A.
-Assuming the GCE already have a Default Service Account associated, then B "sounds good". My problem here is it says "to configure the application with that Service account".

You don't configure the "application" with the service account. What you configure with your Service account is the GCE Instance. That would kill option B and would leave as "the most correct answer" option A.
upvoted 3 times

✉ Jas14 3 months, 4 weeks ago

this makes most sense to me.
upvoted 2 times

✉ Letahrgicbeagle 1 year ago

Selected Answer: B

Never edit the default service account as this may solve the purpose here but may pose a risk for future VMs to be created
upvoted 2 times

✉ theBestStudent 1 year, 1 month ago

Selected Answer: B

B, and never ever modify a default service account of a Compute Instance.
upvoted 2 times

✉ Gregwaw 1 month, 1 week ago

Create VM with custom service account ...
upvoted 1 times

You need to create a Compute Engine instance in a new project that doesn't exist yet. What should you do?

- A. Using the Cloud SDK, create a new project, enable the Compute Engine API in that project, and then create the instance specifying your new project.
- B. Enable the Compute Engine API in the Cloud Console, use the Cloud SDK to create the instance, and then use the --project flag to specify a new project.
- C. Using the Cloud SDK, create the new instance, and use the --project flag to specify the new project. Answer yes when prompted by Cloud SDK to enable the Compute Engine API.
- D. Enable the Compute Engine API in the Cloud Console. Go to the Compute Engine section of the Console to create a new instance, and look for the Create In A New Project option in the creation form.

Correct Answer: B

Community vote distribution

A (100%)

 **ESP_SAP** Highly Voted  3 years, 1 month ago

Correct Answer is (A):

Quickstart: Creating a New Instance Using the Command Line

Before you begin

1. In the Cloud Console, on the project selector page, select or create a Cloud project.
2. Make sure that billing is enabled for your Google Cloud project. Learn how to confirm billing is enabled for your project.

To use the gcloud command-line tool for this quickstart, you must first install and initialize the Cloud SDK:

1. Download and install the Cloud SDK using the instructions given on [Installing Google Cloud SDK](#).

2. Initialize the SDK using the instructions given on [Initializing Cloud SDK](#).

To use gcloud in Cloud Shell for this quickstart, first activate Cloud Shell using the instructions given on [Starting Cloud Shell](#).

<https://cloud.google.com/ai-platform/deep-learning-vm/docs/quickstart-cli#before-you-begin>

upvoted 38 times

 **filco72** Highly Voted  3 years, 1 month ago

I would choose A. Using the Cloud SDK, create a new project, enable the Compute Engine API in that project, and then create the instance specifying your new project.

as first I need to create a project. Instance creation cannot automatically create a project.

upvoted 12 times

 **Captain1212** Most Recent  1 month ago

Selected Answer: A

A is the correct answer, as it follow the right path , to create the compute Engine instance

upvoted 1 times

 **tatyavinchu** 1 month, 3 weeks ago

Correct Answer is A

upvoted 1 times

 **dr1ka** 1 year, 9 months ago

Selected Answer: A

Vote A

upvoted 2 times

 **jaffarali** 1 year, 9 months ago

Selected Answer: A

A is the correct option

upvoted 2 times

 **alaahakim** 1 year, 10 months ago

Selected Answer: A

Correct Ans is : A

upvoted 2 times

 **arsh1916** 2 years, 4 months ago

A is correct

upvoted 2 times

 **kopper2019** 2 years, 5 months ago

A the way to go

upvoted 1 times

 **pondai** 2 years, 6 months ago

Vote A

upvoted 1 times

 **Devgela** 2 years, 6 months ago

A for me

upvoted 1 times

 **EABDAJA** 2 years, 7 months ago

A is correct

upvoted 1 times

 **neerajgoyal** 2 years, 7 months ago

A is the right answer

upvoted 1 times

 **GCP_Student1** 2 years, 7 months ago

A. Using the Cloud SDK, create a new project, enable the Compute Engine API in that project, and then create the instance specifying your new project.

upvoted 2 times

 **JackGlemins** 2 years, 7 months ago

I think is A: <https://cloud.google.com/sdk/gcloud/reference/projects/create>

upvoted 1 times

 **lutoa** 2 years, 7 months ago

yep has to be A, new project must be created first

upvoted 1 times

 **RockAJ** 3 years ago

A for me

upvoted 1 times

Your company runs one batch process in an on-premises server that takes around 30 hours to complete. The task runs monthly, can be performed offline, and must be restarted if interrupted. You want to migrate this workload to the cloud while minimizing cost. What should you do?

- A. Migrate the workload to a Compute Engine Preemptible VM.
- B. Migrate the workload to a Google Kubernetes Engine cluster with Preemptible nodes.
- C. Migrate the workload to a Compute Engine VM. Start and stop the instance as needed.
- D. Create an Instance Template with Preemptible VMs On. Create a Managed Instance Group from the template and adjust Target CPU Utilization. Migrate the workload.

Correct Answer: B

Reference:

<https://cloud.google.com/kubernetes-engine/docs/tutorials/migrating-node-pool>

Community vote distribution

C (100%)

✉  **juliandm**  3 years, 1 month ago

i understand preemptible as a no-go because of "must be restarted if interrupted" here meaning "starting from scratch". So C seems right
upvoted 42 times

✉  **jcloud965** 2 years, 2 months ago

I agree, C.
you won't run 30 hours job on preemptible instances that can be stopped at any time and can't run more than 24 hours.
If the job could be splitted, then preemptible VM is an option.
upvoted 5 times

✉  **Vador** 2 years ago

Preemptible seems fine on batch jobs for at least 24hours, not the case in here
upvoted 1 times

✉  **dttnc1** 1 year, 12 months ago

I agree with C. You can't risk running a processes that take 30 hours on a preemptible VM (Compute Engine always stops preemptible instances after they run for 24 hours). They are good for "short-lived" batch jobs. The scenario is NOT fault tolerant as the whole process restarts if interrupted.

<https://cloud.google.com/compute/docs/instances/preemptible>

upvoted 4 times

✉  **stepkurniawan**  3 years, 1 month ago

Preemptible will be perfect for a batch job that takes less than 24 hours. But it's not in this case.
upvoted 16 times

✉  **Linus11** 2 years, 5 months ago

What if it is a managed group of Pre emptible instances like in D. If one instance stops, another instance will take over.

I choose D.

upvoted 5 times

✉  **sanhoo** 2 years, 4 months ago

is there an option to specify Pre emptible instances while creating template? I couldn't find that. If so then D can't be true
upvoted 1 times

✉  **djgodzilla** 2 years, 3 months ago

Yes under management> Availability policy > preemptibility ON/OFF
upvoted 2 times

✉  **Captain1212**  1 month ago

Selected Answer: C

Option C is correct, bcoz the job is running for more than 30 hours
upvoted 1 times

✉  **tatyavinchu** 1 month, 3 weeks ago

Correct Answer is C

upvoted 1 times

✉  **Naree** 3 months ago

Selected Answer: C

Job runs for 30 hours and must be restarted if interrupted are "indirectly proportional" to "Preemptible"

Ans: C

upvoted 1 times

✉ **ankyt9** 10 months, 1 week ago

Selected Answer: C

Preemptible VMs are cheaper, but they will not be available beyond 24hrs

upvoted 3 times

✉ **Charumathi** 1 year ago

Selected Answer: C

C is the correct answer,

Install the workload in a compute engine VM, start and stop the instance as needed, because as per the question the VM runs for 30 hours, process can be performed offline and should not be interrupted, if interrupted we need to restart the batch process again.

Preemptible VMs are cheaper, but they will not be available beyond 24hrs, and if the process gets interrupted the preemptible VM will restart.

upvoted 2 times

✉ **KapilDhamija** 1 year, 1 month ago

Selected Answer: C

C is the correct answer

upvoted 1 times

✉ **ryumada** 1 year, 2 months ago

Selected Answer: C

The preemptible instance in GKE is same as Compute Engine Instance. They have same behavior that will be last for 24 hours.

Also, see the key here "...and must be restarted if interrupted.". That means the job will start from the scratch again if the preemptible instance terminated. So, you will just waste your preemptible instances because the job will never be finished.

<https://cloud.google.com/kubernetes-engine/docs/how-to/preemptible-vms#overview>

upvoted 1 times

✉ **koko2314** 1 year, 2 months ago

Option D will achieve the goal here. If a preemptible VM goes down in 24 hours a new one will be built for running the batch process through the instance group configuration.

upvoted 1 times

✉ **ryumada** 1 year, 2 months ago

"...and must be restarted if interrupted."

The job will start again from a scratch, then run again for another 24 hours in a new preemptible instance, the job will be terminated again after 24 hours, then start again from a scratch for another 24 hours in a new preemptible instance. This makes the process to be an infinite-loop process with wasted resources.

upvoted 2 times

✉ **AzureDP900** 1 year, 3 months ago

A preemptible VM is an instance that you can create and run at a much lower price than normal instances. However, Compute Engine might stop (preempt) these instances if it requires access to those resources for other tasks. Preemptible instances are excess Compute Engine capacity, so their availability varies with usage.

If your apps are fault-tolerant and can withstand possible instance preemptions then preemptible instances can reduce your Compute Engine costs significantly. For example, batch processing jobs can run on preemptible instances. If some of those instances stop during processing, the job slows but does not completely stop. Preemptible instances complete your batch processing tasks without placing additional workload on your existing instances and without requiring you to pay full price for additional normal instances.

Hence, the correct answer is: Build an instance template configured to launch a Preemptible VM. Provision a managed instance group (MIG) from the template you just created. Adjust the Target CPU Utilization setting.

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

C is correct, since job needs to run 30 hours.

upvoted 1 times

✉ **pfabio** 1 year, 4 months ago

Selected Answer: C

A - D are not possible because: For example, preemptible VMs can only run for up to 24 hours at a time (<https://cloud.google.com/compute/docs/instances/preemptible>)

I agree with C

upvoted 1 times

✉ **rsuresh27** 1 year, 5 months ago

The answer is C. Many people are saying D but this is incorrect. Preemptible instances are finite Compute Engine resources, so they might not always be available. If there is a long running job that must be restarted, it does not make sense to use preemptible VMs as if there are no resources available, we will not be able to even run or restart the job. Thus, C is correct.

upvoted 2 times

✉ **dinesh198728** 1 year, 6 months ago

Selected Answer: C

auto restart not in Prem VM

upvoted 1 times

 **ArunTaneja** 1 year, 8 months ago

Selected Answer: C

It should be C Because a Preemptible VM can't run more than 24 hours

upvoted 1 times

 **shputhan** 1 year, 8 months ago

Answer is C

- Since it is run on premise server, the expectation is to run it on Compute Engine.

- Preemptible VMs are not an option as it runs for 24 hours only continuously. The batch job takes 30 hours and expects it to be restarted incase of interruptions. (Note it is restart & not resume)

upvoted 1 times

 **emv** 1 year, 8 months ago

If your apps are fault-tolerant and can withstand possible instance preemptions, then preemptible instances can reduce your Compute Engine costs significantly. For example, batch processing jobs can run on preemptible instances.

upvoted 1 times

You are developing a new application and are looking for a Jenkins installation to build and deploy your source code. You want to automate the installation as quickly and easily as possible. What should you do?

- A. Deploy Jenkins through the Google Cloud Marketplace.
- B. Create a new Compute Engine instance. Run the Jenkins executable.
- C. Create a new Kubernetes Engine cluster. Create a deployment for the Jenkins image.
- D. Create an instance template with the Jenkins executable. Create a managed instance group with this template.

Correct Answer: C

Reference:

<https://cloud.google.com/solutions/jenkins-on-kubernetes-engine>*Community vote distribution*

A (100%)

filco72 Highly Voted 3 years, 1 month ago

I would choose A. Deploy Jenkins through the Google Cloud Marketplace.
as this is a well known opportunity on the GCP Marketplace

upvoted 28 times

ESP_SAP Highly Voted 3 years, 1 month ago

Correct Answer is (A):

Installing Jenkins

In this section, you use Cloud Marketplace to provision a Jenkins instance. You customize this instance to use the agent image you created in the previous section.

Go to the Cloud Marketplace solution for Jenkins.

Click Launch on Compute Engine.

Change the Machine Type field to 4 vCPUs 15 GB Memory, n1-standard-4.

Machine type selection for Jenkins deployment.

Click Deploy and wait for your Jenkins instance to finish being provisioned. When it is finished, you will see:

Jenkins has been deployed.

https://cloud.google.com/solutions/using-jenkins-for-distributed-builds-on-compute-engine#installing_jenkins
upvoted 18 times

Captain1212 Most Recent 1 month agoSelected Answer: A

Answer is : A as it's the most quick option

upvoted 1 times

tatyavinchu 1 month, 3 weeks ago

Correct Answer is C

upvoted 1 times

sabrinakloud 5 months, 3 weeks agoSelected Answer: A

A is correct

upvoted 1 times

Charumathi 1 year agoSelected Answer: A

A is correct answer,
To quickly deploy Jenkins, deploy it through google cloud marketplace.

upvoted 1 times

KapilDhamija 1 year, 1 month agoSelected Answer: A

remember as quickly as possible, also Google encourage things to be performed in minimal steps so A is the quickest and easiest choice

upvoted 1 times

kiwi123 1 year, 1 month ago

Go for A, the easiest

upvoted 1 times

 **pspandher** 1 year, 3 months ago

This is Repeat Question.

upvoted 2 times

 **AzureDP900** 1 year, 3 months ago

Cloud Market Place is fastest and best .. A is right

upvoted 1 times

 **pfabio** 1 year, 4 months ago

Selected Answer: A

A - ... as quickly and easily as possible

upvoted 1 times

 **RazOr** 1 year, 8 months ago

Selected Answer: A

A is right the rest is nonsense

upvoted 1 times

 **Raju44** 1 year, 9 months ago

As we need to automate installation of Jenkins as pre-requisite (upgrade etc for future), option C coz with Kubernetes you can automate installation using charts (easy to change few parameters)

upvoted 1 times

 **jaffarali** 1 year, 9 months ago

Selected Answer: A

A is right. Using Cloud Marketplace

upvoted 1 times

 **alaahakim** 1 year, 10 months ago

Selected Answer: A

I agree with A

upvoted 1 times

 **maggieli** 1 year, 11 months ago

A. The faster and best solution.

upvoted 1 times

 **jackwillis** 2 years ago

its A.

upvoted 1 times

You have downloaded and installed the gcloud command line interface (CLI) and have authenticated with your Google Account. Most of your Compute Engine instances in your project run in the europe-west1-d zone. You want to avoid having to specify this zone with each CLI command when managing these instances.

What should you do?

- A. Set the europe-west1-d zone as the default zone using the gcloud config subcommand.
- B. In the Settings page for Compute Engine under Default location, set the zone to europe-west1-d.
- C. In the CLI installation directory, create a file called default.conf containing zone=europe-west1-d.
- D. Create a Metadata entry on the Compute Engine page with key compute/zone and value europe-west1-d.

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/gcloud-compute>

Community vote distribution

A (92%) 8%

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (A):

Change your default zone and region in the metadata server

Note: This only applies to the default configuration.

You can change the default zone and region in your metadata server by making a request to the metadata server. For example:

```
gcloud compute project-info add-metadata \
--metadata google-compute-default-region=europe-west1,google-compute-default-zone=europe-west1-b
```

The gcloud command-line tool only picks up on new default zone and region changes after you rerun the gcloud init command. After updating your default metadata, run gcloud init to reinitialize your default configuration.

https://cloud.google.com/compute/docs/gcloud-compute#change_your_default_zone_and_region_in_the_metadata_server
upvoted 27 times

✉  **xtian2900** 3 years ago

does your comment imply that the answer is D ? i'm confused

upvoted 3 times

✉  **mahesh0049** 1 year, 8 months ago

every thing is correct in your explanation but instead of using gcloud compute command they used gcloud config.

upvoted 4 times

✉  **bobthebuilder55110** 1 year, 2 months ago

You can use the gcloud config set command here, https://cloud.google.com/compute/docs/gcloud-compute#set_default_zone_and_region_in_your_local_client

upvoted 4 times

✉  **dtncl** 1 year, 12 months ago

I agree the answer is A.

gcloud config - view and edit Cloud SDK properties

zone

Default zone to use when working with zonal Compute Engine resources.

<https://cloud.google.com/sdk/gcloud/reference/config>

upvoted 3 times

✉  **Examan1** 2 years, 7 months ago

Using gcloud config you can set the zone in your active configuration only. This setting does not apply to other gcloud configurations and does not become the default for the project.

Ref: <https://cloud.google.com/sdk/gcloud/reference/config/set>

So I believe correct answer is B as per <https://cloud.google.com/compute/docs/regions-zones/changing-default-zone-region#console>

In the Cloud Console, go to the Settings page.

From the Zone drop-down menu, select a default zone.

upvoted 3 times

✉  **tavva_prudhvi** 2 years, 6 months ago

bro, it mentioned going into the console settings, not the compute engine settings!

In the Cloud Console, go to the Settings page.

Go to the Settings page

From the Region drop-down menu, select a default region.

From the Zone drop-down menu, select a default zone.

upvoted 1 times

 **jcloud965** 2 years, 2 months ago

This setting in the Cloud Console won't be taken into account for gcloud on your active config

upvoted 2 times

 **SSPC**  3 years, 1 month ago

I would go with the answer A

upvoted 10 times

 **Captain1212**  1 month ago

Selected Answer: A

A is the correct answer, just set default in cloud console in the starting by the cloud shell commands

upvoted 1 times

 **tatyavinchu** 1 month, 3 weeks ago

Correct Answer is A

upvoted 1 times

 **GokulVelusaamy** 8 months, 2 weeks ago

Selected Answer: A

We can set the default zone using the below CLI command,

gcloud config set compute/zone ZONE

Refer : https://cloud.google.com/compute/docs/gcloud-compute#set_default_zone_and_region_in_your_local_client

upvoted 1 times

 **Angel_99** 1 year, 1 month ago

Selected Answer: A

It is clearly mentioned it is to be done via CLI not console

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right

upvoted 1 times

 **Tirthankar17** 1 year, 4 months ago

Selected Answer: A

A, it is clearly mentioned it is to be done via CLI not console.

upvoted 1 times

 **pfabio** 1 year, 4 months ago

Selected Answer: A

Just run this command on CLI

gcloud config set compute/zone ZONE

upvoted 1 times

 **Rukman** 1 year, 6 months ago

Selected Answer: A

Ans: A

"You want to avoid having to specify this zone with each CLI command"

its about CLI not console!

upvoted 2 times

 **injarapu** 1 year, 7 months ago

Selected Answer: B

<https://cloud.google.com/compute/docs/regions-zones/changing-default-zone-region#console>

upvoted 1 times

 **oracle111** 1 year, 7 months ago

Selected Answer: A

<https://cloud.google.com/compute/docs/gcloud-compute>

gcloud config set compute/zone ZONE

upvoted 3 times

 **ElenaL** 1 year, 9 months ago

Selected Answer: A

in GCP quicklabs, this is how they ask you set the zone primarily, so I go with A
upvoted 1 times

 **Raju44** 1 year, 9 months ago

Answer A. There is no file like default.config. In my labs i set default values Zone/Region/project etc default values stored in the following location
cat ~/.config/gcloud/configurations/config_default.
We can also set zone default using " gcloud config set compute/zone <zona_name>"
upvoted 2 times

 **jaffarali** 1 year, 9 months ago

A is the right option
upvoted 1 times

 **Krishna0014** 1 year, 10 months ago

Answer is A: Here we are talking about local CLI client and setting a default zone. It has nothing to do in cloud console any directory. simply can be done with below one command.

As per : https://cloud.google.com/compute/docs/gcloud-compute#set_default_zone_and_region_in_your_local_client

"Set default region and zone in your local client
The default region and zone set in your local client override the default region and zone set in the metadata server."

To set the default region and zone in your local client, run the following commands:

gcloud config set compute/region REGION
gcloud config set compute/zone ZONE
upvoted 2 times

 **wh1t4k3r** 1 year, 10 months ago

I would say A
<https://cloud.google.com/kubernetes-engine/docs/how-to/enable-oslogin?hl=nl#gcloud-config>
upvoted 1 times

The core business of your company is to rent out construction equipment at large scale. All the equipment that is being rented out has been equipped with multiple sensors that send event information every few seconds. These signals can vary from engine status, distance traveled, fuel level, and more. Customers are billed based on the consumption monitored by these sensors. You expect high throughput up to thousands of events per hour per device and need to retrieve consistent data based on the time of the event. Storing and retrieving individual signals should be atomic. What should you do?

- A. Create a file in Cloud Storage per device and append new data to that file.
- B. Create a file in Cloud Filestore per device and append new data to that file.
- C. Ingest the data into Datastore. Store data in an entity group based on the device.
- D. Ingest the data into Cloud Bigtable. Create a row key based on the event timestamp.

Correct Answer: D

Community vote distribution

D (71%) C (29%)

 **hiteshrup**  3 years ago

Answer: D

Keyword need to look for

- "High Throughput",
- "Consistent",
- "Property based data insert/fetch like engine status, distance traveled, fuel level, and more." which can be designed in column,
- "Large Scale Customer Base + Each Customer has multiple sensor which send event in seconds" This will go for pera bytes situation,
- Export data based on the time of the event.
- Atomic

o BigTable will fit all requirement.

o DataStore is not fully Atomic

o CloudStorage is not a option where we can export data based on time of event. We need another solution to do that

o FireStore can be used with MobileSDK.

So go with Option D: Big Table

upvoted 41 times

 **[Removed]** 3 years ago

Its File store not firestore. But still, the argument is correct here as file store will not be automatic like cloud store that no SQL.

If it would be Firestore I would have gone with it, then big data for throughout v/s cost.

upvoted 4 times

 **Hjameel**  3 years, 1 month ago

D is the best answer , Cloud Bigtable

upvoted 8 times

 **har_riy** 2 years, 7 months ago

Simple analogy.

Information every few seconds --> Time Series --> Big Table

upvoted 5 times

 **Captain1212**  1 month ago

Selected Answer: D

D is the right answer, as its can help with Automatically

upvoted 1 times

 **Tosssha** 7 months, 3 weeks ago

ReadModifyWriteRow requests are atomic:

<https://cloud.google.com/bigtable/docs/writes>

upvoted 1 times

 **xaqanik** 7 months, 3 weeks ago

it say any type of information. it means there can be an image file for instance. so, Bigtable is the best fit for this scenario

upvoted 1 times

 **alex000** 9 months, 1 week ago

Selected Answer: C

Answer: C

key work: atomic transaction

<https://cloud.google.com/datastore/docs/concepts/overview>

upvoted 2 times

✉  **[Removed]** 5 months, 4 weeks ago
high throughput == Bigtable, hence the answer is D
upvoted 1 times

✉  **Cornholio_LMC** 1 year ago
had this question today
upvoted 5 times

✉  **tomis2** 1 year, 3 months ago
Selected Answer: D
Timeseries + IoT = Bigtable
upvoted 4 times

✉  **AzureDP900** 1 year, 3 months ago
This is related to IoT , it is no sql means BigTable. D is right
upvoted 2 times

✉  **pfabio** 1 year, 4 months ago
Selected Answer: D
D - Ideal for use cases such as personalization, ad tech, fintech, digital media, and IoT
<https://cloud.google.com/bigtable>
upvoted 3 times

✉  **JelloMan** 1 year, 5 months ago
Selected Answer: D
Big table for the exact reasons as @hitrshrup mentioned. Time related data flowing in and out will point you directly to BigTable without reading anything else
upvoted 1 times

✉  **wael_tn** 1 year, 5 months ago
Selected Answer: D
When we want to store data based on timestamp => Cloud Bigtable
upvoted 1 times

✉  **Akash7** 1 year, 6 months ago
Answer is D, large streaming data = bigtable
upvoted 1 times

✉  **zaxma** 1 year, 6 months ago
Selected Answer: C
I vote c with DataStore (maybe even better with firestore includes pubsub ability)
Looks like the internet agrees with me.
upvoted 2 times

✉  **alaahakim** 1 year, 10 months ago
Answer: D
upvoted 2 times

✉  **CompteBidon91** 2 years, 3 months ago
D,
Keyword : IOT, Database, large scale = Big Table
upvoted 3 times

✉  **vinaybattula** 2 years, 4 months ago
D is correct. If Data is huge, unstructured and related time Bigtable is best option
upvoted 1 times

You are asked to set up application performance monitoring on Google Cloud projects A, B, and C as a single pane of glass. You want to monitor CPU, memory, and disk. What should you do?

- A. Enable API and then share charts from project A, B, and C.
- B. Enable API and then give the metrics.reader role to projects A, B, and C.
- C. Enable API and then use default dashboards to view all projects in sequence.
- D. Enable API, create a workspace under project A, and then add projects B and C.

Correct Answer: C

Community vote distribution

D (100%)

- ✉ **jclaudie** Highly Voted 3 years, 1 month ago
- D. workspaces is made for monitoring multiple projects.
upvoted 36 times
- ✉ **Hjameel** Highly Voted 3 years, 1 month ago
- D , Workspace to monitor multiple projects.
upvoted 11 times
- ✉ **Khoka** 2 years, 11 months ago
- <https://cloud.google.com/monitoring/worksheets>
upvoted 1 times
- ✉ **Captain1212** Most Recent 1 month ago
- Selected Answer: D
Option D is the correct Answer, First create the Workspace under A then add it to the Project B and C
upvoted 1 times
- ✉ **Bobbybash** 7 months, 4 weeks ago
- Selected Answer: D
D. Enable API, create a workspace under project A, and then add projects B and C.
- To monitor multiple Google Cloud projects in a single pane of glass, you can use Google Cloud's operations suite, formerly known as Stackdriver. By enabling the Cloud Monitoring API and creating a workspace under project A, you can add projects B and C to the same workspace. This will allow you to view metrics for CPU, memory, and disk usage for all projects in the same workspace. You can also set up alerting policies to be notified of any potential issues across all projects.
- Enabling the API alone or giving metrics.reader role to the projects will not provide a single pane of glass view of all the projects. Similarly, using default dashboards will not provide a unified view of all projects in a single dashboard.
upvoted 3 times
- ✉ **SathishBandi** 9 months, 3 weeks ago
- Selected Answer: D
In question, mentioned 'as a single pane of glass' and workspace are meant for Monitoring
upvoted 2 times
- ✉ **Charumathi** 1 year ago
- Selected Answer: D
D is the correct answer,
Keep Project A as host project in workspace and Project B and C as Service Project, and monitor the metrics of the Project A for a centralized view.
upvoted 1 times
- ✉ **ale_brd_** 1 year ago
- Selected Answer: D
Stackdriver workspaces are deprecated, now in the monitoring page of the Project you want, you need to select the "Scopes". Anyway he closest answer is D.
- Scopes allow you to monitor multiple projects.
- <https://cloud.google.com/monitoring/settings/multiple-projects>
upvoted 2 times
- ✉ **KapilDhamija** 1 year, 1 month ago
- Selected Answer: D
D should be the correct answer
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

D is correct

upvoted 1 times

 **Rutu_98** 1 year, 4 months ago

Selected Answer: D

D is correct

upvoted 1 times

 **JelloMan** 1 year, 5 months ago

D. One project must be the host (A) in this case and all others can be linked (B&C)

upvoted 2 times

 **Ken701** 1 year, 6 months ago

Selected Answer: D

You can add/link multiple project to a single workspace for monitoring

upvoted 1 times

 **gkb** 1 year, 8 months ago

Wanted to check, if somebody has appeared ACE exam recently and if yes, what is the percentage of questions that come from this site?
Thanks for answering.

upvoted 5 times

 **Eminenza22** 1 year, 8 months ago

D is correct

Workspace can monitor multiple projects but a Google Cloud project can be monitored by exactly 1 Workspace.

upvoted 2 times

 **HansKloss611** 1 year, 8 months ago

Selected Answer: D

D is the correct

upvoted 2 times

 **alex90fin** 1 year, 10 months ago

workspace is mandatory for monitoring

upvoted 1 times

 **alaahakim** 1 year, 10 months ago

i Vote D

upvoted 2 times

You created several resources in multiple Google Cloud projects. All projects are linked to different billing accounts. To better estimate future charges, you want to have a single visual representation of all costs incurred. You want to include new cost data as soon as possible. What should you do?

- A. Configure Billing Data Export to BigQuery and visualize the data in Data Studio.
- B. Visit the Cost Table page to get a CSV export and visualize it using Data Studio.
- C. Fill all resources in the Pricing Calculator to get an estimate of the monthly cost.
- D. Use the Reports view in the Cloud Billing Console to view the desired cost information.

Correct Answer: A

Reference:

<https://cloud.google.com/billing/docs/how-to/visualize-data>*Community vote distribution*

A (75%) D (25%)

Zakaullah Highly Voted 3 years, 1 month ago

Highly recommended website for exam prep.just passed the exam.
upvoted 33 times

IRobot 3 years, 1 month ago

Hope the best for my exam tomorrow
upvoted 6 times

[Removed] 3 years, 1 month ago

how did it go for you? I sit mine tomorrow.
upvoted 3 times

MoreOps 1 year, 11 months ago

Mine is tommorow lol.
upvoted 2 times

Pranad 1 year, 12 months ago

My exam is tomorrow, hope for the best.
upvoted 2 times

moses101 1 year, 11 months ago

how did things go?
upvoted 1 times

arsh1916 Highly Voted 2 years, 4 months ago

A is best answer
upvoted 9 times

SK1990 Most Recent 9 months, 2 weeks ago

Selected Answer: A
A is the best anmwsr.
upvoted 1 times

davidsalomon 10 months ago

It can also be D, as I checked, the console has a nice view now for reporting and this question is old.
upvoted 2 times

PKookNN 11 months, 3 weeks ago

Selected Answer: A
A is more correct as you can show data from multiple billing accounts as well as different projects.
upvoted 3 times

AwesomeGCP 12 months ago

A. Configure Billing Data Export to BigQuery and visualize the data in Data Studio.
upvoted 2 times

Charumathi 1 year ago

Selected Answer: D
D is the correct Answer,
Use the Reports in the Cloud Billing console to view the billing information for the resources in multiple projects.

upvoted 2 times

 **eBookz** 7 months, 4 weeks ago

While D is a possible answer, it is not in this case because "all projects are linked to different billing accounts".

Based on below quoted line from the link you shared, D would have been the answer if all projects were linked to the same billing account:

"On the Account management page, linked projects are listed under Projects linked to this billing account."

upvoted 1 times

 **NikRi** 9 months, 2 weeks ago

here projects are associated to different billing accounts! Reports view provides trends for a single billing account for date range. So ans is A

upvoted 2 times

 **KapilDhamija** 1 year, 1 month ago

Selected Answer: A

Would go with A (Configure Billing Data Export to BigQuery and visualize the data in Data Studio)

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is the best ..

upvoted 1 times

 **Tirthankar17** 1 year, 4 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **EABAJA** 2 years, 7 months ago

A is correct

upvoted 6 times

 **GCP_Student1** 2 years, 7 months ago

A. Configure Billing Data Export to BigQuery and visualize the data in Data Studio.

upvoted 6 times

 **nitinz** 2 years, 7 months ago

A, Any time you see question about billing and has BQ - 99.99% of time correct answer is that has BQ in it.

upvoted 8 times

 **victory108** 2 years, 8 months ago

A - Configure Billing Data Export to BigQuery and visualize the data in Data Studio.

upvoted 3 times

 **sarahf** 2 years, 9 months ago

I was thinking A first.

<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>

"Cloud Billing export to BigQuery enables you to export detailed Google Cloud billing data (such as usage, cost estimates, and pricing data) automatically throughout the day to a BigQuery dataset that you specify."

It also says that:

"Be aware that your BigQuery dataset only reflects Google Cloud billing data incurred from the date you set up Cloud Billing export, and after. That is, Google Cloud billing data is not added retroactively, so you won't see Cloud Billing data from before you enable export."

But I am leaning towards alternative D after reading this page: <https://cloud.google.com/billing/docs/how-to/reports>.

You don't have to export/import to see data and you can see costs from different projects. All you need is the permission: `billing.accounts.getSpendingInformation`. It can forecast future costs too. According to the video at the same page, BigQuery exports to Data Studio is for deeper cost analysis.

upvoted 6 times

 **rramos96** 2 years, 9 months ago

As we can read in the 1st paragraph at <https://cloud.google.com/billing/docs/how-to/reports>, Reports View from Cloud Billing Console shows info about a single Billing Account.

This question says:

- "all projects are linked to different billing accounts".
- "we want to have a single visual representation".

So, in my opinion, the best answer is [A] (BQ+data studio)

upvoted 9 times

 **don_v** 2 years, 10 months ago

I guess it's A.

I'm not sure how to get that info for multiple billing accounts with the billing report as they say "All projects are linked to different billing accounts."

upvoted 1 times

 **Bhagirathi** 2 years, 10 months ago

A or D which one to choose ?

upvoted 2 times

Your company has workloads running on Compute Engine and on-premises. The Google Cloud Virtual Private Cloud (VPC) is connected to your WAN over a Virtual Private Network (VPN). You need to deploy a new Compute Engine instance and ensure that no public Internet traffic can be routed to it. What should you do?

- A. Create the instance without a public IP address.
- B. Create the instance with Private Google Access enabled.
- C. Create a deny-all egress firewall rule on the VPC network.
- D. Create a route on the VPC to route all traffic to the instance over the VPN tunnel.

Correct Answer: B

Get private access to Google services, such as storage, big data, analytics, or machine learning, without having to give your service a public IP address.

Reference:

<https://cloud.google.com/vpc>

Community vote distribution



✉️ [User] **[Removed]** Highly Voted 3 years, 1 month ago

A for sure

B - this allows internal communication, but does nothing to limit public traffic

C - deny all is nice, but it's for egress -- we're looking for ingress

D - this is way to invasive and doesn't explicitly address the issue of preventing public internet traffic from reaching your instance -- if it does, someone let me know how.

upvoted 29 times

✉️ [User] **MohammedGhouse** Highly Voted 3 years, 1 month ago

A: answer looks right

upvoted 12 times

✉️ [User] **Captain1212** Most Recent 1 month ago

Selected Answer: A

A is the correct option, as other not limit the ingress traffic

upvoted 1 times

✉️ [User] **dataengineeruser34** 8 months ago

A for sure

upvoted 1 times

✉️ [User] **SK1990** 9 months, 2 weeks ago

Selected Answer: A

A - for sure

upvoted 1 times

✉️ [User] **SK1990** 9 months, 2 weeks ago

Selected Answer: A

A is the best answer.

upvoted 1 times

✉️ [User] **Nazz1977** 10 months ago

Selected Answer: A

A for sure

upvoted 1 times

✉️ [User] **Sam98845** 11 months, 2 weeks ago

should be A. VMs cannot communicate over the internet without a public IP address. Private Google Access permits access to Google APIs and services in Google's production infrastructure.

<https://cloud.google.com/vpc/docs/private-google-access>

upvoted 1 times

✉️ [User] **kailash** 12 months ago

Selected Answer: A

Elimination

upvoted 1 times

 **Charumathi** 1 year ago

Selected Answer: D

D is the right answer, with private google access for on-premises host,

Private Google Access for on-premises hosts

On-premises hosts with or without external IP addresses. Connect to Google APIs and services, from your on-premises network, through a Cloud VPN tunnel or Cloud Interconnect by using one of the Private Google Access-specific domains and VIPs. The Google services that you can access depend on which Private Google Access-specific domain you use. Use this option to connect to Google APIs and services through a VPC network. This method doesn't require your on-premises hosts to have external IP addresses.

please refer to the link below for more insights,

<https://cloud.google.com/vpc/docs/private-google-access-hybrid>

upvoted 1 times

 **tomis2** 1 year, 3 months ago

Selected Answer: A

Through elimination - A

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right

upvoted 1 times

 **badrik** 1 year, 5 months ago

Selected Answer: A

AAAAAAAAAA

upvoted 1 times

 **HansKloss611** 1 year, 8 months ago

Selected Answer: A

A is correct

upvoted 3 times

 **Yaseed** 1 year, 10 months ago

PRIVATE INSTANCE WITH A PUBLIC LOADBALANCER WOULD GO PUBLIC! IT'S CONFUSING

upvoted 2 times

 **tvinay** 1 year, 10 months ago

Why are you shouting here?

upvoted 11 times

 **chikorita** 5 months, 3 weeks ago

and why is he bringing Load Balancer in middle of nowhere?

upvoted 3 times

 **lxs** 1 year, 11 months ago

The question is about ingress traffic from Internet

A - If the VM does not have public IP it is not routable from Internet. Correct answer

B - it is about how to access Google Services API. It does not tell about ingress Internet traffic

C - It is about egress traffic

D - It could be but we do not know anything about Internet ingress traffic to on prem. What's more default route tells about egress traffic to Internet. Nothing how Internet can access Compute instance.

Correct answer is A.

upvoted 5 times

 **arsh1916** 2 years, 4 months ago

A is correct

upvoted 4 times

Your team maintains the infrastructure for your organization. The current infrastructure requires changes. You need to share your proposed changes with the rest of the team. You want to follow Google's recommended best practices. What should you do?

- A. Use Deployment Manager templates to describe the proposed changes and store them in a Cloud Storage bucket.
- B. Use Deployment Manager templates to describe the proposed changes and store them in Cloud Source Repositories.
- C. Apply the changes in a development environment, run gcloud compute instances list, and then save the output in a shared Storage bucket.
- D. Apply the changes in a development environment, run gcloud compute instances list, and then save the output in Cloud Source Repositories.

Correct Answer: B*Community vote distribution*

B (67%)

A (33%)

ESP_SAP Highly Voted 3 years, 1 month ago

Correct Answer is (A):

Connecting to Cloud Storage buckets

Cloud Storage is a flexible, scalable, and durable storage option for your virtual machine instances. You can read and write files to Cloud Storage buckets from almost anywhere, so you can use buckets as common storage between your instances, App Engine, your on-premises systems, and other cloud services.

<https://cloud.google.com/compute/docs/disks/gcs-buckets>

Why not (B)?

Caution

Cloud Source Repositories are intended to store only the source code for your app and not user or personal data. Don't store any Core App Engine Customer Data (as defined in your License Agreement) in Cloud Source Repositories.

<https://cloud.google.com/source-repositories/docs/features>

upvoted 31 times

mav3r1ck 1 year, 2 months ago

In reality, will you proposed the answer in "A" to your boss? I bet he will show you the exit door. Storing IaC in Cloud Storage? It's like telling to store CloudFormation in S3 instead of CodeCommit.

upvoted 4 times

ESP_SAP 1 year ago

Correction. 2 years later.

Correct Answer is (B):

Using Recommendations for Infrastructure as Code:

This tutorial uses Hashicorp Terraform as the IaC tool, however the architectural patterns and components used in the described automation pipeline can be leveraged even if you are using a different IaC management tool such as Deployment Manager.

Cloud Build

Purpose: Google Cloud Build automates the deployment of infrastructure based on the changes made to the IaC manifests per policy intelligence recommendations.

Access control: The Cloud Build service account must have the right set of permissions to interact with resources in your test project.

See the documentation for configuring a Cloud Build service account.

GitHub

Purpose: The IaC repository uses GitHub for source control. The IaC repository in GitHub is integrated with Cloud Build. When commits are made to the master branch, a Cloud Build job is triggered to run a set of preconfigured tasks.

<https://cloud.google.com/recommender/docs/tutorial-iac#overview>

upvoted 23 times

bobthebuilder55110 1 year, 2 months ago

So, what do you think a deployment manager template will have? It will definitely have the source code for the source code/config for the template and if you are sharing with the team that is even better since that means that now people can collaborate that makes more sense to go with B over A here.

upvoted 1 times

gcpengineer 2 years, 2 months ago

B is the ans

upvoted 3 times

SSPC Highly Voted 3 years, 1 month ago

B is correct. <https://cloud.google.com/source-repositories/docs/features>

upvoted 27 times

 **AmitKM** 3 years, 1 month ago

Using Cloud Storage Repos, you can add comments and describe your changes to the team.Hence this might be a better option.

upvoted 3 times

 **magistrum** 2 years, 9 months ago

I don't see how you can do this when I tried creating:

Add code to your repository

info

Your repository is currently empty. Add some code using a selected method and then refresh your browser. Contents added to this repository can take some time to show up in search results. Learn more.

Select an option to push code to your repository:

Push code from a local Git repository

Clone your repository to a local Git repository

upvoted 2 times

 **Captain1212** Most Recent  1 month ago

Selected Answer: B

B is the correct answer

upvoted 1 times

 **pritampanda1988** 2 months ago

Selected Answer: A

Option B (storing in Cloud Source Repositories) might be suitable for storing application code, but it's not the best practice for storing infrastructure configuration templates.

upvoted 1 times

 **Praxii** 5 months, 1 week ago

Selected Answer: B

Correct answer is B

upvoted 1 times

 **SK1990** 9 months, 2 weeks ago

Selected Answer: B

B is the best answer.

upvoted 1 times

 **davidsalomon** 10 months ago

Selected Answer: B

As 2022 best practice, B

upvoted 1 times

 **Kopy** 10 months, 4 weeks ago

Selected Answer: B

Correction. 2 years later.

Correct Answer is (B):

upvoted 1 times

 **diasporabro** 11 months, 3 weeks ago

Selected Answer: B

Thought it was A for a sec, the realized Cloud Source is similar to Github/CodeCommit... So, B is a better choice

upvoted 1 times

 **AwesomeGCP** 12 months ago

B. Use Deployment Manager templates to describe the proposed changes and store them in Cloud Source Repositories

upvoted 1 times

 **KapilDhamija** 1 year, 1 month ago

Selected Answer: A

A is the right choice, Use Deployment Manager templates to describe the proposed changes and store them in a Cloud Storage bucket

upvoted 1 times

 **tomis2** 1 year, 3 months ago

Selected Answer: B

Don't get confused - it is B. IaC is a matter of version control system like GIT.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is right .. Showing Deployment Manager templates to your team will allow you to define the changes you want to implement in your cloud infrastructure. You can use Cloud Source Repositories to store Deployment Manager templates and collaborate with your team. Cloud Source Repositories are fully-featured, scalable, and private Git repositories you can use to store, manage and track changes to your code.

Hence, the correct answer is: Create Deployment Manager templates to define the proposed changes and save them into Cloud Source Repositories.

upvoted 3 times

 **SergijF** 1 year, 4 months ago

Selected Answer: B

Its B.

upvoted 1 times

 **Rutu_98** 1 year, 4 months ago

Selected Answer: A

A is correct Answer

Connecting to Cloud Storage buckets

Cloud Storage is a flexible, scalable, and durable storage option for your virtual machine instances. You can read and write files to Cloud Storage buckets from almost anywhere, so you can use buckets as common storage between your instances, App Engine, your on-premises systems, and other cloud services.

upvoted 1 times

 **nhadi82** 1 year, 4 months ago

Selected Answer: A

I would go for A as following

[#hosting_external_templates](https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally?hl=en)

upvoted 1 times

 **vruizrob** 1 year, 7 months ago

Correct answer: A. <https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally?hl=en>. You can host external templates in Google Cloud Storage or in a publicly-accessible location, such as GitHub.

upvoted 2 times

You have a Compute Engine instance hosting an application used between 9 AM and 6 PM on weekdays. You want to back up this instance daily for disaster recovery purposes. You want to keep the backups for 30 days. You want the Google-recommended solution with the least management overhead and the least number of services. What should you do?

- A. 1. Update your instances' metadata to add the following value: snapshot "schedule: 0 1 * * * 2. Update your instances' metadata to add the following value: snapshot "retention: 30
- B. 1. In the Cloud Console, go to the Compute Engine Disks page and select your instance's disk. 2. In the Snapshot Schedule section, select Create Schedule and configure the following parameters: - Schedule frequency: Daily - Start time: 1:00 AM 2:00 AM - Autodelete snapshots after: 30 days
- C. 1. Create a Cloud Function that creates a snapshot of your instance's disk. 2. Create a Cloud Function that deletes snapshots that are older than 30 days. 3. Use Cloud Scheduler to trigger both Cloud Functions daily at 1:00 AM.
- D. 1. Create a bash script in the instance that copies the content of the disk to Cloud Storage. 2. Create a bash script in the instance that deletes data older than 30 days in the backup Cloud Storage bucket. 3. Configure the instance's crontab to execute these scripts daily at 1:00 AM.

Correct Answer: B

Community vote distribution

B (100%)

 **ESP_SAP** Highly Voted 3 years, 1 month ago

Correct Answer is (B):

Creating scheduled snapshots for persistent disk

This document describes how to create a snapshot schedule to regularly and automatically back up your zonal and regional persistent disks. Use snapshot schedules as a best practice to back up your Compute Engine workloads. After creating a snapshot schedule, you can apply it to one or more persistent disks.

<https://cloud.google.com/compute/docs/disks/scheduled-snapshots>
upvoted 41 times

 **[Removed]** 3 years, 1 month ago

Definitely B.

With something like this, you should not have to write any custom scripts, custom functions, or cron jobs. This is google's way of saying 'hey, we've already built that stuff in to our snapshot schedules feature.'

upvoted 9 times

 **Ridhanya** Highly Voted 1 year, 10 months ago

it is b. we cannot define snapshot config in instance metadata.

VM instance metadata is used only for:

startup and shutdown scripts

host maintenance

guest attributes

upvoted 5 times

 **Captain1212** Most Recent 1 month ago

Selected Answer: B

B is the correct answer

upvoted 1 times

 **Rajat2309sharma** 8 months, 1 week ago

Selected Answer: B

B is ans

upvoted 1 times

 **s1cvlctetri** 9 months, 2 weeks ago

Selected Answer: B

got this question 2 days ago. B is right.

upvoted 3 times

 **AzureDP900** 1 year, 3 months ago

B is more appropriate

upvoted 1 times

 **nhadi82** 1 year, 4 months ago

Selected Answer: B

Correct Answer B

upvoted 1 times

✉ **gcpengineer** 2 years, 2 months ago

why not C?

upvoted 2 times

✉ **Gianfry** 1 year, 10 months ago

The question calls for "Google-recommended solution with the least management overhead and the least number of services"

upvoted 2 times

✉ **arsh1916** 2 years, 4 months ago

B is correct

upvoted 2 times

✉ **Hi2ALL** 2 years, 6 months ago

B is more realistic approach

upvoted 3 times

✉ **GCP_user** 2 years, 6 months ago

B is the best option so far. However just wonder this: Schedule frequency: Daily " Start time: 1:00 AM " 2:00 AM" Autodelete snapshots: after 30 days; For Saturday and Sunday it will be a waste of resource to create snapshots since the instance is running during weekdays.

upvoted 2 times

✉ **GCP_Student1** 2 years, 7 months ago

B. 1. In the Cloud Console, go to the Compute Engine Disks page and select your instance's disk. 2. In the Snapshot Schedule section, select Create Schedule and configure the following parameters: "" Schedule frequency: Daily "" Start time: 1:00 AM "" 2:00 AM "" Autodelete snapshots after 30 days

upvoted 2 times

✉ **DucSiu** 2 years, 8 months ago

It's B

upvoted 1 times

✉ **Bhagirathi** 2 years, 10 months ago

B for sure, any doubt?

upvoted 1 times

✉ **swatititame** 2 years, 10 months ago

• B. 1. In the Cloud Console, go to the Compute Engine Disks page and select your instance's disk. 2. In the Snapshot Schedule section, select Create Schedule and configure the following parameters: "" Schedule frequency: Daily "" Start time: 1:00 AM "" 2:00 AM "" Autodelete snapshots after 30 days

upvoted 2 times

✉ **hiteshrup** 3 years ago

Has to B...

- scripting which means overhead and maintenance So Option D neglected.

- Cloud function, adding another service which not require.

- MetaData, I am not sure if meta data have something to define as cron job schedule. So not be an option.

upvoted 2 times

✉ **stepkurniawan** 3 years, 1 month ago

Either A or B, I am not sure.

I have tested B and it works, but I dont know the name of the key to do A.

upvoted 1 times

Your existing application running in Google Kubernetes Engine (GKE) consists of multiple pods running on four GKE n1``standard``2 nodes. You need to deploy additional pods requiring n2``highmem``16 nodes without any downtime. What should you do?

- A. Use gcloud container clusters upgrade. Deploy the new services.
- B. Create a new Node Pool and specify machine type n2``highmem``16. Deploy the new pods.
- C. Create a new cluster with n2``highmem``16 nodes. Redeploy the pods and delete the old cluster.
- D. Create a new cluster with both n1``standard``2 and n2``highmem``16 nodes. Redeploy the pods and delete the old cluster.

Correct Answer: B

Community vote distribution



✉️ **GCP_Student1** Highly Voted 2 years, 6 months ago

B is correct answer, read below form google docs;

This tutorial demonstrates how to migrate workloads running on a Google Kubernetes Engine (GKE) cluster to a new set of nodes within the same cluster without incurring downtime for your application. Such a migration can be useful if you want to migrate your workloads to nodes with a different machine type.

Background

A node pool is a subset of machines that all have the same configuration, including machine type (CPU and memory) authorization scopes. Node pools represent a subset of nodes within a cluster; a container cluster can contain one or more node pools.

When you need to change the machine profile of your Compute Engine cluster, you can create a new node pool and then migrate your workloads over to the new node pool.

To migrate your workloads without incurring downtime, you need to:

Mark the existing node pool as unschedulable.

Drain the workloads running on the existing node pool.

Delete the existing node pool.

https://cloud.google.com/kubernetes-engine/docs/tutorials/migrating-node-pool#creating_a_node_pool_with_large_machine_type
upvoted 30 times

✉️ **Captain1212** Most Recent 1 month ago

Selected Answer: B

B is the right answer , if you need the new, and if you want to old one also then its D

upvoted 1 times

✉️ **ashtonez** 7 months ago

Selected Answer: B

B is correct, creating another cluster just doesnt make any sense, node pools are intended for this situations

upvoted 1 times

✉️ **Chiunara** 7 months ago

Selected Answer: B

Answer is obviously B (read @GCP_Student1 and @Bobbybash replies)

upvoted 1 times

✉️ **Bobbybash** 7 months, 3 weeks ago

Selected Answer: B

B. Create a new Node Pool and specify machine type n2``highmem``16. Deploy the new pods.

Creating a new Node Pool with the required machine type is the correct approach to deploy additional pods without any downtime. This approach allows you to scale the cluster horizontally by adding more nodes to the existing cluster. By creating a new Node Pool, you can add n2``highmem``16 nodes to the existing cluster, and deploy new pods on these nodes without affecting the existing services running on the n1``standard``2 nodes. This way, you can ensure high availability and zero downtime during the deployment. Option A (gcloud container clusters upgrade) upgrades the entire cluster, and Option C and D (creating a new cluster) involve deleting the existing cluster, which may cause downtime.

upvoted 2 times

✉️ **BlueJay20** 8 months, 1 week ago

Selected Answer: D

The keyword is "additional". Answer B is good if you want to replace with the new VMs. In this case you want the existing ones as well as the new ones. Therefore D.

upvoted 1 times

✉️ **swa99** 8 months ago

The keyword is "additional", in option D you are deleting the old cluster. SO the answer is B
upvoted 2 times

 **AzureDP900** 1 year, 3 months ago

B makes perfect sense.

upvoted 1 times

 **Tirthankar17** 1 year, 4 months ago

Selected Answer: B

B is correct

upvoted 1 times

 **arsh1916** 2 years, 4 months ago

B is correct

upvoted 4 times

 **Jacky_YO** 2 years, 6 months ago

ANS : B

1. The title did not say to delete four GKE n1.

upvoted 4 times

 **pondai** 2 years, 6 months ago

B, You need to create new node pool for cluster

upvoted 3 times

 **dunhill** 2 years, 7 months ago

I guess it's B. I couldn't find resize parameter under cluster upgrade. C and D are incorrect because it's no need to create new cluster.

upvoted 2 times

 **GCP_Student1** 2 years, 7 months ago

A. Use gcloud container clusters upgrade. Deploy the new services.

upvoted 1 times

 **GCP_Student1** 2 years, 6 months ago

I take it back, the correct answer is "B"

B. Create a new Node Pool and specify machine type n2-highmem16. Deploy the new pods.

upvoted 2 times

 **Jamaal_a** 2 years, 7 months ago

Answer is B - When you need to change the machine profile of your Compute Engine cluster, you can create a new node pool and then migrate your workloads over to the new node pool.

upvoted 3 times

You have an application that uses Cloud Spanner as a database backend to keep current state information about users. Cloud Bigtable logs all events triggered by users. You export Cloud Spanner data to Cloud Storage during daily backups. One of your analysts asks you to join data from Cloud Spanner and Cloud Bigtable for specific users. You want to complete this ad hoc request as efficiently as possible. What should you do?

- A. Create a dataflow job that copies data from Cloud Bigtable and Cloud Storage for specific users.
- B. Create a dataflow job that copies data from Cloud Bigtable and Cloud Spanner for specific users.
- C. Create a Cloud Dataproc cluster that runs a Spark job to extract data from Cloud Bigtable and Cloud Storage for specific users.
- D. Create two separate BigQuery external tables on Cloud Storage and Cloud Bigtable. Use the BigQuery console to join these tables through user fields, and apply appropriate filters.

Correct Answer: B

Community vote distribution

D (67%)

B (33%)

✉  **AmitKM**  3 years, 1 month ago

I think it should be D. <https://cloud.google.com/bigquery/external-data-sources>
upvoted 36 times

✉  **SSPC** 3 years, 1 month ago

The question says: "Join data from Cloud Spanner and Cloud Bigtable for specific users" You can see the Google documentation in the link <https://cloud.google.com/spanner/docs/export>
upvoted 3 times

✉  **Eshkrkrkr** 2 years, 11 months ago

Oh my god, SSPC read you your links!
The process uses Dataflow and exports data to a folder in a Cloud Storage bucket. The resulting folder contains a set of Avro files and JSON manifest files. And what next? I will tell - next you read below: Compute Engine: Before running your export job, you must set up initial quotas for Recommended starting values are:

CPUs: 200

In-use IP addresses: 200

Standard persistent disk: 50 TB

Still think its A?

upvoted 4 times

✉  **punjabishiva123** 2 years, 10 months ago

Hi,

Could u pls mail me all the correct answer to shivaasingh1104@gmail.com,

Really appreciated

Anyone pls help

upvoted 1 times

✉  **ash500** 2 years, 9 months ago

can you pls also share the correct answers with me

upvoted 1 times

✉  **kp0916** 2 years, 7 months ago

Does anyone have list of correct answers, let me know

upvoted 1 times

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (D):

Introduction to external data sources

This page provides an overview of querying data stored outside of BigQuery.

<https://cloud.google.com/bigquery/external-data-sources>

upvoted 27 times

✉  **ESP_SAP** 3 years, 1 month ago

BigQuery offers support for querying data directly from:

Bigtable
Cloud Storage
Google Drive
Cloud SQL (beta)

upvoted 5 times

✉  **ryzior** 1 year, 7 months ago

update:
BigQuery supports the following external data sources:
Bigtable
Cloud Spanner
Cloud SQL
Cloud Storage
Drive
upvoted 6 times

✉  **djgodzilla** 2 years, 3 months ago

but here we're not talking about joining Cloud Storage and Cloud Bigtable external tables.
the join happens between a distributed relational database (Spanner) and key-value NoSQL Database (BigTable). how's converting Spanner to cloud storage an implicit and trivial step.
upvoted 1 times

✉  **djgodzilla** 2 years, 3 months ago

"The Cloud Spanner to Cloud Storage Text template is a batch pipeline that reads in data from a Cloud Spanner table, optionally transforms the data via a JavaScript User Defined Function (UDF) that you provide, and writes it to Cloud Storage as CSV text files."
<https://cloud.google.com/dataflow/docs/guides/templates/provided-batch#cloudspannertogcstext>

"The Dataflow connector for Cloud Spanner lets you read data from and write data to Cloud Spanner in a Dataflow pipeline"
<https://cloud.google.com/spanner/docs/dataflow-connector>
upvoted 3 times

✉  **[Removed]** 3 years ago

As per your comment D is the answer.
I also agree.
But can BigQurey read backed up data? , as we have backup data on Cloud storage, did not get any evidence in the link you shared.
upvoted 2 times

✉  **KC_go_reply** Most Recent 4 months ago

Selected Answer: D

BigQuery is powerful. If you have data in one of the popular sources like Cloud Storage or Bigtable, it is much more efficient - both for cost and computation - to create an external table on those data sources, than to copy their data around.

Besides that, also keep in mind that table clones and snapshots are much more efficient than full table copy etc.
upvoted 2 times

✉  **Praxii** 5 months, 1 week ago

Selected Answer: B

I go for option B. As in option D, the data is backed up data and not the most recent data.
upvoted 1 times

✉  **Bobbybash** 7 months, 3 weeks ago

Selected Answer: B

B. Create a dataflow job that copies data from Cloud Bigtable and Cloud Spanner for specific users.

To join data from Cloud Spanner and Cloud Bigtable for specific users, creating a dataflow job that copies data from both sources is the most efficient option. This approach allows you to process the data in parallel, and you can take advantage of Cloud Dataflow's autoscaling feature to handle large volumes of data. You can use Cloud Dataflow to read data from Cloud Bigtable and Cloud Spanner, join the data based on the user fields, and write the output to a new location or send it to the analyst. Option A (copying data from Cloud Storage) does not provide data from Cloud Spanner, and option C (running a Spark job on a Dataproc cluster) involves higher overhead costs. Option D (using BigQuery external tables) is not efficient for ad hoc requests, as data is exported from Spanner to Cloud Storage during backups, so there may be a delay in data availability.

upvoted 1 times

✉  **anolive** 11 months, 1 week ago

Selected Answer: D

I thinks is D, but not 100% sure, because D does not have any infomation about the specific user like others options.
upvoted 1 times

✉  **Charumathi** 1 year ago

Selected Answer: D

D is the correct answer,
An external data source is a data source that you can query directly from BigQuery, even though the data is not stored in BigQuery storage.

BigQuery supports the following external data sources:

Amazon S3
Azure Storage
Cloud Bigtable
Cloud Spanner
Cloud SQL
Cloud Storage
Drive

upvoted 1 times

✉ **DualCore573** 1 year ago

Selected Answer: D

D makes sense as the BigQuery external tables are made for such use cases. and "efficient" keyword makes sense to use this way as resources used are less.

upvoted 1 times

✉ **soaresleo** 1 year, 2 months ago

First of all, using Dataflow can perhaps be effective, but NOT efficient, specially because of costs.

Second:

"To query Cloud Bigtable data using a permanent external table, you: Create a table definition file (for the API or bq command-line tool); Create a table in BigQuery linked to the external data source; Query the data using the permanent table."

Source: <https://cloud.google.com/bigquery/docs/external-data-bigtable#:~:text=To%20query%20Cloud%20Bigtable%20data,data%20using%20the%20permanent%20table>

Third:

"To query a Cloud Storage external data source, provide the Cloud Storage URI path to your data and create a table that references the data source."

Source:

<https://cloud.google.com/bigquery/docs/external-data-cloud-storage>

Correct answer: D.

upvoted 1 times

✉ **jeffangel128** 1 year, 2 months ago

Selected Answer: D

"efficiently as possible" -> use the least amount of resources and achieve the same result... so I think it's D

upvoted 1 times

✉ **tomis2** 1 year, 3 months ago

Selected Answer: D

Most "cloud" solution is D

upvoted 1 times

✉ **sabbella** 1 year, 6 months ago

Selected Answer: D

option d

upvoted 1 times

✉ **sabbella** 1 year, 6 months ago

Selected Answer: D

Option is D

upvoted 1 times

✉ **rlijhk** 1 year, 8 months ago

Selected Answer: B

I think it is B.

The data in Cloud storage is not up to date as backup window is daily. SO, there are chances is missing one day worth of data. As it is mentioned as "efficiently" instead of quickly, I would choose "B".

upvoted 3 times

✉ **obeythefist** 1 year, 7 months ago

How does this create a "join" between the two tables?

upvoted 1 times

✉ **BigQuery** 1 year, 6 months ago

why do you think one cannot join 2 subsets of data in dataflow Its meant for processing sets of data.

upvoted 1 times

✉ **sasithra** 1 year, 8 months ago

B is correct

upvoted 1 times

✉ **raaj_p** 1 year, 8 months ago

Selected Answer: D

An external data source is a data source that you can query directly from BigQuery, even though the data is not stored in BigQuery storage.

BigQuery supports the following external data sources:

Bigtable

Cloud Spanner

Cloud SQL

Cloud Storage
Drive
upvoted 1 times

✉ **Ridhanya** 1 year, 10 months ago

I think it has to be option A because Dataflow needs to be used for bigtable export and cloud spanner data is already backed up in cloud storage on a daily basis
upvoted 1 times

✉ **Ridhanya** 1 year, 10 months ago

Um, the questions says "join" so D seems right. I take back my previous
upvoted 2 times

You are hosting an application from Compute Engine virtual machines (VMs) in us`central1`a. You want to adjust your design to support the failure of a single Compute Engine zone, eliminate downtime, and minimize cost. What should you do?

- A. Create Compute Engine resources in us`central1`b. Balance the load across both us`central1`a and us`central1`b.
- B. Create a Managed Instance Group and specify us`central1`a as the zone. Configure the Health Check with a short Health Interval.
- C. Create an HTTP(S) Load Balancer. Create one or more global forwarding rules to direct traffic to your VMs.
- D. Perform regular backups of your application. Create a Cloud Monitoring Alert and be notified if your application becomes unavailable. Restore from backups when notified.

Correct Answer: A

Community vote distribution

A (100%)

 **GCP_Student1** Highly Voted 2 years, 7 months ago

A. Create Compute Engine resources in us`central1`b. Balance the load across both us`central1`a and us`central1`b.
upvoted 14 times

 **obeythefist** Highly Voted 1 year, 7 months ago

This seems straightforward. "A" is the only answer that involves putting instances in more than one zone!

- A. Yes, creating instances in another zone and balancing the loads will fix this problem
 - B. Wrong. This keeps all the instances in one zone, but the question says we want to protect against zone failures.
 - C. Wrong. This keeps all the instances in one zone, but the question says we want to protect against zone failures.
 - D. Wrong. This keeps all the instances in one zone, but the question says we want to protect against zone failures.
- upvoted 10 times

 **Captain1212** Most Recent 1 month ago

Selected Answer: A

A seems more right as it help with the Zone failure, all other create the same in same zone
upvoted 1 times

 **DrLegendgun** 6 months, 1 week ago

Selected Answer: A

The Answer is B
upvoted 1 times

 **Angel_99** 1 year, 1 month ago

Selected Answer: A

A is best option
upvoted 1 times

 **abirroy** 1 year, 2 months ago

Selected Answer: A

A is the best option
upvoted 2 times

 **AzureDP900** 1 year, 3 months ago

A is fine.
upvoted 1 times

 **[Removed]** 1 year, 7 months ago

Option A.
Create VMs across more than one region and zone so that you have alternative VMs to point to if a zone or region containing one of your VMs is disrupted. If you host all your VMs in the same zone or region, you won't be able to access any of those VMs if that zone or region becomes unreachable.
<https://cloud.google.com/compute/docs/tutorials/robustsystems#distribute>
upvoted 3 times

 **Ridhanya** 1 year, 10 months ago

A is correct because we have to eliminate single zone failure problem
upvoted 1 times

 **Gianfry** 1 year, 10 months ago

Why not "B" selecting "Regional (multi zone)"?
"Regional (multiple zone) coverage. Regional MIGs let you spread app load across multiple zones. This replication protects against zonal

failures. If that happens, your app can continue serving traffic from instances running in the remaining available zones in the same region."

<https://cloud.google.com/compute/docs/instance-groups/>

upvoted 4 times

 **kimharsh** 1 year, 10 months ago

it should be B , but because it specify the one Zone we can't pick this answer , the closest other option is A

upvoted 1 times

 **arsh1916** 2 years, 4 months ago

A is best option

upvoted 6 times

 **mj98** 2 years, 6 months ago

Can someone explain how A?

upvoted 2 times

 **jabrrJ68w02ond1** 1 year, 10 months ago

Other options do not prepare you for zonal outages

upvoted 1 times

 **tifo16** 2 years, 5 months ago

in order to remediate to the problem of single point of failure, we have to replicate VMs within multiple zones. Only A choice consider this concern

upvoted 10 times

 **nzexamtopics** 2 years, 6 months ago

A? Really? how?

upvoted 1 times

 **NARWAL** 2 years, 6 months ago

A is correct.

upvoted 4 times

 **victory108** 2 years, 7 months ago

A - " Create Compute Engine resources in us-central1-b. " Balance the load across both us-central1-a and us-central1-b.

upvoted 2 times

A colleague handed over a Google Cloud Platform project for you to maintain. As part of a security checkup, you want to review who has been granted the Project Owner role. What should you do?

- A. In the console, validate which SSH keys have been stored as project-wide keys.
- B. Navigate to Identity-Aware Proxy and check the permissions for these resources.
- C. Enable Audit Logs on the IAM & admin page for all resources, and validate the results.
- D. Use the command gcloud projects get iam policy to view the current role assignments.

Correct Answer: A

Reference:

<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>*Community vote distribution*

D (100%)

ESP_SAP Highly Voted 3 years, 1 month ago

Correct Answer is (D):

A simple approach would be to use the command flags available when listing all the IAM policy for a given project. For instance, the following command:

```
'gcloud projects get-iam-policy $PROJECT_ID --flatten="bindings[].members" --format="table(bindings.members)" --filter="bindings.role:roles/owner"'
```

outputs all the users and service accounts associated with the role 'roles/owner' in the project in question.

<https://groups.google.com/g/google-cloud-dev/c/Z6sZs7TvygQ?pli=1>

upvoted 40 times

MohammedGhouse Highly Voted 3 years, 1 month ago

D: is the answer
upvoted 13 times

SSPC 3 years, 1 month ago

D is the correct.
upvoted 3 times

yurstev 2 years, 9 months ago

D IS THE ANSWER
upvoted 4 times

Captain1212 Most Recent 1 month ago

Selected Answer: D
D seems, more correct
upvoted 1 times

tomis2 1 year, 3 months ago

Selected Answer: D
gcloud iam get-iam-policy
upvoted 1 times

AzureDP900 1 year, 3 months ago

D is right
upvoted 1 times

Rutu_98 1 year, 4 months ago

Selected Answer: D
Answer is D
upvoted 2 times

somenick 1 year, 6 months ago

Selected Answer: D
gcloud projects get-iam-policy \$PROJECT_ID
upvoted 1 times

obeythefist 1 year, 7 months ago

I chose D by a process of elimination. Here's my take:

- A. There's more than one way to access an instance than just the SSH keys, and SSH keys have nothing to do with Project Owner role.
- B. Barking up the wrong tree here, Identity-Aware Proxy is more for remotely accessing resources, rather than Project Owner IAM roles.
- C. This will only work if everyone who is a Project Owner accesses the system so you can see them in the logs. What if a Project Owner doesn't access the Project for a while? How long will you wait? Nope.
- D. By elimination, this is the best result.

upvoted 8 times

 **BigQuery** 1 year, 6 months ago

NICE EXPLANATION; WAY TO GO D

upvoted 1 times

 **HansKloss611** 1 year, 8 months ago

Selected Answer: D

D is correct

upvoted 1 times

 **PR0704** 1 year, 10 months ago

how can the admin be so inconsistent throughout with the answers..not good ..its so confusing

upvoted 2 times

 **tvinay** 1 year, 10 months ago

Confusion!! that's the main goal here so that we all go to the docs and Study hard xD

upvoted 1 times

 **mohamedmahmoudf97** 2 years ago

D is the correct answer

upvoted 1 times

 **arsh1916** 2 years, 4 months ago

D is correct

upvoted 2 times

 **GCP_Student1** 2 years, 7 months ago

D. Use the command gcloud projects get""iam""policy to view the current role assignments.

upvoted 3 times

 **Bhagirathi** 2 years, 9 months ago

D 200%

upvoted 2 times

 **Bhagirathi** 2 years, 10 months ago

anyone will be confused - solution says one answer

same time, all you guys have different choices here. what to take from this ?

upvoted 1 times

 **swatititame** 2 years, 10 months ago

D. Use the command gcloud projects get""iam""policy to view the current role assignments.

upvoted 1 times

 **adeyemi5700** 2 years, 12 months ago

C is the likely answer. With D you see the current users with permission. With audit log you see those with prev. and current permission.

upvoted 1 times

You are running multiple VPC-native Google Kubernetes Engine clusters in the same subnet. The IPs available for the nodes are exhausted, and you want to ensure that the clusters can grow in nodes when needed. What should you do?

- A. Create a new subnet in the same region as the subnet being used.
- B. Add an alias IP range to the subnet used by the GKE clusters.
- C. Create a new VPC, and set up VPC peering with the existing VPC.
- D. Expand the CIDR range of the relevant subnet for the cluster.

Correct Answer: D*Community vote distribution*

D (100%)

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (D):

gcloud compute networks subnets expand-ip-range

NAME

gcloud compute networks subnets expand-ip-range - expand the IP range of a Compute Engine subnetwork

<https://cloud.google.com/sdk/gcloud/reference/compute/networks/subnets/expand-ip-range>

upvoted 27 times

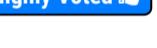
✉  **magistrum** 2 years, 9 months ago

Ok D it is, here's the GKE specific documentation

<https://cloud.google.com/kubernetes-engine/docs/concepts/alias-ips>

Every subnet must have a primary IP address range. You can expand the primary IP address range at any time, even when Google Cloud resources use the subnet; however, you cannot shrink or change a subnet's primary IP address scheme after the subnet has been created. The first two and last two IP addresses of a primary IP address range are reserved by Google Cloud.

upvoted 7 times

✉  **MohammedGhouse**  3 years, 1 month ago

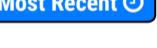
D: is the answer

upvoted 12 times

✉  **SSPC** 3 years, 1 month ago

I agree with you. https://cloud.google.com/vpc/docs/configure-alias-ip-ranges#gcloud_1

upvoted 2 times

✉  **Captain1212**  1 month ago

Selected Answer: D

D is the correct Answer, as you just expand the range

upvoted 1 times

✉  **Bobbybash** 7 months, 3 weeks ago

Selected Answer: D

D. Expand the CIDR range of the relevant subnet for the cluster.

Expanding the CIDR range of the relevant subnet for the cluster would increase the number of available IP addresses and allow the clusters to grow when needed. This can be done by modifying the existing subnet's IP address range in the VPC network settings. Adding a new subnet or VPC peering would not directly address the issue of running out of available IP addresses in the current subnet. Adding an alias IP range to the subnet could provide additional IP addresses, but may not be sufficient for long-term growth.

upvoted 2 times

✉  **AwesomeGCP** 12 months ago

Selected Answer: D

D. Expand the CIDR range of the relevant subnet for the cluster.

upvoted 1 times

✉  **learn_GCP** 1 year ago

Selected Answer: D

D. Expanding CIDR range is enough.

upvoted 1 times

✉  **sonuricky** 1 year, 2 months ago

C is the right answer

upvoted 1 times

✉  **ryumada** 1 year, 2 months ago

Please provide the reason why you choose C as the right answer. ESP_SAP explains clearly about the reason why he choose D as the right answer even he add Google Documentation link too to prove his answer.

upvoted 1 times

 **Bumbah** 1 year, 3 months ago

Selected Answer: D

Correct answer is D:

<https://cloud.google.com/vpc/docs/create-modify-vpc-networks#expand-subnet>

Just expand your subnet.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

D is right

upvoted 1 times

 **GCP_Student1** 2 years, 7 months ago

This might help

Node limiting ranges

The maximum number of Pods and Services for a given GKE cluster is limited by the size of the cluster's secondary ranges. The maximum number of nodes in the cluster is limited by the size of the cluster's subnet's primary IP address range and the cluster's Pod address range.

The Cloud Console shows error messages like the following to indicate that either the subnet's primary IP address range or the cluster's Pod IP address range (the subnet's secondary IP address range for Pods) has been exhausted:

Instance [node name] creation failed: IP space of [cluster subnet] is exhausted

Note: Secondary subnets are not visible in Cloud Console. If you can't find the [cluster subnet] reported by the above error message it means that the error is caused by IP exhaustion in a secondary subnet. In this case check the secondary ranges of the primary subnet.

https://cloud.google.com/kubernetes-engine/docs/concepts/alias-ips#node_limiters

upvoted 6 times

 **GCP_Student1** 2 years, 7 months ago

By the way the answer is;

D. Expand the CIDR range of the relevant subnet for the cluster.

upvoted 3 times

 **Ozymandiax** 2 years, 9 months ago

UHmmm, 1 question. The description of the problem says that the ip's are EXHAUSTED. So, no more IP's available in this subnet.

It also states that we're having a multi-VPC environment... as allways we should not interpret, just take the questions literally.

IF we do not know the actual size of the deployment it can be ANY size, and if IP's are EXHAUSTED, it should BE, BIG as Galactic sized or so....

With all this I wonder if the right answer is not C...

upvoted 3 times

 **akshaym87** 1 year, 8 months ago

Same doubt!

VPC peering seems correct to me.

upvoted 1 times

 **Bhagirathi** 2 years, 10 months ago

D best option to think here.

upvoted 2 times

 **hicham** 2 years, 7 months ago

totally agree

upvoted 1 times

 **swatititame** 2 years, 10 months ago

D. Expand the CIDR range of the relevant subnet for the cluster.

upvoted 1 times

You have a batch workload that runs every night and uses a large number of virtual machines (VMs). It is fault-tolerant and can tolerate some of the VMs being terminated. The current cost of VMs is too high. What should you do?

- A. Run a test using simulated maintenance events. If the test is successful, use preemptible N1 Standard VMs when running future jobs.
- B. Run a test using simulated maintenance events. If the test is successful, use N1 Standard VMs when running future jobs.
- C. Run a test using a managed instance group. If the test is successful, use N1 Standard VMs in the managed instance group when running future jobs.
- D. Run a test using N1 standard VMs instead of N2. If the test is successful, use N1 Standard VMs when running future jobs.

Correct Answer: B

Reference:

<https://cloud.google.com/compute/vm-instance-pricing>

Community vote distribution

A (100%)

  **ESP_SAP** Highly Voted 3 years, 1 month ago

Correct Answer is (A):

Creating and starting a preemptible VM instance

This page explains how to create and use a preemptible virtual machine (VM) instance. A preemptible instance is an instance you can create and run at a much lower price than normal instances. However, Compute Engine might terminate (preempt) these instances if it requires access to those resources for other tasks. Preemptible instances will always terminate after 24 hours. To learn more about preemptible instances, read the preemptible instances documentation.

Preemptible instances are recommended only for fault-tolerant applications that can withstand instance preemptions. Make sure your application can handle preemptions before you decide to create a preemptible instance. To understand the risks and value of preemptible instances, read the preemptible instances documentation.

<https://cloud.google.com/compute/docs/instances/create-start-preemptible-instance>

upvoted 40 times

  **MohammedGhouse** Highly Voted 3 years, 1 month ago

A: is the answer

upvoted 16 times

  **SSPC** 3 years, 1 month ago

"A" is correct

upvoted 3 times

  **juliandm** 3 years, 1 month ago

What about a mixture of preemptible N1 and normal N1 instances? i can't believe just having preemptible is a good practice

upvoted 1 times

  **Ale1973** 3 years, 1 month ago

Good point, in real-world your solution, is the best. For this scenario, the answer is A.

upvoted 8 times

  **akhun** Most Recent 8 months agoSelected Answer: A

It is specific on Batch workload , runs in less than 24 hrs, is fault tolerant. The best candidate for this is job is a preemptible VM
upvoted 2 times

  **Pr44** 10 months, 3 weeks agoSelected Answer: A

Preemptible to save cost and even it is fault tolerant
upvoted 1 times

  **Charumathi** 1 year agoSelected Answer: A

A is correct, preemptible VMs reduce cost, and this is recommended to run batch jobs which run less than 24 hours
upvoted 1 times

  **gcpreviewer** 1 year agoSelected Answer: A

I Vote A as it is clearly correct. Whenever something runs in under 24 hours and is fault tolerant we should be looking at preemptible VMs to save costs.
upvoted 1 times

 **ccieman2016** 1 year, 1 month ago

Selected Answer: A

A is correct

upvoted 1 times

 **gcpj** 1 year, 3 months ago

Selected Answer: A

Answer should be A: preemptible VM instances. Because the workload is fault-tolerant and can tolerate some of the VMs being terminated.

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right ..

upvoted 1 times

 **NBR1** 1 year, 4 months ago

Selected Answer: A

I believe it is A

upvoted 1 times

 **Majkl93** 1 year, 7 months ago

Selected Answer: A

as per the comments

upvoted 1 times

 **maggieli** 1 year, 11 months ago

I vote A. Preempt VM can costdown more.

upvoted 2 times

 **vamgcp** 2 years ago

A is correct because preemptible VMs can provide up to 80% discount over normal VMs if the workloads are fault-tolerant

upvoted 2 times

 **AD_0525** 2 years, 3 months ago

Keyword- Fault tolerant, so answer should be pre emptible VMs, option A

upvoted 2 times

 **GCP_Student1** 2 years, 7 months ago

A. Run a test using simulated maintenance events. If the test is successful, use preemptible N1 Standard VMs when running future jobs

upvoted 3 times

 **swatititame** 2 years, 10 months ago

• A. Run a test using simulated maintenance events. If the test is successful, use preemptible N1 Standard VMs when running future jobs.

upvoted 2 times

You are working with a user to set up an application in a new VPC behind a firewall. The user is concerned about data egress. You want to configure the fewest open egress ports. What should you do?

- A. Set up a low-priority (65534) rule that blocks all egress and a high-priority rule (1000) that allows only the appropriate ports.
- B. Set up a high-priority (1000) rule that pairs both ingress and egress ports.
- C. Set up a high-priority (1000) rule that blocks all egress and a low-priority (65534) rule that allows only the appropriate ports.
- D. Set up a high-priority (1000) rule to allow the appropriate ports.

Correct Answer: C

Community vote distribution

A (87%) 13%

 **ESP_SAP**  3 years, 1 month ago

Correct Answer is (A):

Implied rules

Every VPC network has two implied firewall rules. These rules exist, but are not shown in the Cloud Console:

Implied allow egress rule. An egress rule whose action is allow, destination is 0.0.0.0/0, and priority is the lowest possible (65535) lets any instance send traffic to any destination, except for traffic blocked by Google Cloud. A higher priority firewall rule may restrict outbound access. Internet access is allowed if no other firewall rules deny outbound traffic and if the instance has an external IP address or uses a Cloud NAT instance. For more information, see Internet access requirements.

Implied deny ingress rule. An ingress rule whose action is deny, source is 0.0.0.0/0, and priority is the lowest possible (65535) protects all instances by blocking incoming connections to them. A higher priority rule might allow incoming access. The default network includes some additional rules that override this one, allowing certain types of incoming connections.

https://cloud.google.com/vpc/docs/firewalls#default_firewall_rules

upvoted 39 times

 **patashish** 1 year, 3 months ago

The correct answer is C

upvoted 1 times

 **ryumada** 1 year, 2 months ago

You should visit the documentation link he attached. He's copy those statements from the Google Docs.

upvoted 1 times

 **Roro_Brother** 1 year, 3 months ago

Listen that guy because he is right

upvoted 1 times

 **bobthebuilder55110**  1 year, 2 months ago

Selected Answer: A

Answer is (A) :

First I was going with C but then I read the question again, let's try to understand both options here, the goal is to deny egress and only allow some ports for some functions to perform. If we go with C, lower the number higher the priority (1000) so the rule with this priority 1000 will overwrite (65534), so If we allow only appropriate ports it will be overwritten with the high-priority (1000) rule and all the egress traffic will be blocked.

Remember the goal here is to block egress but not all of it since we still want to configure the fewest open ports and this is statefull meaning for open ports traffic will be both ways.

A fits this condition where it is saying we block all traffic but the required ports are kept open with higher priority which will only allow the required traffic to leave the network.

upvoted 11 times

 **scanner2**  1 month ago

Selected Answer: A

The rule is evaluated on higher priority to lower priority and depends first come first serve basis.

<https://cloud.google.com/firewall/docs/firewall-policies-overview#rule-evaluation>

upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: A

A is the correct answer

upvoted 1 times

 **fragment137** 10 months, 1 week ago

Selected Answer: A

Correct answer is A.

Answer will not be D, because Egress traffic is Allowed by default. You will have to explicitly set the rule blocking outbound traffic.
upvoted 1 times

 **ryumada** 1 year, 2 months ago

Selected Answer: A
Read ESP_SAP comment for the explanation. He explains it clearly.
upvoted 1 times

 **sonuricky** 1 year, 2 months ago

C is the correct answer
upvoted 1 times

 **gscharly** 1 year, 2 months ago

Selected Answer: A
A: is the answer
upvoted 1 times

 **Roro_Brother** 1 year, 3 months ago

Selected Answer: A
Correct answer is A
upvoted 1 times

 **patashish** 1 year, 3 months ago

Correct Answer is C
upvoted 1 times

 **patashish** 1 year, 3 months ago

Hint : All rules are stateful.
VPC firewall rules are stateful. When a connection is allowed through the firewall in either direction, return traffic matching this connection is also allowed. You cannot configure a firewall rule to deny associated response traffic.

As per question , we want to restrict egress traffic.
So focus to restrict egress traffic based on priority of rules.
Allow incoming traffic for appropriate traffic and block all traffic and allow only which are required.

Hence , as per my view C should be correct answer

upvoted 3 times

 **mani098** 1 year, 3 months ago

Selected Answer: D
A incorrect 65534 that blocks all ingress, not egress (except few default ports)
D is correct.
upvoted 3 times

 **patashish** 1 year, 3 months ago

But why D is correct ? Why not C ?

D is more generic , As per question , need to focus on egress traffic
upvoted 1 times

 **pnVino27** 1 year, 10 months ago

Selected Answer: A
Correct Answer is A
upvoted 3 times

 **maggieli** 1 year, 11 months ago

I vote A is correct. Block all port in gress and set low-priority.
upvoted 2 times

 **aamirahal** 1 year, 12 months ago

A is correct
upvoted 2 times

 **vvkds** 2 years, 1 month ago

Correct answer is A. Firewall rules are executed based on the priority.
upvoted 2 times

 **tanito83** 2 years, 4 months ago

The correct answer is A. Please, modify it.
upvoted 3 times

 **kopper2019** 2 years, 5 months ago

A, lower is high priority
Implied rules
Every VPC network has two implied firewall rules. These rules exist, but are not shown in the Cloud Console:

Implied allow egress rule. An egress rule whose action is allow, destination is 0.0.0.0/0, and priority is the lowest possible (65535) lets any

instance send traffic to any destination, except for traffic blocked by Google Cloud. A higher priority firewall rule may restrict outbound access. Internet access is allowed if no other firewall rules deny outbound traffic and if the instance has an external IP address or uses a Cloud NAT instance. For more information, see [Internet access requirements](#).

Implied deny ingress rule. An ingress rule whose action is deny, source is 0.0.0.0/0, and priority is the lowest possible (65535) protects all instances by blocking incoming connections to them. A higher priority rule might allow incoming access. The default network includes some additional rules that override this one, allowing certain types of incoming connections.

upvoted 2 times

Your company runs its Linux workloads on Compute Engine instances. Your company will be working with a new operations partner that does not use Google Accounts. You need to grant access to the instances to your operations partner so they can maintain the installed tooling. What should you do?

- A. Enable Cloud IAP for the Compute Engine instances, and add the operations partner as a Cloud IAP Tunnel User.
- B. Tag all the instances with the same network tag. Create a firewall rule in the VPC to grant TCP access on port 22 for traffic from the operations partner to instances with the network tag.
- C. Set up Cloud VPN between your Google Cloud VPC and the internal network of the operations partner.
- D. Ask the operations partner to generate SSH key pairs, and add the public keys to the VM instances.

Correct Answer: B

Reference:

<https://cloud.google.com/vpc/docs/firewalls>

Community vote distribution

A (95%) 5%

- ✉  **kulikBro**  2 years, 6 months ago
A - <https://cloud.google.com/iap/docs/external-identities>
upvoted 29 times
- ✉  **Bhagirathi**  2 years, 10 months ago
full of confusions for any reader....
You guys all say A, B, C & D but which one is correct ?
upvoted 22 times
- ✉  **yc25744** 2 years, 2 months ago
nothing
upvoted 6 times
- ✉  **Captain1212**  1 month ago
Selected Answer: A
A seems more correct, as to provide the access
upvoted 1 times
- ✉  **Praxii** 5 months, 1 week ago
Selected Answer: A
Answer is A. Although to enable IAP, you do need to create a firewall rule on tcp 22. But if this question wasn't multiple choice then A is correct.
"IAP is a building block toward BeyondCorp, an enterprise security model that enables employees to work from untrusted networks without using a VPN." - So C is not required when A can suffice.
upvoted 1 times
- ✉  **innoculous_chris** 5 months, 2 weeks ago
Selected Answer: C
<https://cloud.google.com/iap/docs/concepts-overview>
IAP is only for google accounts and applies to access to AppEngine, HTTP(s) LB. It explicitly doesn't protect VMs.
upvoted 1 times
- ✉  **innoculous_chris** 5 months, 2 weeks ago
please ignore..Answer should be A..<https://cloud.google.com/iap/docs/external-identities>..this page shows it works for VMs and non-google accounts.
upvoted 2 times
- ✉  **krishna37** 9 months, 3 weeks ago
Selected Answer: A
Please watch this video.
<https://www.youtube.com/watch?v=jZdXyWQuIW0>
upvoted 4 times
- ✉  **[Removed]** 10 months ago
B is the straight forward answer to allow the partner to access via SSH without a Google account. For those suggesting A, carefully read <https://cloud.google.com/iap/docs/external-identities> and you'll notice that external identity isn't available from IAP out of the box and requires Identity Platform.
upvoted 2 times
- ✉  **Pr44** 10 months, 3 weeks ago

Selected Answer: A

Question ask about granting access to new operations partner and that can be done by first option only.

upvoted 1 times

✉ **Charumathi** 1 year ago

Selected Answer: A

A is the correct answer,

IAP controls access to your App Engine apps and Compute Engine VMs running on Google Cloud. It leverages user identity and the context of a request to determine if a user should be allowed access. IAP is a building block toward BeyondCorp, an enterprise security model that enables employees to work from untrusted networks without using a VPN.

By default, IAP uses Google identities and IAM. By leveraging Identity Platform instead, you can authenticate users with a wide range of external identity providers, such as:

Email/password

OAuth (Google, Facebook, Twitter, GitHub, Microsoft, etc.)

SAML

OIDC

Phone number

Custom

Anonymous

This is useful if your application is already using an external authentication system, and migrating your users to Google accounts is impractical.

upvoted 6 times

✉ **AzureDP900** 1 year, 3 months ago

A is right. IAP will allow you to connect compute engine without GCP account .

upvoted 3 times

✉ **pfabio** 1 year, 4 months ago

Selected Answer: A

A - To control which users and groups are allowed to use IAP TCP forwarding and which VM instances they're allowed to connect to, configure Identity and Access Management (IAM) permissions.

How to:

Open the IAP admin page and select the SSH and TCP Resources tab.

Open the IAP admin page

Select the VM instances that you want to configure.

Click Show info panel if the info panel is not visible.

Click Add member and configure the following:

New members: Specify the user or group you want to grant access.

Select a role Select Cloud IAP > IAP-Secured Tunnel User.

<https://cloud.google.com/iap/docs/using-tcp-forwarding#grant-permission>

upvoted 3 times

✉ **SleepyHitman** 1 year, 8 months ago

Selected Answer: A

The answer is A:

Enable Cloud IAP for the Compute Engine instances, and add the operations partner as a Cloud IAP Tunnel User as per [1] and [2].

[1]:<https://cloud.google.com/iap/docs/tcp-forwarding-overview#:~:text=To%20learn%20how%20to%20grant%20principals%20access%20to%20tunneled%20resources%20and%20how%20to%20create%20tunnels%20that%20route%20TCP%20traffic%2C%20see%20Using%20IAP%20for%20TCP%20forwarding>.

[2]:<https://cloud.google.com/iap/docs/tcp-forwarding-overview#:~:text=IAP%27s%20TCP%20forwarding%20feature%20lets%20you%20control%20who%20can%20access%20administrative%20services%20like%20SSH%20and%20RDP%20on%20your%20backends%20from%20the%20public%20internet>.

upvoted 4 times

✉ **emv** 1 year, 8 months ago

IAP controls access to your App Engine apps and Compute Engine VMs running on Google Cloud. It leverages user identity and the context of a request to determine if a user should be allowed access. IAP is a building block toward BeyondCorp, an enterprise security model that enables employees to work from untrusted networks without using a VPN.

By default, IAP uses Google identities and IAM. By leveraging Identity Platform instead, you can authenticate users with a wide range of external identity providers, such as:

Email/password

OAuth (Google, Facebook, Twitter, GitHub, Microsoft, etc.)

SAML

OIDC

Phone number

Custom

Anonymous

This is useful if your application is already using an external authentication system, and migrating your users to Google accounts is impractical.

upvoted 4 times

✉ **Surat** 1 year, 9 months ago

Very confusing quesiton, whats final answer?

upvoted 1 times

 **kimharsh** 1 year, 9 months ago

People THNK SIMPLE

the exam want to test you if you know about the IAP or not , the IAP allow you to access compute engine from the internet without having to have a GCP account ,

the answer is A

upvoted 12 times

 **MariGK** 1 year, 10 months ago

both (CD) combined seems to be the right answer

upvoted 1 times

 **wh1t4k3r** 1 year, 10 months ago

B and C would make the VMs reachable, but it does not mention user access/auth to it, so i would rule them out.
D is viable, but as per google documentation, there is a risk on managing SSH keys manually:

<https://cloud.google.com/compute/docs/instances/access-overview>

"Risks of manual key management

If you create and manage public SSH keys yourself through the Cloud Console, the gcloud command-line tool, or the API, you must keep track of the used keys and delete the public SSH keys for users who no longer have access. For example, if a team member leaves your project, remove their public SSH keys from metadata, so they can't continue to access your instances.

Additionally, specifying your gcloud tool or API calls incorrectly can potentially wipe out all of the public SSH keys in your project or on your instances, which disrupts connections for your project members.

If you aren't sure that you want to manage your own keys, use Compute Engine tools to connect to your instances instead."

A seems to me the best choice in terms of security and administration.

upvoted 1 times

You have created a code snippet that should be triggered whenever a new file is uploaded to a Cloud Storage bucket. You want to deploy this code snippet. What should you do?

- A. Use App Engine and configure Cloud Scheduler to trigger the application using Pub/Sub.
- B. Use Cloud Functions and configure the bucket as a trigger resource.
- C. Use Google Kubernetes Engine and configure a CronJob to trigger the application using Pub/Sub.
- D. Use Dataflow as a batch job, and configure the bucket as a data source.

Correct Answer: B

Community vote distribution

B (100%)

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (B):

Google Cloud Storage Triggers

Cloud Functions can respond to change notifications emerging from Google Cloud Storage. These notifications can be configured to trigger in response to various events inside a bucket—object creation, deletion, archiving and metadata updates.

Note: Cloud Functions can only be triggered by Cloud Storage buckets in the same Google Cloud Platform project.

Event types

Cloud Storage events used by Cloud Functions are based on Cloud Pub/Sub Notifications for Google Cloud Storage and can be configured in a similar way.

Supported trigger type values are:

google.storage.object.finalize

google.storage.object.delete

google.storage.object.archive

google.storage.object.metadataUpdate

Object Finalize

Trigger type value: google.storage.object.finalize

This event is sent when a new object is created (or an existing object is overwritten, and a new generation of that object is created) in the bucket.

https://cloud.google.com/functions/docs/calling/storage#event_types

upvoted 38 times

✉  **francisco_guerra**  3 years, 1 month ago

The answer is B

upvoted 19 times

✉  **SSPC** 3 years, 1 month ago

Sure B? Please you could share the link with the Google documentation

upvoted 1 times

✉  **Ale1973** 3 years, 1 month ago

<https://cloud.google.com/functions/docs/calling/storage>

upvoted 3 times

✉  **scanner2**  1 month ago

Selected Answer: B

Use "Object finalized" event of the Cloud Storage bucket as trigger for the Cloud Functions.

<https://cloud.google.com/functions/docs/calling/storage>

upvoted 1 times

✉  **Captain1212** 1 month ago

Selected Answer: B

B seems the correct option, as we can use the cloud functions as per our requirement for the cloud storage bucket..

upvoted 1 times

✉  **Angel_99** 1 year, 1 month ago

Selected Answer: B

The answer is B

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is correct, it is required on demand when upload happens
upvoted 1 times

 **arvsrv** 1 year, 8 months ago

Selected Answer: B

The answer is B
upvoted 2 times

 **Surat** 1 year, 9 months ago

I vote for B
upvoted 2 times

 **alaahakim** 1 year, 10 months ago

Selected Answer: B

Vote For B
upvoted 4 times

 **vamgcp** 2 years ago

Correct Answer is B - Use Cloud Functions and configure the bucket as a trigger resource.
upvoted 2 times

 **Zimcruza** 2 years, 3 months ago

Question asks:
You want to deploy this code snippet. What should you do?
To me, none of the answers is relevant to DEPLOYMENT - they all are about how you get the trigger to run ...
upvoted 1 times

 **GoCloud** 2 years, 6 months ago

vote for B
upvoted 2 times

 **marialix87** 2 years, 6 months ago

I think is B
upvoted 3 times

 **GCP_Student1** 2 years, 6 months ago

B. Use Cloud Functions and configure the bucket as a trigger resource.
upvoted 3 times

 **lutoa** 2 years, 7 months ago

Answer is B, app engine is for applications, in this case it's just a code snippet which cloud functions is more suitable
upvoted 2 times

 **victory108** 2 years, 8 months ago

B - Use Cloud Functions and configure the bucket as a trigger resource.
upvoted 3 times

 **Bhagirathi** 2 years, 10 months ago

Most of you suggest B.

someone chose A - I will choose B
upvoted 1 times

You have been asked to set up Object Lifecycle Management for objects stored in storage buckets. The objects are written once and accessed frequently for 30 days. After 30 days, the objects are not read again unless there is a special need. The objects should be kept for three years, and you need to minimize cost.

What should you do?

- A. Set up a policy that uses Nearline storage for 30 days and then moves to Archive storage for three years.
- B. Set up a policy that uses Standard storage for 30 days and then moves to Archive storage for three years.
- C. Set up a policy that uses Nearline storage for 30 days, then moves the Coldline for one year, and then moves to Archive storage for two years.
- D. Set up a policy that uses Standard storage for 30 days, then moves to Coldline for one year, and then moves to Archive storage for two years.

Correct Answer: A

Reference:

https://books.google.com.pk/books?id=q0nhDwAAQBAJ&pg=PA52&lpg=PA52&dq=Set+up+a+policy+that+uses+Nearline+storage+for+30+days+and+then+moves+to+Archive+storage+for+three+years.&source=bl&ots=kYLZN1ymA8&sig=ACfU3U2XLmzQ39cmPDwjfWxRbNtDNLc_6g&hl=en&sa=X&ved=2ahUKEwjZmef0pr7qAhVzQkEAHTUgASyQ6AEwAHoECAoQAQ#v=onepage&q=Set%20up%20a%20policy%20that%20uses%20Nearline%20storage%20for%2030%20days%20and%20then%20moves%20to%20Archive%20storage%20for%20three%20years.&f=false

Community vote distribution

B (100%)

ESP_SAP Highly Voted 3 years, 1 month ago

Correct Answer is (B):

The key to understand the requirement is : "The objects are written once and accessed frequently for 30 days"
Standard Storage

Standard Storage is best for data that is frequently accessed ("hot" data) and/or stored for only brief periods of time.

Archive Storage

Archive Storage is the lowest-cost, highly durable storage service for data archiving, online backup, and disaster recovery. Unlike the "coldest" storage services offered by other Cloud providers, your data is available within milliseconds, not hours or days. Archive Storage is the best choice for data that you plan to access less than once a year.

<https://cloud.google.com/storage/docs/storage-classes#standard>
upvoted 50 times

naveedpk00 3 years ago

What if we chose option D to minimize the cost as asked in the question? What do you think?
upvoted 1 times

gcper 2 years, 11 months ago

It doesn't minimize the costs. Check the costs of coldline vs archival
upvoted 5 times

SSPC Highly Voted 3 years, 1 month ago

I think the correct one is B. Because Nearline has a 30-day minimum storage duration.
<https://cloud.google.com/storage/docs/storage-classes>
upvoted 14 times

pepepy 3 years, 1 month ago

The object should be kept for three years, and you need to minimize cost, after 30 days it will be moved to archive, ans A
upvoted 1 times

pepepy 3 years, 1 month ago

Sorry you are right accessed frequently for 30 days, its B
upvoted 6 times

scanner2 Most Recent 1 month ago

Selected Answer: B

Key terms frequently accessed for 30 days -> Standard storage class.
Not accessed unless special need for 3 years -> Archive storage class.

upvoted 1 times

□ **Captain1212** 1 month ago

Selected Answer: B

B is the correct answer, as the data for first 30 days is accessed frequently so for it we can use the standard , and after it to minimize the cost we can use the archive storage for 3 years

upvoted 1 times

□ **ashtonez** 7 months ago

Answer is B, we cannot select A because data is accessed frequently and nearline only allows access once per month (you can access more incurring in additional cost but being not a cost optimized selection)

upvoted 1 times

□ **thaliath** 9 months ago

Answer is A: there is a retrieval fee for data access from nearline. Please check <https://cloud.google.com/storage/docs/storage-classes>. So Standard storage is the cheaper option

upvoted 1 times

□ **Charumathi** 1 year ago

Selected Answer: B

B is the correct Answer,
Frequently accessed data 'Hot Data' should be stored in Standard Storage for 30 days,
Then this can be moved to Archive after 30 days for period of three years which is accessed only when a special need arises, to reduce cost.

upvoted 1 times

□ **taiyi078** 1 year, 3 months ago

<https://cloud.google.com/storage/docs/storage-classes#nearline>

Nearline storage is ideal for data you plan to read or modify on average once per month or less. For example, if you want to continuously add files to Cloud Storage and plan to access those files once a month for analysis, Nearline storage is a great choice.

Nearline storage is also appropriate for data backup, long-tail multimedia content, and data archiving. Note, however, that for data accessed less frequently than once a quarter, Coldline storage or Archive storage are more cost-effective, as they offer lower storage costs.

upvoted 1 times

□ **taiyi078** 1 year, 3 months ago

Nearline storage is a low-cost, highly durable storage service for storing infrequently accessed data. Nearline storage is a better choice than Standard storage in scenarios where slightly lower availability, a 30-day minimum storage duration, and costs for data access are acceptable trade-offs for lowered at-rest storage costs.

upvoted 1 times

□ **AzureDP900** 1 year, 3 months ago

B is right

upvoted 1 times

□ **AzureDP900** 1 year, 3 months ago

I am changing it to D. Set up a policy that uses Standard storage for 30 days, then moves to Coldline for one year, and then moves to Archive storage for two years.

upvoted 2 times

□ **Shweta2jun** 5 months, 1 week ago

Is B right or D?

upvoted 1 times

□ **pluiedust** 1 year, 7 months ago

Selected Answer: B

B for sure

upvoted 1 times

□ **SleepyHitman** 1 year, 8 months ago

Selected Answer: B

The answer is: B

Standard storage description:

[https://cloud.google.com/storage/docs/storage-classes#:~:text=Standard%20Storage%20is%20best%20for%20data%20that%20is%20frequently%20accessed%20\(%22hot%22%20data\)%20and%20stored%20for%20only%20brief%20periods%20of%20time](https://cloud.google.com/storage/docs/storage-classes#:~:text=Standard%20Storage%20is%20best%20for%20data%20that%20is%20frequently%20accessed%20(%22hot%22%20data)%20and%20stored%20for%20only%20brief%20periods%20of%20time)

Nearline storage imp description:

<https://cloud.google.com/storage/docs/storage-classes#:~:text=storage%20service%20for-,storing%20infrequently%20accessed%20data.,-Nearline%20Storage%20is>

Pricing for frequent access:

<https://cloud.google.com/storage/pricing#:~:text=Free%20operations-,Standard%20Storage,Free,-Coldline%20Storage>

upvoted 4 times

□ **Surat** 1 year, 9 months ago

Standard – Frequently access and short period

Nearline - Low cost, highly durable for infrequent data access, lower availability

Coldline -Very low cost, highly durable for infrequently accessed, 90 days minimum storage

Archive - Lowest cost, highly durable for archiving, backup and DR, lower availability

I will also go for B

upvoted 4 times

 **fazalmf** 1 year, 9 months ago

B. Set up a policy that uses Standard storage for 30 days and then moves to Archive storage for three years.

_ clearly mentioned data access frequently for 30days and then object not read again until special need (archive suitable).

upvoted 1 times

 **Naren080914** 1 year, 9 months ago

Selected Answer: B

If you access frequently in Nearline, it will cost you more. There is no retrieval cost for Standard. So for frequently accessed object, go with standard. Also in Standard, there is no minimum period to store the object, so don't get confused by 30 days. Ans is B.

upvoted 1 times

 **alaahakim** 1 year, 10 months ago

Vote For B

upvoted 1 times

 **AD_0525** 2 years, 3 months ago

Answer should be D. object will be frequently access for first 30 days. cost of data access from standard storage is less than Nearline. Then after 30 days the objects can be moved to archive.

upvoted 2 times

 **AD_0525** 2 years, 3 months ago

Typo, answer should e B. Description remains the same.

upvoted 2 times

 **jahnu** 2 years, 5 months ago

My Answer is B. why because Nearline Storage is a low-cost, highly durable storage service for storing infrequently accessed data. so we use standard storage for frequent access.

upvoted 1 times

You are storing sensitive information in a Cloud Storage bucket. For legal reasons, you need to be able to record all requests that read any of the stored data. You want to make sure you comply with these requirements. What should you do?

- A. Enable the Identity Aware Proxy API on the project.
- B. Scan the bucket using the Data Loss Prevention API.
- C. Allow only a single Service Account access to read the data.
- D. Enable Data Access audit logs for the Cloud Storage API.

Correct Answer: D

Reference:

<https://cloud.google.com/storage/docs/audit-logs>*Community vote distribution*

D (100%)

ESP_SAP Highly Voted 3 years, 1 month ago

Correct Answer is (D):

Logged information

Within Cloud Audit Logs, there are two types of logs:

Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.

Data Access logs: Entries for operations that modify objects or read a project, bucket, or object. There are several sub-types of data access logs:

ADMIN_READ: Entries for operations that read the configuration or metadata of a project, bucket, or object.

DATA_READ: Entries for operations that read an object.

DATA_WRITE: Entries for operations that create or modify an object.

<https://cloud.google.com/storage/docs/audit-logs#types>

upvoted 28 times

francisco_guerra Highly Voted 3 years, 1 month ago

D is the correct one

upvoted 19 times

SSPC 3 years, 1 month ago

Yes D is the correct

upvoted 6 times

scanner2 Most Recent 1 month ago**Selected Answer: D**

Enable Data access audit logs for Cloud storage bucket

<https://cloud.google.com/storage/docs/audit-logging>

upvoted 1 times

Captain1212 1 month ago**Selected Answer: D**

D is the correct answer

upvoted 1 times

calm_fox 10 months, 2 weeks ago**Selected Answer: D**

Only logical option

upvoted 1 times

AzureDP900 1 year, 3 months ago

D is right for this use case

upvoted 1 times

Akash7 1 year, 4 months ago

D is correct as Data Access logs pertaining to Cloud Storage operations are not recorded by default. You have to enable them ...

<https://cloud.google.com/storage/docs/audit-logging>

upvoted 2 times

wael_tn 1 year, 5 months ago

Selected Answer: D

I think it's D

upvoted 1 times

 **Surat** 1 year, 9 months ago

I also vote for D

upvoted 2 times

 **Vinoth9289** 2 years, 1 month ago

D is the correct Answer

upvoted 2 times

 **WakandaF** 2 years, 4 months ago

seems that B is the right!

Cloud Data Loss Prevention (DLP) helps you to understand and manage such sensitive data. It provides fast, scalable classification and redaction for sensitive data elements. Using the Data Loss Prevention API and Cloud Functions, you can automatically scan this data before it is uploaded to the shared storage bucket.

upvoted 1 times

 **YAS007** 2 years, 2 months ago

the question doesn't ask you to manage or understand sensitive data :

" you need to be able to record all requests that read any of the stored data"

upvoted 3 times

 **victory108** 2 years, 7 months ago

D - Enable Data Access audit logs for the Cloud Storage API.

upvoted 1 times

 **EABDAJA** 2 years, 7 months ago

D is correct

upvoted 1 times

 **GCP_Student1** 2 years, 7 months ago

D. Enable Data Access audit logs for the Cloud Storage API.

upvoted 2 times

 **swatititame** 2 years, 10 months ago

• D. Enable Data Access audit logs for the Cloud Storage API.

upvoted 1 times

 **RockAJ** 3 years ago

Ans is D

upvoted 2 times

 **pepepy** 3 years, 1 month ago

agree with D

upvoted 2 times

You are the team lead of a group of 10 developers. You provided each developer with an individual Google Cloud Project that they can use as their personal sandbox to experiment with different Google Cloud solutions. You want to be notified if any of the developers are spending above \$500 per month on their sandbox environment. What should you do?

- A. Create a single budget for all projects and configure budget alerts on this budget.
- B. Create a separate billing account per sandbox project and enable BigQuery billing exports. Create a Data Studio dashboard to plot the spending per billing account.
- C. Create a budget per project and configure budget alerts on all of these budgets.
- D. Create a single billing account for all sandbox projects and enable BigQuery billing exports. Create a Data Studio dashboard to plot the spending per project.

Correct Answer: C

Reference:

<https://cloud.google.com/billing/docs/how-to/budgets>

Community vote distribution

C (100%)

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (C):

Set budgets and budget alerts

Overview

Avoid surprises on your bill by creating Cloud Billing budgets to monitor all of your Google Cloud charges in one place. A budget enables you to track your actual Google Cloud spend against your planned spend. After you've set a budget amount, you set budget alert threshold rules that are used to trigger email notifications. Budget alert emails help you stay informed about how your spend is tracking against your budget.

2. Set budget scope

Set the budget Scope and then click Next.

In the Projects field, select one or more projects that you want to apply the budget alert to. To apply the budget alert to all the projects in the Cloud Billing account, choose Select all.

<https://cloud.google.com/billing/docs/how-to/budgets#budget-scop>

upvoted 45 times

✉  **bobthebuilder55110** 1 year, 2 months ago

wait a minute, why not A ?

As you said that

" In the Projects field, select one or more projects that you want to apply the budget alert to. To apply the budget alert to all the projects in the Cloud Billing account, choose Select all. "

As per this I should be able to create single budget for all the projects and should be able to set alert on that, why create separate budget for all 10 projects ?

upvoted 1 times

✉  **Priyanka109** 1 year ago

It will be a combined budget that's why it's C

upvoted 2 times

✉  **dang1986** 1 year, 8 months ago

You're the only answer I take seriously "Thumbs up"

upvoted 3 times

✉  **Hjameel**  3 years, 1 month ago

I think C is the best answer.

upvoted 10 times

✉  **Captain1212**  1 month ago

Selected Answer: C

C is the correct answer, because question demands that which project goes over the 500 per month, to check that you need to create the budget per project

upvoted 1 times

✉  **Kopy** 11 months ago

Selected Answer: C

Correct Answer is (C):

upvoted 1 times

Angel_99 1 year, 1 month ago

Selected Answer: C

Key is anyone goes above \$500 means it requires project level
upvoted 2 times

AzureDP900 1 year, 3 months ago

Key is anyone goes above \$500 means it requires project level so C is right
upvoted 1 times

wael_tn 1 year, 5 months ago

Selected Answer: C

Clearly C is the answer
upvoted 1 times

pondai 2 years, 6 months ago

Does anyone knows Data Studio can be alert to email? If it can't I'll pick C
upvoted 5 times

GCP_Student1 2 years, 7 months ago

C. Create a budget per project and configure budget alerts on all of these budgets.
upvoted 4 times

andregrip 2 years, 8 months ago

I believe is A because of this...
Projects: In the Projects field, select one or more projects that you want to apply the budget alert to. To apply the budget alert to all of the projects in the Cloud Billing account, choose Select all.

Some costs are not related to a project, such as the costs of subscriptions or Support costs.

In the budget's project scope, in the list of projects you can filter on, [Charges not specific to a project] is not an option you can select. If you choose Select all, then the costs in all projects, including Charges not specific to a project, are included in the budget and cost trend chart cost calculations.

If you select one or more projects - but not all projects - then the Charges not specific to a project are not included in the budget and cost trend chart cost calculations.

You can view your costs that are not related to a project in the billing reports. Using the projects filter in the reports page, you can select and view [Charges not specific to a project].

URL: <https://cloud.google.com/billing/docs/how-to/budgets>

upvoted 3 times

Ridhanya 1 year, 10 months ago

but how will you know who crossed the limit. what if the summation of their usage exceeds 500? the corresponding alert would be a false alarm
upvoted 5 times

Bhagirathi 2 years, 10 months ago

which one correct?
upvoted 1 times

swatititame 2 years, 10 months ago

• C. Create a budget per project and configure budget alerts on all of these budgets.
upvoted 1 times

xtian2900 3 years ago

is C, not A
with A, i guess if you create a single budget for all projects, together they can easily beat the \$500 mark and you need to know if "one deveoper" did it.
so one budget per project is the solution.
upvoted 4 times

[Removed] 3 years ago

Yes C is correct, I would have gone with B and D as billing export is the crucial element for billing, but both the option don't talk about notifying about spending.
Also, you don't want to combine the billing for all as each one can spend up to 500 so it will be better if they are individual so A is out.
upvoted 2 times

pepepy 3 years, 1 month ago

You want to be notified if any of the developers are spending above \$500 per month on their sandbox environment, so Answer is C
upvoted 2 times

tothecloud 3 years, 1 month ago

C is correct
upvoted 4 times

SSPC 3 years, 1 month ago

Do you think "C" is correct?
upvoted 2 times

You are deploying a production application on Compute Engine. You want to prevent anyone from accidentally destroying the instance by clicking the wrong button. What should you do?

- A. Disable the flag `Delete boot disk when instance is deleted`.
- B. Enable delete protection on the instance.
- C. Disable Automatic restart on the instance.
- D. Enable Preemptibility on the instance.

Correct Answer: A

Reference:

<https://googlecloudplatform.uservoice.com/forums/302595-compute-engine/suggestions/14227521-set-delete-boot-disk-when-instance-is-deleted-to>

Community vote distribution

B (100%)

 **ESP_SAP** Highly Voted 3 years, 1 month ago

Correct Answer is (B):

Preventing Accidental VM Deletion

This document describes how to protect specific VM instances from deletion by setting the `deletionProtection` property on an Instance resource. To learn more about VM instances, read the Instances documentation.

As part of your workload, there might be certain VM instances that are critical to running your application or services, such as an instance running a SQL server, a server used as a license manager, and so on. These VM instances might need to stay running indefinitely so you need a way to protect these VMs from being deleted.

By setting the `deletionProtection` flag, a VM instance can be protected from accidental deletion. If a user attempts to delete a VM instance for which you have set the `deletionProtection` flag, the request fails. Only a user that has been granted a role with `compute.instances.create` permission can reset the flag to allow the resource to be deleted.

<https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>

upvoted 37 times

 **Naree** 3 months ago

Mr.ESP_SAP, your answers are on the spot and I look forward to your notes on all the questions first.. Appreciate your effort and support for this cloud community.. :)

upvoted 3 times

 **MohammedGhouse** Highly Voted 3 years, 1 month ago

"B" is the answer

upvoted 11 times

 **scanner2** Most Recent 1 month ago

Selected Answer: B

<https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>

upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: B

B is the correct answer , as it helps to prevent critical instance to get deleted

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

This is straight forward question, enable delete protection. B is right

upvoted 1 times

 **Himadhar1997** 1 year, 4 months ago

Selected Answer: B

Preventing Accidental VM Deletion

This document describes how to protect specific VM instances from deletion by setting the `deletionProtection` property on an Instance resource. To learn more about VM instances, read the Instances documentation.

upvoted 2 times

 **Surat** 1 year, 9 months ago

B seems right option

upvoted 3 times

 **kped21** 1 year, 9 months ago

B - on VM Enable delete protection
upvoted 2 times

 **jaffarali** 1 year, 10 months ago

Selected Answer: B

Answer is B. there is an Option in VM instance while creating
upvoted 3 times

 **Sreedharveluru** 2 years, 2 months ago

Option A would not prevent , It can be used only after the damage is done. Hence B
upvoted 1 times

 **NARWAL** 2 years, 6 months ago

B should be the answer.
upvoted 1 times

 **GCP_Student1** 2 years, 7 months ago

B. Enable delete protection on the instance.
upvoted 2 times

 **Lomy** 2 years, 7 months ago

B
The ans is B
upvoted 1 times

 **RockAJ** 3 years ago

B for me
upvoted 2 times

 **muk5658** 3 years ago

Correct Answer is 'B'
upvoted 3 times

 **SSPC** 3 years, 1 month ago

B is correct answer. <https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>
upvoted 5 times

 **francisco_guerra** 3 years, 1 month ago

I think is B
upvoted 3 times

Your company uses a large number of Google Cloud services centralized in a single project. All teams have specific projects for testing and development. The

DevOps team needs access to all of the production services in order to perform their job. You want to prevent Google Cloud product changes from broadening their permissions in the future. You want to follow Google-recommended practices. What should you do?

- A. Grant all members of the DevOps team the role of Project Editor on the organization level.
- B. Grant all members of the DevOps team the role of Project Editor on the production project.
- C. Create a custom role that combines the required permissions. Grant the DevOps team the custom role on the production project.
- D. Create a custom role that combines the required permissions. Grant the DevOps team the custom role on the organization level.

Correct Answer: B*Community vote distribution*

C (100%)

  **ESP_SAP**  3 years, 1 month ago

Correct Answer is (C):

Understanding IAM custom roles

Key Point: Custom roles enable you to enforce the principle of least privilege, ensuring that the user and service accounts in your organization have only the permissions essential to performing their intended functions.

Basic concepts

Custom roles are user-defined, and allow you to bundle one or more supported permissions to meet your specific needs. Custom roles are not maintained by Google; when new permissions, features, or services are added to Google Cloud, your custom roles will not be updated automatically.

When you create a custom role, you must choose an organization or project to create it in. You can then grant the custom role on the organization or project, as well as any resources within that organization or project.

https://cloud.google.com/iam/docs/understanding-custom-roles#basic_concepts

upvoted 51 times

  **SSPC**  3 years, 1 month ago

"You want to prevent Google Cloud product changes from broadening their permissions in the future." then CUSTOM ROLE
upvoted 24 times

  **Rothmansua** 1 year, 11 months ago

Great hint, thanks!

upvoted 2 times

  **rahulrauki**  1 week, 6 days ago**Selected Answer: C**

The giveaway is "prevent google cloud product changes from broadening their permissions". Which means that we need to create a custom role. Also they mentioned all production services and not production projects so C
upvoted 1 times

  **scanner2** 1 month ago**Selected Answer: C**

Custom roles help you enforce the principle of least privilege, because they help to ensure that the principals in your organization have only the permissions that they need.

Custom roles are user-defined, and allow you to bundle one or more supported permissions to meet your specific needs. When you create a custom role, you must choose an organization or project to create it in. You can then grant the custom role on the organization or project, as well as any resources within that organization or project.

Note: You cannot define custom roles at the folder level. If you need to use a custom role within a folder, define the custom role at the organization level.

<https://cloud.google.com/iam/docs/roles-overview#custom>

upvoted 1 times

  **scanner2** 1 month ago**Selected Answer: C**

Custom roles help you enforce the principle of least privilege, because they help to ensure that the principals in your organization have only the permissions that they need.

Custom roles are user-defined, and allow you to bundle one or more supported permissions to meet your specific needs. When you create a custom role, you must choose an organization or project to create it in. You can then grant the custom role on the organization or project, as well as any resources within that organization or project.

Note: You cannot define custom roles at the folder level. If you need to use a custom role within a folder, define the custom role at the

<https://cloud.google.com/iam/docs/roles-overview#custom>

upvoted 1 times

✉️ **Captain1212** 1 month ago

Selected Answer: C

C is the correct answer

upvoted 1 times

✉️ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

✉️ **slcvlctetri** 9 months, 2 weeks ago

Selected Answer: C

Had this question 2 days ago. C is correct.

upvoted 1 times

✉️ **Charumathi** 1 year ago

Selected Answer: C

C is the correct answer, give the devops team the least privileged role, only the required permissions to access the production services, as the question states 'to prevent product changes' for which editor role is not recommended either at Project or organizational level, organizational level access gives broad scope to all the projects in the organization, this role cannot be given to the devops team.

- A. Editor has privilege to change the products, and the scope is broad
- B. Editor has privilege to change the products
- C. Recommended, as this will give only required permission at project level to devops team.
- D. They require only project level access. This gives access to all project in the organization.

upvoted 1 times

✉️ **ale_brd_** 1 year ago

Selected Answer: C

Correct Answer is (C):

Custom roles are user-defined, and allow you to bundle one or more supported permissions to meet your specific needs. Custom roles are not maintained by Google; when new permissions, features, or services are added to Google Cloud, your custom roles will not be updated automatically.

upvoted 1 times

✉️ **sonuricky** 1 year, 2 months ago

correct answer is A

upvoted 2 times

✉️ **Roro_Brother** 1 year, 3 months ago

Selected Answer: C

There is no doubt, the correct answer is C

upvoted 1 times

✉️ **AzureDP900** 1 year, 3 months ago

C is right

upvoted 1 times

✉️ **DaveNZ** 1 year, 7 months ago

Selected Answer: C

C seems to be the popular answer, and it makes sense because the generic roles are not sufficient for these specific requirements. I added this voting comment because the community answers are not currently visible.

upvoted 2 times

✉️ **Surat** 1 year, 9 months ago

I vote for C

upvoted 2 times

✉️ **rachee** 1 year, 10 months ago

I initially thought C. But I think this may be a trick question. "The DevOps team needs access to ALL of the PRODUCTION services..." which are in a "single" project. If "Project Editor" is assigned at on the "production" project it gives them access to "ALL" production services including product changes in the "production" project. A custom role would have to be modified to get access to product changes in the production project that required additional permissions, so the DevOps team would not have access to "ALL" services until the custom role is modified.

I am changing my choice to B.

upvoted 4 times

✉️ **theBestStudent** 1 year, 1 month ago

Your choice doesn't follow the rule of least privilege. So correct answer is C.

upvoted 1 times

✉️ **jahnu** 2 years, 5 months ago

Answer C: custom role enable privileges.

upvoted 2 times

You are building an application that processes data files uploaded from thousands of suppliers. Your primary goals for the application are data security and the expiration of aged data. You need to design the application to:

- * Restrict access so that suppliers can access only their own data.
- * Give suppliers write access to data only for 30 minutes.
- * Delete data that is over 45 days old.

You have a very short development cycle, and you need to make sure that the application requires minimal maintenance. Which two strategies should you use?

(Choose two.)

- A. Build a lifecycle policy to delete Cloud Storage objects after 45 days.
- B. Use signed URLs to allow suppliers limited time access to store their objects.
- C. Set up an SFTP server for your application, and create a separate user for each supplier.
- D. Build a Cloud function that triggers a timer of 45 days to delete objects that have expired.
- E. Develop a script that loops through all Cloud Storage buckets and deletes any buckets that are older than 45 days.

Correct Answer: AE

Community vote distribution

AB (100%)

✉  **ESP_SAP**  3 years, 1 month ago

Correct Answers are: (AB):

(A) Object Lifecycle Management

Delete

The Delete action deletes an object when the object meets all conditions specified in the lifecycle rule.

Exception: In buckets with Object Versioning enabled, deleting the live version of an object causes it to become a noncurrent version, while deleting a noncurrent version deletes that version permanently.

<https://cloud.google.com/storage/docs/lifecycle#delete>

(B) Signed URLs

This page provides an overview of signed URLs, which you use to give time-limited resource access to anyone in possession of the URL, regardless of whether they have a Google account

<https://cloud.google.com/storage/docs/access-control/signed-urls>

upvoted 43 times

✉  **francisco_guerra**  3 years, 1 month ago

AB is the answer

upvoted 14 times

✉  **scanner2**  1 month ago

Selected Answer: AB

Create object lifecycle policy to automatically delete the objects after 45 days. Create signed URLs to temporarily provide the access for specified amount of time.

upvoted 1 times

✉  **Captain1212** 1 month ago

Selected Answer: AB

AB is the correct answer, as A helps to make a lifecycle policy to delete the data after 45 days and b helps the customer to acces their data as per the question requiremnt

upvoted 1 times

✉  **jrls1991** 8 months, 1 week ago

Selected Answer: AB

It's a bit obvious. Cloud functions wouldn't really work well with this and would probably require lots of maintenance just as any other option. AB are the correct ones.

upvoted 2 times

✉  **diasporabro** 11 months, 3 weeks ago

Selected Answer: AB

AB achieves this objective

upvoted 2 times

✉  **olme59** 1 year, 1 month ago

Selected Answer: AB

its clearly AB, life cycle and provider private url
upvoted 1 times

 **Angel_99** 1 year, 1 month ago

Selected Answer: AB

Correct Answer Combo: (AB)
upvoted 1 times

 **abirroy** 1 year, 2 months ago

Selected Answer: AB

Correct Answers are: (AB)
upvoted 1 times

 **patashish** 1 year, 3 months ago

Correct Answers are: A and B
upvoted 1 times

 **Roro_Brother** 1 year, 3 months ago

Selected Answer: AB

A and B is the right answer
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A and B right answer
upvoted 1 times

 **DaveNZ** 1 year, 7 months ago

Selected Answer: AB

The website answer says A and E, but these do the same thing. The question also asks about setting up security, so the other part of the answer must be B
upvoted 1 times

 **jahnu** 2 years, 5 months ago

My Answer: A B
upvoted 2 times

 **GCP_Student1** 2 years, 6 months ago

A. Build a lifecycle policy to delete Cloud Storage objects after 45 days.
B. Use signed URLs to allow suppliers limited time access to store their objects.
upvoted 3 times

 **Bhagirathi** 2 years, 9 months ago

A & B helps.
upvoted 1 times

 **tuta** 2 years, 9 months ago

did you pass exam ?
upvoted 4 times

 **RockAJ** 3 years ago

for me A and B
upvoted 2 times

Your company wants to standardize the creation and management of multiple Google Cloud resources using Infrastructure as Code. You want to minimize the amount of repetitive code needed to manage the environment. What should you do?

- A. Develop templates for the environment using Cloud Deployment Manager.
- B. Use curl in a terminal to send a REST request to the relevant Google API for each individual resource.
- C. Use the Cloud Console interface to provision and manage all related resources.
- D. Create a bash script that contains all requirement steps as gcloud commands.

Correct Answer: A

Reference:

<https://cloud.google.com/deployment-manager/docs/fundamentals>
(see templates)

Community vote distribution

A (100%)

 **jmgf**  2 years, 7 months ago

A

You can use Google Cloud Deployment Manager to create a set of Google Cloud resources and manage them as a unit, called a deployment. For example, if your team's development environment needs two virtual machines (VMs) and a BigQuery database, you can define these resources in a configuration file, and use Deployment Manager to create, change, or delete these resources. You can make the configuration file part of your team's code repository, so that anyone can create the same environment with consistent results.

<https://cloud.google.com/deployment-manager/docs/quickstart>

upvoted 31 times

 **GCP_Student1**  2 years, 6 months ago

A. Develop templates for the environment using Cloud Deployment Manager.

upvoted 11 times

 **scanner2**  1 month ago

Selected Answer: A

Cloud Deployment Manager is an infrastructure deployment service that automates the creation and management of Google Cloud resources. Write flexible template and configuration files and use them to create deployments that have a variety of Google Cloud services, such as Cloud Storage, Compute Engine, and Cloud SQL, configured to work together.

You can use Google Cloud Deployment Manager to create a set of Google Cloud resources and manage them as a unit, called a deployment.

<https://cloud.google.com/deployment-manager/docs>

upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: A

A seems more correct, u can use the deployment manager to create your instances with the same configurstion file

upvoted 1 times

 **N_A** 5 months ago

Selected Answer: A

A. Develop templates for the environment using Cloud Deployment Manager.

Although the preferred IaC tool is Terraform. There no mention of Deployment Manager anymore in the Google on-demand courses but there is an entire course on Terraform.

upvoted 1 times

 **PPP_D** 6 months, 2 weeks ago

I'm going with A

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

A is right

upvoted 1 times

 **POOJA3808** 1 year, 6 months ago

Selected Answer: A

Develop templates for the environment using Cloud Deployment Manager.

upvoted 1 times

 **look1** 1 year, 10 months ago

Selected Answer: A

Templates only

upvoted 2 times

✉️  **kartz14** 1 year, 10 months ago

Its A.

upvoted 2 times

✉️  **TAvenger** 2 years, 6 months ago

A.

According to this document <https://cloud.google.com/solutions/infrastructure-as-code>

IaC (Infrastructure as code) tools for Google Cloud:

Deployment Manager, Terraform, Puppet, Chef ...

upvoted 4 times

✉️  **lio123** 2 years, 7 months ago

I think A

upvoted 2 times

You are performing a monthly security check of your Google Cloud environment and want to know who has access to view data stored in your Google Cloud Project. What should you do?

- A. Enable Audit Logs for all APIs that are related to data storage.
- B. Review the IAM permissions for any role that allows for data access.
- C. Review the Identity-Aware Proxy settings for each resource.
- D. Create a Data Loss Prevention job.

Correct Answer: *B*

Reference:

<https://cloud.google.com/compute/docs/access>

Community vote distribution

B (96%)

 **JelloMan**  1 year, 4 months ago

Selected Answer: B

Only use audit logs to look at history (PAST)

If you need current, up-to-date, info regarding permissions always go to IAM

upvoted 11 times

 **Alejondri**  1 year, 5 months ago

Selected Answer: B

B is the one:

A. Enable Audit Logs for all APIs that are related to data storage. --> That is not the correct answer, if someone with permissions has not accessed or does not access, it will not be listed.

B. Review the IAM permissions for any role that allows for data access. --> That's correct

C. Review the Identity-Aware Proxy settings for each resource. --> Nothing relevant, Proxy? Is configured? The question don't ask or tell something about if it is configured.

D. Create a Data Loss Prevention job. --> Data Loss Prevention nothing to see here.

upvoted 8 times

 **snkhatri**  1 year, 1 month ago

B "WHO HAS ACCESS"

upvoted 2 times

 **Naree** 3 months ago

Yes, that's the catch.. The question here is "Who has access?" and not "Who has accessed?"

Answer is "B".

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is correct

upvoted 1 times

 **akshaychavan7** 1 year, 4 months ago

Selected Answer: B

Without any doubt, it's B.

upvoted 3 times

 **Terzlightyear** 1 year, 5 months ago

Selected Answer: B

B is the one

upvoted 3 times

 **sdflkds** 1 year, 5 months ago

B. 'Audit logs help you answer "who did what, where, and when?"'(from <https://cloud.google.com/logging/docs/audit>). So, not who has access, but rather who accessed.

upvoted 1 times

 **Malib** 1 year, 5 months ago

Selected Answer: A

La réponse A.

upvoted 1 times

Your company has embraced a hybrid cloud strategy where some of the applications are deployed on Google Cloud. A Virtual Private Network (VPN) tunnel connects your Virtual Private Cloud (VPC) in Google Cloud with your company's on-premises network. Multiple applications in Google Cloud need to connect to an on-premises database server, and you want to avoid having to change the IP configuration in all of your applications when the IP of the database changes.

What should you do?

- A. Configure Cloud NAT for all subnets of your VPC to be used when egressing from the VM instances.
- B. Create a private zone on Cloud DNS, and configure the applications with the DNS name.
- C. Configure the IP of the database as custom metadata for each instance, and query the metadata server.
- D. Query the Compute Engine internal DNS from the applications to retrieve the IP of the database.

Correct Answer: A

Community vote distribution

B (100%)

 **kopper2019**  2 years, 5 months ago

B,

Forwarding zones

Cloud DNS forwarding zones let you configure target name servers for specific private zones. Using a forwarding zone is one way to implement outbound DNS forwarding from your VPC network.

A Cloud DNS forwarding zone is a special type of Cloud DNS private zone. Instead of creating records within the zone, you specify a set of forwarding targets. Each forwarding target is an IP address of a DNS server, located in your VPC network, or in an on-premises network connected to your VPC network by Cloud VPN or Cloud Interconnect.

A does not apply, that is to provide internet access to resources

C, does not apply

D, I don't get it

so B

upvoted 29 times

 **MacFreak** 1 year, 1 month ago

"A does not apply, that is to provide internet access to resources" - do you really think NAT is only being used between public and private? Well...it's not! :)

upvoted 1 times

 **meh009** 2 years, 5 months ago

Agreed, It's B although I chose A initially. After some careful consideration and understanding how Cloud NAT works, I'm sticking with B

<https://cloud.google.com/nat/docs/overview>

upvoted 3 times

 **meh009** 2 years, 5 months ago

Further clarification:

"On-premises clients can resolve records in private zones, forwarding zones, and peering zones for which the VPC network has been authorized. On-premises clients use Cloud VPN or Cloud Interconnect to connect to the VPC network."

upvoted 1 times

 **djgodzilla** 2 years, 3 months ago

this is talking about On-premises client resolving nodes outside their network . the question is about how would the application tier within the VPC would resolve the database server . you're confusing the resolution direction my friend

upvoted 3 times

 **djgodzilla** 2 years, 3 months ago

It is still B , but it's rather outbound forward that's needed here :

DNS outbound Forwarding :

- Set up outbound forwarding private zones to query on-premises servers (On-prem Authoritative Zone: corp.example.com)
- In Cloud Router , add a custom route advertisement for GCP DNS proxies range 35.199.192.0/19 to the on-premises environment.

- Make sure inbound DNS traffic from 35.199.192.0/19 is allowed on on-prem firewall

- Cloud Router should be learning on-prem network route from On-prem Router

https://youtu.be/OH_Jw8NhEGU?t=1283

https://cloud.google.com/dns/docs/best-practices#use_forwarding_zones_to_query_on-premises_servers

upvoted 3 times

 **pondai**  2 years, 6 months ago

https://cloud.google.com/dns/docs/best-practices#best_practices_for_dns_forwarding_zones_and_server_policies

Cloud DNS offers DNS forwarding zones and DNS server policies to allow lookups of DNS names between your on-premises and Google Cloud environment. You have multiple options for configuring DNS forwarding. The following section lists best practices for hybrid DNS

setup. These best practices are illustrated in the Reference architectures for hybrid DNS.

So I think B is correct

upvoted 8 times

□  **scanner2** Most Recent 1 month ago

Selected Answer: B

DNS is a hierarchical distributed database that lets you store IP addresses and other data and look them up by name. Cloud DNS lets you publish your zones and records in DNS without the burden of managing your own DNS servers and software.

Cloud DNS offers both public zones and private managed DNS zones. A public zone is visible to the public internet, while a private zone is visible only from one or more Virtual Private Cloud (VPC) networks that you specify.

<https://cloud.google.com/dns/docs/overview>

upvoted 1 times

□  **Captain1212** 1 month ago

Selected Answer: B

B is the correct Answer

upvoted 1 times

□  **jrlsl1991** 8 months, 1 week ago

Selected Answer: B

Based on this - <https://cloud.google.com/dns/docs/overview#dns-forwarding-methods> B must be the best option. I don't think there's a "typo" (or completely wrongly worded answer) in option D (there's comments saying that instead of Compute Engine it should be on-premise). I believe option D is wrong on purpose to create a confusion.

upvoted 1 times

□  **Charumathi** 1 year ago

Selected Answer: B

B is correct answer,

Configure Private Google Access for on-premises hosts,

DNS configuration

Your on-premises network must have DNS zones and records configured so that Google domain names resolve to the set of IP addresses for either private.googleapis.com or restricted.googleapis.com. You can create Cloud DNS managed private zones and use a Cloud DNS inbound server policy, or you can configure on-premises name servers. For example, you can use BIND or Microsoft Active Directory DNS.

<https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid#config-domain>

upvoted 1 times

□  **aforolt** 1 year ago

Ans is D, looks like there is typo

upvoted 1 times

□  **habros** 1 year, 1 month ago

B. DNS works best with dynamic IPs.

upvoted 1 times

□  **patashish** 1 year, 3 months ago

Correct Ans is B

Ref - https://cloud.google.com/dns/docs/best-practices#best_practices_for_private_zones

upvoted 1 times

□  **deadlydeb** 1 year, 3 months ago

Selected Answer: B

B Cloud DNS

upvoted 2 times

□  **AzureDP900** 1 year, 3 months ago

B is correct

upvoted 1 times

□  **deadlydeb** 1 year, 3 months ago

Selected Answer: B

B: DNS

upvoted 1 times

□  **Mike_Malone** 1 year, 9 months ago

ans D

<https://cloud.google.com/compute/docs/internal-dns>

upvoted 3 times

□  **Ridhanya** 1 year, 10 months ago

option B

upvoted 1 times

□  **FernandoJ** 1 year, 10 months ago

Selected Answer: B

<https://cloud.google.com/dns/docs/overview#:~:text>Create%20an%20inbound,the%20VPC%20network.>

upvoted 1 times

 **gloisv** 2 years, 2 months ago

IT's D, because:

- A) Cloud NAT direction will be from the cloud resources to the on prem, an the DB is on prem (It will not work if the IP of the database on prem changes, and you have an VPN you should traffic your data trough it).
- B) If you create a private zone and configure the applications, if your servers doesnt use the dns will not work.
- C) I think is not suitable. I guess you should re-deploy all your Apps with the new conf if the address change.
- D) If you have an A record ip of your DB HOST(wich is on prem) in Engine internal DNS and if it changes, you can update the registry quickly to change it to the new IP address, so it will be the best option for sure.

upvoted 3 times

 **iloveexam** 2 years, 2 months ago

I choose D.

It means setting up a compute engine as internal DNS that contains a domain that resolves IP address of the database.

upvoted 2 times

You have developed a containerized web application that will serve internal colleagues during business hours. You want to ensure that no costs are incurred outside of the hours the application is used. You have just created a new Google Cloud project and want to deploy the application. What should you do?

- A. Deploy the container on Cloud Run for Anthos, and set the minimum number of instances to zero.
- B. Deploy the container on Cloud Run (fully managed), and set the minimum number of instances to zero.
- C. Deploy the container on App Engine flexible environment with autoscaling, and set the value min_instances to zero in the app.yaml.
- D. Deploy the container on App Engine flexible environment with manual scaling, and set the value instances to zero in the app.yaml.

Correct Answer: C

Community vote distribution

B (100%)

 **crysk**  2 years, 7 months ago

I think that is B the correct answer, because Cloud Run can scale to 0:

<https://cloud.google.com/run/docs/about-instance-autoscaling>

And App Engine Flexible can't scale to 0, the minimum instance number is 1:

https://cloud.google.com/appengine/docs/the-appengine-environments#comparing_high-level_features

upvoted 35 times

 **ryumada** 1 year, 2 months ago

No for the App Engine Flexible Environment, but App Engine Standard can also scale to zero.

upvoted 1 times

 **pca2b**  2 years, 6 months ago

B:

not A because Anthos is an add-on to GKE clusters, 'new project' means we don't have a GKE cluster to work with
https://cloud.google.com/kuberun/docs/architecture-overview#components_in_the_default_installation

upvoted 15 times

 **scanner2**  1 month ago

Selected Answer: B

Cloud Run is a managed compute platform that lets you run containers directly on top of Google's scalable infrastructure.

Cloud Run adds and removes instances automatically to handle all incoming requests. If there are no incoming requests to your service, even the last remaining instance will be removed. This behavior is commonly referred to as scale to zero.
<https://cloud.google.com/run/docs/overview/what-is-cloud-run>

upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: B

B is the correct answer, in c, d as App engine flexible environment can't scale to Zero and A in this GKE cluster is used but we have just created a project so it will add extra cost

upvoted 1 times

 **N_A** 5 months ago

Selected Answer: B

C. and D. are wrong answers as only the App Engine standard environment scales down to zero.

Answer A. will incur extra cost as Cloud Run for Anthos runs on Kubernetes, so need to have a k8s cluster available.

B. Is correct, as "Cloud Run automatically scales up or down from zero to N depending on traffic, leveraging container image streaming for a fast startup time." from <https://cloud.google.com/run>

upvoted 4 times

 **Charumathi** 1 year ago

Selected Answer: B

B is the correct answer,

Cloud Functions can scale to zero, whereas App Engine will not be able to scale to zero, it should have at least one instance.

Add-on Info,

App-Engine Standard can scale to zero, whereas App-Engine Flexible couldn't scale down to zero.

upvoted 4 times

 **sylva91** 1 year ago

Selected Answer: B

B is the answer since we can scale to 0 and the other key word is "containerized"

upvoted 2 times

 **Roro_Brother** 1 year, 3 months ago

Selected Answer: B

B is the right answer

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is correct

upvoted 1 times

 **tomnatoli75** 1 year, 4 months ago

<https://cloud.google.com/appengine/docs/standard/python/config/appref>

min_instances

Warning: For this feature to function properly, you must make sure that warmup requests are enabled and that your application handles warmup requests.

Note: This setting applies only if the version of the app defined by this app.yaml file is configured to receive traffic. To learn more about routing traffic to different versions of an app, see Splitting Traffic.

Optional. The minimum number of instances for App Engine to create for this module version. These instances serve traffic when requests arrive, and continue to serve traffic even when additional instances are started up as required to handle traffic.

Specify a value from 0 to 1000. You can set the parameter to the value 0 to allow scaling to 0 instances to lower costs when no requests are being served. Note that you are charged for the number of instances specified whether they are receiving traffic or not.

So C

upvoted 1 times

 **JelloMan** 1 year, 5 months ago

Selected Answer: B

along with the reason that most have stated (only Cloud Run can scale down to 0 instances) another reason is that Cloud Run is pay-per-use. App-engine flexible is paid based on usage of vCPU, memory, and persistent disks, so you will be racking up cost quickly just because your VM's are created (regardless if they're running or not)

upvoted 4 times

 **jblima** 1 year, 5 months ago

Selected Answer: B

B is the correct

upvoted 1 times

 **nshah68** 1 year, 7 months ago

Selected Answer: B

Flex cannot scale down to 0. Standard can but that is not relevant here.

Therefore, Cloud Run is the best answer since it can scale down to 0 when there is no traffic

upvoted 3 times

 **luciorifa** 1 year, 7 months ago

Selected Answer: B

Correct answer is B, Cloud Run is a serverless solution that, same as Cloud Functions, can be activated based on events.

upvoted 1 times

 **DaveNZ** 1 year, 7 months ago

Selected Answer: B

The majority vote here seems to be B so I'm posting this as a voting comment to make that visible. I am not entirely sure, but suspect that the "min-instances" issue is a red herring - that is, I understand why some comments raise this as being a reason why B may not be correct but I think maybe it's not the intention of the question to focus on that rather tricky aspect.

upvoted 1 times

 **mahesh0049** 1 year, 8 months ago

Option B is correct. After Business hours there will be no colleagues to work on that application then cloud run can scale to zero instances

upvoted 2 times

 **gcpengineer** 2 years, 2 months ago

Ans is A. as we seek the solution for internal users, Anthos makes vpn/interconnect to cloud Run, only from internal ip address

upvoted 3 times

You have experimented with Google Cloud using your own credit card and expensed the costs to your company. Your company wants to streamline the billing process and charge the costs of your projects to their monthly invoice. What should you do?

- A. Grant the financial team the IAM role of `Billing Account User` on the billing account linked to your credit card.
- B. Set up BigQuery billing export and grant your financial department IAM access to query the data.
- C. Create a ticket with Google Billing Support to ask them to send the invoice to your company.
- D. Change the billing account of your projects to the billing account of your company.

Correct Answer: D*Community vote distribution*

D (100%)

 **j_mrnr**  2 years, 7 months ago

1000% Ans D

upvoted 25 times

 **rsuresh27**  1 year, 5 months ago

Please do not overthink the question. The question does not mention anything about finance teams, so A cannot be correct. D is the only one that makes sense out of the remaining options.

upvoted 6 times

 **3arle**  1 month, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times

 **N_A** 5 months ago

With A. the financial team can only link the billing account linked to the credit card.

B. C. are wrong (no comment).

D. is the only correct answer even though it doesn't give the exact permission to grant in order to do this. I guess the financial team already has the Billing Project Manager role at a folder or organization level which would allow them to make the change.

upvoted 1 times

 **iamlearning2** 8 months, 2 weeks ago

Selected Answer: D

D is the answer

upvoted 1 times

 **anjanc** 9 months, 2 weeks ago

Selected Answer: D

d seems the answer

upvoted 1 times

 **Cornholio_LMC** 1 year ago

had this one today

upvoted 4 times

 **deadlydeb** 1 year, 3 months ago

Selected Answer: D

D D D D D

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

D is perfect

upvoted 1 times

 **Surat** 1 year, 9 months ago

D seems correct ans

upvoted 2 times

 **walkwolf3** 1 year, 10 months ago

Either A or D are incomplete solution.

Soution A

Grant financial team the IAM role of Project Owner or Project Billing Manager on your project, then let financial team change the billing account.

Solution B

Grant you IAM role of Billing Account Administrator or Billing Account User on company's project, then you can change the billing account.

I hope this question won't be in the exam, or has a more accurate answer in the exam.

upvoted 1 times

 **learn_GCP** 1 year ago

D -- is the answer

can't be A. you don't want to share your personnel credit card data with company billing account

upvoted 1 times

 **MahAli** 6 months ago

the company will let you link their account to your personal project?

upvoted 1 times

 **vamgcp** 2 years ago

correct answer D- Change the billing account of your projects to the billing account of your company

upvoted 2 times

 **arsh1916** 2 years, 4 months ago

D is correct

upvoted 2 times

 **sivre** 2 years, 6 months ago

Does the user have access to the company billing account? Not sure of D, I think more A.

To change the Cloud Billing account for a project, you need to be able to move a project from one Cloud Billing account to another. To accomplish this task, you need permissions adequate to unlink the project from the existing Cloud Billing account AND to link the project to the target Cloud Billing account.

Roles with adequate permissions to perform this task: Project Owner or Project Billing Manager on the project, AND Billing Account Administrator or Billing Account User for the target Cloud Billing account

https://cloud.google.com/billing/docs/how-to/modify-project#change_the_billing_account_for_a_project

upvoted 3 times

 **GoCloud** 2 years, 6 months ago

A would be ok if assigning Project Billing Manager IAM role to the finance team.

So the answer will be D.

upvoted 1 times

 **muneebarshad** 2 years, 3 months ago

"Billing Account User" Role Link projects to billing accounts. Since It does not say anything about user permissions , lets assume if user does not have sufficient permission then Granting the financial team the ""Billing Account User" " role will do the job

upvoted 1 times

 **wh1t4k3r** 1 year, 10 months ago

"Grant the financial team the IAM role of `Billing Account User` on the billing account linked to your credit card."

Giving billing user rights to the financial team to the billing account binded to your credit card wont give them any means to change it or maybe direct/export to the company's billing.

Billing user gives the following permissions on a billing account:

" This role has very restricted permissions, so you can grant it broadly, typically in combination with Project Creator. These two roles allow a user to create new projects linked to the billing account on which the role is granted."

upvoted 1 times

 **wh1t4k3r** 1 year, 10 months ago

Still, you do make a valid point regarding the access. How will i add my projects to the company's without access? Its not possible indeed.

That said, i would say that the best option is B (and when i say best i mean for the given options), since you can export billing data in bigquery:

<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>

upvoted 1 times

 **learn_GCP** 1 year ago

in B option we can get only billing / budget data, but question asks for streamlining of billing process, mean they want to pay bills using company invoice.

upvoted 1 times

 **GCP_Student1** 2 years, 6 months ago

D. Change the billing account of your projects to the billing account of your company.

upvoted 4 times

 **TAvenger** 2 years, 6 months ago

D. Change the billing account of your projects to the billing account of your company.

upvoted 2 times

 **lio123** 2 years, 7 months ago

yes it is D

upvoted 3 times

You are running a data warehouse on BigQuery. A partner company is offering a recommendation engine based on the data in your data warehouse. The partner company is also running their application on Google Cloud. They manage the resources in their own project, but they need access to the BigQuery dataset in your project. You want to provide the partner company with access to the dataset. What should you do?

- A. Create a Service Account in your own project, and grant this Service Account access to BigQuery in your project.
- B. Create a Service Account in your own project, and ask the partner to grant this Service Account access to BigQuery in their project.
- C. Ask the partner to create a Service Account in their project, and have them give the Service Account access to BigQuery in their project.
- D. Ask the partner to create a Service Account in their project, and grant their Service Account access to the BigQuery dataset in your project.

Correct Answer: D

Community vote distribution

D (100%)

- ✉ **GCP_Student1** Highly Voted 2 years, 6 months ago
- D. Ask the partner to create a Service Account in their project, and grant their Service Account access to the BigQuery dataset in your project.
upvoted 31 times
- ✉ **pondai** Highly Voted 2 years, 6 months ago
- BigQuery is in our project, so we need to create a service account and grant it access BigQuery role. That can make partner company to use this account to use it to access our project's BigQuery. So I vote A.
upvoted 13 times
- ✉ **tavva_prudhvi** 2 years, 6 months ago
- See, the ones who want our access needs to create a service account (in our case it's the partner company), then we give access to the service account with the user permissions. Clearly, D says the same thing!
upvoted 8 times
- ✉ **akshaychavan7** 1 year, 4 months ago
- Your understanding is bit wrong here, my friend!
upvoted 3 times
- ✉ **scanner2** Most Recent 1 month ago
- Selected Answer: D**
- Cross project access. Application in Project A want to access a service in project B.
1. Create a service account in project A.
 2. Give the required permission to access the services in project B.
- upvoted 1 times
- ✉ **N_A** 5 months ago
- Selected Answer: D**
- A. Useless if the private key of the Service Account is not shared with the partner (this would not be a good practice in terms of security)
B. Not possible.
C. Useless as the won't have access to the data in our data warehouse on BigQuery.
D. Is the correct answer and follow best practices.
upvoted 1 times
- ✉ **hiromi** 10 months, 3 weeks ago
- Selected Answer: D**
- Should be D
upvoted 2 times
- ✉ **Aninina** 11 months, 1 week ago
- Selected Answer: D**
- "Service accounts are both identities and resources. Because service accounts are identities, you can let a service account access resources in your project by granting it a role, just like you would for any other principal."
upvoted 1 times
- ✉ **Cornholio_LMC** 1 year ago
- had this one today
upvoted 2 times
- ✉ **AzureDP900** 1 year, 3 months ago
- D is right
upvoted 1 times
- ✉ **somenick** 1 year, 6 months ago

Selected Answer: D

<https://gtseries.medium.com/using-service-accounts-across-projects-in-gcp-cf9473fef8f0#:~:text=Go%20to%20the%20destination%20project,Voila!>

upvoted 2 times

✉ **mk1708** 1 year, 7 months ago

I thought it was A. But when I quickly did some research I found this:

"Service accounts are both identities and resources. Because service accounts are identities, you can let a service account access resources in your project by granting it a role, just like you would for any other principal."

Thus, the answer is D.

upvoted 4 times

✉ **wh1t4k3r** 1 year, 10 months ago

D per my understanding: if the need is to authenticate the application to access your dataset, it's the application's service account that will be provided during the authentication, so the service account is to be created at their side to run the application, not the other way around.

upvoted 4 times

✉ **wh1t4k3r** 1 year, 10 months ago

Another insight:

A is too broad. The question states: "provide the partner company with access to the dataset"

A states: "grant this Service Account access to BigQuery in your project"

I think D is a more granular option, given that A would give access to all datasets in your bigquery data warehousing.

upvoted 2 times

✉ **jabrrJ68w02ond1** 1 year, 10 months ago

How is it D? I want to give access to my BigQuery data, so I need to provide the ServiceAccount. I create it, put some decent predefined roles on it, and whenever I stop working with the other company, I either invalidate the JSON key of the SA or I simply delete the SA. For me, it is A.

upvoted 4 times

✉ **JelloMan** 1 year, 5 months ago

Right but how will the other project ever gain access to resources on YOUR project? Key thing here is that you're sharing ACROSS different projects so you need a bridge between them. D provides that bridge by connecting THEIR service account with YOUR resource (big query)

upvoted 2 times

✉ **AD_0525** 2 years, 3 months ago

Answer should be D, as the other company project needs access in your project.

upvoted 6 times

✉ **EmreA** 2 years, 5 months ago

It is A

upvoted 4 times

✉ **kopper2019** 2 years, 5 months ago

D is the answer

upvoted 6 times

✉ **TAvenger** 2 years, 6 months ago

I think it is "D" but I have not found such usecase when you share dataset with another organization via service account

upvoted 3 times

✉ **lio123** 2 years, 7 months ago

it is D

upvoted 4 times

Your web application has been running successfully on Cloud Run for Anthos. You want to evaluate an updated version of the application with a specific percentage of your production users (canary deployment). What should you do?

- A. Create a new service with the new version of the application. Split traffic between this version and the version that is currently running.
- B. Create a new revision with the new version of the application. Split traffic between this version and the version that is currently running.
- C. Create a new service with the new version of the application. Add an HTTP Load Balancer in front of both services.
- D. Create a new revision with the new version of the application. Add an HTTP Load Balancer in front of both revisions.

Correct Answer: A*Community vote distribution***B (100%)**

✉️  **crysk**  2 years, 7 months ago

In my opinion correct answer is B:

https://cloud.google.com/run/docs/rollouts-rollbacks-traffic-migration?utm_campaign=CDR_ahm_aap-severless_cloud-run-faq_&utm_source=external&utm_medium=web

Cloud Run can split traffic between revisions

upvoted 34 times

✉️  **TAvenger** 2 years, 6 months ago

The google doc link is incorrect. You need to specify CloudRun for Anthos

<https://cloud.google.com/kuberun/docs/rollouts-rollbacks-traffic-migration>

Anyway principles for CloudRun and CloundRun for Anthos are the same. Traffic can be split between multiple revisions.

The answer is "B"

upvoted 9 times

✉️  **GCP_Student1**  2 years, 6 months ago

B. Create a new revision with the new version of the application. Split traffic between this version and the version that is currently running.

upvoted 10 times

✉️  **scanner2**  1 month ago

Selected Answer: B

<https://cloud.google.com/anthos/run/docs/rollouts-rollbacks-traffic-migration>

upvoted 1 times

✉️  **kumar262639** 4 months, 3 weeks ago

which answer is correct. the one "Correct Answer" or the Community vote distribution winner ?

upvoted 1 times

✉️  **idontlikeme_3342** 1 month, 1 week ago

I recommend to you to go with the answer that has the most number of Upvotes

upvoted 1 times

✉️  **Emmanski08** 9 months, 1 week ago

Keyword - "Updated Version"

B. Create a new "revision"

upvoted 1 times

✉️  **AzFarid** 9 months, 2 weeks ago

B is ok

upvoted 1 times

✉️  **Untamables** 11 months, 2 weeks ago

Selected Answer: B

The latest Document

<https://cloud.google.com/anthos/run/docs/rollouts-rollbacks-traffic-migration>

upvoted 1 times

✉️  **nonsense** 11 months, 2 weeks ago

Selected Answer: B

B. Create a new revision

upvoted 1 times

✉️  **sylva91** 1 year ago

Selected Answer: B

B is the correct answer

upvoted 1 times

✉  **thimai** 1 year, 1 month ago

Selected Answer: B

i think B

upvoted 1 times

✉  **AzureDP900** 1 year, 3 months ago

B is correct for this scenario, there is no need to create new services

upvoted 2 times

✉  **pfabio** 1 year, 4 months ago

Selected Answer: B

Cloud Run for Anthos allows you to specify which revisions should receive traffic and to specify traffic percentages that are received by a revision

upvoted 1 times

✉  **HansKloss611** 1 year, 8 months ago

Selected Answer: B

If course B

upvoted 1 times

✉  **AkshayKalbhor** 2 years, 3 months ago

Answer should be B.

upvoted 5 times

✉  **RishiAbhishek** 2 years, 3 months ago

Canary deployments are a method of releasing software to a subset of users or servers. The plan is to deliver the update to a small selection of servers first, test it, and then roll it out to the other servers.

So it's a Revision or new update or a new version, not a service.

upvoted 5 times

✉  **nana1995** 2 years, 4 months ago

Its B

<https://cloud.google.com/run/docs/managing/revisions>

upvoted 2 times

✉  **Ashii** 2 years, 4 months ago

Its B, revision or version, not service.

upvoted 3 times

Your company developed a mobile game that is deployed on Google Cloud. Gamers are connecting to the game with their personal phones over the Internet. The game sends UDP packets to update the servers about the gamers' actions while they are playing in multiplayer mode. Your game backend can scale over multiple virtual machines (VMs), and you want to expose the VMs over a single IP address. What should you do?

- A. Configure an SSL Proxy load balancer in front of the application servers.
- B. Configure an Internal UDP load balancer in front of the application servers.
- C. Configure an External HTTP(s) load balancer in front of the application servers.
- D. Configure an External Network load balancer in front of the application servers.

Correct Answer: A

Reference:

<https://cloud.google.com/solutions/connecting-securely>

Community vote distribution

D (100%)

✉ **kopper2019** Highly Voted 2 years, 5 months ago

Answer is D, cell phones are sending UDP packets and the only that can receive that type of traffic is a External Network TCP/UDP
<https://cloud.google.com/load-balancing/docs/network>

upvoted 29 times

✉ **ashrafh** 2 years, 1 month ago

Google Cloud HTTP(S) Load Balancing is a global, proxy-based Layer 7 load balancer that enables you to run and scale your services worldwide behind a single external IP address. External HTTP(S) Load Balancing distributes HTTP and HTTPS traffic to backends hosted on Compute Engine and Google Kubernetes Engine (GKE).

<https://cloud.google.com/load-balancing/docs/https>

upvoted 2 times

✉ **ryumada** 1 year, 2 months ago

All the load balancer products in GCP give you a single IP address for the backend servers you registered to it.

Also, External HTTP(s) load balancer only support the port that used by HTTP which is the port 80 and HTTPS which is the port 443.

And Google Cloud external TCP/UDP Network Load Balancing is referred to as "Network Load Balancing" which supports UDP packets.

- <https://cloud.google.com/load-balancing/docs/load-balancing-overview#about>
- <https://cloud.google.com/load-balancing/docs/network>
- <https://cloud.google.com/load-balancing/docs/https>

upvoted 3 times

✉ **patashish** 1 year, 2 months ago

what you are trying to say ? What is your answer ? A B C D ?

upvoted 3 times

✉ **JH86** Highly Voted 2 years, 4 months ago

Answer is D. there are so many confusion here, from B,C or D. For myself im eliminating all options except B,D due to the traffic type. which leaves me with B or D. Then next the traffic source either external or internal which in this case is an external traffic from the internet, therefore my final answer is D.

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer>

upvoted 12 times

✉ **BobbyFlash** 1 year, 11 months ago

Following the diagram, there's no doubt about D. We have external clients connecting to our gaming service on google cloud that works using UDP traffic that results in using External Network Load Balancing. I feel that it's simple as it is. I also go with D.
upvoted 4 times

✉ **Captain1212** Most Recent 1 month ago

Selected Answer: D

for udp external load balancer, D is the correct answer

upvoted 1 times

✉ **CVGCP** 4 months ago

By elimination

A: SSL proxy LB is for TCP traffic not for UDP, eliminated

B: External LB is required, Eliminated

C: Http LB works at layer 7, here protocol is UDP, eliminated

D: Correct answer

upvoted 3 times

✉ **kumar262639** 4 months, 3 weeks ago

"Correct Answer" says A and community vote says D(100%)
which one is correct?

upvoted 1 times

✉ **PPP_D** 6 months, 2 weeks ago

Going with D
upvoted 1 times

✉ **Andoameda9** 7 months, 3 weeks ago

Selected Answer: D

Ans is D

upvoted 1 times

✉ **fragment137** 10 months, 1 week ago

The question tricked me. I saw UDP and immediately thought it was B.

The correct answer is D, as the LB needs to be External, and SSL\HTTPS are not the right load balancers for this application.
upvoted 2 times

✉ **abirroy** 1 year, 2 months ago

Selected Answer: D

External Network LB used for UDP

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

D seems correct..

upvoted 1 times

✉ **akshaychavan7** 1 year, 4 months ago

Selected Answer: D

I'm dead sure, it's D!

upvoted 1 times

✉ **somenick** 1 year, 6 months ago

Selected Answer: D

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer#lb-decision-tree>

upvoted 2 times

✉ **[Removed]** 1 year, 6 months ago

Selected Answer: D

D - Check <https://cloud.google.com/load-balancing/images/choose-lb.svg>

upvoted 1 times

✉ **Raz0r** 1 year, 8 months ago

Selected Answer: D

D, because:

<https://cloud.google.com/load-balancing/docs/network#:~:text=Google%20Cloud%20external-,TCP/UDP,-Network%20Load%20Balancing>
upvoted 2 times

✉ **exam_war** 1 year, 9 months ago

Answer is C. Players need to access through internet by HTTP(S) load balancing

upvoted 2 times

S00999 1 year, 3 months ago

It is not specified whether the app protocol is HTTP(S) or not, only that it is UDP packets. Internet is not limited to the http protocol.

Answer D

upvoted 1 times

✉ **[Removed]** 1 year, 9 months ago

D. Configure an External Network load balancer in front of the application servers.

"VM over single (external) IP address ->>> getting UDP packets through External LB

"

upvoted 3 times

✉ **Eben01** 1 year, 9 months ago

The key statement we need to watch out for ; "And you want to expose the Vms over a single IP address".

Google Cloud external TCP/UDP Network Load Balancing (after this referred to as Network Load Balancing) is a regional, "pass-through" load balancer. A network load balancer "distributes external traffic among virtual machine (VM) instances" in the same region.

You can configure a network load balancer for TCP, UDP, ESP, and ICMP traffic.

<https://cloud.google.com/load-balancing/docs/network>

upvoted 2 times

You are working for a hospital that stores its medical images in an on-premises data room. The hospital wants to use Cloud Storage for archival storage of these images. The hospital wants an automated process to upload any new medical images to Cloud Storage. You need to design and implement a solution. What should you do?

- A. Create a Pub/Sub topic, and enable a Cloud Storage trigger for the Pub/Sub topic. Create an application that sends all medical images to the Pub/Sub topic.
- B. Deploy a Dataflow job from the batch template, ↗Datastore to Cloud Storage. Schedule the batch job on the desired interval.
- C. Create a script that uses the gsutil command line interface to synchronize the on-premises storage with Cloud Storage. Schedule the script as a cron job.
- D. In the Cloud Console, go to Cloud Storage. Upload the relevant images to the appropriate bucket.

Correct Answer: C

Community vote distribution

C (90%)

10%

✉️ **TAvenger** Highly Voted 2 years, 6 months ago

From the question the key point is "upload ANY NEW medical images to Cloud Storage". So we are not interested in old images. That's why we need some trigger that will upload images. I think option "A" with PubSub is the best

upvoted 23 times

✉️ **dunhill** 2 years, 6 months ago

I am not sure but the question also mentions that "wants to use Cloud Storage for archival storage of these images". It can create an application that sends all medical images to storage and no need via PubSub?

upvoted 1 times

✉️ **pca2b** 2 years, 6 months ago

Pub/Sub will be good for all future files in in-prem data-storage.

we want to sync all + new, so a local on-prem server running a cron job (not GCE CronJob) to run gsutil to transfer files to Cloud Storage would work.

I vote for C

upvoted 5 times

✉️ **yvinisiupacuando** 2 years, 5 months ago

Sorry you are wrong, the question clearly indicates "The hospital wants an automated process to upload ANY NEW medical images to Cloud Storage." It does not mention the need to upload the original stock of images, only the new ones. Then I think the right answer must be A, as you said "Pub/sub will be good for all future files in prem data-storage" which is exactly what the questions is pointing to.

upvoted 6 times

✉️ **Priyanka109** 1 year ago

In option C we are using a cron job, not dragging and dropping the images.

upvoted 1 times

✉️ **gcpengineer** 2 years, 2 months ago

ans is C

upvoted 3 times

✉️ **GCP_Student1** Highly Voted 2 years, 6 months ago

C. Create a script that uses the gsutil command line interface to synchronize the on-premises storage with Cloud Storage. Schedule the script as a cron job.

upvoted 20 times

✉️ **cserra** 1 year, 1 month ago

Where does it say that the on-premises images are already digitized? and even if they are, where does it say that we also keep the old images?

I think the correct answer is "A"

upvoted 1 times

✉️ **theBestStudent** 1 year ago

Tell yo yourself how the images would end up in the pubsub first of all. Also usually the process is in the other way around for pubsub notifications: Once an object lands in GCS the pubsub is notified of it.

Option A makes totally nonsense. Check the flow again.

From the options the only one that "makes more sense" is Option C

upvoted 6 times

 **jkim1708** Most Recent 3 weeks, 1 day ago

Selected Answer: A

I am also for A. Any new data should be send to Cloud Storage. Yes you need to create an application. To send data to pubsub. But for possible migration to cloud you can use the existing setup

upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: C

C is the right answer as they want the automated process to upload any new medical image

upvoted 1 times

 **respawn** 1 month, 1 week ago

Selected Answer: C

C is correct.

A will not work because pub/sub is meant for service to service communication only:

https://cloud.google.com/pubsub/docs/overview#compare_service-to-service_and_service-to-client_communication

Yes C option will sync any new images from onprem to cloud.

upvoted 1 times

 **shreykul** 2 months, 2 weeks ago

Selected Answer: A

New images can use Pub/Sub

upvoted 1 times

 **Charumathi** 1 year ago

Selected Answer: C

C is the correct answer.

Keyword, they require cloud storage for archival and the want to automate the process to upload new medical image to cloud storage, hence we go for gsutil to copy on-prem images to cloud storage and automate the process via cron job. whereas Pub/Sub listens to the changes in the Cloud Storage bucket and triggers the pub/sub topic, which is not required.

upvoted 6 times

 **Naree** 3 months ago

I agree. The requirement is for both history images and future images. So I go with Option "C".

upvoted 1 times

 **zolthar_z** 1 year, 2 months ago

Selected Answer: C

The Hospital wants Cloud storage for archival of old images and also sync the new images, for this logic the answer is C

upvoted 2 times

 **AzureDP900** 1 year, 3 months ago

C is correct

upvoted 1 times

 **jblima** 1 year, 5 months ago

Selected Answer: C

C is correct.

upvoted 1 times

 **wael_tn** 1 year, 5 months ago

Selected Answer: C

Discarding (A) because "Cloud Storage trigger". So for option A the triggering event should be making a change in Cloud Storage, while in the real use case, the triggering should be adding a new medical image to the "on-premises data room"

upvoted 4 times

 **theBestStudent** 1 year ago

Correct, and on top of that, how are they suppose to connect the on premises to the cloud. Nothing is mentioned. But either way, option A does not make sense at all (as you already explained too). Correct option: C

upvoted 1 times

 **somenick** 1 year, 6 months ago

Selected Answer: C

I would go with C

Not A, don't think you can send image files to Pub/Sub. Technically you can do so by converting image to some binary text, but then we don't know the size of the image and there is a limitation on message size. Not recommended.

Not B – there is only this template "Datastore to Cloud Storage Text", as the name implies it is for text,

<https://cloud.google.com/dataflow/docs/guides/templates/provided-batch#datastore-to-cloud-storage-text>, and it reads from datastore which is definitely not where the medical images are stored, from the question "... stores its medical images in an on-premises data room".

Not D – it's not automated

upvoted 3 times

 **lilapause** 1 year, 8 months ago

Selected Answer: C

Pub/Sub could make sense if you ignore the size limit. But the way it is described with the storage trigger would just not be working.
upvoted 1 times

 **exam_war** 1 year, 9 months ago

A is correct. The key is automation whenever there is a new image, it needs to upload to cloud storage. Only pub/sub can make the automation work.
upvoted 1 times

 **[Removed]** 1 year, 9 months ago

I think A(once the hospital receive the new images Cloud pub/sub will act on it) and C(creating a script with rsync command <https://stackoverflow.com/questions/37662416/how-to-sync-a-local-folder-with-a-folder-in-a-google-cloud-platform-bucket>) must be the Correct options
upvoted 1 times

 **RealEL40** 1 year, 9 months ago

" any new medical images" so "A" using Pub\Sub
upvoted 1 times

 **wh1t4k3r** 1 year, 10 months ago

I'll go with A for he following reason:
Yes, the question says that they use onprem to store images and they want those images in gcloud storage from now on, BUT the solution for automation is requested for uploading new images only: "The hospital wants an automated process to upload any new medical images to Cloud Storage"
Using sync (A) for new images implies that you will continue to use your onprem and keep synchronizing it fover... Sync just once for the old images, new images go directly to gcloud via pubsub, and eventually get rid of the onprem.
upvoted 2 times

 **wh1t4k3r** 1 year, 10 months ago

Correcting: "Using sync (C) for new images implies that you will continue to use your onprem and keep synchronizing it fover... Sync just once for the old images, new images go directly to gcloud via pubsub, and eventually get rid of the onprem."
upvoted 1 times

Your auditor wants to view your organization's use of data in Google Cloud. The auditor is most interested in auditing who accessed data in Cloud Storage buckets. You need to help the auditor access the data they need. What should you do?

- A. Turn on Data Access Logs for the buckets they want to audit, and then build a query in the log viewer that filters on Cloud Storage.
- B. Assign the appropriate permissions, and then create a Data Studio report on Admin Activity Audit Logs.
- C. Assign the appropriate permissions, and then use Cloud Monitoring to review metrics.
- D. Use the export logs API to provide the Admin Activity Audit Logs in the format they want.

Correct Answer: D

Reference:

<https://cloud.google.com/storage/docs/audit-logging>

Community vote distribution

A (100%)

✉  **iri_gcp**  2 years, 6 months ago

It should be A.

Data access log are not enabled by default due to the fact that it incurs costs.

So you need to enable it first.

And then you can filter it in the log viewer

upvoted 39 times

✉  **GCP_Student1**  2 years, 6 months ago

A. Turn on Data Access Logs for the buckets they want to audit, and then build a query in the log viewer that filters on Cloud Storage.

upvoted 11 times

✉  **NoCrapEva**  1 month ago

IF Data Access Logs had ALREADY been enabled, then option B would be a good answer

Reason - (1) best practice for cloud auditing - enable Admin Activity audit logs, then set IAM permissions

(ref: <https://cloud.google.com/logging/docs/audit/best-practices>)

and (2) Create a Data Studio (now renamed to Looker) report on Admin Activity Audit Logs

(ref: <https://cloud.google.com/looker/docs/looker-core-audit-logging>)

But you cannot assume from the question that Data Access Logs are enabled (NB: they are NOT by default)

upvoted 1 times

✉  **Captain1212** 1 month ago

Selected Answer: A

A is the right answer as first we need to turn on the data access logs

upvoted 1 times

✉  **anolive** 11 months ago

I have doubts about the answer A, the auditor wants to see the audit logs, and in this answer it is not explicit if he will be allowed to see it.

upvoted 1 times

✉  **Charumathi** 1 year ago

Selected Answer: A

A is the correct answer,

Since the auditor wants to know who accessed the cloud storage data, we need data access logs for cloud storage.

Types of audit logs

Cloud Audit Logs provides the following audit logs for each Cloud project, folder, and organization:

Admin Activity audit logs

Data Access audit logs

System Event audit logs

Policy Denied audit logs

***Data Access audit logs contain API calls that read the configuration or metadata of resources, as well as user-driven API calls that create, modify, or read user-provided resource data.

<https://cloud.google.com/logging/docs/audit#types>

upvoted 1 times

✉  **AzureDP900** 1 year, 3 months ago

A is right

upvoted 1 times

✉  **Jman007** 1 year, 3 months ago

Selected Answer: A

question says auditor is most interested in who accessED data in Cloud Storage. im not sure how auditing is done for those who answered A but this means they want the logs for past users who accessed the data from a sepecified time. Turning on the feature now is kind of too late. poorly written question and answers. No point in an auditor coming in and giving the company all the exact questions they are going to ask and come back and ask them in a few months time. A seems like the better choices though

upvoted 2 times

akshaychavan7 1 year, 4 months ago

If it's A then how will we assign the permission for the auditor to view the logs?

I had chosen option A on the first place, but later changed it considering that the auditor won't have the access to view the logs.

upvoted 1 times

peugeotdude 1 year, 5 months ago

Selected Answer: A

Based on how I read the question-

We want Data Access log, not Admin Activity Audit Logs.

upvoted 1 times

somenick 1 year, 6 months ago

Selected Answer: A

Data access log are not enabled by default due to the fact that it incurs costs.

So you need to enable it first.

And then you can filter it in the log viewer

upvoted 1 times

[Removed] 1 year, 6 months ago

<https://cloud.google.com/logging/docs/audit#data-access>

Cloud Storage: When Cloud Storage usage logs are enabled, Cloud Storage writes usage data to the Cloud Storage bucket, which generates Data Access audit logs for the bucket. The generated Data Access audit log has its caller identity redacted.

upvoted 1 times

DaveNZ 1 year, 7 months ago

Selected Answer: A

The majority vote here is A, despite some confusion around the wording of the question. I tend to agree because it's the solution that most closely reflects the requirements of the question (buckets, cloud storage).

upvoted 1 times

wh1t4k3r 1 year, 10 months ago

A. I could not find a way to enable audit logs in specific buckets, only on the whole storage level:
<https://cloud.google.com/logging/docs/audit/services>

B. Admin activity audit logs cover admin actions, such as metada or config changes:
<https://cloud.google.com/logging/docs/audit#admin-activity>

C. Cloud monitoring is not for auditing: <https://cloud.google.com/monitoring>

D. Again, Admin Activity Audit Logs should not be used to audit data access, specially from bukets.

My conclusion: all these answers are wrong. My assumption: A is badly written. Specific buckets were not to be mentioned. I Vote A, but i think this Q&A is messed up. Maybe a correction? or deletion.

upvoted 3 times

MarcoDipa 1 year, 10 months ago

Actually, there is a different service named User Logs that permits to focus on a single bucket.

Refer to google page:

<https://cloud.google.com/storage/docs/access-logs>

Usage logs provide information for all of the requests made on a specified bucket

upvoted 1 times

obeythefist 1 year, 7 months ago

The question just says "buckets" and hints that the audit should cover all org data, so I don't think there is any need to overanalyse, you are correct in choosing A

upvoted 1 times

ericyev 1 year, 10 months ago

I choose D. reason is here: Cloud Audit Logs generates the following audit logs for operations in Cloud Storage:

Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.

Data Access logs: Entries for operations that modify objects or read a project, bucket, or object. There are several sub-types of data access logs:

ADMIN_READ: Entries for operations that read the configuration or metadata of a project, bucket, or object.

DATA_READ: Entries for operations that read an object.

DATA_WRITE: Entries for operations that create or modify an object.

upvoted 4 times

kimharsh 1 year, 10 months ago

Also A because it's the only one that mention DATA ACCESS LOGS, which is the one that Logs objects access , t

Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.

Data Access logs: Entries for operations that modify objects or read a project, bucket, or object. There are several sub-types of data access logs:

ADMIN_READ: Entries for operations that read the configuration or metadata of a project, bucket, or object.

DATA_READ: Entries for operations that read an object.

DATA_WRITE: Entries for operations that create or modify an object.

<https://cloud.google.com/storage/docs/audit-logging>

upvoted 1 times

 **jackwillis** 1 year, 12 months ago

Question is about user activities log not about Data Access log.

upvoted 1 times

You received a JSON file that contained a private key of a Service Account in order to get access to several resources in a Google Cloud project. You downloaded and installed the Cloud SDK and want to use this private key for authentication and authorization when performing gcloud commands. What should you do?

- A. Use the command gcloud auth login and point it to the private key.
- B. Use the command gcloud auth activate-service-account and point it to the private key.
- C. Place the private key file in the installation directory of the Cloud SDK and rename it to `credentials.json`.
- D. Place the private key file in your home directory and rename it to `GOOGLE_APPLICATION_CREDENTIALS`.

Correct Answer: B

Reference:

<https://cloud.google.com/sdk/docs/authorizing>

Community vote distribution

B (100%)

✉  **GCP_Student1**  2 years, 6 months ago

- B. Use the command gcloud auth activate-service-account and point it to the private key.

Authorizing with a service account

gcloud auth activate-service-account authorizes access using a service account. As with gcloud init and gcloud auth login, this command saves the service account credentials to the local system on successful completion and sets the specified account as the active account in your Cloud SDK configuration.

https://cloud.google.com/sdk/docs/authorizing#authorizing_with_a_service_account

upvoted 42 times

✉  **TAvenger**  2 years, 6 months ago

- B.
gcloud auth activate-service-account --help

NAME)

gcloud auth activate-service-account - authorize access to Google Cloud Platform with a service account

SYNOPSIS

gcloud auth activate-service-account [ACCOUNT] --key-file=KEY_FILE
[-password-file=PASSWORD_FILE | --prompt-for-password]
[GCLOUD_WIDE_FLAG ...]

DESCRIPTION

To allow gcloud (and other tools in Cloud SDK) to use service account credentials to make requests, use this command to import these credentials from a file that contains a private authorization key, and activate them for use in gcloud. gcloud auth activate-service-account serves the same function as gcloud auth login but uses a service account rather than Google user credentials.

upvoted 19 times

✉  **eBookz** 7 months, 3 weeks ago

See below information suggesting that service account can be used to authorize with the command "gcloud auth login". Not sure if this is a recent update:

"The gcloud auth login command authorizes access by using workload identity federation, which provides access to external workloads, or by using a service account key."

"To activate your service account, run gcloud auth login with the --cred-file flag:

gcloud auth login --cred-file=CONFIGURATION_OR_KEY_FILE
Replace CONFIGURATION_OR_KEY_FILE with the path to one of the following:

A credential configuration file for workload identity federation
A service account key file"

https://cloud.google.com/sdk/docs/authorizing#authorize_with_a_service_account

upvoted 1 times

✉  **itsimranmalik** 1 month ago

As per google - gcloud auth activate-service-account serves the same function as gcloud auth login but uses a service account rather than Google user credentials.

Ref: <https://cloud.google.com/sdk/gcloud/reference/auth/activate-service-account>

upvoted 1 times

✉ **scanner2** Most Recent 1 month ago

Selected Answer: B

<https://cloud.google.com/sdk/gcloud/reference/auth/activate-service-account>

upvoted 1 times

✉ **N_A** 5 months ago

D. This method is for application default credentials. See: <https://cloud.google.com/docs/authentication/application-default-credentials>

A. This method is to obtain credentials for a user account.

C. This does nothing. Useless.

B. Is the correct answer. See: <https://cloud.google.com/sdk/gcloud/reference/auth/activate-service-account>

upvoted 2 times

✉ **Abhi00754** 6 months, 1 week ago

Selected Answer: B

<https://cloud.google.com/sdk/gcloud/reference/auth/activate-service-account>

B

upvoted 1 times

✉ **abirroy** 1 year, 2 months ago

Selected Answer: B

Use the command gcloud auth activate-service-account and point it to the private key

upvoted 1 times

✉ **skrjha20** 1 year, 3 months ago

ANswer is B

To activate your service account, run gcloud auth activate-service-account:

gcloud auth activate-service-account [ACCOUNT] --key-file=[KEY_FILE]

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

I will go with B

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

B is right Please refer

<https://cloud.google.com/storage/docs/authentication>

upvoted 1 times

✉ **Rukman** 1 year, 6 months ago

Selected Answer: B

Ans: B

upvoted 1 times

✉ **KunK** 2 years ago

B, really straightforward

upvoted 5 times

✉ **ravixkumar** 2 years, 7 months ago

Ans : B

gcloud auth activate-service-account --key-file=/test-service-account.json

upvoted 11 times

You are working with a Cloud SQL MySQL database at your company. You need to retain a month-end copy of the database for three years for audit purposes.

What should you do?

- A. Set up an export job for the first of the month. Write the export file to an Archive class Cloud Storage bucket.
- B. Save the automatic first-of-the-month backup for three years. Store the backup file in an Archive class Cloud Storage bucket.
- C. Set up an on-demand backup for the first of the month. Write the backup to an Archive class Cloud Storage bucket.
- D. Convert the automatic first-of-the-month backup to an export file. Write the export file to a Coldline class Cloud Storage bucket.

Correct Answer: B

Community vote distribution

A (100%)

 **TAvenger**  2 years, 6 months ago

<https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

not B: Automatic backups are made EVERY SINGLE DAY. You can set only the number of backups up to 365. Also you cannot choose your Archival storage as destination

not C: You cannot setup "on-demand" backup. User would have to make backups manually every month. Also you cannot choose your Archival storage as destination

not D: You cannot convert backup to export file. Also Coldline class is less cost-effective than Archival class.

The only option left is "A"

You can set up your job with any date/time schedule. You can export file to any storage with any storage class.

upvoted 46 times

 **djgodzilla** 2 years, 3 months ago

from the same link :

Can I export a backup?

No, you can't export a backup. You can only export instance data. See Exporting data from Cloud SQL to a dump in Cloud storage.

upvoted 5 times

 **JieHeng**  2 years, 3 months ago

First need to understand backup vs export, two different concepts. - <https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

A – yes, you can export data from Cloud SQL to Cloud Storage- <https://cloud.google.com/sql/docs/mysql/import-export/exporting#cloud-sql>

Not B, C, D – be it automatic or on-demand backup, according to the doc "No, you can't export a backup. You can only export instance data."

upvoted 16 times

 **santhush**  3 weeks, 3 days ago

<https://www.exam-answer.com/retain-month-end-copy-cloud-sql-mysql-database-three-years> B is the correct answer.. I am not sure why people are posting wrong answers here.

upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: A

A is the correct answer, as in you can export as per your requirement and then move it to archive class , but in b,c,d you can't do that

upvoted 1 times

 **akkinepallyn** 5 months ago

Option B is the best choice because it allows you to leverage the automatic first-of-the-month backup feature that is provided by Cloud SQL. Cloud SQL provides automated backups that can be configured to run at specific times, including the first of the month. By retaining the first-of-the-month backup for three years, you can be sure that you have a complete copy of the database for that month.

upvoted 1 times

 **Vismaya** 7 months, 1 week ago

Answer A

upvoted 1 times

 **researched_answer_boi** 8 months, 3 weeks ago

Answer A is the correct one according to "<https://cloud.google.com/blog/topics/developers-practitioners/scheduling-cloud-sql-exports-using-cloud-functions-and-cloud-scheduler>" and "https://cloud.google.com/sql/docs/mysql/backup-recovery/backups#backups_versus_exports".

upvoted 1 times

 **Kopy** 11 months ago

Selected Answer: A

So Answer is A.

You can't export back up. Very clear.

https://cloud.google.com/sql/docs/mysql/backup-recovery/backups#can_i_export_a_backup

upvoted 1 times

✉ **Charumathi** 1 year ago

Selected Answer: A

A is correct answer,

Export the SQL month end data as a CSV file to cloud storage bucket, and move the data to Archival Storages for 3 years for audit purpose.

<https://cloud.google.com/sql/docs/mysql/import-export?authuser=1>

upvoted 2 times

✉ **learn_GCP** 1 year ago

Selected Answer: A

A. Although Cloud SQL doesn't provide a built-in way to automate database exports, you can build your own automation tool using several Google Cloud components.

https://cloud.google.com/sql/docs/mysql/import-export#automating_export_operations

upvoted 1 times

✉ **Cornholio_LMC** 1 year ago

had this question today

upvoted 4 times

✉ **butki** 1 year ago

B is correct

upvoted 1 times

✉ **snkhatri** 1 year, 1 month ago

A seems right to me key word: "month-end copy"

upvoted 2 times

✉ **patashish** 1 year, 2 months ago

Correct Ans - C

On demand backup

You can create a backup at any time (Here we need backup a month-end copy of the database for three years). You can create on-demand backups for any instance, whether the instance has automatic backups enabled or not.

Reason :

1) you can't export a backup. You can only export instance data so export option A is out from answer.

Backups encompass the entire database. Exports can select specific contents.

As per question You need to retain a month-end copy of the database not specific contents.

2) Automated backups are taken daily, within a 4-hour backup window. Up to seven most recent backups are retained, by default.
*Cost** to store all backups ..

3) Option D not applicable **Coldline class** *Cost*

upvoted 4 times

✉ **mk1471** 1 year, 8 months ago

why not D? it is the only one that doesn't store it as Archive class, and since it is for Audit purposes this can't be used as Archive allows LESS than one access per year.

upvoted 2 times

✉ **look1** 1 year, 10 months ago

Selected Answer: A

I would go with export not backup. Question stated that this is copy. In addition restore of three years old backup might be not possible in newer version. In addition you can retain max 365 backups.

upvoted 4 times

✉ **jackwillis** 1 year, 12 months ago

As question itself states that we need month end backup, automatic backup runs daily which is required as per the question.

<https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

upvoted 3 times

You are monitoring an application and receive user feedback that a specific error is spiking. You notice that the error is caused by a Service Account having insufficient permissions. You are able to solve the problem but want to be notified if the problem recurs. What should you do?

- A. In the Log Viewer, filter the logs on severity 'Error' and the name of the Service Account.
- B. Create a sink to BigQuery to export all the logs. Create a Data Studio dashboard on the exported logs.
- C. Create a custom log-based metric for the specific error to be used in an Alerting Policy.
- D. Grant Project Owner access to the Service Account.

Correct Answer: A

Reference:

<https://cloud.google.com/logging/docs/view/advanced-queries>

Community vote distribution

C (100%)

-  **GCP_Student1** Highly Voted 2 years, 6 months ago
C. Create a custom log-based metrics for the specific error to be used in an Alerting Policy.
upvoted 23 times
-  **greatsam321** Highly Voted 2 years, 7 months ago
C seems to be the right answer.
upvoted 11 times
-  **Captain1212** Most Recent 1 month ago
Selected Answer: C
User wants to check the if problem recurs, that can be only possible by Alert, C is the correct option
upvoted 1 times
-  **_F4LLEN_** 5 months, 3 weeks ago
C. The keyword here is "want to be notified" that means an alert.
upvoted 2 times
-  **ranjidk** 10 months, 2 weeks ago
Can someone send me the full dumps with correct answers to this mail :- ranjith.deepak93@gmail.com
upvoted 1 times
-  **Charumathi** 1 year ago
Selected Answer: C
C is the correct answer,
Since the problem is resolved, We need to monitor if the error recurs, hence we create a custom log based metrics to monitor only the particular service account.
upvoted 2 times
-  **snkhatri** 1 year, 1 month ago
C as Keyword "want to be notified if the problem recurs"
upvoted 1 times
-  **AzureDP900** 1 year, 3 months ago
C right
upvoted 1 times
-  **AzureDP900** 1 year, 3 months ago
C is correct.
upvoted 1 times
-  **PAUGURU** 1 year, 5 months ago
Selected Answer: C
C - the only answer that outputs a notification
upvoted 3 times
-  **RazOr** 1 year, 8 months ago
Selected Answer: C
"C" is right, the only answer that includes setting up an alert!
upvoted 3 times
-  **Wolf13ts** 1 year, 8 months ago
Selected Answer: C

upvoted 1 times

 **alaahakim** 1 year, 10 months ago

Ans: C

upvoted 2 times

 **AD_0525** 2 years, 3 months ago

You want to be alerted next time, so only option C meets that criteria.

upvoted 4 times

 **Enzo** 2 years, 4 months ago

C is correct : You are able to solve the problem but want to be notified if the problem recurs.

upvoted 4 times

 **arsh1916** 2 years, 4 months ago

A is correct

upvoted 1 times

 **kopper2019** 2 years, 5 months ago

You are managing a project for the Business Intelligence (BI) department in your company. A data pipeline ingests data into BigQuery via streaming. You want the users in the BI department to be able to run the custom SQL queries against the latest data in BigQuery. What should you do?

- A. Create a Data Studio dashboard that uses the related BigQuery tables as a source and give the BI team view access to the Data Studio dashboard.
- B. Create a Service Account for the BI team and distribute a new private key to each member of the BI team.
- C. Use Cloud Scheduler to schedule a batch Dataflow job to copy the data from BigQuery to the BI team's internal data warehouse.
- D. Assign the IAM role of BigQuery User to a Google Group that contains the members of the BI team.

it's A

upvoted 2 times

 **mj98** 2 years, 5 months ago

I think it's D. Can anyone confirm?

upvoted 2 times

 **kopper2019** 2 years, 5 months ago

Question 178

upvoted 1 times

 **Petza** 2 years, 5 months ago

A. "Every time the dashboard is refreshed, it pulls new data from the view, which in turn dynamically reflects the latest data in BigQuery". Data Science on the Google Cloud Platform: Implementing End-to-End Real-Time (C)

upvoted 2 times

 **lxgywil** 2 years, 5 months ago

The question is saying "want to be notified if the problem recurs", I don't see how A meets that requirement.

upvoted 1 times

You are developing a financial trading application that will be used globally. Data is stored and queried using a relational structure, and clients from all over the world should get the exact identical state of the data. The application will be deployed in multiple regions to provide the lowest latency to end users. You need to select a storage option for the application data while minimizing latency. What should you do?

- A. Use Cloud Bigtable for data storage.
- B. Use Cloud SQL for data storage.
- C. Use Cloud Spanner for data storage.
- D. Use Firestore for data storage.

Correct Answer: C

Reference:

<https://cloud.google.com/solutions/best-practices-compute-engine-region-selection>*Community vote distribution*C (94%) 6%**JieHeng**  2 years, 3 months ago

C, Cloud Spanner, keywords are globally, relational structure and lastly "clients from all over the world should get the exact identical state of the data" which implies strong consistency is needed.

upvoted 15 times

Captain1212  1 month ago

Selected Answer: C
question demands, exact state of data and minimum latency to users , for this cloud spanner is the only option

upvoted 1 times

sabrinakloud 5 months, 3 weeks ago

Selected Answer: C
financial trading application
relational structure
multiple regions
upvoted 1 times

ashtonez 7 months ago

Selected Answer: C
C, always you need to select BBDD check for data analysis bigquery, something very big or fast bigtable, something with HA cloud sql, and something globally available cloud spanner, the key here is globaly available
upvoted 3 times

Charumathi 1 year ago

Selected Answer: C
C is the correct answer,
Keywords, Financial data (large data) used globally, data stored and queried using relational structure (SQL), clients should get exact identical copies(Strong Consistency), Multiple region, low latency to end user, select storage option to minimize latency.
upvoted 3 times

Charumathi 1 year ago

Spanner powers business-critical applications in retail, financial services, gaming, media and entertainment, technology, healthcare and more.

Use cases for Cloud Spanner
<https://www.youtube.com/watch?v=1b4flZwAQfM&t=1s>
upvoted 1 times

ale_brd_ 1 year ago

Selected Answer: C
it's C 100%
Guys come on, it's a pretty straight forward scenario.
if you have the keywords "relational DB" and the word "Globally" in a sentence always go for Cloud Spanner.
upvoted 3 times

learn_GCP 1 year ago

Selected Answer: C
C. is the answer
upvoted 1 times

sri333 1 year ago

With a Bigtable you can have a distributed database and with a

upvoted 1 times

 **tony_2023** 8 months, 4 weeks ago

BigTable is not a relational database. Everything else is true for it but it a noSQL non Relational Database.

upvoted 1 times

 **zellck** 1 year ago

Selected Answer: C

C is the answer.

Cloud Spanner is a global relational database.

upvoted 1 times

 **Mr_MIXER007** 1 year, 1 month ago

Selected Answer: C

Should be the correct answer.

upvoted 1 times

 **GayuSundar** 1 year, 1 month ago

C, Cloud Spanner. Globally and trading (tend to receive 1000s Records per second) are key here.

upvoted 1 times

 **VietmanOfficiel** 1 year, 1 month ago

Selected Answer: B

Why not CLoud Sql with replicas in multiple region to serve "global" ?

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

This is straight forward question, Answer is C.

upvoted 1 times

 **PAUGURU** 1 year, 5 months ago

Selected Answer: C

C is correct - Relational and Global

upvoted 1 times

 **alaahakim** 1 year, 10 months ago

Choose - C

upvoted 2 times

 **associatecloudexamuser** 2 years, 2 months ago

Yes. C is right answer

upvoted 3 times

 **Fidget_** 2 years, 5 months ago

C) Spanner - Global, low latency, relational

upvoted 4 times

You are about to deploy a new Enterprise Resource Planning (ERP) system on Google Cloud. The application holds the full database in-memory for fast data access, and you need to configure the most appropriate resources on Google Cloud for this application. What should you do?

- A. Provision preemptible Compute Engine instances.
- B. Provision Compute Engine instances with GPUs attached.
- C. Provision Compute Engine instances with local SSDs attached.
- D. Provision Compute Engine instances with M1 machine type.

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/disks/local-ssd>

Community vote distribution

D (100%)

Rightsaidfred Highly Voted 2 years, 5 months ago

Yes D, M1 Machine types for ERP i.e. SAP-HANA:
<https://cloud.google.com/compute/docs/machine-types>

upvoted 27 times

Captain1212 Most Recent 1 month ago**Selected Answer: D**

D is the right answer as M1 one are best for the databases i.e SAP hana
upvoted 1 times

Bobbybash 7 months, 3 weeks ago

C. Provision Compute Engine instances with local SSDs attached.

The best option for an ERP system that holds the full database in-memory for fast data access is to provision Compute Engine instances with local SSDs attached. Local SSDs offer high input/output operations per second (IOPS) and low latency, which can significantly improve the performance of in-memory databases. Preemptible Compute Engine instances are designed for short-lived and fault-tolerant workloads and are not recommended for a critical system like an ERP. GPUs are typically used for specialized compute-intensive workloads like machine learning and deep learning. M1 machine type is a general-purpose machine type and may not provide enough performance for an in-memory database.

upvoted 3 times

Charumathi 1 year ago**Selected Answer: D**

D is the correct answer,
M1 machine series
Medium in-memory databases such as SAP HANA
Tasks that require intensive use of memory with higher memory-to-vCPU ratios than the general-purpose high-memory machine types.
In-memory databases and in-memory analytics, business warehousing (BW) workloads, genomics analysis, SQL analysis services.
Microsoft SQL Server and similar databases.
upvoted 2 times

snkhatri 1 year, 1 month ago

D, keyword "Full database in-memory"
upvoted 1 times

ryumada 1 year, 2 months ago**Selected Answer: D**

Vote for D as the right answer. M1 machine type is the one of two Memory-Optimized machine types in GCP.
<https://cloud.google.com/compute/docs/machine-types>

upvoted 2 times

ryumada 1 year, 2 months ago

Read this also to see the difference of the two.

<https://cloud.google.com/compute/docs/memory-optimized-machines>
upvoted 1 times

abirroy 1 year, 2 months ago**Selected Answer: D**

D: M1 Machine types for ERP i.e. SAP-HANA

Medium-large in-memory databases such as SAP HANA

In-memory databases and in-memory analytics
Microsoft SQL Server and similar databases
upvoted 2 times

✉ **AzureDP900** 1 year, 3 months ago

D is right choice, when answer selected by author is C doesn't make any sense. User also need to understand the services well before attempting exam.
upvoted 1 times

✉ **akshaychavan7** 1 year, 4 months ago

I chose option D, because the first three didn't make any sense :D
upvoted 1 times

✉ **Raz0r** 1 year, 8 months ago

D!!!
The "M1" VM type is right, it offers between 1.4TB and 3.75TB of RAM.
upvoted 3 times

✉ **[Removed]** 1 year, 9 months ago

<https://cloud.google.com/compute/docs/machine-types#:~:text=databases%20such%20as-,SAP%20HANA,-In%2Dmemory%20databases>
<https://www.sap.com/india/products/hana.html#:~:text=is%20SAP%20HANA-,in%2Dmemory,-database%3F>
upvoted 1 times

✉ **jaffarali** 1 year, 10 months ago

Selected Answer: D
Correct Answer is D
upvoted 1 times

✉ **alaahakim** 1 year, 10 months ago

D is the Answer
upvoted 1 times

✉ **maggieli** 1 year, 11 months ago

"m1-megamem-96" can attach local SSD. Correct is D.
upvoted 1 times

✉ **jackdbd** 2 years, 2 months ago

Note that VM instances m1-megamem-96 are both from the M1 family AND can have local SSDs attached to them.
https://cloud.google.com/compute/docs/memory-optimized-machines#m1_vms
upvoted 2 times

✉ **associatecloudexamuser** 2 years, 2 months ago

Answer is D.
Applications of Memory optimized VMs are,
1. Medium-large in-memory databases such as SAP HANA
2. In-memory databases and in-memory analytics
3. Microsoft SQL Server and similar databases
upvoted 4 times

✉ **JieHeng** 2 years, 3 months ago

D, "The application holds the full database in-memory for fast data access", so it'll be more appropriate to use memory-optimized machine types - https://cloud.google.com/compute/docs/machine-types#m1_machine_types
upvoted 4 times

You have developed an application that consists of multiple microservices, with each microservice packaged in its own Docker container image. You want to deploy the entire application on Google Kubernetes Engine so that each microservice can be scaled individually. What should you do?

- A. Create and deploy a Custom Resource Definition per microservice.
- B. Create and deploy a Docker Compose File.
- C. Create and deploy a Job per microservice.
- D. Create and deploy a Deployment per microservice.

Correct Answer: D*Community vote distribution*

D (100%)

✉  **obeythefist** Highly Voted 1 year, 7 months ago

I was a little unsure about this question, here's how I understand why D is the best answer

- A. Custom Resource Definition... we have docker containers already, which is an established kind of resource for Kubernetes. We don't need to create a whole new type of resource, so this is wrong.
- B. Docker Compose is a wholly different tool from Kubernetes.
- C. A Kubernetes job describes a specific "task" which involves a bunch of pods and things. It makes no sense to have one job per microservice, a "Job" would be a bunch of different microservices executing together.
- D. is the leftover, correct answer. You can add scaling to each Deployment, an important aspect of the question.

upvoted 19 times

✉  **akshaychavan7** 1 year, 4 months ago

Thanks for your insights! Makes sense.

upvoted 2 times

✉  **Kollipara** Highly Voted 2 years, 5 months ago

D is the correct answer

upvoted 19 times

✉  **Captain1212** Most Recent 1 month ago

Selected Answer: D

D is the correct answer

upvoted 1 times

✉  **scanner2** 1 month ago

Selected Answer: D

D is the correct answer.

upvoted 1 times

✉  **Captain1212** 1 month ago

Selected Answer: D

D seems more correct , as in A , we already have the own docker container image , so no need to create ,
b is completely differnet tool,
c is also of no use

upvoted 1 times

✉  **snkhatri** 1 year, 1 month ago

D, keyword "each microservice can be scaled individually"!

upvoted 1 times

✉  **abirroy** 1 year, 2 months ago

Selected Answer: D

D is the correct answer

upvoted 1 times

✉  **AzureDP900** 1 year, 3 months ago

D is the best answer among other choices.

upvoted 1 times

✉  **Raz0r** 1 year, 8 months ago

D is right!

It's one of Googles main ideas to distribute a complex system into microservices. They do it as well and encourage customers to do the same.

upvoted 2 times

✉  **alaahakim** 1 year, 10 months ago

Selected Answer: D

D is the Answer
upvoted 2 times

 **alaahakim** 1 year, 10 months ago

Ans: D
upvoted 2 times

 **associatecloudexamuser** 2 years, 2 months ago

Yes. D is correct. Can deploy each service through
kubectl apply -f <deployment_config.yaml>
upvoted 5 times

 **pacman_user** 2 years, 2 months ago

D is the correct answer
upvoted 1 times

 **AD_0525** 2 years, 3 months ago

D is correct
upvoted 2 times

 **norrec9** 2 years, 5 months ago

D is the answer
upvoted 6 times

You will have several applications running on different Compute Engine instances in the same project. You want to specify at a more granular level the service account each instance uses when calling Google Cloud APIs. What should you do?

- A. When creating the instances, specify a Service Account for each instance.
- B. When creating the instances, assign the name of each Service Account as instance metadata.
- C. After starting the instances, use gcloud compute instances update to specify a Service Account for each instance.
- D. After starting the instances, use gcloud compute instances update to assign the name of the relevant Service Account as instance metadata.

Correct Answer: C*Community vote distribution*

A (100%)

✉  **GoCloud**  2 years, 5 months ago

A.

upvoted 22 times

✉  **JieHeng**  2 years, 3 months ago

A, when you create an instance using the gcloud command-line tool or the Google Cloud Console, you can specify which service account the instance uses when calling Google Cloud APIs - https://cloud.google.com/compute/docs/access/service-accounts#associating_a_service_account_to_an_instance

upvoted 16 times

✉  **Captain1212**  1 month ago

Selected Answer: A

Option A is correct , when you create the instance , that time itself you can specify the service account of each instance

upvoted 1 times

✉  **Bobbybash** 7 months, 3 weeks ago

Selected Answer: A

A. When creating the instances, specify a Service Account for each instance.

To specify a more granular level of service account for each Compute Engine instance, you should specify a Service Account for each instance when you create it. This can be done through the Compute Engine API or the Cloud Console. By doing so, the specified Service Account will be used when calling Google Cloud APIs from that instance.

Option B, assigning the name of each Service Account as instance metadata, is not the best solution as metadata can be accessed by anyone with access to the instance, which could potentially lead to security issues.

Options C and D, using gcloud compute instances update to specify a Service Account or assign the name of a Service Account as instance metadata after starting the instances, can also be done, but it is a less efficient approach as it requires additional steps and can lead to human error if not properly documented.

upvoted 3 times

✉  **VarunGo** 7 months ago

used chatgpt

upvoted 3 times

✉  **ryumada** 1 year, 2 months ago

Selected Answer: A

Vote for A, because there is no instance running yet. "You will have several applications running..."

upvoted 2 times

✉  **Roro_Brother** 1 year, 3 months ago

Selected Answer: A

A, there is no instance running yet

upvoted 1 times

✉  **AzureDP900** 1 year, 3 months ago

A is good option for given scenario.

upvoted 1 times

✉  **somenick** 1 year, 6 months ago

Selected Answer: A

You can set/update the service account only when the instance is not running

upvoted 3 times

✉ **MeikJ02** 1 year, 7 months ago

Selected Answer: A

A - the instances are not running yet
upvoted 2 times

 **RazOr** 1 year, 8 months ago

A: you can define which GCP service account is associated with a Compute Engine instance when creating one. It is still possible to change the service account later.
Link to the GCP docs: <https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#using>
upvoted 1 times

 **liyux21** 1 year, 11 months ago

A is correct. You can change the assigned service account, use gcloud compute instances set-service-account, not the update
upvoted 3 times

 **techabhi2_0** 2 years ago

Answer is C. The instances are already running. So you need to change the Service account
upvoted 2 times

 **techabhi2_0** 2 years ago

My bad - "you will have" . Correct answer - A
upvoted 8 times

 **kaes** 3 months, 3 weeks ago

Even if they would be running, I don't think it's possible to change the service account with the "update" command. You need to use "set-service-account" appropriately: <https://cloud.google.com/sdk/gcloud/reference/compute/instances/set-service-account>
upvoted 1 times

 **AD_0525** 2 years, 3 months ago

A should be correct
upvoted 4 times

 **norrec9** 2 years, 5 months ago

A is the answer
upvoted 4 times

 **Biju1** 2 years, 5 months ago

It should be A
upvoted 3 times

You are creating an application that will run on Google Kubernetes Engine. You have identified MongoDB as the most suitable database system for your application and want to deploy a managed MongoDB environment that provides a support SLA. What should you do?

- A. Create a Cloud Bigtable cluster, and use the HBase API.
- B. Deploy MongoDB Atlas from the Google Cloud Marketplace.
- C. Download a MongoDB installation package, and run it on Compute Engine instances.
- D. Download a MongoDB installation package, and run it on a Managed Instance Group.

Correct Answer: C*Community vote distribution*

B (89%)

11%

✉  **arsh1916**  2 years, 4 months ago

Simple it's B
upvoted 14 times

✉  **Ixgywil**  2 years, 5 months ago

MongoDB Atlas is actually managed and supported by third-party service providers.

<https://console.cloud.google.com/marketplace/details/gc-launcher-for-mongodb-atlas/mongodb-atlas>
upvoted 11 times

✉  **Ixgywil** 2 years, 5 months ago

I think that's it. The answer is B
upvoted 6 times

✉  **rahulrauki**  1 week, 6 days ago

Selected Answer: B

The keyword is managed MongoDB environment, both C and D are managed by users, A is irrelevant, So B
upvoted 1 times

✉  **Captain1212** 1 month ago

Selected Answer: B
B seems more correct , just deploy it from the Market place.
upvoted 1 times

✉  **WendyLC** 3 months, 3 weeks ago

Selected Answer: B
Answer is B
upvoted 1 times

✉  **PKookNN** 11 months, 2 weeks ago

Selected Answer: B
the best answer is B
upvoted 1 times

✉  **nonsense** 11 months, 2 weeks ago

Selected Answer: B
b. fast and simple
upvoted 1 times

✉  **11kc03** 11 months, 3 weeks ago

Selected Answer: C
Answer is B
upvoted 1 times

✉  **learn_GCP** 1 year ago

Selected Answer: B
B. is the answer
upvoted 1 times

✉  **snkhatri** 1 year, 1 month ago

B, keyword " support SLA"
upvoted 1 times

✉  **AzureDP900** 1 year, 3 months ago

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is right

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

B is correct

upvoted 1 times

 **rsuresh27** 1 year, 5 months ago

Anytime the question mentions a third party software, always use the cloud marketplace. Answer is B.

upvoted 6 times

 **dishum** 1 year, 8 months ago

the question says "want to deploy a managed MongoDB environment" which means it should be managed by something, i.e compute engine or MIG.

Near Ans is C or D

I choose C - becoz no need of Mongo db running on MIG, GKE can easily handle mongoDb on compute engine.

upvoted 1 times

 **fazalmf** 1 year, 9 months ago

Deploy MongoDB Atlas its free Tier - does free tier provides a Support SLA?

upvoted 1 times

 **jaffarali** 1 year, 9 months ago

Selected Answer: B

B is the right option to use managed services.

upvoted 2 times

 **alaahakim** 1 year, 10 months ago

I Pick B

upvoted 1 times

You are managing a project for the Business Intelligence (BI) department in your company. A data pipeline ingests data into BigQuery via streaming. You want the users in the BI department to be able to run the custom SQL queries against the latest data in BigQuery. What should you do?

- A. Create a Data Studio dashboard that uses the related BigQuery tables as a source and give the BI team view access to the Data Studio dashboard.
- B. Create a Service Account for the BI team and distribute a new private key to each member of the BI team.
- C. Use Cloud Scheduler to schedule a batch Dataflow job to copy the data from BigQuery to the BI team's internal data warehouse.
- D. Assign the IAM role of BigQuery User to a Google Group that contains the members of the BI team.

Correct Answer: C*Community vote distribution*

D (100%)

ApaMokus Highly Voted 2 years, 5 months ago

D is correct

roles/bigquery.user

When applied to a dataset, this role provides the ability to read the dataset's metadata and list tables in the dataset.

When applied to a project, this role also provides the ability to run jobs, including queries, within the project. A member with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project. Additionally, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets.

upvoted 26 times

blan_ak Highly Voted 2 years, 1 month ago

Why on the earth would the answer be C? It has no relevance to the question. The answer is D, hands down

upvoted 9 times

Captain1212 Most Recent 1 month ago**Selected Answer: D**

D is the right answer, just assign them the role by IAM and they will be able to use BQ

upvoted 1 times

ankyt9 10 months, 1 week ago**Selected Answer: D**

D is correct

upvoted 1 times

anolive 11 months, 1 week ago**Selected Answer: D**

makes mor sense

upvoted 1 times

sylva91 1 year ago**Selected Answer: D**

D is correct because google recommendations are always to privilege groups to individual accounts and this is what can make the users query the database unlike the Data Studio

upvoted 1 times

snkhatri 1 year, 1 month ago**Selected Answer: D**

D is right

upvoted 1 times

snkhatri 1 year, 1 month ago

D is correct

upvoted 1 times

patashish 1 year, 2 months ago

D is the answer

Hint - to **run the custom SQL queries*** against the latest data in BigQuery

upvoted 2 times

TaniaMalfoy 1 year, 3 months ago

Is correct, data pipeline is the key:

To create this sample batch data pipeline, you must have access to the following resources in your project:

A Cloud Storage bucket to store input and output files

A BigQuery dataset where you will create a table.

upvoted 1 times

✉️ **S00999** 1 year, 3 months ago

Selected Answer: D

Answer: D

The simplest.

It is not requested to automate the query. The BI team may also need to modify their query or have several different ones to meet the needs.

upvoted 1 times

✉️ **AzureDP900** 1 year, 3 months ago

D is correct

<https://cloud.google.com/bigquery/docs/access-control>

upvoted 1 times

✉️ **RazOr** 1 year, 8 months ago

Selected Answer: D

D sounds perfect with minimal steps.

Quote from the GCP docs: "BigQuery User
(roles/bigquery.user)

When applied to a dataset, this role provides the ability to read the dataset's metadata and list tables in the dataset.

When applied to a project, this role also provides the ability to run jobs, including queries, within the project. A principal with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project. Additionally, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets."

upvoted 1 times

✉️ **dishum** 1 year, 8 months ago

Answer is C

Option C says, there is a probability of an internal BI datawarehouse. Before providing the iam permissions, it is better to copy data to internal BI.

My view

upvoted 1 times

✉️ **arvsrv** 1 year, 8 months ago

Selected Answer: D

agree With D

upvoted 2 times

✉️ **gioresin1** 1 year, 9 months ago

maybe D is not sufficient because, as per documentation:

"Note: For a user to be able to query the tables in a dataset, it is not sufficient for the user to have access to the dataset. A user must also have permission to run a query job in a project. If you want to give a user permission to run a query from your project, give the user the bigquery.jobs.create permission for the project. You can do this by assigning the user the roles/bigquery.jobUser role for your project. For more information, see Access control examples".

upvoted 1 times

✉️ **alaahakim** 1 year, 10 months ago

I agree With D

upvoted 2 times

Your company is moving its entire workload to Compute Engine. Some servers should be accessible through the Internet, and other servers should only be accessible over the internal network. All servers need to be able to talk to each other over specific ports and protocols. The current on-premises network relies on a demilitarized zone (DMZ) for the public servers and a Local Area Network (LAN) for the private servers. You need to design the networking infrastructure on Google Cloud to match these requirements. What should you do?

- A. 1. Create a single VPC with a subnet for the DMZ and a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for the DMZ.
- B. 1. Create a single VPC with a subnet for the DMZ and a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public egress traffic for the DMZ.
- C. 1. Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for the DMZ.
- D. 1. Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public egress traffic for the DMZ.

Correct Answer: A

Community vote distribution

A (100%)

 **perdigiorno**  2 years, 3 months ago

Passed the test today. About 80% of the questions are here.
upvoted 25 times

 **associatecloudexamuser** 2 years, 2 months ago

Congratulations!
upvoted 3 times

 **sumantrao1** 1 year, 12 months ago

you got same questions from this examtopics
upvoted 2 times

 **yvinisiupacuando**  2 years, 5 months ago

A is the Right answer. You can discard B and C because they lack the need of creating Network Peering to communicate the DMZ VPC with the LAN VPC (LAN VPC is not exposed to public so they need to communicate via private addresses which cannot be achieved with 2 VPCs without Network Peering). Plus, you can discard B, as you don't need to enable the egress traffic, you always need to enable the ingress traffic as this is never enabled by default.

upvoted 23 times

 **Alela** 2 years, 4 months ago

A is wrong. You don't need to set up firewall rules between subnets of the same VPC. C is the answer
upvoted 12 times

 **Ashii** 2 years, 4 months ago

C is Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for the DMZ. Without peering 2 VPC's how this this be done ?

upvoted 5 times

 **BenKenGo6** 1 year, 1 month ago

and where do you have the VPC peering to communicate both VPCs?
upvoted 1 times

 **gcpengineer** 2 years, 2 months ago

You need fw rules
upvoted 1 times

 **demnok_lannik** 1 year, 8 months ago

of course you do
upvoted 2 times

 **Captain1212**  1 month ago

Selected Answer: A

A is the correct answer, as it meet the question requirement
upvoted 1 times

 **diasporabro** 11 months, 3 weeks ago

Selected Answer: A

A is the right choice
upvoted 1 times

 **snkhatri** 1 year, 1 month ago

Selected Answer: A

A seems right
upvoted 1 times

 **an0nym0us1** 1 year, 1 month ago

hi All what is the ans
upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

1 VPC enough for LAN and DMZ , Need to open appropriate firewall rules. A is right.
upvoted 1 times

 **S00999** 1 year, 3 months ago

Selected Answer: A

Vote for A

By default traffic between subnets on a VPC network is not allowed (except on the "default" network).
(This blocks traffic between all instances, not just traffic between subnets => FW rules must be defined to allow communications between all instances, regardless the subnets)

2 VPC will not work without peering.

upvoted 2 times

 **somenick** 1 year, 6 months ago

Selected Answer: A

You can't explicitly create a FW rule for the subnet, but connections are allowed or denied on a per-instance basis. You can think of the VPC firewall rules as existing not only between your instances and other networks, but also between individual instances within the same network.

C will not work without peering...

upvoted 2 times

 **akshaym87** 1 year, 8 months ago

Guys i cleared my exam last week. This question bank is must. 80% questions were from here.
upvoted 6 times

 **Boumer** 1 year, 2 months ago

congratulation, i need to pass exam in end month .Can you give me some advise please?
upvoted 1 times

 **HansKloss611** 1 year, 8 months ago

Selected Answer: A

A - my vote. Two different vpc need vpc peering.
upvoted 2 times

 **gioresin1** 1 year, 9 months ago

I don't understand why you say that the answer is A. If you have 2 subnets in the same network you won't have firewall between the 2 subnets. So you can't have a DMZ that can communicate with a private network. So the answer should be C.
upvoted 1 times

 **[Removed]** 2 years ago

A is correct
upvoted 1 times

 **maan2935** 2 years ago

Hi, I have an exam today. Are the questions still to some extent valid?
upvoted 3 times

 **jackwillis** 1 year, 12 months ago

Yes it is. I passed my exam on 3rd Oct received certificate on 7th Oct. Exam topics and a study course in Udemy which I bought during their discount sale, helped me.
upvoted 1 times

 **gerhardbl** 2 years, 1 month ago

Textbook example of DMZ and private subnet topology, hence answer A. Anyone who thinks C or multiple VPCs or whatever I strongly suggest you do CCNA before coming here.
upvoted 4 times

 **associatecloudexamuser** 2 years, 2 months ago

Yes. Correct answer is A. No need to complicate the setup by creating two different VPC networks.
upvoted 2 times

 **shayanahmed** 2 years, 3 months ago

All questions are still valid. I cleared my paper yesterday (shayan18@live.com)
upvoted 5 times

 **associatecloudexamuser** 2 years, 2 months ago

Congratulations!

upvoted 1 times

 **JieHeng** 2 years, 3 months ago

Should be A.

Not B, allow public egress traffic for DMZ won't help anything (also by default there is already this allow public egress traffic rule)

Not C & D, Network Peering is needed to allow internal IP address connectivity across two Virtual Private Cloud (VPC)

<https://cloud.google.com/vpc/docs/vpc-peering>

upvoted 5 times

You have just created a new project which will be used to deploy a globally distributed application. You will use Cloud Spanner for data storage. You want to create a Cloud Spanner instance. You want to perform the first step in preparation of creating the instance. What should you do?

- A. Enable the Cloud Spanner API.
- B. Configure your Cloud Spanner instance to be multi-regional.
- C. Create a new VPC network with subnetworks in all desired regions.
- D. Grant yourself the IAM role of Cloud Spanner Admin.

Correct Answer: B

Reference:

<https://cloud.google.com/spanner/docs/instances>

Community vote distribution

A (76%) 14% 10%

 **AzureDP900**  1 year, 3 months ago

A is right

<https://cloud.google.com/spanner/docs/getting-started/set-up>

upvoted 8 times

 **pfabio**  1 year, 4 months ago

Selected Answer: A

If you click on Create instance, the message is show in bottom: Cloud Spanner API for your project has been enabled.

upvoted 5 times

 **joao_01**  2 weeks, 4 days ago

Selected Answer: A

Its definitely A. Link: <https://cloud.google.com/spanner/docs/quickstart-console>

upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: A

Answer A is correct before you do anything first you needto enable the API of that particlur service

upvoted 1 times

 **Capability** 8 months, 4 weeks ago

Selected Answer: A

https://cloud.google.com/spanner/docs/quickstart-console?_ga=2.68426577.-1975890344.1661276010&_gac=1.161955406.1673625078.Cj0KCQiAn4SeBhCwARIsANEf9DjxfolckwcRZqaOS7Rem2pzXW_GmlBaLlxK4hHSe3YZ4DtE5oHHKVMQaArPUEALw_wcB#:~:text=see%20Pricing.-,Before%20you%20begin,Enable%20the%20Cloud%20Spanner%20API,-Create%20an%20instance

upvoted 1 times

 **Capability** 8 months, 4 weeks ago

A is right

https://cloud.google.com/spanner/docs/quickstart-console?_ga=2.68426577.-1975890344.1661276010&_gac=1.161955406.1673625078.Cj0KCQiAn4SeBhCwARIsANEf9DjxfolckwcRZqaOS7Rem2pzXW_GmlBaLlxK4hHSe3YZ4DtE5oHHKVMQaArPUEALw_wcB#:~:text=see%20Pricing.-,Before%20you%20begin,Enable%20the%20Cloud%20Spanner%20API,-Create%20an%20instance

upvoted 1 times

 **raaad** 9 months, 4 weeks ago

Selected Answer: A

Try the scenario yourself. Its A

upvoted 1 times

 **snkhatri** 1 year, 1 month ago

Selected Answer: A

A seems right

upvoted 1 times

 **bobthebuilder55110** 1 year, 2 months ago

Selected Answer: B

Answer must be B , here is why?

I was confused between A and B but I tested this by creating a new project, when you go to spanner and click on create a spanner instance it automatically enables the API for you and you can all see this activity on the notification panel on the top right along with the visibility of this message clearly on that instance page as well, First it will auto-enable the API and then It will give you an option to select

multi-region, Now the questions say your first step that has to be multi-region since enabling the API was done by google automatically and none other options makes sense here.

upvoted 3 times

 **bobthebuilder55110** 1 year, 2 months ago

Correction : It's A

Since it does not specify if we are using command line tool or UI, if you are using command line tool then you will have to enable this.

upvoted 6 times

 **mani098** 1 year, 3 months ago

Selected Answer: C

because Api auto enabled when you click create new instance on cloud spanner UI

upvoted 2 times

 **mav3r1ck** 1 year, 2 months ago

This is TRUE.

Verified this:

- Go to your GCP project and verify that "Cloud Spanner API" is NOT enabled.
- Go to Cloud Spanner. Click "Create Instance".
- Check back again the "Cloud Spanner API" , you will see that status is "API Enabled".

upvoted 3 times

 **Rutu_98** 1 year, 4 months ago

Selected Answer: A

Answer : A --> Tested

upvoted 3 times

 **bobthebuilder55110** 1 year, 2 months ago

How ? I tested this as well and It auto enables the API. Unless this is a new feature by google and the question is old then I am not sure.

upvoted 1 times

 **adarsh4503** 1 year ago

If you used CLI then the API would have to be enabled manually.

upvoted 1 times

 **Random_Mane** 1 year, 5 months ago

Selected Answer: A

Enabling API is the first step

upvoted 2 times

 **jblima** 1 year, 5 months ago

Selected Answer: A

A - I tested...

upvoted 1 times

 **parag09** 1 year, 5 months ago

A:-tested it need to Enable Api first

upvoted 1 times

 **sivasan** 1 year, 5 months ago

shouldnt the ans be enable cloud spanner API.Option A

upvoted 1 times

 **PAUGURU** 1 year, 5 months ago

A before anything else

upvoted 2 times

You have created a new project in Google Cloud through the gcloud command line interface (CLI) and linked a billing account. You need to create a new Compute Engine instance using the CLI. You need to perform the prerequisite steps. What should you do?

- A. Create a Cloud Monitoring Workspace.
- B. Create a VPC network in the project.
- C. Enable the compute googleapis.com API.
- D. Grant yourself the IAM role of Computer Admin.

Correct Answer: D

Reference:

<https://isb-cancer-genomics-cloud.readthedocs.io/en/latest/sections/gcp-info/gcp-info2/LaunchVM.html>*Community vote distribution*

C (100%)

✉️ **michalmrozik** 8 months, 3 weeks ago

Why not B? Can you create Compute Engine instance without assigning it to VPC?

upvoted 1 times

✉️ **Mike_SG** 4 months, 4 weeks ago

When you create a new project on the GCP, a default VPC network is automatically created for you.

upvoted 1 times

✉️ **Kyle1776** 4 months ago

Yeah, but who uses the default VPC and CIDR ranges? Technically you could, but it's not best practice and RARELY would fit in with a company's existing infrastructure.

upvoted 1 times

✉️ **lummy** 6 months, 2 weeks ago

I believe you can make use of the default vpc

upvoted 2 times

✉️ **roaming_panda** 9 months, 3 weeks ago

api > iam role .i vote for C !!

upvoted 2 times

✉️ **sylva91** 1 year ago**Selected Answer: C**

nothing can be done before activating the API

upvoted 2 times

✉️ **zellick** 1 year ago**Selected Answer: C**

C is the obvious answer.

upvoted 1 times

✉️ **III_bbb** 1 year, 1 month ago**Selected Answer: C**

api first

upvoted 2 times

✉️ **snkhatri** 1 year, 1 month ago**Selected Answer: C**

C the compute googleapis.com API

upvoted 1 times

✉️ **Nishanth222** 1 year, 1 month ago

Must be C

upvoted 1 times

Your company has developed a new application that consists of multiple microservices. You want to deploy the application to Google Kubernetes Engine (GKE), and you want to ensure that the cluster can scale as more applications are deployed in the future. You want to avoid manual intervention when each new application is deployed. What should you do?

- A. Deploy the application on GKE, and add a HorizontalPodAutoscaler to the deployment.
- B. Deploy the application on GKE, and add a VerticalPodAutoscaler to the deployment.
- C. Create a GKE cluster with autoscaling enabled on the node pool. Set a minimum and maximum for the size of the node pool.
- D. Create a separate node pool for each application, and deploy each application to its dedicated node pool.

Correct Answer: C

Community vote distribution

C (56%)

A (44%)

✉ **WendyLC** 3 months, 2 weeks ago

Selected Answer: C

C is the right choice... See this for reference https://cloud.google.com/architecture/best-practices-for-running-cost-effective-kubernetes-applications-on-gke#fine-tune_gke_autoscaling

A- HorizontalPodAutoscaler - it is best suited for stateless workers that can spin up quickly to react to usage spikes, and shut down gracefully to avoid workload instability.

upvoted 1 times

✉ **efar_cloud** 4 months ago

Answer is C

The key point is "ensure that the CLUSTER can scale"

A- HorizontalPodAutoscaler - ensures to scale the number of pods

while

C- Create a GKE cluster with autoscaling enabled on the node pool. Set a minimum and maximum for the size of the node pool. ensures to scale the number of nodes in the cluster.

So the answer is C.

upvoted 2 times

✉ **MrJkr** 4 months, 1 week ago

Selected Answer: A

Its A,

When you first deploy your workload to a Kubernetes cluster, you may not be sure about its resource requirements and how those requirements might change depending on usage patterns, external dependencies, or other factors. Horizontal Pod autoscaling helps to ensure that your workload functions consistently in different situations, and allows you to control costs by only paying for extra capacity when you need it.

upvoted 2 times

✉ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: A

i think it is A "you want to ensure that the cluster can scale as more applications are deployed in the future."

upvoted 2 times

✉ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: C

option C

upvoted 2 times

✉ **sabrinakloud** 5 months, 3 weeks ago

option A*

upvoted 2 times

✉ **dobberzoon** 6 months, 1 week ago

Selected Answer: A

Not knowing how many pods, and like nooneknows said, chatgpt... A is correct.

upvoted 1 times

✉ **nooneknows** 6 months, 1 week ago

Chat GPT said A is the Answer!

upvoted 2 times

✉ **lummy** 6 months, 2 weeks ago

i believe A is the answer, you cant figure out how many nodes you will need in the future...how you gotta set a maximum

upvoted 2 times

 **abdelsha** 7 months, 3 weeks ago

How do you set the maximum number of nodes and you do not know how your app will scale in the future? I think A is more accurate here.
upvoted 2 times

 **xmh5025** 9 months, 2 weeks ago

Selected Answer: A

less manual intervention
upvoted 4 times

 **diasporabro** 11 months, 3 weeks ago

Selected Answer: C

C is the right choice... See this for reference <https://cloud.google.com/kubernetes-engine/docs/concepts/cluster-autoscaler>
upvoted 1 times

 **Charumathi** 1 year ago

Selected Answer: C

C is the correct answer, you can enable the cluster autoscaling in node pool by specifying min and max node size.

https://cloud.google.com/kubernetes-engine/docs/how-to/cluster-autoscaler#adding_a_node_pool_with_autoscaling
upvoted 1 times

 **ale_brd_** 1 year ago

Selected Answer: C

it's mentioning "the cluster can scale" the answer is C
upvoted 2 times

 **hanskristian** 1 year ago

I Think it should be C
upvoted 1 times

 **zellck** 1 year ago

Selected Answer: C

C is my answer. You need to scale the cluster, not the pod.
upvoted 4 times

 **manjtrade2** 1 year ago

Selected Answer: A

A
<https://cloud.google.com/kubernetes-engine/docs/concepts/horizontalpodautoscaler>
upvoted 2 times

 **manjtrade2** 1 year ago

ITts A
<https://cloud.google.com/kubernetes-engine/docs/concepts/horizontalpodautoscaler>
upvoted 2 times

You need to manage a third-party application that will run on a Compute Engine instance. Other Compute Engine instances are already running with default configuration. Application installation files are hosted on Cloud Storage. You need to access these files from the new instance without allowing other virtual machines (VMs) to access these files. What should you do?

- A. Create the instance with the default Compute Engine service account. Grant the service account permissions on Cloud Storage.
- B. Create the instance with the default Compute Engine service account. Add metadata to the objects on Cloud Storage that matches the metadata on the new instance.
- C. Create a new service account and assign this service account to the new instance. Grant the service account permissions on Cloud Storage.
- D. Create a new service account and assign this service account to the new instance. Add metadata to the objects on Cloud Storage that matches the metadata on the new instance.

Correct Answer: A

Reference:

<https://cloud.google.com/compute/docs/access/service-accounts>

Community vote distribution

C (100%)

 **scanner2** 1 month ago

Selected Answer: C

C is correct.

upvoted 1 times

 **gcpreviewer** 1 year ago

Selected Answer: C

C is the clear choice. Want to create a new service account instead of using the default and grant it permissions in cloud storage. Straightforward C.

upvoted 1 times

 **manjtrade2** 1 year ago

Selected Answer: C

C is right

upvoted 1 times

 **snkhatri** 1 year, 1 month ago

Selected Answer: C

C seems right to me

upvoted 1 times

 **VietmanOfficiel** 1 year, 1 month ago

Selected Answer: C

"without allowing other instances" , the other instances are created with default compute engine service account. So you must create a new independant service account

upvoted 3 times

 **AzureDP900** 1 year, 3 months ago

C

<https://cloud.google.com/iam/docs/best-practices-for-using-and-managing-service-accounts>

If an application uses third-party or custom identities and needs to access a resource, such as a BigQuery dataset or a Cloud Storage bucket, it must perform a transition between principals. Because Google Cloud APIs don't recognize third-party or custom identities, the application can't propagate the end-user's identity to BigQuery or Cloud Storage. Instead, the application has to perform the access by using a different Google identity.

upvoted 1 times

 **KRIV_1** 1 year, 4 months ago

Although C is the correct answer notice that, as Google recommend, you first need to grant the service account the required permission before attach it to a resource.

upvoted 1 times

 **JelloMan** 1 year, 5 months ago

Selected Answer: C

C all the way. Restricts access to other VMs since they won't have the new service account you have associated with your new VM

upvoted 4 times

 **amindbesideitself** 1 year, 5 months ago

Selected Answer: C

C, other VMs will run as default service account.

upvoted 2 times

 **Akash7** 1 year, 5 months ago

C is correct as the other vms have default service accounts.

upvoted 2 times

 **PAUGURU** 1 year, 5 months ago

Selected Answer: C

C, using Default account makes the storage visible to other machines

upvoted 2 times

You need to configure optimal data storage for files stored in Cloud Storage for minimal cost. The files are used in a mission-critical analytics pipeline that is used continually. The users are in Boston, MA (United States). What should you do?

- A. Configure regional storage for the region closest to the users. Configure a Nearline storage class.
- B. Configure regional storage for the region closest to the users. Configure a Standard storage class.
- C. Configure dual-regional storage for the dual region closest to the users. Configure a Nearline storage class.
- D. Configure dual-regional storage for the dual region closest to the users. Configure a Standard storage class.

Correct Answer: B

Community vote distribution

B (51%)

D (49%)

✉  **akshaychavan7**  1 year, 4 months ago

Selected Answer: D

Mission critical is the keyword here which specifies that we need to have a multi-regional backup of the data to survive any regional failures.

So option D is the correct choice here.

upvoted 22 times

✉  **mav3r1ck** 1 year, 2 months ago

Keywords: minimal cost and mission-critical

Looks like people are just looking to be on the cost side. You need to meet both.

In this case, it needs to be "dual-region". This is much cheaper than storage in "multi-region" which is obviously not in the choices.

upvoted 3 times

✉  **Aninina** 11 months, 2 weeks ago

Dual region is expensive than multi-region. (Also mentioned in the documentation:

<https://cloud.google.com/storage/docs/locations>)

When we set objects to be multi-regional, we get to decide/shuffle the data around at will to meet our storage needs. When you take that control away from us, it reduces the flexibility of our systems, making it more expensive to operate.

upvoted 3 times

✉  **ryumada** 1 year, 2 months ago

At the first point in this documentation says that dual-regional storage is used for business continuity and disaster recovery. Disaster can affect to a regional architecture. I think it's make sense to use dual-regional storage for this case. Also, dual-regional storage is cheaper than multi-regional.

<https://cloud.google.com/storage/docs/dual-regions#use-dual-region-storage>

upvoted 1 times

✉  **JelloMan**  1 year, 5 months ago

Selected Answer: B

Continuous access to data means Standard since all of the other options are for infrequently accessed storage (Nearline, Coldline, Archive). Since no other regions are mentioned, single region is best in this case

upvoted 18 times

✉  **KRIV_1** 1 year, 4 months ago

And because single region is "cost-effective".

upvoted 7 times

✉  **SinghAnc**  6 days, 17 hours ago

Selected Answer: B

Correct answer is B

Requirement is low cost and mission critical. It needs low latency which by using region storage will be achieved.

Dual Storage is more expensive than multi regional. There is no need for dual storage in this case. If there will be a zonal failure it can be handled by regional storage.

Using Dual storage here will increase cost.

upvoted 1 times

✉  **scanner2** 1 month ago

Selected Answer: B

B is correct answer. Cloud storage provides 11 9's durability which is safe for mission critical files. Plus the question provides the users location to be a single region, and minimal cost should be there.

upvoted 2 times

✉  **NoCrapEva** 1 month ago

Selected Answer: B

Because Google Cloud Storage is 11 Nines durable (99.99999999%) its pretty safe for mission critical data (ref: <https://cloud.google.com/blog/products/storage-data-transfer/understanding-cloud-storage-11-9s-durability-target>). BUT the AVAILABILITY SLA on single region is 99.9% and for dual region its 99.95% (ONLY 0.05% better)... For the additional cost - dual region offers MINIMAL additional availability (ref: <https://cloud.google.com/storage/docs/storage-classes#standard>)

upvoted 1 times

✉ **gpais** 1 month, 2 weeks ago

Selected Answer: B

https://cloud.google.com/storage/docs/locations#location_recommendations

upvoted 1 times

✉ **Mo73w** 4 months, 3 weeks ago

Selected Answer: B

C is not correct because it would be more expensive than B. Dual-regional storage is designed for data that needs to be available in multiple regions. It is more expensive than regional storage, but it offers higher availability.

In this case, the files are used in a mission-critical analytics pipeline that is used continually. This means that the files need to be available with low latency. However, the users are only in Boston, MA (United States). Therefore, there is no need for the files to be available in multiple regions.

Therefore, the best option is to configure regional storage for the region closest to the users (us-central1) and configure a Standard storage class. This will be less expensive than dual-regional storage and will still provide the required low latency.

upvoted 1 times

✉ **Jelly_Wang** 5 months ago

Selected Answer: B

Answer is B. If you go with D, you must be a newbie trying to pass the exam and then hopefully get some practice with your next employer. Regional storage doesn't mean no automated failover and fallback. It does have automated failover and fallback on zonal failure, which is more than enough to cover a single city (Boston). Plus, I work for a provincial government who do business in a single province. We use regional storage. Man if you work for a company that run business in a single city and you go with multi regional storage, finger cross for yourself.

upvoted 3 times

✉ **vivekvj** 5 months, 2 weeks ago

Selected Answer: D

D. because of mission critical so high availability must be considered as well.

upvoted 1 times

✉ **Technobie** 5 months, 3 weeks ago

D :

Dual-region storage is a 'best of both worlds' solution, as it provides the ability to scale to TB per second (like regional), but also provides a second copy of data in a second region, protecting against regional outages. Similar to regional storage, dual-region storage provides customers with an environment to drive high-throughput analytical workloads by co-locating compute and storage in two regions of their choice.

upvoted 1 times

✉ **dobberzoon** 6 months, 1 week ago

Selected Answer: D

Mission critical made me opt D instead of B. You want availability.

upvoted 2 times

✉ **ashtonez** 7 months ago

Selected Answer: D

The answer is D, key here is we need standard , so B or D, but its a mission critical that must run continuously, so it should be always running, even if there is some region disaster, and the only option to avoid going down on one region disaster is D with dual region standard

upvoted 2 times

✉ **JC0926** 7 months, 3 weeks ago

Selected Answer: B

Boston, MA (United States).

upvoted 1 times

✉ **nigdyniezapomne** 8 months ago

Selected Answer: D

standard & dual-region

upvoted 1 times

✉ **JoniMONI** 8 months, 3 weeks ago

B. Configure regional storage for the region closest to the users. Configure a Standard storage class.

Regional storage stores data in a specific geographic region, which minimizes latency for users in that region. The Standard storage class is a good option for files that are used frequently, as it provides high availability and performance. Nearline storage is a lower-cost option that is best suited for data that is accessed less frequently. Configuring dual-regional storage and a Nearline storage class would not be appropriate for a mission-critical analytics pipeline that is used continually.

upvoted 2 times

✉ **robertozga** 8 months, 3 weeks ago

Answer is B:

Dual-region storage lets users specify two regions within the same continent. The users are here in one region. The mission-critical app refers to the autoscaling, which can be done in Cloud Spanner.

Please refer to: <https://cloud.google.com/storage/docs/dual-regions>

upvoted 2 times

 **Nazz1977** 8 months, 3 weeks ago

Selected Answer: B

I think it is B

upvoted 1 times

You are developing a new web application that will be deployed on Google Cloud Platform. As part of your release cycle, you want to test updates to your application on a small portion of real user traffic. The majority of the users should still be directed towards a stable version of your application. What should you do?

- A. Deploy the application on App Engine. For each update, create a new version of the same service. Configure traffic splitting to send a small percentage of traffic to the new version.
- B. Deploy the application on App Engine. For each update, create a new service. Configure traffic splitting to send a small percentage of traffic to the new service.
- C. Deploy the application on Kubernetes Engine. For a new release, update the deployment to use the new version.
- D. Deploy the application on Kubernetes Engine. For a new release, create a new deployment for the new version. Update the service to use the new deployment.

Correct Answer: B

Reference:

<https://cloud.google.com/appengine/docs/admin-api/migrating-splitting-traffic>*Community vote distribution* A (100%)

✉  **scanner2** 1 month ago

Selected Answer: A

Answer is A.

upvoted 1 times

✉  **gary_gary** 4 months, 3 weeks ago

Similar questions seem to appear multiple times.

upvoted 3 times

✉  **urcloudpartner** 9 months, 2 weeks ago

some of these questions, the default by examtopics is completely different why so, why cannot they fix it once a real answer is known.
upvoted 2 times

✉  **Charumathi** 1 year ago

Selected Answer: A

A is correct answer,

Keyword, Version, traffic splitting, App Engine supports traffic splitting for versions before releasing.

upvoted 3 times

✉  **snkhatri** 1 year, 1 month ago

Selected Answer: A

A obvious choice

upvoted 1 times

✉  **KapilDhamija** 1 year, 1 month ago

Selected Answer: A

Vote goes to A

upvoted 1 times

✉  **AzureDP900** 1 year, 3 months ago

It is no brainer questions, It is A.

upvoted 3 times

✉  **Tirthankar17** 1 year, 4 months ago

A obviously. No need to create a new service.

upvoted 1 times

✉  **pluiedust** 1 year, 4 months ago

Selected Answer: A

Obviously A

upvoted 1 times

✉  **JelloMan** 1 year, 5 months ago

Selected Answer: A

All the way. When deploying new versions you can split traffic for A/B testing to see how user base reacts to changes.

upvoted 4 times

 **amindbesideitself** 1 year, 5 months ago

Selected Answer: A

A is correct, there's no need to create a new service for each update.

upvoted 1 times

 **aabbcc321** 1 year, 5 months ago

ans is A

upvoted 1 times

 **sivasan** 1 year, 5 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **PAUGURU** 1 year, 5 months ago

Selected Answer: A

A is correct

upvoted 3 times

You need to add a group of new users to Cloud Identity. Some of the users already have existing Google accounts. You want to follow one of Google's recommended practices and avoid conflicting accounts. What should you do?

- A. Invite the user to transfer their existing account.
- B. Invite the user to use an email alias to resolve the conflict.
- C. Tell the user that they must delete their existing account.
- D. Tell the user to remove all personal email from the existing account.

Correct Answer: A

Reference:

<https://support.google.com/cloudidentity/answer/7332836?hl=en>

Community vote distribution

A (94%) 6%

✉ **ggupton1** Highly Voted 1 year, 5 months ago

Selected Answer: A

<https://cloud.google.com/architecture/identity/assessing-existing-user-accounts>

If you want to maintain the access rights and some of the data associated with the Gmail account, you can ask the owner to remove Gmail from the user account so that you can then migrate them to Cloud Identity or Google Workspace.

upvoted 6 times

✉ **scanner2** Most Recent 1 month ago

Selected Answer: A

Answer is A.

upvoted 1 times

✉ **snkhatri** 1 year, 1 month ago

Selected Answer: A

A obvious choice

upvoted 1 times

✉ **bobthebuilder55110** 1 year, 2 months ago

Selected Answer: A

Correct Answer: A

upvoted 1 times

✉ **bobthebuilder55110** 1 year, 2 months ago

Here is why ?

Question states "Some of the users already have existing Google accounts." Meaning they have personal account or any google account and what Option B is saying is to use aliases, as per google documentation this is only helpful when we want someone to receive emails in one inbox with 2 email names, meaning x@google.com and y@google.com goes to the same inbox BUT what you can't do is to have personal@google.com and company@google.com since the company wouldn't add you to their domain as that is not google recommended practice.

https://support.google.com/a/answer/33327?hl=en#when_to_use

upvoted 2 times

✉ **ryumada** 1 year, 2 months ago

Selected Answer: A

Vote for A as the right answer. The docs in this link:

<https://cloud.google.com/architecture/identity/migrating-consumer-accounts>

as provided by PAUGURU in his comment explains clearly about resolving account conflict. In the doc says nothing about to change email alias to resolve the conflict. So, following the documentation in that link means you are following the Googles Recommended Practices.

upvoted 4 times

✉ **zolthar_z** 1 year, 2 months ago

Selected Answer: A

A is the answer, for security reasons google best practices recommend transfer the account

upvoted 1 times

✉ **sai_learner** 1 year, 2 months ago

Selected Answer: B

Answer is B

<https://support.google.com/cloudidentity/answer/7062710>

upvoted 1 times

 **ryumada** 1 year, 2 months ago

I am not sure if the link you provide explains the reason of the reason of your chosen answer. As in the documentation stated the email alias after what happens if you rename the email. Also, the documentation doesn't explain about account conflict.

Better with the docs link provided by PAUGURU: <https://cloud.google.com/architecture/identity/migrating-consumer-accounts>

It's explains clearly about the conflicting email and best practices too.

upvoted 1 times

 **ryumada** 1 year, 2 months ago

*explains the reason of your chosen answer.

sorry, messed up sentence

upvoted 1 times

 **bobthebuilder55110** 1 year, 2 months ago

It should be A, I was confused with this as well but B is not relevant in this use case. Look at my Above comments.

upvoted 1 times

 **ramss** 1 year, 3 months ago

As per my understanding, B is the correct answer.

upvoted 2 times

 **AzureDP900** 1 year, 3 months ago

A is right and followed GCP documentation to get more information.

upvoted 1 times

 **PAUGURU** 1 year, 5 months ago

Selected Answer: A

A looks right

<https://cloud.google.com/architecture/identity/migrating-consumer-accounts>

upvoted 3 times

You need to manage a Cloud Spanner instance for best query performance. Your instance in production runs in a single Google Cloud region. You need to improve performance in the shortest amount of time. You want to follow Google best practices for service configuration. What should you do?

- A. Create an alert in Cloud Monitoring to alert when the percentage of high priority CPU utilization reaches 45%. If you exceed this threshold, add nodes to your instance.
- B. Create an alert in Cloud Monitoring to alert when the percentage of high priority CPU utilization reaches 45%. Use database query statistics to identify queries that result in high CPU usage, and then rewrite those queries to optimize their resource usage.
- C. Create an alert in Cloud Monitoring to alert when the percentage of high priority CPU utilization reaches 65%. If you exceed this threshold, add nodes to your instance.
- D. Create an alert in Cloud Monitoring to alert when the percentage of high priority CPU utilization reaches 65%. Use database query statistics to identify queries that result in high CPU usage, and then rewrite those queries to optimize their resource usage.

Correct Answer: C

Reference:

<https://cloud.google.com/spanner/docs/instances>

Community vote distribution

C (100%)

✉ **PAUGURU** Highly Voted 1 year, 5 months ago

Selected Answer: C

C looks correct, increase instances on single region if CPU above 65%
<https://cloud.google.com/spanner/docs/cpu-utilization#recommended-max>
upvoted 9 times

✉ **Ahmed_Y** Most Recent 1 month, 2 weeks ago

Selected Answer: C

I was keep thinking of A until I get to the link that thanks for @rsuresh27 provided bellow. the 45% is for the multi region.
upvoted 2 times

✉ **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: C

Metric Maximum for single-region instances Maximum per region for multi-region instances
High priority total 65% 45%
24-hour smoothed aggregate 90% 90%
upvoted 3 times

✉ **dobberzoon** 6 months, 1 week ago

Selected Answer: C

C makes sense.
upvoted 1 times

✉ **Aninina** 11 months, 2 weeks ago

Selected Answer: C

<https://cloud.google.com/spanner/docs/cpu-utilization>
upvoted 1 times

✉ **snkhatri** 1 year, 1 month ago

Selected Answer: C

C looks correct
upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

shortest timeframe is key here , I am going with C as my answer.
upvoted 1 times

✉ **Tirthankar17** 1 year, 4 months ago

Selected Answer: C

C is correct
upvoted 1 times

✉ **pfabio** 1 year, 4 months ago

Selected Answer: C

Metric Maximum for single-region instances Maximum per region for multi-region instances

High priority total> 65% 45%

24-hour smoothed aggregate> 90% 90%

upvoted 1 times

 **rsuresh27** 1 year, 5 months ago

Selected Answer: C

The correct answer is C

<https://cloud.google.com/spanner/docs/cpu-utilization#recommended-max>

upvoted 3 times

Your company has an internal application for managing transactional orders. The application is used exclusively by employees in a single physical location. The application requires strong consistency, fast queries, and ACID guarantees for multi-table transactional updates. The first version of the application is implemented in PostgreSQL, and you want to deploy it to the cloud with minimal code changes. Which database is most appropriate for this application?

- A. BigQuery
- B. Cloud SQL
- C. Cloud Spanner
- D. Cloud Datastore

Correct Answer: C

Reference:

<https://cloud.google.com/spanner/docs/transactions>*Community vote distribution*

B (58%)	C (42%)
---------	---------

✉  **peugeotduke**  1 year, 4 months ago

Read the question :

The application is used exclusively by employees in a single physical location.

upvoted 16 times

✉  **theBestStudent** 1 year ago

Correct. That is the key thing. I have no idea why some people ended up thinking cloud spanner is better. Definitely is alternative A.

upvoted 4 times

✉  **theBestStudent** 1 year ago

Sorry I meant B, I had a typo.

upvoted 4 times

✉  **PAUGURU**  1 year, 5 months ago

Selected Answer: B

B -> minimal code changes

upvoted 11 times

✉  **cheeseburger12388** 1 year, 4 months ago

Cloud SQL for PostgreSQL is a fully-managed database service that helps you set up, maintain, manage, and administer your PostgreSQL relational databases on Google Cloud Platform.

<https://cloud.google.com/sql/docs/postgres>

upvoted 4 times

✉  **Dino0411**  2 months ago

Selected Answer: C

C. Select Cloud Spanner.

Cloud Spanner offers strong consistency, fast queries, and importantly, ACID guarantees for updates in multi-table transactions. Cloud Spanner is well suited for large transactional databases that require horizontal scaling and offers relational database semantics. Even if the first version was PostgreSQL, Cloud Spanner is the best choice for this kind of application with strict requirements for ACID transactions.

Cloud SQL is also a relational database service, and while some database engines offer ACID transactions, it is not designed like Cloud Spanner for the strict requirements of multi-table transactional updates.

Reference link: Google Cloud - Cloud Spanner: <https://cloud.google.com/spanner>

Reference link: Google Cloud - ACID Transactions in Cloud Spanner: <https://cloud.google.com/spanner/docs/transactions>

upvoted 1 times

✉  **Mo73w** 4 months, 3 weeks ago

Selected Answer: B

the best choice for this application is Cloud SQL for PostgreSQL. It offers the required strong consistency, fast queries, and ACID guarantees for multi-table transactional updates. It is also a good choice for applications that are implemented in PostgreSQL and that you want to deploy to the cloud with minimal code changes.

upvoted 1 times

✉  **N_A** 5 months ago

Selected Answer: B

B. instead of C. because it's use in a single physical location and is compatible with PostgreSQL.

upvoted 1 times

 **dobberzoon** 6 months, 1 week ago

B makes sense.
upvoted 1 times

 **ashtonez** 7 months ago

Selected Answer: B
Its B because of single location, usually cloud sql is for non global or/and HA sql ddbb, spanner is for global sql ddbb
upvoted 1 times

 **JC0926** 7 months, 3 weeks ago

Selected Answer: B
key: single physical location
upvoted 1 times

 **Superflip314** 8 months ago

The answer is B.
One of the major advantages of Spanner is that the data is replicated over the Globe, since this data will only be accessed at one location Cloud SQL will be sufficient.
upvoted 1 times

 **Nazz1977** 8 months ago

Selected Answer: B
Not only Cloud Spanner has the ACID guarantees. SQL and others can be.....now I go to B (I have changed my vote)
upvoted 1 times

 **Di4sa** 8 months, 1 week ago

Selected Answer: C
C is the right answer
Only Cloud Spanner has the ACID guarantees and String consistency as mentioned in the question,
for those who are saying B because of minimal changes with PostgreSQL check all the feature of CloudSpanner with below link
<https://cloud.google.com/spanner#all-features>
<https://cloud.google.com/spanner/docs/postgresql-interface>
upvoted 2 times

 **Sammydp2020** 8 months, 1 week ago

Selected Answer: C
Reference: <https://cloud.google.com/blog/topics/developers-practitioners/your-google-cloud-database-options-explained>
upvoted 1 times

 **Nazz1977** 8 months, 3 weeks ago

Selected Answer: C
Only Cloud Spanner has the ACID guarantees.....I go to C
upvoted 1 times

 **FeaRoX** 8 months, 1 week ago

no, not only spanner has acid. ACID is characteristic of basically any RDBMS, including postgres or mysql.
upvoted 3 times

 **Nazz1977** 8 months ago

You are right my friend. I have changed my vote afte your answer. Thanks alot. Now it is B.
upvoted 1 times

 **alex000** 9 months ago

Selected Answer: B
Read the test:
- from a single phisical location
- with minimal changes: onpremise Postgres -> Cloud SQL (Postgres)
upvoted 2 times

 **davidsalomon** 10 months ago

Selected Answer: B
No need for global scaling
upvoted 1 times

 **ChristN** 10 months, 1 week ago

Selected Answer: C
C: Cloud spanner is the answer.
I was confused in the beginning because of this single location but after following this link (https://cloud.google.com/spanner/docs/instances?&_ga=2.57108579.-838482844.1646871720#regional_configurations) you will see that spanner can be either regional or multi-regional: " When you create an instance, you must configure it as either regional (that is, all the resources are contained within a single Google Cloud region) or multi-region (that is, the resources span more than one region). You make this choice by selecting an instance configuration, which determines where your data is stored for that instance."
and all the other conditions for a cloud spanner db are met ==> Transactional, strong consistency, fast queries, ACID, postgreSQL, minimal code change.
upvoted 2 times

 **penillo** 10 months, 1 week ago

Selected Answer: C

Voting for Spanner:

"It is a unique database that combines ACID transactions, SQL queries, and relational structure with the scalability that you typically associate with non-relational or NoSQL databases. As a result, Spanner is best used for applications such as gaming, payment solutions, global financial ledgers, retail banking and inventory management that require ability to scale limitlessly with strong-consistency and high-availability. "

<https://cloud.google.com/blog/topics/developers-practitioners/your-google-cloud-database-options-explained>

upvoted 2 times

You are assigned to maintain a Google Kubernetes Engine (GKE) cluster named 'dev' that was deployed on Google Cloud. You want to manage the GKE configuration using the command line interface (CLI). You have just downloaded and installed the Cloud SDK. You want to ensure that future CLI commands by default address this specific cluster. What should you do?

- A. Use the command gcloud config set container/cluster dev.
- B. Use the command gcloud container clusters update dev.
- C. Create a file called gke.default in the ~/.gcloud folder that contains the cluster name.
- D. Create a file called defaults.json in the ~/.gcloud folder that contains the cluster name.

Correct Answer: A

Reference:

<https://cloud.google.com/kubernetes-engine/docs/how-to/cluster-access-for-kubectl#gcloud-config>

Community vote distribution

A (95%) 5%

✉ **73173v2** Highly Voted 1 year, 1 month ago

Selected Answer: A

To set a default cluster for gcloud commands, run the following command:

<https://cloud.google.com/kubernetes-engine/docs/how-to/managing-clusters>

upvoted 7 times

✉ **scanner2** Most Recent 1 month ago

Selected Answer: A

Answer = A

upvoted 1 times

✉ **snkhatri** 1 year, 1 month ago

Selected Answer: A

A looks right to me

upvoted 1 times

✉ **AzureDP900** 1 year, 3 months ago

A is right

To set a default cluster for gcloud commands, run the following command:

gcloud config set container/cluster CLUSTER_NAME

upvoted 3 times

✉ **ggupton1** 1 year, 5 months ago

Selected Answer: A

Set a default cluster for gcloud

To set a default cluster for commands gcloud, run the following command:

Per <https://cloud.google.com/kubernetes-engine/docs/how-to/managing-clusters?hl=fr>
gcloud config set container/cluster CLUSTER_NAME

upvoted 3 times

✉ **Akash7** 1 year, 5 months ago

Answer is A,

To set a default cluster for gcloud commands, run the following command:

gcloud config set container/cluster CLUSTER_NAME

<https://cloud.google.com/kubernetes-engine/docs/how-to/managing-clusters?hl=en>

upvoted 1 times

✉ **MadMikedD** 1 year, 5 months ago

Selected Answer: A

To set a default cluster for gcloud commands, run the following command:

gcloud config set container/cluster CLUSTER_NAME

upvoted 3 times

✉ **cheeseburger12388** 1 year, 4 months ago

https://cloud.google.com/kubernetes-engine/docs/how-to/managing-clusters#default_cluster_kubectl

upvoted 1 times

✉ **aswinachu** 1 year, 5 months ago

Selected Answer: B

Correct Ans B

<https://cloud.google.com/sdk/gcloud/reference/container/clusters/update>

upvoted 1 times

 **VympelGRU** 1 year, 5 months ago

B sounds right.

upvoted 2 times

 **Malib** 1 year, 5 months ago

Selected Answer: A

<https://cloud.google.com/kubernetes-engine/docs/how-to/managing-clusters?hl=fr>

upvoted 3 times

The sales team has a project named Sales Data Digest that has the ID acme-data-digest. You need to set up similar Google Cloud resources for the marketing team but their resources must be organized independently of the sales team. What should you do?

- A. Grant the Project Editor role to the Marketing team for acme-data-digest.
- B. Create a Project Lien on acme-data-digest and then grant the Project Editor role to the Marketing team.
- C. Create another project with the ID acme-marketing-data-digest for the Marketing team and deploy the resources there.
- D. Create a new project named Marketing Data Digest and use the ID acme-data-digest. Grant the Project Editor role to the Marketing team.

Correct Answer: A*Community vote distribution*

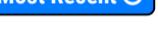
C (100%)

✉  **gcpj**  1 year, 3 months ago

Selected Answer: C

Answer should be C because the resources for the marketing team should be independent from the Sales team. Resources are tied and separated by projects.

upvoted 8 times

✉  **fdelacortina**  20 hours, 7 minutes ago

Why not D?

upvoted 1 times

✉  **scanner2** 1 month ago

Selected Answer: C

Answer = C

upvoted 1 times

✉  **Dmosh** 3 months, 2 weeks ago

Selected Answer: C

No more a technical exam ;(

upvoted 1 times

✉  **Priyanka109** 1 year ago

Sales and Marketing resource has to be in separate project so C.

upvoted 1 times

✉  **snkhatri** 1 year, 1 month ago

Selected Answer: C

C looks right

upvoted 1 times

✉  **[Removed]** 1 year, 2 months ago

WhatsApp : +1(956)-520-4006 to obtain PMP, CISM, CCNA, CEH, PRINCE2, CISCO, ISTQB, PRINCE2, AWS/Azure/Sale force/ITIL Foundation/EC- COUNCIL...

Get Certified with 100% pass guarantee. Pay after exam.

all CISCO, ISACA & EC- COUNCIL certifications

For the Below certificates we offer 100% pass guarantee:

1. AWS Certification
2. Sales force
3. Scrum Master
4. Oracle Certification: OCA, OCP
5. Cisco Certification: CCNA, CCNP, CCIE
6. ITIL Foundation & Intermediate
7. Prince 2 Foundation and Practitioner
8. VMWARE Certification
9. Check Point Certification
10. EC-COUNCIL Certification (CEH V-9, CCISO, CND)
11. Cloud Certification
12. IBM Certification
13. HP Certification
14. Citrix Certification
15. Juniper certification
16. Azure
17. Skype 70-333/34
18. PMI (PMP/CAPM/ACP/PBA)
19. ISTQB
20. SAP
21. ISACA (CISA, CISM, CRISC, CGEIT, COBIT)

 **Gautam_Thampy** 1 year ago

what the heck is this?... good job mods

upvoted 3 times

 **XDcloud** 1 year, 2 months ago

Selected Answer: C

C for sure

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

C is right answer based on given scenario..

upvoted 2 times

 **JelloMan** 1 year, 5 months ago

C is straight to the point and addresses exactly all of the concerns

upvoted 3 times

 **jblima** 1 year, 5 months ago

C is the better choice

upvoted 3 times

 **amindbesideitself** 1 year, 5 months ago

Selected Answer: C

C, other options seem obviously wrong

upvoted 2 times

 **Terzlightyear** 1 year, 5 months ago

Selected Answer: C

C seems the right answer

upvoted 3 times

 **PAUGURU** 1 year, 5 months ago

Selected Answer: C

Looks more like C

upvoted 2 times

You have deployed multiple Linux instances on Compute Engine. You plan on adding more instances in the coming weeks. You want to be able to access all of these instances through your SSH client over the internet without having to configure specific access on the existing and new instances. You do not want the Compute Engine instances to have a public IP. What should you do?

- A. Configure Cloud Identity-Aware Proxy for HTTPS resources.
- B. Configure Cloud Identity-Aware Proxy for SSH and TCP resources
- C. Create an SSH keypair and store the public key as a project-wide SSH Key.
- D. Create an SSH keypair and store the private key as a project-wide SSH Key.

Correct Answer: B*Community vote distribution*

B (75%) C (19%) 6%

✉  **Akash7**  1 year, 5 months ago

B is correct as question say no public IP on the instance.
upvoted 10 times

✉  **Akash7** 1 year, 5 months ago

Use IAP TCP to enable access to VM instances that do not have external IP addresses or do not permit direct access over the internet.
<https://cloud.google.com/iap/docs/using-tcp-forwarding>
upvoted 12 times

✉  **sabrinakloud**  5 months, 3 weeks ago

Selected Answer: B
B is correct
upvoted 1 times

✉  **Untamables** 11 months, 2 weeks ago

Selected Answer: B
Absolutely B
https://cloud.google.com/iap/docs/using-tcp-forwarding#tunneling_ssh_connections
upvoted 4 times

✉  **Gautam_Thampy** 1 year ago

Selected Answer: B
b is right
upvoted 1 times

✉  **snkhatri** 1 year, 1 month ago

Selected Answer: B
B looks right
upvoted 1 times

✉  **AzureDP900** 1 year, 3 months ago

B is correct, With TCP forwarding, IAP can protect SSH and RDP access to your VMs hosted on Google Cloud. Your VM instances don't even need public IP addresses.
upvoted 4 times

✉  **Rutu_98** 1 year, 4 months ago

Selected Answer: B
B is correct
upvoted 2 times

✉  **lixamec** 1 year, 4 months ago

Selected Answer: B
I think it is B
<https://medium.com/google-cloud/how-to-ssh-into-your-gce-machine-without-a-public-ip-4d78bd23309e>
upvoted 1 times

✉  **CloudAce7890** 1 year, 4 months ago

Selected Answer: B
B is correct as it uses IAP
upvoted 2 times

✉  **JelloMan** 1 year, 5 months ago

Selected Answer: C

IAP lets you establish a central authorization layer for applications accessed by HTTPS. This statement immediately eliminates A & B since you would need to publicly access the instance. C is also incorrect because it uses a private SSH key. Private keys are only meant for the user themselves as proof of their identity. Public SSH keys are meant to be used for access within applications, so that is the most applicable in this case. D, final answer.

upvoted 1 times

✉ **JelloMan** 1 year, 5 months ago

Selected Answer: D

IAP lets you establish a central authorization layer for applications accessed by HTTPS. This statement immediately eliminates A & B since you would need to publicly access the instance. C is also incorrect because it uses a private SSH key. Private keys are only meant for the user themselves as proof of their identity. Public SSH keys are meant to be used for access within applications, so that is the most applicable in this case. D, final answer.

upvoted 1 times

✉ **jeffangel28** 1 year, 2 months ago

<https://cloud.google.com/iap/docs/using-tcp-forwarding>
read pls!!

upvoted 1 times

✉ **JelloMan** 1 year, 5 months ago

Scratch this. Made a mistake. C is the correct answer (thought they were flipped)

upvoted 1 times

✉ **dark_3k03r** 1 year, 4 months ago

The question states "You do not want the Compute Engine instances to have a public IP", so that knocks out C and D as both options require public access.

IAP however supports port forwarding for your client so the instances are never exposed. That leaves A and B. With the question explicitly stating "your ssh client", that means you need to configure the Cloud Identity Aware proxy for ssh (i.e. port forwarding). Touching the proxies is not "configure specific access on the existing and new instances" as all this occurs on IAP and not the compute engine api.

upvoted 5 times

✉ **PAUGURU** 1 year, 5 months ago

Selected Answer: C

C looks better

upvoted 2 times

You have created an application that is packaged into a Docker image. You want to deploy the Docker image as a workload on Google Kubernetes Engine. What should you do?

- A. Upload the image to Cloud Storage and create a Kubernetes Service referencing the image.
- B. Upload the image to Cloud Storage and create a Kubernetes Deployment referencing the image.
- C. Upload the image to Container Registry and create a Kubernetes Service referencing the image.
- D. Upload the image to Container Registry and create a Kubernetes Deployment referencing the image.

Correct Answer: B

Reference:

<https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app>*Community vote distribution*

D (100%)

✉  **Aninina**  11 months, 2 weeks ago

Selected Answer: D

A deployment is responsible for keeping a set of pods running. A service is responsible for enabling network access to a set of pods.
upvoted 11 times

✉  **kaes**  3 months, 3 weeks ago

Selected Answer: D

Keep in mind, in the new exam, it would be rather "Artifact Registry"
upvoted 4 times

✉  **Priyanka109** 1 year ago

Upload your docker image on container registry then give a ref while creating deployment. So D!
upvoted 1 times

✉  **rimjhim09** 1 year ago

Selected Answer: D

I also vote for D. I passed my exam today and this question was there.
upvoted 4 times

✉  **nurai** 10 months ago

You said passed the exam, How do I know which answer to take? Voting and the actual answer never match. You passed the exam using this question? Should I use community answer or most voted? Pls, help. I have an exam in 3 days.Thanks
upvoted 1 times

✉  **nanakn** 9 months, 3 weeks ago

how was your exam? what's your answer to this question that you asked, now that you have given it.
upvoted 1 times

✉  **Gautam_Thampy** 1 year ago

its D , A and B are obviously incorrect
upvoted 2 times

✉  **pkmdb66** 1 year ago

Selected Answer: D

It's D
upvoted 1 times

✉  **snkhatri** 1 year, 1 month ago

Selected Answer: D

D is right
upvoted 1 times

✉  **Sakhi1234** 1 year, 1 month ago

Selected Answer: D

I have hands on expereince to answer this question.
upvoted 1 times

✉  **learn_GCP** 1 year, 1 month ago

Selected Answer: D

D is the Answer
upvoted 1 times

 **Bootshale** 1 year, 3 months ago

Selected Answer: D

D - not even a debate!

upvoted 1 times

 **AzureDP900** 1 year, 3 months ago

D perfect answer for given scenario.

upvoted 1 times

 **akshaychavan7** 1 year, 4 months ago

Selected Answer: D

D!!!!!!

upvoted 1 times

 **JelloMan** 1 year, 5 months ago

Selected Answer: D

You can only create GKE instances for docker containers through container registry. D all the way

upvoted 4 times

 **amindbesideitself** 1 year, 5 months ago

Selected Answer: D

D is best for storing and deploying containers to GKE

upvoted 2 times

 **Akash7** 1 year, 5 months ago

Correct answer is D.

upvoted 3 times

You are using Data Studio to visualize a table from your data warehouse that is built on top of BigQuery. Data is appended to the data warehouse during the day.

At night, the daily summary is recalculated by overwriting the table. You just noticed that the charts in Data Studio are broken, and you want to analyze the problem. What should you do?

- A. Review the Error Reporting page in the Cloud Console to find any errors.
- B. Use the BigQuery interface to review the nightly job and look for any errors.
- C. Use Cloud Debugger to find out why the data was not refreshed correctly.
- D. In Cloud Logging, create a filter for your Data Studio report.

Correct Answer: C

Community vote distribution

B (45%) D (32%) C (16%) 8%

 **JoniMONI**  8 months, 3 weeks ago

Selected Answer: B

B. Use the BigQuery interface to review the nightly job and look for any errors.

Since the problem is related to the data in the data warehouse, it would be useful to check the status of the nightly job that recalculates the data and overwrites the table. By reviewing the job in the BigQuery interface, you can see if it completed successfully and if there were any errors that may have caused the charts in Data Studio to break. Reviewing the Error Reporting page in the Cloud Console, using Cloud Debugger and creating a filter in Cloud Logging may not be directly related to the problem with the data.

upvoted 8 times

 **on2it**  3 months ago

Selected Answer: B

It needs to be B because it is the only way to proper investigate the issue.

upvoted 4 times

 **on2it** 3 months, 1 week ago

Selected Answer: B

The correct answer in this scenario would be option B: Use the BigQuery interface to review the nightly job and look for any errors.

When the charts in Data Studio are broken after the nightly recalculation of the table, it indicates that there might be an issue with the data refresh or update process in BigQuery. By reviewing the nightly job in the BigQuery interface, you can check for any errors or anomalies that occurred during the recalculation process.

upvoted 2 times

 **N_A** 5 months ago

Selected Answer: B

I am sorry to say that there is no filter for Data Studio in Cloud Logging (see: <https://cloud.google.com/logging/docs/api/v2/resource-list#resource-indexes>). So the correct answer is B.

upvoted 1 times

 **Jelly_Wang** 5 months ago

You are correct. However you can use Cloud Logging to troubleshoot issues with Data Studio by looking at the logs of the data sources that are being used by your report. I think this is what D mean to be.

upvoted 1 times

 **MahAli** 6 months ago

Selected Answer: D

I'll select D intuitively since I'm not using Data Studio and bigQuery that much

upvoted 1 times

 **dobberzoon** 6 months, 1 week ago

Selected Answer: D

D makes sense.

upvoted 1 times

 **cuong11111212** 6 months, 4 weeks ago

Selected Answer: D

i think the answer is D

upvoted 1 times

 **JC0926** 7 months, 3 weeks ago

Selected Answer: D

My ans: D

upvoted 1 times

✉ **Sabarno** 9 months ago

Selected Answer: D

There might be many issues in Google Data Studio, the fastest process might be to check Cloud Logging. Hence option D should be correct.

upvoted 1 times

✉ **Tanmay1178** 9 months, 2 weeks ago

Selected Answer: C

It is C

upvoted 1 times

✉ **Untamables** 11 months, 2 weeks ago

Selected Answer: A

I vote A.

<https://cloud.google.com/error-reporting/docs/check-app-crash>

upvoted 2 times

✉ **PKookNN** 11 months, 2 weeks ago

Selected Answer: D

not sure but I think D is the most appropriate

upvoted 1 times

✉ **btimmons** 11 months, 3 weeks ago

The answer is D

upvoted 1 times

✉ **Erie** 11 months, 3 weeks ago

Selected Answer: D

D because cloud logging is enabled by default. Cloud Debugger is something we need to configure manually in our code.

upvoted 2 times

✉ **diasporabro** 11 months, 3 weeks ago

Selected Answer: A

Hmmm.... Between A and D, I'll just select A and move on...

upvoted 1 times

✉ **diasporabro** 11 months, 3 weeks ago

I take that back, will select D... seems to be a more direct approach to analyze...

upvoted 1 times

✉ **AwesomeGCP** 12 months ago

Selected Answer: B

B. Use the BigQuery interface to review the nightly job and look for any errors.

upvoted 2 times

✉ **theBestStudent** 1 year ago

Selected Answer: D

On top of being almost deprecated, I'm not even sure if cloud debugger would be suitable for bigquery or cloud data studio stuff.

Error reporting I'm not even sure if it supports error reporting for data studio stuff.

For my sounds more simpler to just create a filter to see the logs in

"cloud logging" seems simpler, faster are more suitable solution. I go with D.

upvoted 4 times

You have been asked to set up the billing configuration for a new Google Cloud customer. Your customer wants to group resources that share common IAM policies. What should you do?

- A. Use labels to group resources that share common IAM policies.
- B. Use folders to group resources that share common IAM policies.
- C. Set up a proper billing account structure to group IAM policies.
- D. Set up a proper project naming structure to group IAM policies.

Correct Answer: B*Community vote distribution***B (100%)**

 **thimai**  1 year, 1 month ago

Selected Answer: B

B for me

"Folders are used to group resources that share common IAM policies"

<https://cloud.google.com/resource-manager/docs/creating-managing-folders>

upvoted 10 times

 **ExamsFR**  3 weeks, 6 days ago

Selected Answer: B

B is correct Answer,

upvoted 1 times

 **diasporabro** 11 months, 3 weeks ago

Selected Answer: B

"Folders are used to group resources that share common IAM policies."

upvoted 3 times

 **Charumathi** 1 year ago

Selected Answer: B

B is correct Answer,

Folders are nodes in the Cloud Platform Resource Hierarchy. A folder can contain projects, other folders, or a combination of both. Organizations can use folders to group projects under the organization node in a hierarchy. For example, your organization might contain multiple departments, each with its own set of Google Cloud resources. Folders allow you to group these resources on a per-department basis. Folders are used to group resources that share common IAM policies. While a folder can contain multiple folders or resources, a given folder or resource can have exactly one parent.

<https://cloud.google.com/resource-manager/docs/creating-managing-folders>

upvoted 4 times

 **zellick** 1 year ago

Selected Answer: B

B is the answer.

<https://cloud.google.com/resource-manager/docs/access-control-folders#best-practices-folders-iam>

upvoted 3 times

 **RockingRohit6** 1 year ago

Folders are used to group resources that share common IAM policies

upvoted 2 times

 **snkhatri** 1 year, 1 month ago

Selected Answer: B

B seems right to me

upvoted 1 times

You have been asked to create robust Virtual Private Network (VPN) connectivity between a new Virtual Private Cloud (VPC) and a remote site. Key requirements include dynamic routing, a shared address space of 10.19.0.1/22, and no overprovisioning of tunnels during a failover event. You want to follow Google-recommended practices to set up a high availability Cloud VPN. What should you do?

- A. Use a custom mode VPC network, configure static routes, and use active/passive routing.
- B. Use an automatic mode VPC network, configure static routes, and use active/active routing.
- C. Use a custom mode VPC network, use Cloud Router border gateway protocol (BGP) routes, and use active/passive routing.
- D. Use an automatic mode VPC network, use Cloud Router border gateway protocol (BGP) routes, and configure policy-based routing.

Correct Answer: D

Community vote distribution

C (100%)

✉ **theBestStudent** Highly Voted 1 year ago

Selected Answer: C

we need custom mode vpc so subnets are not created automatically (the ip range is mentioned in the question) also we will need active/passive HA VPN (as it is not mentioned we will have to use more than one HA VPN gateway).

Links : <https://cloud.google.com/network-connectivity/docs/vpn/concepts/best-practices>
<https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview#active>
<https://cloud.google.com/vpc/docs/vpc#subnet-ranges>

upvoted 7 times

✉ **theBestStudent** 1 year ago

Also for dynamic routing we need HA VPN

Link: <https://cloud.google.com/network-connectivity/docs/vpn/concepts/choosing-networks-routing#dynamic-routing>

upvoted 2 times

✉ **Captain1212** Most Recent 1 month ago

Selected Answer: C

C is the correct answer as we need to make sure that the subnets are not being created automatically

upvoted 1 times

✉ **Charumathi** 1 year ago

Selected Answer: C

C . Choose a Cloud VPN gateway that uses dynamic routing and the Border Gateway Protocol (BGP). Google recommends using HA VPN and deploying on-premises devices that support BGP.

Choose the appropriate tunnel configuration

Choose the appropriate tunnel configuration based on the number of HA VPN gateways:

If you have a single HA VPN gateway, use an active/passive tunnel configuration.

If you have more than one HA VPN gateway, use an active/active tunnel configuration.

<https://cloud.google.com/network-connectivity/docs/vpn/concepts/best-practices>

upvoted 4 times

✉ **ale_brd_** 1 year ago

Selected Answer: C

c is the correct one

upvoted 2 times

✉ **Arulkumar** 1 year ago

Selected Answer: C

Google Cloud Router

On Google Cloud, dynamic routing can be established using Cloud Router. It exchanges network topology information through Border Gateway Protocol (BGP). Cloud Router advertises subnets from its VPC network to another router or gateway via BGP. This is great for setting up VPN between the cloud and on-prem, as topology changes automatically propagate with no manual intervention and higher redundancy for your systems.

You now have:

Discovery of remote networks

Maintaining up-to-date routing information

Choosing the best path to destination networks

Ability to find a new best path if the current path is no longer available

And a great side effect can be lower latency because Cloud Router learns routes through BGP which allows for optimal data paths to reach

its destination, whether that be another network or a VPN gateway to on-premise. Cloud Router is also how Dedicated Interconnect can give you 10 gbp/s bandwidth between your cloud VPC and your peered on-premise data center.

upvoted 4 times

 **manjtrade2** 1 year ago

Selected Answer: C

C might be right

upvoted 1 times

 **snkhatri** 1 year, 1 month ago

Selected Answer: C

I think it should be C as there is too much customisation.

upvoted 1 times

You are running multiple microservices in a Kubernetes Engine cluster. One microservice is rendering images. The microservice responsible for the image rendering requires a large amount of CPU time compared to the memory it requires. The other microservices are workloads that are optimized for n1-standard machine types. You need to optimize your cluster so that all workloads are using resources as efficiently as possible. What should you do?

- A. Assign the pods of the image rendering microservice a higher pod priority than the other microservices.
- B. Create a node pool with compute-optimized machine type nodes for the image rendering microservice. Use the node pool with general-purpose machine type nodes for the other microservices.
- C. Use the node pool with general-purpose machine type nodes for the image rendering microservice. Create a node pool with compute-optimized machine type nodes for the other microservices.
- D. Configure the required amount of CPU and memory in the resource requests specification of the image rendering microservice deployment. Keep the resource requests for the other microservices at the default.

Correct Answer: B

Community vote distribution

B (100%)

✉  **calm_fox** 10 months, 2 weeks ago

Selected Answer: B

B is right

upvoted 1 times

✉  **diasporabro** 11 months, 3 weeks ago

Selected Answer: B

B looks like the right choice here

upvoted 1 times

✉  **learn_GCP** 1 year ago

Selected Answer: B

B. is the Answer

upvoted 1 times

✉  **adarsh4503** 1 year ago

I agree B is the answer.

upvoted 2 times

✉  **zellick** 1 year ago

Selected Answer: B

B is the answer.

upvoted 1 times

✉  **Gautam_Thampy** 1 year ago

Selected Answer: B

C is not correct coz general purpose machine types will not suffice for image rendering.

B is the most suitable answer.

upvoted 3 times

✉  **osanchez** 1 year ago

B is correct

upvoted 1 times

Your organization has three existing Google Cloud projects. You need to bill the Marketing department for only their Google Cloud services for a new initiative within their group. What should you do?

- A. 1. Verify that you are assigned the Billing Administrator IAM role for your organization's Google Cloud Project for the Marketing department.
2. Link the new project to a Marketing Billing Account.
- B. 1. Verify that you are assigned the Billing Administrator IAM role for your organization's Google Cloud account. 2. Create a new Google Cloud Project for the Marketing department. 3. Set the default key-value project labels to department:marketing for all services in this project.
- C. 1. Verify that you are assigned the Organization Administrator IAM role for your organization's Google Cloud account. 2. Create a new Google Cloud Project for the Marketing department. 3. Link the new project to a Marketing Billing Account.
- D. 1. Verify that you are assigned the Organization Administrator IAM role for your organization's Google Cloud account. 2. Create a new Google Cloud Project for the Marketing department. 3. Set the default key-value project labels to department:marketing for all services in this project.

Correct Answer: A

Community vote distribution

A (64%) B (20%) C (16%)

 **gcpreviewer**  1 year ago

Selected Answer: A

I understand that the question implies the creation of a new project, however neither of the roles listed have that functionality. If you chose B you are choosing an answer that has a direct contradiction because the Billing Account Admin does not have the permissions to create a new project. Thus, I think it is better to assume the new initiative/project is already created or being created by someone else and your job is simply to link the project to the account which you do have the appropriate permissions to perform.

A is my choice.

upvoted 14 times

 **moitsu**  1 year, 1 month ago

Selected Answer: B

Between A& B, Billing Administrator IAM role is either at the organisation level not project level. Hence A is out. C & D doesn't make sense.

upvoted 5 times

 **Gautam_Thampy** 1 year ago

The billing account administrator role can also be given at the project level. A is correct. Refer to this doc:
<https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 1 times

 **Gautam_Thampy** 1 year ago

Correction I meant the billing account admin role can be given at the organisation or the billing account level.

upvoted 1 times

 **pritampanda1988**  2 months, 1 week ago

Selected Answer: C

Option C: This option is accurate. As an Organization Administrator, you have the authority to create new projects within the organization's Google Cloud account. After creating the new Google Cloud Project for the Marketing department, you can link the project to a specific Marketing Billing Account to ensure that only their services are billed under their initiative.

upvoted 1 times

 **[Removed]** 3 months, 3 weeks ago

Selected Answer: A

The ans is A: Billing Admin cannot create a new project.

upvoted 1 times

 **N_A** 5 months ago

Selected Answer: A

C. doesn't provide the necessary permission to link a billing account to a project.

B. and D. don't provide the necessary permission to create a new project.

Furthermore, labels are used to filter resources in the billing breakdown. (See: <https://cloud.google.com/compute/docs/labeling-resources#what-are-labels>). Also, billing is set at a project level (See: <https://cloud.google.com/billing/docs/how-to/modify-project#overview>): "A Cloud Billing account is used to define who pays for a given set of resources, and it can be linked to one or more projects. Project usage is charged to the linked Cloud Billing account."

Answer A. is the only correct option.

upvoted 1 times

 **sabrinakloud** 5 months, 3 weeks ago

Selected Answer: A

Option A: no need to create a new project

upvoted 2 times

✉ **N1N0** 5 months, 4 weeks ago

C. 1. Verify that you are assigned the Organization Administrator IAM role for your organization's Google Cloud account. 2. Create a new Google Cloud Project for the Marketing department. 3. Link the new project to a Marketing Billing Account.

This approach ensures that you have the necessary permissions (Organization Administrator IAM role) to create a new project for the Marketing department. After creating the project, you'll link it to a Marketing Billing Account so that only their Google Cloud services are billed to their department. This way, you can separate costs for the Marketing department's new initiative from other projects in your organization.

upvoted 2 times

✉ **dobberzoon** 6 months, 1 week ago

Selected Answer: A

A makes sense.

upvoted 1 times

✉ **Nazz1977** 8 months, 2 weeks ago

Selected Answer: A

"Your organization has three existing Google Cloud projects".

Don't need to create anything....

I go to A

upvoted 3 times

✉ **researched_answer_boi** 8 months, 3 weeks ago

D is da way to go, I think.

An Organization Administrator can create projects. Owner of a new project is its creator. Owner can add labels to the project, and this is exactly the requirement here: selective billing (filtering) via labels.

upvoted 1 times

✉ **JoniMONI** 8 months, 3 weeks ago

Selected Answer: C

C. 1. Verify that you are assigned the Organization Administrator IAM role for your organization's Google Cloud account. 2. Create a new Google Cloud Project for the Marketing department. 3. Link the new project to a Marketing Billing Account.

You will need to be an Organization administrator to create new projects, and linking the new project to a Marketing Billing account will ensure that the Marketing department is only billed for their usage of Google Cloud services within that project.

Additionally, labels can be used to organize and categorize resources within a project, but they do not have any direct impact on billing. The Billing Administrator IAM role allows you to manage billing settings and view billing reports, but it does not allow you to create new projects or link them to specific billing accounts.

upvoted 4 times

✉ **GokulVelusaamy** 8 months, 3 weeks ago

Selected Answer: B

Billing Account Administrator is at Organization or billing account Level. Also, Labels are used to group the resources.

"A label is a key-value pair that helps you organize your Google Cloud instances. You can attach a label to each resource, then filter the resources based on their labels. Information about labels is forwarded to the billing system, so you can break down your billed charges by label."

upvoted 2 times

✉ **HiddenClouds** 9 months, 2 weeks ago

Selected Answer: A

The answer is A job is to link the project to the account with the appropriate permissions to do so.

upvoted 1 times

✉ **sureshsays** 10 months, 1 week ago

Selected Answer: C

with billing Admin role, you cannot create a Project. so, A & B cannot be true.

upvoted 2 times

✉ **Viggy1212** 10 months, 3 weeks ago

Answer: A

Option B, C and D talks about creating a new project. Both Billing Admin IAM and Organization Admin don't have access to create a new project.

upvoted 2 times

✉ **gneknurd** 11 months, 2 weeks ago

This question was on the test which I passed. I chose A

upvoted 4 times

✉ **PKookNN** 11 months, 2 weeks ago

Selected Answer: B

it was either A or B but since Billing Admin is only granted at org level or billing level so A is out. Correct A (C & D are not correct)

upvoted 1 times

You deployed an application on a managed instance group in Compute Engine. The application accepts Transmission Control Protocol (TCP) traffic on port 389 and requires you to preserve the IP address of the client who is making a request. You want to expose the application to the internet by using a load balancer. What should you do?

- A. Expose the application by using an external TCP Network Load Balancer.
- B. Expose the application by using a TCP Proxy Load Balancer.
- C. Expose the application by using an SSL Proxy Load Balancer.
- D. Expose the application by using an internal TCP Network Load Balancer.

Correct Answer: B*Community vote distribution*

A (75%)

B (25%)

joao_01 3 weeks, 5 days ago

Its A, for sure
upvoted 1 times

shreykul 2 months, 2 weeks ago

Selected Answer: A
<https://cloud.google.com/load-balancing/docs/choosing-load-balancer#:~:text=Proxy%20load%20balancers%20do%20not%20preserve%20client%20IP>
upvoted 2 times

geeroylenkins 2 months, 3 weeks ago

Selected Answer: A
I am going with A as the client IP needs to be preserved. Not sure with on2it votes once for A and once for B with the same comment especially because you need a *pass-through* load balancer to preserve the client IP as stated here: <https://cloud.google.com/load-balancing/docs/choosing-load-balancer#proxy-pass-through>

"You'd choose a passthrough Network Load Balancer to preserve client source IP addresses"
upvoted 1 times

on2it 3 months, 1 week ago

Selected Answer: B
The correct answer in this scenario would be option B: Expose the application by using a TCP Proxy Load Balancer.

A TCP Proxy Load Balancer is suitable for preserving the client's IP address when accepting TCP traffic on a specific port, such as port 389 in this case. When a client makes a request to the load balancer, the load balancer maintains the client's source IP address and forwards the traffic to the appropriate backend instances in the managed instance group. This allows the application to see the original client IP address and respond accordingly.

Option A, using an external TCP Network Load Balancer, does not preserve the client's IP address. The load balancer's IP address is seen as the source IP by the application, which may not meet the requirement.
upvoted 2 times

on2it 3 months, 1 week ago

Selected Answer: A
The correct answer in this scenario would be option B: Expose the application by using a TCP Proxy Load Balancer.

A TCP Proxy Load Balancer is suitable for preserving the client's IP address when accepting TCP traffic on a specific port, such as port 389 in this case. When a client makes a request to the load balancer, the load balancer maintains the client's source IP address and forwards the traffic to the appropriate backend instances in the managed instance group. This allows the application to see the original client IP address and respond accordingly.

Option A, using an external TCP Network Load Balancer, does not preserve the client's IP address. The load balancer's IP address is seen as the source IP by the application, which may not meet the requirement.
upvoted 1 times

TomFoot 3 months, 1 week ago

Selected Answer: A
Ansver is A!
If you are using a TCP/UDP network load balancer that preserves the client IP address (AWS Network Load Balancer, GCP External Network Load Balancer, Azure Load Balancer) or you are using Round-Robin DNS, then you can use the externalTrafficPolicy: Local setting to also preserve the client IP inside Kubernetes by bypassing kube-proxy and preventing it from sending traffic to other nodes.
upvoted 2 times

JaimieMS 3 months, 1 week ago

Selected Answer: A

The answer has to be A: external TCP Network Load Balancer.

From the Google doc "Choose a load balancer" (<https://cloud.google.com/load-balancing/docs/choosing-load-balancer>) it clearly states: "You'd choose a passthrough (passthrough = non-proxy) Network Load Balancer to preserve client source IP addresses (...)".

I also found a source that shows how to preserves a client's IP address in the TCP payload in a Proxy TCP load balancer (<https://medium.com/google-cloud/preserving-client-ips-through-google-clouds-global-tcp-and-ssl-proxy-load-balancers-3697d76feeb1>)... with A LOT of effort... but with such a clear statement from the google doc, I would stick to A.

By the way, I think that Google made a real mess with the load balancer types. This is the only thing more simple in AWS than in GCP.
upvoted 3 times

✉ **MrJkr** 3 months, 2 weeks ago

Selected Answer: B

In this scenario, using a TCP Proxy Load Balancer would be the most appropriate choice. A TCP Proxy Load Balancer operates at the transport layer (Layer 4) of the OSI model and can preserve the client's IP address while load balancing the TCP traffic to the backend instances.

External TCP Network Load Balancer (option A) is designed for network-level load balancing, but it does not have built-in support for preserving the client's IP address. Therefore, it may not be suitable for this specific requirement.

upvoted 1 times

✉ **gpa1s** 3 months, 1 week ago

As part of the requirements it requires to perserve the ip address which is incompatible with the TCP proxy. Leaving option A as the only viable option: <https://googlecloudarchitect.us/types-of-load-balancers-gcp/>

upvoted 1 times

✉ **gpa1s** 3 months, 2 weeks ago

A we need to perserve the client ip addresses: <https://cloud.in28minutes.com/gcp-certification-google-cloud-load-balancers>

upvoted 1 times

You are building a multi-player gaming application that will store game information in a database. As the popularity of the application increases, you are concerned about delivering consistent performance. You need to ensure an optimal gaming performance for global users, without increasing the management complexity. What should you do?

- A. Use Cloud SQL database with cross-region replication to store game statistics in the EU, US, and APAC regions.
- B. Use Cloud Spanner to store user data mapped to the game statistics.
- C. Use BigQuery to store game statistics with a Redis on Memorystore instance in the front to provide global consistency.
- D. Store game statistics in a Bigtable database partitioned by username.

Correct Answer: B

Community vote distribution

B (100%)

 **gpais** 1 month, 2 weeks ago

Selected Answer: B

<https://cloud.google.com/solutions/databases/games>
upvoted 1 times

 **tatyavinchu** 1 month, 3 weeks ago

global users = Cloud Spanner
Correct Answer is B
upvoted 2 times

 **3arle** 2 months ago

Selected Answer: B

Spanner should meet expectation
upvoted 1 times

 **MrJkr** 3 months, 2 weeks ago

Selected Answer: B

Among the options provided, the better answer for ensuring optimal gaming performance for global users without increasing management complexity would be option B

Cloud Spanner is a globally distributed, horizontally scalable database service provided by Google Cloud Platform. It offers strong consistency guarantees, high availability, and automatic scaling.

It offers the necessary features to ensure optimal gaming performance, global scalability, strong consistency, and automatic scaling, making it a suitable choice for storing user data mapped to game statistics.

upvoted 1 times

You are building an application that stores relational data from users. Users across the globe will use this application. Your CTO is concerned about the scaling requirements because the size of the user base is unknown. You need to implement a database solution that can scale with your user growth with minimum configuration changes. Which storage solution should you use?

- A. Cloud SQL
- B. Firestore
- C. Cloud Spanner
- D. Bigtable

Correct Answer: C

Community vote distribution

C (100%)

✉️  **3arle** 2 months ago

Selected Answer: C

and scales horizontally
upvoted 1 times

✉️  **Husni_adam** 2 months, 3 weeks ago

Selected Answer: C

Cloud Spanner because Relational database, scale across regions for workloads that have more stringent availability requirements,
Handles large amounts of data and for high transactional consistency

<https://cloud.google.com/blog/topics/developers-practitioners/databases-google-cloud-part-2-options-glance/>
upvoted 1 times

✉️  **mapcio123** 3 months, 1 week ago

Selected Answer: C

spanner- relational and global
upvoted 1 times

✉️  **_F4LLEN_** 3 months, 1 week ago

Spanner
upvoted 1 times

✉️  **gpais** 3 months, 2 weeks ago

I vote option C
upvoted 1 times

Your company has multiple projects linked to a single billing account in Google Cloud. You need to visualize the costs with specific metrics that should be dynamically calculated based on company-specific criteria. You want to automate the process. What should you do?

- A. In the Google Cloud console, visualize the costs related to the projects in the Reports section.
- B. In the Google Cloud console, visualize the costs related to the projects in the Cost breakdown section.
- C. In the Google Cloud console, use the export functionality of the Cost table. Create a Looker Studio dashboard on top of the CSV export.
- D. Configure Cloud Billing data export to BigQuery for the billing account. Create a Looker Studio dashboard on top of the BigQuery export.

Correct Answer: D

Community vote distribution

D (100%)

✉  **scanner2** 1 month ago

Selected Answer: D

Cloud Billing export to BigQuery enables you to export detailed Google Cloud billing data (such as usage, cost estimates, and pricing data) automatically throughout the day to a BigQuery dataset that you specify. Then you can access your Cloud Billing data from BigQuery for detailed analysis, or use a tool like Looker Studio to visualize your data.

upvoted 2 times

✉  **scanner2** 1 month ago

<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>

upvoted 1 times

✉  **mzchelino** 1 month, 2 weeks ago

Selected Answer: D

D!

Keys here are "dynamically calculated" and "automate".

upvoted 1 times

✉  **Shubha1** 2 months ago

why not A? can someone pls explain

upvoted 1 times

✉  **Husni_adam** 2 months, 3 weeks ago

Selected Answer: D

D is the right answer because you can automate process with export billing data to bigquery and then using the bigquery export to create the Looker studio dashboard

upvoted 2 times

✉  **Husni_adam** 2 months, 3 weeks ago

<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>

upvoted 2 times

✉  **MrJkr** 3 months, 2 weeks ago

Selected Answer: D

Option D closely aligns with the requirements mentioned in the question.

By configuring Cloud Billing data export to BigQuery, you can automate the process of exporting billing data to a BigQuery dataset. You can then use Looker Studio, a data visualization and exploration platform, to create a dashboard on top of the BigQuery export. This allows you to visualize costs with specific metrics that can be dynamically calculated based on company-specific criteria.

upvoted 4 times

You have an application that runs on Compute Engine VM instances in a custom Virtual Private Cloud (VPC). Your company's security policies only allow the use of internal IP addresses on VM instances and do not let VM instances connect to the internet. You need to ensure that the application can access a file hosted in a Cloud Storage bucket within your project. What should you do?

- A. Enable Private Service Access on the Cloud Storage Bucket.
- B. Add storage.googleapis.com to the list of restricted services in a VPC Service Controls perimeter and add your project to the list of protected projects.
- C. Enable Private Google Access on the subnet within the custom VPC.
- D. Deploy a Cloud NAT instance and route the traffic to the dedicated IP address of the Cloud Storage bucket.

Correct Answer: A*Community vote distribution*

C (100%)

✉️ 🚩 Captain1212 1 month ago**Selected Answer: C**

C is the correct Answer as Private Google Access allows you to connect on the internal networks, A is incorrect because Cloud Storage buckets don't have such services to connect to Private Access.

upvoted 1 times

✉️ 🚩 scanner2 1 month ago**Selected Answer: C**

Cloud Storage is not a supported service for Private Service Access. Hence, A cannot be the answer.
<https://cloud.google.com/vpc/docs/private-services-access#private-services-supported-services>

VM instances that only have internal IP addresses (no external IP addresses) can use Private Google Access. They can reach the external IP addresses of Google APIs and services. If you disable Private Google Access, the VM instances can no longer reach Google APIs and services; they can only send traffic within the VPC network.

<https://cloud.google.com/vpc/docs/private-google-access>

upvoted 2 times

✉️ 🚩 juliorevk 2 months, 1 week ago**Selected Answer: C**

Private Google Access lets you connect VM instances to GCP services without external IP addresses and only internal. A is wrong because even though Private Services Access lets you also access GCP and other services through internal IPs, it also allows the VMs to have external IPs.

<https://cloud.google.com/vpc/docs/private-google-access>

upvoted 4 times

✉️ 🚩 FJ82 3 months ago**Selected Answer: C**

Private Google Access is a VPC feature

upvoted 3 times

✉️ 🚩 TomFoot 3 months, 1 week ago**Selected Answer: C**

C allows access to Google services & API's

upvoted 1 times

✉️ 🚩 georgesouzafarias 3 months, 2 weeks ago**Selected Answer: C**

Right answer.

upvoted 3 times

Your company completed the acquisition of a startup and is now merging the IT systems of both companies. The startup had a production Google Cloud project in their organization. You need to move this project into your organization and ensure that the project is billed to your organization. You want to accomplish this task with minimal effort. What should you do?

- A. Use the projects.move method to move the project to your organization. Update the billing account of the project to that of your organization.
- B. Ensure that you have an Organization Administrator Identity and Access Management (IAM) role assigned to you in both organizations. Navigate to the Resource Manager in the startup's Google Cloud organization, and drag the project to your company's organization.
- C. Create a Private Catalog for the Google Cloud Marketplace, and upload the resources of the startup's production project to the Catalog. Share the Catalog with your organization, and deploy the resources in your company's project.
- D. Create an infrastructure-as-code template for all resources in the project by using Terraform, and deploy that template to a new project in your organization. Delete the project from the startup's Google Cloud organization.

Correct Answer: D*Community vote distribution*

A (100%)

✉️  **MrJkr**  3 months, 2 weeks ago**Selected Answer: A**

Option A is correct as it suggests using the "projects.move" method provided by Google Cloud to move the project from the startup's organization to your organization. This method allows you to transfer the ownership and control of a project to another organization. By moving the project, you can ensure that it is under your organization's management. While the other options contain elements that may be relevant in certain scenarios, they do not directly address the requirement of moving the project and ensuring billing to your organization.

upvoted 6 times

✉️  **Captain1212**  1 month ago**Selected Answer: A**

A is correct Answer you can use the project move method to move the project to your organization

upvoted 1 times

✉️  **scanner2** 1 month ago**Selected Answer: A**

Here, A is the only choice because neither B nor C nor D be the answer.

You can move a project within an organization. But you need to migrate a project across organizations.

<https://cloud.google.com/resource-manager/docs/project-migration>

<https://medium.com/google-cloud/migrating-a-project-from-one-organization-to-another-gcp-4b37a86dd9e6>

upvoted 2 times

✉️  **georgesouzafarias** 3 months, 2 weeks ago**Selected Answer: A**

Don't overthink it.

<https://cloud.google.com/resource-manager/docs/project-migration-checklist>

upvoted 4 times

All development (dev) teams in your organization are located in the United States. Each dev team has its own Google Cloud project. You want to restrict access so that each dev team can only create cloud resources in the United States (US). What should you do?

- A. Create a folder to contain all the dev projects. Create an organization policy to limit resources in US locations.
- B. Create an organization to contain all the dev projects. Create an Identity and Access Management (IAM) policy to limit the resources in US regions.
- C. Create an Identity and Access Management (IAM) policy to restrict the resources locations in the US. Apply the policy to all dev projects.
- D. Create an Identity and Access Management (IAM) policy to restrict the resources locations in all dev projects. Apply the policy to all dev roles.

Correct Answer: C*Community vote distribution*

A (74%)

C (26%)

  **joao_01** 2 weeks, 4 days ago**Selected Answer: A**

Its A.

"Organization Policy" does not indicate that it will be ONLY applied to a organization, it can be applied to any resource within a organization to restrict and add conditions. This policy focus on WHAT and not WHO (IAM). So, since in this case we want to restrict to VMs in US, its clearly the option A.

Link: <https://cloud.google.com/resource-manager/docs/organization-policy/overview>

"Identity and Access Management focuses on who, and lets the administrator authorize who can take action on specific resources based on permissions."

"Organization Policy focuses on what, and lets the administrator set restrictions on specific resources to determine how they can be configured."

upvoted 2 times

  **ExamsFR** 3 weeks, 6 days ago**Selected Answer: A**

<https://cloud.google.com/resource-manager/docs/organization-policy/defining-locations?hl=fr#gcloud>

upvoted 1 times

  **Captain1212** 1 month ago**Selected Answer: A**

A is the correct answer as you can just add the projects to the same folder and just apply the Organization policy

upvoted 1 times

  **DocOck** 1 month ago

The organization policy is set on an organization, folder, or project resource to enforce the constraint on that resource and any child resources. An organization policy contains one or more rules that specify how, and whether, to enforce the constraint.

A

upvoted 1 times

  **scanner2** 1 month ago**Selected Answer: A**

Restricting resources creation in specific locations can be achieved by Organization policy. The Organization Policy Service gives you centralized and programmatic control over your organization's cloud resources.

You can use "Google Cloud Platform - Resource Location Restriction" Organization policy constraint for the solution asked in the question.
<https://cloud.google.com/resource-manager/docs/organization-policy/overview>

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints>

upvoted 2 times

  **scanner2** 1 month ago

Option A is mentioning to create a folder that contains all dev projects.

upvoted 1 times

  **LONGBOW_RA** 1 month, 2 weeks ago**Selected Answer: C**

<https://cloud.google.com/resource-manager/docs/access-control-folders>

A is wrong, The documentation says it needs a folder admin role to apply folder's IAM policy to limit resources access. A said organization level policies.

upvoted 1 times

  **LONGBOW RA** 1 month, 2 weeks ago

OK, I changed my mind. A folder contains projects. Then mange the folder's IAM policies maybe can be considered as "organization level".

WTF, I feel I am getting an English certification.

upvoted 1 times

✉️ **Jerica_** 1 month, 3 weeks ago

Selected Answer: C

c as it cannot be A as it will apply to all folders

upvoted 2 times

✉️ **3arle** 1 month, 3 weeks ago

Selected Answer: A

Organization policies are made up of constraints that allow you to:

Limit resource sharing based on domain.

Limit the usage of Identity and Access Management service accounts.

Restrict the physical location of newly created resources!!!!!!

<https://cloud.google.com/resource-manager/docs/organization-policy/overview>

upvoted 2 times

✉️ **lululala007** 2 months ago

Selected Answer: C

C is the correct Answer,

It CANNOT be A as setting an organization policy will restrict every single project in the organization and not only the dev projects in the folder.

It CANNOT be B either, because projects can only even be created if an organization already exists.

upvoted 2 times

✉️ **shreykul** 2 months, 2 weeks ago

Selected Answer: A

Option A

upvoted 1 times

✉️ **MrJkr** 3 months, 2 weeks ago

Selected Answer: A

Option A is the most suitable answer among the provided choices.

By creating a folder to contain all the dev projects, you can organize them in a logical structure within your organization. Then, you can apply an organization policy to limit the resources in US locations. This policy can be configured to restrict the creation of cloud resources outside the United States. It provides a centralized approach to enforce the restriction across all the dev projects within the folder.

upvoted 3 times

✉️ **georgesouzafarias** 3 months, 2 weeks ago

Selected Answer: A

You need to use "Google Cloud Platform - Resource Location Restriction" organization policy.

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints>

upvoted 2 times

You are configuring Cloud DNS. You want to create DNS records to point home.mydomain.com, mydomain.com, and www.mydomain.com to the IP address of your Google Cloud load balancer. What should you do?

- A. Create one CNAME record to point mydomain.com to the load balancer, and create two A records to point WWW and HOME to mydomain.com respectively.
- B. Create one CNAME record to point mydomain.com to the load balancer, and create two AAAA records to point WWW and HOME to mydomain.com respectively.
- C. Create one A record to point mydomain.com to the load balancer, and create two CNAME records to point WWW and HOME to mydomain.com respectively.
- D. Create one A record to point mydomain.com to the load balancer, and create two NS records to point WWW and HOME to mydomain.com respectively.

Correct Answer: B*Community vote distribution*

C (100%)

  **Captain1212** 1 month ago**Selected Answer: C**

C is the correct answer, as for this you need to create One A record and two CNAME records
upvoted 1 times

  **scanner2** 1 month ago**Selected Answer: C**

Record name | Type | Value
mydomain.com | A | load balancer IP
www.mydomain.com | CNAME | mydomain.com
home.mydomain.com | CNAME | mydomain.com
https://cloud.google.com/dns/docs/records-overview#supported_dns_record_types
upvoted 2 times

  **3arle** 2 months ago**Selected Answer: C**

A record is always for load balancer IP, then CNAME
upvoted 1 times

  **juliorevk** 2 months, 1 week ago**Selected Answer: C**

C because the A record which points a domain name to an IPv4 address is used for the main domain (mydomain.com), the other two for www and home are aliases (CNAME) that can be pointed to the main domain of mydomain.com.
upvoted 2 times

  **MrJkr** 3 months, 2 weeks ago**Selected Answer: C**

Option A suggests creating one CNAME record to point mydomain.com to the load balancer, which is incorrect because CNAME records cannot coexist with other record types on the same domain/subdomain. In this case, you need to use an A record instead.

Option B suggests creating two AAAA records, which are used for IPv6 addresses. Unless you specifically have an IPv6 address for your load balancer, using AAAA records would not be appropriate.

Option D suggests creating two NS records, which are used for specifying the authoritative name servers for a domain. NS records are not used to point subdomains to IP addresses or load balancers.

Therefore, option C is the correct answer, as it correctly suggests creating one A record to point mydomain.com to the load balancer, and two CNAME records to point WWW and HOME to mydomain.com respectively.

upvoted 4 times

  **georgesouzafarias** 3 months, 2 weeks ago**Selected Answer: C**

You can only associate A(IP) record to a domain.

https://cloud.google.com/dns/docs/set-up-dns-records-domain-name#create_a_record_to_point_the_domain_to_an_external_ip_address
upvoted 1 times

You have two subnets (subnet-a and subnet-b) in the default VPC. Your database servers are running in subnet-a. Your application servers and web servers are running in subnet-b. You want to configure a firewall rule that only allows database traffic from the application servers to the database servers. What should you do?

- A. • Create service accounts sa-app and sa-db.
 - Associate service account sa-app with the application servers and the service account sa-db with the database servers.
 - Create an ingress firewall rule to allow network traffic from source service account sa-app to target service account sa-db.
- B. • Create network tags app-server and db-server.
 - Add the app-server tag to the application servers and the db-server tag to the database servers.
 - Create an egress firewall rule to allow network traffic from source network tag app-server to target network tag db-server.
- C. • Create a service account sa-app and a network tag db-server.
 - Associate the service account sa-app with the application servers and the network tag db-server with the database servers.
 - Create an ingress firewall rule to allow network traffic from source VPC IP addresses and target the subnet-a IP addresses.
- D. • Create a network tag app-server and service account sa-db.
 - Add the tag to the application servers and associate the service account with the database servers.
 - Create an egress firewall rule to allow network traffic from source network tag app-server to target service account sa-db.

Correct Answer: B*Community vote distribution*

A (57%)	B (43%)
---------	---------

✉  **Captain1212** 1 month ago

Selected Answer: A

Answer A is correct as question demands traffic to application to database which can be only be achieved by the ingreess rule
upvoted 1 times

✉  **scanner2** 1 month ago

Selected Answer: A

Both service accounts and network tags can be used for creating a Cloud Firewall rule. The prime word is "to allow network traffic from app server to database server" which is achievable by inbound/ingress rule and not egress rule.
https://cloud.google.com/firewall/docs/firewalls#rule_assignment
upvoted 1 times

✉  **tatyavinchu** 1 month, 3 weeks ago

Correct Answer is A
upvoted 1 times

✉  **3arle** 2 months ago

Selected Answer: A

From the TomFoot link
'for example, allow my "application x" servers to access my "database y."'
upvoted 1 times

✉  **juliorevk** 2 months, 1 week ago

Selected Answer: B

Even though you could use service accounts for firewall rules, why is B wrong? It seems to do what the question requests and is the standard method.
upvoted 3 times

✉  **Ahmed_Y** 1 month, 2 weeks ago

Because we need an ingress firewall.
upvoted 1 times

✉  **_F4LLEN_** 3 months, 1 week ago

Selected Answer: A

Service accs can be used for firewall management.
upvoted 1 times

✉  **TomFoot** 3 months, 1 week ago

Selected Answer: A

You can use service for firewall rules.

<https://cloud.google.com/blog/products/gcp/simplify-cloud-vpc-firewall-management-with-service-accounts>

upvoted 2 times

 **gpais** 3 months, 1 week ago

Selected Answer: A

A seems to be the most appropriate: <https://cloud.google.com/firewall/docs/firewalls>
upvoted 2 times

 **gpais** 3 months, 1 week ago

A seems to be the most appropriate: <https://cloud.google.com/firewall/docs/firewalls>
upvoted 2 times

 **georgesouzafarias** 3 months, 2 weeks ago

Selected Answer: B

Service account? It doesn't make any sense.
It's clearly a firewall solution.

upvoted 3 times

 **TomFoot** 3 months, 1 week ago

You can use service for firewall rules.

<https://cloud.google.com/blog/products/gcp/simplify-cloud-vpc-firewall-management-with-service-accounts>
upvoted 2 times

Question #207

Topic 1

Your team wants to deploy a specific content management system (CMS) solution to Google Cloud. You need a quick and easy way to deploy and install the solution. What should you do?

- A. Search for the CMS solution in Google Cloud Marketplace. Use gcloud CLI to deploy the solution.
- B. Search for the CMS solution in Google Cloud Marketplace. Deploy the solution directly from Cloud Marketplace.
- C. Search for the CMS solution in Google Cloud Marketplace. Use Terraform and the Cloud Marketplace ID to deploy the solution with the appropriate parameters.
- D. Use the installation guide of the CMS provider. Perform the installation through your configuration management system.

Correct Answer: C

Community vote distribution

B (100%)

 **juliorevk** 2 months, 1 week ago

Selected Answer: B

Fastest and easiest way to deploy a solution straight from the marketplace
upvoted 2 times

 **on2it** 3 months, 1 week ago

Selected Answer: B

Indeed directly from Cloud Marketplace
upvoted 1 times

 **_F4LLEN_** 3 months, 1 week ago

Selected Answer: B

We can deploy it directly from Cloud Marketplace.
upvoted 1 times

 **TomFoot** 3 months, 1 week ago

Selected Answer: B

I think B would be simple and fast.
upvoted 1 times

 **georgesouzafarias** 3 months, 2 weeks ago

Selected Answer: B

Key Words: Quick and Easy
upvoted 2 times

You are working for a startup that was officially registered as a business 6 months ago. As your customer base grows, your use of Google Cloud increases. You want to allow all engineers to create new projects without asking them for their credit card information. What should you do?

- A. Create a Billing account, associate a payment method with it, and provide all project creators with permission to associate that billing account with their projects.
- B. Grant all engineers permission to create their own billing accounts for each new project.
- C. Apply for monthly invoiced billing, and have a single invoice for the project paid by the finance team.
- D. Create a billing account, associate it with a monthly purchase order (PO), and send the PO to Google Cloud.

Correct Answer: A

Community vote distribution

A (100%)

 **scanner2** 1 month ago

Selected Answer: A

Cloud Billing accounts pay for Google Cloud projects and Google Maps Platform projects.

You might be eligible to switch your account type to monthly invoiced billing if your business meets certain requirements. These requirements include, but aren't limited to the following:

- You must be registered as a business for a minimum of one year.
- You expect to spend a minimum of \$40,000 a year on Google Cloud.
- Invoiced billing must be available in your country.

<https://cloud.google.com/billing/docs/how-to/invoiced-billing>

So A is the correct answer here.

upvoted 1 times

 **_F4LLEN_** 3 months, 1 week ago

Selected Answer: A

A seems the best among the rest.

upvoted 1 times

 **MrJkr** 3 months, 2 weeks ago

Selected Answer: A

Option A is the better answer for the given scenario. It allows you to centralize billing and payment management while providing flexibility to project creators. By creating a billing account and associating a payment method with it, you establish a central source for billing and payment for all projects.

Granting project creators permission to associate the billing account with their projects ensures that they can create projects without the need for their individual credit card information. This approach streamlines the process and avoids the hassle of collecting credit card details from each engineer.

Additionally, this option allows for easy monitoring and management of project costs through a single billing account, making it simpler to track expenses and allocate resources effectively.

upvoted 2 times

 **gpais** 3 months, 2 weeks ago

Selected Answer: A

It's option A

upvoted 1 times

 **georgesouzafarias** 3 months, 2 weeks ago

Selected Answer: A

Don't overthink.

upvoted 1 times

Your continuous integration and delivery (CI/CD) server can't execute Google Cloud actions in a specific project because of permission issues. You need to validate whether the used service account has the appropriate roles in the specific project.

What should you do?

- A. Open the Google Cloud console, and check the Identity and Access Management (IAM) roles assigned to the service account at the project or inherited from the folder or organization levels.
- B. Open the Google Cloud console, and check the organization policies.
- C. Open the Google Cloud console, and run a query to determine which resources this service account can access.
- D. Open the Google Cloud console, and run a query of the audit logs to find permission denied errors for this service account.

Correct Answer: A

Community vote distribution

A (80%)

C (20%)

 **joao_01** 3 weeks, 5 days ago

A it's really straightforward
upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: A
A is the correct Answer
upvoted 1 times

 **happydays** 2 months, 1 week ago

Selected Answer: A
It's A from GPT
upvoted 1 times

 **TungTrinhSenacor** 2 months, 1 week ago

Selected Answer: A
I go for A
upvoted 2 times

 **ankitb4u** 2 months, 1 week ago

Selected Answer: C
C should be the correct answer here
upvoted 1 times

 **TungTrinhSenacor** 2 months, 1 week ago

Why not A?
upvoted 1 times

Your team is using Linux instances on Google Cloud. You need to ensure that your team logs in to these instances in the most secure and cost efficient way. What should you do?

- A. Attach a public IP to the instances and allow incoming connections from the internet on port 22 for SSH.
- B. Use the gcloud compute ssh command with the --tunnel-through-iap flag. Allow ingress traffic from the IP range 35.235.240.0/20 on port 22.
- C. Use a third party tool to provide remote access to the instances.
- D. Create a bastion host with public internet access. Create the SSH tunnel to the instance through the bastion host.

Correct Answer: D*Community vote distribution*

✉️ **Captain1212** 1 month ago

Selected Answer: B

Common sense B is the correct answer , must safer from using the third party apps or the public id addresses
upvoted 2 times

✉️ **3arle** 2 months ago

Selected Answer: B

You can use Bastion if
"You have a specific use case, like session recording, and you can't use IAP."
<https://cloud.google.com/compute/docs/connect/ssh-internal-ip>
upvoted 2 times

✉️ **Ahmed_Y** 1 month, 1 week ago

Thanks for that link but I think it is C, Although totally agree that Bastion comes 2nd in that table, no way all the user would have IP within this range 35.235.240.0/20!

upvoted 1 times

✉️ **itsimranmalik** 1 month ago

"allows ingress traffic from the IP range `35.235.240.0/20`. This range contains all IP addresses that IAP uses for TCP forwarding"
<https://cloud.google.com/iap/docs/using-tcp-forwarding#create-firewall-rule>

upvoted 1 times

✉️ **qannik** 2 months ago

Selected Answer: D

But the question states "You need to ensure that your team logs in to these instances in the most secure and cost efficient way"
Bastion is more secure than IAP but I'm not sure is more cost effective...
Hard to choose

upvoted 1 times

✉️ **juliorevk** 2 months, 1 week ago

Understood about IAP being a secure way to SSH but where did the "Allow ingress traffic from the IP range 35.235.240.0/20 on port 22." come from and how does that fit in? The question had no details about it and the IP range seemed to come out of nowhere.
upvoted 4 times

✉️ **Husni_adam** 2 months, 3 weeks ago

Selected Answer: B

<https://cloud.google.com/compute/docs/connect/ssh-using-iap#gcloud>
according the documentation the correct answer is B
upvoted 3 times

✉️ **techsteph** 2 months, 3 weeks ago

Selected Answer: B

<https://cloud.google.com/compute/docs/connect/ssh-using-iap#gcloud>
upvoted 2 times

An external member of your team needs list access to compute images and disks in one of your projects. You want to follow Google-recommended practices when you grant the required permissions to this user. What should you do?

- A. Create a custom role, and add all the required compute.disks.list and compute.images.list permissions as includedPermissions. Grant the custom role to the user at the project level.
- B. Create a custom role based on the Compute Image User role. Add the compute.disks.list to the includedPermissions field. Grant the custom role to the user at the project level.
- C. Create a custom role based on the Compute Storage Admin role. Exclude unnecessary permissions from the custom role. Grant the custom role to the user at the project level.
- D. Grant the Compute Storage Admin role at the project level.

Correct Answer: A

Community vote distribution

A (62%)	B (23%)	C (15%)
---------	---------	---------

 **rahulrauki** 1 week, 6 days ago

Selected Answer: A

You can't give B because, Image user will be able to use the Image to create resources. Only give list access
upvoted 1 times

 **joao_01** 3 weeks, 5 days ago

Its A. Give user ONLY the required permission.
upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: A

A is the correct answer , just create the custom role add all the required permissoons , give to the user
upvoted 1 times

 **demoro86** 1 month, 2 weeks ago

Selected Answer: A

I have successfully created a custom role with compute.disks.list and compute.image.list permissions. I have also tried creating it based on the Compute Storage Admin role. However, you still need to select compute.disks.list and compute.image.list individually; all permissions are unchecked by default. So A fits fine.
upvoted 1 times

 **3arle** 2 months ago

Selected Answer: A

The key word is "needs list access", so only A meets this requirement
upvoted 2 times

 **juliorevk** 2 months, 1 week ago

Selected Answer: A

<https://cloud.google.com/iam/docs/custom-roles-permissions-support> - Both compute.disks.list and compute.images.list are available as permissions for custom roles. Makes more sense to make a new custom role than going off an admin one then adjusting it.
upvoted 1 times

 **shreykul** 2 months, 2 weeks ago

Selected Answer: B

Option B allows you to create a custom role that is based on the existing Compute Image User role, which already includes the necessary permissions for accessing compute images. Then, you add the compute.disks.list permission to the custom role's includedPermissions field to grant the user list access to compute disks as well. This ensures that the user has precisely the permissions needed for their specific tasks and nothing more, following the principle of least privilege.
upvoted 3 times

 **shreykul** 2 months, 2 weeks ago

Selected Answer: A

<https://cloud.google.com/sdk/gcloud/reference/compute/images/list>

<https://cloud.google.com/compute/docs/reference/rest/v1/disks/list>
upvoted 1 times

 **fatanu88** 2 months, 2 weeks ago

Answer is B: Compute image user role provide permission to list and read images without having other permissions on the image. Granting this role at the project level gives users the ability to list all images in the project and create resources, such as instances and persistent disks, based on images in the project. Adding the compute.disks.list then meet all the question requirements

upvoted 2 times

✉️ **FJ82** 2 months, 3 weeks ago

Selected Answer: C

Tried this, could not find those permissions when I tried to create custom role directly, you need to create from the role
upvoted 2 times

✉️ **techsteph** 2 months, 2 weeks ago

You're right, changing my answer to C.

upvoted 1 times

✉️ **techsteph** 2 months, 3 weeks ago

Selected Answer: A

<https://cloud.google.com/compute/docs/access/iam>

upvoted 1 times

You are running a web application on Cloud Run for a few hundred users. Some of your users complain that the initial web page of the application takes much longer to load than the following pages. You want to follow Google's recommendations to mitigate the issue. What should you do?

- A. Set the minimum number of instances for your Cloud Run service to 3.
- B. Set the concurrency number to 1 for your Cloud Run service.
- C. Set the maximum number of instances for your Cloud Run service to 100.
- D. Update your web application to use the protocol HTTP/2 instead of HTTP/1.1.

Correct Answer: A*Community vote distribution*

A (71%)	D (29%)
---------	---------

✉  **joao_01** 3 weeks, 5 days ago

Its A. Look at this link: https://cloud.google.com/run/docs/tips/general#optimize_performance

You'll see that one of Google's recommendations for improve performance and reduce "cold starts", is to set a minimum of instances.
upvoted 1 times

✉  **DannSecurity** 1 month ago

I checked the google recommendations for cloud run but they never mention HTTP/2. I am going with A
upvoted 1 times

✉  **Cherrycardo** 2 months ago

Selected Answer: A

<https://cloud.google.com/functions/docs/configuring/min-instances>

Even though the initial # of instances present in the VM is not stated, setting a min amount of instances "can further help you avoid cold starts and reduce application latency".

upvoted 1 times

✉  **3arle** 2 months ago

Selected Answer: A

min. # of instances should solve latency issue

upvoted 1 times

✉  **qannik** 2 months ago

Selected Answer: A

I chose A too.

The question should be related to GCP products and how can you best configure them.

I shouldn't be forced to change my app to use HTTP/2

upvoted 2 times

✉  **qannik** 2 months ago

<https://cloud.google.com/blog/topics/developers-practitioners/3-ways-optimize-cloud-run-response-times>

upvoted 1 times

✉  **Vzero2333** 2 months, 1 week ago

I chosse A

notice "some of user" and "initial web page slowly than other page"

because the some of user login the web without load and the cloud run scale the instance to 0 so that the users had to waiting the startup new instance sometimes.

upvoted 1 times

✉  **geeroylenkins** 2 months, 3 weeks ago

Selected Answer: D

<https://www.cloudflare.com/learning/performance/http2-vs-http1.1/>

I'm going with D too.

B wouldn't help: <https://cloud.google.com/run/docs/about-concurrency#concurrency-1>

C wouldn't help - setting a max won't increase speed ever.

A would not necessarily help - there's no indication that the initial page is taking much longer just because there are too few instances. However, D would improve how things load, as per the first link I posted.

upvoted 2 times

✉  **Speridian** 2 months, 3 weeks ago

It should be D. HTTP/2 helps to speed up app performance.

upvoted 1 times

✉  **techsteph** 2 months, 3 weeks ago

Selected Answer: A

<https://cloud.google.com/run/docs/configuring/min-instances>

upvoted 1 times

Question #213

Topic 1

You are building a data lake on Google Cloud for your Internet of Things (IoT) application. The IoT application has millions of sensors that are constantly streaming structured and unstructured data to your backend in the cloud. You want to build a highly available and resilient architecture based on Google-recommended practices. What should you do?

- A. Stream data to Pub/Sub, and use Dataflow to send data to Cloud Storage.
- B. Stream data to Pub/Sub, and use Storage Transfer Service to send data to BigQuery.
- C. Stream data to Dataflow, and use Dataprep by Trifacta to send data to Bigtable.
- D. Stream data to Dataflow, and use Storage Transfer Service to send data to BigQuery.

Correct Answer: A

Community vote distribution

A (100%)

✉  **joao_01** 3 weeks, 5 days ago

Selected Answer: A

Its A, for sure.

upvoted 1 times

✉  **Captain1212** 1 month ago

A is the correct answer as there is both unstructured and structured data

upvoted 1 times

✉  **3arle** 2 months ago

Selected Answer: A

according to <https://cloud.google.com/architecture/optimized-large-scale-analytics-ingestion>

upvoted 2 times

✉  **shreykul** 2 months, 2 weeks ago

Selected Answer: A

A. Streaming data to Pub/Sub allows you to decouple the ingestion of data from the processing and storage, providing a scalable and reliable message queue that can handle the high volume of data coming from millions of sensors.

Using Dataflow to consume data from Pub/Sub and send it to Cloud Storage allows for real-time data processing and storage. Dataflow is a fully managed service for processing data in real-time or batch mode, making it an ideal choice for handling the constant stream of data from IoT sensors.

Storing data in Cloud Storage offers high durability and availability, providing a robust foundation for building a data lake. Cloud Storage is a scalable object storage service that can handle large volumes of structured and unstructured data, making it well-suited for the IoT application's data requirements.

upvoted 2 times

You are running out of primary internal IP addresses in a subnet for a custom mode VPC. The subnet has the IP range 10.0.0.0/20, and the IP addresses are primarily used by virtual machines in the project. You need to provide more IP addresses for the virtual machines. What should you do?

- A. Add a secondary IP range 10.1.0.0/20 to the subnet.
- B. Change the subnet IP range from 10.0.0.0/20 to 10.0.0.0/18.
- C. Change the subnet IP range from 10.0.0.0/20 to 10.0.0.0/22.
- D. Convert the subnet IP range from IPv4 to IPv6.

Correct Answer: A*Community vote distribution*

B (70%)

A (30%)

joao_01 2 weeks, 3 days ago**Selected Answer: A**

This one is tricky. First i was going with B, then i did some search. Option A and B can indeed add more IPs. However, i think the option is A because between those 2 options the option A we will add IPs without changing the any ours VMs configurations that we currently have. If we choose B might need to change our current VMs configuration in order to reflect the new IP range expanded. You guys understand what i mean?

Link: <https://cloud.google.com/vpc/docs/create-modify-vpc-networks#expand-subnet>

"If you expand the primary IPv4 range of a subnet, you might need to modify other configurations that are assuming this IP address range"

upvoted 2 times

Captain1212 1 month ago

B is the correct answer

upvoted 1 times

scanner2 1 month ago**Selected Answer: B**

<https://cloud.google.com/vpc/docs/create-modify-vpc-networks#expand-subnet>

<https://cloud.google.com/sdk/gcloud/reference/compute/networks/subnets/expand-ip-range>

upvoted 1 times

vinodthakur49 1 month, 3 weeks ago**Selected Answer: A**

The ask is to add more IPs, it's not about to keep the VM ins same subnet or to keep communication open between existing and new VMs, so the answer should be A...

upvoted 1 times

joao_01 3 weeks, 5 days ago

By changing the IP range to ..18 it will add more IPs to the subnet. So the answer B is correct.

upvoted 1 times

3arle 2 months ago**Selected Answer: B**

<https://cloud.google.com/vpc/docs/create-modify-vpc-networks#expand-subnet>

upvoted 1 times

shreykul 2 months, 2 weeks ago**Selected Answer: B**

Clearly B

upvoted 1 times

shreykul 2 months, 2 weeks ago**Selected Answer: B**

Should be B for subnet expansion

upvoted 2 times

geeroylenkins 2 months, 3 weeks ago**Selected Answer: B**

<https://cloud.google.com/sdk/gcloud/reference/compute/networks/subnets/expand-ip-range>

<https://cloud.google.com/vpc/docs/create-modify-vpc-networks#expand-subnet>

<https://techlibrary.hpe.com/docs/otlink-wo/CIDR-Conversion-Table.html>

Expanding the subnet to a smaller prefix will allow more primary internal IP addresses. I don't know why gw2100 "withdraws" his vote for B and changes it to A - it's possible to add a secondary IP range as in answer A, but expanding the initial range as in B would work just fine.

C is shrinking the range - cannot work.

Not sure about D

upvoted 2 times

 **gw2100** 2 months, 3 weeks ago

should be B. you expand the ip with a bigger network

upvoted 1 times

 **gw2100** 2 months, 3 weeks ago

I withdraw my vote. It should be A.

upvoted 1 times

Your company requires all developers to have the same permissions, regardless of the Google Cloud project they are working on. Your company's security policy also restricts developer permissions to Compute Engine, Cloud Functions, and Cloud SQL. You want to implement the security policy with minimal effort. What should you do?

- A. • Create a custom role with Compute Engine, Cloud Functions, and Cloud SQL permissions in one project within the Google Cloud organization.
 - Copy the role across all projects created within the organization with the gcloud iam roles copy command.
 - Assign the role to developers in those projects.
- B. • Add all developers to a Google group in Google Groups for Workspace.
 - Assign the predefined role of Compute Admin to the Google group at the Google Cloud organization level.
- C. • Add all developers to a Google group in Cloud Identity.
 - Assign predefined roles for Compute Engine, Cloud Functions, and Cloud SQL permissions to the Google group for each project in the Google Cloud organization.
- D. • Add all developers to a Google group in Cloud Identity.
 - Create a custom role with Compute Engine, Cloud Functions, and Cloud SQL permissions at the Google Cloud organization level.
 - Assign the custom role to the Google group.

Correct Answer: D

Community vote distribution

D (75%)

C (25%)

 **joao_01** 3 weeks, 5 days ago

Selected Answer: D

I vote for D
upvoted 1 times

 **joao_01** 3 weeks, 5 days ago

I vote for D
upvoted 1 times

 **Captain1212** 1 month ago

d is the correct answer
upvoted 1 times

 **scanner2** 1 month ago

Selected Answer: D

Permissions provided at the Organization level are inherited to the folder level and project level.
upvoted 1 times

 **gpais** 1 month, 3 weeks ago

Selected Answer: C

Use predefined roles: Use predefined roles, such as "Editor" or "Viewer", instead of creating custom roles. This makes it easier to understand the level of access associated with a role.
Use custom roles: Create custom roles when predefined roles do not meet the specific needs of your organization.

In the link below:

<https://cloud.google.com/iam/docs/roles-overview#custom>

When to use custom roles

In most situations, you should be able to use predefined roles instead of custom roles. Predefined roles are maintained by Google, and are updated automatically when new permissions, features, or services are added to Google Cloud. In contrast, custom roles are not maintained by Google; when Google Cloud adds new permissions, features, or services, your custom roles will not be updated automatically.

upvoted 2 times

 **3arle** 2 months ago

Selected Answer: D

only D meets best practices
upvoted 1 times

 **shreykul** 2 months, 2 weeks ago

Selected Answer: D

<https://www.cloudskillsboost.google/focuses/1035?parent=catalog#:~:text=custom%20role%20at%20the%20organization%20level>
upvoted 3 times

You are working for a hospital that stores its medical images in an on-premises data room. The hospital wants to use Cloud Storage for archival storage of these images. The hospital wants an automated process to upload any new medical images to Cloud Storage. You need to design and implement a solution. What should you do?

- A. Create a Pub/Sub topic, and enable a Cloud Storage trigger for the Pub/Sub topic. Create an application that sends all medical images to the Pub/Sub topic.
- B. Create a script that uses the gcloud storage command to synchronize the on-premises storage with Cloud Storage. Schedule the script as a cron job.
- C. Create a Pub/Sub topic, and create a Cloud Function connected to the topic that writes data to Cloud Storage. Create an application that sends all medical images to the Pub/Sub topic.
- D. In the Google Cloud console, go to Cloud Storage. Upload the relevant images to the appropriate bucket.

Correct Answer: B*Community vote distribution*

B (70%)

C (30%)

joao_01 3 weeks, 5 days ago**Selected Answer: B**

This question already appear before. But before the B option was with "gsutil" command, now is with the "gcloud" storage command. Despite this difference i believe the answer is still B cause this new command is more or less the same as the "gsutil". Look at these link:
<https://cloud.google.com/blog/products/storage-data-transfer/new-gcloud-storage-cli-for-your-data-transfers>

upvoted 1 times

kavi444 1 month, 3 weeks ago

B is not correct , as it states using gcloud command. One must use gsutil instead.
IMO, C is correct.

upvoted 1 times

markoniz 1 month, 3 weeks ago

I believe that you could use "gcloud storage cp" command as alternative to gsutil
<https://cloud.google.com/sdk/gcloud/reference/storage/cp>
<https://cloud.google.com/blog/products/storage-data-transfer/new-gcloud-storage-cli-for-your-data-transfers>

upvoted 1 times

3arle 2 months ago**Selected Answer: B**

Repeated question
upvoted 1 times

happydays 2 months, 1 week ago**Selected Answer: B**

THIS IS B
upvoted 3 times

gpais 2 months, 1 week ago**Selected Answer: B**

We are dealing with image files so google cloud storage seems more appropriate
upvoted 2 times

shreykul 2 months, 2 weeks ago**Selected Answer: C**

Option C is more robust and utilises the GCP functionalities correctly.
upvoted 3 times

qannik 2 months ago

You can not send images to a Pub/Sub topic.
upvoted 2 times

Your company has an internal application for managing transactional orders. The application is used exclusively by employees in a single physical location. The application requires strong consistency, fast queries, and ACID guarantees for multi-table transactional updates. The first version of the application is implemented in PostgreSQL, and you want to deploy it to the cloud with minimal code changes. Which database is most appropriate for this application?

- A. Bigtable
- B. BigQuery
- C. Cloud SQL
- D. Firestore

Correct Answer: C

Community vote distribution

C (100%)

✉️  **xhilmi** 3 weeks, 2 days ago

So what about this another question?

Your company has an internal application for managing transactional orders. The application is used exclusively by employees in a single physical location. The application requires strong consistency, fast queries, and ACID guarantees for multi-table transactional updates. The first version of the application is implemented in PostgreSQL, and you want to deploy it to the cloud with minimal code changes. Which database is most appropriate for this application?

- A. BigQuery
- B. Cloud SQL
- C. Cloud Spanner
- D. Cloud Datastore

It should be Cloud Spanner or Cloud SQL ?

upvoted 1 times

✉️  **joao_01** 2 weeks, 3 days ago

Cloud SQL

upvoted 2 times

✉️  **3arle** 2 months ago

Selected Answer: C

ACID and strong consistency are in C or D, but Firestore is for documents and in question we have "multi-table updates" so there left C
upvoted 2 times

✉️  **rahulrauki** 1 week, 6 days ago

You missed the main "postgres" part but yeah

upvoted 1 times

✉️  **pritampanda1988** 2 months, 1 week ago

Selected Answer: C

Cloud SQL is the most appropriate choice for deploying the application with the required characteristics while minimizing code changes and maintaining strong consistency, fast queries, and ACID guarantees for multi-table transactional updates.

upvoted 1 times

Your company runs one batch process in an on-premises server that takes around 30 hours to complete. The task runs monthly, can be performed offline, and must be restarted if interrupted. You want to migrate this workload to the cloud while minimizing cost. What should you do?

- A. Create an Instance Template with Spot VMs On. Create a Managed Instance Group from the template and adjust Target CPU Utilization.
Migrate the workload.
- B. Migrate the workload to a Compute Engine VM. Start and stop the instance as needed.
- C. Migrate the workload to a Google Kubernetes Engine cluster with Spot nodes.
- D. Migrate the workload to a Compute Engine Spot VM.

Correct Answer: A*Community vote distribution***B (100%)**

✉  **joao_01** 3 weeks, 5 days ago

Selected Answer: B

Must be B. Stop instances are cost-effective yes and are for batch jobs. HOWEVER, after 24h teh instances will be change and, thus, the job will restart again and again after 24h. The only option that dont use Stop instances are the B. So, my answer is B.

upvoted 1 times

✉  **rahulrauki** 1 week, 6 days ago

Only preemptive VMs have that 24h time limit, the newer version Spot VMs don't have the time limit. But answer is still B because, they are not fault tolerant

upvoted 1 times

✉  **happydays** 2 months, 1 week ago

Selected Answer: B

B is answer

upvoted 1 times

✉  **fdelacortina** 2 months, 1 week ago

Why don't to use Spot? Is the most cost-effective option here due they said "minimizing the cost"

upvoted 2 times

✉  **Gregwaw** 1 month ago

Spot VMs are highly affordable compute instances suitable for batch jobs and fault-tolerant workloads.

In the case from question jobs need to be restarted and last 30 hours, so it's not fault-tolerant.

upvoted 1 times

✉  **ChungLe** 2 months, 1 week ago

Selected Answer: B

B is correct

upvoted 1 times

✉  **shreykul** 2 months, 2 weeks ago

Selected Answer: B

B. Migrating the workload to a Compute Engine VM and starting and stopping the instance as needed allows you to control when the task runs. This approach provides flexibility in terms of when to initiate the batch process, and it can be easily scheduled to run monthly. By stopping the instance when the task is not running, you can save on compute costs.

upvoted 4 times

You are planning to migrate the following on-premises data management solutions to Google Cloud:

- One MySQL cluster for your main database
- Apache Kafka for your event streaming platform
- One Cloud SQL for PostgreSQL database for your analytical and reporting needs

You want to implement Google-recommended solutions for the migration. You need to ensure that the new solutions provide global scalability and require minimal operational and infrastructure management. What should you do?

- A. Migrate from MySQL to Cloud SQL, from Kafka to Pub/Sub, and from Cloud SQL for PostgreSQL to BigQuery.
- B. Migrate from MySQL to Cloud Spanner, from Kafka to Pub/Sub, and from Cloud SQL for PostgreSQL to BigQuery.
- C. Migrate from MySQL to Cloud Spanner, from Kafka to Memorystore, and from Cloud SQL for PostgreSQL to Cloud SQL.
- D. Migrate from MySQL to Cloud SQL, from Kafka to Memorystore, and from Cloud SQL for PostgreSQL to Cloud SQL.

Correct Answer: A

Community vote distribution

B (100%)

 **ExamsFR** 2 weeks, 5 days ago

Selected Answer: B

B is the correct answer.
upvoted 1 times

 **joao_01** 3 weeks, 5 days ago

Selected Answer: B

Its B, makes sense
upvoted 1 times

 **Captain1212** 1 month ago

B is the correct answer as question demands global scalability for it cloud spanner , for PostgreSQL to BIG query
upvoted 1 times

 **scanner2** 1 month ago

Selected Answer: B

For event streaming -> Cloud Pub/Sub is a good choice.
<https://cloud.google.com/pubsub/docs/overview>

Database for analytics and reporting -> BigQuery is a good choice.
<https://cloud.google.com/bigquery/docs/introduction>

For MySQL, we have two options Cloud SQL or Cloud Spanner. Cloud SQL has MySQL flavored database, while Cloud Spanner provides Google Standard SQL (aka GoogleSQL) which supports standard SQL queries. In addition, the question is asking about "global scalability and require minimal operational and infrastructure management", where Cloud Spanner wins a score.

Hence, B is the correct answer.

<https://cloud.google.com/spanner>

<https://cloud.google.com/spanner/docs/create-query-database-console#create-instance>

<https://cloud.google.com/spanner/docs/create-query-database-console#create-database>

upvoted 1 times

 **3arle** 2 months ago

Selected Answer: B

agree B
upvoted 1 times

 **Speridian** 2 months, 1 week ago

B should be the answer. Cloud Spanner - Global support.

upvoted 1 times

 **[Removed]** 2 months, 1 week ago

Selected Answer: B

B should be right
upvoted 3 times

 **ankitb4u** 2 months, 1 week ago

B should be the answer as cloud spanner provides scalability

upvoted 4 times

During a recent audit of your existing Google Cloud resources, you discovered several users with email addresses outside of your Google Workspace domain. You want to ensure that your resources are only shared with users whose email addresses match your domain. You need to remove any mismatched users, and you want to avoid having to audit your resources to identify mismatched users. What should you do?

- A. Create a Cloud Scheduler task to regularly scan your projects and delete mismatched users.
- B. Create a Cloud Scheduler task to regularly scan your resources and delete mismatched users.
- C. Set an organizational policy constraint to limit identities by domain to automatically remove mismatched users.
- D. Set an organizational policy constraint to limit identities by domain, and then retroactively remove the existing mismatched users

Correct Answer: B

Community vote distribution

D (100%)

 **joao_01** 3 weeks, 5 days ago

Selected Answer: D

Its D. "The domain restriction constraint is not retroactive. Once a domain restriction is set, this limitation will apply to IAM policy changes made from that point forward, and not to any previous changes.". Link: <https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains>

upvoted 1 times

 **scanner2** 1 month ago

Selected Answer: D

D seems to be most appropriate. You can use organization policy constraint to limit the identities by domain. Once the organization policy is set, you can remove the leftover users that mismatched the conditions.

<https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains>

upvoted 1 times

 **Cherrycardo** 2 months ago

Selected Answer: D

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints>

This list constraint defines the set of domains that email addresses added to Essential Contacts can have.

By default, email addresses with any domain can be added to Essential Contacts.

The allowed/denied list must specify one or more domains of the form @example.com. If this constraint is active and configured with allowed values, only email addresses with a suffix matching one of the entries from the list of allowed domains can be added in Essential Contacts.

This constraint has no effect on updating or removing existing contacts.

`constraints/essentialcontacts.allowedContactDomains`

upvoted 1 times

 **3arle** 2 months ago

Selected Answer: D

In order to define an organization policy, you choose a constraint, which is a particular type of restriction

upvoted 1 times

 **juliorevk** 2 months ago

Selected Answer: D

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints> - Domain restricted sharing

If this constraint is active, only principals that belong to the allowed customer IDs can be added to IAM policies. It doesn't specifically say, but I think it doesn't get rid of existing principals.

upvoted 2 times

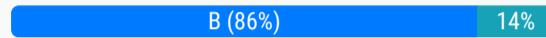
 **Speridian** 2 months, 1 week ago

Should be D. Organization policy does not remove users automatically.

upvoted 1 times

Your application is running on Google Cloud in a managed instance group (MIG). You see errors in Cloud Logging for one VM that one of the processes is not responsive. You want to replace this VM in the MIG quickly. What should you do?

- A. Use the gcloud compute instances update command with a REFRESH action for the VM.
- B. Use the gcloud compute instance-groups managed recreate-instances command to recreate the VM.
- C. Select the MIG from the Compute Engine console and, in the menu, select Replace VMs.
- D. Update and apply the instance template of the MIG.

Correct Answer: D*Community vote distribution* B (86%) 14%

✉️  **Ben_oso** 2 weeks, 1 day ago

Selected Answer: C

I think its C, because the question talk of "Replace the VM".
In B you recreate the same VM, so i think the C is more sense.
upvoted 1 times

✉️  **Captain1212** 1 month ago

Selected Answer: B

B is the correct answer , as the question demands in MIG managed instance group. B
upvoted 1 times

✉️  **scanner2** 1 month ago

Selected Answer: B

You can recreate specified instance(s) in a managed instance group.
<https://cloud.google.com/sdk/gcloud/reference/compute/instance-groups/managed/recreate-instances>
upvoted 1 times

✉️  **qannik** 2 months ago

Selected Answer: B

<https://cloud.google.com/sdk/gcloud/reference/compute/instance-groups/managed/recreate-instances>
upvoted 2 times

✉️  **Speridian** 2 months, 1 week ago

should be B.
upvoted 1 times

✉️  **gpais** 2 months, 1 week ago

Selected Answer: B

Following the document below is B:<https://cloud.google.com/sdk/gcloud/reference/compute/instance-groups/managed/recreate-instances>
upvoted 2 times

You want to permanently delete a Pub/Sub topic managed by Config Connector in your Google Cloud project. What should you do?

- A. Use kubectl to create the label deleted-by-cnrm and to change its value to true for the topic resource.
- B. Use kubectl to delete the topic resource.
- C. Use gcloud CLI to delete the topic.
- D. Use gcloud CLI to update the topic label managed-by-cnrm to false.

Correct Answer: D*Community vote distribution*

✉️ **Captain1212** 1 month ago

Selected Answer: B

B is the right answer, if the resource is managed by the config connector you can just use the kubectl to use delete or update the topic resource

upvoted 1 times

✉️ **scanner2** 1 month ago

Selected Answer: B

If a resource is managed by the Config Connector, you can update/delete it through kubectl command.
https://cloud.google.com/config-connector/docs/how-to/getting-started#before_you_begin

upvoted 2 times

✉️ **3arle** 2 months ago

Selected Answer: B

created by kubectl should be removed by it

upvoted 2 times

✉️ **qannik** 2 months ago

Selected Answer: B

https://cloud.google.com/config-connector/docs/how-to/getting-started#deleting_a_resource

upvoted 2 times

✉️ **Speridian** 2 months, 1 week ago

should be C.

upvoted 1 times

✉️ **happydays** 2 months, 1 week ago

Selected Answer: B

THIS IS B

upvoted 1 times

✉️ **gpais** 2 months, 1 week ago

Selected Answer: C

<https://cloud.google.com/sdk/gcloud/reference/pubsub/topics/delete>

upvoted 1 times

✉️ **qannik** 2 months ago

Read the question. It's a Pub/Sub topic managed by Config Connector.

https://cloud.google.com/config-connector/docs/how-to/getting-started#deleting_a_resource

upvoted 1 times

Your company is using Google Workspace to manage employee accounts. Anticipated growth will increase the number of personnel from 100 employees to 1,000 employees within 2 years. Most employees will need access to your company's Google Cloud account. The systems and processes will need to support 10x growth without performance degradation, unnecessary complexity, or security issues. What should you do?

- A. Migrate the users to Active Directory. Connect the Human Resources system to Active Directory. Turn on Google Cloud Directory Sync (GCDS) for Cloud Identity. Turn on Identity Federation from Cloud Identity to Active Directory.
- B. Organize the users in Cloud Identity into groups. Enforce multi-factor authentication in Cloud Identity.
- C. Turn on identity federation between Cloud Identity and Google Workspace. Enforce multi-factor authentication for domain wide delegation.
- D. Use a third-party identity provider service through federation. Synchronize the users from Google Workplace to the third-party provider in real time.

Correct Answer: C

Community vote distribution

C (50%)	B (30%)	A (20%)
---------	---------	---------

 **Ben_oso** 2 weeks, 1 day ago

Selected Answer: A

I think its A

upvoted 2 times

 **Gregwaw** 4 weeks, 1 day ago

Selected Answer: B

There is nothing like "identity federation between Cloud Identity and Google Workspace". You can only add Cloud Identity Free to your existing account.

You can manage users using Cloud Identity without adding all of them to Google Workspace.

Groups will help to manage the increased number of users.

Therefore B answer is correct.

upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: C

c is right

upvoted 1 times

 **NoCrapEva** 1 month ago

Selected Answer: C

(a) If you already administer a Google Workspace account and want to enable more users to use Google Cloud, you might not want to assign all users a Google Workspace license. In this case, add Cloud Identity Free to your existing account. You can then onboard more users without additional charge and decide which users should have access to Google Workspace by assigning them a Google Workspace license..

(b) To let your users collaborate by using Google Workspace, and to minimize administrative overhead, it's best to manage all users through a single Cloud Identity or Google Workspace account and provide a single user account to each individual. This approach helps ensure that settings such as password policies, single sign-on, and two-step verification are consistently applied to all users.

ref: <https://cloud.google.com/architecture/identity/best-practices-for-planning>

upvoted 1 times

 **Shivamwalia** 1 month, 1 week ago

Why not A?

upvoted 1 times

 **rahulrauki** 1 week, 6 days ago

They need reduced complexity, while this works, we skip one step (active directory) by linking cloud identity with workspace

upvoted 1 times

 **ee_23** 1 month, 3 weeks ago

Selected Answer: B

<https://cloud.google.com/architecture/identity/best-practices-for-planning>

upvoted 2 times

 **gpais** 2 months, 1 week ago

Selected Answer: C

I think C is the best

upvoted 3 times

You want to host your video encoding software on Compute Engine. Your user base is growing rapidly, and users need to be able to encode their videos at any time without interruption or CPU limitations. You must ensure that your encoding solution is highly available, and you want to follow Google-recommended practices to automate operations. What should you do?

- A. Deploy your solution on multiple standalone Compute Engine instances, and increase the number of existing instances when CPU utilization on Cloud Monitoring reaches a certain threshold.
- B. Deploy your solution on multiple standalone Compute Engine instances, and replace existing instances with high-CPU instances when CPU utilization on Cloud Monitoring reaches a certain threshold.
- C. Deploy your solution to an instance group, and increase the number of available instances whenever you see high CPU utilization in Cloud Monitoring.
- D. Deploy your solution to an instance group, and set the autoscaling based on CPU utilization.

Correct Answer: A*Community vote distribution*

D (100%)

✉️ 🚩 Captain1212 1 month ago

D seems more correct, just set the autoscaling based on the CPU utilization
upvoted 1 times

✉️ 🚩 scanner2 1 month, 2 weeks ago**Selected Answer: D**

The answer is D.
You can create a managed instance group with autoscaling enabled based on CPU utilization. This way appropriate number of instances can be added or removed based on the CPU metrics.
<https://cloud.google.com/compute/docs/instance-groups/create-mig-with-basic-autoscaling>
<https://cloud.google.com/sdk/gcloud/reference/compute/instance-groups/managed/set-autoscaling>
upvoted 2 times

✉️ 🚩 Jerica_ 1 month, 3 weeks ago**Selected Answer: D**

D seems most appropriate
upvoted 1 times

✉️ 🚩 3arle 2 months ago**Selected Answer: D**

definitely D
upvoted 2 times

✉️ 🚩 gpais 2 months, 1 week ago**Selected Answer: D**

<https://cloud.google.com/sdk/gcloud/reference/compute/instance-groups/managed/set-autoscaling>
upvoted 3 times

Your managed instance group raised an alert stating that new instance creation has failed to create new instances. You need to solve the instance creation problem. What should you do?

- A. Create an instance template that contains valid syntax which will be used by the instance group. Delete any persistent disks with the same name as instance names.
- B. Create an instance template that contains valid syntax that will be used by the instance group. Verify that the instance name and persistent disk name values are not the same in the template.
- C. Verify that the instance template being used by the instance group contains valid syntax. Delete any persistent disks with the same name as instance names. Set the disks.autoDelete property to true in the instance template.
- D. Delete the current instance template and replace it with a new instance template. Verify that the instance name and persistent disk name values are not the same in the template. Set the disks.autoDelete property to true in the instance template.

Correct Answer: A

Community vote distribution

A (83%) C (17%)

✉  **scanner2** 1 month, 2 weeks ago

Selected Answer: A

C is eliminated because you cannot update or modify an existing instance template. It is immutable.
https://cloud.google.com/compute/docs/instance-templates#how_to_update_instance_templates

D is eliminated because you cannot delete an instance template if a managed instance group references it.
https://cloud.google.com/compute/docs/instance-templates/get-list-delete-instance-templates#delete_an_instance_template

B is eliminated because you cannot set different/custom name for persistent disk in an instance template.
<https://cloud.google.com/sdk/gcloud/reference/compute/instance-templates/create#--create-disk>

So we left with A. Therefore A is the correct answer. Although, it should set the disks.autoDelete property to true for completeness.
upvoted 1 times

✉  **ee_23** 1 month, 3 weeks ago

Selected Answer: A

<https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs>
upvoted 2 times

✉  **3arle** 2 months ago

Selected Answer: C

Existing instance group has already template so no requirement to create new one. Problem is with persistent disk so delete it and configure autodelete solves problem now and in future
upvoted 1 times

✉  **3arle** 1 month, 3 weeks ago

i was wrong, templates are immutable so you cannot update them, so option A is valid
upvoted 2 times

✉  **ee_23** 1 month, 3 weeks ago

https://cloud.google.com/compute/docs/instance-templates#how_to_update_instance_templates
Instance templates are designed to create instances with identical configurations. So you cannot update an existing instance template or change an instance template after you create it.

If you need to make changes to the configuration, create a new instance template.
upvoted 1 times

✉  **happydays** 2 months, 1 week ago

Selected Answer: A

THIS IS A
upvoted 2 times

You have created an application that is packaged into a Docker image. You want to deploy the Docker image as a workload on Google Kubernetes Engine. What should you do?

- A. Upload the image to Cloud Storage and create a Kubernetes Service referencing the image.
- B. Upload the image to Cloud Storage and create a Kubernetes Deployment referencing the image.
- C. Upload the image to Artifact Registry and create a Kubernetes Service referencing the image.
- D. Upload the image to Artifact Registry and create a Kubernetes Deployment referencing the image.

Correct Answer: D

Community vote distribution

D (100%)

✉️  **scanner2** 1 month, 2 weeks ago

Selected Answer: D

A and B are eliminated because Cloud Storage is not a preferred place for storing docker images.
<https://cloud.google.com/artifact-registry/docs/docker/store-docker-container-images>

C is eliminated because Kubernetes Service is responsible for networking and connectivity between pods and external entities.

Kubernetes Deployment is responsible for deploying and managing an application (running in pods) on your GKE cluster.
https://cloud.google.com/kubernetes-engine/docs/how-to/deploying-workloads-overview#stateless_applications

upvoted 2 times

✉️  **3arle** 2 months ago

Selected Answer: D

best choice

upvoted 1 times

✉️  **pritampanda1988** 2 months, 1 week ago

Selected Answer: D

Artifact Registry is a fully managed container registry that integrates seamlessly with Google Kubernetes Engine and other Google Cloud services. By uploading the Docker image to Artifact Registry, you can create a Kubernetes Deployment that references the image stored in Artifact Registry. This ensures that Kubernetes can pull the image from a trusted and managed source, while the Deployment manages the deployment and scaling of the application pods based on the image.

upvoted 2 times

You are using Looker Studio to visualize a table from your data warehouse that is built on top of BigQuery. Data is appended to the data warehouse during the day. At night, the daily summary is recalculated by overwriting the table. You just noticed that the charts in Looker Studio are broken, and you want to analyze the problem. What should you do?

- A. In Cloud Logging, create a filter for your Looker Studio report.
- B. Use the open source CLI tool, Snapshot Debugger, to find out why the data was not refreshed correctly.
- C. Review the Error Reporting page in the Google Cloud console to find any errors.
- D. Use the BigQuery interface to review the nightly job and look for any errors.

Correct Answer: A

Community vote distribution

D (100%)

 **elvskimutai** 2 weeks, 3 days ago

D is the answer
upvoted 1 times

 **joao_01** 3 weeks, 4 days ago

Selected Answer: D

This question appear before. I would go with D, althoug the question before had opinions between the Cloud Logging and Bigquery options. Despite this i think that D is most appropriate for this case.
upvoted 3 times

 **Mallu_Mounika** 1 month, 4 weeks ago

Selected Answer: D
D is right
upvoted 3 times

 **happydays** 2 months, 1 week ago

Selected Answer: D
D is correct
upvoted 3 times

You have a batch workload that runs every night and uses a large number of virtual machines (VMs). It is fault-tolerant and can tolerate some of the VMs being terminated. The current cost of VMs is too high. What should you do?

- A. Run a test using simulated maintenance events. If the test is successful, use Spot N2 Standard VMs when running future jobs.
- B. Run a test using simulated maintenance events. If the test is successful, use N2 Standard VMs when running future jobs.
- C. Run a test using a managed instance group. If the test is successful, use N2 Standard VMs in the managed instance group when running future jobs.
- D. Run a test using N1 standard VMs instead of N2. If the test is successful, use N1 Standard VMs when running future jobs.

Correct Answer: D

Community vote distribution

A (100%)

 **joao_01** 3 weeks, 4 days ago

Selected Answer: A

Its A, makes sense
upvoted 1 times

 **ExamsFR** 3 weeks, 5 days ago

Selected Answer: A

A is the correct answer
upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: A

A is the correct answer as Spot VMs are highly affordable
upvoted 1 times

 **scanner2** 1 month, 2 weeks ago

Selected Answer: A

Here, the keywords are batch workloads, large number of VMs, tolerate some VMs being terminated, and addressing cost of VMs. Spot VMs have significant discounts than on-demand VMs, therefore A is the correct answer.
<https://cloud.google.com/compute/docs/instances/spot>
upvoted 1 times

 **3arle** 2 months ago

Selected Answer: A

Spot VMs are highly affordable compute instances suitable for batch jobs and fault-tolerant workloads. Spot VMs offer the same machine types, options, and performance as regular compute instances. If your applications are fault tolerant and can withstand possible instance preemptions, then Spot instances can reduce your Compute Engine costs by up to 91%!
upvoted 3 times

 **gpais** 2 months, 1 week ago

Selected Answer: A

A definitely
upvoted 3 times

 **Alberto06** 2 months, 1 week ago

Please , put the reference, thanks.
upvoted 1 times

You created several resources in multiple Google Cloud projects. All projects are linked to different billing accounts. To better estimate future charges, you want to have a single visual representation of all costs incurred. You want to include new cost data as soon as possible. What should you do?

- A. Fill all resources in the Pricing Calculator to get an estimate of the monthly cost.
- B. Use the Reports view in the Cloud Billing Console to view the desired cost information.
- C. Visit the Cost Table page to get a CSV export and visualize it using Looker Studio.
- D. Configure Billing Data Export to BigQuery and visualize the data in Looker Studio.

Correct Answer: D

Community vote distribution

D (100%)

 **joao_01** 3 weeks, 4 days ago

Selected Answer: D

Its D, most appropriate
upvoted 1 times

 **ExamsFR** 3 weeks, 5 days ago

Selected Answer: D

D is the correct answer
"single visual representation of all costs incurred"
upvoted 1 times

 **scanner2** 1 month, 2 weeks ago

Selected Answer: D

A and C are straightaway eliminated.
Cloud Billing Reports displays a chart that plots usage costs for all projects linked to a Cloud Billing account, but projects should be linked to a single Cloud Billing account. So, C is also eliminated.
<https://cloud.google.com/billing/docs/how-to/reports>

You can combine Cloud Billing data export to BigQuery with Looker Studio to stay up to date on your Google Cloud costs.
<https://cloud.google.com/billing/docs/how-to/visualize-data>

upvoted 2 times

 **gpais** 2 months ago

Selected Answer: D

We want to aggregate the costs for multiple billing accounts
upvoted 1 times

Your company has a large quantity of unstructured data in different file formats. You want to perform ETL transformations on the data. You need to make the data accessible on Google Cloud so it can be processed by a Dataflow job. What should you do?

- A. Upload the data to BigQuery using the bq command line tool.
- B. Upload the data to Cloud Storage using the gcloud storage command.
- C. Upload the data into Cloud SQL using the import function in the Google Cloud console.
- D. Upload the data into Cloud Spanner using the import function in the Google Cloud console.

Correct Answer: B

Community vote distribution

B (100%)

✉  **joao_01** 3 weeks, 4 days ago

Selected Answer: B

Its B, non structure data
upvoted 1 times

✉  **ExamsFR** 3 weeks, 5 days ago

Selected Answer: B

"unstructured data in different file formats"
upvoted 1 times

✉  **Captain1212** 1 month ago

Selected Answer: B

For unstrucutred data cloud striooage is the right answer
upvoted 1 times

✉  **scanner2** 1 month, 2 weeks ago

Selected Answer: B

The question is asking about unstructured data. So, options A, C & E are eliminated.
B is the correct answer.
upvoted 1 times

✉  **Cherrycardo** 2 months ago

Selected Answer: B

Unstructured is the keyword in this questions. All possible answers are structured, but Cloud Storage.
upvoted 1 times

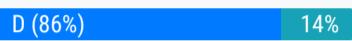
✉  **3arle** 2 months ago

Selected Answer: B

only B is for unstructured data
upvoted 1 times

You have deployed an application on a single Compute Engine instance. The application writes logs to disk. Users start reporting errors with the application. You want to diagnose the problem. What should you do?

- A. Navigate to Cloud Logging and view the application logs.
- B. Configure a health check on the instance and set a “consecutive successes” Healthy threshold value of 1.
- C. Connect to the instance’s serial console and read the application logs.
- D. Install and configure the Ops agent and view the logs from Cloud Logging.

Correct Answer: D*Community vote distribution*

✉ **joao_01** 3 weeks, 4 days ago

Selected Answer: D

We need to install the agent. Its D
upvoted 2 times

✉ **scanner2** 1 month, 2 weeks ago

The Ops Agent is the primary agent for collecting telemetry from your Compute Engine instances.
Answer is D

<https://cloud.google.com/logging/docs/agent/ops-agent>
<https://cloud.google.com/logging/docs/agent/ops-agent/configuration>
upvoted 1 times

✉ **3arle** 2 months ago

Selected Answer: D

<https://cloud.google.com/logging/docs/logging-gce-quickstart>
upvoted 2 times

✉ **qannik** 2 months ago

Selected Answer: D

I would go with D.
By default there is no logs agent installed on a compute instance.
So first you will have to install the Ops Agent and after a few minutes the logs will be visible in Cloud logging
upvoted 2 times

✉ **happydays** 2 months, 1 week ago

Selected Answer: A

Chat Gpt says it's A
upvoted 1 times

You recently received a new Google Cloud project with an attached billing account where you will work. You need to create instances, set firewalls, and store data in Cloud Storage. You want to follow Google-recommended practices. What should you do?

- A. Use the gcloud CLI services enable cloudresourcemanager.googleapis.com command to enable all resources.
- B. Use the gcloud services enable compute.googleapis.com command to enable Compute Engine and the gcloud services enable storage-api.googleapis.com command to enable the Cloud Storage APIs.
- C. Open the Google Cloud console and enable all Google Cloud APIs from the API dashboard.
- D. Open the Google Cloud console and run gcloud init --project in a Cloud Shell.

Correct Answer: B

Community vote distribution

B (100%)

 **joao_01** 3 weeks, 4 days ago

Selected Answer: B

Enable correspondent APIs. Its B
upvoted 2 times

 **Captain1212** 1 month ago

Selected Answer: B

B is the correct answer
upvoted 1 times

 **3arle** 1 month, 3 weeks ago

Selected Answer: B

At first you need to enable API
upvoted 2 times

Your application development team has created Docker images for an application that will be deployed on Google Cloud. Your team does not want to manage the infrastructure associated with this application. You need to ensure that the application can scale automatically as it gains popularity. What should you do?

- A. Create an instance template with the container image, and deploy a Managed Instance Group with Autoscaling.
- B. Upload Docker images to Artifact Registry, and deploy the application on Google Kubernetes Engine using Standard mode.
- C. Upload Docker images to the Cloud Storage, and deploy the application on Google Kubernetes Engine using Standard mode.
- D. Upload Docker images to Artifact Registry, and deploy the application on Cloud Run.

Correct Answer: B*Community vote distribution*

D (100%)

joao_01 3 weeks, 4 days ago**Selected Answer: D**

Its D. GKE standard the nodes are managed by user. So D is correct.

upvoted 1 times

nnecode 3 weeks, 4 days ago**Selected Answer: D**

D. Upload Docker images to Artifact Registry and deploy the application on Cloud Run.

upvoted 1 times

Captain1212 1 month ago**Selected Answer: D**

D is the correct answer , as question says your team dont want to manage the infrastrucutre associated withthe application this is offered by the cloudrun

upvoted 1 times

AkshayJangwal 1 month, 1 week agoKey Hint : Your team does not want to manage the infrastructure associated with this application
This is offered by Cloud Run, hence option D.

upvoted 1 times

scanner2 1 month, 2 weeks ago**Selected Answer: D**

Answer is D.

Cloud Run is container-as-a-service offering from Google Cloud. You can deploy the containerized application directly on top of Google's infrastructure, and only when a request comes.

<https://cloud.google.com/run/docs/overview/what-is-cloud-run>

upvoted 1 times

scanner2 1 month, 2 weeks ago

You don't need to worry about underlying infrastructure.

A, C and E requires resources, virtual instance, GKE cluster provisioning and maintaining. so these are eliminated.

upvoted 1 times

3arle 2 months ago**Selected Answer: D**

GKE Standard mode: You manage the underlying infrastructure, including configuring the individual nodes.

Instance group - you manage the infrastructure as well

so after elimination A,B,C stays D

upvoted 1 times

happydays 2 months ago**Selected Answer: D**

IT'S D

upvoted 2 times

gpais 2 months, 1 week ago**Selected Answer: D**

Option D

upvoted 3 times

You are migrating a business critical application from your local data center into Google Cloud. As part of your high-availability strategy, you want to ensure that any data used by the application will be immediately available if a zonal failure occurs. What should you do?

- A. Store the application data on a zonal persistent disk. Create a snapshot schedule for the disk. If an outage occurs, create a new disk from the most recent snapshot and attach it to a new VM in another zone.
- B. Store the application data on a zonal persistent disk. If an outage occurs, create an instance in another zone with this disk attached.
- C. Store the application data on a regional persistent disk. Create a snapshot schedule for the disk. If an outage occurs, create a new disk from the most recent snapshot and attach it to a new VM in another zone.
- D. Store the application data on a regional persistent disk. If an outage occurs, create an instance in another zone with this disk attached.

Correct Answer: A*Community vote distribution*

D (83%)	C (17%)
---------	---------

✉ **joao_01** 3 weeks, 4 days ago

Selected Answer: D

I thought first it was C. Then i saw Google documentation and for sure the answer is D.

upvoted 1 times

✉ **Captain1212** 1 month ago

Selected Answer: D

D is the correct answer

upvoted 1 times

✉ **DannSecurity** 1 month ago

C,

If you are designing robust systems or high availability services on Compute Engine, use regional Persistent Disk combined with other best practices such as backing up your data using snapshots.

upvoted 2 times

✉ **Speridian** 1 month, 3 weeks ago

It should be D.

upvoted 1 times

✉ **gfalconia** 1 month, 3 weeks ago

Selected Answer: D

D, C doesn't make sense because article explicitly says that: 'During the failover, the regional persistent disk that is synchronously replicated to the secondary zone is force attached to the standby VM by the application control plane, and all traffic is directed to that VM based on health check signals.'

<https://cloud.google.com/compute/docs/disks/high-availability-regional-persistent-disk#failover>

upvoted 2 times

✉ **_claudio_** 1 month, 3 weeks ago

Selected Answer: D

The benefit of regional persistent disks is that in the event of a zonal outage, where your virtual machine (VM) instance might become unavailable, you can usually force attach a regional persistent disk to a VM instance in a secondary zone in the same region.

upvoted 3 times

✉ **3arle** 2 months ago

Selected Answer: D

<https://cloud.google.com/compute/docs/disks/high-availability-regional-persistent-disk>

upvoted 2 times

✉ **Vignesshvar** 2 months ago

CHATGPT says C

upvoted 2 times

✉ **hoai_nam_1512** 2 months ago

correct: C

upvoted 2 times

✉ **Speridian** 2 months, 1 week ago

It should be C.

upvoted 2 times

✉ **Husni_adam** 2 months, 1 week ago

I think c, for high availability use regional instead of zonal, and you can create snapshot to protect against data loss
https://cloud.google.com/compute/docs/disks#reliability_2:~:text=Similar%20to%20zonal%20Persistent%20Disk%2C%20you%20can%20create%20snapshots%20of%20Persistent%20Disk%20to%20protect%20against%20data%20loss%20due%20to%20user%20error.%20Snapshots%20are%20incremental%2C%20and%20take%20only%20minutes%20to%20create%20even%20if%20you%20snapshot%20disks%20that%20are%20attached%20to%20running%20instances.

upvoted 2 times

 **happydays** 2 months, 1 week ago

Selected Answer: D

Chat Gpt sayd "D"

upvoted 1 times

The DevOps group in your organization needs full control of Compute Engine resources in your development project. However, they should not have permission to create or update any other resources in the project. You want to follow Google's recommendations for setting permissions for the DevOps group. What should you do?

- A. Grant the basic role roles/viewer and the predefined role roles/compute.admin to the DevOps group.
- B. Create an IAM policy and grant all compute.instanceAdmin.* permissions to the policy. Attach the policy to the DevOps group.
- C. Create a custom role at the folder level and grant all compute.instanceAdmin.* permissions to the role. Grant the custom role to the DevOps group.
- D. Grant the basic role roles/editor to the DevOps group.

Correct Answer: B*Community vote distribution*

A (53%)	C (27%)	B (20%)
---------	---------	---------

✉ **joao_01** 3 weeks, 3 days ago

Selected Answer: B

For me its B, until anyone says the contrary and why. It give ONLY the permissions required. No more or less.
upvoted 1 times

✉ **DannSecurity** 4 weeks ago

Answer A
Compute Admin

(roles/compute.admin)
Full control of all Compute Engine resources.
upvoted 2 times

✉ **Captain1212** 1 month ago

Selected Answer: A

A is th correct answer as it provied all the required access
upvoted 2 times

✉ **NoCrapEva** 1 month ago

Selected Answer: A

Compute Admin (roles/compute.admin) = Full control of all Compute Engine resources.
The only permission to have full control of Computer Engine Resources (as required in question)
ref: <https://cloud.google.com/iam/docs/understanding-roles#compute.admin>

Compute.instanceAdmin does NOT allow FULL control of Compute Engine, only
Permissions to create, modify, and delete virtual machine instances. This includes permissions to create, modify, and delete disks, and
also to configure Shielded VM settings.
upvoted 2 times

✉ **Az900Exam2021** 1 month ago

Selected Answer: C

C meets the requirement of permission with least privilege
upvoted 1 times

✉ **tejasshellar2022** 1 month, 1 week ago

Selected Answer: B

I Think answer is B
A- roles/compute.admin can give required access , but role roles/viewer is additional access given which is not mentioned in problem statement that group should have viewer access on all resources.

C - its creating custom role at folder level and folder roles can be applied to multiple project, problem statement stating that it should be on development project only.

D- out of scope as its editor role
upvoted 2 times

✉ **joao_01** 3 weeks, 4 days ago

I think its B as well. Gives the group ONLY the permissions the problem said to give it to. About your explanation, its accurate cause it was what i thought as well. I dont get why people are not consider this one.
upvoted 1 times

✉ **Ahmed_Y** 1 month, 2 weeks ago

Selected Answer: A

I think it is A for two reasons:

1- "they should not have permission to create or update any other resources in the project" so they should have the view access and A is providing that.

2- the answer C is on folder level while we want to give that only to development project.

upvoted 1 times

✉ **tczorro** 1 month, 3 weeks ago

Selected Answer: C

C is the answer as compute.instanceAdmin contains all the required permission

upvoted 1 times

✉ **gpais** 2 months ago

Selected Answer: A

seems the logic choice based on the requirement of giving full control of all compute resources

upvoted 1 times

✉ **gpais** 1 month, 3 weeks ago

I change my mind I think C based on the link: <https://cloud.google.com/iam/docs/using-iam-securely>
Least privilege



Basic roles include thousands of permissions across all Google Cloud services. In production environments, do not grant basic roles unless there is no alternative. Instead, grant the most limited predefined roles or custom roles that meet your needs.

If you need to replace a basic role, you can use role recommendations to determine which roles to grant instead. You can also use the Policy Simulator to ensure that changing the role won't affect the principal's access.

It might be appropriate to grant basic roles in the following cases:

When the Google Cloud service does not provide a predefined role. See the predefined roles table for a list of all available predefined roles.

When you want to grant broader permissions for a project. This often happens when you're granting permissions in development or test environments.

When you work in a small team where the team members don't need granular permissions.

upvoted 1 times

✉ **joao_01** 3 weeks, 4 days ago

With C you will give the permissions to all projects inside a folder. So, indeed you will give to the project in the problem. However, you'll also give to other projects that are under that folder. So, it cannot be C according to the least privilege rule.

upvoted 1 times

✉ **3arle** 2 months ago

Selected Answer: A

only A gives 'Full control of all Compute Engine resources.' (not instances)
<https://cloud.google.com/iam/docs/understanding-roles>

upvoted 2 times

✉ **happydays** 2 months, 1 week ago

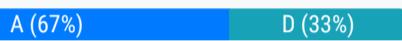
Selected Answer: C

IT'S C

upvoted 2 times

Your team is running an on-premises ecommerce application. The application contains a complex set of microservices written in Python, and each microservice is running on Docker containers. Configurations are injected by using environment variables. You need to deploy your current application to a serverless Google Cloud cloud solution. What should you do?

- A. Use your existing CI/CD pipeline. Use the generated Docker images and deploy them to Cloud Run. Update the configurations and the required endpoints.
- B. Use your existing continuous integration and delivery (CI/CD) pipeline. Use the generated Docker images and deploy them to Cloud Function. Use the same configuration as on-premises.
- C. Use the existing codebase and deploy each service as a separate Cloud Function. Update the configurations and the required endpoints.
- D. Use your existing codebase and deploy each service as a separate Cloud Run. Use the same configurations as on-premises.

Correct Answer: A*Community vote distribution*

✉ **joao_01** 3 weeks, 3 days ago

Selected Answer: A

I think its A

upvoted 1 times

✉ **nnecode** 3 weeks, 4 days ago

Selected Answer: D

I vote D

upvoted 1 times

✉ **Captain1212** 1 month ago

Selected Answer: A

A is the correct answer , use your existing CI/cd pipeline and update the configuratons and the required endpoints

upvoted 2 times

✉ **Cherrycardo** 2 months ago

Selected Answer: D

<https://cloud.google.com/run/docs/configuring/services/environment-variables>

"The environment variables defined in the container runtime contract are reserved and cannot be set. In particular, the PORT environment variable is injected inside your container by Cloud Run. You should not set it yourself."

Hence, by "using the same configurations of on-premise" you are just using the environment variables already present on the container.

upvoted 1 times

✉ **Cherrycardo** 2 months ago

I was wrong. The right answer is A. The current approach to load Docker images into Artifact Registry (formerly Container Registry), is by using CI/CD Pipelines.

upvoted 2 times

✉ **3arle** 2 months ago

Selected Answer: A

app was written for docker image, it likely should be rewritten for cloud run

upvoted 1 times

✉ **qannik** 2 months ago

Selected Answer: A

I think it's A.

It can't be D because you can not use the same configuration as on-premise.

upvoted 1 times

✉ **qannik** 2 months ago

Selected Answer: D

I vote for D

upvoted 1 times

✉ **qannik** 2 months ago

I vote for D

upvoted 1 times

✉ **happydays** 2 months, 1 week ago

Selected Answer: A

Question #237

Topic 1

You are running multiple microservices in a Kubernetes Engine cluster. One microservice is rendering images. The microservice responsible for the image rendering requires a large amount of CPU time compared to the memory it requires. The other microservices are workloads that are optimized for n2-standard machine types. You need to optimize your cluster so that all workloads are using resources as efficiently as possible. What should you do?

- A. Assign the pods of the image rendering microservice a higher pod priority than the other microservices.
- B. Create a node pool with compute-optimized machine type nodes for the image rendering microservice. Use the node pool with general-purpose machine type nodes for the other microservices.
- C. Use the node pool with general-purpose machine type nodes for the image rendering microservice. Create a node pool with compute-optimized machine type nodes for the other microservices.
- D. Configure the required amount of CPU and memory in the resource requests specification of the image rendering microservice deployment. Keep the resource requests for the other microservices at the default.

Correct Answer: B

Community vote distribution

B (100%)

✉  **joao_01** 3 weeks, 3 days ago

Selected Answer: B

Its B, question appear before

upvoted 2 times

✉  **3arle** 2 months ago

Selected Answer: B

repeated question

upvoted 1 times

✉  **qannik** 2 months ago

Selected Answer: B

B has logic

upvoted 1 times

You are working in a team that has developed a new application that needs to be deployed on Kubernetes. The production application is business critical and should be optimized for reliability. You need to provision a Kubernetes cluster and want to follow Google-recommended practices.

What should you do?

- A. Create a GKE Autopilot cluster. Enroll the cluster in the rapid release channel.
- B. Create a GKE Autopilot cluster. Enroll the cluster in the stable release channel.
- C. Create a zonal GKE standard cluster. Enroll the cluster in the stable release channel.
- D. Create a regional GKE standard cluster. Enroll the cluster in the rapid release channel.

Correct Answer: B

Community vote distribution

B (100%)

 **joao_01** 3 weeks, 3 days ago

Selected Answer: B

I selected B
upvoted 1 times

 **Captain1212** 1 month ago

Selected Answer: B

Autopilot cluster is more reliable and gives more time to fix
upvoted 1 times

 **3arle** 2 months ago

Selected Answer: B

Autopilot is more reliable and stable release gives more time to fix issues in new version of GKE
upvoted 1 times

 **qannik** 2 months ago

Selected Answer: B

Autopilot cluster is recommended by Google
upvoted 1 times

You are responsible for a web application on Compute Engine. You want your support team to be notified automatically if users experience high latency for at least 5 minutes. You need a Google-recommended solution with no development cost. What should you do?

- A. Export Cloud Monitoring metrics to BigQuery and use a Looker Studio dashboard to monitor your web application's latency.
- B. Create an alert policy to send a notification when the HTTP response latency exceeds the specified threshold.
- C. Implement an App Engine service which invokes the Cloud Monitoring API and sends a notification in case of anomalies.
- D. Use the Cloud Monitoring dashboard to observe latency and take the necessary actions when the response latency exceeds the specified threshold.

Correct Answer: C

Community vote distribution

B (100%)

✉  **rahulrauki** 1 week, 6 days ago

Selected Answer: B

"No development cost" : BigQuery and GAE out
"Automatically" : Option D out ("take necessary actions")
Left with option B
upvoted 1 times

✉  **ExamsFR** 3 weeks, 5 days ago

Selected Answer: B

"You need a Google-recommended solution with no development cost"
upvoted 1 times

✉  **3arle** 2 months ago

Selected Answer: B

<https://cloud.google.com/monitoring/alerts#alerting-example>
upvoted 3 times

✉  **gpais** 2 months, 1 week ago

Selected Answer: B

B seems to be the best answer
upvoted 1 times

You have an on-premises data analytics set of binaries that processes data files in memory for about 45 minutes every midnight. The sizes of those data files range from 1 gigabyte to 16 gigabytes. You want to migrate this application to Google Cloud with minimal effort and cost. What should you do?

- A. Create a container for the set of binaries. Use Cloud Scheduler to start a Cloud Run job for the container.
- B. Create a container for the set of binaries. Deploy the container to Google Kubernetes Engine (GKE) and use the Kubernetes scheduler to start the application.
- C. Upload the code to Cloud Functions. Use Cloud Scheduler to start the application.
- D. Lift and shift to a VM on Compute Engine. Use an instance schedule to start and stop the instance.

Correct Answer: A*Community vote distribution*

D (67%)	A (33%)
---------	---------

✉  **taylz876** 3 days, 5 hours ago

Selected Answer: D

Here's why option D is the most appropriate:

->Compute Engine: Compute Engine provides virtual machines (VMs) that closely resemble traditional on-premises servers. It allows you to migrate your existing application as-is to the Google Cloud platform.

->Instance Scheduling: You can schedule the VM instance to start and stop at specific times, such as midnight, to align with your existing processing schedule. This ensures that the application runs at the required time, similar to the on-premises setup.

->Minimal Effort and Cost: The "lift and shift" approach minimizes the need for code modifications or containerization, reducing migration complexity. It also allows you to use the same binaries and configurations as your on-premises setup, saving development effort. You only pay for the VM's compute resources when it's running, making it cost-effective.

upvoted 1 times

✉  **SinghAnc** 1 week ago

Selected Answer: D

D is the correct answer.

upvoted 1 times

✉  **joao_01** 3 weeks, 3 days ago

Selected Answer: D

This one is a tough one. Ill consider this in my answers:

Cost --> Both are more the same (the process will run at the same frequency)

Effort --> Create the image in A takes more effort then to option D.

With this in mind ill choose D. (before i was choosing A).

upvoted 1 times

✉  **Captain1212** 1 month ago

Selected Answer: D

D is the correct answer as it requires the minimal effort

upvoted 1 times

✉  **scanner2** 1 month ago

Selected Answer: A

A is correct answer.

upvoted 1 times

✉  **tczorro** 1 month, 3 weeks ago

Selected Answer: D

minimal effort is to lift and shift to VM

minimal cost is to use schedule to start and stop the instance

upvoted 2 times

✉  **rsvd** 1 month, 4 weeks ago

Selected Answer: A

Things to consider:

1. can cloud functions run binary file?

2. Data file size varies from 1 to 16 GB. so how to determine the VM spec?

upvoted 1 times

✉  **3arle** 2 months ago

i've made mistake, should be D
upvoted 2 times

 **3arle** 2 months ago

Selected Answer: A
lift and shift is less effort and cost, you don't have to pay for refactoring or creating image.
upvoted 1 times

 **qannik** 2 months ago

Selected Answer: D
I would choose D because I want to migrate this application to Google Cloud with minimal effort and cost.
Cloud Run requires to create a container image and this means some kind of development and testing.
upvoted 2 times

 **gpa1s** 2 months, 1 week ago

Selected Answer: A
A seems to be the less disruptive solution with lower costs
upvoted 1 times

You used the gcloud container clusters command to create two Google Cloud Kubernetes (GKE) clusters: prod-cluster and dev-cluster.

- prod-cluster is a standard cluster.
- dev-cluster is an auto-pilot cluster.

When you run the kubectl get nodes command, you only see the nodes from prod-cluster. Which commands should you run to check the node status for dev-cluster?

- A. gcloud container clusters get-credentials dev-cluster
kubectl get nodes
- B. gcloud container clusters update -generate-password dev-cluster kubectl get nodes
- C. kubectl config set-context dev-cluster
kubectl cluster-info
- D. kubectl config set-credentials dev-cluster
kubectl cluster-info

Correct Answer: D

Community vote distribution

A (100%)

✉  **joao_01** 3 weeks, 3 days ago

Selected Answer: A

Its the A

upvoted 1 times

✉  **Captain1212** 1 month ago

Selected Answer: A

A is the correct answer as it , updated the config file

upvoted 1 times

✉  **pritampanda1988** 1 month, 2 weeks ago

Selected Answer: A

gcloud container clusters get-credentials updates a kubeconfig file with appropriate credentials and endpoint information to point kubectl at a specific cluster in Google Kubernetes Engine

upvoted 1 times

✉  **3arle** 2 months ago

Selected Answer: A

gcloud container clusters get-credentials updates a kubeconfig file with appropriate credentials and endpoint information to point kubectl at a specific cluster in Google Kubernetes Engine

upvoted 2 times

✉  **qannik** 2 months ago

Selected Answer: A

The gcloud container clusters get-credentials command sets the Kubernetes context to the specified cluster (in this case, dev-cluster). This ensures that the subsequent kubectl commands will be executed against the dev-cluster.

After setting the context, the kubectl get nodes command is used to retrieve the node status for the dev-cluster, showing the list of nodes in the cluster.

upvoted 2 times

✉  **Speridian** 2 months, 1 week ago

It should be A

upvoted 2 times

You recently discovered that your developers are using many service account keys during their development process. While you work on a long term improvement, you need to quickly implement a process to enforce short-lived service account credentials in your company. You have the following requirements:

- All service accounts that require a key should be created in a centralized project called pj-sa.
- Service account keys should only be valid for one day.

You need a Google-recommended solution that minimizes cost. What should you do?

- A. Implement a Cloud Run job to rotate all service account keys periodically in pj-sa. Enforce an org policy to deny service account key creation with an exception to pj-sa.
- B. Implement a Kubernetes CronJob to rotate all service account keys periodically. Disable attachment of service accounts to resources in all projects with an exception to pj-sa.
- C. Enforce an org policy constraint allowing the lifetime of service account keys to be 24 hours. Enforce an org policy constraint denying service account key creation with an exception on pj-sa.
- D. Enforce a DENY org policy constraint over the lifetime of service account keys for 24 hours. Disable attachment of service accounts to resources in all projects with an exception to pj-sa.

Correct Answer: D

Community vote distribution

C (100%)

 **joao_01** 3 weeks, 3 days ago

Selected Answer: C

Its C, makes sense
upvoted 1 times

 **Captain1212** 1 month ago

c is the coorect answer
upvoted 1 times

 **scanner2** 1 month ago

Selected Answer: C

C is correct.
https://cloud.google.com/resource-manager/docs/organization-policy/restricting-service-accounts#limit_key_expiry
https://cloud.google.com/resource-manager/docs/organization-policy/restricting-service-accounts#disable_service_account_key_creation
upvoted 1 times

 **3arle** 2 months ago

Selected Answer: C

it should be C
upvoted 3 times

 **qannik** 2 months ago

Selected Answer: C

You can use an org policy to enforce a 24-hour lifetime for service account keys.
You can use an org policy to deny service account key creation, with an exception for the pj-sa project.
This is a Google-recommended solution and it is relatively inexpensive.
upvoted 2 times

 **niedobry** 2 months, 1 week ago

Answer is C. Constraint: constraints/iam.serviceAccountKeyExpiryHours does not accept DENY values so D can not be correct.
upvoted 2 times

Your company is running a three-tier web application on virtual machines that use a MySQL database. You need to create an estimated total cost of cloud infrastructure to run this application on Google Cloud instances and Cloud SQL. What should you do?

- A. Create a Google spreadsheet with multiple Google Cloud resource combinations. On a separate sheet, import the current Google Cloud prices and use these prices for the calculations within formulas.
- B. Use the Google Cloud Pricing Calculator and select the Cloud Operations template to define your web application with as much detail as possible.
- C. Implement a similar architecture on Google Cloud, and run a reasonable load test on a smaller scale. Check the billing information, and calculate the estimated costs based on the real load your system usually handles.
- D. Use the Google Cloud Pricing Calculator to determine the cost of every Google Cloud resource you expect to use. Use similar size instances for the web server, and use your current on-premises machines as a comparison for Cloud SQL.

Correct Answer: D

Community vote distribution



✉ **joao_01** 3 weeks, 3 days ago

Selected Answer: D

Its D, try to simulate using what we have
upvoted 1 times

✉ **scanner2** 1 month ago

Selected Answer: D

D is correct.
upvoted 1 times

✉ **ptapia_el** 1 month, 2 weeks ago

es la D
upvoted 1 times

✉ **rsvd** 1 month, 4 weeks ago

Selected Answer: D

There is no such thing called "Cloud Operations template"
upvoted 1 times

✉ **3arle** 2 months ago

Selected Answer: D

it's D
upvoted 1 times

✉ **qannik** 2 months ago

Selected Answer: D

Google Cloud Pricing Calculator, is the recommended approach for creating an estimated total cost of cloud infrastructure. By selecting the relevant Google Cloud resources (such as instances for web servers and Cloud SQL for the database), and specifying similar sizes and configurations, you can obtain a more accurate estimation of the costs.
upvoted 2 times

✉ **happydays** 2 months, 1 week ago

Selected Answer: B

Use calculation and perations template // it's B
upvoted 1 times

✉ **joao_01** 2 weeks, 3 days ago

Dafuq, that doesn't make any sense
upvoted 1 times

You have a Bigtable instance that consists of three nodes that store personally identifiable information (PII) data. You need to log all read or write operations, including any metadata or configuration reads of this database table, in your company's Security Information and Event Management (SIEM) system. What should you do?

- A. • Navigate to Cloud Monitoring in the Google Cloud console, and create a custom monitoring job for the Bigtable instance to track all changes.
 - Create an alert by using webhook endpoints, with the SIEM endpoint as a receiver.
- B. • Navigate to the Audit Logs page in the Google Cloud console, and enable Admin Write logs for the Bigtable instance.
 - Create a Cloud Functions instance to export logs from Cloud Logging to your SIEM.
- C. • Navigate to the Audit Logs page in the Google Cloud console, and enable Data Read, Data Write and Admin Read logs for the Bigtable instance.
 - Create a Pub/Sub topic as a Cloud Logging sink destination, and add your SIEM as a subscriber to the topic.
- D. • Install the Ops Agent on the Bigtable instance during configuration.
 - Create a service account with read permissions for the Bigtable instance.
 - Create a custom Dataflow job with this service account to export logs to the company's SIEM system.

Correct Answer: C*Community vote distribution*

✉ **taylz876** 3 days, 5 hours ago

Selected Answer: C

Option C is the most appropriate choice for capturing audit and data access logs from a Bigtable instance and sending them to your SIEM system.

- 1) Enabling Data Read, Data Write, and Admin Read logs for the Bigtable instance ensures that you capture the relevant operations, including read and write operations, as well as administrative reads, in the audit logs.
- 2) Creating a Pub/Sub topic as a Cloud Logging sink destination allows you to export the logs from Cloud Logging to Pub/Sub. This is a common approach for sending logs to external systems, including SIEMs.
- 3) Adding your SIEM as a subscriber to the Pub/Sub topic ensures that the logs are forwarded to your SIEM system, allowing you to monitor and analyze them for security and compliance purposes.

NB: A Cloud Logging sink destination is a configuration that specifies where logs collected by Google Cloud's Cloud Logging service should be sent or exported. It allows you to control the destination of logs generated by various Google Cloud services, such as Compute Engine, Cloud Storage, BigQuery, and more.

upvoted 1 times

✉ **joao_01** 3 weeks, 3 days ago

Selected Answer: C

Its C!

upvoted 1 times

✉ **Captain1212** 1 month ago

Selected Answer: C

C is the correct answer, as it helps you to read and write

upvoted 1 times

✉ **scanner2** 1 month ago

Selected Answer: C

C is correct.

upvoted 1 times

✉ **ee_23** 1 month, 3 weeks ago

Selected Answer: C

<https://cloud.google.com/bigtable/docs/audit-logging#available-logs>

B: Admin write logs are already enabled by default

upvoted 1 times

✉ **3arle** 2 months ago

Selected Answer: C

Data Access audit logs—except for BigQuery—are disabled by default and you need to enable them

upvoted 2 times

✉ **qannik** 2 months ago

Selected Answer: B

Enabling Admin Write logs for the Bigtable instance in Cloud Logging will capture administrative write actions on the Bigtable instance. This includes any configuration changes and metadata reads related to the Bigtable instance.

Creating a Cloud Functions instance and configuring it to export logs from Cloud Logging to your SIEM allows you to take the captured logs and route them to your SIEM system in a format that your SIEM can understand. Cloud Functions can act as a serverless function to process and forward the logs to your SIEM using an appropriate method, such as sending them via an API or message queue.

upvoted 1 times

Question #245

Topic 1

You want to set up a Google Kubernetes Engine cluster. Verifiable node identity and integrity are required for the cluster, and nodes cannot be accessed from the internet. You want to reduce the operational cost of managing your cluster, and you want to follow Google-recommended practices. What should you do?

- A. Deploy a private autopilot cluster.
- B. Deploy a public autopilot cluster.
- C. Deploy a standard public cluster and enable shielded nodes.
- D. Deploy a standard private cluster and enable shielded nodes.

Correct Answer: A

Community vote distribution

A (78%) D (22%)

✉  **scanner2** 1 month ago

Selected Answer: A

In a private cluster, nodes only have internal IP addresses, which means that nodes and Pods are isolated from the internet by default.
<https://cloud.google.com/kubernetes-engine/docs/how-to/private-clusters>

Shielded GKE Nodes provide strong, verifiable node identity and integrity to increase the security of Google Kubernetes Engine (GKE) nodes.

Note: For GKE Autopilot clusters, the Shielded GKE Nodes feature is enabled by default and cannot be overridden.
<https://cloud.google.com/kubernetes-engine/docs/how-to/shielded-gke-nodes>

upvoted 1 times

✉  **rsvd** 1 month, 4 weeks ago

Selected Answer: A

Note: For GKE Autopilot clusters, the Shielded GKE Nodes feature is enabled by default and cannot be overridden.

upvoted 3 times

✉  **Cherrycardo** 2 months ago

Selected Answer: A

<https://cloud.google.com/kubernetes-engine/docs/how-to/shielded-gke-nodes>

"For GKE Autopilot clusters, the Shielded GKE Nodes feature is enabled by default and cannot be overridden"

upvoted 1 times

✉  **3arle** 2 months ago

Selected Answer: A

The Shielded GKE node feature is enabled by default for all Autopilot clusters and is impossible to disable manually.

<https://www.googlecloudcommunity.com/gc/Architecture-Framework-Community/Manage-GKE-Cluster-Security-with-Autopilot-Mode/ba-p/396435>

upvoted 2 times

✉  **qannik** 2 months ago

Selected Answer: D

<https://cloud.google.com/kubernetes-engine/docs/how-to/shielded-gke-nodes>

upvoted 1 times

✉  **gpais** 2 months, 1 week ago

Selected Answer: D

Shielded GKE Nodes provide strong, verifiable node identity and integrity to increase the security of GKE nodes and should be enabled on all GKE clusters.: <https://cloud.google.com/kubernetes-engine/docs/how-to/hardening-your-cluster>

upvoted 1 times

Your company wants to migrate their on-premises workloads to Google Cloud. The current on-premises workloads consist of:

- A Flask web application
- A backend API
- A scheduled long-running background job for ETL and reporting

You need to keep operational costs low. You want to follow Google-recommended practices to migrate these workloads to serverless solutions on Google Cloud. What should you do?

- Migrate the web application to App Engine and the backend API to Cloud Run. Use Cloud Tasks to run your background job on Compute Engine.
- Migrate the web application to App Engine and the backend API to Cloud Run. Use Cloud Tasks to run your background job on Cloud Run.
- Run the web application on a Cloud Storage bucket and the backend API on Cloud Run. Use Cloud Tasks to run your background job on Cloud Run.
- Run the web application on a Cloud Storage bucket and the backend API on Cloud Run. Use Cloud Tasks to run your background job on Compute Engine.

Correct Answer: D

Community vote distribution

B (100%)

 **joao_01** 3 weeks, 1 day ago

Selected Answer: B

Its B, it's serveless and low cost
upvoted 1 times

 **scanner2** 1 month ago

Selected Answer: B

Since the question is asking about serverless solutions. Here, B is the correct answer.
Migrate web application to the App Engine.
Migrate backend API to Cloud Run.
Migrate scheduled long-running job to Cloud Task that will run the background job using Cloud Run.
upvoted 1 times

 **ee_23** 1 month, 3 weeks ago

Selected Answer: B

Serverless
upvoted 3 times

 **3arle** 2 months ago

Selected Answer: B

B is most reasonable
upvoted 1 times

 **qannik** 2 months ago

Selected Answer: B

<https://cloud.google.com/architecture/migration-to-gcp-deploying-your-workloads>
upvoted 1 times

 **gpais** 2 months, 1 week ago

Selected Answer: B

B seems the best option
upvoted 1 times